

Alex's Anthology of Algorithms
Common Code for Contests in Concise C++
v1.0, July 26, 2026

Alex Li

© 2026. All rights reserved.

This page is intentionally left blank.

Preface

Browse the latest version online at: algorithms.alexli.ca

Contributions are always welcome on GitHub: github.com/alxli/algorithm-anthology

Introduction

Welcome to a practical collection of concise C++ implementations of algorithms, data structures, and contest utilities. This is not a first textbook; it is a codebook about implementation details, interfaces, edge cases, and the choices that make familiar algorithms usable in real code. Use it:

- as a reference when you know an algorithm conceptually but want a clean implementation;
- as a printed codebook for contests that allow prepared material, including ICPC-style events;
- as a checklist for edge cases, complexity, overflow, indexing, and API choices;
- as a map of related techniques: slower versions, faster versions, variants, and common reductions sit near one another.

For learning, do not start by copying. Study the idea, prove the invariant or recurrence, implement it yourself, then compare; the differences reveal the practical choices. Each section gives only enough theory to orient the implementation: problem, interface, assumptions, and complexity. Source files are organized as documentation, reusable implementation, then example usage. The reusable part ends at the **Example Usage** marker; everything after exists only so the file compiles and demonstrates itself. When adapting a section, drop the example block and keep the implementation above it.

Guiding principles:

- *Clarity*: A reader familiar with the algorithm should follow the code without reverse-engineering it. Tricks, non-obvious invariants, overflow risks, and language assumptions are documented.
- *Concision*: Contest code is found, read, adapted, and typed under pressure. Short code helps when it removes noise, not when it hides the idea. So the code here is often more verbose than compact ICPC snippets; the extra lines buy clearer names, configurable types, safer defaults, reset functions, explicit handling of variants, and more.
- *Efficiency*: Expected asymptotic behavior and reasonable constants. Both simpler and more robust versions are often included, but educational-only sections are labeled so.
- *Adaptability*: Interfaces make common changes easy, such as swapping a numeric type, changing a combine function, exposing an edge ID, or choosing whether boundaries are included.
- *Consistency*: Similar structures expose similar operations, with repeated naming and example styles, so moving between chapters is easy.

- *Portability*: Targets modern C++17 and avoids unnecessary dependencies; the few contest-useful exceptions are noted at each use.

This began as a personal contest codebook, so it favors single-file snippets, predictable APIs, compact examples, and low constants. Use it actively: annotate it, delete what you dislike, rename functions to match your habits, and stress-test anything you intend to trust. Cheers!

Portability Notes

Examples target a modern GCC or Clang environment with C++17. A typical compile command:

```
g++ -std=c++17 -O2 -Wall -pedantic
```

Most code assumes common contest-platform type sizes: 8-bit `char`, 32-bit `int` and `float`, 64-bit `double` and `long long`. `long double` width and precision are implementation-dependent; sections relying on it say so.

The code targets ISO C++17 unless noted. The most common exceptions:

- GCC/Clang integer extensions like `__int128` or `__uint128_t`, usually guarded with a portable fallback when overflow safety matters.
- Compiler builtins like `__builtin_popcount()` and `__builtin_clz()` — convenient and widely available, though C++20's `<bit>` header standardizes many of them. See <https://gcc.gnu.org/onlinedocs/gcc/Other-Builtins.html> for GCC's reference.
- GNU policy-based data structures like `__gnu_pbds::tree`, used in a few sections to show useful non-standard components.
- A few documented representation assumptions, such as `signbit_()` in the math utilities assuming an IEEE-like floating-point sign bit layout.

Contents

Preface	i
1 Elementary Algorithms	1
1.1 Array Transformations	1
1.1.1 Sorting Algorithms	1
1.1.2 Partitioning and Selection	7
1.1.3 Array Rotation	9
1.1.4 Coordinate Compression	11
1.1.5 Counting Inversions	13
1.2 Arrays and Dynamic Programming	15
1.2.1 Prefix Sums	15
1.2.2 Difference Array	17
1.2.3 Two Pointers and Sliding Window	19
1.2.4 Monotonic Stack	23
1.2.5 Aggregate Queue	25
1.2.6 Maximum Subarray Sum (Kadane)	27
1.2.7 Longest Increasing Subsequence	31
1.2.8 Minimum Excluded Value	32
1.2.9 Iterated Function Cycle Detection	34
1.2.10 0-1 Knapsack	36
1.2.11 Unbounded Knapsack and Coin Change	37
1.2.12 Subset Sum	39
1.2.13 Grid Dynamic Programming	42
1.2.14 Matrix Chain Multiplication	44
1.3 Greedy and Scheduling	45
1.3.1 Fractional Knapsack	45
1.3.2 Interval Scheduling	46
1.3.3 Weighted Interval Scheduling	47
1.3.4 Interval Partitioning	49
1.3.5 Deadline Scheduling	50

1.3.6	Minimizing Late Jobs (Moore-Hodgson)	52
1.4	Streaming and Randomized Algorithms	53
1.4.1	Fisher-Yates Shuffle	53
1.4.2	Reservoir Sampling	54
1.4.3	Frequent Elements	55
1.4.4	Online Statistics	57
1.5	Bit Manipulation	60
1.5.1	Bit Operations and Masks	60
1.5.2	Subset and Submask Iteration	62
1.5.3	Bitmask Dynamic Programming	64
1.5.4	Gray Code	67
1.5.5	XOR Basis	68
1.5.6	Binary Trie	70
1.5.7	Bitwise Convolution (FWHT)	73
2	Data Structures	77
2.1	Linked Lists	77
2.1.1	Singly Linked List	77
2.1.2	Doubly Linked List	79
2.1.3	LRU Cache	82
2.2	Heaps and Priority Queues	83
2.2.1	Binary Heap	83
2.2.2	Leftist Heap	85
2.2.3	Skew Heap	88
2.2.4	Pairing Heap	90
2.3	Dictionaries and Ordered Sets	93
2.3.1	Treap	93
2.3.2	AVL Tree	98
2.3.3	Red-Black Tree	102
2.3.4	Splay Tree	107
2.3.5	Size Balanced Tree	111
2.3.6	Interval Treap	115
2.3.7	Hash Table (Chaining)	119
2.3.8	Hash Table (Open Addressing)	123
2.3.9	Skip List	129
2.4	Range Queries in One Dimension	132
2.4.1	Sparse Table	132
2.4.2	Disjoint Sparse Table	133
2.4.3	Square Root Decomposition	135

2.4.4	Mo's Algorithm	137
2.4.5	Segment Tree (Point Update)	139
2.4.6	Segment Tree (Range Update)	142
2.4.7	Iterative Segment Tree (Point Update)	146
2.4.8	Iterative Segment Tree (Range Update)	149
2.4.9	Segment Tree Beats	153
2.4.10	Sparse Segment Tree	157
2.4.11	Persistent Segment Tree	161
2.4.12	Merge Sort Tree	163
2.4.13	Wavelet Tree	165
2.4.14	Cartesian Tree	167
2.4.15	Implicit Treap	169
2.5	Range Queries in Two Dimensions	174
2.5.1	2D Segment Tree	174
2.5.2	Sparse 2D Segment Tree	177
2.5.3	2D Range Tree	181
2.5.4	K-d Tree for 2D Range Query	184
2.5.5	K-d Tree for Nearest Neighbor	186
2.6	Fenwick Trees	188
2.6.1	Fenwick Tree (Point Update, Range Query)	188
2.6.2	Fenwick Tree (Range Update, Point Query)	190
2.6.3	Fenwick Tree (Range Update, Range Query)	191
2.6.4	Sparse Fenwick Tree (Range Update, Range Query)	193
2.6.5	2D Fenwick Tree	195
2.6.6	Sparse 2D Fenwick Tree	197
2.6.7	Offline 2D Fenwick Tree	199
2.7	Disjoint Sets and Tree Structures	202
2.7.1	Disjoint Set Union	202
2.7.2	Sparse Disjoint Set Union	203
2.7.3	Rollback Disjoint Set Union	205
2.7.4	Lowest Common Ancestor (Sparse Table)	207
2.7.5	Lowest Common Ancestor (Segment Tree)	209
2.7.6	Virtual Tree	212
2.7.7	Weighted Tree Path Queries (Binary Lifting)	215
2.7.8	Heavy Light Decomposition	218
2.7.9	Euler Tour Tree	223
2.7.10	Link-Cut Tree	227
2.7.11	Top Tree	233

3	Strings	243
3.1	String Utilities	243
3.2	Hashing	250
3.2.1	Hash Functions	250
3.2.2	Rolling Hash	255
3.2.3	Rolling Hash Deque	258
3.2.4	Dynamic Rolling Hash	261
3.2.5	Zobrist Hashing	263
3.3	Pattern Matching	265
3.3.1	String Searching (KMP)	265
3.3.2	String Searching (Z Algorithm)	266
3.3.3	String Searching (Aho-Corasick)	268
3.3.4	Palindrome Searching (Manacher)	270
3.3.5	Lyndon Factorization (Duval)	272
3.3.6	De Bruijn Sequence	273
3.4	Suffix Arrays and LCP	275
3.4.1	Suffix Array and LCP (Manber-Myers)	275
3.4.2	Suffix Array and LCP (Counting Sort)	277
3.4.3	Suffix Array and LCP (DC3)	279
3.5	String Data Structures	282
3.5.1	Trie	282
3.5.2	Radix Tree	286
3.5.3	Suffix Tree (Ukkonen)	291
3.5.4	Suffix Automaton	294
3.5.5	Eertree	297
3.6	Sequence Dynamic Programming	299
3.6.1	Longest Common Substring	299
3.6.2	Longest Common Subsequence	300
3.6.3	Edit Distance and Sequence Alignment	302
3.6.4	Palindrome Dynamic Programming	305
3.7	Expression Parsing	307
3.7.1	Recursive Descent Parsing (Simple)	307
3.7.2	Recursive Descent Parsing	308
3.7.3	Shunting Yard Parsing	312
3.8	Encoding and Compression	316
3.8.1	Entropy and Frequency Tables	316
3.8.2	Run-Length Encoding	318
3.8.3	Base64 Encoding	319

3.8.4	Huffman Coding	320
3.8.5	Burrows-Wheeler Transform	322
4	Graphs	325
4.1	DFS and Tree Algorithms	325
4.1.1	Graph Class and DFS	325
4.1.2	Topological Sort	328
4.1.3	Cycle Finding	330
4.1.4	Tree Centers and Diameter	334
4.1.5	Tree Rerooting DP	336
4.1.6	Subtree Small-to-Large Merging	339
4.1.7	Centroid Decomposition	341
4.1.8	Prufer Code	343
4.1.9	Eulerian Trails (Hierholzer)	345
4.2	Connectivity	351
4.2.1	Connected Components	351
4.2.2	Strongly Connected Components (Kosaraju)	352
4.2.3	Strongly Connected Components (Tarjan)	354
4.2.4	Dominator Tree (Lengauer-Tarjan)	356
4.2.5	Articulation Points, Biconnected Components, Block-Cut Forest	358
4.2.6	Bridges, 2-Edge-Connected Components, Bridge Forest	361
4.2.7	Online Bridge Finding	363
4.2.8	2-SAT	365
4.2.9	Functional Graph	367
4.2.10	Offline Dynamic Connectivity	370
4.3	Shortest Paths	373
4.3.1	Unweighted Shortest Path (BFS)	373
4.3.2	0-1 Shortest Path (Deque BFS)	375
4.3.3	Shortest Path (Dijkstra)	376
4.3.4	Shortest Path (Bellman-Ford)	378
4.3.5	Shortest Path (SPFA)	380
4.3.6	DAG Shortest Path (Topological DFS)	381
4.3.7	All-Pairs Shortest Path (Floyd-Warshall)	383
4.3.8	All-Pairs Shortest Path (Johnson)	385
4.4	Spanning Trees	387
4.4.1	Minimum Spanning Tree (Prim)	387
4.4.2	Minimum Spanning Tree (Kruskal)	389
4.4.3	Kruskal Reconstruction Tree	390
4.4.4	Minimum Spanning Arborescence (Edmonds)	393

4.5	Flows and Cuts	394
4.5.1	Maximum Flow (Ford-Fulkerson)	394
4.5.2	Maximum Flow (Edmonds-Karp)	396
4.5.3	Maximum Flow and Decomposition (Dinic)	398
4.5.4	Maximum Flow (Push-Relabel)	402
4.5.5	Flow with Lower Bounds	405
4.5.6	Minimum-Cost Maximum-Flow (SPFA)	408
4.5.7	Minimum-Cost Maximum-Flow (Potentials)	411
4.5.8	Gomory-Hu Tree	414
4.5.9	Global Minimum Cut (Stoer-Wagner)	417
4.6	Matching and Assignment	419
4.6.1	Maximum Bipartite Matching (Kuhn)	419
4.6.2	Maximum Bipartite Matching (Hopcroft-Karp)	420
4.6.3	Minimum Vertex Cover (Konig)	422
4.6.4	Maximum Graph Matching (Blossom)	424
4.6.5	Maximum Weight Graph Matching (Weighted Blossom)	427
4.6.6	Assignment (Hungarian Algorithm)	433
4.6.7	Stable Marriage (Gale-Shapley)	435
4.7	Exponential Graph Problems	436
4.7.1	Maximum Clique (Bron-Kerbosch)	436
4.7.2	Graph Coloring (Backtracking)	438
4.7.3	Edge Coloring (Misra-Gries)	440
4.7.4	Shortest Hamiltonian Path and Cycle (Held-Karp)	443
5	Optimization	446
5.1	Monotone Predicate Search	446
5.1.1	Binary Search	446
5.1.2	Exponential Search	448
5.1.3	Local Minimum (Binary Search)	448
5.2	Unimodal and Continuous Search	449
5.2.1	Ternary and Golden Section Search	449
5.2.2	2D Hill Climbing	452
5.2.3	Simulated Annealing	453
5.3	Dynamic Programming Tricks	455
5.3.1	Divide-and-Conquer DP Optimization	455
5.3.2	SMAWK Row Minima	456
5.3.3	Knuth DP Optimization	458
5.3.4	Lagrangian Relaxation (Aliens Trick)	459
5.3.5	Convex Hull Trick (Semi-Dynamic)	461

5.3.6	Convex Hull Trick (Fully-Dynamic)	462
5.3.7	Li Chao Tree	464
5.4	Numerical Methods	466
5.4.1	Root Finding (Bracketing)	466
5.4.2	Root Finding (Iteration)	468
5.4.3	Polynomial Root Finding (Differentiation)	469
5.4.4	Polynomial Root Finding (Laguerre)	471
5.4.5	Polynomial Root Finding (Ehrlich-Aberth)	473
5.4.6	Integration (Simpson)	477
5.4.7	Linear Programming (Simplex)	478
6	Mathematics	482
6.1	Math Utilities	482
6.2	Combinatorics	488
6.2.1	Combinatorial Calculations	488
6.2.2	Enumerating Arrangements	492
6.2.3	Enumerating Permutations	494
6.2.4	Enumerating Combinations	498
6.2.5	Enumerating Partitions	502
6.2.6	Enumerating Generic Combinatorial Sequences	505
6.3	Number Theory	509
6.3.1	GCD, LCM, Mod Inverse, Chinese Remainder	509
6.3.2	Modular Arithmetic	514
6.3.3	Primitive Roots	518
6.3.4	Discrete Logarithm (Baby-Step Giant-Step)	520
6.3.5	Modular Square Root (Tonelli-Shanks)	522
6.3.6	Jacobi Symbol	524
6.3.7	Prime Generation	525
6.3.8	Prime Tests and Factorization	527
6.3.9	Euler's Totient Function	535
6.3.10	Mobius Function and Inclusion-Exclusion	537
6.4	Arbitrary Precision Arithmetic	540
6.4.1	Big Integers (Simple)	540
6.4.2	Big Integers	543
6.4.3	Rational Numbers	554
6.5	Linear Algebra	557
6.5.1	Matrix Utilities	557
6.5.2	Row Reduction	563
6.5.3	Determinant and Inverse	565

6.5.4	LU Decomposition	569
6.5.5	Characteristic Polynomial	572
6.5.6	Matrix-Tree Theorem	574
6.5.7	Sparse Matrix	576
6.6	Polynomials	582
6.6.1	Number Theoretic Transform	582
6.6.2	Fast Fourier Transform	584
6.6.3	Polynomial Operations	586
6.6.4	Berlekamp-Massey	590
6.6.5	Linear Recurrence (Kitamasa)	591
6.6.6	Polynomial Interpolation	593
6.7	Game Theory	595
6.7.1	Nim	595
6.7.2	Nim Product	596
6.7.3	Sprague-Grundy	598
6.7.4	Grundy Numbers on DAG	599
6.7.5	Minimax and Alpha-Beta Pruning	600
7	Geometry	603
7.1	Geometry Primitives	603
7.1.1	Point	603
7.1.2	Line	607
7.1.3	Circle	611
7.1.4	Triangle	614
7.1.5	Rectangle	617
7.2	Angles, Distances, and Intersections	619
7.2.1	Angles	619
7.2.2	Distances	621
7.2.3	Line Intersection	624
7.2.4	Circle Intersection	629
7.2.5	Circle Tangents	633
7.3	Polygons and Point Sets	634
7.3.1	Polygon Sorting and Area	634
7.3.2	Point-in-Polygon	637
7.3.3	Convex Hull and Diametral Pair	638
7.3.4	Convex Polygon Queries	641
7.3.5	Minimum Enclosing Circle	644
7.3.6	Closest Pair	646
7.3.7	Multiple Segment Intersection	648

7.3.8	Maximum Collinear Points	652
7.3.9	Manhattan MST	653
7.4	Advanced Planar Geometry	655
7.4.1	Convex Polygon Cut	655
7.4.2	Half-Plane Intersection	657
7.4.3	Minkowski Sum	659
7.4.4	Polygon Intersection and Union	662
7.4.5	Delaunay Triangulation (Simple)	666
7.4.6	Delaunay Triangulation (Guibas-Stolfi)	668
7.5	3D Geometry	673
7.5.1	3D Point	673
7.5.2	Polyhedron Volume	674
7.5.3	3D Convex Hull	675
8	Miscellany	678
8.1	Contest Utilities	678
8.2	Fast IO	681
8.3	Debugging	685
8.4	Timing and Benchmarking	688
8.5	Stress Testing	689
8.6	Policy-Based Data Structures	690
8.7	Multithreading	693
8.8	Memory Allocation	694

Chapter 1

Elementary Algorithms

1.1 Array Transformations

1.1.1 Sorting Algorithms

The following functions are equivalent to `std::sort()`, taking random-access iterators as a half-open range `[lo, hi)` to be sorted. The range is sorted into ascending order after the function call. Optionally, a comparison function object specifying a strict weak ordering may be specified to replace the default `operator<`.

- `quicksort(lo, hi, comp = std::less<>)` sorts the range using quicksort.
- `mergesort(lo, hi, comp = std::less<>)` sorts the range using merge sort, which is stable.
- `heapsort(lo, hi, comp = std::less<>)` sorts the range using heapsort.
- `insertion_sort(lo, hi, comp = std::less<>)` sorts the range using insertion sort, which is stable.
- `combsort(lo, hi, comp = std::less<>)` sorts the range using comb sort.
- `radix_sort(lo, hi)` sorts the range using least-significant-byte radix sort.

`radix_sort()` is the exception to the shared interface above: it takes no comparator and requires an integer value type. These functions are not meant to compete with standard library implementations in terms of speed. Instead, they are meant to demonstrate how common sorting algorithms can be concisely implemented in C++.

```
#include <algorithm>
#include <functional>
#include <iterator>
#include <type_traits>
#include <vector>
```

Quicksort repeatedly selects a pivot and partitions the range so that elements comparing less than the pivot precede it, elements comparing equal stay in the middle, and elements comparing greater follow it. Divide and conquer is then applied to the two outer ranges until the original range is sorted. Despite having a worst case of $O(n^2)$, quicksort is often faster in practice than merge sort and heapsort, which both have a worst case time complexity of $O(n \log n)$.

The pivot chosen in this implementation is always a middle element of the range to be sorted. To reduce the likelihood of encountering the worst case, the pivot can be chosen in better ways (e.g. randomly, or using the “median of three” technique).

Time Complexity (Average): $O(n \log n)$.

Time Complexity (Worst): $O(n^2)$.

Space Complexity: $O(\log n)$ auxiliary stack space.

Stable?: No.

```
template<typename It, typename Compare = std::less<>>
void quicksort(It lo, It hi, Compare comp = Compare()) {
    while (hi - lo >= 2) {
        auto pivot = *(lo + (hi - lo) / 2);
        It lt = lo, mid = lo, gt = hi;
        while (mid != gt) {
            if (comp(*mid, pivot)) {
                std::iter_swap(lt++, mid++);
            } else if (comp(pivot, *mid)) {
                std::iter_swap(mid, --gt);
            } else {
                ++mid;
            }
        }
        if (lt - lo < hi - gt) {
            quicksort(lo, lt, comp);
            lo = gt;
        } else {
            quicksort(gt, hi, comp);
            hi = lt;
        }
    }
}
```

Merge sort first divides a list into n sublists of one element each, then recursively merges the sublists into sorted order until only a single sorted sublist remains. Merge sort is a stable sort, meaning that it preserves the relative order of elements which compare equal by `operator<` or the custom comparator given.

An analogous function in the C++ standard library is `std::stable_sort()`, except that the implementation here requires sufficient memory to be available. When $O(n)$ auxiliary memory is not available, `std::stable_sort()` falls back to a time complexity of $O(n \log^2 n)$ whereas the implementation here will simply fail.

Time Complexity (Average): $O(n \log n)$.

Time Complexity (Worst): $O(n \log n)$.

Space Complexity: $O(\log n)$ auxiliary stack space and $O(n)$ auxiliary heap space.

Stable?: Yes.

```
template<typename It, typename Compare = std::less<>>
void mergesort(It lo, It hi, Compare comp = Compare()) {
    if (hi - lo < 2) {
        return;
    }
    It mid = lo + (hi - lo - 1) / 2, a = lo, c = mid + 1;
    mergesort(lo, mid + 1, comp);
    mergesort(mid + 1, hi, comp);
    using T = typename std::iterator_traits<It>::value_type;
    std::vector<T> merged;
    merged.reserve(hi - lo);
    while (a <= mid && c < hi) {
        merged.push_back(comp(*c, *a) ? *c++ : *a++);
    }
    merged.insert(merged.end(), a, mid + 1);
    merged.insert(merged.end(), c, hi);
    std::copy(merged.begin(), merged.end(), lo);
}
```

Heapsort first rearranges an array to satisfy the max-heap property. Then, it repeatedly pops the max element of the heap (the left, unsorted subrange), moving it to the beginning of the right, sorted subrange until the entire

range is sorted. Heapsort has a better worst case time complexity than quicksort and also a better space complexity than merge sort.

The C++ standard library equivalent is calling `std::make_heap(lo, hi)`, followed by `std::sort_heap(lo, hi)`.

The loop below has two modes controlled by whether `i > lo`. While `i > lo`, it is heapifying (Floyd's $O(n)$ bottom-up heap build: sift down each non-leaf from right to left). Once `i == lo`, the heap is built, and the loop switches to extraction: repeatedly swap the root to the back of the unsorted region (`*j`) and sift down the new root.

Time Complexity (Average): $O(n \log n)$.

Time Complexity (Worst): $O(n \log n)$.

Space Complexity: $O(1)$ auxiliary.

Stable?: No.

```
template<typename It, typename Compare = std::less<>>
void heapsort(It lo, It hi, Compare comp = Compare()) {
    if (hi - lo < 2) {
        return;
    }
    using T = typename std::iterator_traits<It>::value_type;
    T tmp = *lo;
    It i = lo + (hi - lo) / 2, j = hi, parent, child;
    while (true) {
        if (i <= lo) {
            if (--j == lo) {
                return;
            }
            tmp = *j;
            *j = *lo;
        } else {
            tmp = *(--i);
        }
        parent = i;
        child = lo + 2 * (i - lo) + 1;
        while (child < j) {
            if (child + 1 < j && comp(*child, *(child + 1))) {
                child++;
            }
            if (!comp(tmp, *child)) {
                break;
            }
            *parent = *child;
            parent = child;
            child = lo + 2 * (parent - lo) + 1;
        }
        *parent = tmp;
    }
}
```

Insertion sort builds the sorted range one element at a time. It scans left to right, and for each element shifts the larger elements of the already-sorted prefix one position to the right to open a slot where the element belongs. Although its average and worst cases are quadratic, it is simple, stable, in-place, and online (it can sort a stream as elements arrive).

Its strength is being adaptive: on an already-sorted or nearly-sorted range, each element travels only a short distance, giving $O(n)$ best-case time and $O(n + d)$ in general for d total inversions. This is why it outperforms the asymptotically faster sorts on small or nearly-sorted ranges, and why hybrid sorts such as introsort and Timsort fall back to it once a subrange is small enough.

The comparison `comp(key, *(j - 1))` is strict, so an element never moves past an earlier element it compares equal to, which keeps the sort stable.

Time Complexity (Best): $O(n)$ on an already-sorted range.

Time Complexity (Average): $O(n^2)$.

Time Complexity (Worst): $O(n^2)$.

Space Complexity: $O(1)$ auxiliary.

Stable?: Yes.

```
template<typename It, typename Compare = std::less<>>
void insertion_sort(It lo, It hi, Compare comp = Compare()) {
    using T = typename std::iterator_traits<It>::value_type;
    for (It i = lo; i != hi; ++i) {
        T key = *i;
        It j = i;
        while (j != lo && comp(key, *(j - 1))) {
            *j = *(j - 1);
            --j;
        }
        *j = key;
    }
}
```

Comb sort is an improved bubble sort. While bubble sort compares only adjacent elements, comb sort compares elements separated by a fixed gap, decreasing that gap after every pass. Large gaps move small values near the end toward the front quickly, avoiding the slow movement of such values in bubble sort. The shrink factor 1.3 is a common empirical choice.

Time Complexity (Worst): $O(n^2)$.

Space Complexity: $O(1)$ auxiliary.

Stable?: No.

```
template<typename It, typename Compare = std::less<>>
void combsort(It lo, It hi, Compare comp = Compare()) {
    int gap = static_cast<int>(hi - lo);
    bool swapped = true;
    while (gap > 1 || swapped) {
        if (gap > 1) {
            gap = gap * 10 / 13;
        }
        swapped = false;
        for (It it = lo; it + gap < hi; ++it) {
            if (comp(*(it + gap), *it)) {
                std::iter_swap(it, it + gap);
                swapped = true;
            }
        }
    }
}
```

Radix sort is used to sort integer elements with a constant number of bits in linear time. This implementation works on ranges pointing to any signed or unsigned integer primitive. Signed values are handled by sorting on a key that flips the sign bit, which maps the two's-complement order onto the unsigned order so the most negative value sorts first.

Digits are processed from least significant to most significant, and each counting-sort pass is stable. After one pass, the values are sorted by the processed digit; stability ensures that a later pass preserves the ordering already established by all less-significant digits. Inductively, after the final pass the values are sorted by the entire key.

This implementation uses one byte per digit, so digits can be extracted with shifts and masks and the counting table has $2^8 = 256$ entries.

Time Complexity: $O(n \cdot w)$ for n integers of w bits each.

Space Complexity: $O(n + 2^b)$ auxiliary for a radix of b bits, i.e. $O(n)$ for constant b .

```
template<typename It>
void radix_sort(It lo, It hi) {
    if (hi - lo < 2) {
        return;
    }
    const int radix_bits = 8;
    const int radix_base = 1 << radix_bits; // e.g. 2^8 = 256
    const int radix_mask = radix_base - 1; // e.g. 2^8 - 1 = 0xFF
    using T = typename std::iterator_traits<It>::value_type;
    static_assert(std::is_integral<T>::value && !std::is_same<T, bool>::value);
    using U = typename std::make_unsigned<T>::type;
    int num_bits = 8 * sizeof(T); // 8 bits per byte
    // Sort on an unsigned key. For signed types, flipping the sign bit sends the most negative value
    // to 0, mapping the signed order onto the unsigned order; logical shifts then extract each digit.
    auto key = [](T x) -> U {
        U u = static_cast<U>(x);
        return std::is_signed<T>::value ? (u ^ (U(1) << (8 * sizeof(T) - 1))) : u;
    };
    std::vector<T> buf(hi - lo);
    for (int pos = 0; pos < num_bits; pos += radix_bits) {
        std::vector<int> count(radix_base);
        for (It it = lo; it != hi; ++it) {
            count[(key(*it) >> pos) & radix_mask]++;
        }
        std::vector<T *> bucket(radix_base);
        T *curr = buf.data();
        for (int i = 0; i < radix_base; curr += count[i++]) {
            bucket[i] = curr;
        }
        for (It it = lo; it != hi; ++it) {
            *bucket[(key(*it) >> pos) & radix_mask]++ = *it;
        }
        std::copy(buf.begin(), buf.end(), lo);
    }
}
```

Example Usage

```
#include <cassert>
#include <cstdlib>
#include <ctime>
#include <iomanip>
#include <iostream>
#include <random>
#include <vector>
using namespace std;

template<typename It>
void print_range(It lo, It hi) {
    while (lo != hi) {
        cout << *lo++ << " ";
    }
    cout << endl;
}

int main() {
    { // Can be used to sort arrays like std::sort().
        int a[] {32, 71, 12, 45, 26, 80, 53, 33};
        quicksort(begin(a), end(a));
        assert(is_sorted(begin(a), end(a)));
    }
    { // STL containers work too.
```

```

vector<int> a{32, 71, 12, 45, 26, 80, 53, 33};
quicksort(a.begin(), a.end());
assert(is_sorted(a.begin(), a.end()));
}
{ // Reverse iterators work as expected.
vector<int> a{32, 71, 12, 45, 26, 80, 53, 33};
heapsort(a.rbegin(), a.rend());
assert(is_sorted(a.rbegin(), a.rend()));
}
{ // We can sort doubles just as well.
vector<double> a{1.1, -5.0, 6.23, 4.123, 155.2};
combsort(a.begin(), a.end());
assert(is_sorted(a.begin(), a.end()));
}
{ // Insertion sort is adaptive: an already-sorted range is left untouched in O(n).
vector<int> a{12, 26, 32, 33, 45, 53, 71, 80};
insertion_sort(a.begin(), a.end());
assert(is_sorted(a.begin(), a.end()));
}
{ // radix_sort() handles signed integers (including negatives), unlike a plain counting sort.
vector<int> a{32, -71, 12, -45, 26, -80, 53, 33};
radix_sort(a.begin(), a.end());
assert(is_sorted(a.begin(), a.end()));
}
{ // Empty and singleton ranges are valid no-ops.
vector<int> empty, single{42};
quicksort(empty.begin(), empty.end());
mergesort(empty.begin(), empty.end());
heapsort(empty.begin(), empty.end());
insertion_sort(empty.begin(), empty.end());
combsort(empty.begin(), empty.end());
radix_sort(empty.begin(), empty.end());
mergesort(single.begin(), single.end());
assert(empty.empty() && single[0] == 42);
}

// Example from: http://www.cplusplus.com/reference/algorithm/stable\_sort
const vector<double> a{3.14, 1.41, 2.72, 4.67, 1.73, 1.32, 1.62, 2.58};
{
vector<double> v(a);
cout << "mergesort() with default comparisons: ";
mergesort(v.begin(), v.end());
print_range(v.begin(), v.end());
}
{
vector<double> v(a);
cout << "mergesort() with integer comparisons: ";
mergesort(v.begin(), v.end(), [](double i, double j) {
return static_cast<int>(i) < static_cast<int>(j);
});
print_range(v.begin(), v.end());
}
cout << "-----" << endl;

std::mt19937 rng(1234567); // Fixed seed for reproducibility.
vector<int> data;
for (int i = 0; i < 5000000; i++) {
data.push_back(static_cast<int>(rng()));
}
cout << "Sorting five million integers..." << endl;
cout.precision(3);

// Lambdas resolve each function template once benchmark() supplies the iterator types.
auto benchmark = [&](const string &name, auto sort) {
vector<int> v = data;
clock_t start = clock();
sort(v.begin(), v.end());
double t = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;

```

```

    cout << setw(14) << left << name + "() : " << fixed << t << "s" << endl;
    assert(is_sorted(v.begin(), v.end()));
};
benchmark("std::sort", [](auto lo, auto hi) { std::sort(lo, hi); });
benchmark("quicksort", [](auto lo, auto hi) { quicksort(lo, hi); });
benchmark("mergesort", [](auto lo, auto hi) { mergesort(lo, hi); });
benchmark("heapsort", [](auto lo, auto hi) { heapsort(lo, hi); });
benchmark("combsort", [](auto lo, auto hi) { combsort(lo, hi); });
benchmark("radix_sort", [](auto lo, auto hi) { radix_sort(lo, hi); });
return 0;
}

```

Example Output

```

mergesort() with default comparisons: 1.32 1.41 1.62 1.73 2.58 2.72 3.14 4.67
mergesort() with integer comparisons: 1.41 1.73 1.32 1.62 2.72 2.58 3.14 4.67
-----
Sorting five million integers...
std::sort(): 0.223s
quicksort(): 0.253s
mergesort(): 0.621s
heapsort(): 0.439s
combsort(): 0.365s
radix_sort(): 0.016s

```

1.1.2 Partitioning and Selection

Partition a range around a pivot, or select the element with a given sorted rank without fully sorting the range. The comparator `comp` defines the ordering and defaults to `std::less<>`. Three-way partitioning, also known as the Dutch National Flag algorithm, rearranges the elements into consecutive groups that precede, are equivalent to, and follow the pivot. A single pass maintains those three regions around an unclassified middle region. An element that precedes the pivot is swapped with `*lt` and advances both `lt` and `mid`; an equivalent element advances only `mid`; and a following element is swapped with `*gt` without advancing `mid`, since the incoming element is still unclassified.

Quickselect repeatedly applies this partition and continues only into the group containing the desired rank. This is the same task as `std::nth_element()`: after the call, `*nth` is the value that would appear there in comparator order, no earlier value follows it, and no later value precedes it. Choosing pivots uniformly at random gives expected linear time, while the three-way split avoids unnecessary work on duplicate-heavy inputs.

- `partition_three_way(lo, hi, pivot, comp = std::less<>())` rearranges `[lo, hi)` in-place and returns a pair of iterators (`mid1`, `mid2`). The three resulting ranges contain elements that precede, are equivalent to, and follow `pivot` according to `comp`, respectively.
- `nth_element2(lo, nth, hi, comp = std::less<>())` rearranges `[lo, hi)` in-place around the 0-based rank represented by iterator `nth` according to `comp`. This requires random-access iterators.

Time Complexity:

- $O(n)$ per call to `partition_three_way()`, where n is the range length.
- $O(n)$ expected per call to `nth_element2()`, and $O(n^2)$ in the worst case.

Space Complexity: $O(1)$ auxiliary.

```

#include <algorithm>
#include <functional>
#include <random>
#include <utility>

template<typename It, typename T, typename Compare = std::less<>>
std::pair<It, It> partition_three_way(It lo, It hi, const T &pivot, Compare comp = Compare()) {
    It lt = lo, mid = lo, gt = hi;
    while (mid != gt) {

```

```

    if (comp(*mid, pivot)) {
        std::iter_swap(lt++, mid++);
    } else if (comp(pivot, *mid)) {
        std::iter_swap(mid, --gt);
    } else {
        ++mid;
    }
}
return {lt, gt};
}

template<typename It, typename Compare = std::less<>>
void nth_element2(It lo, It nth, It hi, Compare comp = Compare()) {
    static std::mt19937 rng(std::random_device{}());
    while (hi - lo > 1) {
        std::uniform_int_distribution<int> dist(0, hi - lo - 1);
        auto pivot = *(lo + dist(rng));
        auto [lt, gt] = partition_three_way(lo, hi, pivot, comp);
        if (nth < lt) {
            hi = lt;
        } else if (nth >= gt) {
            lo = gt;
        } else {
            return;
        }
    }
}

```

Example Usage

```

#include <cassert>
#include <vector>
using namespace std;

int main() {
    vector<int> values{2, 0, 1, 2, 1, 0, 1};
    // On a range containing only 0, 1, and 2, partitioning around 1 sorts the range.
    partition_three_way(values.begin(), values.end(), 1);
    assert(std::is_sorted(values.begin(), values.end()));

    vector<int> b{4, 2, 5, 3, 3, 1};
    auto [mid1, mid2] = partition_three_way(b.begin(), b.end(), 3);
    for (auto it = b.begin(); it != mid1; ++it) {
        assert(*it < 3);
    }
    for (auto it = mid1; it != mid2; ++it) {
        assert(*it == 3);
    }
    for (auto it = mid2; it != b.end(); ++it) {
        assert(*it > 3);
    }

    vector<int> a{5, 6, 4, 3, 2, 6, 7, 9, 3};
    int n = static_cast<int>(a.size());
    nth_element2(a.begin(), a.begin() + n / 2, a.end());
    assert(a[n / 2] == 5);
    // Values left of the median are <=, and values right are >= (the exact order within each side is
    // randomized, since the pivot is chosen at random).
    for (int i = 0; i < n / 2; i++) {
        assert(a[i] <= a[n / 2]);
    }
    for (int i = n / 2 + 1; i < n; i++) {
        assert(a[i] >= a[n / 2]);
    }

    vector<int> descending{1, 4, 2, 3};
}

```

```

nth_element2(descending.begin(), descending.begin() + 1, descending.end(), greater<int>());
assert(descending[1] == 3);
return 0;
}

```

1.1.3 Array Rotation

Rotates the half-open iterator range $[lo, hi)$ left around the split point mid , where $lo \leq mid \leq hi$. Writing the adjacent subranges $[lo, mid)$ and $[mid, hi)$ as **A** and **B**, respectively, the operation transforms **A B** into **B A** while preserving the relative order within each subrange. The standard library function `std::rotate(lo, mid, hi)` performs the same rearrangement.

All three versions below operate in place, using different algorithms and iterator capabilities.

- `rotate1(lo, mid, hi)` requires ForwardIterators and repeatedly swaps the next elements of the two unfinished subranges. When the right iterator reaches **hi**, it wraps to the current **mid**; when the left iterator reaches **mid**, that boundary advances to the right iterator.
- `rotate2(lo, mid, hi)` requires BidirectionalIterators, applying a trick with three reversals. Writing the input subranges as **A B**, two initial reversals yield `reverse(A) reverse(B)`; then the whole range is reversed to yield **B A**, restoring the order within each subrange.
- `rotate3(lo, mid, hi)` requires random-access iterators, applying a juggling algorithm which first divides the range into `gcd(hi - lo, mid - lo)` sets and then rotates the corresponding elements in each set. The method follows the permutation that sends each position to its rotated destination. These form that many disjoint cycles, so walking each cycle once moves every element to its final position.

Time Complexity: $O(n)$ per call to all versions, where n is the distance between **lo** and **hi**.

Space Complexity: $O(1)$ auxiliary for all versions.

```

#include <algorithm>
#include <numeric>

template<typename It>
void rotate1(It lo, It mid, It hi) {
    if (lo == mid || mid == hi) {
        return;
    }
    It next = mid;
    while (lo != next) {
        std::iter_swap(lo++, next++);
        if (next == hi) {
            next = mid;
        } else if (lo == mid) {
            mid = next;
        }
    }
}

template<typename It>
void rotate2(It lo, It mid, It hi) {
    if (lo == mid || mid == hi) {
        return;
    }
    std::reverse(lo, mid);
    std::reverse(mid, hi);
    std::reverse(lo, hi);
}

template<typename It>
void rotate3(It lo, It mid, It hi) {
    if (lo == mid || mid == hi) {
        return;
    }

```

```

}
int n = static_cast<int>(hi - lo);
int jump = static_cast<int>(mid - lo);
int g = std::gcd(jump, n), cycle = n / g;
for (int i = 0; i < g; i++) {
    int curr = i, next;
    for (int j = 0; j < cycle - 1; j++) {
        next = curr + jump;
        if (next >= n) {
            next -= n;
        }
        std::iter_swap(lo + curr, lo + next);
        curr = next;
    }
}
}
}

```

Example Usage

```

#include <algorithm>
#include <cassert>
#include <vector>
using namespace std;

int main() {
    vector<int> a0, a1, a2, a3;
    for (int i = 0; i < 10000; i++) {
        a0.push_back(i);
    }
    a1 = a2 = a3 = a0;
    int mid = 5678;
    std::rotate(a0.begin(), a0.begin() + mid, a0.end());
    rotate1(a1.begin(), a1.begin() + mid, a1.end());
    rotate2(a2.begin(), a2.begin() + mid, a2.end());
    rotate3(a3.begin(), a3.begin() + mid, a3.end());
    assert(a0 == a1 && a0 == a2 && a0 == a3);

    vector<int> no_op{1, 2, 3};
    rotate1(no_op.begin(), no_op.begin(), no_op.end());
    assert((no_op == vector<int>{1, 2, 3}));
    rotate2(no_op.begin(), no_op.end(), no_op.end());
    assert((no_op == vector<int>{1, 2, 3}));
    rotate3(no_op.begin(), no_op.begin(), no_op.end());
    assert((no_op == vector<int>{1, 2, 3}));

    // Example from: http://en.cppreference.com/w/cpp/algorithm/rotate
    vector<int> a{2, 4, 2, 0, 5, 10, 7, 3, 7, 1};

    // Insertion sort.
    for (auto i = a.begin(); i != a.end(); ++i) {
        rotate1(std::upper_bound(a.begin(), i, *i), i, i + 1);
    }
    assert((a == vector<int>{0, 1, 2, 2, 3, 4, 5, 7, 7, 10}));

    // Simple rotation to the left.
    rotate2(a.begin(), a.begin() + 1, a.end());
    assert((a == vector<int>{1, 2, 2, 3, 4, 5, 7, 7, 10, 0}));

    // Simple rotation to the right.
    rotate3(a.rbegin(), a.rbegin() + 1, a.rend());
    assert((a == vector<int>{0, 1, 2, 2, 3, 4, 5, 7, 7, 10}));
    return 0;
}

```

1.1.4 Coordinate Compression

Given a range of n numerical elements, reassign each element to an integer in the domain $[0, k)$, where k is the number of distinct elements in the original range, while preserving the initial relative ordering of elements. That is, if a is the array of original values and b is the array of compressed values, then every pair of indices i, j in $[0, n)$ shall satisfy $a[i] < a[j]$ if and only if $b[i] < b[j]$.

All implementations below take ranges as ForwardIterators. The comparator `comp` defines the value ordering: `comp(a, b)` is true when `a` precedes `b`. The two `compress()` functions rewrite a range in place and then discard the mapping.

- `compress1(lo, hi, comp = std::less<>())` performs the compression by sorting the array, removing duplicates, and binary searching for the position of each original value.
- `compress2(lo, hi, comp = std::less<>())` achieves the same result by inserting all values into an `std::map`, which automatically removes duplicate values and supports efficient lookups of the compressed values.

`CoordinateCompressor` is a class version that retains the sorted table of distinct values so that arbitrary values can be mapped to and from compressed ranks long after construction (e.g. for offline queries arriving separately from the array being compressed). It uses `std::less<T>` by default; to customize the ordering, instantiate `CoordinateCompressor<T, Compare>` and pass the comparator to either constructor.

- `CoordinateCompressor<T>()` constructs an empty compressor. Register values with `add(x)`, then call `build()` once before the first query.
- `CoordinateCompressor<T>(lo, hi)` constructs a compressor from the half-open iterator range `[lo, hi)`.
- `size()` returns the number of distinct registered values k .
- `value(r)` returns the original value with rank `r`, inverting `rank()`.
- `rank(x)` returns the compressed value (rank) of `x` in $[0, k)$. `x` must have been registered.
- `contains(x)` returns whether `x` is a registered value.

Time Complexity:

- $O(n \log n)$ per call to `compress1(lo, hi)`, `compress2(lo, hi)`, and `CoordinateCompressor(lo, hi)`, where n is the distance between `lo` and `hi`.
- $O(1)$ amortized per call to `add()`, and $O(m \log m)$ per call to `build()`, where m is the total number of values registered.
- $O(\log k)$ per call to `rank()` and `contains()`, where k is the number of distinct values.
- $O(1)$ per call to `size()` and `value()`.

Space Complexity:

- $O(n)$ auxiliary for `compress1()` and `compress2()`.
- $O(k)$ storage for `CoordinateCompressor`.

```
#include <algorithm>
#include <cassert>
#include <functional>
#include <iterator>
#include <map>
#include <vector>

template<typename It, typename Compare = std::less<>>
void compress1(It lo, It hi, Compare comp = Compare()) {
    using T = typename std::iterator_traits<It>::value_type;
    std::vector<T> v(lo, hi);
    std::sort(v.begin(), v.end(), comp);
    v.resize(
        std::unique(
            v.begin(), v.end(), [&](const T &a, const T &b) { return !comp(a, b) && !comp(b, a); }
        ) -
        v.begin()
    );
    for (It it = lo; it != hi; ++it) {
        *it = static_cast<int>(std::lower_bound(v.begin(), v.end(), *it, comp) - v.begin());
    }
}
```



```

    }
}

template<typename It, typename Compare = std::less<>>
void compress2(It lo, It hi, Compare comp = Compare()) {
    using T = typename std::iterator_traits<It>::value_type;
    std::map<T, int, Compare> m(comp);
    for (It it = lo; it != hi; ++it) {
        m[*it] = 0;
    }
    int id = 0;
    for (auto &[key, val] : m) {
        val = id++;
    }
    for (It it = lo; it != hi; ++it) {
        *it = m[*it];
    }
}

template<typename T, typename Compare = std::less<T>>
class CoordinateCompressor {
    Compare comp;
    std::vector<T> values;

public:
    explicit CoordinateCompressor(Compare comp = Compare()) : comp(std::move(comp)) {}

    template<typename It>
    CoordinateCompressor(It lo, It hi, Compare comp = Compare())
        : comp(std::move(comp)), values(lo, hi) {
        build();
    }

    int size() const { return static_cast<int>(values.size()); }
    const T &value(int r) const { return values[r]; }
    void add(const T &x) { values.push_back(x); }

    void build() {
        std::sort(values.begin(), values.end(), comp);
        values.resize(
            std::unique(
                values.begin(), values.end(),
                [&](const T &a, const T &b) { return !comp(a, b) && !comp(b, a); }
            ) -
            values.begin()
        );
    }

    int rank(const T &x) const {
        int r =
            static_cast<int>(std::lower_bound(values.begin(), values.end(), x, comp) - values.begin());
        assert(r < size() && !comp(x, values[r])); // x must be registered.
        return r;
    }

    bool contains(const T &x) const {
        return std::binary_search(values.begin(), values.end(), x, comp);
    }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {

```

```

{
    vector<int> a{1, 30, 30, 7, 9, 8, 99, 99};
    compress1(a.begin(), a.end());
    assert((a == vector<int>{0, 4, 4, 1, 3, 2, 5, 5}));
}
{
    vector<int> a{1, 30, 30, 7, 9, 8, 99, 99};
    compress2(a.begin(), a.end());
    assert((a == vector<int>{0, 4, 4, 1, 3, 2, 5, 5}));
}
{ // Non-integral types work too, as long as ints can be assigned to them.
    vector<double> a{0.5, -1.0, 3, -1.0, 20, 0.5};
    compress1(a.begin(), a.end());
    assert((a == vector<double>{1, 0, 2, 0, 3, 1}));
}
{
    vector<int> a{10, 5, 30, 5, 20};
    CoordinateCompressor<int> cc(a.begin(), a.end());
    assert(cc.size() == 4);
    assert(cc.rank(5) == 0);
    assert(cc.rank(30) == 3);
    assert(cc.value(2) == 20);
    assert(cc.contains(10) && !cc.contains(7));
}
{ // Register values incrementally, then build once before querying.
    CoordinateCompressor<double> cc;
    cc.add(0.5);
    cc.add(-1.0);
    cc.add(0.5);
    cc.build();
    assert(cc.size() == 2);
    assert(cc.rank(0.5) == 1);
    assert(cc.value(0) == -1.0);
}
{
    vector<int> a{10, 30, 20, 10};
    CoordinateCompressor<int, greater<int>> cc(a.begin(), a.end());
    assert(cc.rank(30) == 0 && cc.rank(10) == 2);
    assert(cc.value(1) == 20);
    compress1(a.begin(), a.end(), greater<int>());
    assert((a == vector<int>{2, 0, 1, 2}));
}
return 0;
}

```

1.1.5 Counting Inversions

The number of inversions for a sequence a is the number of ordered pairs (i, j) such that $i < j$ and $a[i] > a[j]$. This is roughly how “close” an array is to being sorted, but is *not* the minimum number of swaps required to sort the array. If the array is sorted, then the inversion count is 0. If the array is sorted in decreasing order, then the inversion count is maximal. The following two functions are each techniques to efficiently count inversions. In the merge sort approach, whenever the merge step emits an element from the right half, that element jumps ahead of every unmerged left-half element, and exactly that many inversions are added to the count.

- `inversions(lo, hi, comp = std::less<>)` uses merge sort to return the number of inversions given two random-access iterators as a range $[lo, hi)$. The input range will be sorted after the function call. Optionally, a comparison function object specifying a strict weak ordering may be specified to replace the default `operator<`.
- `inversions(a)` uses coordinate compression and a Fenwick tree to return the number of inversions in an integer vector without modifying it.

Time Complexity:

- $O(n \log n)$ per call to `inversions(lo, hi)`, where n is the distance between `lo` and `hi`.
- $O(n \log n)$ per call to `inversions(a)`.

Space Complexity:

- $O(\log n)$ auxiliary stack space and $O(n)$ auxiliary heap space for `inversions(lo, hi)`.
- $O(n)$ auxiliary heap space for `inversions(a)`.

```
#include <algorithm>
#include <cstdint>
#include <functional>
#include <iterator>
#include <vector>

template<typename It, typename Compare = std::less<>>
int64_t inversions(It lo, It hi, Compare comp = Compare()) {
    int n = static_cast<int>(hi - lo);
    if (n < 2) {
        return 0;
    }
    It mid = lo + (n - 1) / 2, a = lo, c = mid + 1;
    int64_t res = inversions(lo, mid + 1, comp) + inversions(mid + 1, hi, comp);
    using T = typename std::iterator_traits<It>::value_type;
    std::vector<T> merged;
    merged.reserve(n);
    while (a <= mid && c < hi) {
        if (comp(*c, *a)) {
            merged.push_back(*(c++));
            res += (mid - a) + 1;
        } else {
            merged.push_back(*(a++));
        }
    }
    if (a > mid) {
        for (It it = c; it != hi; ++it) {
            merged.push_back(*it);
        }
    } else {
        for (It it = a; it <= mid; ++it) {
            merged.push_back(*it);
        }
    }
    for (It it = lo; it != hi; ++it) {
        *it = merged[it - lo];
    }
    return res;
}

int64_t inversions(const std::vector<int> &a) {
    int n = static_cast<int>(a.size());
    std::vector<int> values(a);
    std::sort(values.begin(), values.end());
    values.resize(std::unique(values.begin(), values.end()) - values.begin());
    std::vector<int> bit(values.size() + 1, 0);
    int64_t res = 0;
    for (int i = n - 1; i >= 0; i--) {
        int id =
            static_cast<int>(std::lower_bound(values.begin(), values.end(), a[i]) - values.begin()) + 1;
        for (int j = id - 1; j > 0; j -= j & -j) {
            res += bit[j];
        }
        for (int j = id; j < static_cast<int>(bit.size()); j += j & -j) {
            bit[j]++;
        }
    }
    return res;
}
```

```
}

```

Example Usage

```
#include <cassert>

int main() {
    {
        std::vector<int> a{6, 9, 1, 14, 8, 12, 3, 2};
        assert(inversions(a.begin(), a.end()) == 16);
        assert(std::is_sorted(a.begin(), a.end()));
    }
    {
        std::vector<int> a{6, 9, 1, 14, 8, 12, 3, 2};
        assert(inversions(a) == 16);
        assert((a == std::vector<int>{6, 9, 1, 14, 8, 12, 3, 2}));
    }
    {
        std::vector<int> a{2, 2, 1};
        assert(inversions(a) == 2); // Equal elements are not inversions.
    }
    return 0;
}
```

1.2 Arrays and Dynamic Programming

1.2.1 Prefix Sums

Precomputes prefix sums so range-sum queries can be answered in constant time. This is one of the most common array preprocessing tools, and also appears as a building block for subarray sums, difference arrays, and two-dimensional grids. Each table entry stores the sum of all elements before it, so any range sum is the difference of two entries; in two dimensions, rectangle sums combine four entries by inclusion-exclusion.

Prefix sums can also be combined with a frequency table to count subarrays having a target sum. When the current prefix sum is s , every earlier prefix sum equal to $s - t$ identifies a subarray with sum t . Recording frequencies rather than only whether a prefix exists counts duplicate sums correctly.

- `prefix_sums(a)` returns the prefix sum array `pref` of length $n + 1$, with `pref[0] = 0` and `pref[i + 1]` equal to the sum of `a[0]` through `a[i]`.
- `range_sum(pref, lo, hi)` returns the sum of range $[lo, hi]$.
- `prefix_sums_2d(a)` returns a two-dimensional prefix sum table for matrix `a`.
- `rectangle_sum(pref, r1, c1, r2, c2)` returns the sum of the rectangle with rows $[r1, r2]$ and columns $[c1, c2]$.
- `count_subarrays_with_sum(a, target)` returns the number of contiguous subarrays whose sum is `target`.

All prefix sums, target differences, inclusion-exclusion results, and returned counts must fit in `int64_t`.

Time Complexity:

- $O(n)$ per call to `prefix_sums(a)`, where n is the array size.
- $O(m \cdot n)$ per call to `prefix_sums_2d(a)`, where m and n are the number of rows and columns of `a`, respectively.
- $O(1)$ per range or rectangle query.
- $O(n)$ expected per call to `count_subarrays_with_sum(a, target)`.

Space Complexity:

- $O(n)$ for the array returned by `prefix_sums(a)`.
- $O(m \cdot n)$ for the table returned by `prefix_sums_2d(a)`.

- $O(1)$ auxiliary for `range_sum()` and `rectangle_sum()`.
- $O(n)$ auxiliary for `count_subarrays_with_sum(a, target)`.

```
#include <cassert>
#include <cstdint>
#include <unordered_map>
#include <vector>

std::vector<int64_t> prefix_sums(const std::vector<int> &a) {
    std::vector<int64_t> pref(a.size() + 1, 0);
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        pref[i + 1] = pref[i] + a[i];
    }
    return pref;
}

int64_t range_sum(const std::vector<int64_t> &pref, int lo, int hi) {
    assert(0 <= lo && lo <= hi && hi < static_cast<int>(pref.size()) - 1);
    return pref[hi + 1] - pref[lo];
}

std::vector<std::vector<int64_t>> prefix_sums_2d(const std::vector<std::vector<int>> &a) {
    int rows = static_cast<int>(a.size());
    int cols = a.empty() ? 0 : static_cast<int>(a[0].size());
    std::vector<std::vector<int64_t>> pref(rows + 1, std::vector<int64_t>(cols + 1, 0));
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            pref[i + 1][j + 1] = a[i][j] + pref[i][j + 1] + pref[i + 1][j] - pref[i][j];
        }
    }
    return pref;
}

int64_t rectangle_sum(
    const std::vector<std::vector<int64_t>> &pref, int r1, int c1, int r2, int c2
) {
    assert(
        !pref.empty() && !pref[0].empty() && 0 <= r1 && r1 <= r2 &&
        r2 < static_cast<int>(pref.size()) - 1 && 0 <= c1 && c1 <= c2 &&
        c2 < static_cast<int>(pref[0].size()) - 1
    );
    return pref[r2 + 1][c2 + 1] - pref[r1][c2 + 1] - pref[r2 + 1][c1] + pref[r1][c1];
}

int64_t count_subarrays_with_sum(const std::vector<int> &a, int64_t target) {
    std::unordered_map<int64_t, int64_t> count{{0, 1}};
    int64_t sum = 0, result = 0;
    for (int x : a) {
        sum += x;
        int64_t needed = sum - target;
        if (auto it = count.find(needed); it != count.end()) {
            result += it->second;
        }
        count[sum]++;
    }
    return result;
}
```

Example Usage

```
using namespace std;

int main() {
    vector<int> a{3, -1, 4, 1, 5};
    auto pref = prefix_sums(a);
    assert(range_sum(pref, 0, 4) == 12); // Whole array.
```

```

assert(range_sum(pref, 1, 3) == 4);    // -1 + 4 + 1.
assert(range_sum(pref, 2, 2) == 4);    // Single element.
assert(count_subarrays_with_sum(a, 4) == 2);
assert(count_subarrays_with_sum(vector<int>{0, 0, 0}, 0) == 6);

vector<vector<int>> grid{
    {1, 2, 3},
    {4, 5, 6}
};
auto pre2 = prefix_sums_2d(grid);
assert(rectangle_sum(pre2, 0, 1, 1, 2) == 16); // Rows 0-1 and columns 1-2.
assert(rectangle_sum(pre2, 0, 0, 1, 2) == 21); // Whole grid.
assert(rectangle_sum(pre2, 1, 1, 1, 1) == 5);  // Single cell.
return 0;
}

```

1.2.2 Difference Array

Applies many range-add updates to an array in linear total time using a difference array. Instead of immediately adding `delta` to every element of range `[lo, hi]`, we can add `delta` at `diff[lo]` and subtract it at `diff[hi + 1]`; a final prefix sum reconstructs the updated values. This idea extends to two dimensions: a rectangle update touches just the four corner cells of a difference grid by inclusion-exclusion, and a final two-dimensional prefix sum reconstructs the grid.

- `DifferenceArray(n)` constructs an initially zero array of size `n`.
- `add(lo, hi, delta)` adds `delta` to every position in range `[lo, hi]`.
- `build()` returns the final array after all range updates.
- `DifferenceArray2D(rows, cols)` constructs an initially zero grid with `rows` rows and `cols` columns.
- `add(r1, c1, r2, c2, delta)` adds `delta` to every cell of the rectangle with rows `[r1, r2]` and columns `[c1, c2]`.
- `build()` returns the final grid after all rectangle updates.

All stored differences and reconstructed values must fit in `int64_t`.

Time Complexity:

- $O(1)$ per call to `add()` of either version.
- $O(n)$ per call to `build()` of `DifferenceArray`, where n is the array size.
- $O(m \cdot n)$ per call to `build()` of `DifferenceArray2D`, where m and n are the number of rows and columns, respectively.

Space Complexity:

- $O(n)$ for storage and $O(n)$ for the array returned by `DifferenceArray::build()`.
- $O(m \cdot n)$ for storage and $O(m \cdot n)$ for the grid returned by `DifferenceArray2D::build()`.

```

#include <cassert>
#include <cstdint>
#include <vector>

class DifferenceArray {
    std::vector<int64_t> diff;

public:
    explicit DifferenceArray(int n) {
        assert(n >= 0);
        diff.assign(n + 1, 0);
    }

    void add(int lo, int hi, int64_t delta) {
        assert(0 <= lo && lo <= hi && hi + 1 < static_cast<int>(diff.size()));
        diff[lo] += delta; // Overflow warning.
        diff[hi + 1] -= delta;
    }
};

```

```

    }

    std::vector<int64_t> build() const {
        std::vector<int64_t> res(static_cast<int>(diff.size()) - 1);
        int64_t cur = 0;
        for (int i = 0; i < static_cast<int>(res.size()); i++) {
            cur += diff[i]; // Overflow warning.
            res[i] = cur;
        }
        return res;
    }
};

class DifferenceArray2D {
    std::vector<std::vector<int64_t>> diff;

public:
    DifferenceArray2D(int rows, int cols) {
        assert(rows >= 0 && cols >= 0);
        diff.assign(rows + 1, std::vector<int64_t>(cols + 1, 0));
    }

    void add(int r1, int c1, int r2, int c2, int64_t delta) {
        int rows = static_cast<int>(diff.size()) - 1;
        int cols = static_cast<int>(diff[0].size()) - 1;
        assert(0 <= r1 && r1 <= r2 && r2 < rows);
        assert(0 <= c1 && c1 <= c2 && c2 < cols);
        diff[r1][c1] += delta; // Overflow warning.
        diff[r1][c2 + 1] -= delta;
        diff[r2 + 1][c1] -= delta;
        diff[r2 + 1][c2 + 1] += delta;
    }

    std::vector<std::vector<int64_t>> build() const {
        int rows = static_cast<int>(diff.size()) - 1;
        int cols = static_cast<int>(diff[0].size()) - 1;
        std::vector<std::vector<int64_t>> res(rows, std::vector<int64_t>(cols));
        for (int i = 0; i < rows; i++) {
            for (int j = 0; j < cols; j++) {
                int64_t up = i > 0 ? res[i - 1][j] : 0;
                int64_t left = j > 0 ? res[i][j - 1] : 0;
                int64_t diag = i > 0 && j > 0 ? res[i - 1][j - 1] : 0;
                res[i][j] = diff[i][j] + up + left - diag; // Overflow warning.
            }
        }
        return res;
    }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    DifferenceArray d(5);
    d.add(0, 2, 2); // Add 2 to indices 0-2.
    d.add(1, 4, 4); // Add 4 to indices 1-4.
    d.add(2, 3, -1); // Subtract 1 from indices 2-3.
    d.add(4, 4, 3); // Single index update.
    vector<int64_t> a = d.build();
    vector<int64_t> expected{2, 6, 5, 3, 7};
    for (int i = 0; i < static_cast<int>(expected.size()); i++) {
        assert(a[i] == expected[i]);
    }
}

```

```

DifferenceArray2D d2(3, 4);
d2.add(0, 0, 1, 2, 1); // Add 1 to rows 0-1 and columns 0-2.
d2.add(1, 1, 2, 3, 2); // Add 2 to rows 1-2 and columns 1-3.
d2.add(2, 0, 2, 0, 5); // Single cell update.
vector<vector<int64_t>> g = d2.build();
vector<vector<int64_t>> expected2{
    {1, 1, 1, 0},
    {1, 3, 3, 2},
    {5, 2, 2, 2}
};
assert(g == expected2);
return 0;
}

```

1.2.3 Two Pointers and Sliding Window

Uses two moving indices to avoid trying all pairs of positions. This technique has two common forms: opposite-end pointers over sorted arrays, and a sliding window over contiguous subarrays. For sorted 2-sum, if the current sum is too small, pairing the left value with any smaller right value also fails, so the left pointer can safely advance; if the sum is too large, the symmetric argument lets the right pointer retreat. Each step therefore discards an entire row or column of candidate pairs.

A sliding window applies the same idea when extending and shrinking a range changes its state monotonically. With nonnegative values, extending a window cannot decrease its sum and removing its leftmost value cannot increase it, so no shorter feasible window is skipped. For a bound such as at most k distinct values, the left pointer advances only until the window becomes valid again. Since neither endpoint moves backward, each element enters and leaves the window at most once.

For fixed-length windows, a monotonic deque of indices can report every minimum or maximum in linear total time. It stores only candidates that could still become the answer, ordered from the current extreme at the front to the least useful value at the back. A new element removes every older value that it dominates, while expired indices are removed from the front. Each index is again pushed and popped at most once.

The examples below cover sorted 2-sum/3-sum variants, finding the shortest subarray with sum at least a target when all values are nonnegative, and finding the longest subarray with at most k distinct values. They also cover the minimum or maximum of every fixed-length window and an online dynamic programming recurrence.

- `two_sum_sorted(a, target)` returns two indices whose values sum to `target`, or `(-1, -1)` if no such pair exists. The input array must be sorted.
- `three_sum(a, target)` returns three original indices whose values sum to `target`, or `(-1, -1, -1)` if none exist.
- `min_length_at_least(a, target)` returns a tuple `(length, lo, hi)`, the minimum length and inclusive endpoints of a contiguous subarray with sum at least `target`, or length `-1` if none exists. Values in `a` must be nonnegative. If `target` is nonpositive, it returns the empty subarray.
- `longest_at_most_k_distinct(a, k)` returns a tuple `(length, lo, hi)`, the maximum length and inclusive endpoints of a contiguous subarray containing at most `k` distinct values. If `k` is nonpositive, it returns the empty subarray.
- `sliding_window_extrema(a, k, comp = std::less<>)` returns the extreme value in each window of length `k`. With the default `less<>` comparator it returns minimums; passing `greater<>` returns maximums.

Time Complexity:

- $O(n)$ per call to `two_sum_sorted()` and `min_length_at_least()`, where n is the array size.
- $O(n^2)$ per call to `three_sum()`.
- $O(n)$ expected per call to `longest_at_most_k_distinct()`.
- $O(n)$ per call to `sliding_window_extrema()`.

Space Complexity:

- $O(1)$ auxiliary for `two_sum_sorted(a, target)`.

- $O(n)$ auxiliary for `three_sum(a, target)`.
- $O(1)$ auxiliary for `min_length_at_least(a, target)`.
- $O(k)$ auxiliary for `longest_at_most_k_distinct(a, k)`.
- $O(k)$ auxiliary and $O(n)$ for the result of `sliding_window_extrema()`.

```

#include <algorithm>
#include <array>
#include <cassert>
#include <cstdint>
#include <deque>
#include <functional>
#include <tuple>
#include <unordered_map>
#include <utility>
#include <vector>

std::pair<int, int> two_sum_sorted(const std::vector<int> &a, int64_t target) {
    for (int l = 0, r = static_cast<int>(a.size()) - 1; l < r;) {
        int64_t sum = static_cast<int64_t>(a[l]) + a[r];
        if (sum == target) {
            return {l, r};
        } else if (sum < target) {
            l++;
        } else {
            r--;
        }
    }
    return {-1, -1};
}

std::array<int, 3> three_sum(const std::vector<int> &a, int64_t target) {
    int n = static_cast<int>(a.size());
    std::vector<std::pair<int, int>> sorted;
    sorted.reserve(a.size());
    for (int i = 0; i < n; i++) {
        sorted.push_back({a[i], i});
    }
    std::sort(sorted.begin(), sorted.end());
    for (int i = 0; i < n; i++) {
        int l = i + 1, r = n - 1;
        while (l < r) {
            int64_t sum = static_cast<int64_t>(sorted[i].first) + sorted[l].first + sorted[r].first;
            if (sum == target) {
                return {sorted[i].second, sorted[l].second, sorted[r].second};
            } else if (sum < target) {
                l++;
            } else {
                r--;
            }
        }
    }
    return {-1, -1, -1};
}

std::tuple<int, int, int> min_length_at_least(const std::vector<int> &a, int64_t target) {
    if (target <= 0) {
        return {0, 0, -1};
    }
    int best = static_cast<int>(a.size()) + 1, best_lo = 0, best_hi = -1;
    int64_t sum = 0;
    for (int l = 0, r = 0; r < static_cast<int>(a.size()); r++) {
        sum += a[r]; // Overflow warning.
        while (sum >= target) {
            if (r - l + 1 < best) {
                best = r - l + 1;
                best_lo = l;
                best_hi = r;
            }
            sum -= a[l];
            l++;
        }
    }
    return {best_lo, best_hi, best};
}

```

```

    }
    sum -= a[l++];
}
}
if (best == static_cast<int>(a.size()) + 1) {
    return {-1, 0, -1};
}
return {best, best_lo, best_hi};
}

std::tuple<int, int, int> longest_at_most_k_distinct(const std::vector<int> &a, int k) {
    if (k <= 0) {
        return {0, 0, -1};
    }
    std::unordered_map<int, int> count;
    int best = 0, best_lo = 0, best_hi = -1;
    for (int l = 0, r = 0; r < static_cast<int>(a.size()); r++) {
        count[a[r]]++;
        while (static_cast<int>(count.size()) > k) {
            if (--count[a[l]] == 0) {
                count.erase(a[l]);
            }
            l++;
        }
        if (best < r - l + 1) {
            best = r - l + 1;
            best_lo = l;
            best_hi = r;
        }
    }
    return {best, best_lo, best_hi};
}

template<typename T, typename Compare = std::less<>>
std::vector<T> sliding_window_extrema(const std::vector<T> &a, int k, Compare comp = Compare()) {
    int n = static_cast<int>(a.size());
    assert(1 <= k && k <= n);
    std::deque<int> window;
    std::vector<T> res;
    res.reserve(n - k + 1);
    for (int i = 0; i < n; i++) {
        while (!window.empty() && !comp(a[window.back()], a[i])) {
            window.pop_back();
        }
        window.push_back(i);
        if (window.front() <= i - k) {
            window.pop_front();
        }
        if (i >= k - 1) {
            res.push_back(a[window.front()]);
        }
    }
    return res;
}

```

A monotone queue can be used to maintain the minimum or maximum value in an online sliding window. This is useful for dynamic programming recurrences where each transition may only come from one of the last w states, such as $dp(i) = a_i + \min_{j \in [i-w, i-1]} dp(j)$.

The queue stores candidate (`index`, `value`) pairs in monotone order. Expired indices are removed from the front, and dominated values are removed from the back before inserting a new candidate.

- `MonotoneQueue<T>()` constructs an empty queue for minimum queries. Instantiate `MonotoneQueue<T, std::greater<T>>` for maximum queries.
- `push(index, value)` inserts candidate `value` at position `index`, removing dominated candidates from the back.
- `expire(first_valid)` removes candidates with index less than `first_valid`.

- `top()` returns the best active (`index`, `value`) pair. The queue must be nonempty.
- `empty()` returns whether the queue has no active candidates.

Time Complexity:

- $O(n)$ for any sequence of n calls to `push(index, value)` and `expire(first_valid)`.
- $O(1)$ amortized per call to `push(index, value)` and `expire(first_valid)`.
- $O(1)$ worst-case per call to `top()` and `empty()`.

Space Complexity: $O(w)$ for storage of a sliding window of width w .

```
template<typename T, typename Compare = std::less<T>>
class MonotoneQueue {
    std::deque<std::pair<int, T>> q;
    Compare better;

public:
    bool empty() const { return q.empty(); }

    void push(int index, const T &value) {
        while (!q.empty() && !better(q.back().second, value)) {
            q.pop_back();
        }
        q.emplace_back(index, value);
    }

    void expire(int first_valid) {
        while (!q.empty() && q.front().first < first_valid) {
            q.pop_front();
        }
    }

    std::pair<int, T> top() const {
        assert(!q.empty());
        return q.front();
    }
};
```

Example Usage

```
using namespace std;

int main() {
    vector<int> c{-3, -1, 2, 4, 8, 11};
    assert(two_sum_sorted(c, 10) == make_pair(1, 5)); // -1 + 11.
    assert(two_sum_sorted(c, 7) == make_pair(1, 4)); // -1 + 8.
    assert(two_sum_sorted(c, 100) == make_pair(-1, -1));

    vector<int> d{4, -1, 8, 2, -3, 11};
    array<int, 3> triple = three_sum(d, 9);
    assert(d[triple[0]] + d[triple[1]] + d[triple[2]] == 9);
    assert((three_sum(d, 30) == array<int, 3>{-1, -1, -1}));

    vector<int> a{2, 3, 1, 2, 4, 3};
    auto [length, lo, hi] = min_length_at_least(a, 7);
    assert(length == 2);
    assert(lo == 4 && hi == 5); // [4, 3].
    auto [empty_length, empty_lo, empty_hi] = min_length_at_least(a, 0);
    assert(empty_length == 0);
    assert(empty_lo == 0 && empty_hi == -1);

    vector<int> b{1, 2, 1, 3, 4, 3, 5};
    auto [longest, longest_lo, longest_hi] = longest_at_most_k_distinct(b, 2);
    assert(longest == 3);
    assert(longest_lo == 0 && longest_hi == 2); // [1, 2, 1].
```

```

vector<int> e{1, 3, -1, -3, 5, 3, 6, 7};
assert((sliding_window_extrema(e, 3) == vector<int>{-1, -3, -3, -3, 3, 3}));
assert((sliding_window_extrema(e, 3, greater<>()) == vector<int>{3, 3, 5, 5, 6, 7}));

// Toy DP: dp[i] = f[i] + min(dp[j]) for j in [i - max_jump, i - 1].
// Scanning the window costs O(max_jump) per state; the monotone queue makes it O(1) amortized.
vector<int> f{4, 2, 7, 1, 3, 6};
int max_jump = 3;
vector<int> dp(f.size());
MonotoneQueue<int> best;
dp[0] = f[0];
best.push(0, dp[0]);
for (int i = 1; i < static_cast<int>(f.size()); i++) {
    best.expire(i - max_jump);
    dp[i] = f[i] + best.top().second;
    best.push(i, dp[i]);
}
assert(dp[5] == 11);
return 0;
}

```

1.2.4 Monotonic Stack

A monotonic stack maintains its elements in sorted order as values are pushed, popping away any entries that would violate the ordering. Scanning an array once while keeping such a stack of indices answers, for every position, where the nearest smaller or larger neighbor lies. These “nearest smaller/greater” relationships underlie many array problems, the most famous being the largest rectangle in a histogram. Each index is pushed and popped at most once, so a full scan runs in linear time. A largest all-zero submatrix can then be found by sweeping rows, turning each row into a histogram of consecutive zeros ending at that row.

For any histogram bar, the nearest strictly lower bars to its left and right are the first positions that prevent a rectangle at the current bar’s height from extending farther. The widest such rectangle therefore spans the positions strictly between those boundaries. For a binary matrix, the height at each column counts consecutive zeros ending in the current row; every all-zero rectangle has some bottom row, so solving one histogram per row considers every possible rectangle.

All functions below take a vector `a` of n comparable values and return a vector of n indices. A “less” query uses a strictly smaller neighbor and a “greater” query uses a strictly larger one; changing the comparison from strict to non-strict (e.g. `>=` to `>`) toggles how ties are handled.

- `previous_less(a)` returns, for each i , the largest index $j < i$ with $a[j] < a[i]$, or -1 if there’s no such index.
- `next_less(a)` returns, for each i , the smallest index $j > i$ with $a[j] < a[i]$, or n if there’s no such index.
- `previous_greater(a)` returns, for each i , the largest index $j < i$ with $a[j] > a[i]$, or -1 if there’s no such index.
- `next_greater(a)` returns, for each i , the smallest index $j > i$ with $a[j] > a[i]$, or n if there’s no such index.
- `largest_histogram_rectangle(heights)` returns the maximum area of an axis-aligned rectangle that fits under the given histogram, given an array of `heights` where each bar has width 1.
- `largest_zero_submatrix(a)` returns the area of the largest all-zero rectangular submatrix of the 0/1 matrix `a`.

Histogram areas must fit in the height value type.

Time Complexity:

- $O(n)$ per call to all one-dimensional functions, where n is the size of the input.
- $O(r \cdot c)$ per call to `largest_zero_submatrix()`, where r and c are the matrix dimensions.

Space Complexity:

- $O(n)$ auxiliary and $O(n)$ for the returned indices from each neighbor query.
- $O(n)$ auxiliary for `largest_histogram_rectangle()`.
- $O(c)$ auxiliary for `largest_zero_submatrix()`.

```

#include <algorithm>
#include <cassert>
#include <stack>
#include <vector>

template<typename T>
std::vector<int> previous_less(const std::vector<T> &a) {
    int n = static_cast<int>(a.size());
    std::vector<int> res(n);
    std::stack<int> s;
    for (int i = 0; i < n; i++) {
        while (!s.empty() && a[s.top()] >= a[i]) {
            s.pop();
        }
        res[i] = s.empty() ? -1 : s.top();
        s.push(i);
    }
    return res;
}

template<typename T>
std::vector<int> next_less(const std::vector<T> &a) {
    int n = static_cast<int>(a.size());
    std::vector<int> res(n);
    std::stack<int> s;
    for (int i = n - 1; i >= 0; i--) {
        while (!s.empty() && a[s.top()] >= a[i]) {
            s.pop();
        }
        res[i] = s.empty() ? n : s.top();
        s.push(i);
    }
    return res;
}

template<typename T>
std::vector<int> previous_greater(const std::vector<T> &a) {
    int n = static_cast<int>(a.size());
    std::vector<int> res(n);
    std::stack<int> s;
    for (int i = 0; i < n; i++) {
        while (!s.empty() && a[s.top()] <= a[i]) {
            s.pop();
        }
        res[i] = s.empty() ? -1 : s.top();
        s.push(i);
    }
    return res;
}

template<typename T>
std::vector<int> next_greater(const std::vector<T> &a) {
    int n = static_cast<int>(a.size());
    std::vector<int> res(n);
    std::stack<int> s;
    for (int i = n - 1; i >= 0; i--) {
        while (!s.empty() && a[s.top()] <= a[i]) {
            s.pop();
        }
        res[i] = s.empty() ? n : s.top();
        s.push(i);
    }
    return res;
}

template<typename T>
T largest_histogram_rectangle(const std::vector<T> &heights) {
    int n = static_cast<int>(heights.size());

```

```

std::vector<int> left = previous_less(heights), right = next_less(heights);
T best = 0;
for (int i = 0; i < n; i++) {
    T area = heights[i] * (right[i] - left[i] - 1); // Overflow warning.
    if (area > best) {
        best = area;
    }
}
return best;
}

int largest_zero_submatrix(const std::vector<std::vector<char>> &a) {
    if (a.empty()) {
        return 0;
    }
    int rows = static_cast<int>(a.size()), cols = static_cast<int>(a[0].size());
    int best = 0;
    std::vector<int> height(cols, 0);
    for (int i = 0; i < rows; i++) {
        for (int j = 0; j < cols; j++) {
            height[j] = a[i][j] ? 0 : height[j] + 1;
        }
        best = std::max(best, largest_histogram_rectangle(height));
    }
    return best;
}

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{2, 1, 4, 3};
    assert((previous_less(a) == vector<int>{-1, -1, 1, 1}));
    assert((next_less(a) == vector<int>{1, 4, 3, 4}));
    assert((previous_greater(a) == vector<int>{-1, 0, -1, 2}));
    assert((next_greater(a) == vector<int>{2, 2, 4, 4}));

    vector<int> hist{2, 1, 5, 6, 2, 3};
    // The best histogram rectangle uses bars 2 and 3: min height 5 times width 2.
    assert(largest_histogram_rectangle(hist) == 10);

    vector<vector<char>> grid{
        {1, 0, 1, 1, 0, 0},
        {1, 0, 0, 1, 0, 0},
        {0, 0, 0, 0, 0, 1},
        {1, 0, 0, 1, 0, 0},
        {1, 0, 1, 0, 0, 1}
    };
    // The largest zero rectangle has 3 rows and 2 columns.
    assert(largest_zero_submatrix(grid) == 6);
    return 0;
}

```

1.2.5 Aggregate Queue

Maintain a first-in, first-out queue that additionally answers, in $O(1)$ amortized time, the fold of an associative operation over all elements currently in the queue. This is the “sliding window aggregation” (SWAG) technique: by pushing new elements and popping old ones as a window advances, it reports an aggregate of every window in linear total time.

Unlike the monotonic deque used for sliding-window extrema (see 1.2.3), the operation need only be associative, not order-based, so the same structure handles min, max, sum, gcd, matrix products, and other monoids. It is sometimes called a “monotonic queue,” but that term usually denotes the monotonic deque of 1.2.3 (which handles only min/max); the two are different structures. The queue is built from two stacks: a back stack receives pushes while a front stack serves pops, and each stack stores running folds so the two halves combine in $O(1)$. Whenever the front stack empties, the back stack is flipped onto it; this costs $O(1)$ amortized per element, since each element is moved between stacks at most once.

The fold is defined by an associative `combine(a, b)` function. The default code below returns the “min” of the elements; for “sum”, `combine(a, b)` should return `a + b`. The fold is taken in queue order (front to back), so non-commutative operations such as matrix products are also supported.

- `AggregateQueue<T>()` constructs an empty queue.
- `empty()` returns whether the queue is empty.
- `size()` returns the number of elements in the queue.
- `push(x)` appends `x` to the back of the queue.
- `pop()` removes the element at the front of the queue, which must be non-empty.
- `front()` returns the element at the front of the queue, which must be non-empty.
- `fold()` returns `combine()` applied to all elements in queue order; the queue must be non-empty.

Time Complexity:

- $O(1)$ per call to the constructor, `empty()`, `size()`, `front()`, and `fold()`.
- $O(1)$ amortized per call to `push()` and `pop()`.

Space Complexity:

- $O(n)$ for storage of the queue elements, where n is the number of elements.
- $O(1)$ auxiliary per operation, amortized (a `pop()` may flip the back stack onto the front stack).

```
#include <algorithm>
#include <cassert>
#include <utility>
#include <vector>

template<typename T>
class AggregateQueue {
    static T combine(const T &a, const T &b) { return std::min(a, b); }

    // Each stack entry is (value, fold), where fold combines the value with all entries between it
    // and its end of the queue (the front for front_stack, the back for back_stack).
    std::vector<std::pair<T, T>> front_stack, back_stack;

public:
    bool empty() const { return front_stack.empty() && back_stack.empty(); }
    int size() const { return static_cast<int>(front_stack.size() + back_stack.size()); }

    void push(const T &x) {
        T acc = back_stack.empty() ? x : combine(back_stack.back().second, x);
        back_stack.emplace_back(x, acc);
    }

    void pop() {
        assert(!empty());
        if (front_stack.empty()) {
            // Flip the back stack onto the front stack, reversing order so the oldest element ends up on
            // top, and recompute the front folds in queue order.
            while (!back_stack.empty()) {
                T x = back_stack.back().first;
                back_stack.pop_back();
                T acc = front_stack.empty() ? x : combine(x, front_stack.back().second);
                front_stack.emplace_back(x, acc);
            }
        }
        front_stack.pop_back();
    }
};
```

```

}

T front() const {
    assert(!empty());
    return front_stack.empty() ? back_stack.front().first : front_stack.back().first;
}

T fold() const {
    assert(!empty());
    if (front_stack.empty()) {
        return back_stack.back().second;
    }
    if (back_stack.empty()) {
        return front_stack.back().second;
    }
    return combine(front_stack.back().second, back_stack.back().second);
}
};

```

Example Usage

```

using namespace std;

int main() {
    // Used directly as a FIFO queue with an O(1) min query.
    AggregateQueue<int> q;
    assert(q.empty());
    q.push(5);
    q.push(3);
    q.push(8);
    assert(!q.empty());
    assert(q.front() == 5 && q.size() == 3);
    assert(q.fold() == 3); // min(5, 3, 8)
    q.pop(); // Removes 5.
    assert(q.fold() == 3); // min(3, 8)
    q.pop(); // Removes 3.
    assert(q.fold() == 8);
    q.pop(); // Removes 8.
    assert(q.empty());

    // Canonical use: the minimum of every width-3 window, advancing the window by push/pop.
    vector<int> a{4, 2, 5, 1, 3, 6};
    AggregateQueue<int> w;
    vector<int> mins;
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        w.push(a[i]);
        if (w.size() > 3) {
            w.pop();
        }
        if (w.size() == 3) {
            mins.push_back(w.fold());
        }
    }
    assert((mins == vector<int>{2, 1, 1, 1})); // min of [4,2,5], [2,5,1], [5,1,3], [1,3,6]
    return 0;
}

```

1.2.6 Maximum Subarray Sum (Kadane)

Given an array of numbers, determine the maximum possible sum of any contiguous subarray. Kadane's algorithm scans the array while maintaining a nonnegative running sum. At each position, it adds the current value, records the largest sum seen, and resets the running sum to zero if it becomes negative. Discarding a negative sum is safe

because it could only reduce every subarray extending it. This can be adapted to compute the maximal submatrix sum as well.

Two related scans maintain slightly different states. For a circular array, the best wrapped subarray is the whole-array sum minus the minimum-sum subarray. For a maximum-product subarray, a negative value swaps the roles of the minimum and maximum products ending at the previous position, so both products must be tracked.

- `max_subarray_sum(lo, hi)` returns the sum and inclusive endpoints (`sum`, `begin`, `end`) for the maximal-sum subarray of $[lo, hi)$, where `lo` and `hi` are random-access iterators to numeric types. This implementation requires operators `+` and `<` on the iterators' value type. By convention, the empty subarray is allowed, so an input range containing only negative values returns sum 0 and endpoints $[0, -1]$.
- `max_circular_subarray_sum(lo, hi)` returns a tuple (`sum`, `begin`, `end`) containing the maximal circular-subarray sum and its inclusive endpoints. If `begin` $\leq$ `end`, the result is an ordinary range; if `begin` $>$ `end`, it wraps from `begin` through the end of the input and continues through `end`. The empty subarray is allowed under the same convention as `max_subarray_sum()`.
- `max_product_subarray(lo, hi)` returns a tuple (`product`, `begin`, `end`) containing the product and inclusive endpoints of a nonempty maximal-product subarray, or product 0 and endpoints $[0, -1]$ for an empty input range.
- `max_submatrix_sum(a)` returns the sum and inclusive boundaries (`sum`, `r1`, `c1`, `r2`, `c2`) for the largest rectangular submatrix of a matrix `a`. This implementation requires operators `+` and `<` on the matrix value type. By convention, the empty submatrix is allowed, so a matrix containing only negative values returns sum 0 with empty row and column intervals.

Overflow warning: All resulting sums and products must fit in the value type.

Time Complexity:

- $O(n)$ per call to `max_subarray_sum()`, `max_circular_subarray_sum()`, and `max_product_subarray()`, where n is the distance between `lo` and `hi`.
- $O(m \cdot n^2)$ per call to `max_submatrix_sum()`, where m is the number of rows and n is the number of columns in the matrix.

Space Complexity:

- $O(1)$ auxiliary for each subarray function.
- $O(m)$ auxiliary for `max_submatrix_sum()`, where m is the number of rows in the matrix.

```
#include <iterator>
#include <tuple>
#include <utility>
#include <vector>

template<typename It>
auto max_subarray_sum(It lo, It hi) {
    using T = typename std::iterator_traits<It>::value_type;
    if (lo == hi) {
        return std::make_tuple(T(), 0, -1);
    }
    int curr_begin = 0, begin = 0, end = -1;
    T sum = 0, max_sum = 0;
    for (It it = lo; it != hi; ++it) {
        sum += *it; // Overflow warning.
        if (sum < 0) {
            sum = 0;
            curr_begin = (it - lo) + 1;
        } else if (max_sum < sum) {
            max_sum = sum;
            begin = curr_begin;
            end = it - lo;
        }
    }
    return std::make_tuple(max_sum, begin, end);
}
```

```

template<typename It>
auto max_circular_subarray_sum(It lo, It hi) {
    using T = typename std::iterator_traits<It>::value_type;
    auto [max_sum, begin, end] = max_subarray_sum(lo, hi);
    if (lo == hi) {
        return std::make_tuple(max_sum, begin, end);
    }
    int curr_begin = 0, min_begin = 0, min_end = -1;
    T total = 0, min_sum = 0, curr_sum = 0;
    for (It it = lo; it != hi; ++it) {
        total += *it;    // Overflow warning.
        curr_sum += *it; // Overflow warning.
        if (0 < curr_sum) {
            curr_sum = 0;
            curr_begin = (it - lo) + 1;
        } else if (curr_sum < min_sum) {
            min_sum = curr_sum;
            min_begin = curr_begin;
            min_end = it - lo;
        }
    }
    T wrapped_sum = total - min_sum;
    if (max_sum < wrapped_sum) {
        max_sum = wrapped_sum;
        int n = static_cast<int>(hi - lo);
        begin = (min_end + 1) % n;
        end = (min_begin + n - 1) % n;
    }
    return std::make_tuple(max_sum, begin, end);
}

template<typename It>
auto max_product_subarray(It lo, It hi) {
    using T = typename std::iterator_traits<It>::value_type;
    if (lo == hi) {
        return std::make_tuple(T(), 0, -1);
    }
    T max_ending = *lo, min_ending = *lo, best = *lo;
    int max_begin = 0, min_begin = 0, begin = 0, end = 0;
    for (It it = lo + 1; it != hi; ++it) {
        int i = it - lo;
        if (*it < 0) {
            std::swap(max_ending, min_ending);
            std::swap(max_begin, min_begin);
        }
        T next_max = max_ending * *it; // Overflow warning.
        T next_min = min_ending * *it; // Overflow warning.
        if (next_max < *it) {
            max_ending = *it;
            max_begin = i;
        } else {
            max_ending = next_max;
        }
        if (*it < next_min) {
            min_ending = *it;
            min_begin = i;
        } else {
            min_ending = next_min;
        }
        if (best < max_ending) {
            best = max_ending;
            begin = max_begin;
            end = i;
        }
    }
    return std::make_tuple(best, begin, end);
}

```

```

template<typename T>
std::tuple<T, int, int, int, int> max_submatrix_sum(const std::vector<std::vector<T>> &a) {
    if (a.empty() || a[0].empty()) {
        return {T(), 0, 0, -1, -1};
    }
    int rows = static_cast<int>(a.size()), cols = static_cast<int>(a[0].size());
    int r1 = 0, c1 = 0, r2 = -1, c2 = -1;
    T max_sum = 0;
    for (int clo = 0; clo < cols; clo++) {
        std::vector<T> sums(rows, 0);
        for (int chi = clo; chi < cols; chi++) {
            for (int i = 0; i < rows; i++) {
                sums[i] += a[i][chi]; // Overflow warning.
            }
            auto [sum, rlo, rhi] = max_subarray_sum(sums.begin(), sums.end());
            if (max_sum < sum) {
                max_sum = sum;
                r1 = rlo;
                c1 = clo;
                r2 = rhi;
                c2 = chi;
            }
        }
    }
    return {max_sum, r1, c1, r2, c2};
}

```

Example Usage

```

#include <cassert>
#include <iostream>
using namespace std;

int main() {
    {
        vector<int> a{-2, -1, -3, 4, -1, 2, 1, -5, 4};
        auto [neg_sum, neg_begin, neg_end] = max_subarray_sum(a.begin(), a.begin() + 3);
        assert(neg_sum == 0 && neg_begin == 0 && neg_end == -1);
        auto [sum, begin, end] = max_subarray_sum(a.begin(), a.end());
        assert(sum == 6);
        cout << "Maximal sum subarray:" << endl;
        for (int i = begin; i <= end; i++) {
            cout << a[i] << " ";
        }
        cout << endl;

        vector<int> circular{5, -3, 5};
        auto [circular_sum, circular_begin, circular_end] =
            max_circular_subarray_sum(circular.begin(), circular.end());
        assert(circular_sum == 10 && circular_begin == 2 && circular_end == 0);

        vector<int> b{2, 3, -2, 4};
        auto [product, product_begin, product_end] = max_product_subarray(b.begin(), b.end());
        assert(product == 6 && product_begin == 0 && product_end == 1);
    }
    {
        vector<vector<int>> a{{0, -2, -7, 0, 5},
                             {9, 2, -6, 2, -4},
                             {-4, 1, -4, 1, 0},
                             {-1, 8, 0, -2, 3}};
        auto [sum, r1, c1, r2, c2] = max_submatrix_sum(a);
        assert(sum == 15);
        cout << "\nMaximal sum submatrix:" << endl;
        for (int i = r1; i <= r2; i++) {
            for (int j = c1; j <= c2; j++) {
                cout << a[i][j] << " ";
            }
        }
    }
}

```

```

    }
    cout << endl;
}
}
return 0;
}

```

Example Output

Maximal sum subarray:
4 -1 2 1

Maximal sum submatrix:
9 2
-4 1
-1 8

1.2.7 Longest Increasing Subsequence

Given a range of elements, determine a longest subsequence where all elements are in strictly increasing order. The subsequence is not necessarily contiguous or unique, so only one such answer will be found. The answer is computed by scanning left to right while maintaining, for each length, the smallest possible tail value of an increasing subsequence of that length. Each element extends or improves one of these tails, located in $O(\log n)$ by binary search; predecessor links then reconstruct the subsequence. The comparator `comp` defines the ordering and defaults to `std::less<>`. This implementation requires random-access iterators.

- `longest_increasing_subsequence(lo, hi, comp = std::less<>())` returns the indices relative to `lo` of one longest strictly increasing subsequence of `[lo, hi)` according to `comp`, in order.

Time Complexity: $O(n \log n)$ per call, where n is the distance between `lo` and `hi`.

Space Complexity: $O(n)$ auxiliary and $O(n)$ for the returned indices.

```

#include <algorithm>
#include <functional>
#include <vector>

template<typename It, typename Compare = std::less<>>
std::vector<int> longest_increasing_subsequence(It lo, It hi, Compare comp = Compare()) {
    int n = static_cast<int>(hi - lo);
    if (n == 0) {
        return {};
    }
    int len = 0;
    // prev[i] = index of the predecessor of element i in its LIS (or -1 for the first LIS element).
    // tail[i] = index (into [lo, hi)) of the last element of the best known LIS of length i + 1.
    std::vector<int> prev(n), tail(n);
    for (int i = 0; i < n; i++) {
        // Find the tail to extend or improve. Comparing by value through the stored indices keeps tail
        // index-based for reconstruction; use an upper-bound search to allow a non-decreasing result.
        int pos = static_cast<int>(
            std::lower_bound(
                tail.begin(), tail.begin() + len, i,
                [&](int t, int x) { return comp(*(lo + t), *(lo + x)); }
            ) -
            tail.begin()
        );
        if (pos == len) {
            len++;
        }
        prev[i] = pos > 0 ? tail[pos - 1] : -1;
        tail[pos] = i;
    }
    std::vector<int> res(len);
}

```

```

    for (int i = tail[len - 1]; i != -1; i = prev[i]) {
        res[--len] = i;
    }
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> a{-2, -5, 1, 9, 10, 8, 11, 10, 13, 11};
    vector<int> res = longest_increasing_subsequence(a.begin(), a.end());
    assert((res == vector<int>{1, 2, 3, 4, 6, 8}));

    vector<int> b{1, 4, 3, 2};
    assert(
        (longest_increasing_subsequence(b.begin(), b.end(), greater<int>()) == vector<int>{1, 2, 3})
    );
    return 0;
}

```

1.2.8 Minimum Excluded Value

Given a set of integers, find the minimum excluded nonnegative value (MEX), that is, the smallest integer $x \geq 0$ that does not appear in the set. MEX often appears in array problems, game theory with Grundy numbers, and dynamic programming states. Since the MEX of n values never exceeds n , the one-shot version tracks seen values in $[0, n]$ before scanning to find the first missing.

- `mex(lo, hi)` returns the MEX of the values in $[lo, hi)$.

The MEX of a dynamic multiset can be computed using counts of present values and an ordered set of currently missing candidates.

- `DynamicMex()` maintains the MEX of a multiset of nonnegative integers.
- `add(x)` inserts one copy of value `x` if `x` ≥ 0 .
- `remove(x)` removes one copy of value `x` if present.
- `get()` returns the current MEX.

Time Complexity:

- $O(n)$ per call to `mex(lo, hi)`, where n is the number of values.
- $O(\log n)$ expected per call to `add(x)` and `remove(x)` for `DynamicMex`, and $O(n)$ in the collision-heavy worst case for the hash table.
- $O(1)$ per call to `get()` for `DynamicMex`.

Space Complexity:

- $O(n)$ auxiliary for `mex(lo, hi)`.
- $O(n)$ storage for `DynamicMex`.

```

#include <iterator>
#include <set>
#include <unordered_map>
#include <vector>

template<typename It>
int mex(It lo, It hi) {
    int n = std::distance(lo, hi);
    std::vector<char> seen(n + 1, false);
}

```

```

    for (It it = lo; it != hi; ++it) {
        if (0 <= *it && *it <= n) {
            seen[*it] = true;
        }
    }
    for (int x = 0; x <= n; x++) {
        if (!seen[x]) {
            return x;
        }
    }
    return n + 1;
}

class DynamicMex {
    std::unordered_map<int, int> count;
    std::set<int> missing; // Missing candidates in [0, size].
    int size = 0;

public:
    DynamicMex() { missing.insert(0); }

    void add(int x) {
        if (x < 0) {
            return;
        }
        if (!count.count(size + 1)) {
            missing.insert(size + 1);
        }
        if (count[x]++ == 0) {
            missing.erase(x);
        }
        size++;
    }

    void remove(int x) {
        auto it = count.find(x);
        if (it == count.end()) {
            return;
        }
        bool removed_last = --it->second == 0;
        if (removed_last) {
            count.erase(it);
        }
        size--;
        missing.erase(size + 1);
        if (removed_last && x <= size) {
            missing.insert(x);
        }
    }

    int get() const { return *missing.begin(); }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> a{0, 1, 4, 2, 1};
    assert(mex(a.begin(), a.end()) == 3);
    vector<int> b{-1, 0, 1, 2};
    assert(mex(b.begin(), b.end()) == 3); // Negative values are ignored.

    DynamicMex m;
    m.add(0);
}

```

```

m.add(1);
m.add(1);
m.add(3);
assert(m.get() == 2);
m.add(2);
assert(m.get() == 4);
m.remove(1);
assert(m.get() == 4); // One copy of 1 remains.
m.remove(1);
assert(m.get() == 1);
m.remove(100);
assert(m.get() == 1);
return 0;
}

```

1.2.9 Iterated Function Cycle Detection

Given a function f mapping a finite set to itself and a starting value x_0 , return the entry value, position, and length of the cycle reached by repeatedly applying f . Formally, the sequence $x_0, x_1 = f(x_0), x_2 = f(x_1), \dots, x_n = f(x_{n-1}), \dots$ must eventually repeat some value. If $x_i = x_j$ for $i < j$, then the sequence from x_i to x_{j-1} repeats forever. This is useful for detecting cycles in functional graphs (see 4.2.9), pseudo-random generators, Pollard rho style algorithms, degenerate linked lists, and arrays where each index points to the next index. For example, if an array $a_0, \dots, a_{n+1}$ contains only values in $[1, n]$, with exactly one value occurring more than once, then $f(i) = a_i$ forms a functional graph; the duplicate value is the entry point of the cycle reached from 0.

Floyd’s cycle-finding algorithm, a.k.a. the “tortoise and the hare algorithm”, is a space-efficient default choice: it moves two pointers through the sequence at different speeds, finds a meeting point inside the cycle, then locates the cycle start. Brent’s algorithm is a compact variant that usually calls f fewer times by only resetting the tortoise at powers of two; prefer it when evaluating f is expensive and constant factors matter. In the detection phase, Brent advances one pointer with one call to f per step, while Floyd advances the tortoise once and the hare twice, using three calls to f per step.

For Floyd’s algorithm, suppose the tail has length m and the cycle has length n . At a meeting, the hare has traveled twice as far as the tortoise, so the tortoise’s distance is a multiple of n . Consequently, the distance from the meeting point around the cycle to its entry is congruent to m modulo n . Resetting one pointer to x_0 and advancing both one step at a time makes them meet at the cycle entry after m steps. Brent first determines n , then advances one pointer n steps; moving both together from that fixed separation likewise makes them meet at the entry.

- `find_cycle_floyd(f, x0)` repeatedly applies `f` starting from `x0` and returns the reached cycle as a tuple `(entry, start, length)`, where `entry` is the first cycle value, `start` is its index in the sequence, and `length` is the number of distinct values in the cycle.
- `find_cycle_brent(f, x0)` does the same with fewer calls to `f` in many cases.

Time Complexity: $O(m + n)$ per call, where m is the first index in the cycle and n is the cycle length.

Space Complexity: $O(1)$ auxiliary.

```

#include <tuple>

template<typename Fn, typename T>
std::tuple<T, int, int> find_cycle_floyd(Fn f, T x0) {
    // Find a meeting point inside the cycle.
    T tortoise = f(x0), hare = f(f(x0));
    while (tortoise != hare) {
        tortoise = f(tortoise);
        hare = f(f(hare));
    }
    // Find the cycle entry.
}

```

```

int start = 0;
tortoise = x0;
while (tortoise != hare) {
    tortoise = f(tortoise);
    hare = f(hare);
    start++;
}
// Measure the cycle length.
int length = 1;
hare = f(tortoise);
while (tortoise != hare) {
    hare = f(hare);
    length++;
}
return {tortoise, start, length};
}

template<typename Fn, typename T>
std::tuple<T, int, int> find_cycle_brent(Fn f, T x0) {
    // Find the cycle length by doubling the search block.
    int power = 1, length = 1;
    T tortoise = x0, hare = f(x0);
    while (tortoise != hare) {
        if (power == length) {
            tortoise = hare;
            power *= 2;
            length = 0;
        }
        hare = f(hare);
        length++;
    }
    // Separate the pointers by one cycle length.
    hare = x0;
    for (int i = 0; i < length; i++) {
        hare = f(hare);
    }
    // Find the cycle entry.
    int start = 0;
    tortoise = x0;
    while (tortoise != hare) {
        tortoise = f(tortoise);
        hare = f(hare);
        start++;
    }
    return {tortoise, start, length};
}

```

Example Usage

```

#include <cassert>
#include <set>
#include <vector>
using namespace std;

int f(int x) {
    return (123 * x * x + 4567890) % 1337;
}

void verify(int x0, int start, int length) {
    set<int> s;
    int x = x0;
    for (int i = 0; i < start; i++) {
        assert(!s.count(x));
        s.insert(x);
        x = f(x);
    }
}

```



```

int startx = x;
s.clear();
for (int i = 0; i < length; i++) {
    assert(!s.count(x));
    s.insert(x);
    x = f(x);
}
assert(startx == x);
}

// Given n + 1 integers within [1, n], where one value in [1, n] is present 2 or more times (but
// other values in [1, n] are present at most once), find the duplicate without modifying the array.
// Interpret each value as a pointer to the next index; the duplicate is the cycle entry.
int find_duplicate_integer(const vector<int> &a) {
    return get<0>(find_cycle_floyd([&](int i) { return a[i]; }, 0));
}

int main() {
    int x0 = 0;
    auto [floyd_entry, floyd_start, floyd_len] = find_cycle_floyd(f, x0);
    assert(floyd_start == 4 && floyd_len == 2);
    verify(x0, floyd_start, floyd_len);

    auto [brent_entry, brent_start, brent_len] = find_cycle_brent(f, x0);
    assert(brent_entry == floyd_entry && brent_start == floyd_start && brent_len == floyd_len);

    assert(find_duplicate_integer({1, 3, 4, 2, 2}) == 2);
    assert(find_duplicate_integer({3, 1, 3, 4, 2}) == 3);
    return 0;
}

```

1.2.10 0-1 Knapsack

Solves the 0-1 knapsack problem: given items with nonnegative integer weights and values, choose each item at most once to maximize total value without exceeding a capacity limit.

The one-dimensional dynamic programming array `dp[w]` stores the best value that can be achieved with total weight at most `w` after processing some prefix of the items. Iterating weights in decreasing order is what enforces the 0-1 constraint; it prevents the same item from being reused during the current item update.

- `knapsack_01(weight, value, capacity)` returns a pair `(best_value, items)` containing the maximum value and the selected item indices in increasing order, where item `i` has weight `weight[i]` and value `value[i]`. All weights must be nonnegative integers. Items with weight 0 are allowed and can each be taken once. The take/skip decision for every item and capacity is stored to reconstruct the optimal subset.

Time Complexity: $O(n \cdot W)$ per call, where n is the number of items and W is the capacity.

Space Complexity: $O(n \cdot W)$ auxiliary.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <utility>
#include <vector>

std::pair<int64_t, std::vector<int>> knapsack_01(
    const std::vector<int> &weight, const std::vector<int64_t> &value, int capacity
) {
    int n = static_cast<int>(weight.size());
    assert(static_cast<int>(value.size()) == n && capacity >= 0);
    for (int i = 0; i < n; i++) {
        assert(weight[i] >= 0 && value[i] >= 0);
    }
    std::vector<int64_t> dp(capacity + 1, 0);

```

```

std::vector<std::vector<char>> take(n, std::vector<char>(capacity + 1, false));
for (int i = 0; i < n; i++) {
    for (int w = capacity; w >= weight[i]; w--) {
        int64_t candidate = dp[w - weight[i]] + value[i]; // Overflow warning.
        if (dp[w] < candidate) {
            dp[w] = candidate;
            take[i][w] = true;
        }
    }
}
// Optional: reconstruct the selected items.
std::vector<int> items;
for (int i = n - 1, w = capacity; i >= 0; i--) {
    if (take[i][w]) {
        items.push_back(i);
        w -= weight[i];
    }
}
std::reverse(items.begin(), items.end());
return {dp[capacity], items};
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> weight{3, 4, 5, 9};
    vector<int64_t> value{4, 5, 7, 10};

    auto [best_value, items] = knapsack_01(weight, value, 8);
    assert(best_value == 11);
    assert((items == vector<int>{0, 2}));
    return 0;
}

```

1.2.11 Unbounded Knapsack and Coin Change

Solves unbounded knapsack and two related coin change problems, where each item or coin may be used any number of times. In the maximum-value version, each item has a weight and value, and the goal is to maximize total value subject to a capacity limit. In coin change variants, the goal is either to count unordered ways to form a target sum or to minimize the number of coins used.

For unbounded maximum-value knapsack, capacities are processed in increasing order. Thus, when computing `dp[w]`, every smaller-capacity state `dp[w - weight[i]]` is already final and may itself contain item `i`, allowing unlimited reuse. For unordered coin change, coins are the outer loop so each multiset of coins is counted once, independent of ordering. For minimum coins, each target sum chooses the best previous sum `sum - coin`, and the predecessor coin reconstructs one optimal multiset.

- `unbounded_knapsack(weight, value, capacity)` returns a pair `(best_value, count)` containing the maximum value achievable with total weight at most `capacity` and `count[i]`, the number of copies selected of each item `i`.
- `count_coin_change(coins, target)` returns the number of unordered ways to make sum `target` using the given coin denominations.
- `min_coin_change(coins, target)` returns a pair `(count, used)`, where `count` is the minimum number of coins needed to make sum `target`, and `used[i]` is the number of copies of `coins[i]` in one optimal multiset. If `target` cannot be formed, `count` is `-1` and `used` is empty.

`capacity` and `target` must be nonnegative integers. Weights and coin denominations must be positive. Coin denominations must also be distinct when counting ways; duplicate denominations are treated as different coin types and overcount identical multisets.

Time Complexity:

- $O(n \cdot W)$ per call to `unbounded_knapsack()`, where n is the number of items and W is `capacity`.
- $O(c \cdot t)$ per call to `count_coin_change()` and `min_coin_change()`, where c is the number of coin denominations and t is the target.

Space Complexity:

- $O(W)$ auxiliary and $O(n)$ for the returned `count` from `unbounded_knapsack()`.
- $O(t)$ auxiliary for `count_coin_change()`.
- $O(t)$ auxiliary and $O(c)$ for the returned `used` from `min_coin_change()`.

```
#include <cassert>
#include <cstdlib>
#include <utility>
#include <vector>

std::pair<int64_t, std::vector<int>> unbounded_knapsack(
    const std::vector<int> &weight, const std::vector<int64_t> &value, int capacity
) {
    int n = static_cast<int>(weight.size());
    assert(static_cast<int>(value.size()) == n && capacity >= 0);
    for (int w : weight) {
        assert(w > 0);
    }
    std::vector<int64_t> dp(capacity + 1, 0);
    std::vector<int> choice(capacity + 1, -1);
    for (int w = 1; w <= capacity; w++) {
        for (int i = 0; i < n; i++) {
            if (weight[i] <= w) {
                int64_t candidate = dp[w - weight[i]] + value[i]; // Overflow warning.
                if (dp[w] < candidate) {
                    dp[w] = candidate;
                    choice[w] = i;
                }
            }
        }
    }
    std::vector<int> count(n, 0);
    for (int w = capacity; choice[w] != -1; w -= weight[choice[w]]) {
        count[choice[w]]++;
    }
    return {dp[capacity], count};
}

int64_t count_coin_change(const std::vector<int> &coins, int target) {
    assert(target >= 0);
    std::vector<int64_t> dp(target + 1, 0);
    dp[0] = 1;
    for (int coin : coins) {
        assert(coin > 0);
        for (int sum = coin; sum <= target; sum++) {
            dp[sum] += dp[sum - coin]; // Overflow warning.
        }
    }
    return dp[target];
}

std::pair<int, std::vector<int>> min_coin_change(const std::vector<int> &coins, int target) {
    assert(target >= 0);
    for (int coin : coins) {
        assert(coin > 0);
    }
}
```

```

int inf = target + 1;
std::vector<int> dp(target + 1, inf), choice(target + 1, -1);
dp[0] = 0;
for (int sum = 1; sum <= target; sum++) {
    for (int i = 0; i < static_cast<int>(coins.size()); i++) {
        if (coins[i] <= sum && dp[sum] > dp[sum - coins[i]] + 1) {
            dp[sum] = dp[sum - coins[i]] + 1;
            choice[sum] = i;
        }
    }
}
if (dp[target] == inf) {
    return {-1, {}};
}
std::vector<int> used(coins.size(), 0);
for (int sum = target; sum > 0; sum -= coins[choice[sum]]) {
    used[choice[sum]]++;
}
return {dp[target], used};
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> weight{3, 4, 5};
    vector<int64_t> value{4, 5, 7};
    auto [best_value, count] = unbounded_knapsack(weight, value, 10);
    assert(best_value == 14);
    assert((count == vector<int>{0, 0, 2}));

    vector<int> coins{1, 2, 5};
    assert(count_coin_change(coins, 5) == 4); // 5, 2+2+1, 2+1+1+1, 1*5.

    auto [coin_count, used] = min_coin_change(coins, 11);
    assert(coin_count == 3);
    assert((used == vector<int>{1, 0, 2}));

    auto [impossible, empty_used] = min_coin_change(vector<int>{4, 7}, 6);
    assert(impossible == -1 && empty_used.empty());
    return 0;
}

```

1.2.12 Subset Sum

Given a collection of integers and a target v , the subset sum problem asks for the maximum sum of any subset that is at most v . Two algorithms below solve this problem in complementary regimes, and which one wins depends on whether the bottleneck is the number of items or the size of the weights.

The meet-in-the-middle method splits the collection in half, enumerates every subset sum for each half, sorts one side, and binary searches for the best compatible partner. It makes no assumption about the sign or size of the values, so it handles negative inputs, and it is the method of choice when the number of items n is small (up to roughly 40).

Pisinger's algorithm targets the opposite regime: nonnegative weights whose maximum value is small, with possibly many items. Let W be the largest weight. It greedily takes a maximal prefix with sum s that fits the target; because the next weight is at most W and does not fit, s is greater than $\text{target} - W$. The prefix itself is feasible, so the optimum lies between s and target and only that width- W window matters. A balanced sweep explores exchanges that add later items to sums below the target and remove earlier items from sums above it.

For each sum, it retains the farthest prefix boundary reached, which dominates smaller boundaries because it permits at least the same future exchanges. This gives a running time linear in the number of items times W , independent of the target beyond that window. This routine is sometimes called “fast knapsack,” a slight misnomer: it maximizes a subset sum rather than packing items with distinct values, so it is really the bounded-weight case of subset sum rather than the 0-1 value knapsack.

- `max_subset_sum_at_most(lo, hi, v)` returns a pair `(sum, items)` containing that maximum sum and the selected item indices relative to `lo`. Values may be negative. The range is supplied as random-access iterators, and `v` must be nonnegative.
- `max_subset_sum_bounded(lo, hi, target)` returns the maximum sum of any subset of the range `[lo, hi)` that does not exceed `target`. All weights and `target` must be nonnegative integers. This is faster than meet-in-the-middle when the largest weight is small relative to $2^{n/2}$.

All subset sums and intermediate differences in `max_subset_sum_at_most()` must fit in `int64_t`.

Time Complexity:

- $O(2^{n/2})$ per call to `max_subset_sum_at_most()`, where n is the distance between `lo` and `hi`.
- $O(n \cdot W)$ per call to `max_subset_sum_bounded()`, where n is the number of items and W is the largest weight.

Space Complexity:

- $O(2^{n/2})$ auxiliary for `max_subset_sum_at_most()`.
- $O(W)$ auxiliary for `max_subset_sum_bounded()`.

```
#include <algorithm>
#include <cassert>
#include <climits>
#include <cstdint>
#include <utility>
#include <vector>

template<typename It>
std::pair<int64_t, std::vector<int>> max_subset_sum_at_most(It lo, It hi, int64_t v) {
    int n = static_cast<int>(hi - lo);
    assert(v >= 0 && n >= 0 && n / 2 < 63 && n - n / 2 < 63);
    int64_t llen = 1LL << (n / 2), hlen = 1LL << (n - n / 2);
    std::vector<int64_t> lsum(llen);
    std::vector<std::pair<int64_t, int64_t>> hsum(hlen);
    for (int64_t mask = 1; mask < llen; mask++) {
        int bit = __builtin_ctzll(mask);
        lsum[mask] = lsum[mask ^ (1LL << bit)] + *(lo + bit); // Overflow warning.
    }
    for (int64_t mask = 1; mask < hlen; mask++) {
        int bit = __builtin_ctzll(mask);
        hsum[mask].first = hsum[mask ^ (1LL << bit)].first + *(lo + n / 2 + bit);
        hsum[mask].second = mask;
    }
    std::sort(hsum.begin(), hsum.end());
    int64_t best = INT64_MIN, lmask = 0, hmask = 0;
    for (int64_t mask = 0; mask < llen; mask++) {
        int64_t limit = v - lsum[mask]; // Overflow warning.
        auto it = std::upper_bound(hsum.begin(), hsum.end(), std::make_pair(limit, INT64_MAX));
        if (it != hsum.begin()) {
            --it;
            int64_t candidate = lsum[mask] + it->first;
            if (best < candidate) {
                best = candidate;
                lmask = mask;
                hmask = it->second;
            }
        }
    }
    std::vector<int> items;
    for (int i = 0; i < n / 2; i++) {
        if ((lmask >> i) & 1) {
```

```

        items.push_back(i);
    }
}
for (int i = 0; i < n - n / 2; i++) {
    if ((hmask >> i) & 1) {
        items.push_back(n / 2 + i);
    }
}
return {best, items};
}

template<typename It>
int64_t max_subset_sum_bounded(It lo, It hi, int target) {
    int n = static_cast<int>(hi - lo);
    assert(target >= 0 && n >= 0);
    for (It it = lo; it != hi; ++it) {
        int64_t weight = *it;
        assert(0 <= weight && weight <= INT_MAX);
    }
    // Greedily take a prefix while it still fits; the optimum then lies within one weight of target.
    int sum = 0, b = 0;
    while (b < n && static_cast<int>(*(lo + b)) <= target - sum) {
        sum += static_cast<int>(*(lo + b++));
    }
    if (b == n) {
        return sum; // Every item fits, so the whole collection is the best subset.
    }
    int m = 0;
    for (It it = lo; it != hi; ++it) {
        m = std::max(m, static_cast<int>(*it));
    }
    // reach[s - target + m] = largest prefix boundary that can realize a balanced subset of sum s.
    // Only sums within m of target are tracked, which is where the optimum must lie.
    std::vector<int> reach(2 * m, -1);
    reach[sum - target + m] = b;
    for (int i = b; i < n; i++) {
        std::vector<int> prev = reach;
        int wi = static_cast<int>(*(lo + i));
        for (int x = 0; x < m; x++) { // Add item i to sums currently below target.
            reach[x + wi] = std::max(reach[x + wi], prev[x]);
        }
        for (int x = 2 * m - 1; x > m; x--) { // Drop an earlier item from sums above target.
            for (int j = std::max(0, prev[x]); j < reach[x]; j++) {
                int wj = static_cast<int>(*(lo + j));
                reach[x - wj] = std::max(reach[x - wj], j);
            }
        }
    }
    int s = target;
    while (reach[s - target + m] < 0) {
        s--;
    }
    return s;
}

```

Example Usage

```

#include <cassert>

int main() {
    std::vector<int> a{9, 1, 5, 0, 1, 11, 5};
    auto [sum, items] = max_subset_sum_at_most(a.begin(), a.end(), 8);
    int64_t selected_sum = 0;
    for (int i : items) {
        selected_sum += a[i];
    }
}

```

```

assert(sum == 7 && selected_sum == sum);
std::vector<int> b{-7, -3, -2, 5, 8};
assert(max_subset_sum_at_most(b.begin(), b.end(), 0).first == 0);

// The bounded-weight method agrees with meet-in-the-middle on nonnegative inputs.
assert(max_subset_sum_bounded(a.begin(), a.end(), 8) == 7);
std::vector<int> c{3, 34, 4, 12, 5, 2};
assert(max_subset_sum_bounded(c.begin(), c.end(), 9) == 9);    // 4 + 5.
assert(max_subset_sum_bounded(c.begin(), c.end(), 10) == 10);  // 3 + 5 + 2.
assert(max_subset_sum_bounded(c.begin(), c.end(), 100) == 60); // All items fit.
return 0;
}

```

1.2.13 Grid Dynamic Programming

Solves basic dynamic programming problems on a rectangular grid. In grid DP, each cell represents a state and transitions usually come from already-computed neighboring cells, most commonly the cell above and the cell to the left. Processing rows from top to bottom and columns from left to right therefore makes every predecessor available when a state is computed.

For path-counting, `dp[r][c]` stores the number of ways to reach cell (r, c) . For minimum-cost paths, `dp[r][c]` stores the cheapest cost to reach that cell, and a predecessor direction reconstructs one optimal path. The helpers below assume movement only down or right.

- `count_grid_paths(blocked)` returns the number of paths from the upper-left cell to the lower-right cell, moving only down or right and avoiding blocked cells marked nonzero.
- `min_grid_path_cost(cost)` returns a pair $(\text{sum}, \text{path})$ containing the minimum cost of a down/right path from the upper-left cell to the lower-right cell and the cells of one optimal path in order. If the grid is empty, `sum` is 0 and `path` is empty.

Path counts and path costs must fit in `int64_t`.

Time Complexity: $O(R \cdot C)$ per call to `count_grid_paths(blocked)` and `min_grid_path_cost(cost)`, where R and C are the grid dimensions.

Space Complexity:

- $O(C)$ auxiliary for `count_grid_paths(blocked)`.
- $O(R \cdot C)$ auxiliary and $O(R + C)$ for the returned path from `min_grid_path_cost(cost)`.

```

#include <algorithm>
#include <cstdlib>
#include <utility>
#include <vector>

int64_t count_grid_paths(const std::vector<std::vector<char>> &blocked) {
    int rows = static_cast<int>(blocked.size());
    int cols = rows == 0 ? 0 : static_cast<int>(blocked[0].size());
    if (rows == 0 || cols == 0 || blocked[0][0] || blocked[rows - 1][cols - 1]) {
        return 0;
    }
    std::vector<int64_t> dp(cols, 0);
    dp[0] = 1;
    for (int r = 0; r < rows; r++) {
        for (int c = 0; c < cols; c++) {
            if (blocked[r][c]) {
                dp[c] = 0;
            } else if (c > 0) {
                dp[c] += dp[c - 1];    // Overflow warning.
            }
        }
    }
    return dp[cols - 1];
}

```

```

}

std::pair<int64_t, std::vector<std::pair<int, int>>> min_grid_path_cost(
    const std::vector<std::vector<int>>> &cost
) {
    int rows = static_cast<int>(cost.size());
    int cols = rows == 0 ? 0 : static_cast<int>(cost[0].size());
    if (rows == 0 || cols == 0) {
        return {0, {}};
    }
    std::vector<std::vector<int64_t>> dp(rows, std::vector<int64_t>(cols));
    std::vector<std::vector<char>> parent(rows, std::vector<char>(cols, 0));
    dp[0][0] = cost[0][0];
    for (int r = 0; r < rows; r++) {
        for (int c = 0; c < cols; c++) {
            if (r == 0 && c == 0) {
                continue;
            }
            if (c == 0 || (r > 0 && dp[r - 1][c] <= dp[r][c - 1])) {
                dp[r][c] = dp[r - 1][c] + cost[r][c]; // Overflow warning.
                parent[r][c] = 'U';
            } else {
                dp[r][c] = dp[r][c - 1] + cost[r][c];
                parent[r][c] = 'L';
            }
        }
    }
    std::vector<std::pair<int, int>> path;
    for (int r = rows - 1, c = cols - 1;;) {
        path.emplace_back(r, c);
        if (r == 0 && c == 0) {
            break;
        }
        if (parent[r][c] == 'U') {
            r--;
        } else {
            c--;
        }
    }
    std::reverse(path.begin(), path.end());
    return {dp[rows - 1][cols - 1], path};
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<vector<char>> blocked{
        {0, 0, 0},
        {0, 1, 0},
        {0, 0, 0},
    };
    assert(count_grid_paths(blocked) == 2);

    vector<vector<int>> cost{
        {1, 3, 1},
        {1, 5, 1},
        {4, 2, 1},
    };
    auto [sum, path] = min_grid_path_cost(cost);
    assert(sum == 7);
    assert((path == vector<pair<int, int>>{{0, 0}, {0, 1}, {0, 2}, {1, 2}, {2, 2}}));
    return 0;
}

```


1.2.14 Matrix Chain Multiplication

Given a fixed-order chain of compatible matrices, finds a parenthesization that minimizes the number of scalar multiplications needed to compute their product. The algorithm examines only the matrix dimensions; it does not multiply matrix entries. Matrix multiplication is associative, but multiplying an a -by- b matrix with a b -by- c matrix costs abc scalar multiplications, so different parenthesizations can have very different costs. This is a canonical example of interval dynamic programming.

Write d_i for `dimensions[i]`, so matrix i has d_i rows and d_{i+1} columns. For a chain of matrices, $dp(l, r)$ is the minimum cost of multiplying matrices l through r . Trying every final split k gives the recurrence $dp(l, r) = \min_{l \leq k < r} (dp(l, k) + dp(k + 1, r) + d_l d_{k+1} d_{r+1})$. The final product multiplies a d_l -by- d_{k+1} matrix with a d_{k+1} -by- d_{r+1} matrix. Processing intervals in increasing order of length ensures that both subintervals have already been solved. The same pattern applies whenever a contiguous interval is formed by combining two smaller intervals.

- `matrix_chain_order(dimensions)` returns a pair (`operations`, `parenthesization`) containing the minimum number of scalar multiplications and one optimal parenthesization. Matrix `i` has dimensions `dimensions[i]` by `dimensions[i + 1]`; all dimensions must be positive.

Overflow warning: The minimum cost and all intermediate products must fit in `int64_t`.

Time Complexity: $O(n^3)$ per call, where n is the number of matrices. The dimension values affect the returned cost but not the running time.

Space Complexity: $O(n^2)$ auxiliary for the dynamic programming and split tables.

```
#include <cassert>
#include <cstdint>
#include <string>
#include <utility>
#include <vector>

std::string matrix_chain_parenthesization(
    const std::vector<std::vector<int>>> &split, int lo, int hi
) {
    if (lo == hi) {
        return "A" + std::to_string(lo);
    }
    int k = split[lo][hi];
    return "(" + matrix_chain_parenthesization(split, lo, k) + " * " +
        matrix_chain_parenthesization(split, k + 1, hi) + ")";
}

std::pair<int64_t, std::string> matrix_chain_order(const std::vector<int> &dimensions) {
    assert(dimensions.size() >= 2);
    for (int d : dimensions) {
        assert(d > 0);
    }
    int n = static_cast<int>(dimensions.size()) - 1;
    std::vector<std::vector<int64_t>> dp(n, std::vector<int64_t>(n, 0));
    std::vector<std::vector<int>> split(n, std::vector<int>(n, 0));
    for (int length = 2; length <= n; length++) {
        for (int lo = 0; lo + length <= n; lo++) {
            int hi = lo + length - 1;
            for (int k = lo; k < hi; k++) {
                int64_t candidate = dp[lo][k] + dp[k + 1][hi] +
                    static_cast<int64_t>(dimensions[lo]) * dimensions[k + 1] *
                    dimensions[hi + 1]; // Overflow warning.
                if (k == lo || candidate < dp[lo][hi]) {
                    dp[lo][hi] = candidate;
                    split[lo][hi] = k;
                }
            }
        }
    }
}
```

```
    return {dp[0][n - 1], matrix_chain_parenthesization(split, 0, n - 1)};
}
```

Example Usage

```
using namespace std;

int main() {
    auto [operations, parenthesization] = matrix_chain_order({40, 20, 30, 10, 30});
    assert(operations == 26000);
    assert(parenthesization == "((A0 * (A1 * A2)) * A3)");

    auto [single_operations, single_parenthesization] = matrix_chain_order({3, 5});
    assert(single_operations == 0 && single_parenthesization == "A0");
    return 0;
}
```

1.3 Greedy and Scheduling

1.3.1 Fractional Knapsack

Solves the fractional knapsack problem: given items with positive integer weights and nonnegative integer values, choose any fraction of each item to maximize total value without exceeding a capacity limit.

The greedy algorithm processes items in decreasing order of value per unit weight, taking each item in full until only part of the next item fits. This is optimal because replacing any weight taken from a lower-density item with the same weight from a higher-density item cannot decrease the total value. Unlike 0-1 knapsack, allowing fractions makes each such exchange feasible.

- `fractional_knapsack(weight, value, capacity)` returns a pair `(best_value, fraction)` containing the maximum value and the fraction of each input item taken. Item `i` has weight `weight[i]` and value `value[i]`, and `fraction[i]` is in `[0, 1]`.

Time Complexity: $O(n \log n)$ per call due to sorting, where n is the number of items.

Space Complexity: $O(n)$ auxiliary and $O(n)$ for the returned fractions.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <numeric>
#include <utility>
#include <vector>

std::pair<double, std::vector<double>> fractional_knapsack(
    const std::vector<int> &weight, const std::vector<int> &value, int capacity
) {
    int n = static_cast<int>(weight.size());
    assert(static_cast<int>(value.size()) == n && capacity >= 0);
    std::vector<int> order(n);
    std::iota(order.begin(), order.end(), 0);
    for (int i = 0; i < n; i++) {
        assert(weight[i] > 0 && value[i] >= 0);
    }
    std::sort(order.begin(), order.end(), [&](int i, int j) {
        int64_t lhs = static_cast<int64_t>(value[i]) * weight[j];
        int64_t rhs = static_cast<int64_t>(value[j]) * weight[i];
        return lhs != rhs ? lhs > rhs : i < j;
    });
    double best_value = 0;
```

```

std::vector<double> fraction(n, 0);
for (int i : order) {
    int taken = std::min(capacity, weight[i]);
    fraction[i] = static_cast<double>(taken) / weight[i];
    best_value += fraction[i] * value[i];
    capacity -= taken;
    if (capacity == 0) {
        break;
    }
}
return {best_value, fraction};
}

```

Example Usage

```

#include <cmath>
using namespace std;

int main() {
    vector<int> weight{10, 20, 30};
    vector<int> value{60, 100, 120};
    auto [best_value, fraction] = fractional_knapsack(weight, value, 50);
    assert(fabs(best_value - 240) < 1e-9);
    assert((fraction == vector<double>{1, 1, 2.0 / 3}));
    return 0;
}

```

1.3.2 Interval Scheduling

Selects the maximum number of non-overlapping intervals using the classic earliest-finish-time greedy algorithm. This is the unweighted interval scheduling problem, also known as activity selection: every interval has the same value, so the objective is to choose as many compatible intervals as possible.

The greedy choice is safe because among all intervals that could be chosen first, taking one with the earliest finish time leaves the most room for the remaining intervals. In any optimal solution, the first chosen interval can be exchanged for an earliest-finishing compatible interval without reducing the number of intervals selected. Repeating this argument after each choice proves the greedy algorithm optimal. The weighted version in 1.3.3 needs dynamic programming instead.

Intervals are represented as half-open ranges $[\text{start}, \text{finish})$, so two intervals are compatible if the next interval's `start` is at least the previous interval's `finish`.

- `schedule_intervals(intervals)` returns one maximum-size compatible subset of intervals, in the order they are selected, as original indices into an input vector of `Interval` with fields `start` and `finish`. Every interval must satisfy `start < finish`.

Time Complexity: $O(n \log n)$ per call due to sorting, where n is the number of intervals.

Space Complexity: $O(n)$ auxiliary and $O(n)$ for the returned indices.

```

#include <algorithm>
#include <cassert>
#include <climits>
#include <numeric>
#include <vector>

struct Interval {
    int start, finish;
};

std::vector<int> schedule_intervals(const std::vector<Interval> &intervals) {

```

```

int n = static_cast<int>(intervals.size());
std::vector<int> order(n);
std::iota(order.begin(), order.end(), 0);
for (const auto &iv : intervals) {
    assert(iv.start < iv.finish);
}
std::sort(order.begin(), order.end(), [&](int i, int j) {
    return intervals[i].finish != intervals[j].finish ? intervals[i].finish < intervals[j].finish
        : intervals[i].start < intervals[j].start;
});
std::vector<int> selected;
int last_finish = INT_MIN;
for (int i : order) {
    const auto &iv = intervals[i];
    if (iv.start >= last_finish) {
        selected.push_back(i);
        last_finish = iv.finish;
    }
}
return selected;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<Interval> intervals{{1, 4}, {3, 5}, {0, 6}, {5, 7}, {8, 9}};
    vector<int> selected = schedule_intervals(intervals);
    // Earliest-finish greedy chooses original indices 0, 3, 4.
    assert((selected == vector<int>{0, 3, 4}));

    vector<Interval> touching{{0, 2}, {2, 4}, {4, 5}};
    assert(schedule_intervals(touching).size() == 3);
    return 0;
}

```

1.3.3 Weighted Interval Scheduling

Selects a maximum-weight subset of non-overlapping intervals using dynamic programming and binary search. Unlike unweighted interval scheduling, the earliest finish-time greedy choice is not sufficient when intervals have weights. With intervals sorted by finish time, the best total for the first i intervals either skips interval i or adds its weight to the best total over intervals finishing no later than its start, with that predecessor located by binary search.

Intervals are represented as half-open ranges $[start, finish)$, so two intervals are compatible if the next interval's `start` is at least the previous interval's `finish`.

- `select_weighted_intervals(intervals)` returns a pair $(weight, selected)$ containing that maximum weight and the selected intervals as original input indices in execution order, from an input vector of `WeightedInterval` with fields `start`, `finish`, and `weight`. `dp[i]` stores the best answer using the first i intervals after sorting by finish time. Each interval must satisfy `start < finish`. The empty subset is allowed, so nonpositive weights need not be selected.

Accumulated weights must fit in `int64_t`.

Time Complexity: $O(n \log n)$ per call due to sorting and binary searching compatible intervals.

Space Complexity: $O(n)$ auxiliary and $O(n)$ for the returned indices.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <numeric>
#include <utility>
#include <vector>

struct WeightedInterval {
    int start, finish;
    int64_t weight;
};

std::pair<int64_t, std::vector<int>> select_weighted_intervals(
    const std::vector<WeightedInterval> &intervals
) {
    int n = static_cast<int>(intervals.size());
    std::vector<int> order(n), finish(n), prev(n);
    std::iota(order.begin(), order.end(), 0);
    for (const auto &iv : intervals) {
        assert(iv.start < iv.finish);
    }
    std::sort(order.begin(), order.end(), [&](int i, int j) {
        return intervals[i].finish != intervals[j].finish ? intervals[i].finish < intervals[j].finish
            : intervals[i].start < intervals[j].start;
    });
    for (int i = 0; i < n; i++) {
        finish[i] = intervals[order[i]].finish;
    }
    std::vector<int64_t> dp(n + 1, 0);
    std::vector<char> take(n + 1, false);
    for (int i = 1; i <= n; i++) {
        int j =
            std::upper_bound(finish.begin(), finish.begin() + i - 1, intervals[order[i - 1]].start) -
            finish.begin();
        prev[i - 1] = j;
        int64_t candidate = dp[j] + intervals[order[i - 1]].weight; // Overflow warning.
        if (dp[i - 1] < candidate) {
            dp[i] = candidate;
            take[i] = true;
        } else {
            dp[i] = dp[i - 1];
        }
    }
    std::vector<int> selected;
    for (int i = n; i > 0; i--) {
        if (take[i]) {
            selected.push_back(order[i - 1]);
            i = prev[i - 1];
        } else {
            i--;
        }
    }
    std::reverse(selected.begin(), selected.end());
    return {dp[n], selected};
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<WeightedInterval> intervals{{1, 3, 5}, {2, 5, 6}, {4, 6, 5}, {6, 7, 4}, {5, 8, 11}};
    auto [weight, selected] = select_weighted_intervals(intervals);
    // Taking ids 1 and 4 beats the earliest-finish unweighted-looking choices.
    assert(weight == 17);
}

```

```

assert((selected == vector<int>{1, 4}));

vector<WeightedInterval> touching{{0, 2, 5}, {2, 4, 6}, {1, 3, 7}};
auto [touching_weight, touching_selected] = select_weighted_intervals(touching);
assert(touching_weight == 11);
assert((touching_selected == vector<int>{0, 1}));
return 0;
}

```

1.3.4 Interval Partitioning

Assigns intervals to the minimum number of groups so that no two intervals in the same group overlap. This is the classic interval partitioning problem, also known as finding the minimum number of rooms needed for meetings.

The greedy algorithm sorts intervals by start time and reuses the group whose current finish time is earliest whenever possible. Otherwise, it creates a new group. This is optimal because the number of groups must be at least the maximum number of intervals overlapping at any time. When the earliest-finishing active group still overlaps the next interval, every active group overlaps it, so creating a new group is unavoidable. If that group is free, reusing it preserves as many later options as possible.

Intervals are treated as half-open ranges $[start, finish)$, so one interval may reuse a room that another interval vacates at the same time.

- `partition_intervals(intervals)` returns a pair `(rooms, room)`, where `rooms` is the minimum number of rooms and `room[i]` is the assigned room for input interval `i`. Each `Interval` has fields `start` and `finish`, which must satisfy `start < finish`.

Time Complexity: $O(n \log n)$ per call due to sorting and priority queue operations.

Space Complexity: $O(n)$ auxiliary and $O(n)$ for the returned assignment.

```

#include <algorithm>
#include <cassert>
#include <functional>
#include <numeric>
#include <queue>
#include <utility>
#include <vector>

struct Interval {
    int start, finish;
};

std::pair<int, std::vector<int>>> partition_intervals(const std::vector<Interval> &intervals) {
    int n = static_cast<int>(intervals.size());
    std::vector<int> order(n);
    std::iota(order.begin(), order.end(), 0);
    for (const auto &iv : intervals) {
        assert(iv.start < iv.finish);
    }
    std::sort(order.begin(), order.end(), [&](int i, int j) {
        return intervals[i].start != intervals[j].start ? intervals[i].start < intervals[j].start
            : intervals[i].finish < intervals[j].finish;
    });
    std::vector<int> room(n);
    std::priority_queue<std::pair<int, int>, std::vector<std::pair<int, int>>, std::greater<>> pq;
    int rooms = 0;
    for (int i : order) {
        if (!pq.empty() && pq.top().first <= intervals[i].start) {
            int r = pq.top().second;
            pq.pop();
            room[i] = r;
        } else {

```

```

        room[i] = rooms++;
    }
    pq.emplace(intervals[i].finish, room[i]);
}
return {rooms, room};
}

```

Example Usage

```

#include <algorithm>
#include <cassert>
using namespace std;

int main() {
    vector<Interval> intervals{{0, 30}, {5, 10}, {15, 20}};
    auto [rooms, room] = partition_intervals(intervals);
    // The long interval overlaps both short intervals, but the short intervals can share a room.
    assert(rooms == 2);
    assert(room[1] == room[2]);
    assert(room[0] != room[1]);

    vector<Interval> touching{{0, 2}, {2, 4}, {4, 5}};
    auto [touching_rooms, touching_room] = partition_intervals(touching);
    assert(touching_rooms == 1 && *max_element(touching_room.begin(), touching_room.end()) == 0);
    assert(partition_intervals({}).first == 0);
    return 0;
}

```

1.3.5 Deadline Scheduling

Schedules unit-time jobs with deadlines and profits to maximize total profit. Each job takes one time slot and must be scheduled at or before its deadline. The greedy algorithm processes jobs in decreasing profit order and places each chosen job in the latest available slot before its deadline.

The latest-slot choice is safe because it uses the current high-profit job while leaving earlier slots available for jobs with tighter deadlines. A disjoint-set forest over time slots makes this efficient: `find_root(t)` returns the latest still-free slot at or before `t`, and occupying slot `t` links it to the next candidate slot `t - 1`.

- `select_deadline_jobs(jobs)` returns a pair (`profit`, `slot`) containing that maximum profit and `slot[i]`, the assigned time slot of input job `i`, or `-1` if that job is not selected, for a vector of `Job` with fields `deadline` and `profit`. Deadlines are positive integers, and assigned slots are numbered starting from 1. Jobs with nonpositive profit are not selected.

The total profit must fit in `int64_t`.

Time Complexity: $O(n \log n)$ per call due to sorting, with near-constant disjoint-set operations.

Space Complexity: $O(n)$ auxiliary and $O(n)$ for the returned slot assignment.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <numeric>
#include <utility>
#include <vector>

struct Job {
    int deadline;
    int64_t profit;
};

class SlotDSU {

```

```

    std::vector<int> root;

public:
    explicit SlotDSU(int n) : root(n + 1) { std::iota(root.begin(), root.end(), 0); }

    int find_root(int u) {
        if (root[u] != u) {
            root[u] = find_root(root[u]);
        }
        return root[u];
    }

    void occupy(int u) { root[u] = find_root(u - 1); }
};

std::pair<int64_t, std::vector<int>> select_deadline_jobs(const std::vector<Job> &jobs) {
    int n = static_cast<int>(jobs.size());
    std::vector<int> order(n);
    std::iota(order.begin(), order.end(), 0);
    for (const auto &job : jobs) {
        assert(job.deadline > 0);
    }
    std::sort(order.begin(), order.end(), [&](int i, int j) {
        return jobs[i].profit != jobs[j].profit ? jobs[i].profit > jobs[j].profit
            : jobs[i].deadline < jobs[j].deadline;
    });
    int max_deadline = 0;
    for (const auto &j : jobs) {
        max_deadline = std::max(max_deadline, j.deadline);
    }
    max_deadline = std::min(max_deadline, n);
    SlotDSU slots(max_deadline);
    std::vector<int> assigned(n, -1);
    int64_t res = 0;
    for (int i : order) {
        if (jobs[i].profit <= 0) {
            continue;
        }
        int slot = slots.find_root(std::min(jobs[i].deadline, max_deadline));
        if (slot > 0) {
            res += jobs[i].profit; // Overflow warning.
            assigned[i] = slot;
            slots.occupy(slot);
        }
    }
    return {res, assigned};
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<Job> jobs{{2, 100}, {1, 19}, {2, 27}, {1, 25}, {3, 15}};
    auto [profit, slot] = select_deadline_jobs(jobs);
    // Jobs 0, 2, 4 fill slots 2, 1, 3 for total profit 100 + 27 + 15.
    assert(profit == 142);
    assert((slot == vector<int>{2, -1, 1, -1, 3}));
    return 0;
}

```


1.3.6 Minimizing Late Jobs (Moore-Hodgson)

Selects a maximum-size subset of jobs that can all be completed by their deadlines. Each job has a processing time and deadline, and all jobs are available at time 0. The Moore-Hodgson algorithm sorts by deadline and keeps the accepted jobs in a max-heap by processing time; whenever the schedule becomes late, it removes the longest accepted job.

After considering jobs up to some deadline, if the accepted set no longer fits, at least one accepted job must be removed. Removing the longest one leaves the most remaining time while keeping the same number of removed jobs, so it is never worse than removing a shorter accepted job. Repeating this repair after each deadline leaves a largest feasible accepted set.

- `select_on_time_jobs(jobs)` returns the selected jobs as original input indices in a feasible earliest-deadline-first execution order, for an input vector of `TimedJob` with fields `duration` and `deadline`. Durations and deadlines must be nonnegative integers.

The total duration must fit in `int64_t`.

Time Complexity: $O(n \log n)$ per call due to sorting and heap operations.

Space Complexity: $O(n)$ auxiliary and $O(n)$ for the returned indices.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <numeric>
#include <queue>
#include <utility>
#include <vector>

struct TimedJob {
    int duration, deadline;
};

std::vector<int> select_on_time_jobs(const std::vector<TimedJob> &jobs) {
    int n = static_cast<int>(jobs.size());
    std::vector<int> order(n);
    std::iota(order.begin(), order.end(), 0);
    for (const auto &job : jobs) {
        assert(job.duration >= 0 && job.deadline >= 0);
    }
    std::sort(order.begin(), order.end(), [&](int i, int j) {
        return jobs[i].deadline != jobs[j].deadline ? jobs[i].deadline < jobs[j].deadline
            : jobs[i].duration < jobs[j].duration;
    });
    std::priority_queue<std::pair<int, int>> accepted;
    std::vector<char> selected(n, false);
    int64_t time = 0;
    for (int i : order) {
        time += jobs[i].duration; // Overflow warning.
        accepted.push({jobs[i].duration, i});
        selected[i] = true;
        if (time > jobs[i].deadline) {
            time -= accepted.top().first;
            selected[accepted.top().second] = false;
            accepted.pop();
        }
    }
    std::vector<int> res;
    for (int i : order) {
        if (selected[i]) {
            res.push_back(i);
        }
    }
    return res;
}
```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    vector<TimedJob> jobs{{3, 4}, {2, 3}, {1, 2}, {2, 7}};
    assert((select_on_time_jobs(jobs) == vector<int>{2, 1, 3}));
    return 0;
}
```

1.4 Streaming and Randomized Algorithms

1.4.1 Fisher-Yates Shuffle

Randomly permutes a range in-place using the Fisher-Yates shuffle. Each possible permutation is produced with equal probability, assuming the random number source is uniform.

The algorithm scans from right to left. At each position i , it chooses a random position j from $[0, i]$ and swaps the two elements. This fixes one uniformly random remaining element at a time.

- `fisher_yates_shuffle(lo, hi)` randomly shuffles the range $[lo, hi)$ using an internal random engine. The range must provide random-access iterators.

Time Complexity: $O(n)$ per call, where n is the distance between `lo` and `hi`.

Space Complexity: $O(1)$ auxiliary.

```
#include <algorithm>
#include <random>

template<typename It>
void fisher_yates_shuffle(It lo, It hi) {
    static std::mt19937 rng(std::random_device{}());
    int n = static_cast<int>(hi - lo);
    for (int i = n - 1; i > 0; i--) {
        std::uniform_int_distribution<int> dist(0, i);
        int j = dist(rng);
        std::iter_swap(lo + i, lo + j);
    }
}
```

Example Usage

```
#include <algorithm>
#include <cassert>
#include <vector>
using namespace std;

int main() {
    vector<int> a{1, 2, 3, 4, 5}, original(a);
    fisher_yates_shuffle(a.begin(), a.end());
    sort(a.begin(), a.end());
    assert(a == original);
    return 0;
}
```

1.4.2 Reservoir Sampling

Selects uniformly random samples from a stream whose length may be unknown in advance. This is useful when values arrive online or when storing the full dataset is unnecessary.

Reservoir sampling keeps only the chosen sample or samples in memory while processing each element once. The first arrivals fill the reservoir, and the i -th element then replaces a uniformly chosen occupant with probability $1/i$ (or k/i when keeping k samples), which by induction leaves every element seen so far equally likely to be in the sample.

Each class maintains its reservoir incrementally; call `add(x)` once per stream element in any order, then call `get()` to retrieve the result.

- `ReservoirSampleOne<T>()` constructs a single-element sampler.
- `ReservoirSampleK<T>(k)` constructs a k -element sampler.
- `add(x)` incorporates one more stream element.
- `get()` returns the current sample. For `ReservoirSampleOne`, `get()` requires at least one `add()` call; for `ReservoirSampleK`, it returns the full reservoir, which may contain fewer than k elements if the stream was shorter.
- `count()` returns the number of stream elements seen so far.

Time Complexity:

- $O(1)$ per call to `ReservoirSampleOne::add()` and $O(1)$ amortized per call to `ReservoirSampleK::add()`.
- $O(n)$ to process a range of n elements.

Space Complexity:

- $O(1)$ storage for `ReservoirSampleOne`.
- $O(k)$ storage for `ReservoirSampleK`.

```
#include <cassert>
#include <optional>
#include <random>
#include <vector>

int rand_int(int lo, int hi) {
    static std::mt19937 rng(std::random_device{}());
    return std::uniform_int_distribution<int>(lo, hi)(rng);
}

template<typename T>
class ReservoirSampleOne {
    std::optional<T> sample;
    int seen = 0;

public:
    void add(const T &x) {
        seen++;
        if (rand_int(1, seen) == 1) {
            sample = x;
        }
    }

    const T &get() const {
        assert(sample.has_value());
        return *sample;
    }

    int count() const { return seen; }
};

template<typename T>
class ReservoirSampleK {
    int k, seen = 0;
```

```

std::vector<T> reservoir;

public:
explicit ReservoirSampleK(int k) : k(k) { assert(k >= 0); }

void add(const T &x) {
    ++seen;
    if (k == 0) {
        return;
    }
    if (static_cast<int>(reservoir.size()) < k) {
        reservoir.push_back(x);
    } else {
        int j = rand_int(0, seen - 1);
        if (j < k) {
            reservoir[j] = x;
        }
    }
}

const std::vector<T> &get() const { return reservoir; }
int count() const { return seen; }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> a{10, 20, 30, 40, 50};

    // Streaming interface: feed elements one at a time.
    ReservoirSampleOne<int> s1;
    for (int x : a) {
        s1.add(x);
    }
    assert(s1.count() == 5);
    int one = s1.get();
    assert(one == 10 || one == 20 || one == 30 || one == 40 || one == 50);

    ReservoirSampleK<int> sk(3);
    for (int x : a) {
        sk.add(x);
    }
    assert(sk.count() == 5);
    assert(sk.get().size() == 3);

    ReservoirSampleK<int> large(10);
    for (int x : a) {
        large.add(x);
    }
    assert(large.count() == 5);
    assert(large.get().size() == 5);
    return 0;
}

```

1.4.3 Frequent Elements

Finds frequent elements in a range using cancellation algorithms that retain only a small set of candidates. Boyer-Moore finds a possible majority element, while Misra-Gries generalizes the same idea to frequencies above any threshold n/k .

The Boyer-Moore voting algorithm maintains one candidate and counter, incrementing on a match and decrementing on a mismatch. Each decrement cancels two different values. Removing such pairs cannot eliminate an element occurring more than $\lfloor n/2 \rfloor$ times, so a true majority survives as the final candidate. A second pass verifies that candidate.

Misra-Gries keeps at most $k - 1$ candidates. A tracked value increments its counter, an untracked value claims a free counter, and otherwise every counter is decremented. Each full decrement cancels k distinct occurrences, so any value occurring more than $\lfloor n/k \rfloor$ times must remain a candidate. These candidates require a second pass for exact verification because the algorithm does not retain the stream.

- `majority_element(lo, hi)` returns an iterator to the first occurrence of the majority element of $[lo, hi)$, or `hi` if no majority element exists. The range must provide ForwardIterators, and its value type must support equality.
- `frequent_candidates(lo, hi, k)` returns a hash table of candidate values to their residual counters. `k` must be at least 2, and the value type must support equality and `std::hash`. Use `k = 2` for the unverified Boyer-Moore candidate phase.

Time Complexity:

- $O(n)$ per call to `majority_element(lo, hi)`, where n is the range length.
- $O(n)$ expected per call to `frequent_candidates(lo, hi, k)`: a full-table decrement cancels k occurrences, so all decrement sweeps take $O(n)$ total.
- $O(n \cdot k)$ per call to `frequent_candidates()` in the collision-heavy worst case.

Space Complexity:

- $O(1)$ auxiliary for `majority_element()`.
- $O(k)$ for the candidates returned by `frequent_candidates()`.

```
#include <cassert>
#include <iterator>
#include <unordered_map>

template<typename It>
It majority_element(It lo, It hi) {
    int count = 0;
    It candidate = lo;
    for (It it = lo; it != hi; ++it) {
        if (count == 0) {
            candidate = it;
            count = 1;
        } else if (*it == *candidate) {
            count++;
        } else {
            count--;
        }
    }
    // Second pass to verify. If a majority is guaranteed, skip and return candidate directly instead.
    int n = 0;
    count = 0;
    It first = hi;
    for (It it = lo; it != hi; ++it) {
        n++;
        if (*it == *candidate) {
            if (first == hi) {
                first = it;
            }
            count++;
        }
    }
    if (count <= n / 2) {
        return hi;
    }
    return first;
}
```

```

template<typename It>
auto frequent_candidates(It lo, It hi, int k) {
    using T = typename std::iterator_traits<It>::value_type;
    assert(k >= 2);
    std::unordered_map<T, int> count;
    for (It it = lo; it != hi; ++it) {
        if (auto found = count.find(*it); found != count.end()) {
            found->second++;
        } else if (static_cast<int>(count.size()) < k - 1) {
            count[*it] = 1;
        } else {
            for (auto jt = count.begin(); jt != count.end(); ) {
                if (--jt->second == 0) {
                    jt = count.erase(jt);
                } else {
                    ++jt;
                }
            }
        }
    }
    return count;
}

```

Example Usage

```

#include <vector>

int main() {
    std::vector<int> a{3, 2, 3, 1, 3};
    assert(majority_element(a.begin(), a.end()) == a.begin());
    std::vector<int> b{2, 3, 3, 3, 2, 1};
    assert(majority_element(b.begin(), b.end()) == b.end());

    std::vector<int> c{1, 2, 1, 3, 1, 2, 1, 4, 2, 2, 2};
    auto candidates = frequent_candidates(c.begin(), c.end(), 3);
    assert(candidates.count(1));
    assert(candidates.count(2));
    assert(!candidates.count(3));
    return 0;
}

```

1.4.4 Online Statistics

Maintains summary statistics as values arrive without storing the stream externally. Welford's algorithm computes the mean and variance using constant storage, while a pair of heaps maintains the exact median using linear storage.

The naive variance formula $\text{Var}(X) = E[X^2] - E[X]^2$ can lose precision when the two terms are large and close together. Welford instead maintains the current mean $\bar{x}_n$ (stored as `avg`) and the unnormalized second central moment $M_{2,n} = \sum_{i=1}^n (x_i - \bar{x}_n)^2$ (stored as `m2`). Each new value shifts the running mean by its deviation divided by the new count, then adds the product of its deviations from the old and new means to `m2`.

- `OnlineStatistics()` constructs an empty summary.
- `add(x)` incorporates one more value `x` into the summary.
- `count()` returns the number of values seen so far.
- `mean()` returns the current arithmetic mean, or 0 if the summary is empty.
- `variance_population()` returns population variance, dividing by n , or 0 if the summary is empty.
- `variance_sample()` returns sample variance, dividing by $n - 1$, or 0 if fewer than two values have been added.

For the median, a max-heap stores the lower half of the values and a min-heap stores the upper half. The heaps are balanced so that the lower heap has either the same number of values as the upper heap or one extra value. Their roots are therefore the middle value or values of the stream.

- `OnlineMedian<T>()` constructs an empty median tracker for an ordered numeric type `T`.
- `add(x)` inserts `x` into the median tracker.
- `count()` and `empty()` return the number of values and whether the tracker is empty.
- `lower_median()` returns the lower of the two middle values, or the middle value for an odd count. The tracker must be nonempty.
- `median()` returns the middle value for an odd count or the arithmetic mean of the two middle values for an even count. The tracker must be nonempty.

Time Complexity:

- $O(1)$ per call to `OnlineStatistics::add()` and each `OnlineStatistics` query.
- $O(\log n)$ per call to `OnlineMedian::add()`, where n is the number of values stored.
- $O(1)$ per `OnlineMedian` query.

Space Complexity:

- $O(1)$ for `OnlineStatistics`.
- $O(n)$ for `OnlineMedian`.

```
#include <cassert>
#include <functional>
#include <queue>
#include <vector>

class OnlineStatistics {
    int n;
    double avg, m2;

public:
    OnlineStatistics() : n(0), avg(0), m2(0) {}

    void add(double x) {
        n++;
        double delta = x - avg;
        avg += delta / n;
        double delta2 = x - avg;
        m2 += delta * delta2;
    }

    int count() const { return n; }
    double mean() const { return avg; }
    double variance_population() const { return n == 0 ? 0 : m2 / n; }
    double variance_sample() const { return n <= 1 ? 0 : m2 / (n - 1); }
};

template<typename T>
class OnlineMedian {
    std::priority_queue<T> lower;
    std::priority_queue<T, std::vector<T>, std::greater<T>> upper;

public:
    bool empty() const { return lower.empty(); }
    int count() const { return static_cast<int>(lower.size() + upper.size()); }

    void add(const T &x) {
        if (lower.empty() || x <= lower.top()) {
            lower.push(x);
        } else {
            upper.push(x);
        }
        if (lower.size() < upper.size()) {
            lower.push(upper.top());
        }
    }
};
```

```

    upper.pop();
} else if (lower.size() > upper.size() + 1) {
    upper.push(lower.top());
    lower.pop();
}
}

const T &lower_median() const {
    assert(!empty());
    return lower.top();
}

double median() const {
    assert(!empty());
    if (lower.size() != upper.size()) {
        return static_cast<double>(lower.top());
    }
    return static_cast<double>(lower.top()) / 2 + static_cast<double>(upper.top()) / 2;
}
};

```

Example Usage

```

#include <cmath>

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-9;
}

int main() {
    OnlineStatistics empty;
    assert(empty.count() == 0);
    assert(EQ(empty.mean(), 0.0));
    assert(EQ(empty.variance_population(), 0.0));
    assert(EQ(empty.variance_sample(), 0.0));

    OnlineStatistics singleton;
    singleton.add(7);
    assert(singleton.count() == 1);
    assert(EQ(singleton.mean(), 7.0));
    assert(EQ(singleton.variance_population(), 0.0));
    assert(EQ(singleton.variance_sample(), 0.0));

    OnlineStatistics stats;
    for (int x = 1; x <= 5; x++) {
        stats.add(x);
    }
    assert(stats.count() == 5);
    assert(EQ(stats.mean(), 3.0));
    assert(EQ(stats.variance_population(), 2.0));
    assert(EQ(stats.variance_sample(), 2.5));

    OnlineStatistics shifted;
    shifted.add(1e12 + 1);
    shifted.add(1e12 + 2);
    shifted.add(1e12 + 3);
    assert(EQ(shifted.mean(), 1e12 + 2));
    assert(EQ(shifted.variance_population(), 2.0 / 3));
    assert(EQ(shifted.variance_sample(), 1.0));

    OnlineMedian<int> medians;
    assert(medians.empty() && medians.count() == 0);
    medians.add(5);
    assert(EQ(medians.median(), 5));
    medians.add(1);
    assert(EQ(medians.median(), 3));
}

```



```

medians.add(9);
assert(EQ(medians.median(), 5));
medians.add(3);
assert(medians.lower_median() == 3);
assert(EQ(medians.median(), 4));
return 0;
}

```

1.5 Bit Manipulation

1.5.1 Bit Operations and Masks

Integers are all fixed-size sets of bits, where the i -th bit (counting from the least significant bit where $i = 0$) is “present” when set to 1. Treating an `int` as a bitmask makes set membership, insertion, deletion, and iteration into single machine instructions, which is why bitmasks are the backbone of subset dynamic programming and many low-level tricks. The helpers below operate on `Mask`, which is `uint32_t` by default; change its alias to `uint64_t` to use 64-bit masks. The loop implementations below exist for education only; production code should prefer the matching GCC built-ins when available, or on C++20 systems the generic `constexpr` equivalents from the `<bit>` header, which are additionally well-defined at 0.

- `test_bit(x, i)` returns whether the i -th bit of x is set, where $0 \leq i < \text{MASK_BITS}$.
- `set_bit(x, i)`, `clear_bit(x, i)`, and `toggle_bit(x, i)` return x with the i -th respectively forced to 1, forced to 0, or flipped, where $0 \leq i < \text{MASK_BITS}$.
- `lowest_set_bit(x)` returns the value of the lowest set bit of x (a power of two), or 0 if x is 0. `clear_lowest_set_bit(x)` returns x with its lowest set bit removed.
- `popcount(x)` returns the number of set bits, analogous to C++20’s `std::popcount()`.
- `parity(x)` returns 1 if the number of set bits is odd and 0 otherwise.
- `ctz(x)` returns the number of trailing 0-bits for $x > 0$, analogous to C++20’s `std::countr_zero()`.
- `ffs(x)` returns one plus the position of the lowest 1-bit of x , or 0 if x is 0.
- `clz(x)` returns the number of leading 0-bits for $x > 0$, analogous to C++20’s `std::countl_zero()`.
- `is_pow2(x)` returns whether x has exactly one set bit, analogous to C++20’s `std::has_single_bit()`.
- `floor_pow2(x)` returns the largest power of two that is $\leq x$ (for $x > 0$), analogous to C++20’s `std::bit_floor()`.
- `ceil_pow2(x)` returns the smallest power of two that is $\geq x$, for $0 < x \leq 2^{b-1}$ where b is `MASK_BITS`, analogous to C++20’s `std::bit_ceil()`.
- `for_each_set_bit(x, f)` calls `f(i)` once for each set bit position i of x , in increasing order.

Time Complexity:

- $O(b)$ worst case per call to `popcount()`, `parity()`, `ctz()`, `ffs()`, `clz()`, and `for_each_set_bit()`, where b is `MASK_BITS`.
- $O(1)$ per call to every other function.

Space Complexity: $O(1)$ auxiliary for all operations.

```

#include <cassert>
#include <climits>
#include <cstdint>

using Mask = uint32_t;
const int MASK_BITS = sizeof(Mask) * CHAR_BIT;

bool test_bit(Mask x, int i) { return (x >> i) & 1; }
Mask set_bit(Mask x, int i) { return x | (Mask{1} << i); }
Mask clear_bit(Mask x, int i) { return x & ~(Mask{1} << i); }
Mask toggle_bit(Mask x, int i) { return x ^ (Mask{1} << i); }
Mask lowest_set_bit(Mask x) { return x & -x; }

```

```

Mask clear_lowest_set_bit(Mask x) { return x & (x - 1); }

// popcount(), parity(), ctz(), ffs(), and clz() are spelled out here for educational purposes.

int popcount(Mask x) { // std::popcount(x) in C++20.
    int count = 0;
    for (; x != 0; x = clear_lowest_set_bit(x)) {
        count++;
    }
    return count;
}

int parity(Mask x) {
    return popcount(x) & 1;
}

int ctz(Mask x) { // std::countr_zero(x) in C++20.
    assert(x != 0);
    int count = 0;
    for (; (x & 1) == 0; x >>= 1) {
        count++;
    }
    return count;
}

int ffs(Mask x) {
    return x == 0 ? 0 : ctz(x) + 1;
}

int clz(Mask x) { // std::countl_zero(x) in C++20.
    assert(x != 0);
    int count = 0;
    for (Mask mask = Mask{1} << (MASK_BITS - 1); (x & mask) == 0; mask >>= 1) {
        count++;
    }
    return count;
}

bool is_pow2(Mask x) { // std::has_single_bit(x) in C++20.
    return x != 0 && (x & (x - 1)) == 0;
}

Mask floor_pow2(Mask x) { // std::bit_floor(x) in C++20.
    assert(x != 0);
    return Mask{1} << (MASK_BITS - 1 - clz(x));
}

Mask ceil_pow2(Mask x) { // std::bit_ceil(x) in C++20.
    assert(0 < x && x <= (Mask{1} << (MASK_BITS - 1)));
    return is_pow2(x) ? x : Mask{1} << (MASK_BITS - clz(x));
}

template<typename Fn>
void for_each_set_bit(Mask x, Fn f) {
    while (x != 0) {
        f(ctz(x));
        x = clear_lowest_set_bit(x);
    }
}

```

Example Usage

```

#include <cassert>
#include <vector>
using namespace std;

```

```

int main() {
    Mask x = 0b101100;
    assert(test_bit(x, 2) == true);
    assert(test_bit(x, 0) == false);
    assert(set_bit(x, 0) == 0b101101u);
    assert(clear_bit(x, 2) == 0b101000u);
    assert(toggle_bit(x, 3) == 0b100100u);

    assert(lowest_set_bit(x) == 0b100u);
    assert(clear_lowest_set_bit(x) == 0b101000u);
    assert(lowest_set_bit(0) == 0u);
    assert(clear_lowest_set_bit(0) == 0u);
    assert(popcount(x) == 3);
    assert(popcount(0) == 0);
    assert(parity(x) == 1);
    assert(parity(0b101101u) == 0);
    assert(parity(0) == 0);
    assert(ctz(x) == 2);
    assert(ffs(x) == 3);
    assert(ffs(0) == 0);
    assert(clz(x) == MASK_BITS - 6);

    assert(is_pow2(16) == true);
    assert(is_pow2(24) == false);
    assert(floor_pow2(20) == 16u);
    assert(ceil_pow2(20) == 32u);
    assert(ceil_pow2(16) == 16u);
    assert(ceil_pow2((Mask{1} << (MASK_BITS - 1)) - 1) == (Mask{1} << (MASK_BITS - 1)));

    vector<int> bits;
    for_each_set_bit(x, [&](int b) { bits.push_back(b); });
    assert((bits == vector<int>{2, 3, 5}));
    return 0;
}

```

1.5.2 Subset and Submask Iteration

Many bitmask algorithms must visit related masks in a structured order: every submask of a fixed mask, every mask with a fixed number of set bits, or every mask while accumulating contributions from its submasks. Each pattern has a short, well-known idiom.

The classic submask loop `for (int s = m; s > 0; s = (s - 1) & m)` walks every nonzero submask of `m` in decreasing order; subtracting 1 borrows through the zero bits of `m`, and the `& m` masks them back off. Summed over all masks `m` of `n` bits, the total work is $\sum_k \binom{n}{k} 2^k = 3^n$, since each of the `n` bit positions is independently absent, present in `m` only, or present in both `m` and `s`.

Gosper’s hack advances to the next mask with the same popcount. It isolates the lowest set bit and adds it to move the lowest movable 1 one position left, clearing the block of 1-bits beneath it. The remaining expression counts those cleared bits and packs them into the least-significant positions, producing the smallest larger integer with the same number of set bits.

- `submasks(m)` returns all submasks of `m`, including `m` itself and 0, in decreasing order.
- `next_popcount(x)` returns the smallest integer greater than `x` with the same number of set bits (Gosper’s hack), for `x > 0` when that next mask is representable in `Mask`.
- `masks_with_popcount(n, k)` returns all `k`-bit subsets of an `n`-bit universe as masks, in increasing order, where $0 \leq k \leq n < \text{MASK_BITS}$.
- `subset_sum_transform(f)` overwrites `f` (indexed by mask over `n` bits) so that `f[m]` becomes the sum of the original `f[s]` over all submasks `s` of `m`. This “sum over subsets” (SOS) DP is the bitmask analog of a prefix sum, accumulating one bit dimension at a time. The value type must support `+=`, and all intermediate sums must be representable.

Time Complexity:

- $O(2^p)$ per call to `submasks(m)`, where p is `popcount(m)`.
- $O(1)$ per call to `next_popcount()`.
- $O(\binom{n}{k})$ per call to `masks_with_popcount(n, k)`.
- $O(n \cdot 2^n)$ per call to `subset_sum_transform(f)`, where `f` has 2^n entries.

Space Complexity:

- $O(1)$ auxiliary for `next_popcount()` and `subset_sum_transform()`.
- $O(1)$ auxiliary and $O(s)$ for the returned masks from `submasks(m)` and `masks_with_popcount(n, k)`, where s is the result size.

```
#include <cassert>
#include <climits>
#include <cstdint>
#include <vector>

using Mask = uint32_t;
const int MASK_BITS = sizeof(Mask) * CHAR_BIT;

std::vector<Mask> submasks(Mask m) {
    std::vector<Mask> res;
    Mask s = m;
    while (true) {
        res.push_back(s);
        if (s == 0) {
            break;
        }
        s = (s - 1) & m;
    }
    return res;
}

Mask next_popcount(Mask x) {
    Mask c = x & (Mask{0} - x), r = x + c;
    assert(x != 0 && r != 0);
    return r | ((x ^ r) >> 2) / c;
}

std::vector<Mask> masks_with_popcount(int n, int k) {
    assert(0 <= k && k <= n && n < MASK_BITS);
    std::vector<Mask> res;
    if (k == 0) {
        return {0};
    }
    Mask limit = Mask{1} << n;
    for (Mask x = (Mask{1} << k) - 1; x < limit; x = next_popcount(x)) {
        res.push_back(x);
    }
    return res;
}

template<typename T>
void subset_sum_transform(std::vector<T> &f) {
    int size = static_cast<int>(f.size());
    assert(size > 0 && (size & (size - 1)) == 0);
    int n = __builtin_ctz(static_cast<unsigned>(size)); // size = 2^n.
    for (int b = 0; b < n; b++) {
        for (int m = 0; m < size; m++) {
            if (m & (1u << b)) {
                f[m] += f[m ^ (1u << b)];
            }
        }
    }
}
```

Example Usage

```
using namespace std;

int main() {
    assert((submasks(0b101) == vector<Mask>{0b101, 0b100, 0b001, 0b000}));

    assert(next_popcount(0b0110) == 0b1001u);
    assert(
        (masks_with_popcount(4, 2) == vector<Mask>{0b0011, 0b0101, 0b0110, 0b1001, 0b1010, 0b1100})
    );

    // f[m] = 1 for every mask; after the transform f[m] counts submasks of m = 2^popcount(m).
    vector<int> f(1 << 3, 1);
    subset_sum_transform(f);
    assert(f[0b000] == 1);
    assert(f[0b101] == 4);
    assert(f[0b111] == 8);
    return 0;
}
```

1.5.3 Bitmask Dynamic Programming

Bitmask dynamic programming uses a bitmask to represent a small set or frontier, then treats that mask as the DP state. This is useful when one dimension is small enough for 2^n states: assignment problems can store which jobs are already used, set cover can store which elements are covered, submask DP can split a set into groups, and grid tiling can store which cells on the current frontier are already occupied.

For assignment, `dp[mask]` is the minimum cost after assigning the first `popcount(mask)` workers to the selected jobs. The next worker tries every unused job. For set cover, `dp[mask]` is the minimum number of chosen sets needed to cover exactly the elements in `mask`; adding one set moves to the union mask. For set partitioning, iterating all submasks gives the recurrence `dp[mask] = min(dp[mask ^ sub] + cost[sub])`. Over all masks, this has $O(3^n)$ transitions. For domino tiling, the mask is a profile of the current row; this technique is often called DP on a broken profile or plug DP. A domino covers exactly two edge-adjacent cells, so each piece is either horizontal (1×2) or vertical (2×1).

- `assignment_min_cost(cost)` returns a pair `(sum, job)`, where `sum` is the minimum cost of assigning each worker to a distinct job, and `job[i]` is the job assigned to worker `i`. The input is a square matrix with one row per worker and one column per job.
- `minimum_set_cover(sets, universe_size)` returns a pair `(count, chosen)`, where `count` is the minimum number of sets needed to cover all elements in $[0, u)$, where `u` is `universe_size`, and `chosen` contains one optimal list of set indices. If no cover exists, `count` is `-1` and `chosen` is empty. Each input set is represented as a bitmask.
- `partition_min_cost(group_cost)` returns the minimum total cost to partition all elements into disjoint nonempty groups, where `group_cost[mask]` is the cost of taking `mask` as one group.
- `count_domino_tilings(rows, cols)` returns the number of ways to tile a `rows` by `cols` rectangle with 1×2 and 2×1 dominoes.

All accumulated costs and tiling counts must fit in `int64_t`.

Time Complexity:

- $O(n^2 \cdot 2^n)$ per call to `assignment_min_cost()`, where n is the number of workers and jobs.
- $O(s \cdot 2^u)$ per call to `minimum_set_cover()`, where s is the number of sets and u is `universe_size`.
- $O(3^n)$ per call to `partition_min_cost()`, where n is the number of elements.
- $O(r \cdot c \cdot 2^c)$ per call to `count_domino_tilings()`, where r and c are the number of rows and columns, respectively.

Space Complexity:

- $O(2^n)$ auxiliary and $O(n)$ for the returned assignment from `assignment_min_cost()`.
- $O(2^u)$ auxiliary and $O(s)$ for the returned indices from `minimum_set_cover()`.
- $O(2^n)$ auxiliary for `partition_min_cost()`.
- $O(2^c)$ auxiliary for `count_domino_tilings()`.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <utility>
#include <vector>

std::pair<int64_t, std::vector<int>> assignment_min_cost(
    const std::vector<std::vector<int64_t>> &cost
) {
    int n = static_cast<int>(cost.size());
    assert(n < 31);
    int states = 1 << n;
    std::vector<int64_t> dp(states);
    std::vector<int> parent(states, -1);
    for (int mask = 0; mask < states; mask++) {
        int worker = __builtin_popcount(static_cast<unsigned>(mask));
        if (worker == n) {
            continue;
        }
        for (int job = 0; job < n; job++) {
            if ((mask & (1 << job)) == 0) {
                int next = mask | (1 << job);
                int64_t candidate = dp[mask] + cost[worker][job]; // Overflow warning.
                if (parent[next] == -1 || dp[next] > candidate) {
                    dp[next] = candidate;
                    parent[next] = job;
                }
            }
        }
    }
    std::vector<int> job(n);
    for (int mask = states - 1, worker = n - 1; worker >= 0; worker--) {
        job[worker] = parent[mask];
        mask ^= 1 << job[worker];
    }
    return {dp[states - 1], job};
}

std::pair<int, std::vector<int>> minimum_set_cover(
    const std::vector<int> &sets, int universe_size
) {
    assert(0 <= universe_size && universe_size < 31);
    int n = static_cast<int>(sets.size());
    int states = 1 << universe_size;
    int full = states - 1;
    for (int subset : sets) {
        assert((subset & ~full) == 0);
    }
    int inf = n + 1;
    std::vector<int> dp(states, inf), parent(states, -1), prev(states, -1);
    dp[0] = 0;
    for (int mask = 0; mask < states; mask++) {
        if (dp[mask] == inf) {
            continue;
        }
        for (int i = 0; i < n; i++) {
            int next = mask | sets[i];
            if (dp[next] > dp[mask] + 1) {
                dp[next] = dp[mask] + 1;
                parent[next] = i;
            }
        }
    }

```

```

        prev[next] = mask;
    }
}
}
if (dp[full] == inf) {
    return {-1, {}};
}
std::vector<int> chosen;
for (int mask = full; mask != 0; mask = prev[mask]) {
    int i = parent[mask];
    chosen.push_back(i);
}
std::reverse(chosen.begin(), chosen.end());
return {dp[full], chosen};
}

int64_t partition_min_cost(const std::vector<int64_t> &group_cost) {
    int states = static_cast<int>(group_cost.size());
    assert(states > 0 && (states & (states - 1)) == 0);
    std::vector<int64_t> dp(states);
    for (int mask = 1; mask < states; mask++) {
        dp[mask] = group_cost[mask];
        for (int sub = (mask - 1) & mask; sub > 0; sub = (sub - 1) & mask) {
            int64_t candidate = dp[mask ^ sub] + group_cost[sub]; // Overflow warning.
            dp[mask] = std::min(dp[mask], candidate);
        }
    }
    return dp[states - 1];
}

int64_t count_domino_tilings(int rows, int cols) {
    assert(rows >= 0 && cols >= 0);
    if (rows < cols) {
        return count_domino_tilings(cols, rows);
    }
    assert(cols < 31);
    int states = 1 << cols;
    std::vector<int64_t> dp(states, 0), next(states, 0);
    dp[0] = 1;
    for (int r = 0; r < rows; r++) {
        for (int c = 0; c < cols; c++) {
            next.assign(states, 0);
            for (int mask = 0; mask < states; mask++) {
                if (dp[mask] == 0) {
                    continue;
                }
                if (mask & (1 << c)) {
                    next[mask ^ (1 << c)] += dp[mask]; // Overflow warning.
                } else {
                    if (c + 1 < cols && (mask & (1 << (c + 1))) == 0) {
                        next[mask | (1 << (c + 1))] += dp[mask];
                    }
                    if (r + 1 < rows) {
                        next[mask | (1 << c)] += dp[mask];
                    }
                }
            }
        }
        dp.swap(next);
    }
    return dp[0];
}

```

Example Usage

```
#include <cassert>

int main() {
    std::vector<std::vector<int64_t>> cost{{9, 2, 7}, {6, 4, 3}, {5, 8, 1}};
    auto [assignment_cost, job] = assignment_min_cost(cost);
    assert(assignment_cost == 9);
    assert((job == std::vector<int>{1, 0, 2}));

    std::vector<int> sets{0b0011, 0b0110, 0b1100, 0b1000};
    auto [cover_count, chosen] = minimum_set_cover(sets, 4);
    assert(cover_count == 2);
    assert((chosen == std::vector<int>{0, 2}));

    auto [impossible_count, impossible_chosen] = minimum_set_cover(std::vector<int>{0b001, 0b010}, 3);
    assert(impossible_count == -1 && impossible_chosen.empty());

    std::vector<int64_t> group_cost{
        0,          // Empty group is unused.
        4, 6, 7,    // Singletons.
        6, 2, 3,    // Pairs.
        9,          // All three together.
    };
    assert(partition_min_cost(group_cost) == 7); // Groups {0} and {1, 2}.

    assert(count_domino_tilings(0, 5) == 1);
    assert(count_domino_tilings(2, 3) == 3);
    assert(count_domino_tilings(3, 3) == 0);
    assert(count_domino_tilings(4, 4) == 36);
    return 0;
}
```

1.5.4 Gray Code

The binary reflected Gray code is an ordering of all 2^n bitmasks in which consecutive masks differ in exactly one bit. It is useful whenever the cost of moving between adjacent states depends on a single bit flip: enumerating subsets while incrementally maintaining a value, hardware where only one bit may toggle at a time, and constructions like Hamiltonian paths on the hypercube.

The k -th Gray code is obtained from the ordinary binary value k by $k \oplus (k \gg 1)$. Incrementing k flips one binary bit and every lower bit, but XORing each bit with its higher neighbor cancels all but the highest of those changes, so consecutive Gray codes differ in exactly one bit. The map is a bijection: each binary bit is the prefix XOR of the Gray bits at and above it. The inverse loop computes those prefix XORs in doubling steps to recover the original rank.

- `gray_code(k)` returns the k -th mask in Gray code order.
- `inverse_gray_code(g)` returns the rank k such that `gray_code(k)` equals `g`.
- `gray_sequence(n)` returns all 2^n masks in Gray code order; consecutive entries (and the last with the first when `n > 0`) differ in exactly one bit, where $0 \leq n < \text{MASK_BITS}$.

Time Complexity:

- $O(1)$ per call to `gray_code()`.
- $O(\log b)$ per call to `inverse_gray_code()`, where b is `MASK_BITS`.
- $O(2^n)$ per call to `gray_sequence(n)`.

Space Complexity:

- $O(1)$ auxiliary for `gray_code()` and `inverse_gray_code()`.
- $O(1)$ auxiliary and $O(2^n)$ for the returned sequence from `gray_sequence(n)`.


```

#include <cassert>
#include <climits>
#include <cstdint>
#include <vector>

using Mask = uint32_t;
const int MASK_BITS = sizeof(Mask) * CHAR_BIT;

Mask gray_code(Mask k) {
    return k ^ (k >> 1);
}

Mask inverse_gray_code(Mask g) {
    for (int shift = 1; shift < MASK_BITS && (g >> shift) != 0; shift <= 1) {
        g ^= g >> shift;
    }
    return g;
}

std::vector<Mask> gray_sequence(int n) {
    assert(0 <= n && n < MASK_BITS);
    std::vector<Mask> res(Mask{1} << n);
    for (Mask k = 0; k < static_cast<Mask>(res.size()); k++) {
        res[k] = gray_code(k);
    }
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    assert(gray_code(0) == 0b000u);
    assert(gray_code(1) == 0b001u);
    assert(gray_code(2) == 0b011u);
    assert(gray_code(3) == 0b010u);

    for (Mask k = 0; k < 256; k++) {
        assert(inverse_gray_code(gray_code(k)) == k);
    }
    assert(inverse_gray_code(gray_code(0xdeadbeefu)) == 0xdeadbeefu);

    vector<Mask> seq = gray_sequence(3);
    assert(seq.size() == 8);
    for (int i = 0; i < static_cast<int>(seq.size()); i++) {
        Mask diff = seq[i] ^ seq[(i + 1) % seq.size()];
        assert((diff & (diff - 1)) == 0); // Cyclically, one bit changes per step.
    }
    return 0;
}

```

1.5.5 XOR Basis

An XOR basis (also called a linear basis over the field $\text{GF}(2)$) is the analog of a vector space basis for integers under the bitwise XOR operation. Each integer is treated as a vector of bits, where XOR plays the role of vector addition. Inserting a set of integers one at a time and keeping only those that are linearly independent yields a basis whose XOR-combinations (subset XORs) span exactly the same set of values as the original integers. From a basis of size r , exactly 2^r distinct values are reachable, and queries such as the maximum attainable XOR become a simple greedy scan.

The basis is kept in row-echelon form: the i -th basis element is either 0 or a value whose highest set bit is bit i , and no two basis elements share the same highest bit. Reducing a value means XORing it against the basis elements from the highest bit down, eliminating each set bit that a basis element covers.

- `XorBasis()` constructs an empty basis over `Mask`, which is `uint64_t` by default.
- `insert(x)` adds `x` to the basis, returning `true` if `x` was linearly independent of the current basis (increasing its size), or `false` if `x` was already representable as a subset XOR.
- `contains(x)` returns whether `x` is representable as the XOR of some subset of the inserted values, i.e. whether it reduces to 0 against the basis.
- `max_xor(base = 0)` returns the maximum value of `base` XORed with some subset XOR of the basis. With the default `base` of 0, this is the maximum subset XOR.
- `size()` returns the number of elements in the basis (its rank).

Time Complexity:

- $O(b)$ per call to `insert()`, `contains()`, and `max_xor()`, where b is the number of bits.
- $O(1)$ per call to `size()`.

Space Complexity:

- $O(b)$ for storage of the basis.
- $O(1)$ auxiliary for all operations.

```
#include <array>
#include <climits>
#include <cstdint>

using Mask = uint64_t;
const int MASK_BITS = sizeof(Mask) * CHAR_BIT;

class XorBasis {
    std::array<Mask, MASK_BITS> basis{}; // basis[i] != 0 has its highest set bit at bit i.
    int sz = 0;

public:
    bool insert(Mask x) {
        for (int i = MASK_BITS - 1; i >= 0; i--) {
            if (!((x >> i) & 1)) {
                continue;
            }
            if (basis[i] == 0) {
                basis[i] = x;
                sz++;
                return true;
            }
            x ^= basis[i];
        }
        return false;
    }

    bool contains(Mask x) const {
        for (int i = MASK_BITS - 1; i >= 0; i--) {
            if (((x >> i) & 1) && basis[i] != 0) {
                x ^= basis[i];
            }
        }
        return x == 0;
    }

    Mask max_xor(Mask base = 0) const {
        for (int i = MASK_BITS - 1; i >= 0; i--) {
            if ((base ^ basis[i]) > base) {
                base ^= basis[i];
            }
        }
        return base;
    }
};
```

```

    }

    int size() const { return sz; }
};

```

Example Usage

```

#include <cassert>

int main() {
    XorBasis b;
    assert(b.insert(1)); // 001
    assert(b.insert(2)); // 010
    assert(!b.insert(3)); // 011 = 1 ^ 2, already representable.
    assert(b.size() == 2);

    assert(b.contains(3));
    assert(b.contains(0));
    assert(!b.contains(4));

    assert(b.max_xor() == 3); // Best subset XOR over {0, 1, 2, 3}.
    assert(b.max_xor(4) == 7); // 4 ^ 3.
    return 0;
}

```

1.5.6 Binary Trie

Maintain a multiset of nonnegative b -bit integers as a binary trie, storing each value as a root-to-leaf path of its bits from the most significant (bit $b - 1$) down to the least significant. Every node keeps a count of the stored values passing through it, which supports deletion and counting queries, and lets the trie answer XOR-extremal queries by a greedy walk down the bits.

The classic application is the maximum-XOR query: to maximize $\mathbf{x} \oplus \mathbf{y}$ over all stored $\mathbf{y}$, walk from the most significant bit. At each step, descend toward the child whose bit differs from that of $\mathbf{x}$ whenever such a branch exists, since setting a higher bit of the result always dominates any combination of lower bits. The minimum-XOR query is the mirror image, preferring the matching bit.

This is a multiset (a value may be inserted more than once), and is distinct from the XOR basis: the basis spans subset XORs of a fixed set, whereas this answers XOR queries against the stored elements themselves with support for insertion and deletion.

Values must lie in $[0, 2^b)$ for bit width b ; the template parameters select the unsigned word type `U` and the bit width `BITS`, which must be less than the width of `U` (for example, 30 with `uint32_t` or 62 with `uint64_t`). Nodes are kept in a pool indexed by integer, where index 0 is the root and also serves as the null child.

- `BinaryTrie<U, BITS>()` constructs an empty multiset.
- `size()` returns the number of stored values, and `empty()` returns whether there are none.
- `insert(x)` adds one occurrence of $\mathbf{x}$.
- `erase(x)` removes one occurrence of $\mathbf{x}$, returning whether one was present.
- `count(x)` returns the number of occurrences of $\mathbf{x}$.
- `max_xor(x)` and `min_xor(x)` return the maximum and minimum of $\mathbf{x} \oplus \mathbf{y}$ over all stored $\mathbf{y}$. The trie must be non-empty.
- `count_xor_less(x, k)` returns the number of stored $\mathbf{y}$ (with multiplicity) such that $\mathbf{x} \oplus \mathbf{y}$ is strictly less than $\mathbf{k}$.

Time Complexity:

- $O(b)$ per call to `insert()`, `erase()`, `count()`, `max_xor()`, `min_xor()`, and `count_xor_less()`, where b is the bit width `BITS`.

- $O(1)$ per call to `size()` and `empty()`.

Space Complexity:

- $O(n \cdot b)$ for storage, where n is the number of distinct values inserted so far.
- $O(1)$ auxiliary per operation.

```
#include <array>
#include <cassert>
#include <climits>
#include <cstdint>
#include <vector>

template<typename U = uint32_t, int BITS = 30>
class BinaryTrie {
    static_assert(0 < BITS && BITS < static_cast<int>(sizeof(U) * CHAR_BIT));

    std::vector<std::array<int, 2>> child; // child[node][bit], where 0 means no child.
    std::vector<int> cnt;                // Number of stored values passing through each node.

    static void assert_fits(U x) { assert((x >> BITS) == 0); }

    int new_node() {
        child.push_back({0, 0});
        cnt.push_back(0);
        return static_cast<int>(cnt.size()) - 1;
    }

public:
    BinaryTrie() { new_node(); } // Node 0 is the root (and doubles as the null child).

    int size() const { return cnt[0]; }
    bool empty() const { return cnt[0] == 0; }

    void insert(U x) {
        assert_fits(x);
        int node = 0;
        cnt[node]++;
        for (int b = BITS - 1; b >= 0; b--) {
            int bit = (x >> b) & 1;
            if (child[node][bit] == 0) {
                int next = new_node();
                child[node][bit] = next;
            }
            node = child[node][bit];
            cnt[node]++;
        }
    }

    int count(U x) const {
        assert_fits(x);
        int node = 0;
        for (int b = BITS - 1; b >= 0; b--) {
            int next = child[node][(x >> b) & 1];
            if (next == 0) {
                return 0;
            }
            node = next;
        }
        return cnt[node];
    }

    bool erase(U x) {
        if (count(x) == 0) {
            return false;
        }
        int node = 0;
        cnt[node]--;
```

```

    for (int b = BITS - 1; b >= 0; b--) {
        node = child[node][(x >> b) & 1];
        cnt[node]--;
    }
    return true;
}

U max_xor(U x) const {
    assert(!empty());
    assert_fits(x);
    int node = 0;
    U res = 0;
    for (int b = BITS - 1; b >= 0; b--) {
        int bit = (x >> b) & 1;
        int opp = child[node][bit ^ 1];
        if (opp != 0 && cnt[opp] > 0) {
            res |= U(1) << b;
            node = opp;
        } else {
            node = child[node][bit];
        }
    }
    return res;
}

U min_xor(U x) const {
    assert(!empty());
    assert_fits(x);
    int node = 0;
    U res = 0;
    for (int b = BITS - 1; b >= 0; b--) {
        int bit = (x >> b) & 1;
        int same = child[node][bit];
        if (same != 0 && cnt[same] > 0) {
            node = same;
        } else {
            res |= U(1) << b;
            node = child[node][bit ^ 1];
        }
    }
    return res;
}

int count_xor_less(U x, U k) const {
    assert_fits(x);
    if ((k >> BITS) != 0) {
        return cnt[0]; // Every XOR result is below 2^BITS, which is at most k.
    }
    int node = 0, res = 0;
    for (int b = BITS - 1; b >= 0; b--) {
        int xb = (x >> b) & 1, kb = (k >> b) & 1;
        if (kb == 1) {
            // Stored values matching x at bit b give XOR bit 0 here, hence are already below k.
            int below = child[node][xb];
            if (below != 0) {
                res += cnt[below];
            }
            node = child[node][xb ^ 1]; // Stay tied with k by taking XOR bit 1.
        } else {
            node = child[node][xb]; // XOR bit must be 0 to remain below k.
        }
        if (node == 0) {
            break;
        }
    }
    return res;
}
};

```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    BinaryTrie<> trie;
    assert(trie.empty());
    for (int x : {3, 10, 5, 25, 2}) {
        trie.insert(x);
    }
    assert(!trie.empty());
    assert(trie.size() == 5);

    // x XOR y over the stored y: {26, 19, 28, 0, 27} for x = 25.
    assert(trie.max_xor(25) == 28);           // 25 XOR 5.
    assert(trie.min_xor(25) == 0);           // 25 XOR 25.
    assert(trie.count_xor_less(25, 20) == 2); // XOR values 0 and 19 are below 20.
    assert(trie.count_xor_less(25, uint32_t(1) << 30) == 5); // Bound at/above 2^BITS counts all.

    assert(trie.count(10) == 1);
    trie.insert(10);
    assert(trie.count(10) == 2); // Multiset: two copies of 10.
    assert(trie.erase(10) && trie.count(10) == 1);
    assert(trie.erase(25) && trie.count(25) == 0);
    assert(!trie.erase(25));      // Already gone.
    assert(trie.max_xor(0) == 10); // Largest remaining value is 10.
    return 0;
}
```

1.5.7 Bitwise Convolution (FWHT)

Bitwise convolution combines two arrays indexed by masks, grouping pairs of masks by a bitwise operation instead of by ordinary addition. For arrays `a` and `b` over an n -bit universe, the XOR convolution `c[m]` is the sum of `a[x]*b[y]` over all $x \hat{=} y = m$. OR and AND convolution replace $x \hat{=} y$ with $x | y$ or $x \& y$. These appear in subset DP, counting pairs by bitwise result, and algebra over the Boolean hypercube.

The fast Walsh-Hadamard transform (FWHT) diagonalizes XOR convolution in the same way that the FFT diagonalizes polynomial convolution: transform both arrays, multiply pointwise, and transform back. The OR and AND versions are Zeta/Mobius transforms over the subset lattice, but they have the same butterfly-shaped implementation and the same $O(n \log n)$ runtime for arrays of length $n = 2^k$.

- `xor_transform(f, invert)` applies the XOR FWHT to `f` in place. The inverse transform uses `invert = true` and divides every coefficient by `f.size()`.
- `or_transform(f, invert)` applies the OR Zeta transform in place; the inverse is Mobius inversion.
- `and_transform(f, invert)` applies the AND Zeta transform in place; the inverse is Mobius inversion.
- `xor_convolve(a, b)` returns the XOR convolution of `a` and `b`.
- `or_convolve(a, b)` returns the OR convolution of `a` and `b`.
- `and_convolve(a, b)` returns the AND convolution of `a` and `b`.

The in-place transforms require a nonempty array whose length is a power of two. A convolution returns empty if either input is empty; otherwise, both inputs are padded with zeros to the smallest power of two at least their maximum length, which is also the output length. For exact integer XOR convolution, the inverse divisions are exact; for modular arithmetic, replace the division by multiplication with the modular inverse of the transform length. All intermediate sums, differences, and products must be representable in the value type.

Time Complexity: $O(n \log n)$ per transform or convolution, where n is the padded power-of-two length.

Space Complexity:

- $O(1)$ auxiliary for the in-place transforms.

- $O(n)$ auxiliary and $O(n)$ for the returned array from each convolution.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <vector>

template<typename T>
void xor_transform(std::vector<T> &f, bool invert) {
    int n = static_cast<int>(f.size());
    assert(n > 0 && (n & (n - 1)) == 0);
    for (int len = 1; len < n; len <= 1) {
        for (int i = 0; i < n; i += len <= 1) {
            for (int j = 0; j < len; j++) {
                T u = f[i + j], v = f[i + j + len];
                f[i + j] = u + v; // Overflow warning.
                f[i + j + len] = u - v;
            }
        }
    }
    if (invert) {
        for (T &x : f) {
            x /= n;
        }
    }
}

template<typename T>
void or_transform(std::vector<T> &f, bool invert) {
    int n = static_cast<int>(f.size());
    assert(n > 0 && (n & (n - 1)) == 0);
    for (int bit = 1; bit < n; bit <= 1) {
        for (int mask = 0; mask < n; mask++) {
            if (mask & bit) {
                if (invert) {
                    f[mask] -= f[mask ^ bit]; // Overflow warning.
                } else {
                    f[mask] += f[mask ^ bit];
                }
            }
        }
    }
}

template<typename T>
void and_transform(std::vector<T> &f, bool invert) {
    int n = static_cast<int>(f.size());
    assert(n > 0 && (n & (n - 1)) == 0);
    for (int bit = 1; bit < n; bit <= 1) {
        for (int mask = 0; mask < n; mask++) {
            if ((mask & bit) == 0) {
                if (invert) {
                    f[mask] -= f[mask ^ bit]; // Overflow warning.
                } else {
                    f[mask] += f[mask ^ bit];
                }
            }
        }
    }
}

template<typename T, typename Transform>
std::vector<T> bitwise_convolve(std::vector<T> a, std::vector<T> b, Transform transform) {
    if (a.empty() || b.empty()) {
        return {};
    }
    int needed = static_cast<int>(std::max(a.size(), b.size()));
    int n = 1;

```

```

while (n < needed) {
    n <= 1;
}
a.resize(n);
b.resize(n);
transform(a, false);
transform(b, false);
for (int i = 0; i < n; i++) {
    a[i] *= b[i]; // Overflow warning.
}
transform(a, true);
return a;
}

template<typename T>
std::vector<T> xor_convolve(std::vector<T> a, std::vector<T> b) {
    return bitwise_convolve(a, b, xor_transform<T>);
}

template<typename T>
std::vector<T> or_convolve(std::vector<T> a, std::vector<T> b) {
    return bitwise_convolve(a, b, or_transform<T>);
}

template<typename T>
std::vector<T> and_convolve(std::vector<T> a, std::vector<T> b) {
    return bitwise_convolve(a, b, and_transform<T>);
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int64_t> a{1, 2, 3, 4};
    vector<int64_t> b{5, 6, 7, 8};

    {
        vector<int64_t> f{1, 2, 3, 4}, original(f);
        xor_transform(f, false);
        xor_transform(f, true);
        assert(f == original);
    }
    {
        vector<int64_t> f{1, 2, 3, 4}, original(f);
        or_transform(f, false);
        or_transform(f, true);
        assert(f == original);
    }
    {
        vector<int64_t> f{1, 2, 3, 4}, original(f);
        and_transform(f, false);
        and_transform(f, true);
        assert(f == original);
    }

    // XOR: c[m] sums a[x] * b[y] over pairs with x ^ y == m.
    assert((xor_convolve(a, b) == vector<int64_t>{70, 68, 62, 60}));

    // OR: c[m] sums a[x] * b[y] over pairs with x | y == m.
    assert((or_convolve(a, b) == vector<int64_t>{5, 28, 43, 184}));

    // AND: c[m] sums a[x] * b[y] over pairs with x & y == m.
    assert((and_convolve(a, b) == vector<int64_t>{103, 52, 73, 32}));
}

```



```
vector<int64_t> one_mask{0, 1, 0, 0};  
vector<int64_t> two_mask{0, 0, 1};  
assert((xor_convolve(one_mask, two_mask) == vector<int64_t>{0, 0, 0, 1}));  
assert((or_convolve(one_mask, two_mask) == vector<int64_t>{0, 0, 0, 1}));  
assert((and_convolve(one_mask, two_mask) == vector<int64_t>{1, 0, 0, 0}));  
return 0;  
}
```

Chapter 2

Data Structures

2.1 Linked Lists

2.1.1 Singly Linked List

Implements common singly linked list operations using raw nodes. Linked lists are rarely the best contest choice compared with arrays, vectors, and `std::list`, but the pointer patterns are useful for interview-style problems and for understanding constant-time splicing.

For production C++, prefer containers such as `std::list`, `std::forward_list`, or `std::vector` when they fit the problem. Use the raw-node style below when the problem statement already gives nodes, or when manual pointer manipulation is the point of the exercise.

- `reverse_list(head)` reverses a singly linked list and returns the new head.
- `split_half(head, &second)` cuts a list into two halves, returning the first half and storing the second half in `second`.
- `merge_sorted_lists(a, b)` merges two sorted lists using a dummy node and returns the merged head.
- `splice_after(pos, before)` moves the node after `before` so it appears after `pos`.
- `splice_range_after(pos, before_first, last)` moves the half-open range (`before_first`, `last`) so it appears after `pos`.

The splicing helpers use the same “after” convention as `std::forward_list::splice_after()`. They work naturally with a dummy head node, which makes insertions and removals at the beginning of a list match all other positions. For `splice_range_after()`, `pos` must not lie inside the moved range. For doubly linked list splicing, use the next section.

For cycle detection, see the variant on iterated functions in chapter 1, which can be easily adapted for linked lists.

Time Complexity:

- $O(n)$ per call to `reverse_list()` and `split_half()`.
- $O(n + m)$ per call to `merge_sorted_lists()`.
- $O(1)$ per call to `splice_after()`.
- $O(k)$ per call to `splice_range_after()`, where k is the number of moved nodes. The actual pointer splicing is $O(1)$, but this compact version walks to the end of the moved range.

Space Complexity: $O(1)$ auxiliary for all operations.

```
struct ListNode {
    int value;
    ListNode *next;

    explicit ListNode(int value = 0) : value(value), next(nullptr) {}
};
```

```

};

ListNode *reverse_list(ListNode *head) {
    ListNode *prev = nullptr;
    while (head != nullptr) {
        ListNode *next = head->next;
        head->next = prev;
        prev = head;
        head = next;
    }
    return prev;
}

ListNode *split_half(ListNode *head, ListNode **second) {
    if (head == nullptr || head->next == nullptr) {
        *second = nullptr;
        return head;
    }
    ListNode *slow = head, *fast = head->next;
    while (fast != nullptr && fast->next != nullptr) {
        slow = slow->next;
        fast = fast->next->next;
    }
    *second = slow->next;
    slow->next = nullptr;
    return head;
}

ListNode *merge_sorted_lists(ListNode *a, ListNode *b) {
    ListNode dummy;
    ListNode *tail = &dummy;
    while (a != nullptr && b != nullptr) {
        if (b->value < a->value) {
            tail->next = b;
            b = b->next;
        } else {
            tail->next = a;
            a = a->next;
        }
        tail = tail->next;
    }
    tail->next = a != nullptr ? a : b;
    return dummy.next;
}

void splice_after(ListNode *pos, ListNode *before) {
    ListNode *node = before->next;
    if (node == nullptr || pos == before || pos == node) {
        return;
    }
    before->next = node->next;
    node->next = pos->next;
    pos->next = node;
}

void splice_range_after(ListNode *pos, ListNode *before_first, ListNode *last) {
    ListNode *first = before_first->next;
    if (first == last) {
        return;
    }
    ListNode *tail = first;
    while (tail->next != last) {
        tail = tail->next;
    }
    before_first->next = last;
    tail->next = pos->next;
    pos->next = first;
}

```

Example Usage

```
#include <cassert>
#include <vector>
using namespace std;

vector<int> values(ListNode *head) {
    vector<int> res;
    for (; head != nullptr; head = head->next) {
        res.push_back(head->value);
    }
    return res;
}

int main() {
    vector<ListNode> a{ListNode(1), ListNode(3), ListNode(5)};
    vector<ListNode> b{ListNode(2), ListNode(4), ListNode(6)};
    for (int i = 0; i + 1 < static_cast<int>(a.size()); i++) {
        a[i].next = &a[i + 1];
        b[i].next = &b[i + 1];
    }
    ListNode *merged = merge_sorted_lists(&a[0], &b[0]);
    vector<int> merged_values = values(merged);
    assert((merged_values == vector<int>{1, 2, 3, 4, 5, 6}));

    ListNode *second = nullptr;
    ListNode *first = split_half(merged, &second);
    assert(values(first).size() == 3);
    assert(values(second).size() == 3);
    assert(values(reverse_list(first))[0] == 3);

    ListNode dummy(0), w(1), x(2), y(3), z(4);
    dummy.next = &w;
    w.next = &x;
    x.next = &y;
    y.next = &z;
    splice_after(&dummy, &y); // Move 4 to the front: 4, 1, 2, 3.
    assert((values(dummy.next) == vector<int>{4, 1, 2, 3}));
    splice_range_after(&z, &w, nullptr); // Move 2, 3 after 4: 4, 2, 3, 1.
    assert((values(dummy.next) == vector<int>{4, 2, 3, 1}));
    return 0;
}
```

2.1.2 Doubly Linked List

Implements common doubly linked list operations using raw intrusive nodes. A node is intrusive when the `prev` and `next` pointers live inside the stored object itself, rather than inside a separate container allocation. This is useful when a problem already gives node pointers, when values must not be copied, or when a map needs to jump directly to a list node in $O(1)$ time, as in an LRU cache.

The functions below use a circular sentinel node. When the list is empty, `sentinel.next` and `sentinel.prev` both point back to `sentinel`. This removes special cases for inserting or erasing at the front and back of the list.

- `init_list(sentinel)` initializes an empty circular list around `sentinel`.
- `insert_after(pos, node)` inserts detached `node` immediately after `pos`.
- `insert_before(pos, node)` inserts detached `node` immediately before `pos`.
- `erase(node)` removes `node` from its current list and leaves it detached.
- `push_front(sentinel, node)` inserts detached `node` at the front of the list.
- `push_back(sentinel, node)` inserts detached `node` at the back of the list.
- `move_to_front(sentinel, node)` moves an already-linked `node` to the front of the list.
- `empty(sentinel)` returns whether the list has no data nodes.
- `splice(pos, node)` moves an already-linked `node` so it appears immediately before `pos`.

- `splice_range(pos, first, last)` moves the half-open range `[first, last)` so it appears immediately before `pos`.

The data nodes passed to insertion functions must be detached, meaning their `prev` and `next` pointers are both null. The node passed to `erase()` or `move_to_front()` must already be linked into a list and must not be the sentinel. For `splice_range(pos, first, last)`, the `pos` node must not lie inside the moved range.

Time Complexity: $O(1)$ per call to all operations.

Space Complexity: $O(1)$ auxiliary for all operations.

```
struct DListNode {
    int value;
    DListNode *prev, *next;

    explicit DListNode(int value = 0) : value(value), prev(nullptr), next(nullptr) {}
};

void init_list(DListNode *sentinel) {
    sentinel->prev = sentinel;
    sentinel->next = sentinel;
}

bool empty(DListNode *sentinel) {
    return sentinel->next == sentinel;
}

void insert_after(DListNode *pos, DListNode *node) {
    node->prev = pos;
    node->next = pos->next;
    pos->next->prev = node;
    pos->next = node;
}

void insert_before(DListNode *pos, DListNode *node) {
    insert_after(pos->prev, node);
}

DListNode *erase(DListNode *node) {
    node->prev->next = node->next;
    node->next->prev = node->prev;
    node->prev = nullptr;
    node->next = nullptr;
    return node;
}

void push_front(DListNode *sentinel, DListNode *node) {
    insert_after(sentinel, node);
}

void push_back(DListNode *sentinel, DListNode *node) {
    insert_before(sentinel, node);
}

void move_to_front(DListNode *sentinel, DListNode *node) {
    push_front(sentinel, erase(node));
}

void splice(DListNode *pos, DListNode *node) {
    if (pos == node || pos == node->next) {
        return;
    }
    insert_before(pos, erase(node));
}

void splice_range(DListNode *pos, DListNode *first, DListNode *last) {
    if (first == last || pos == first) {
        return;
    }
}
```

```

DListNode *before_first = first->prev;
DListNode *tail = last->prev;
// Detach [first, last) from its current list.
before_first->next = last;
last->prev = before_first;
// Insert the detached range immediately before pos.
DListNode *before_pos = pos->prev;
before_pos->next = first;
first->prev = before_pos;
tail->next = pos;
pos->prev = tail;
}

```

Example Usage

```

#include <cassert>
#include <vector>
using namespace std;

// Returns the data values from front to back (walks the circular list until it loops to sentinel).
vector<int> values(DListNode *sentinel) {
    vector<int> res;
    for (DListNode *p = sentinel->next; p != sentinel; p = p->next) {
        res.push_back(p->value);
    }
    return res;
}

int main() {
    // dummy is the sentinel; a, b, c, d start detached. push_back/push_front link at back/front.
    DListNode dummy, a(1), b(2), c(3), d(4);
    init_list(&dummy);
    assert(empty(&dummy));

    push_back(&dummy, &a);
    push_back(&dummy, &b);
    push_back(&dummy, &c);
    push_front(&dummy, &d);
    assert((values(&dummy) == vector<int>{4, 1, 2, 3}));

    move_to_front(&dummy, &b); // list: 2 4 1 3
    assert((values(&dummy) == vector<int>{2, 4, 1, 3}));

    erase(&d); // detach d; list: 2 1 3
    assert((values(&dummy) == vector<int>{2, 1, 3}));

    insert_before(&dummy, &d); // before the sentinel == the back; list: 2 1 3 4
    assert((values(&dummy) == vector<int>{2, 1, 3, 4}));

    splice(&d, &a); // move a to just before d; list: 2 3 1 4
    assert((values(&dummy) == vector<int>{2, 3, 1, 4}));

    splice_range(&dummy, &a, &dummy); // range {a, d} already at the back: no-op
    assert((values(&dummy) == vector<int>{2, 3, 1, 4}));

    splice_range(&b, &a, &dummy); // move {a, d} to before b; list: 1 4 2 3
    assert((values(&dummy) == vector<int>{1, 4, 2, 3}));
    return 0;
}

```

2.1.3 LRU Cache

Maintains a least-recently-used cache with fixed capacity. The cache supports lookups and updates by key, evicting the least recently used entry whenever an insertion would exceed capacity.

This implementation combines a doubly linked list (`std::list`) with a map from keys to list iterators. The front of the list stores the most recently used item. The operation `items.splice(items.begin(), items, it)` is the standard-library way to move an existing list node to the front in $O(1)$ time. The key type `K` must support `operator==`, and `std::hash<K>` must be defined.

To implement the same cache without `std::list`, use the intrusive doubly linked list operations from section 2.1.2. Store a `Node*` in the map instead of a list iterator, move hits to the front with `move_to_front(&sentinel, node)`, and evict `sentinel.prev`.

- `LRUCache<K, V>(capacity)` constructs an empty cache holding at most `capacity` keys.
- `get(key, &value)` returns whether `key` is present. If present, it stores the value in `value` and marks the key as most recently used.
- `put(key, value)` inserts or updates `key`, marking it as most recently used and evicting the least recently used key if needed.
- `size()` returns the number of keys currently stored.

Time Complexity: $O(1)$ expected amortized per call to `get()` and `put()`, and $O(n)$ in the collision-heavy worst case for the hash table.

Space Complexity: $O(n)$, where n is the cache capacity.

```
#include <cassert>
#include <list>
#include <unordered_map>
#include <utility>

// Replacing std::list with a manual doubly-linked list would look something like:
// struct DListNode { K key; V value; DListNode *prev, *next; };
// std::unordered_map<K, DListNode*> where;
//
// On get:
// move_to_front(&sentinel, where[key]);
//
// On eviction:
// DListNode *old = sentinel.prev;
// erase(old);
// where.erase(old->key).

template<typename K, typename V>
class LRUCache {
    int cap;
    std::list<std::pair<K, V>> items;
    std::unordered_map<K, decltype(items.begin())> where;

public:
    explicit LRUCache(int capacity) : cap(capacity) { assert(capacity >= 0); }

    int size() const { return static_cast<int>(items.size()); }

    bool get(const K &key, V *value) {
        assert(value != nullptr);
        auto it = where.find(key);
        if (it == where.end()) {
            return false;
        }
        items.splice(items.begin(), items, it->second);
        it->second = items.begin();
        *value = it->second->second;
        return true;
    }
}
```

```

void put(const K &key, const V &value) {
    auto it = where.find(key);
    if (it != where.end()) {
        it->second->second = value;
        items.splice(items.begin(), items, it->second);
        it->second = items.begin();
        return;
    }
    if (cap == 0) {
        return;
    }
    if (static_cast<int>(items.size()) == cap) {
        where.erase(items.back().first);
        items.pop_back();
    }
    items.push_front({key, value});
    where[key] = items.begin();
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    LRUCache<int, int> cache(2);
    int value = 0;
    cache.put(1, 10);
    cache.put(2, 20);
    assert(cache.get(1, &value) && value == 10);
    cache.put(3, 30); // Evicts key 2.
    assert(!cache.get(2, &value));
    assert(cache.get(3, &value) && value == 30);
    cache.put(1, 11);
    assert(cache.get(1, &value) && value == 11);

    LRUCache<int, int> disabled(0);
    disabled.put(1, 10);
    assert(disabled.size() == 0);
    return 0;
}

```

2.2 Heaps and Priority Queues

2.2.1 Binary Heap

Maintains a priority queue, returning the minimum element by default. The comparator `comp` defines priority: `comp(a, b)` is true when `a` has higher priority than `b`, so `std::greater<T>` instead creates a max-priority queue. A binary heap stores a heap-ordered binary tree implicitly in an array, where the children of the node at index i live at indices $2i + 1$ and $2i + 2$. Insertion appends the new element and bubbles it up toward the root, while extraction moves the last element to the root and sifts it down. To customize the ordering, instantiate `BinaryHeap<T, Compare>` and pass the comparator to either constructor.

- `BinaryHeap<T>()` constructs an empty priority queue.
- `BinaryHeap<T>(lo, hi)` constructs a priority queue from the elements in the half-open iterator range `[lo, hi)`.
- `size()` returns the size of the priority queue.
- `empty()` returns whether the priority queue is empty.

- `push(v)` inserts the value `v` into the priority queue.
- `pop()` removes the highest-priority element from the priority queue.
- `top()` returns the highest-priority element in the priority queue.

Time Complexity:

- $O(1)$ per call to the first constructor, `size()`, `empty()`, and `top()`.
- $O(\log n)$ per call to `push()` and `pop()`, where n is the number of elements in the priority queue.
- $O(n)$ per call to the second constructor on the distance between `lo` and `hi` (bottom-up heapify).

Space Complexity:

- $O(n)$ for storage of the priority queue elements.
- $O(1)$ auxiliary for all operations.

```
#include <cassert>
#include <functional>
#include <utility>
#include <vector>

template<typename T, typename Compare = std::less<T>>
class BinaryHeap {
    std::vector<T> heap;
    Compare comp;

    void sift_down(int i) {
        while (true) {
            int child = 2 * i + 1;
            if (child >= static_cast<int>(heap.size())) {
                break;
            }
            if (child + 1 < static_cast<int>(heap.size()) && comp(heap[child + 1], heap[child])) {
                child++;
            }
            if (comp(heap[child], heap[i])) {
                std::swap(heap[i], heap[child]);
                i = child;
            } else {
                break;
            }
        }
    }

public:
    explicit BinaryHeap(Compare comp = Compare()) : comp(std::move(comp)) {}

    template<typename It>
    BinaryHeap(It lo, It hi, Compare comp = Compare()) : heap(lo, hi), comp(std::move(comp)) {
        for (int i = static_cast<int>(heap.size()) / 2 - 1; i >= 0; i--) {
            sift_down(i);
        }
    }

    int size() const { return static_cast<int>(heap.size()); }
    bool empty() const { return heap.empty(); }

    void push(const T &v) {
        heap.push_back(v);
        int i = static_cast<int>(heap.size()) - 1;
        while (i > 0) {
            int parent = (i - 1) / 2;
            if (!comp(heap[i], heap[parent])) {
                break;
            }
            std::swap(heap[i], heap[parent]);
            i = parent;
        }
    }
}
```

```

}

void pop() {
    assert(!heap.empty());
    heap[0] = heap.back();
    heap.pop_back();
    if (!heap.empty()) {
        sift_down(0);
    }
}

const T &top() const {
    assert(!heap.empty());
    return heap[0];
}
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{0, 5, -1, 12};
    BinaryHeap<int> h(a.begin(), a.end());
    assert(h.size() == 4);
    assert(h.top() == -1);
    h.push(10);
    vector<int> popped;
    while (!h.empty()) {
        popped.push_back(h.top());
        h.pop();
    }
    assert((popped == vector<int>{-1, 0, 5, 10, 12}));
    assert(h.empty());

    BinaryHeap<int, greater<int>> max_heap;
    max_heap.push(1);
    max_heap.push(3);
    assert(max_heap.top() == 3);
    return 0;
}

```

2.2.2 Leftist Heap

Maintains a mergeable priority queue, returning the minimum element by default. The comparator `comp` defines priority: `comp(a, b)` is true when `a` has higher priority than `b`, so `std::greater<T>` instead creates a max-priority queue. A leftist heap stores at each node a rank (the length of the shortest path to a missing descendant) and keeps every left child's rank at least that of its right child, so the right spine has length $O(\log n)$. Merging walks only this right spine, swapping children where needed to restore the invariant. Every other operation reduces to this merge: insertion merges a one-node heap, and extraction merges the root's two subtrees. Unlike the self-adjusting skew heap, the explicit rank bound makes every operation $O(\log n)$ in the worst case rather than amortized. To customize the ordering, instantiate `LeftistHeap<T, Compare>` and pass the comparator to either constructor.

- `LeftistHeap<T>()` constructs an empty priority queue.
- `LeftistHeap<T>(lo, hi)` constructs a priority queue from the elements in the half-open iterator range `[lo, hi)`.
- `size()` returns the size of the priority queue.
- `empty()` returns whether the priority queue is empty.
- `push(v)` inserts the value `v` into the priority queue.
- `pop()` removes the highest-priority element from the priority queue.

- `top()` returns the highest-priority element in the priority queue.
- `absorb(h)` inserts every value from the distinct heap `h` and sets `h` to the empty priority queue. Both heaps must use equivalent comparators.

Time Complexity:

- $O(1)$ per call to the first constructor, `size()`, `empty()`, and `top()`.
- $O(\log n)$ per call to `push()`, `pop()`, and `absorb()`, where n is the number of elements in the priority queue.
- $O(n \log n)$ per call to the second constructor, where n is the distance between `lo` and `hi`.

Space Complexity:

- $O(n)$ for storage of the priority queue elements.
- $O(\log n)$ auxiliary stack space for `push()`, `pop()`, and `absorb()`.
- $O(n)$ auxiliary stack space in the worst case for destruction.
- $O(1)$ auxiliary for all other operations.

```
#include <cassert>
#include <functional>
#include <utility>

template<typename T, typename Compare = std::less<T>>
class LeftistHeap {
    struct Node {
        T value;
        int rank;
        Node *left, *right;

        explicit Node(const T &v) : value(v), rank(1), left(nullptr), right(nullptr) {}
    } *root;

    int num_nodes;
    Compare comp;

    static int rank(Node *n) { return (n == nullptr) ? 0 : n->rank; }

    Node *merge(Node *a, Node *b) {
        if (a == nullptr) {
            return b;
        }
        if (b == nullptr) {
            return a;
        }
        if (comp(b->value, a->value)) {
            std::swap(a, b);
        }
        a->right = merge(a->right, b);
        if (rank(a->left) < rank(a->right)) {
            std::swap(a->left, a->right);
        }
        a->rank = rank(a->right) + 1;
        return a;
    }

    static void clean_up(Node *n) {
        if (n != nullptr) {
            clean_up(n->left);
            clean_up(n->right);
            delete n;
        }
    }

public:
    explicit LeftistHeap(Compare comp = Compare())
        : root(nullptr), num_nodes(0), comp(std::move(comp)) {}

    template<typename It>
```

```

LeftistHeap(It lo, It hi, Compare comp = Compare())
: root(nullptr), num_nodes(0), comp(std::move(comp)) {
    while (lo != hi) {
        push(*(lo++));
    }
}

~LeftistHeap() { clean_up(root); }
LeftistHeap(const LeftistHeap &) = delete;
LeftistHeap &operator=(const LeftistHeap &) = delete;
int size() const { return num_nodes; }
bool empty() const { return root == nullptr; }

void push(const T &v) {
    root = merge(root, new Node(v));
    num_nodes++;
}

void pop() {
    assert(!empty());
    Node *tmp = root;
    root = merge(root->left, root->right);
    delete tmp;
    num_nodes--;
}

const T &top() const {
    assert(!empty());
    return root->value;
}

void absorb(LeftistHeap &h) {
    assert(this != &h);
    root = merge(root, h.root);
    num_nodes += h.num_nodes;
    h.root = nullptr;
    h.num_nodes = 0;
}
};

```

Example Usage

```

#include <vector>
using namespace std;

int main() {
    LeftistHeap<int> h, h2;
    assert(h.empty());
    h.push(12);
    h.push(10);
    h2.push(5);
    h2.push(-1);
    h2.push(0);
    h.absorb(h2);
    assert(h.size() == 5);
    assert(h2.empty());
    vector<int> popped;
    while (!h.empty()) {
        popped.push_back(h.top());
        h.pop();
    }
    assert((popped == vector<int>{-1, 0, 5, 10, 12}));

    LeftistHeap<int, greater<int>> max_heap;
    max_heap.push(1);
    max_heap.push(3);
}

```

```

    assert(max_heap.top() == 3);
    return 0;
}

```

2.2.3 Skew Heap

Maintains a mergeable priority queue, returning the minimum element by default. The comparator `comp` defines priority: `comp(a, b)` is true when `a` has higher priority than `b`, so `std::greater<T>` instead creates a max-priority queue. A skew heap attempts to maintain balance by unconditionally swapping all nodes in the merge path when merging. Every other operation reduces to this merge: insertion merges a one-node heap, and extraction merges the root's two subtrees. To customize the ordering, instantiate `SkewHeap<T, Compare>` and pass the comparator to either constructor.

- `SkewHeap<T>()` constructs an empty priority queue.
- `SkewHeap<T>(lo, hi)` constructs a priority queue from the elements in the half-open iterator range `[lo, hi)`.
- `size()` returns the size of the priority queue.
- `empty()` returns whether the priority queue is empty.
- `push(v)` inserts the value `v` into the priority queue.
- `pop()` removes the highest-priority element from the priority queue.
- `top()` returns the highest-priority element in the priority queue.
- `absorb(h)` inserts every value from the distinct heap `h` and sets `h` to the empty priority queue. Both heaps must use equivalent comparators.

Time Complexity:

- $O(1)$ per call to the first constructor, `size()`, `empty()`, and `top()`.
- $O(\log n)$ amortized per call to `push()`, `pop()`, and `absorb()`, where n is the number of elements in the priority queue.
- $O(n \log n)$ amortized per call to the second constructor, where n is the distance between `lo` and `hi`.

Space Complexity:

- $O(n)$ for storage of the priority queue elements.
- $O(n)$ auxiliary stack space in the worst case for `push()`, `pop()`, `absorb()`, and destruction.
- $O(1)$ auxiliary for all other operations.

```

#include <cassert>
#include <functional>
#include <utility>

template<typename T, typename Compare = std::less<T>>
class SkewHeap {
    struct Node {
        T value;
        Node *left, *right;

        explicit Node(const T &v) : value(v), left(nullptr), right(nullptr) {}
    } *root;

    int num_nodes;
    Compare comp;

    Node *merge(Node *a, Node *b) {
        if (a == nullptr) {
            return b;
        }
        if (b == nullptr) {
            return a;
        }
        if (comp(b->value, a->value)) {
            std::swap(a, b);
        }
    }
};

```

```

    }
    std::swap(a->left, a->right);
    a->left = merge(b, a->left);
    return a;
}

static void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
    explicit SkewHeap(Compare comp = Compare())
        : root(nullptr), num_nodes(0), comp(std::move(comp)) {}

    template<typename It>
    SkewHeap(It lo, It hi, Compare comp = Compare())
        : root(nullptr), num_nodes(0), comp(std::move(comp)) {
        while (lo != hi) {
            push(*(lo++));
        }
    }

    ~SkewHeap() { clean_up(root); }
    SkewHeap(const SkewHeap &) = delete;
    SkewHeap &operator=(const SkewHeap &) = delete;
    int size() const { return num_nodes; }
    bool empty() const { return root == nullptr; }

    void push(const T &v) {
        root = merge(root, new Node(v));
        num_nodes++;
    }

    void pop() {
        assert(!empty());
        Node *tmp = root;
        root = merge(root->left, root->right);
        delete tmp;
        num_nodes--;
    }

    const T &top() const {
        assert(!empty());
        return root->value;
    }

    void absorb(SkewHeap &h) {
        assert(this != &h);
        root = merge(root, h.root);
        num_nodes += h.num_nodes;
        h.root = nullptr;
        h.num_nodes = 0;
    }
};

```

Example Usage

```

#include <vector>
using namespace std;

int main() {
    SkewHeap<int> h, h2;

```

```

assert(h.empty());
h.push(12);
h.push(10);
h2.push(5);
h2.push(-1);
h2.push(0);
h.absorb(h2);
assert(h.size() == 5);
assert(h2.empty());
vector<int> popped;
while (!h.empty()) {
    popped.push_back(h.top());
    h.pop();
}
assert((popped == vector<int>{-1, 0, 5, 10, 12}));

SkewHeap<int, greater<int>> max_heap;
max_heap.push(1);
max_heap.push(3);
assert(max_heap.top() == 3);
return 0;
}

```

2.2.4 Pairing Heap

Maintains a mergeable priority queue, returning the minimum element by default. The comparator `comp` defines priority: `comp(a, b)` is true when `a` has higher priority than `b`, so `std::greater<T>` instead creates a max-priority queue. A pairing heap is a heap-ordered multi-way tree, using a two-pass merge to self-adjust during each deletion. Unlike the other heaps in this section, it also supports `improve_key()` (the usual decrease-key operation for a min-heap) and `erase()` via handles returned by `push()`, the feature pairing heaps are best known for.

To customize the ordering, instantiate `PairingHeap<T, Compare>` and pass the comparator to either constructor.

- `PairingHeap<T>()` constructs an empty priority queue.
- `PairingHeap<T>(lo, hi)` constructs a priority queue from the elements in the half-open iterator range `[lo, hi)`.
- `size()` returns the size of the priority queue.
- `empty()` returns whether the priority queue is empty.
- `push(v)` inserts the value `v` and returns a handle to the new element. The handle remains valid until that element is removed by `pop()` or `erase()`.
- `pop()` removes the highest-priority element from the priority queue.
- `top()` returns the highest-priority element in the priority queue.
- `absorb(h)` inserts every value from the distinct heap `h` and sets `h` to the empty priority queue. Both heaps must use equivalent comparators.
- `improve_key(n, v)` replaces the value at handle `n` with a value `v` of equal or higher priority.
- `erase(n)` removes the element at handle `n` from the priority queue.

In practice, `improve_key()` runs in near-constant time, making the pairing heap the usual practical substitute for a Fibonacci heap. When handles are not needed, the lazy-deletion idiom (push updated entries and skip stale ones on `pop()`, as the Dijkstra section does) is simpler and often faster. On GNU C++ judges, `__gnu_pbds::priority_queue` with `pairing_heap_tag` provides a shorter handled and meldable heap; see 8.6 for the contest wrapper.

Time Complexity:

- $O(1)$ per call to the first constructor, `size()`, `empty()`, `push()`, `top()`, and `absorb()`.
- $O(\log n)$ amortized per call to `pop()` and `erase()`.
- $O(\log n)$ amortized per call to `improve_key()`, a safe bound to budget for. The true amortized cost is sub-logarithmic, though the long-conjectured $O(1)$ bound is provably impossible.
- $O(n)$ per call to the second constructor on the distance between `lo` and `hi`.

Space Complexity:

- $O(n)$ for storage of the priority queue elements.
- $O(n)$ auxiliary stack space in the worst case for `pop()`, `erase()`, and destruction.
- $O(1)$ auxiliary for all other operations.

```

#include <cassert>
#include <functional>
#include <utility>

template<typename T, typename Compare = std::less<T>>
class PairingHeap {
    struct Node {
        T value;
        Node *left, *next, *prev;

        explicit Node(const T &v) : value(v), left(nullptr), next(nullptr), prev(nullptr) {}

        // prev points to the previous sibling, or to the parent when this is the first child, so any
        // node can be cut from its child list in O(1).
    void add_child(Node *n) {
        n->prev = this;
        n->next = left;
        if (left != nullptr) {
            left->prev = n;
        }
        left = n;
    }
} *root;

int num_nodes;
Compare comp;

Node *merge(Node *a, Node *b) {
    if (a == nullptr) {
        return b;
    }
    if (b == nullptr) {
        return a;
    }
    if (comp(a->value, b->value)) {
        a->add_child(b);
        a->prev = a->next = nullptr;
        return a;
    }
    b->add_child(a);
    b->prev = b->next = nullptr;
    return b;
}

// Detaches non-root node n from its parent's child list.
static void cut(Node *n) {
    if (n->prev->left == n) {
        n->prev->left = n->next;
    } else {
        n->prev->next = n->next;
    }
    if (n->next != nullptr) {
        n->next->prev = n->prev;
    }
    n->prev = n->next = nullptr;
}

Node *merge_pairs(Node *n) {
    if (n == nullptr || n->next == nullptr) {
        return n;
    }
}

```



```

    Node *a = n, *b = n->next, *c = n->next->next;
    a->next = b->next = nullptr;
    return merge(merge(a, b), merge_pairs(c));
}

static void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->next);
        delete n;
    }
}

public:
    explicit PairingHeap(Compare comp = Compare())
        : root(nullptr), num_nodes(0), comp(std::move(comp)) {}

    template<typename It>
    PairingHeap(It lo, It hi, Compare comp = Compare())
        : root(nullptr), num_nodes(0), comp(std::move(comp)) {
        while (lo != hi) {
            push(*(lo++));
        }
    }

    ~PairingHeap() { clean_up(root); }
    PairingHeap(const PairingHeap &) = delete;
    PairingHeap &operator=(const PairingHeap &) = delete;
    int size() const { return num_nodes; }
    bool empty() const { return root == nullptr; }

    Node *push(const T &v) {
        Node *n = new Node(v);
        root = merge(root, n);
        num_nodes++;
        return n;
    }

    void pop() {
        assert(!empty());
        Node *tmp = root;
        root = merge_pairs(root->left);
        if (root != nullptr) {
            root->prev = root->next = nullptr;
        }
        delete tmp;
        num_nodes--;
    }

    const T &top() const {
        assert(!empty());
        return root->value;
    }

    void absorb(PairingHeap &h) {
        assert(this != &h);
        root = merge(root, h.root);
        num_nodes += h.num_nodes;
        h.root = nullptr;
        h.num_nodes = 0;
    }

    void improve_key(Node *n, const T &v) {
        assert(!comp(n->value, v));
        n->value = v;
        if (n != root) {
            cut(n);
            root = merge(root, n);
        }
    }

```

```

    }
}

void erase(Node *n) {
    if (n == root) {
        pop();
        return;
    }
    cut(n);
    root = merge(root, merge_pairs(n->left));
    delete n;
    num_nodes--;
}
};

```

Example Usage

```

#include <vector>
using namespace std;

int main() {
    PairingHeap<int> h, h2;
    assert(h.empty());
    h.push(12);
    h.push(10);
    h2.push(5);
    h2.push(-1);
    h2.push(0);
    h.absorb(h2);
    assert(h.size() == 5);
    assert(h2.empty());

    auto handle = h.push(100);
    h.improve_key(handle, -5); // 100 -> -5, the new minimum.
    assert(h.top() == -5);
    h.erase(handle); // Remove it again, restoring the original five elements.
    assert(h.size() == 5);

    vector<int> popped;
    while (!h.empty()) {
        popped.push_back(h.top());
        h.pop();
    }
    assert((popped == vector<int>{-1, 0, 5, 10, 12}));

    PairingHeap<int, greater<int>> max_heap;
    auto max_handle = max_heap.push(1);
    max_heap.improve_key(max_handle, 3);
    assert(max_heap.top() == 3);
    return 0;
}

```

2.3 Dictionaries and Ordered Sets

2.3.1 Treap

Maintain an ordered map, that is, an ordered collection of key-value pairs such that each possible key appears at most once in the collection. A binary search tree (a.k.a. BST) implements this map by inserting and deleting keys into a binary tree while upholding the BST property: for each node in the BST, every node in its left subtree has a lesser key and every node in its right subtree has a greater key.

A treap is a BST where each node also holds a randomly assigned priority, simultaneously maintaining both the BST property on keys and the min-heap property on priorities (lower priority value is closer to root). Since priorities are random, this keeps the tree balanced with high probability, making insertions and deletions run in $O(\log n)$ on average.

The comparator `comp` defines the key ordering: `comp(a, b)` is true when `a` precedes `b`. It defaults to `std::less<K>`; to customize the ordering, instantiate `Treap<K, V, Compare>` and pass the comparator to the constructor.

- `Treap<K, V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to a const value associated with key `k`, or `nullptr` if the key was not found.
- `entries()` returns all key-value entries in comparator order.
- `min()` and `max()` return the first and last entries in comparator order, or `std::nullopt` if the map is empty.
- `lower_bound(k)`, `upper_bound(k)`, `prev(k)`, and `next(k)` return the matching entry according to the comparator, or `std::nullopt` if no such entry exists. These navigation routines depend only on the BST property and can be adapted to the other BST variants later in this chapter.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `empty()`.
- $O(\log n)$ on average per call to `insert()`, `erase()`, `find()`, `min()`, `max()`, `lower_bound()`, `upper_bound()`, `prev()`, and `next()`, where n is the number of entries currently in the map.
- $O(n)$ per call to `entries()`.

Space Complexity:

- $O(n)$ for storage of the map elements.
- $O(\log n)$ auxiliary stack space on average for `insert()`, `erase()`, `entries()`, and destruction, and $O(n)$ in the worst case.
- $O(n)$ for the vector returned by `entries()`.
- $O(1)$ auxiliary for all other operations.

```
#include <cstdint>
#include <functional>
#include <optional>
#include <utility>
#include <vector>

template<typename K, typename V, typename Compare = std::less<K>>
class Treap {
    struct Node {
        static uint32_t rand32() {
            static uint32_t x = 123456789;
            x ^= x << 13;
            x ^= x >> 17;
            x ^= x << 5;
            return x;
        }
    };

    K key;
    V value;
    uint32_t priority;
    Node *left, *right;

    Node(const K &k, const V &v)
        : key(k), value(v), priority(rand32()), left(nullptr), right(nullptr) {}
} *root;
```

```

int num_nodes;
Compare comp;

static void rotate_left(Node *&n) {
    Node *tmp = n;
    n = n->right;
    tmp->right = n->left;
    n->left = tmp;
}

static void rotate_right(Node *&n) {
    Node *tmp = n;
    n = n->left;
    tmp->left = n->right;
    n->right = tmp;
}

bool insert(Node *&n, const K &k, const V &v) {
    if (n == nullptr) {
        n = new Node(k, v);
        num_nodes++;
        return true;
    }
    if (comp(k, n->key) && insert(n->left, k, v)) {
        if (n->left->priority < n->priority) {
            rotate_right(n);
        }
        return true;
    }
    if (comp(n->key, k) && insert(n->right, k, v)) {
        if (n->right->priority < n->priority) {
            rotate_left(n);
        }
        return true;
    }
    return false;
}

bool erase(Node *&n, const K &k) {
    if (n == nullptr) {
        return false;
    }
    if (comp(k, n->key)) {
        return erase(n->left, k);
    } else if (comp(n->key, k)) {
        return erase(n->right, k);
    }
    if (n->left != nullptr && n->right != nullptr) {
        if (n->left->priority < n->right->priority) {
            rotate_right(n);
            return erase(n->right, k);
        }
        rotate_left(n);
        return erase(n->left, k);
    }
    Node *tmp = (n->left != nullptr) ? n->left : n->right;
    delete n;
    n = tmp;
    num_nodes--;
    return true;
}

static void collect_entries(Node *n, std::vector<std::pair<K, V>> &res) {
    if (n != nullptr) {
        collect_entries(n->left, res);
        res.emplace_back(n->key, n->value);
        collect_entries(n->right, res);
    }
}

```

```

    }
}

static std::optional<std::pair<K, V>> entry(Node *n) {
    if (n == nullptr) {
        return std::nullopt;
    }
    return std::pair<K, V>{n->key, n->value};
}

static void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
    explicit Treap(Compare comp = Compare()) : root(nullptr), num_nodes(0), comp(std::move(comp)) {}

    ~Treap() { clean_up(root); }
    Treap(const Treap &) = delete;
    Treap &operator=(const Treap &) = delete;
    int size() const { return num_nodes; }
    bool empty() const { return root == nullptr; }
    bool insert(const K &k, const V &v) { return insert(root, k, v); }
    bool erase(const K &k) { return erase(root, k); }

    const V *find(const K &k) const {
        Node *n = root;
        while (n != nullptr) {
            if (comp(k, n->key)) {
                n = n->left;
            } else if (comp(n->key, k)) {
                n = n->right;
            } else {
                return &(n->value);
            }
        }
        return nullptr;
    }

    std::vector<std::pair<K, V>> entries() const {
        std::vector<std::pair<K, V>> res;
        res.reserve(num_nodes);
        collect_entries(root, res);
        return res;
    }

    std::optional<std::pair<K, V>> min() const {
        Node *n = root;
        if (n == nullptr) {
            return std::nullopt;
        }
        while (n->left != nullptr) {
            n = n->left;
        }
        return std::pair<K, V>{n->key, n->value};
    }

    std::optional<std::pair<K, V>> max() const {
        Node *n = root;
        if (n == nullptr) {
            return std::nullopt;
        }
        while (n->right != nullptr) {
            n = n->right;
        }
    }

```

```

    }
    return std::pair<K, V>(n->key, n->value);
}

std::optional<std::pair<K, V>> lower_bound(const K &k) const {
    Node *n = root, *best = nullptr;
    while (n != nullptr) {
        if (!comp(n->key, k)) {
            best = n;
            n = n->left;
        } else {
            n = n->right;
        }
    }
    return entry(best);
}

std::optional<std::pair<K, V>> upper_bound(const K &k) const {
    Node *n = root, *best = nullptr;
    while (n != nullptr) {
        if (comp(k, n->key)) {
            best = n;
            n = n->left;
        } else {
            n = n->right;
        }
    }
    return entry(best);
}

std::optional<std::pair<K, V>> prev(const K &k) const {
    Node *n = root, *best = nullptr;
    while (n != nullptr) {
        if (comp(n->key, k)) {
            best = n;
            n = n->right;
        } else {
            n = n->left;
        }
    }
    return entry(best);
}

std::optional<std::pair<K, V>> next(const K &k) const { return upper_bound(k); }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    Treap<int, char> t;
    t.insert(2, 'b');
    t.insert(1, 'a');
    t.insert(3, 'c');
    t.insert(5, 'e');
    assert(t.insert(4, 'd'));
    assert(*t.find(4) == 'd');
    assert(!t.insert(4, 'd'));
    assert(t.min()->first == 1 && t.min()->second == 'a');
    assert(t.max()->first == 5 && t.max()->second == 'e');
    assert(t.lower_bound(4)->first == 4);
    assert(t.upper_bound(4)->first == 5);
    assert(t.prev(4)->first == 3);
    assert(t.next(4)->first == 5);
}

```

```

assert(!t.prev(1));
assert(!t.next(5));
assert(
    (t.entries() == vector<pair<int, char>>{{1, 'a'}, {2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}})
);
assert(t.erase(1));
assert(!t.erase(1));
assert(t.find(1) == nullptr);
assert((t.entries() == vector<pair<int, char>>{{2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}}));

Treap<int, char, greater<int>> descending;
for (int key : {2, 1, 3}) {
    descending.insert(key, '0' + key);
}
assert((descending.entries() == vector<pair<int, char>>{{3, '3'}, {2, '2'}, {1, '1'}}));
assert(descending.min()->first == 3 && descending.max()->first == 1);
assert(descending.lower_bound(2)->first == 2);
assert(descending.upper_bound(2)->first == 1);
assert(descending.prev(2)->first == 3 && descending.next(2)->first == 1);
assert(descending.erase(2) && descending.find(2) == nullptr);
return 0;
}

```

2.3.2 AVL Tree

Maintain an ordered map, that is, an ordered collection of key-value pairs such that each possible key appears at most once in the collection. An AVL tree is a binary search tree balanced by height, guaranteeing $O(\log n)$ worst-case running time in insertions and deletions by making sure that the heights of the left and right subtrees at every node differ by at most 1. Whenever an insertion or deletion breaks this invariant, it is repaired with one or two rotations at each affected node along the search path.

The comparator `comp` defines the key ordering: `comp(a, b)` is true when `a` precedes `b`. It defaults to `std::less<K>`; to customize the ordering, instantiate `AVLTree<K, V, Compare>` and pass the comparator to the constructor.

- `AVLTree<K, V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to a const value associated with key `k`, or `nullptr` if the key was not found.
- `entries()` returns all key-value entries in comparator order.

The comparator-aware navigation routines `min()`, `max()`, `lower_bound(k)`, `upper_bound(k)`, `prev(k)`, and `next(k)` from the treap in 2.3.1 depend only on the BST property and may be adapted here as needed.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `empty()`.
- $O(\log n)$ per call to `insert()`, `erase()`, and `find()`, where n is the number of entries currently in the map.
- $O(n)$ per call to `entries()`.

Space Complexity:

- $O(n)$ for storage of the map elements.
- $O(\log n)$ auxiliary stack space for `insert()`, `erase()`, `entries()`, and destruction.
- $O(n)$ for the vector returned by `entries()`.
- $O(1)$ auxiliary for all other operations.

```

#include <algorithm>
#include <functional>
#include <utility>
#include <vector>

template<typename K, typename V, typename Compare = std::less<K>>
class AVLTree {
    struct Node {
        K key;
        V value;
        int height;
        Node *left, *right;

        Node(const K &k, const V &v) : key(k), value(v), height(1), left(nullptr), right(nullptr) {}
    } *root;

    int num_nodes;
    Compare comp;

    static int height(Node *n) { return (n != nullptr) ? n->height : 0; }

    static void update_height(Node *n) {
        if (n != nullptr) {
            n->height = 1 + std::max(height(n->left), height(n->right));
        }
    }

    static void rotate_left(Node *&n) {
        Node *tmp = n;
        n = n->right;
        tmp->right = n->left;
        n->left = tmp;
        update_height(tmp);
        update_height(n);
    }

    static void rotate_right(Node *&n) {
        Node *tmp = n;
        n = n->left;
        tmp->left = n->right;
        n->right = tmp;
        update_height(tmp);
        update_height(n);
    }

    static int balance_factor(Node *n) {
        return (n != nullptr) ? (height(n->left) - height(n->right)) : 0;
    }

    static void rebalance(Node *&n) {
        if (n == nullptr) {
            return;
        }
        update_height(n);
        int bf = balance_factor(n);
        if (bf > 1 && balance_factor(n->left) >= 0) {
            rotate_right(n);
        } else if (bf > 1 && balance_factor(n->left) < 0) {
            rotate_left(n->left);
            rotate_right(n);
        } else if (bf < -1 && balance_factor(n->right) <= 0) {
            rotate_left(n);
        } else if (bf < -1 && balance_factor(n->right) > 0) {
            rotate_right(n->right);
            rotate_left(n);
        }
    }
}

```



```

bool insert(Node *&n, const K &k, const V &v) {
    if (n == nullptr) {
        n = new Node(k, v);
        num_nodes++;
        return true;
    }
    if ((comp(k, n->key) && insert(n->left, k, v)) || (comp(n->key, k) && insert(n->right, k, v))) {
        rebalance(n);
        return true;
    }
    return false;
}

bool erase(Node *&n, const K &k) {
    if (n == nullptr) {
        return false;
    }
    if (!(comp(k, n->key) || comp(n->key, k))) {
        if (n->left != nullptr && n->right != nullptr) {
            Node *tmp = n->right;
            while (tmp->left != nullptr) {
                tmp = tmp->left;
            }
            K successor_key = tmp->key;
            n->key = successor_key;
            n->value = tmp->value;
            if (!erase(n->right, successor_key)) {
                return false;
            }
        }
        else {
            Node *tmp = (n->left != nullptr) ? n->left : n->right;
            delete n;
            n = tmp;
            num_nodes--;
        }
        rebalance(n);
        return true;
    }
    if ((comp(k, n->key) && erase(n->left, k)) || (comp(n->key, k) && erase(n->right, k))) {
        rebalance(n);
        return true;
    }
    return false;
}

static void collect_entries(Node *n, std::vector<std::pair<K, V>> &res) {
    if (n != nullptr) {
        collect_entries(n->left, res);
        res.emplace_back(n->key, n->value);
        collect_entries(n->right, res);
    }
}

static void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
explicit AVLTree(Compare comp = Compare()) : root(nullptr), num_nodes(0), comp(std::move(comp)) {}

~AVLTree() { clean_up(root); }
AVLTree(const AVLTree &) = delete;
AVLTree &operator=(const AVLTree &) = delete;
int size() const { return num_nodes; }

```

```

bool empty() const { return root == nullptr; }
bool insert(const K &k, const V &v) { return insert(root, k, v); }
bool erase(const K &k) { return erase(root, k); }

const V *find(const K &k) const {
    Node *n = root;
    while (n != nullptr) {
        if (comp(k, n->key)) {
            n = n->left;
        } else if (comp(n->key, k)) {
            n = n->right;
        } else {
            return &(n->value);
        }
    }
    return nullptr;
}

std::vector<std::pair<K, V>> entries() const {
    std::vector<std::pair<K, V>> res;
    res.reserve(num_nodes);
    collect_entries(root, res);
    return res;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    AVLTree<int, char> t;
    assert(t.empty());
    t.insert(2, 'b');
    t.insert(1, 'a');
    t.insert(3, 'c');
    t.insert(5, 'e');
    assert(t.insert(4, 'd'));
    assert(!t.empty() && t.size() == 5);
    assert(*t.find(4) == 'd');
    assert(!t.insert(4, 'd'));
    assert(t.size() == 5);
    assert(
        (t.entries() == vector<pair<int, char>>{{1, 'a'}, {2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}})
    );
    assert(t.erase(1));
    assert(!t.erase(1));
    assert(t.find(1) == nullptr);
    assert(t.size() == 4);
    assert((t.entries() == vector<pair<int, char>>{{2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}}));

    AVLTree<int, int> deep_successor;
    for (int key : {20, 10, 30, 25, 40, 22}) {
        deep_successor.insert(key, key);
    }
    assert(deep_successor.erase(20));
    assert(
        (deep_successor.entries() ==
         vector<pair<int, int>>{{10, 10}, {22, 22}, {25, 25}, {30, 30}, {40, 40}})
    );

    AVLTree<int, char, greater<int>> descending;
    for (int key : {2, 1, 3}) {
        descending.insert(key, '0' + key);
    }
}

```

```

assert((descending.entries() == vector<pair<int, char>>{{3, '3'}, {2, '2'}, {1, '1'}}));
assert(*descending.find(2) == '2');
assert(descending.erase(2) && descending.find(2) == nullptr);
return 0;
}

```

2.3.3 Red-Black Tree

Maintain an ordered map, that is, an ordered collection of key-value pairs such that each possible key appears at most once in the collection. A red-black tree is a binary search tree balanced by coloring its nodes red or black, then constraining node colors on any simple path from the root to a leaf. Specifically, a red node may never have a red child, and every path from the root to a null leaf must pass through the same number of black nodes, which together bound the tree's height to $O(\log n)$. Insertions and deletions repair these invariants by recoloring nodes and performing rotations.

The comparator `comp` defines the key ordering: `comp(a, b)` is true when `a` precedes `b`. The comparator defaults to `std::less<K>`; to customize the ordering, instantiate `RedBlackTree<K, V, Compare>` and pass the comparator to the constructor. The sentinel node also requires `K` and `V` to be default constructible.

- `RedBlackTree<K, V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to a const value associated with key `k`, or `nullptr` if the key was not found.
- `entries()` returns all key-value entries in comparator order.

The comparator-aware navigation routines `min()`, `max()`, `lower_bound(k)`, `upper_bound(k)`, `prev(k)`, and `next(k)` from the treap in 2.3.1 depend only on the BST property and may be adapted here as needed.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `empty()`.
- $O(\log n)$ per call to `insert()`, `erase()`, and `find()`, where n is the number of entries currently in the map.
- $O(n)$ per call to `entries()`.

Space Complexity:

- $O(n)$ for storage of the map elements.
- $O(\log n)$ auxiliary stack space for `entries()` and destruction.
- $O(n)$ for the vector returned by `entries()`.
- $O(1)$ auxiliary for all other operations.

```

#include <algorithm>
#include <functional>
#include <utility>
#include <vector>

template<typename K, typename V, typename Compare = std::less<K>>
class RedBlackTree {
    enum Color { RED, BLACK };
    struct Node {
        K key;
        V value;
        Color color;
        Node *left, *right, *parent;
    };

```

```

Node(const K &k, const V &v, Color c)
    : key(k), value(v), color(c), left(nullptr), right(nullptr), parent(nullptr) {}
} *root, *LEAF_NIL;

int num_nodes;
Compare comp;

void rotate_left(Node *n) {
    Node *tmp = n->right;
    if ((n->right = tmp->left) != LEAF_NIL) {
        n->right->parent = n;
    }
    if ((tmp->parent = n->parent) == LEAF_NIL) {
        root = tmp;
    } else if (n->parent->left == n) {
        n->parent->left = tmp;
    } else {
        n->parent->right = tmp;
    }
    tmp->left = n;
    n->parent = tmp;
}

void rotate_right(Node *n) {
    Node *tmp = n->left;
    if ((n->left = tmp->right) != LEAF_NIL) {
        n->left->parent = n;
    }
    if ((tmp->parent = n->parent) == LEAF_NIL) {
        root = tmp;
    } else if (n->parent->right == n) {
        n->parent->right = tmp;
    } else {
        n->parent->left = tmp;
    }
    tmp->right = n;
    n->parent = tmp;
}

void insert_fix(Node *n) {
    while (n->parent->color == RED) {
        Node *parent = n->parent;
        Node *grandparent = n->parent->parent;
        if (parent == grandparent->left) {
            Node *uncle = grandparent->right;
            if (uncle->color == RED) {
                grandparent->color = RED;
                parent->color = BLACK;
                uncle->color = BLACK;
                n = grandparent;
            } else {
                if (n == parent->right) {
                    rotate_left(parent);
                    n = parent;
                    parent = n->parent;
                }
                rotate_right(grandparent);
                std::swap(parent->color, grandparent->color);
                n = parent;
            }
        } else if (parent == grandparent->right) {
            Node *uncle = grandparent->left;
            if (uncle->color == RED) {
                grandparent->color = RED;
                parent->color = BLACK;
                uncle->color = BLACK;
                n = grandparent;
            } else {

```

```

        if (n == parent->left) {
            rotate_right(parent);
            n = parent;
            parent = n->parent;
        }
        rotate_left(grandparent);
        std::swap(parent->color, grandparent->color);
        n = parent;
    }
}
root->color = BLACK;
}

void replace(Node *n, Node *replacement) {
    if (n->parent == LEAF_NIL) {
        root = replacement;
    } else if (n == n->parent->left) {
        n->parent->left = replacement;
    } else {
        n->parent->right = replacement;
    }
    replacement->parent = n->parent;
}

void erase_fix(Node *n) {
    while (n != root && n->color == BLACK) {
        Node *parent = n->parent;
        if (n == parent->left) {
            Node *sibling = parent->right;
            if (sibling->color == RED) {
                sibling->color = BLACK;
                parent->color = RED;
                rotate_left(parent);
                sibling = parent->right;
            }
            if (sibling->left->color == BLACK && sibling->right->color == BLACK) {
                sibling->color = RED;
                n = parent;
            } else {
                if (sibling->right->color == BLACK) {
                    sibling->left->color = BLACK;
                    sibling->color = RED;
                    rotate_right(sibling);
                    sibling = parent->right;
                }
                sibling->color = parent->color;
                parent->color = BLACK;
                sibling->right->color = BLACK;
                rotate_left(parent);
                n = root;
            }
        } else {
            Node *sibling = parent->left;
            if (sibling->color == RED) {
                sibling->color = BLACK;
                parent->color = RED;
                rotate_right(parent);
                sibling = parent->left;
            }
            if (sibling->left->color == BLACK && sibling->right->color == BLACK) {
                sibling->color = RED;
                n = parent;
            } else {
                if (sibling->left->color == BLACK) {
                    sibling->right->color = BLACK;
                    sibling->color = RED;
                    rotate_left(sibling);
                }
            }
        }
    }
}

```

```

        sibling = parent->left;
    }
    sibling->color = parent->color;
    parent->color = BLACK;
    sibling->left->color = BLACK;
    rotate_right(parent);
    n = root;
}
}
}
n->color = BLACK;
}

void collect_entries(Node *n, std::vector<std::pair<K, V>> &res) const {
    if (n != LEAF_NIL) {
        collect_entries(n->left, res);
        res.push_back({n->key, n->value});
        collect_entries(n->right, res);
    }
}

void clean_up(Node *n) {
    if (n != LEAF_NIL) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
explicit RedBlackTree(Compare comp = Compare()) : num_nodes(0), comp(std::move(comp)) {
    root = LEAF_NIL = new Node(K(), V(), BLACK);
    LEAF_NIL->left = LEAF_NIL->right = LEAF_NIL;
}

~RedBlackTree() {
    clean_up(root);
    delete LEAF_NIL;
}

RedBlackTree(const RedBlackTree &) = delete;
RedBlackTree &operator=(const RedBlackTree &) = delete;
int size() const { return num_nodes; }
bool empty() const { return num_nodes == 0; }

bool insert(const K &k, const V &v) {
    Node *curr = root, *prev = LEAF_NIL;
    while (curr != LEAF_NIL) {
        prev = curr;
        if (comp(k, curr->key)) {
            curr = curr->left;
        } else if (comp(curr->key, k)) {
            curr = curr->right;
        } else {
            return false;
        }
    }
    Node *n = new Node(k, v, RED);
    n->parent = prev;
    if (prev == LEAF_NIL) {
        root = n;
    } else if (comp(k, prev->key)) {
        prev->left = n;
    } else {
        prev->right = n;
    }
    n->left = n->right = LEAF_NIL;
    insert_fix(n);
}

```

```

    num_nodes++;
    return true;
}

bool erase(const K &k) {
    Node *n = root;
    while (n != LEAF_NIL) {
        if (comp(k, n->key)) {
            n = n->left;
        } else if (comp(n->key, k)) {
            n = n->right;
        } else {
            break;
        }
    }
    if (n == LEAF_NIL) {
        return false;
    }
    Color color = n->color;
    Node *replacement;
    if (n->left == LEAF_NIL) {
        replacement = n->right;
        replace(n, n->right);
    } else if (n->right == LEAF_NIL) {
        replacement = n->left;
        replace(n, n->left);
    } else {
        Node *tmp = n->right;
        while (tmp->left != LEAF_NIL) {
            tmp = tmp->left;
        }
        color = tmp->color;
        replacement = tmp->right;
        if (tmp->parent == n) {
            replacement->parent = tmp;
        } else {
            replace(tmp, tmp->right);
            tmp->right = n->right;
            tmp->right->parent = tmp;
        }
        replace(n, tmp);
        tmp->left = n->left;
        tmp->left->parent = tmp;
        tmp->color = n->color;
    }
    delete n;
    if (color == BLACK) {
        erase_fix(replacement);
    }
    num_nodes--;
    return true;
}

const V *find(const K &k) const {
    Node *n = root;
    while (n != LEAF_NIL) {
        if (comp(k, n->key)) {
            n = n->left;
        } else if (comp(n->key, k)) {
            n = n->right;
        } else {
            return &(n->value);
        }
    }
    return nullptr;
}

std::vector<std::pair<K, V>> entries() const {

```

```

    std::vector<std::pair<K, V>> res;
    res.reserve(num_nodes);
    collect_entries(root, res);
    return res;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    RedBlackTree<int, char> t;
    assert(t.empty());
    t.insert(2, 'b');
    t.insert(1, 'a');
    t.insert(3, 'c');
    t.insert(5, 'e');
    assert(t.insert(4, 'd'));
    assert(!t.empty() && t.size() == 5);
    assert(*t.find(4) == 'd');
    assert(!t.insert(4, 'd'));
    assert(t.size() == 5);
    assert(
        (t.entries() == vector<pair<int, char>>{{1, 'a'}, {2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}})
    );
    assert(t.erase(1));
    assert(!t.erase(1));
    assert(t.find(1) == nullptr);
    assert(t.size() == 4);
    assert((t.entries() == vector<pair<int, char>>{{2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}}));
    for (int key : {2, 3, 4, 5}) {
        assert(t.erase(key));
    }
    assert(t.empty() && t.size() == 0);

    RedBlackTree<int, char, greater<int>> descending;
    for (int key : {2, 1, 3}) {
        descending.insert(key, '0' + key);
    }
    assert((descending.entries() == vector<pair<int, char>>{{3, '3'}, {2, '2'}, {1, '1'}}));
    assert(*descending.find(2) == '2');
    assert(descending.erase(2) && descending.find(2) == nullptr);
    return 0;
}

```

2.3.4 Splay Tree

Maintain an ordered map, that is, an ordered collection of key-value pairs such that each possible key appears at most once in the collection. A splay tree is a balanced binary search tree with the additional property that recently accessed elements are quick to access again. Every operation “splays” its target node up to the root through a series of rotations, reshaping the tree so that frequently accessed keys settle near the top. A single operation may degrade to $O(n)$, but any sequence of operations averages out to $O(\log n)$ amortized each.

The comparator `comp` defines the key ordering: `comp(a, b)` is true when `a` precedes `b`. It defaults to `std::less<K>`; to customize the ordering, instantiate `SplayTree<K, V, Compare>` and pass the comparator to the constructor.

- `SplayTree<K, V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.

- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to a const value associated with key `k`, or `nullptr` if the key was not found.
- `entries()` returns all key-value entries in comparator order.

The comparator-aware navigation routines `min()`, `max()`, `lower_bound(k)`, `upper_bound(k)`, `prev(k)`, and `next(k)` from the treap in 2.3.1 depend only on the BST property and may be adapted here as needed.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `empty()`.
- $O(\log n)$ amortized per call to `insert()`, `erase()`, and `find()`, where n is the number of entries currently in the map.
- $O(n)$ per call to `entries()`.

Space Complexity:

- $O(n)$ for storage of the map elements.
- $O(n)$ auxiliary stack space in the worst case for `insert()`, `erase()`, `find()`, `entries()`, and destruction.
- $O(n)$ for the vector returned by `entries()`.
- $O(1)$ auxiliary for `size()` and `empty()`.

```
#include <functional>
#include <utility>
#include <vector>

template<typename K, typename V, typename Compare = std::less<K>>
class SplayTree {
    struct Node {
        K key;
        V value;
        Node *left, *right;

        Node(const K &k, const V &v) : key(k), value(v), left(nullptr), right(nullptr) {}
    } *root;

    int num_nodes;
    Compare comp;

    static void rotate_left(Node *&n) {
        Node *tmp = n;
        n = n->right;
        tmp->right = n->left;
        n->left = tmp;
    }

    static void rotate_right(Node *&n) {
        Node *tmp = n;
        n = n->left;
        tmp->left = n->right;
        n->right = tmp;
    }

    void splay(Node *&n, const K &k) {
        if (n == nullptr) {
            return;
        }
        if (comp(k, n->key) && n->left != nullptr) {
            if (comp(k, n->left->key)) {
                splay(n->left->left, k);
                rotate_right(n);
            } else if (comp(n->left->key, k)) {
```

```

        splay(n->left->right, k);
        if (n->left->right != nullptr) {
            rotate_left(n->left);
        }
    }
    if (n->left != nullptr) {
        rotate_right(n);
    }
} else if (comp(n->key, k) && n->right != nullptr) {
    if (comp(k, n->right->key)) {
        splay(n->right->left, k);
        if (n->right->left != nullptr) {
            rotate_right(n->right);
        }
    } else if (comp(n->right->key, k)) {
        splay(n->right->right, k);
        rotate_left(n);
    }
    if (n->right != nullptr) {
        rotate_left(n);
    }
}
}

bool insert(Node *&n, const K &k, const V &v) {
    if (n == nullptr) {
        n = new Node(k, v);
        num_nodes++;
        return true;
    }
    splay(n, k);
    if (comp(k, n->key)) {
        Node *tmp = new Node(k, v);
        tmp->left = n->left;
        tmp->right = n;
        n->left = nullptr;
        n = tmp;
    } else if (comp(n->key, k)) {
        Node *tmp = new Node(k, v);
        tmp->left = n;
        tmp->right = n->right;
        n->right = nullptr;
        n = tmp;
    } else {
        return false;
    }
    num_nodes++;
    return true;
}

bool erase(Node *&n, const K &k) {
    if (n == nullptr) {
        return false;
    }
    splay(n, k);
    if (comp(k, n->key) || comp(n->key, k)) {
        return false;
    }
    Node *tmp = n;
    if (n->left == nullptr) {
        n = n->right;
    } else {
        splay(n->left, k);
        n = n->left;
        n->right = tmp->right;
    }
    delete tmp;
    num_nodes--;
}

```

```

    return true;
}

static void collect_entries(Node *n, std::vector<std::pair<K, V>> &res) {
    if (n != nullptr) {
        collect_entries(n->left, res);
        res.emplace_back(n->key, n->value);
        collect_entries(n->right, res);
    }
}

static void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
    explicit SplayTree(Compare comp = Compare())
        : root(nullptr), num_nodes(0), comp(std::move(comp)) {}

    ~SplayTree() { clean_up(root); }
    SplayTree(const SplayTree &) = delete;
    SplayTree &operator=(const SplayTree &) = delete;
    int size() const { return num_nodes; }
    bool empty() const { return root == nullptr; }
    bool insert(const K &k, const V &v) { return insert(root, k, v); }
    bool erase(const K &k) { return erase(root, k); }

    const V *find(const K &k) {
        splay(root, k);
        if (root == nullptr) {
            return nullptr;
        }
        return (comp(k, root->key) || comp(root->key, k)) ? nullptr : &(root->value);
    }

    std::vector<std::pair<K, V>> entries() const {
        std::vector<std::pair<K, V>> res;
        res.reserve(num_nodes);
        collect_entries(root, res);
        return res;
    }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    SplayTree<int, char> t;
    assert(t.empty());
    t.insert(2, 'b');
    t.insert(1, 'a');
    t.insert(3, 'c');
    t.insert(5, 'e');
    assert(t.insert(4, 'd'));
    assert(!t.empty() && t.size() == 5);
    assert(*t.find(4) == 'd');
    assert(!t.insert(4, 'd'));
    assert(
        (t.entries() == vector<pair<int, char>>{{1, 'a'}, {2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}})
    );
}

```

```

assert(t.erase(1));
assert(!t.erase(1));
assert(t.find(1) == nullptr);
assert(t.size() == 4);
assert((t.entries() == vector<pair<int, char>>{{2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}}));

SplayTree<int, char, greater<int>> descending;
for (int key : {2, 1, 3}) {
    descending.insert(key, '0' + key);
}
assert((descending.entries() == vector<pair<int, char>>{{3, '3'}, {2, '2'}, {1, '1'}}));
assert(*descending.find(2) == '2');
assert(descending.erase(2) && descending.find(2) == nullptr);
return 0;
}

```

2.3.5 Size Balanced Tree

Maintain an ordered map, that is, an ordered collection of key-value pairs such that each possible key appears at most once in the collection. In addition, support queries for keys given their ranks as well as queries for the ranks of given keys. A size balanced tree augments each node with the size of its subtree, using it to maintain balance and compute order statistics. After each update, rotations restore the invariant that every subtree is at least as large as each of its sibling's child subtrees, keeping the height logarithmic.

The comparator `comp` defines the key ordering: `comp(a, b)` is true when `a` precedes `b`. It defaults to `std::less<K>`; to customize the ordering, instantiate `SBTree<K, V, Compare>` and pass the comparator to the constructor.

- `SBTree<K, V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to a const value associated with key `k`, or `nullptr` if the key was not found.
- `find_by_order(k)` returns a key-value pair of the node with a key of 0-based rank `k`, which must lie in the range `[0, size())`.
- `order_of_key(x)` returns the number of keys that precede `x` in comparator order. The key does not need to be present in the map.
- `entries()` returns all key-value entries in comparator order.

The comparator-aware navigation routines `min()`, `max()`, `lower_bound(k)`, `upper_bound(k)`, `prev(k)`, and `next(k)` from the treap in 2.3.1 depend only on the BST property and may be adapted here as needed. For contest use on GNU C++ judges, PBDS ordered trees provide the same order-statistic operations with much less code; see 8.6. This implementation is useful when PBDS is unavailable or when customizing tree internals.

The order-statistic API matches GNU PBDS naming and 0-based rank conventions.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `empty()`.
- $O(\log n)$ per call to `insert()`, `erase()`, `find()`, `find_by_order()`, and `order_of_key()`, where n is the number of entries currently in the map.
- $O(n)$ per call to `entries()`.

Space Complexity:

- $O(n)$ for storage of the map elements.
- $O(\log n)$ auxiliary stack space for `insert()`, `erase()`, `entries()`, and destruction.

- $O(n)$ for the vector returned by `entries()`.
- $O(1)$ auxiliary for all other operations.

```

#include <cassert>
#include <functional>
#include <utility>
#include <vector>

template<typename K, typename V, typename Compare = std::less<K>>
class SBTTree {
    struct Node {
        K key;
        V value;
        int size;
        Node *left, *right;

        Node(const K &k, const V &v) : key(k), value(v), size(1), left(nullptr), right(nullptr) {}

        inline Node *&child(int c) { return (c == 0) ? left : right; }

        void update() {
            size = 1;
            if (left != nullptr) {
                size += left->size;
            }
            if (right != nullptr) {
                size += right->size;
            }
        }
    } *root;

    Compare comp;

    static inline int size(Node *n) { return (n == nullptr) ? 0 : n->size; }

    static void rotate(Node *&n, int c) {
        Node *tmp = n->child(c);
        n->child(c) = tmp->child(!c);
        tmp->child(!c) = n;
        n->update();
        tmp->update();
        n = tmp;
    }

    static void maintain(Node *&n, int c) {
        if (n == nullptr || n->child(c) == nullptr) {
            return;
        }
        Node *&tmp = n->child(c);
        if (size(tmp->child(c)) > size(n->child(!c))) {
            rotate(n, c);
        } else if (size(tmp->child(!c)) > size(n->child(!c))) {
            rotate(tmp, !c);
            rotate(n, c);
        } else {
            return;
        }
        maintain(n->left, 0);
        maintain(n->right, 1);
        maintain(n, 0);
        maintain(n, 1);
    }

    bool insert(Node *&n, const K &k, const V &v) {
        if (n == nullptr) {
            n = new Node(k, v);
            return true;
        }
    }

```

```

    }
    bool found;
    if (comp(k, n->key)) {
        found = insert(n->left, k, v);
        maintain(n, 0);
    } else if (comp(n->key, k)) {
        found = insert(n->right, k, v);
        maintain(n, 1);
    } else {
        return false;
    }
    n->update();
    return found;
}

bool erase(Node *&n, const K &k) {
    if (n == nullptr) {
        return false;
    }
    bool found;
    int c = comp(k, n->key);
    if (comp(k, n->key)) {
        found = erase(n->left, k);
    } else if (comp(n->key, k)) {
        found = erase(n->right, k);
    } else {
        if (n->right == nullptr || n->left == nullptr) {
            Node *tmp = n;
            n = (n->right == nullptr) ? n->left : n->right;
            delete tmp;
            return true;
        }
        Node *p = n->right;
        while (p->left != nullptr) {
            p = p->left;
        }
        K successor_key = p->key;
        n->key = successor_key;
        n->value = p->value;
        found = erase(n->right, successor_key);
    }
    maintain(n, c);
    n->update();
    return found;
}

static std::pair<K, V> find_by_order(Node *n, int k) {
    int left_size = size(n->left);
    if (k < left_size) {
        return find_by_order(n->left, k);
    } else if (k > left_size) {
        return find_by_order(n->right, k - left_size - 1);
    }
    return {n->key, n->value};
}

int order_of_key(Node *n, const K &x) const {
    if (n == nullptr) {
        return 0;
    }
    if (comp(x, n->key)) {
        return order_of_key(n->left, x);
    } else if (comp(n->key, x)) {
        return order_of_key(n->right, x) + size(n->left) + 1;
    }
    return size(n->left);
}

```

```

static void collect_entries(Node *n, std::vector<std::pair<K, V>> &res) {
    if (n != nullptr) {
        collect_entries(n->left, res);
        res.emplace_back(n->key, n->value);
        collect_entries(n->right, res);
    }
}

static void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
    explicit SBTTree(Compare comp = Compare()) : root(nullptr), comp(std::move(comp)) {}

    ~SBTTree() { clean_up(root); }
    SBTTree(const SBTTree &) = delete;
    SBTTree &operator=(const SBTTree &) = delete;
    int size() const { return size(root); }
    bool empty() const { return root == nullptr; }
    bool insert(const K &k, const V &v) { return insert(root, k, v); }
    bool erase(const K &k) { return erase(root, k); }
    int order_of_key(const K &x) const { return order_of_key(root, x); }

    const V *find(const K &k) const {
        Node *n = root;
        while (n != nullptr) {
            if (comp(k, n->key)) {
                n = n->left;
            } else if (comp(n->key, k)) {
                n = n->right;
            } else {
                return &(n->value);
            }
        }
        return nullptr;
    }

    std::pair<K, V> find_by_order(int k) const {
        assert(0 <= k && k < size(root));
        return find_by_order(root, k);
    }

    std::vector<std::pair<K, V>> entries() const {
        std::vector<std::pair<K, V>> res;
        res.reserve(size(root));
        collect_entries(root, res);
        return res;
    }
};

```

Example Usage

```

using namespace std;

int main() {
    SBTTree<int, char> t;
    t.insert(2, 'b');
    t.insert(1, 'a');
    t.insert(3, 'c');
    t.insert(5, 'e');
    assert(t.insert(4, 'd'));
}

```

```

assert(*t.find(4) == 'd');
assert(!t.insert(4, 'd'));
assert(
    (t.entries() == vector<pair<int, char>>{{1, 'a'}, {2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}})
);
assert(t.erase(1));
assert(!t.erase(1));
assert(t.find(1) == nullptr);
assert((t.entries() == vector<pair<int, char>>{{2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}}));
assert(t.order_of_key(2) == 0);
assert(t.order_of_key(3) == 1);
assert(t.order_of_key(5) == 3);
assert(t.order_of_key(6) == 4);
assert(t.find_by_order(0).first == 2);
assert(t.find_by_order(1).first == 3);
assert(t.find_by_order(2).first == 4);

SBTree<int, char> replacement;
replacement.insert(2, 'b');
replacement.insert(1, 'a');
replacement.insert(3, 'c');
assert(replacement.erase(2));
assert(*replacement.find(3) == 'c');

SBTree<int, char, greater<int>> descending;
for (int key : {2, 1, 3}) {
    descending.insert(key, '0' + key);
}
assert((descending.entries() == vector<pair<int, char>>{{3, '3'}, {2, '2'}, {1, '1'}}));
assert(descending.order_of_key(2) == 1);
assert(descending.find_by_order(0).first == 3);
assert(descending.erase(2) && descending.find(2) == nullptr);
return 0;
}

```

2.3.6 Interval Treap

Maintain an ordered map from closed, one-dimensional intervals to values while supporting efficient reporting of any or all entries that intersect with a given query interval. A treap is used to process the entries, where keys are compared lexicographically as pairs. Each node is additionally augmented with the maximum upper endpoint among the intervals in its subtree, letting intersection queries skip any subtree whose maximum falls below the query interval's lower bound.

This implementation uses `std::pair` to represent intervals and requires `operator<` on the numeric key type $\mathbb{K}$. Every interval must satisfy `!(hi < lo)`.

- `IntervalTreap<K, V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(lo, hi, v)` adds an entry with key `[lo, hi]` and value `v` to the map, returning `true` if a new interval was added or `false` if the interval already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(lo, hi)` removes the entry with key `[lo, hi]` from the map, returning `true` if the removal was successful or `false` if the interval was not found.
- `find_key(lo, hi)` returns a pointer to a const `std::pair` representing the key of some interval in the map which intersects with `[lo, hi]`, or `nullptr` if no such entry was found.
- `find_value(lo, hi)` returns a pointer to a const value of some entry in the map with a key that intersects with `[lo, hi]`, or `nullptr` if no such entry was found.
- `find_all(lo, hi, f)` calls the function `f(lo, hi, v)` on each entry in the map that overlaps with `[lo, hi]`, in lexicographically ascending order of intervals.

- `entries()` returns all intervals and their values in lexicographically ascending order of intervals.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `empty()`.
- $O(\log n)$ on average per call to `insert()`, `erase()`, `find_key()`, and `find_value()`, where n is the number of intervals currently in the set.
- $O(\log n + m)$ on average per call to `find_all()`, where m is the number of intersecting intervals that are reported.
- $O(n)$ per call to `entries()`.

Space Complexity:

- $O(n)$ for storage of the map elements.
- $O(1)$ auxiliary for `size()` and `empty()`.
- $O(\log n)$ auxiliary stack space on average for all other operations and destruction, and $O(n)$ in the worst case.
- $O(n)$ for the vector returned by `entries()`.

```
#include <cassert>
#include <cstdlib>
#include <tuple>
#include <utility>
#include <vector>

template<typename K, typename V>
class IntervalTreap {
    using Interval = std::pair<K, K>;

    struct Node {
        static uint32_t rand32() {
            static uint32_t x = 123456789;
            x ^= x << 13;
            x ^= x >> 17;
            x ^= x << 5;
            return x;
        }

        Interval interval;
        V value;
        K max;
        uint32_t priority;
        Node *left, *right;

        Node(const Interval &i, const V &v)
            : interval(i), value(v), max(i.second), priority(rand32()), left(nullptr), right(nullptr) {}

        void update() {
            max = interval.second;
            if (left != nullptr && max < left->max) {
                max = left->max;
            }
            if (right != nullptr && max < right->max) {
                max = right->max;
            }
        }
    };

    Node *root;

    int num_nodes;

    static void rotate_left(Node *&n) {
        Node *tmp = n;
        n = n->right;
        tmp->right = n->left;
        n->left = tmp;
        tmp->update();
    }
};
```

```

}

static void rotate_right(Node *&n) {
    Node *tmp = n;
    n = n->left;
    tmp->left = n->right;
    n->right = tmp;
    tmp->update();
}

bool insert(Node *&n, const Interval &i, const V &v) {
    if (n == nullptr) {
        n = new Node(i, v);
        num_nodes++;
        return true;
    }
    if (i < n->interval && insert(n->left, i, v)) {
        if (n->left->priority < n->priority) {
            rotate_right(n);
        }
        n->update();
        return true;
    }
    if (n->interval < i && insert(n->right, i, v)) {
        if (n->right->priority < n->priority) {
            rotate_left(n);
        }
        n->update();
        return true;
    }
    return false;
}

bool erase(Node *&n, const Interval &i) {
    if (n == nullptr) {
        return false;
    }
    if (i < n->interval) {
        bool res = erase(n->left, i);
        if (res) {
            n->update();
        }
        return res;
    }
    if (n->interval < i) {
        bool res = erase(n->right, i);
        if (res) {
            n->update();
        }
        return res;
    }
    if (n->left != nullptr && n->right != nullptr) {
        bool res;
        if (n->left->priority < n->right->priority) {
            rotate_right(n);
            res = erase(n->right, i);
        } else {
            rotate_left(n);
            res = erase(n->left, i);
        }
        n->update();
        return res;
    }
    Node *tmp = (n->left != nullptr) ? n->left : n->right;
    delete n;
    n = tmp;
    num_nodes--;
    return true;
}

```

```

}

static Node *find_any(Node *n, const Interval &i) {
    if (n == nullptr) {
        return nullptr;
    }
    if (!(i.second < n->interval.first) && !(n->interval.second < i.first)) {
        return n;
    }
    if (n->left != nullptr && !(n->left->max < i.first)) {
        return find_any(n->left, i);
    }
    return find_any(n->right, i);
}

template<typename Fn>
static void find_all(Node *n, const Interval &i, Fn &f) {
    if (n == nullptr || n->max < i.first) {
        return;
    }
    find_all(n->left, i, f);
    if (!(i.second < n->interval.first) && !(n->interval.second < i.first)) {
        f(n->interval.first, n->interval.second, n->value);
    }
    if (!(i.second < n->interval.first)) {
        find_all(n->right, i, f);
    }
}

static void collect_entries(Node *n, std::vector<std::tuple<K, K, V>> &res) {
    if (n != nullptr) {
        collect_entries(n->left, res);
        res.emplace_back(n->interval.first, n->interval.second, n->value);
        collect_entries(n->right, res);
    }
}

static void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
    IntervalTreap() : root(nullptr), num_nodes(0) {}

    ~IntervalTreap() { clean_up(root); }
    IntervalTreap(const IntervalTreap &) = delete;
    IntervalTreap &operator=(const IntervalTreap &) = delete;
    int size() const { return num_nodes; }
    bool empty() const { return root == nullptr; }

    bool insert(const K &lo, const K &hi, const V &v) {
        assert(!(hi < lo));
        return insert(root, Interval{lo, hi}, v);
    }

    bool erase(const K &lo, const K &hi) {
        assert(!(hi < lo));
        return erase(root, Interval{lo, hi});
    }

    const Interval *find_key(const K &lo, const K &hi) const {
        assert(!(hi < lo));
        Node *n = find_any(root, Interval{lo, hi});
        return (n == nullptr) ? nullptr : &(n->interval);
    }

```

```

}

const V *find_value(const K &lo, const K &hi) const {
    assert(!(hi < lo));
    Node *n = find_any(root, Interval{lo, hi});
    return (n == nullptr) ? nullptr : &(n->value);
}

template<typename Fn>
void find_all(const K &lo, const K &hi, Fn f) const {
    assert(!(hi < lo));
    find_all(root, Interval{lo, hi}, f);
}

std::vector<std::tuple<K, K, V>> entries() const {
    std::vector<std::tuple<K, K, V>> res;
    res.reserve(num_nodes);
    collect_entries(root, res);
    return res;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    IntervalTreap<int, char> t;
    t.insert(15, 20, 'a');
    t.insert(10, 30, 'b');
    t.insert(17, 19, 'c');
    t.insert(5, 20, 'd');
    t.insert(12, 15, 'e');
    t.insert(10, 40, 'f');
    assert(t.size() == 6);
    assert(!t.insert(5, 20, 'x'));
    t.erase(17, 19);
    assert(t.size() == 5);
    assert(*t.find_key(3, 9) == (pair<int, int>{5, 20}));
    assert(*t.find_value(3, 9) == 'd');
    vector<pair<int, int>> intersections;
    t.find_all(16, 20, [&](int lo, int hi, char) { intersections.emplace_back(lo, hi); });
    assert((intersections == vector<pair<int, int>>{{5, 20}, {10, 30}, {10, 40}, {15, 20}}));
    assert(
        (t.entries() == vector<tuple<int, int, char>>{
            {5, 20, 'd'}, {10, 30, 'b'}, {10, 40, 'f'}, {12, 15, 'e'}, {15, 20, 'a'}
        })
    );
    return 0;
}

```

2.3.7 Hash Table (Chaining)

Maintain an unordered map: a collection of key-value pairs in which each key appears at most once. Keys are hashed into buckets, and this implementation resolves collisions by chaining the entries in each bucket into a linked list. The hash and key-equality policies follow the conventions of `std::unordered_map`.

Compared with the open-addressing version (2.3.8), chaining keeps each entry at a stable address: a rehash never moves existing nodes, so pointers from `find()` and references from `operator[]` stay valid, and load factors above 1 are tolerated gracefully. The costs are a separate allocation per entry and cache-unfriendly pointer chasing during traversal.

- `ChainingHashMap<K, V>(buckets = 128)` constructs an empty map with the positive number of buckets given by `buckets`, using `std::hash<K>` and `std::equal_to<K>`.
- `ChainingHashMap<K, V, Hash, KeyEqual>(buckets, hash, equal)` instead stores the supplied hash and equality policies.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to the value associated with key `k`, or `nullptr` if the key was not found.
- `operator[k]` returns a reference to key `k`'s associated value (which may be modified), or if necessary, inserts and returns a new entry with the default constructed value if key `k` was not originally found.
- `entries()` returns all key-value entries in no guaranteed order.

This is an educational implementation; in practice prefer one of the standard options:

- `std::unordered_map` is the portable standard-library hash map. It is also chaining-based, so this class is essentially a transparent version of it. Pass a custom hash (see 3.2.1) to harden it against adversarial inputs, as its default integer hash is effectively the identity.
- `__gnu_pbds::gp_hash_table` (from GCC's policy-based library) is an open-addressing table that is typically several times faster than `std::unordered_map`, but is a non-portable GNU extension. Prefer it on GCC-only judges when hashing is the bottleneck; see 8.6 for `HashMap`/`HashSet` wrappers with a randomized integer hash.

Time Complexity:

- $O(b)$ per call to the constructor, where b is the initial number of buckets.
- $O(1)$ per call to `size()` and `empty()`.
- $O(1)$ expected amortized per call to `insert()`, `erase()`, `find()`, and `operator[]`, and $O(n)$ in the collision-heavy worst case.
- $O(n + b)$ per call to `entries()`, where n is the number of entries and b is the number of buckets.

Space Complexity:

- $O(n + b)$ for storage of the map elements and buckets.
- $O(n + b)$ auxiliary during a rehash.
- $O(n)$ for the vector returned by `entries()`.
- $O(1)$ auxiliary for all other operations.

```
#include <cassert>
#include <cstddef>
#include <cstdint>
#include <functional>
#include <list>
#include <utility>
#include <vector>

template<typename K, typename V, typename Hash = std::hash<K>, typename KeyEqual = std::equal_to<K>>
class ChainingHashMap {
    struct HashNode {
        K key;
        V value;
        HashNode(const K &k, const V &v) : key(k), value(v) {}
    };

    std::vector<std::list<HashNode>> table;
    int table_size, num_entries;
    Hash hash;
    KeyEqual equal;

    int bucket(const K &k) const {
```

```

    return static_cast<int>(hash(k) % static_cast<std::size_t>(table_size));
}

void grow_and_rehash() {
    auto old = std::move(table);
    table_size = 2 * table_size;
    table.assign(table_size, std::list<HashNode>());
    for (auto &chain : old) {
        while (!chain.empty()) {
            auto it = chain.begin();
            int i = bucket(it->key);
            table[i].splice(table[i].end(), chain, it);
        }
    }
}

public:
explicit ChainingHashMap(int buckets = 128, Hash hash = Hash(), KeyEqual equal = KeyEqual())
    : table_size(buckets), num_entries(0), hash(std::move(hash)), equal(std::move(equal)) {
    assert(buckets > 0);
    table.resize(buckets);
}

int size() const { return num_entries; }
bool empty() const { return num_entries == 0; }

bool insert(const K &k, const V &v) {
    if (find(k) != nullptr) {
        return false;
    }
    if (num_entries >= table_size) {
        grow_and_rehash();
    }
    int i = bucket(k);
    table[i].emplace_back(k, v);
    num_entries++;
    return true;
}

bool erase(const K &k) {
    int i = bucket(k);
    auto it = table[i].begin();
    while (it != table[i].end() && !equal(it->key, k)) {
        ++it;
    }
    if (it == table[i].end()) {
        return false;
    }
    table[i].erase(it);
    num_entries--;
    return true;
}

V *find(const K &k) {
    int i = bucket(k);
    auto it = table[i].begin();
    while (it != table[i].end() && !equal(it->key, k)) {
        ++it;
    }
    if (it == table[i].end()) {
        return nullptr;
    }
    return &(it->value);
}

V &operator[](const K &k) {
    if (V *ptr = find(k); ptr != nullptr) {
        return *ptr;
    }

```

```

    }
    if (num_entries >= table_size) {
        grow_and_rehash();
    }
    int i = bucket(k);
    table[i].emplace_back(k, V());
    num_entries++;
    return table[i].back().value;
}

std::vector<std::pair<K, V>> entries() const {
    std::vector<std::pair<K, V>> res;
    res.reserve(num_entries);
    for (int i = 0; i < table_size; i++) {
        for (const auto &[k, v] : table[i]) {
            res.emplace_back(k, v);
        }
    }
    return res;
}
};

```

Example Usage

```

#include <algorithm>
#include <cassert>
#include <chrono>
#include <string>
using namespace std;

// Example key hashers. For more hash algorithms and overloads, see 3.2.1. To defend against
// adversarially crafted collisions in open-hacking environments, a random seed is added to input
// keys before mixing (as 3.2.1's IntHasher does).
struct Hasher {
    inline static const uint64_t RAND_SEED =
        std::chrono::steady_clock::now().time_since_epoch().count();

    // Signed -> unsigned delegates.
    uint32_t operator()(int k) const { return Hasher()(static_cast<uint32_t>(k)); }
    uint32_t operator()(int64_t k) const { return Hasher()(static_cast<uint64_t>(k)); }

    // Knuth's multiplicative method. Fast, but affine: an additive RAND_SEED only shifts all buckets
    // uniformly and won't stop crafted collisions (unlike the non-linear hashers below). To harden
    // it, randomize the odd multiplier and take the high bits instead.
    uint32_t operator()(uint32_t k) const {
        return k * 2654435761u; // Or just return k.
    }

    // SplitMix64 mixer (see 3.2.1's mix64).
    uint32_t operator()(uint64_t k) const {
        k += RAND_SEED;
        k = (k ^ (k >> 30)) * 0xbf58476d1ce4e5b9ULL;
        k = (k ^ (k >> 27)) * 0x94d049bb133111ebULL;
        return static_cast<uint32_t>(k ^ (k >> 31));
    }

    // Jenkins's one-at-a-time hash.
    uint32_t operator()(const std::string &k) const {
        uint32_t hash = RAND_SEED;
        for (char c : k) {
            hash += ((hash + static_cast<unsigned char>(c)) << 10);
            hash ^= (hash >> 6);
        }
        hash += (hash << 3);
        hash ^= (hash >> 11);
        return hash + (hash << 15);
    }
};

```

```

    }
};

struct AbsHash {
    int salt;
    explicit AbsHash(int salt = 0) : salt(salt) {}
    size_t operator()(int x) const { return (x < 0 ? -static_cast<int64_t>(x) : x) + salt; }
};

struct AbsEqual {
    bool operator()(int a, int b) const { return AbsHash()(a) == AbsHash()(b); }
};

int main() {
    ChainingHashMap<string, char, Hasher> m;
    assert(m.empty());
    m["foo"] = 'a';
    assert(m.insert("bar", 'b'));
    assert(!m.insert("bar", 'z'));
    assert(!m.empty() && m.size() == 2);
    assert(m["foo"] == 'a');
    assert(m["bar"] == 'b');
    assert(m.find("foo") != nullptr && *m.find("foo") == 'a');
    assert(m.find("qux") == nullptr);
    assert(m["baz"] == '\0');
    m["baz"] = 'c';

    string vals;
    for (const auto &[k, v] : m.entries()) {
        vals += v;
    }
    sort(vals.begin(), vals.end());
    assert(vals == "abc");
    assert(m.erase("foo"));
    assert(m.size() == 2);
    assert(m["foo"] == '\0');
    assert(m.size() == 3);

    ChainingHashMap<string, char, Hasher> stable(1);
    stable["x"] = 'x';
    char *ptr = stable.find("x");
    stable["y"] = 'y'; // Rehashes without relocating the existing list node.
    assert(ptr == stable.find("x") && *ptr == 'x');

    ChainingHashMap<string, int> defaults;
    defaults["answer"] = 42;
    assert(*defaults.find("answer") == 42);

    ChainingHashMap<int, char, AbsHash, AbsEqual> absolute(128, AbsHash(7), AbsEqual());
    assert(absolute.insert(-2, 'x'));
    assert(!absolute.insert(2, 'y') && *absolute.find(2) == 'x');
    return 0;
}

```

2.3.8 Hash Table (Open Addressing)

Maintain an unordered map: a collection of key-value pairs in which each key appears at most once. Keys are hashed into table slots, and this implementation resolves collisions using open addressing with linear probing. Deletions leave tombstones behind so that probe chains remain searchable. The hash and key-equality policies follow the conventions of `std::unordered_map`.

Linear probing checks slots i , $i + 1$, $i + 2$, ... (modulo the table size) until the desired key or an empty slot is found. Other probe-step schemes include quadratic probing and double hashing, but these require additional

table-sizing and coprimality conditions to guarantee correct coverage of the table. The `next_bucket()` helper includes a commented quadratic-probing alternative for power-of-two table sizes.

Compared with the chaining version (2.3.7), open addressing stores entries contiguously, so it is more cache-friendly and needs no per-entry allocation. The costs: deletions must leave tombstones, the load factor must stay well below 1, and, unlike chaining or `std::unordered_map`, a pointer from `find()` or reference from `operator[]` is invalidated as soon as a later insertion reshapes and relocates the entries.

- `ProbingHashMap<K, V>(buckets = 128)` constructs an empty map with the positive number of slots given by `buckets`, using `std::hash<K>` and `std::equal_to<K>`.
- `ProbingHashMap<K, V, Hash, KeyEqual>(buckets, hash, equal)` instead stores the supplied hash and equality policies.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to the value associated with key `k`, or `nullptr` if the key was not found.
- `operator[k]` returns a reference to key `k`'s associated value (which may be modified), or if necessary, inserts and returns a new entry with the default constructed value if key `k` was not originally found.
- `entries()` returns all key-value entries in no guaranteed order.

This is an educational implementation; in practice prefer one of the standard options:

- `std::unordered_map` is the portable standard-library hash map. Unlike this class, it is chaining-based rather than open-addressed. Pass a custom hash (see 3.2.1) to harden it against adversarial inputs, as its default integer hash is effectively the identity.
- `__gnu_pbds::gp_hash_table` (from GCC's policy-based library) is an open-addressing table that is typically several times faster than `std::unordered_map`, but is a non-portable GNU extension. Prefer it on GCC-only judges when hashing is the bottleneck; see 8.6 for `HashMap`/`HashSet` wrappers with a randomized integer hash.

Time Complexity:

- $O(b)$ per call to the constructor, where b is the initial number of slots.
- $O(1)$ per call to `size()` and `empty()`.
- $O(1)$ expected amortized per call to `insert()`, `erase()`, `find()`, and `operator[]`, and $O(n)$ in the collision-heavy worst case.
- $O(n + b)$ per call to `entries()`, where n is the number of entries and b is the number of slots.

Space Complexity:

- $O(n + b)$ for storage of the map elements and slots.
- $O(n + b)$ auxiliary during a rehash.
- $O(n)$ for the vector returned by `entries()`.
- $O(1)$ auxiliary for all other operations.

```
#include <cassert>
#include <cstddef>
#include <cstdint>
#include <functional>
#include <optional>
#include <utility>
#include <vector>

template<typename K, typename V, typename Hash = std::hash<K>, typename KeyEqual = std::equal_to<K>>
class ProbingHashMap {
    enum class State { EMPTY, OCCUPIED, DELETED };

    struct HashNode {
```

```

    K key;
    V value;
    HashNode(const K &k, const V &v) : key(k), value(v) {}
};

std::vector<std::optional<HashNode>> table;
std::vector<State> state;
int table_size, num_entries, num_tombstones;
Hash hash;
KeyEqual equal;

int bucket(const K &k) const {
    return static_cast<int>(hash(k) % static_cast<std::size_t>(table_size));
}

int next_bucket(int i, int probes) const {
    return (i + 1) % table_size; // Linear probing.

    // Quadratic probing for power-of-two table sizes (cumulative offsets 1, 3, 6, 10, ...):
    // return (i + probes + 1) & (table_size - 1);
}

void rehash(int new_size) {
    auto old_table = std::move(table);
    auto old_state = std::move(state);
    table_size = new_size;
    table.assign(table_size, std::nullopt);
    state.assign(table_size, State::EMPTY);
    num_entries = 0;
    num_tombstones = 0;
    for (int i = 0; i < static_cast<int>(old_table.size()); i++) {
        if (old_state[i] == State::OCCUPIED) {
            insert(old_table[i]->key, old_table[i]->value);
        }
    }
}

void grow_if_needed() {
    // Keep total non-empty slots below 50%. Tombstones count because they lengthen probe chains.
    if (2 * (num_entries + num_tombstones) >= table_size) {
        rehash(2 * num_entries >= table_size ? 2 * table_size : table_size);
    }
}

public:
explicit ProbingHashMap(int buckets = 128, Hash hash = Hash(), KeyEqual equal = KeyEqual())
    : table_size(buckets),
      num_entries(0),
      num_tombstones(0),
      hash(std::move(hash)),
      equal(std::move(equal)) {
    assert(buckets > 0);
    // Also require this when using the quadratic alternative in next_bucket():
    // assert((buckets & (buckets - 1)) == 0);
    table.resize(buckets);
    state.assign(buckets, State::EMPTY);
}

int size() const { return num_entries; }
bool empty() const { return num_entries == 0; }

bool insert(const K &k, const V &v) {
    if (2 * (num_entries + num_tombstones) >= table_size && find(k) != nullptr) {
        return false;
    }
    grow_if_needed();
    int first_deleted = table_size;
    int i = bucket(k);

```

```

for (int probes = 0; probes < table_size; probes++) {
    if (state[i] == State::OCCUPIED) {
        if (equal(table[i]->key, k)) {
            return false;
        }
    } else if (state[i] == State::DELETED) {
        if (first_deleted == table_size) {
            first_deleted = i;
        }
    } else {
        int dest = first_deleted == table_size ? i : first_deleted;
        table[dest].emplace(k, v);
        if (state[dest] == State::DELETED) {
            num_tombstones--;
        }
        state[dest] = State::OCCUPIED;
        num_entries++;
        return true;
    }
    i = next_bucket(i, probes);
}
// Should be rare because grow_if_needed keeps slack, but safe.
rehash(2 * table_size);
return insert(k, v);
}

bool erase(const K &k) {
    int i = bucket(k);
    for (int probes = 0; probes < table_size; probes++) {
        if (state[i] == State::EMPTY) {
            return false;
        }
        if (state[i] == State::OCCUPIED && equal(table[i]->key, k)) {
            table[i].reset();
            state[i] = State::DELETED;
            num_entries--;
            num_tombstones++;
            return true;
        }
        i = next_bucket(i, probes);
    }
    return false;
}

V *find(const K &k) {
    int i = bucket(k);
    for (int probes = 0; probes < table_size; probes++) {
        if (state[i] == State::EMPTY) {
            return nullptr;
        }
        if (state[i] == State::OCCUPIED && equal(table[i]->key, k)) {
            return &table[i]->value;
        }
        i = next_bucket(i, probes);
    }
    return nullptr;
}

V &operator[](const K &k) {
    if (V *ptr = find(k); ptr != nullptr) {
        return *ptr;
    }
    grow_if_needed();
    int first_deleted = table_size;
    int i = bucket(k);
    for (int probes = 0; probes < table_size; probes++) {
        if (state[i] == State::DELETED) {
            if (first_deleted == table_size) {

```

```

        first_deleted = i;
    }
} else if (state[i] == State::EMPTY) {
    int dest = first_deleted == table_size ? i : first_deleted;
    table[dest].emplace(k, V());
    if (state[dest] == State::DELETED) {
        num_tombstones--;
    }
    state[dest] = State::OCCUPIED;
    num_entries++;
    return table[dest]->value;
}
i = next_bucket(i, probes);
}
rehash(2 * table_size);
return (*this)[k];
}

std::vector<std::pair<K, V>> entries() const {
    std::vector<std::pair<K, V>> res;
    res.reserve(num_entries);
    for (int i = 0; i < table_size; i++) {
        if (state[i] == State::OCCUPIED) {
            res.emplace_back(table[i]->key, table[i]->value);
        }
    }
    return res;
}
};

```

Example Usage

```

#include <algorithm>
#include <cassert>
#include <chrono>
#include <string>
using namespace std;

// Example key hashers. For more hash algorithms and overloads, see 3.2.1. To defend against
// adversarially crafted collisions in open-hacking environments, a random seed is added to input
// keys before mixing (as 3.2.1's IntHasher does).
struct Hasher {
    inline static const uint64_t RAND_SEED =
        std::chrono::steady_clock::now().time_since_epoch().count();

    // Signed -> unsigned delegates.
    uint32_t operator()(int k) const { return Hasher()(static_cast<uint32_t>(k)); }
    uint32_t operator()(int64_t k) const { return Hasher()(static_cast<uint64_t>(k)); }

    // Knuth's multiplicative method. Fast, but affine: an additive RAND_SEED only shifts all buckets
    // uniformly and won't stop crafted collisions (unlike the non-linear hashers below). To harden
    // it, randomize the odd multiplier and take the high bits instead.
    uint32_t operator()(uint32_t k) const {
        return k * 2654435761u; // Or just return k.
    }

    // SplitMix64 mixer (see 3.2.1's mix64).
    uint32_t operator()(uint64_t k) const {
        k += RAND_SEED;
        k = (k ^ (k >> 30)) * 0xbf58476d1ce4e5b9ULL;
        k = (k ^ (k >> 27)) * 0x94d049bb133111ebULL;
        return static_cast<uint32_t>(k ^ (k >> 31));
    }

    // Jenkins's one-at-a-time hash.
    uint32_t operator()(const std::string &k) const {

```

```

    uint32_t hash = RAND_SEED;
    for (char c : k) {
        hash += ((hash + static_cast<unsigned char>(c)) << 10);
        hash ^= (hash >> 6);
    }
    hash += (hash << 3);
    hash ^= (hash >> 11);
    return hash + (hash << 15);
}
};

struct AbsHash {
    int salt;
    explicit AbsHash(int salt = 0) : salt(salt) {}
    size_t operator()(int x) const { return (x < 0 ? -static_cast<int64_t>(x) : x) + salt; }
};

struct AbsEqual {
    bool operator()(int a, int b) const { return AbsHash()(a) == AbsHash()(b); }
};

int main() {
    ProbingHashMap<string, char, Hasher> m;
    assert(m.empty());
    m["foo"] = 'a';
    assert(m.insert("bar", 'b'));
    assert(!m.insert("bar", 'z'));
    assert(!m.empty() && m.size() == 2);
    assert(m["foo"] == 'a');
    assert(m["bar"] == 'b');
    assert(m.find("foo") != nullptr && *m.find("foo") == 'a');
    assert(m.find("qux") == nullptr);
    assert(m["baz"] == '\0');
    m["baz"] = 'c';

    string vals;
    for (const auto &[k, v] : m.entries()) {
        vals += v;
    }
    sort(vals.begin(), vals.end());
    assert(vals == "abc");
    assert(m.erase("foo"));
    assert(m.size() == 2);
    assert(m["foo"] == '\0');
    assert(m.size() == 3);

    ProbingHashMap<string, char, Hasher> duplicate(1);
    duplicate.insert("x", 'x');
    char *ptr = duplicate.find("x");
    assert(!duplicate.insert("x", 'z'));
    assert(ptr == duplicate.find("x")); // A rejected duplicate does not trigger a rehash.

    ProbingHashMap<string, int> defaults;
    defaults["answer"] = 42;
    assert(*defaults.find("answer") == 42);

    ProbingHashMap<int, char, AbsHash, AbsEqual> absolute(128, AbsHash(7), AbsEqual());
    assert(absolute.insert(-2, 'x'));
    assert(!absolute.insert(2, 'y') && *absolute.find(2) == 'x');
    return 0;
}

```

2.3.9 Skip List

Maintain an ordered map, that is, an ordered collection of key-value pairs such that each possible key appears at most once in the collection. The comparator `comp` defines the key ordering: `comp(a, b)` is true when `a` precedes `b`. A skip list maintains a linked hierarchy of sorted subsequences with each successive subsequence skipping over fewer elements than the previous one. Each new node joins a number of levels decided by repeated coin flips, and a search starts at the sparsest level, moving forward until the next key would overshoot and then dropping down a level, which makes operations take $O(\log n)$ with high probability. The comparator defaults to `std::less<K>`; to customize the ordering, instantiate `SkipList<K, V, Compare>` and pass the comparator to the constructor.

- `SkipList<K, V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(k, v)` adds an entry with key `k` and value `v` to the map, returning `true` if a new entry was added or `false` if the key already exists (in which case the map is unchanged and the old value associated with the key is preserved).
- `erase(k)` removes the entry with key `k` from the map, returning `true` if the removal was successful or `false` if the key to be removed was not found.
- `find(k)` returns a pointer to the const value associated with key `k`, or `nullptr` if the key was not found.
- `operator[k]` returns a reference to key `k`'s associated value (which may be modified), or if necessary, inserts and returns a new entry with the default constructed value if key `k` was not originally found.
- `entries()` returns all key-value entries in comparator order.

The sentinel node requires `K` and `V` to be default constructible.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `empty()`.
- $O(\log n)$ on average per call to `insert()`, `erase()`, `find()`, and `operator[]`, where n is the number of entries currently in the map.
- $O(n)$ per call to `entries()`.

Space Complexity:

- $O(n)$ for storage of the map elements.
- $O(1)$ auxiliary for all operations because the maximum number of levels is fixed at 32.
- $O(n)$ for the vector returned by `entries()`.

```
#include <algorithm>
#include <functional>
#include <random>
#include <utility>
#include <vector>

template<typename K, typename V, typename Compare = std::less<K>>
class SkipList {
    static const int MAX_LEVELS = 32; // log2(max possible keys)

    struct Node {
        K key;
        V value;
        std::vector<Node *> next;

        Node(const K &k, const V &v, int levels) : key(k), value(v), next(levels, nullptr) {}
    } *head;

    int num_nodes;
    Compare comp;

    static int random_level() {
        static std::mt19937 rng(std::random_device{}());
        static std::uniform_int_distribution<int> coin(0, 1);
        int level = 1;
```

```

    while (coin(rng) && level < MAX_LEVELS) {
        level++;
    }
    return level;
}

static int node_level(const std::vector<Node *> &v) {
    int i = 0;
    while (i < static_cast<int>(v.size()) && v[i] != nullptr) {
        i++;
    }
    return std::max(1, i);
}

Node *find_node(const K &k) const {
    Node *n = head;
    for (int i = node_level(n->next); i-- > 0;) {
        while (n->next[i] != nullptr && comp(n->next[i]->key, k)) {
            n = n->next[i];
        }
    }
    n = n->next[0];
    return (n != nullptr && !comp(k, n->key) && !comp(n->key, k)) ? n : nullptr;
}

public:
explicit SkipList(Compare comp = Compare())
    : head(new Node(K(), V(), MAX_LEVELS)), num_nodes(0), comp(std::move(comp)) {
    for (auto &ptr : head->next) {
        ptr = nullptr;
    }
}

~SkipList() {
    Node *n = head;
    while (n != nullptr) {
        Node *next = n->next[0];
        delete n;
        n = next;
    }
}

SkipList(const SkipList &) = delete;
SkipList &operator=(const SkipList &) = delete;
int size() const { return num_nodes; }
bool empty() const { return num_nodes == 0; }

bool insert(const K &k, const V &v) {
    std::vector<Node *> update(head->next);
    int curr_level = node_level(update);
    Node *n = head;
    for (int i = curr_level; i-- > 0;) {
        while (n->next[i] != nullptr && comp(n->next[i]->key, k)) {
            n = n->next[i];
        }
        update[i] = n;
    }
    n = n->next[0];
    if (n != nullptr && !comp(k, n->key) && !comp(n->key, k)) {
        return false;
    }
    int new_level = random_level();
    if (new_level > curr_level) {
        for (int i = curr_level; i < new_level; i++) {
            update[i] = head;
        }
    }
    n = new Node(k, v, new_level);

```

```

    for (int i = 0; i < new_level; i++) {
        n->next[i] = update[i]->next[i];
        update[i]->next[i] = n;
    }
    num_nodes++;
    return true;
}

bool erase(const K &k) {
    std::vector<Node *> update(head->next);
    Node *n = head;
    for (int i = node_level(update); i-- > 0;) {
        while (n->next[i] != nullptr && comp(n->next[i]->key, k)) {
            n = n->next[i];
        }
        update[i] = n;
    }
    n = n->next[0];
    if (n != nullptr && !comp(k, n->key) && !comp(n->key, k)) {
        for (int i = 0; i < static_cast<int>(n->next.size()); i++) {
            update[i]->next[i] = n->next[i];
        }
        delete n;
        num_nodes--;
        return true;
    }
    return false;
}

const V *find(const K &k) const {
    Node *n = find_node(k);
    return n == nullptr ? nullptr : &(n->value);
}

V &operator[](const K &k) {
    Node *n = find_node(k);
    if (n != nullptr) {
        return n->value;
    }
    insert(k, V());
    return find_node(k)->value;
}

std::vector<std::pair<K, V>> entries() const {
    std::vector<std::pair<K, V>> res;
    res.reserve(num_nodes);
    Node *n = head->next[0];
    while (n != nullptr) {
        res.emplace_back(n->key, n->value);
        n = n->next[0];
    }
    return res;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    SkipList<int, char> l;
    assert(l.empty());
    l.insert(2, 'b');
    l.insert(1, 'a');
    l.insert(3, 'c');
}

```



```

l.insert(5, 'e');
assert(l.insert(4, 'd'));
assert(!l.empty() && l.size() == 5);
assert(*l.find(4) == 'd');
assert(!l.insert(4, 'd'));
assert(l.size() == 5);
assert(
    (l.entries() == vector<pair<int, char>>{{1, 'a'}, {2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}})
);
assert(l.erase(1));
assert(!l.erase(1));
assert(l.find(1) == nullptr);
assert(l.size() == 4);
assert((l.entries() == vector<pair<int, char>>{{2, 'b'}, {3, 'c'}, {4, 'd'}, {5, 'e'}}));

Skiplist<int, char, greater<int>> descending;
for (int key : {2, 1, 3}) {
    descending.insert(key, '0' + key);
}
assert((descending.entries() == vector<pair<int, char>>{{3, '3'}, {2, '2'}, {1, '1'}}));
assert(*descending.find(2) == '2');
assert(descending.erase(2) && descending.find(2) == nullptr);
return 0;
}

```

2.4 Range Queries in One Dimension

2.4.1 Sparse Table

Given a static array, precompute a table that answers range minimum queries in constant time. More generally, the sparse table works for any range query whose combining operation is both associative and idempotent (i.e. `combine(x, x)` must equal `x`). The canonical examples are “min”, “max”, “gcd”, and the bitwise “and” and “or”.

For each level j , the entry `dp[j][i]` holds the result of `combine()` over the half-open range $[i, i + 2^j)$, built by combining two ranges of length 2^{j-1} . A query over $[l_0, h_i]$ is answered by combining the two overlapping ranges of length 2^j that are anchored at l_0 and at h_i , where 2^j is the largest power of two not exceeding the length of the range. Because these two ranges overlap whenever the query length is not a power of two, the elements in the overlap are folded in twice; this is harmless only when `combine()` is idempotent. To support a non-idempotent but associative operation such as “sum” or “XOR” in constant time, use the disjoint sparse table in the next section instead, whose ranges meet at a center and never overlap.

The query operation is defined by an associative and idempotent function `combine(a, b)`. The default code below returns the “min” of the range; for “gcd”, `combine(a, b)` should return `gcd(a, b)`.

- `SparseTable<T>(l0, hi)` builds the table over the half-open iterator range $[l_0, h_i]$.
- `size()` returns the size of the array.
- `query(l0, hi)` returns `combine()` applied to all indices in $[l_0, h_i]$.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the size of the array.
- $O(1)$ per call to `size()` and `query()`.

Space Complexity:

- $O(n \log n)$ for storage of the table.
- $O(1)$ auxiliary for `query()`.

```

#include <algorithm>
#include <cassert>

```

```

#include <utility>
#include <vector>

template<typename T>
class SparseTable {
    static T combine(const T &a, const T &b) { return std::min(a, b); }

    int len;
    std::vector<int> log2; // Can also use 31 - __builtin_clz(n) instead of precomputing log2[n].
    std::vector<std::vector<T>> dp; // dp[j][i] = combine() over [i, i + 2^j).

public:
    template<typename It>
    SparseTable(It lo, It hi) {
        std::vector<T> a(lo, hi);
        len = static_cast<int>(a.size());
        log2.assign(len + 1, 0);
        for (int i = 2; i <= len; i++) {
            log2[i] = log2[i >> 1] + 1;
        }
        dp.assign(log2[len] + 1, std::vector<T>(len));
        dp[0] = std::move(a);
        for (int j = 1; (1 << j) <= len; j++) {
            for (int i = 0; i + (1 << j) <= len; i++) {
                dp[j][i] = combine(dp[j - 1][i], dp[j - 1][i + (1 << (j - 1))]);
            }
        }
    }

    int size() const { return len; }

    T query(int lo, int hi) const {
        assert(0 <= lo && lo <= hi && hi < len);
        int j = log2[hi - lo + 1];
        return combine(dp[j][lo], dp[j][hi - (1 << j) + 1]);
    }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10};
    SparseTable<int> t(a.begin(), a.end());
    assert(t.size() == 5);
    assert(t.query(0, 3) == -2);
    assert(t.query(2, 4) == 1);
    assert(t.query(3, 3) == 8);
    return 0;
}

```

2.4.2 Disjoint Sparse Table

Given a static array, precompute a table that answers range queries for any associative operation in constant time. Unlike the ordinary sparse table for range minimum queries, the disjoint sparse table does not require the operation to be idempotent, so it also handles sums, products, and matrix products where overlapping the two query halves would double-count.

The array is divided into blocks whose size doubles at each of the $O(\log n)$ levels. At the level whose block size is 2^{k+1} , every block is split at its center, and the table stores, for each position, the fold of the contiguous run from

that position up to (or down to) the center. To answer a query $[lo, hi]$, find the level at which lo and hi first fall on opposite sides of a center: this is the position of the highest set bit of $lo \sim hi$. At that level, lo lies in the block's left half and hi in its right half, so the answer is the stored suffix fold ending at the center combined with the stored prefix fold beginning at the center.

The query operation is defined by an associative function `combine(a, b)`. The default code below returns the “min” of the range; for “sum”, `combine(a, b)` should return `a + b`.

- `DisjointSparseTable<T>(lo, hi)` builds the table over the half-open iterator range $[lo, hi)$.
- `size()` returns the size of the array.
- `query(lo, hi)` returns `combine()` applied to all indices in $[lo, hi]$.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the size of the array.
- $O(1)$ per call to `size()` and `query()`.

Space Complexity:

- $O(n \log n)$ for storage of the table.
- $O(1)$ auxiliary for `query()`.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <vector>

template<typename T>
class DisjointSparseTable {
    static T combine(const T &a, const T &b) { return std::min(a, b); }

    int len;
    std::vector<T> a;
    std::vector<std::vector<T>> fold;

public:
    template<typename It>
    DisjointSparseTable(It lo, It hi) : a(lo, hi) {
        len = static_cast<int>(a.size());
        if (len == 0) {
            return;
        }
        int levels = 1;
        while ((1 << levels) < len) {
            levels++;
        }
        fold.assign(levels, std::vector<T>(len));
        for (int level = 0; level < levels; level++) {
            int range = 1 << (level + 1), half = 1 << level;
            for (int center = half; center < len + half; center += range) {
                int left = center - half, right = std::min(len, center + half);
                int top = std::min(center, len) - 1; // Suffix folds over [i, center).
                fold[level][top] = a[top];
                for (int i = top - 1; i >= left; i--) {
                    fold[level][i] = combine(a[i], fold[level][i + 1]);
                }
                if (center < len) { // Prefix folds over [center, i].
                    fold[level][center] = a[center];
                    for (int i = center + 1; i < right; i++) {
                        fold[level][i] = combine(fold[level][i - 1], a[i]);
                    }
                }
            }
        }
    }

    int size() const { return len; }
```

```

T query(int lo, int hi) const {
    assert(0 <= lo && lo <= hi && hi < len);
    if (lo == hi) {
        return a[lo];
    }
    int level = 31 - __builtin_clz(static_cast<uint32_t>(lo ^ hi));
    return combine(fold[level][lo], fold[level][hi]);
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10, -5, 3};
    DisjointSparseTable<int> t(a.begin(), a.end());
    assert(t.size() == 7);
    assert(t.query(0, 6) == -5);
    assert(t.query(1, 3) == -2);
    assert(t.query(4, 4) == 10);
    assert(t.query(2, 4) == 1);

    vector<int> empty;
    DisjointSparseTable<int> empty_table(empty.begin(), empty.end());
    assert(empty_table.size() == 0);
    return 0;
}

```

2.4.3 Square Root Decomposition

Maintain a fixed-size array while supporting point updates and contiguous range aggregate queries. Square root decomposition partitions the array into contiguous blocks of about $\sqrt{n}$ elements, each caching the aggregate of its block. A range query combines the cached aggregates of the whole blocks contained in the range with the individual elements of the partial blocks at either end, while a point update refreshes only the single block containing the changed index.

The query operation is defined by an associative aggregate function `combine(a, b)`. The default definition below assumes a numerical array type, supporting queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`.

The point update operation is defined by `apply_delta(v, d)`, which returns the new value at a single updated index. The default definition below supports updates that “set” the chosen array index to a new value. Another possible update operation is “increment”, in which case `apply_delta(v, d)` should return `v + d`.

- `SqrtDecomposition<T>(n, v = T())` constructs an array of size `n` with indices `[0, n)`, and all values initialized to `v`.
- `SqrtDecomposition<T>(lo, hi)` constructs an array from the half-open random-access iterator range `[lo, hi)`.
- `size()` returns the size of the array.
- `at(i)` returns the value at index `i`.
- `query(lo, hi)` returns the result of `combine()` applied to all indices in `[lo, hi]`.
- `update(i, d)` assigns the value `v` at index `i` to `apply_delta(v, d)`.

The supported operations are identical to those of the point-update segment tree in this section.

Time Complexity:

- $O(n)$ per call to both constructors, where n is the size of the array.

- $O(1)$ per call to `size()` and `at()`.
- $O(\sqrt{n})$ per call to `query()` and `update()`.

Space Complexity:

- $O(n)$ for storage of the array elements.
- $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cassert>
#include <cmath>
#include <vector>

template<typename T>
class SqrtDecomposition {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d) { return d; }

    int len, blocklen;
    std::vector<T> value, block;

    void init() {
        blocklen = std::max(1, static_cast<int>(std::sqrt(len)));
        int nblocks = (len + blocklen - 1) / blocklen;
        for (int i = 0; i < nblocks; i++) {
            T blockval = value[i * blocklen];
            int blockhi = std::min(len, (i + 1) * blocklen);
            for (int j = i * blocklen + 1; j < blockhi; j++) {
                blockval = combine(blockval, value[j]);
            }
            block.push_back(blockval);
        }
    }

public:
    explicit SqrtDecomposition(int n, const T &v = T()) : len(n), value(n, v) { init(); }

    template<typename It>
    SqrtDecomposition(It lo, It hi) : len(hi - lo), value(lo, hi) {
        init();
    }

    int size() const { return len; }

    T at(int i) const {
        assert(0 <= i && i < len);
        return value[i];
    }

    T query(int lo, int hi) const {
        assert(0 <= lo && lo <= hi && hi < len);
        int blocklo = (lo + blocklen - 1) / blocklen, blockhi = (hi + 1) / blocklen - 1;
        if (blocklo > blockhi) {
            T res = value[lo];
            for (int i = lo + 1; i <= hi; i++) {
                res = combine(res, value[i]);
            }
            return res;
        }
        T res = block[blocklo];
        for (int i = blocklo + 1; i <= blockhi; i++) {
            res = combine(res, block[i]);
        }
        for (int i = lo; i < blocklo * blocklen; i++) {
            res = combine(res, value[i]);
        }
        for (int i = (blockhi + 1) * blocklen; i <= hi; i++) {
            res = combine(res, value[i]);
        }
    }
};
```

```

    }
    return res;
}

void update(int i, const T &d) {
    assert(0 <= i && i < len);
    value[i] = apply_delta(value[i], d);
    int b = i / blocklen;
    int blockhi = std::min(len, (b + 1) * blocklen);
    block[b] = value[b * blocklen];
    for (int j = b * blocklen + 1; j < blockhi; j++) {
        block[b] = combine(block[b], value[j]);
    }
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10};
    SqrtDecomposition<int> sd(a.begin(), a.end());
    sd.update(2, 4);
    vector<int> expected{6, -2, 4, 8, 10};
    for (int i = 0; i < sd.size(); i++) {
        assert(sd.at(i) == expected[i]);
    }
    assert(sd.query(0, 3) == -2);
    return 0;
}

```

2.4.4 Mo's Algorithm

Mo's algorithm answers a batch of range queries offline by reordering them so that the current window $[\text{cur}_l, \text{cur}_r]$ can be transformed into the next query's range with as few single-element insertions and deletions as possible. It applies whenever there is no fast direct formula, but extending or shrinking the window by one element updates the answer cheaply, as with the number of distinct values in a range or the sum of squares of element frequencies.

The trick is the query order. Sort queries by the block of their left endpoint, where each block spans about $n/\sqrt{q}$ indices, breaking ties by right endpoint. The left pointer then travels $O(n/\sqrt{q})$ per query, for $O(n\sqrt{q})$ moves in all, while within each block the right pointer sweeps monotonically across the array, contributing $O(n)$ per block over the $\sqrt{q}$ blocks, again $O(n\sqrt{q})$. So the pointers move $O(n\sqrt{q})$ times in total. Sorting the right endpoint in alternating directions between consecutive blocks (a “boustrophedon” order) avoids resetting the right pointer at every block boundary.

- `mos_algorithm(n, queries, add, remove, current)` answers the given inclusive 0-based ranges over an array of size `n`. The caller supplies three callbacks that maintain an arbitrary window state: `add(i)` and `remove(i)` insert or delete index `i` from the current window, and `current()` returns the answer for the window in its present state. The result is a vector of answers in the original query order, with element type matching the return type of `current()`.

Time Complexity: $O(n\sqrt{q})$ calls to `add()` and `remove()` in total, plus $O(q \log q)$ to sort, where n is the array size and q is the number of queries.

Space Complexity: $O(q)$ auxiliary and $O(q)$ for the returned answers, excluding whatever window state the callbacks maintain.

```

#include <algorithm>
#include <cassert>
#include <cmath>
#include <numeric>
#include <utility>
#include <vector>

template<typename AddFn, typename RemoveFn, typename CurrentFn>
std::vector<decltype(std::declval<CurrentFn>())> mos_algorithm(
    int n, const std::vector<std::pair<int, int>> &queries, AddFn add, RemoveFn remove,
    CurrentFn current
) {
    assert(n >= 0);
    for (auto [lo, hi] : queries) {
        assert(0 <= lo && lo <= hi && hi < n);
    }
    int q = static_cast<int>(queries.size());
    int block = std::max(1, static_cast<int>(n / std::max(1.0, std::sqrt(static_cast<double>(q)))));
    std::vector<int> order(q);
    std::iota(order.begin(), order.end(), 0);
    std::sort(order.begin(), order.end(), [&](int x, int y) {
        int bx = queries[x].first / block, by = queries[y].first / block;
        if (bx != by) {
            return bx < by;
        }
        return (bx & 1) ? queries[x].second > queries[y].second : queries[x].second < queries[y].second;
    });
    std::vector<decltype(current())> res(q);
    int cur_l = 0, cur_r = -1; // Window [cur_l, cur_r] is initially empty.
    for (int idx : order) {
        int l = queries[idx].first, r = queries[idx].second;
        while (cur_r < r) add(++cur_r);
        while (cur_l > l) add(--cur_l);
        while (cur_r > r) remove(cur_r--);
        while (cur_l < l) remove(cur_l++);
        res[idx] = current();
    }
    return res;
}

```

Example Usage

```

using namespace std;

int main() {
    // Answer "number of distinct values in [l, r]" for several ranges.
    vector<int> a{1, 1, 2, 1, 3, 2, 3};
    vector<int> freq(4, 0);
    int distinct = 0;
    auto add = [&](int i) {
        if (freq[a[i]]++ == 0) distinct++;
    };
    auto remove = [&](int i) {
        if (--freq[a[i]] == 0) distinct--;
    };
    auto current = [&]() { return distinct; };

    int n = static_cast<int>(a.size());
    vector<pair<int, int>> queries{{0, 6}, {0, 2}, {3, 5}, {1, 1}, {2, 4}};
    vector<int> ans = mos_algorithm(n, queries, add, remove, current);

    assert((ans == vector<int>{3, 2, 3, 1, 3}));
    return 0;
}

```

2.4.5 Segment Tree (Point Update)

Maintain a fixed-size array while supporting point updates and range aggregate queries. A segment tree stores one aggregate value for each recursively split interval, so an update only rebuilds the $O(\log n)$ intervals on the path from one leaf to the root.

The query operation is defined by an associative aggregate function `combine(a, b)`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`.

The point update operation is defined by `apply_delta(v, d)`, which returns the new value at a single updated index. The default definition below supports updates that “set” the chosen array index to a new value. Another possible update operation is “increment”, in which case `apply_delta(v, d)` should return `v + d`.

- `SegTree<T>(n, v = T())` constructs an array of size `n` with indices $[0, n)$, and all values initialized to `v`.
- `SegTree<T>(lo, hi)` constructs an array from the half-open random-access iterator range $[lo, hi)$.
- `size()` returns the size of the array.
- `at(i)` returns the value at index `i`.
- `query(lo, hi)` returns the result of `combine()` applied to all indices in $[lo, hi]$. If `lo == hi`, then the single specified value is returned.
- `update(i, d)` assigns the value `v` at index `i` to `apply_delta(v, d)`.
- `max_right(lo, pred)` returns the largest boundary `hi` such that the aggregate over the half-open range $[lo, hi)$ satisfies the monotonic predicate `pred()`. It returns `size()` if `pred()` remains true to the end.
- `min_left(hi, pred)` returns the smallest boundary `lo` such that the aggregate over the half-open range $[lo, hi)$ satisfies the monotonic predicate `pred()`. It returns 0 if `pred()` remains true to the beginning.

For the boundary-search functions, `pred()` takes aggregate `T` values of candidate ranges. As a range grows, `pred()` may change from true to false but never back to true; the empty range is considered valid. E.g. for `combine = min`, use `pred(mn) = (mn > x)` to find the first value $\leq x$, or for `combine = sum`, use `pred(sum) = (sum <= x)` with nonnegative values to find the longest range within the limit `x`.

Time Complexity:

- $O(n)$ per call to both constructors, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log n)$ per call to `at()`, `query()`, `update()`, `max_right()`, and `min_left()`.

Space Complexity:

- $O(n)$ for storage of the array elements.
- $O(\log n)$ auxiliary stack space for `query()`, `update()`, `max_right()`, and `min_left()`.
- $O(1)$ auxiliary for `size()`.

```
#include <algorithm>
#include <cassert>
#include <optional>
#include <vector>

template<typename T>
class SegTree {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d) { return d; }

    int len;
    std::vector<T> value;

    template<typename Gen>
    void build(int i, int lo, int hi, const Gen &gen) {
        if (lo == hi) {
            value[i] = gen(lo);
            return;
        }
        int mid = lo + (hi - lo) / 2;
```



```

    build(i * 2 + 1, lo, mid, gen);
    build(i * 2 + 2, mid + 1, hi, gen);
    value[i] = combine(value[i * 2 + 1], value[i * 2 + 2]);
}

T query(int i, int lo, int hi, int tgt_lo, int tgt_hi) const {
    if (lo == tgt_lo && hi == tgt_hi) {
        return value[i];
    }
    int mid = lo + (hi - lo) / 2;
    if (tgt_lo <= mid && mid < tgt_hi) {
        return combine(
            query(i * 2 + 1, lo, mid, tgt_lo, std::min(tgt_hi, mid)),
            query(i * 2 + 2, mid + 1, hi, std::max(tgt_lo, mid + 1), tgt_hi)
        );
    }
    if (tgt_lo <= mid) {
        return query(i * 2 + 1, lo, mid, tgt_lo, std::min(tgt_hi, mid));
    }
    return query(i * 2 + 2, mid + 1, hi, std::max(tgt_lo, mid + 1), tgt_hi);
}

void update(int i, int lo, int hi, int target, const T &d) {
    if (lo == hi) {
        value[i] = apply_delta(value[i], d);
        return;
    }
    int mid = lo + (hi - lo) / 2;
    if (target <= mid) {
        update(i * 2 + 1, lo, mid, target, d);
    } else {
        update(i * 2 + 2, mid + 1, hi, target, d);
    }
    value[i] = combine(value[i * 2 + 1], value[i * 2 + 2]);
}

template<typename Pred>
int max_right(int i, int lo, int hi, int tgt_lo, const Pred &pred, std::optional<T> &acc) const {
    if (hi < tgt_lo) {
        return -1;
    }
    if (tgt_lo <= lo) {
        T next = acc ? combine(*acc, value[i]) : value[i];
        if (pred(next)) {
            acc = next;
            return -1;
        }
        if (lo == hi) {
            return lo;
        }
    }
    int mid = lo + (hi - lo) / 2;
    int res = max_right(i * 2 + 1, lo, mid, tgt_lo, pred, acc);
    return res != -1 ? res : max_right(i * 2 + 2, mid + 1, hi, tgt_lo, pred, acc);
}

template<typename Pred>
int min_left(int i, int lo, int hi, int tgt_hi, const Pred &pred, std::optional<T> &acc) const {
    if (tgt_hi <= lo) {
        return -1;
    }
    if (hi < tgt_hi) {
        T next = acc ? combine(value[i], *acc) : value[i];
        if (pred(next)) {
            acc = next;
            return -1;
        }
    }
    if (lo == hi) {

```

```

        return lo + 1;
    }
}
int mid = lo + (hi - lo) / 2;
int res = min_left(i * 2 + 2, mid + 1, hi, tgt_hi, pred, acc);
return res != -1 ? res : min_left(i * 2 + 1, lo, mid, tgt_hi, pred, acc);
}

public:
explicit SegTree(int n, const T &v = T()) : len(n), value(4 * len) {
    assert(len > 0);
    build(0, 0, len - 1, [&](int) { return v; });
}

template<typename It>
SegTree(It lo, It hi) : len(hi - lo), value(4 * len) {
    assert(len > 0);
    build(0, 0, len - 1, [&](int i) { return *(lo + i); });
}

int size() const { return len; }

T at(int i) const {
    assert(0 <= i && i < len);
    return query(i, i);
}

T query(int lo, int hi) const {
    assert(0 <= lo && lo <= hi && hi < len);
    return query(0, 0, len - 1, lo, hi);
}

void update(int i, const T &d) {
    assert(0 <= i && i < len);
    update(0, 0, len - 1, i, d);
}

template<typename Pred>
int max_right(int lo, const Pred &pred) const {
    assert(0 <= lo && lo <= len);
    std::optional<T> acc;
    int res = max_right(0, 0, len - 1, lo, pred, acc);
    return res == -1 ? len : res;
}

template<typename Pred>
int min_left(int hi, const Pred &pred) const {
    assert(0 <= hi && hi <= len);
    std::optional<T> acc;
    int res = min_left(0, 0, len - 1, hi, pred, acc);
    return res == -1 ? 0 : res;
}
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10};
    SegTree<int> t(a.begin(), a.end());
    t.update(2, 4);
    vector<int> expected{6, -2, 4, 8, 10};
    for (int i = 0; i < t.size(); i++) {
        assert(t.at(i) == expected[i]);
    }
}

```

```

assert(t.query(0, 3) == -2);

// Boundary search by accumulated aggregate: stop before including the -2 at index 1.
assert(t.max_right(0, [](int mn) { return mn > -2; }) == 1);
assert(t.min_left(5, [](int mn) { return mn > -2; }) == 2);
assert(t.max_right(2, [](int mn) { return mn >= 4; }) == t.size());
assert(t.min_left(2, [](int mn) { return mn >= 0; }) == 2);
return 0;
}

```

2.4.6 Segment Tree (Range Update)

Maintain a fixed-size array while supporting range updates and range aggregate queries. Lazy propagation lets a segment tree defer work below intervals covered by an update: each node stores an aggregate value and a pending update that is pushed to its children only when needed.

The query operation is defined by an associative aggregate function `combine(a, b)`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`.

Range updates are defined by `apply_delta(v, d, len)`, which applies an update delta `d` to an aggregate summary `v` representing `len` array values, and by `compose_deltas(old, d)`, which combines a pending older delta with a newer delta in that order. These functions do not support arbitrary combinations: applying a delta to a combined segment must be equivalent to applying it to each child segment and then combining the results, and composed deltas must be equivalent to performing their updates sequentially. The default code below defines range assignment. For range increment, `compose_deltas(old, d)` should return `old + d`; `apply_delta(v, d, len)` should return `v + d` for range-min/range-max queries, and `v + d * len` for range-sum queries.

- `LazySegTree<T>(n, v = T())` constructs an array of size `n` with indices `[0, n)`, and all values initialized to `v`.
- `LazySegTree<T>(lo, hi)` constructs an array from the half-open random-access iterator range `[lo, hi)`.
- `size()` returns the size of the array.
- `at(i)` returns the value at index `i`.
- `query(lo, hi)` returns the result of `combine()` applied to all indices in `[lo, hi]`. If `lo == hi`, then the single specified value is returned.
- `update(i, d)` modifies the value at index `i` by applying the delta `d`.
- `update(lo, hi, d)` modifies the value at each array index in `[lo, hi]` by applying the delta `d` to each value.
- `max_right(lo, pred)` returns the largest boundary `hi` such that the aggregate over the half-open range `[lo, hi)` satisfies the monotonic predicate `pred()`. It returns `size()` if `pred()` remains true to the end.
- `min_left(hi, pred)` returns the smallest boundary `lo` such that the aggregate over the half-open range `[lo, hi)` satisfies the monotonic predicate `pred()`. It returns 0 if `pred()` remains true to the beginning.

For the boundary-search functions, `pred()` takes aggregate `T` values of candidate ranges. As a range grows, `pred()` may change from true to false but never back to true; the empty range is considered valid. E.g. for `combine = min`, use `pred(mn) = (mn > x)` to find the first value `<= x`, or for `combine = sum`, use `pred(sum) = (sum <= x)` with nonnegative values to find the longest range within the limit `x`.

Time Complexity:

- $O(n)$ per call to both constructors, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log n)$ per call to `at()`, `query()`, `update()`, `max_right()`, and `min_left()`.

Space Complexity:

- $O(n)$ for storage of the array elements.
- $O(\log n)$ auxiliary stack space for `query()`, `update()`, `max_right()`, and `min_left()`.
- $O(1)$ auxiliary for `size()`.

```
#include <algorithm>
```

```

#include <cassert>
#include <cstdint>
#include <optional>
#include <vector>

template<typename T>
class LazySegTree {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d, int64_t len) { return d; }
    static T compose_deltas(const T &d1, const T &d2) { return d2; }

    int len;
    std::vector<T> value, delta;
    std::vector<char> pending;

    template<typename Gen>
    void build(int i, int lo, int hi, const Gen &gen) {
        if (lo == hi) {
            value[i] = gen(lo);
            return;
        }
        int mid = lo + (hi - lo) / 2;
        build(i * 2 + 1, lo, mid, gen);
        build(i * 2 + 2, mid + 1, hi, gen);
        value[i] = combine(value[i * 2 + 1], value[i * 2 + 2]);
    }

    void push_delta(int i, int lo, int hi) {
        if (pending[i]) {
            value[i] = apply_delta(value[i], delta[i], hi - lo + 1);
            if (lo != hi) {
                int l = 2 * i + 1, r = 2 * i + 2;
                delta[l] = pending[l] ? compose_deltas(delta[l], delta[i]) : delta[i];
                delta[r] = pending[r] ? compose_deltas(delta[r], delta[i]) : delta[i];
                pending[l] = pending[r] = true;
            }
            pending[i] = false;
        }
    }

    T query(int i, int lo, int hi, int tgt_lo, int tgt_hi) {
        push_delta(i, lo, hi);
        if (lo == tgt_lo && hi == tgt_hi) {
            return value[i];
        }
        int mid = lo + (hi - lo) / 2;
        if (tgt_lo <= mid && mid < tgt_hi) {
            return combine(
                query(i * 2 + 1, lo, mid, tgt_lo, std::min(tgt_hi, mid)),
                query(i * 2 + 2, mid + 1, hi, std::max(tgt_lo, mid + 1), tgt_hi)
            );
        }
        if (tgt_lo <= mid) {
            return query(i * 2 + 1, lo, mid, tgt_lo, std::min(tgt_hi, mid));
        }
        return query(i * 2 + 2, mid + 1, hi, std::max(tgt_lo, mid + 1), tgt_hi);
    }

    void update(int i, int lo, int hi, int tgt_lo, int tgt_hi, const T &d) {
        push_delta(i, lo, hi);
        if (hi < tgt_lo || lo > tgt_hi) {
            return;
        }
        if (tgt_lo <= lo && hi <= tgt_hi) {
            delta[i] = d;
            pending[i] = true;
            push_delta(i, lo, hi);
            return;
        }

```

```

    }
    update(2 * i + 1, lo, lo + (hi - lo) / 2, tgt_lo, tgt_hi, d);
    update(2 * i + 2, lo + (hi - lo) / 2 + 1, hi, tgt_lo, tgt_hi, d);
    value[i] = combine(value[2 * i + 1], value[2 * i + 2]);
}

template<typename Pred>
int max_right(int i, int lo, int hi, int tgt_lo, const Pred &pred, std::optional<T> &acc) {
    push_delta(i, lo, hi);
    if (hi < tgt_lo) {
        return -1;
    }
    if (tgt_lo <= lo) {
        T next = acc ? combine(*acc, value[i]) : value[i];
        if (pred(next)) {
            acc = next;
            return -1;
        }
        if (lo == hi) {
            return lo;
        }
    }
    int mid = lo + (hi - lo) / 2;
    int res = max_right(i * 2 + 1, lo, mid, tgt_lo, pred, acc);
    return res != -1 ? res : max_right(i * 2 + 2, mid + 1, hi, tgt_lo, pred, acc);
}

template<typename Pred>
int min_left(int i, int lo, int hi, int tgt_hi, const Pred &pred, std::optional<T> &acc) {
    push_delta(i, lo, hi);
    if (tgt_hi <= lo) {
        return -1;
    }
    if (hi < tgt_hi) {
        T next = acc ? combine(value[i], *acc) : value[i];
        if (pred(next)) {
            acc = next;
            return -1;
        }
        if (lo == hi) {
            return lo + 1;
        }
    }
    int mid = lo + (hi - lo) / 2;
    int res = min_left(i * 2 + 2, mid + 1, hi, tgt_hi, pred, acc);
    return res != -1 ? res : min_left(i * 2 + 1, lo, mid, tgt_hi, pred, acc);
}

public:
    explicit LazySegTree(int n, const T &v = T())
        : len(n), value(4 * len), delta(4 * len), pending(4 * len, false) {
        assert(len > 0);
        build(0, 0, len - 1, [&](int) { return v; });
    }

    template<typename It>
    LazySegTree(It lo, It hi)
        : len(hi - lo), value(4 * len), delta(4 * len), pending(4 * len, false) {
        assert(len > 0);
        build(0, 0, len - 1, [&](int i) { return *(lo + i); });
    }

    int size() const { return len; }

    T at(int i) {
        assert(0 <= i && i < len);
        return query(i, i);
    }
}

```

```

T query(int lo, int hi) {
    assert(0 <= lo && lo <= hi && hi < len);
    return query(0, 0, len - 1, lo, hi);
}

void update(int i, const T &d) {
    assert(0 <= i && i < len);
    update(0, 0, len - 1, i, i, d);
}

void update(int lo, int hi, const T &d) {
    assert(0 <= lo && lo <= hi && hi < len);
    update(0, 0, len - 1, lo, hi, d);
}

template<typename Pred>
int max_right(int lo, const Pred &pred) {
    assert(0 <= lo && lo <= len);
    std::optional<T> acc;
    int res = max_right(0, 0, len - 1, lo, pred, acc);
    return res == -1 ? len : res;
}

template<typename Pred>
int min_left(int hi, const Pred &pred) {
    assert(0 <= hi && hi <= len);
    std::optional<T> acc;
    int res = min_left(0, 0, len - 1, hi, pred, acc);
    return res == -1 ? 0 : res;
}
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10};
    LazySegTree<int> t(a.begin(), a.end());
    t.update(2, 4);
    vector<int> expected{6, -2, 4, 8, 10};
    for (int i = 0; i < t.size(); i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.query(0, 3) == -2);
    t.update(0, 4, 5);
    t.update(3, 2);
    t.update(3, 1);
    expected = {5, 5, 5, 1, 5};
    for (int i = 0; i < t.size(); i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.query(0, 3) == 1);

    // Boundary search works through lazy updates too; values are now {5, 5, 5, 1, 5}.
    assert(t.max_right(0, [](int mn) { return mn > 1; }) == 3);
    assert(t.min_left(5, [](int mn) { return mn > 1; }) == 4);
    return 0;
}

```

2.4.7 Iterative Segment Tree (Point Update)

Maintain a fixed-size array while supporting point updates and range aggregate queries using an iterative segment tree. The leaves begin at a power-of-two offset in a flat array, and each internal node stores the aggregate of its two children. Point updates rebuild ancestors by walking upward, while queries scan a decomposition of the requested range. Compared with the recursive segment tree, these flat loops avoid recursive stack usage and can reduce constant factors.

The query operation is defined by an associative aggregate function `combine(a, b)`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`.

The point update operation is defined by `apply_delta(v, d)`, which returns the new value at a single updated index. The default definition below supports updates that “set” the chosen array index to a new value. Another possible update operation is “increment”, in which case `apply_delta(v, d)` should return `v + d`.

- `IterativeSegTree<T>(n, v = T())` constructs an array of size `n` with indices `[0, n)`, and all values initialized to `v`.
- `IterativeSegTree<T>(lo, hi)` constructs an array from the half-open random-access iterator range `[lo, hi)`.
- `size()` returns the size of the array.
- `at(i)` returns the value at index `i`.
- `query(lo, hi)` returns the result of `combine()` applied to all indices in `[lo, hi]`. If `lo == hi`, then the single specified value is returned.
- `update(i, d)` assigns the value `v` at index `i` to `apply_delta(v, d)`.
- `max_right(lo, pred)` returns the largest boundary `hi` such that the aggregate over the half-open range `[lo, hi)` satisfies the monotonic predicate `pred()`. It returns `size()` if `pred()` remains true to the end.
- `min_left(hi, pred)` returns the smallest boundary `lo` such that the aggregate over the half-open range `[lo, hi)` satisfies the monotonic predicate `pred()`. It returns 0 if `pred()` remains true to the beginning.

For the boundary-search functions, `pred()` takes aggregate `T` values of candidate ranges. As a range grows, `pred()` may change from true to false but never back to true; the empty range is considered valid. E.g. for `combine = min`, use `pred(mn) = (mn > x)` to find the first value `<= x`, or for `combine = sum`, use `pred(sum) = (sum <= x)` with nonnegative values to find the longest range within the limit `x`.

Both boundary searches are iterative to demonstrate how the flat-array layout can avoid recursion. They decompose the search range into $O(\log n)$ canonical nodes, scan those nodes in order while accumulating their aggregates, then descend inside the first node that makes `pred` false. A recursive descent like the one in 2.4.5 is equally valid and may be easier to adapt.

Time Complexity:

- $O(n)$ per call to both constructors, where n is the size of the array.
- $O(1)$ per call to `size()` and `at()`.
- $O(\log n)$ per call to `query()`, `update()`, `max_right()`, and `min_left()`.

Space Complexity:

- $O(n)$ for storage of the array elements.
- $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cassert>
#include <optional>
#include <vector>

template<typename T>
class IterativeSegTree {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d) { return d; }

    int len, base;
    std::vector<T> value;
```

```

template<typename Gen>
void build(const Gen &gen) {
    base = 1;
    while (base < len) {
        base <= 1;
    }
    value.assign(2 * base, T());
    for (int i = 0; i < len; i++) {
        value[base + i] = gen(i);
    }
    for (int i = base - 1; i > 0; i--) {
        value[i] = combine(value[i << 1], value[i << 1 | 1]);
    }
}

public:
explicit IterativeSegTree(int n, const T &v = T()) : len(n) {
    assert(len > 0);
    build([&](int) { return v; });
}

template<typename It>
IterativeSegTree(It lo, It hi) : len(hi - lo) {
    assert(len > 0);
    build([&](int i) { return *(lo + i); });
}

int size() const { return len; }

T at(int i) const {
    assert(0 <= i && i < len);
    return value[base + i];
}

T query(int lo, int hi) const {
    assert(0 <= lo && lo <= hi && hi < len);
    std::optional<T> left, right;
    for (lo += base, hi += base + 1; lo < hi; lo >>= 1, hi >>= 1) {
        if (lo & 1) {
            left = left ? combine(*left, value[lo++]) : value[lo++];
        }
        if (hi & 1) {
            right = right ? combine(value[--hi], *right) : value[--hi];
        }
    }
    return !left ? *right : (!right ? *left : combine(*left, *right));
}

void update(int i, const T &d) {
    assert(0 <= i && i < len);
    i += base;
    value[i] = apply_delta(value[i], d);
    for (i >>= 1; i > 0; i >>= 1) {
        value[i] = combine(value[i << 1], value[i << 1 | 1]);
    }
}

template<typename Pred>
int max_right(int lo, const Pred &pred) const {
    assert(0 <= lo && lo <= len);
    std::optional<T> acc;
    for (int pos = lo; pos < len; pos++) {
        int width = pos == 0 ? base : pos & -pos;
        while (width > len - pos) {
            width >>= 1;
        }
        int i = (base + pos) / width;
    }
}

```



```

    T next = acc ? combine(*acc, value[i]) : value[i];
    if (pred(next)) {
        acc = next;
        pos += width;
        continue;
    }
    while (i < base) {
        i <= 1;
        next = acc ? combine(*acc, value[i]) : value[i];
        if (pred(next)) {
            acc = next;
            i++;
        }
    }
    return i - base;
}
return len;
}

template<typename Pred>
int min_left(int hi, const Pred &pred) const {
    assert(0 <= hi && hi <= len);
    std::optional<T> acc;
    for (int pos = hi; pos > 0;) {
        int width = pos & -pos;
        int i = (base + pos) / width - 1;
        T next = acc ? combine(value[i], *acc) : value[i];
        if (pred(next)) {
            acc = next;
            pos -= width;
            continue;
        }
        while (i < base) {
            i = i << 1 | 1;
            next = acc ? combine(value[i], *acc) : value[i];
            if (pred(next)) {
                acc = next;
                i--;
            }
        }
        return i - base + 1;
    }
    return 0;
}
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10};
    IterativeSegTree<int> t(a.begin(), a.end());
    t.update(2, 4);
    vector<int> expected{6, -2, 4, 8, 10};
    for (int i = 0; i < t.size(); i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.query(0, 3) == -2);

    assert(t.max_right(0, [](int mn) { return mn > -2; }) == 1);
    assert(t.min_left(5, [](int mn) { return mn > -2; }) == 2);
    return 0;
}

```

2.4.8 Iterative Segment Tree (Range Update)

Maintain a fixed-size array while supporting range updates and range aggregate queries using an iterative segment tree with lazy propagation. The leaves begin at a power-of-two offset in a flat array, and each internal node stores an aggregate value plus a pending update. Range updates and queries push pending updates only along paths that inspect children, then scan a decomposition of the requested range. Compared with the recursive lazy segment tree, these flat loops avoid recursive stack usage and can reduce constant factors, though the recursive version is usually easier to adapt when the update or search logic becomes unusual.

The query operation is defined by an associative aggregate function `combine(a, b)`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`.

Range updates are defined by `apply_delta(v, d, len)`, which applies an update delta `d` to an aggregate summary `v` representing `len` array values, and by `compose_deltas(old, d)`, which combines a pending older delta with a newer delta in that order. These functions do not support arbitrary combinations: applying a delta to a combined segment must be equivalent to applying it to each child segment and then combining the results, and composed deltas must be equivalent to performing their updates sequentially. The default code below defines range assignment. For range increment, `compose_deltas(old, d)` should return `old + d`; `apply_delta(v, d, len)` should return `v + d` for range-min/range-max queries, and `v + d * len` for range-sum queries.

- `IterativeLazySegTree<T>(n, v = T())` constructs an array of size `n` with indices `[0, n)`, and all values initialized to `v`.
- `IterativeLazySegTree<T>(lo, hi)` constructs an array from the half-open random-access iterator range `[lo, hi)`.
- `size()` returns the size of the array.
- `at(i)` returns the value at index `i`.
- `query(lo, hi)` returns the result of `combine()` applied to all indices in `[lo, hi)`. If `lo == hi`, then the single specified value is returned.
- `update(i, d)` assigns the value at index `i` by applying the delta `d`.
- `update(lo, hi, d)` modifies the value at each array index in `[lo, hi)` by applying the delta `d` to each value.
- `max_right(lo, pred)` returns the largest boundary `hi` such that the aggregate over the half-open range `[lo, hi)` satisfies the monotonic predicate `pred()`. It returns `size()` if `pred()` remains true to the end.
- `min_left(hi, pred)` returns the smallest boundary `lo` such that the aggregate over the half-open range `[lo, hi)` satisfies the monotonic predicate `pred()`. It returns 0 if `pred()` remains true to the beginning.

For the boundary-search functions, `pred()` takes aggregate `T` values of candidate ranges. As a range grows, `pred()` may change from true to false but never back to true; the empty range is considered valid. E.g. for `combine = min`, use `pred(mn) = (mn > x)` to find the first value `<= x`, or for `combine = sum`, use `pred(sum) = (sum <= x)` with nonnegative values to find the longest range within the limit `x`.

Both boundary searches are iterative to demonstrate how the flat-array layout can avoid recursion. They push lazy deltas along the endpoints, decompose the search range into $O(\log n)$ canonical nodes, scan those nodes in order while accumulating their aggregates, then push and descend inside the first node that makes `pred` false. A recursive descent like the one in 2.4.6 is equally valid and may be easier to adapt.

Time Complexity:

- $O(n)$ per call to both constructors, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log n)$ per call to `at()`, `query()`, `update()`, `max_right()`, and `min_left()`.

Space Complexity:

- $O(n)$ for storage of the array elements and lazy deltas.
- $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cassert>
#include <stdint>
```

```

#include <optional>
#include <vector>

template<typename T>
class IterativeLazySegTree {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d, int64_t len) { return d; }
    static T compose_deltas(const T &d1, const T &d2) { return d2; }

    int len, base, height;
    std::vector<T> value, delta;
    std::vector<int64_t> seg_len;
    std::vector<char> pending;

    template<typename Gen>
    void build(const Gen &gen) {
        base = 1;
        height = 0;
        while (base < len) {
            base <= 1;
            height++;
        }
        value.assign(2 * base, T());
        delta.assign(2 * base, T());
        pending.assign(2 * base, false);
        seg_len.assign(2 * base, 1);
        for (int i = base - 1; i > 0; i--) {
            seg_len[i] = seg_len[i << 1] + seg_len[i << 1 | 1];
        }
        for (int i = 0; i < len; i++) {
            value[base + i] = gen(i);
        }
        for (int i = base - 1; i > 0; i--) {
            pull(i);
        }
    }

    void pull(int i) { value[i] = combine(value[i << 1], value[i << 1 | 1]); }

    void apply(int i, const T &d) {
        value[i] = apply_delta(value[i], d, seg_len[i]);
        if (i < base) {
            delta[i] = pending[i] ? compose_deltas(delta[i], d) : d;
            pending[i] = true;
        }
    }

    void push(int i) {
        if (pending[i]) {
            apply(i << 1, delta[i]);
            apply(i << 1 | 1, delta[i]);
            pending[i] = false;
        }
    }

    void push_path(int i) {
        for (int h = height; h > 0; h--) {
            push(i >> h);
        }
    }

    void pull_after_update(int l, int r) {
        for (int h = 1; h <= height; h++) {
            if (((l >> h) << h) != 1) {
                pull(l >> h);
            }
            if (((r >> h) << h) != r) {
                pull((r - 1) >> h);
            }
        }
    }
};

```

```

    }
}

public:
explicit IterativeLazySegTree(int n, const T &v = T()) : len(n) {
    assert(len > 0);
    build([&](int) { return v; });
}

template<typename It>
IterativeLazySegTree(It lo, It hi) : len(hi - lo) {
    assert(len > 0);
    build([&](int i) { return *(lo + i); });
}

int size() const { return len; }

T at(int i) {
    assert(0 <= i && i < len);
    i += base;
    push_path(i);
    return value[i];
}

T query(int lo, int hi) {
    assert(0 <= lo && lo <= hi && hi < len);
    int l = lo + base, r = hi + 1 + base;
    push_path(l);
    push_path(r - 1);
    std::optional<T> left, right;
    for (; l < r; l >>= 1, r >>= 1) {
        if (l & 1) {
            left = left ? combine(*left, value[l++]) : value[l++];
        }
        if (r & 1) {
            right = right ? combine(value[--r], *right) : value[--r];
        }
    }
    return !left ? *right : (!right ? *left : combine(*left, *right));
}

void update(int lo, int hi, const T &d) {
    assert(0 <= lo && lo <= hi && hi < len);
    int l = lo + base, r = hi + 1 + base;
    push_path(l);
    push_path(r - 1);
    int l0 = l, r0 = r;
    for (; l < r; l >>= 1, r >>= 1) {
        if (l & 1) {
            apply(l++, d);
        }
        if (r & 1) {
            apply(--r, d);
        }
    }
    pull_after_update(l0, r0);
}

void update(int i, const T &d) {
    assert(0 <= i && i < len);
    update(i, i, d);
}

template<typename Pred>
int max_right(int lo, const Pred &pred) {
    assert(0 <= lo && lo <= len);
    if (lo == len) {

```

```

    return len;
}
push_path(lo + base);
push_path(len - 1 + base);
std::optional<T> acc;
for (int pos = lo; pos < len;) {
    int width = pos == 0 ? base : pos & -pos;
    while (width > len - pos) {
        width >>= 1;
    }
    int i = (base + pos) / width;
    T next = acc ? combine(*acc, value[i]) : value[i];
    if (pred(next)) {
        acc = next;
        pos += width;
        continue;
    }
    while (i < base) {
        push(i);
        i <<= 1;
        next = acc ? combine(*acc, value[i]) : value[i];
        if (pred(next)) {
            acc = next;
            i++;
        }
    }
    return i - base;
}
return len;
}

template<typename Pred>
int min_left(int hi, const Pred &pred) {
    assert(0 <= hi && hi <= len);
    if (hi == 0) {
        return 0;
    }
    push_path(base);
    push_path(hi - 1 + base);
    std::optional<T> acc;
    for (int pos = hi; pos > 0;) {
        int width = pos & -pos;
        int i = (base + pos) / width - 1;
        T next = acc ? combine(value[i], *acc) : value[i];
        if (pred(next)) {
            acc = next;
            pos -= width;
            continue;
        }
        while (i < base) {
            push(i);
            i = i << 1 | 1;
            next = acc ? combine(value[i], *acc) : value[i];
            if (pred(next)) {
                acc = next;
                i--;
            }
        }
        return i - base + 1;
    }
    return 0;
}
};

```

Example Usage

```
using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10};
    IterativeLazySegTree<int> t(a.begin(), a.end());
    t.update(2, 4);
    vector<int> expected{6, -2, 4, 8, 10};
    for (int i = 0; i < t.size(); i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.query(0, 3) == -2);
    t.update(0, 4, 5);
    t.update(3, 2);
    t.update(3, 1);
    expected = {5, 5, 5, 1, 5};
    for (int i = 0; i < t.size(); i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.query(0, 3) == 1);

    // Boundary search works through lazy updates too; values are now {5, 5, 5, 1, 5}.
    assert(t.max_right(0, [](int mn) { return mn > 1; }) == 3);
    assert(t.min_left(5, [](int mn) { return mn > 1; }) == 4);
    return 0;
}
```

2.4.9 Segment Tree Beats

Maintain a fixed-size array under the range “chmin” update $a[i] = \min(a[i], t)$ for all i in a range, while answering range-sum and range-maximum queries. An ordinary lazy segment tree cannot represent a clamp like this with a single composable tag, because clamping affects only the entries that currently exceed t . Segment tree beats (a.k.a. the Ji Driver tree) resolves this by storing, for each node, the strict maximum, the second-largest distinct value, and a count of entries equal to the maximum. A clamp then has three cases at each node: if $t \geq \text{max1}$, the clamp changes nothing and the recursion stops (the prune case); if $\text{max2} < t < \text{max1}$, every entry equal to max1 drops to t and all others are untouched, so the node updates in $O(1)$, the sum falls by $(\text{max1} - t) * \text{count}$, and max1 becomes t (the break, or tag, case); otherwise ($t \leq \text{max2}$) the node is too coarse to update directly, so the clamp recurses into both children and re-pulls. The lazy tag pushes down to a child only when the child’s maximum is larger.

The break condition is what makes this efficient: a potential-function argument on the number of distinct values along root-to-leaf paths shows the total work over any sequence of clamps is amortized $O(\log^2 n)$ each, even though a single clamp can momentarily descend deep. The same scheme extends to range “chmax” by tracking the symmetric minimum trio, and to mixed clamp-and-add workloads by carrying an additional add tag; only the chmin / sum / max core is shown here. Unlike the generic lazy segment tree, this structure cannot be reduced to a pluggable `combine` and `apply`, because a clamp acts on only the maximal entries of a node rather than uniformly on all of them, so each operation set must be hand-written and its amortized bound argued separately.

- `SegTreeBeats<T>(n, v = T())` constructs an array of size n , indices $[0, n)$, all equal to v .
- `SegTreeBeats<T>(lo, hi)` constructs an array from the half-open random-access iterator range $[lo, hi)$.
- `size()` returns the size of the array.
- `chmin(lo, hi, t)` replaces $a[i]$ with $\min(a[i], t)$ for every i in $[lo, hi]$.
- `query_sum(lo, hi)` returns the sum of $a[i]$ over i in $[lo, hi]$.
- `query_max(lo, hi)` returns the maximum of $a[i]$ over i in $[lo, hi]$.
- `at(i)` returns the value at index i .

Time Complexity:

- $O(n)$ per call to both constructors, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log^2 n)$ amortized per call to `chmin()`.
- $O(\log n)$ per call to `query_sum()`, `query_max()`, and `at()`.

Space Complexity:

- $O(n)$ for storage of the tree.
- $O(\log n)$ auxiliary stack space per `chmin()` and query call.

```
#include <algorithm>
#include <cassert>
#include <cstdlib>
#include <limits>
#include <vector>

template<typename T>
class SegTreeBeats {
    static constexpr T NEG_INF = std::numeric_limits<T>::lowest();

    int len;
    std::vector<T> max1, max2, sum;
    std::vector<int> cnt; // Number of entries in the node equal to max1.

    void pull(int i) {
        int l = 2 * i + 1, r = 2 * i + 2;
        sum[i] = sum[l] + sum[r];
        if (max1[l] == max1[r]) {
            max1[i] = max1[l];
            cnt[i] = cnt[l] + cnt[r];
            max2[i] = std::max(max2[l], max2[r]);
        } else if (max1[l] > max1[r]) {
            max1[i] = max1[l];
            cnt[i] = cnt[l];
            max2[i] = std::max(max2[l], max1[r]);
        } else {
            max1[i] = max1[r];
            cnt[i] = cnt[r];
            max2[i] = std::max(max1[l], max2[r]);
        }
    }

    // Lower the node's maximum to t. Valid only when max2 < t < max1, so exactly the cnt entries
    // equal to max1 change and the second maximum is left untouched.
    void apply_chmin(int i, const T &t) {
        sum[i] -= (max1[i] - t) * static_cast<T>(cnt[i]); // Overflow warning.
        max1[i] = t;
    }

    void push(int i) {
        for (int c : {2 * i + 1, 2 * i + 2}) {
            if (max1[c] > max1[i]) {
                apply_chmin(c, max1[i]);
            }
        }
    }

    template<typename Gen>
    void build(int i, int lo, int hi, const Gen &gen) {
        if (lo == hi) {
            max1[i] = sum[i] = gen(lo);
            max2[i] = NEG_INF;
            cnt[i] = 1;
            return;
        }
    }
}
```

```

    int mid = lo + (hi - lo) / 2;
    build(i * 2 + 1, lo, mid, gen);
    build(i * 2 + 2, mid + 1, hi, gen);
    pull(i);
}

void chmin(int i, int lo, int hi, int tlo, int thi, const T &t) {
    if (thi < lo || hi < tlo || max1[i] <= t) {
        return; // Prune: outside the range, or every entry is already at most t.
    }
    if (tlo <= lo && hi <= thi && (lo == hi || max2[i] < t)) {
        apply_chmin(i, t); // Break: only the maximal entries are affected.
        return;
    }
    push(i);
    int mid = lo + (hi - lo) / 2;
    chmin(i * 2 + 1, lo, mid, tlo, thi, t);
    chmin(i * 2 + 2, mid + 1, hi, tlo, thi, t);
    pull(i);
}

T query_sum(int i, int lo, int hi, int tlo, int thi) {
    if (thi < lo || hi < tlo) {
        return 0;
    }
    if (tlo <= lo && hi <= thi) {
        return sum[i];
    }
    push(i);
    int mid = lo + (hi - lo) / 2;
    return query_sum(i * 2 + 1, lo, mid, tlo, thi) + query_sum(i * 2 + 2, mid + 1, hi, tlo, thi);
}

T query_max(int i, int lo, int hi, int tlo, int thi) {
    if (thi < lo || hi < tlo) {
        return NEG_INF;
    }
    if (tlo <= lo && hi <= thi) {
        return max1[i];
    }
    push(i);
    int mid = lo + (hi - lo) / 2;
    return std::max(
        query_max(i * 2 + 1, lo, mid, tlo, thi), query_max(i * 2 + 2, mid + 1, hi, tlo, thi)
    );
}

public:
    explicit SegTreeBeats(int n, const T &v = T())
        : len(n), max1(4 * n), max2(4 * n), sum(4 * n), cnt(4 * n) {
        assert(len > 0);
        build(0, 0, len - 1, [&](int) { return v; });
    }

    template<typename It>
    SegTreeBeats(It lo, It hi)
        : len(static_cast<int>(hi - lo)), max1(4 * len), max2(4 * len), sum(4 * len), cnt(4 * len) {
        assert(len > 0);
        build(0, 0, len - 1, [&](int i) { return *(lo + i); });
    }

    int size() const { return len; }

    void chmin(int lo, int hi, const T &t) {
        assert(0 <= lo && lo <= hi && hi < len);
        chmin(0, 0, len - 1, lo, hi, t);
    }

```



```

T query_sum(int lo, int hi) {
    assert(0 <= lo && lo <= hi && hi < len);
    return query_sum(0, 0, len - 1, lo, hi);
}

T query_max(int lo, int hi) {
    assert(0 <= lo && lo <= hi && hi < len);
    return query_max(0, 0, len - 1, lo, hi);
}

T at(int i) {
    assert(0 <= i && i < len);
    return query_sum(0, 0, len - 1, i, i);
}
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int64_t> a{5, 4, 3, 2, 1};
    SegTreeBeats<int64_t> t(a.begin(), a.end());
    assert(t.query_sum(0, 4) == 15);
    assert(t.query_max(0, 4) == 5);

    t.chmin(0, 2, 3); // {3, 3, 3, 2, 1}: only the 5 and 4 are clamped down to 3.
    assert(t.at(0) == 3);
    assert(t.at(1) == 3);
    assert(t.query_sum(0, 4) == 12);
    assert(t.query_max(0, 4) == 3);

    t.chmin(0, 4, 2); // {2, 2, 2, 2, 1}.
    assert(t.query_sum(0, 4) == 9);
    assert(t.query_max(0, 3) == 2);

    SegTreeBeats<int> lowest(1, numeric_limits<int>::lowest() + 1);
    lowest.chmin(0, 0, numeric_limits<int>::lowest());
    assert(lowest.at(0) == numeric_limits<int>::lowest());

    // Cross-check against a brute-force array over random clamps and queries.
    vector<int64_t> b{7, 1, 4, 9, 2, 6, 5, 8, 3, 0};
    SegTreeBeats<int64_t> s(b.begin(), b.end());
    int clamps[][3] = {{0, 9, 6}, {2, 5, 3}, {4, 8, 7}, {0, 4, 2}, {3, 9, 5}};
    for (auto &c : clamps) {
        for (int i = c[0]; i <= c[1]; i++) {
            b[i] = min(b[i], static_cast<int64_t>(c[2]));
        }
        s.chmin(c[0], c[1], c[2]);
        for (int lo = 0; lo < 10; lo++) {
            int64_t want_sum = 0, want_max = INT64_MIN;
            for (int hi = lo; hi < 10; hi++) {
                want_sum += b[hi];
                want_max = max(want_max, b[hi]);
                assert(s.query_sum(lo, hi) == want_sum);
                assert(s.query_max(lo, hi) == want_max);
            }
        }
    }
    return 0;
}

```

2.4.10 Sparse Segment Tree

A sparse segment tree (also commonly called a dynamic or implicit segment tree) maintains an array over a large index range while supporting both dynamic queries and updates of contiguous subarrays via the lazy propagation technique. This implementation uses lazy initialization of nodes to conserve memory: only the nodes covering touched indices are ever allocated, so indices up to N are supported without preallocating the whole tree.

The query operation is defined by an associative aggregate function `combine(a, b)`. Since untouched nodes are implicit, `repeat_value(v, len)` must return the aggregate summary of `len` copies of the initial value `v`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. For range-sum queries, `combine(a, b)` should return `a + b` and `repeat_value(v, len)` should return `v * len`.

Range updates are defined by `apply_delta(v, d, len)`, which applies an update delta `d` to an aggregate summary `v` representing `len` array values, and by `compose_deltas(old, d)`, which combines a pending older delta with a newer delta in that order. These functions do not support arbitrary combinations: applying a delta to a combined segment must be equivalent to applying it to each child segment and then combining the results, and composed deltas must be equivalent to performing their updates sequentially. The default code below defines range assignment. For range increment, `compose_deltas(old, d)` should return `old + d`; `apply_delta(v, d, len)` should return `v + d` for range-min/range-max queries, and `v + d * len` for range-sum queries.

- `SparseSegTree<T, N>(v = T())` constructs an array over indices $[0, N]$, with every value implicitly initialized to `v`. Nodes are allocated lazily as indices are touched.
- `at(i)` returns the value at index `i`, where `i` must be in $[0, N]$.
- `query(lo, hi)` returns the result of `combine()` applied to all indices in $[lo, hi]$. If `lo == hi`, then the single specified value is returned.
- `update(i, d)` assigns the value `v` at index `i` to `apply_delta(v, d)`.
- `update(lo, hi, d)` modifies the value at each array index in $[lo, hi]$ by applying the delta `d` to each value.
- `max_right(lo, pred)` returns the largest boundary `hi` such that the aggregate over the half-open range $[lo, hi)$ satisfies the monotonic predicate `pred()`. It returns $N + 1$ if `pred()` remains true to the end.
- `min_left(hi, pred)` returns the smallest boundary `lo` such that the aggregate over the half-open range $[lo, hi)$ satisfies the monotonic predicate `pred()`. It returns 0 if `pred()` remains true to the beginning.

For the boundary-search functions, `pred()` takes aggregate `T` values of candidate ranges. As a range grows, `pred()` may change from true to false but never back to true; the empty range is considered valid. E.g. for `combine = min`, use `pred(mn) = (mn > x)` to find the first value $\leq x$, or for `combine = sum`, use `pred(sum) = (sum <= x)` with nonnegative values to find the longest range within the limit `x`.

Time Complexity:

- $O(1)$ per call to the constructor.
- $O(\log N)$ per call to `at()`, `query()`, `update()`, `max_right()`, and `min_left()`.

Space Complexity:

- $O(q \log N)$ for storage after q updates.
- $O(\log N)$ auxiliary stack space for `query()`, `update()`, `max_right()`, and `min_left()`.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <limits>
#include <optional>
#include <vector>

template<typename T, int N = 1000000000>
class SparseSegTree {
    static_assert(0 <= N && N < std::numeric_limits<int>::max());

    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T repeat_value(const T &v, int64_t len) { return v; }
    static T apply_delta(const T &v, const T &d, int64_t len) { return d; }
    static T compose_deltas(const T &d1, const T &d2) { return d2; }
```

```

struct Node {
    T value, delta;
    bool pending;
    Node *left, *right;

    explicit Node(const T &v) : value(v), pending(false), left(nullptr), right(nullptr) {}
} *root;

T init;

void update_delta(Node *&n, const T &d, int64_t len) {
    if (n == nullptr) {
        n = new Node(repeat_value(init, len));
    }
    n->delta = n->pending ? compose_deltas(n->delta, d) : d;
    n->pending = true;
}

void push_delta(Node *n, int lo, int hi) {
    if (n == nullptr) {
        return;
    }
    if (n->pending) {
        n->value = apply_delta(n->value, n->delta, hi - lo + 1);
        if (lo != hi) {
            int mid = lo + (hi - lo) / 2;
            update_delta(n->left, n->delta, mid - lo + 1);
            update_delta(n->right, n->delta, hi - mid);
        }
    }
    n->pending = false;
}

T query(Node *n, int lo, int hi, int tgt_lo, int tgt_hi) {
    if (n == nullptr) {
        return repeat_value(init, hi - lo + 1);
    }
    push_delta(n, lo, hi);
    if (lo == tgt_lo && hi == tgt_hi) {
        return n->value;
    }
    int mid = lo + (hi - lo) / 2;
    if (tgt_lo <= mid && mid < tgt_hi) {
        return combine(
            query(n->left, lo, mid, tgt_lo, std::min(tgt_hi, mid)),
            query(n->right, mid + 1, hi, std::max(tgt_lo, mid + 1), tgt_hi)
        );
    }
    if (tgt_lo <= mid) {
        return query(n->left, lo, mid, tgt_lo, std::min(tgt_hi, mid));
    }
    return query(n->right, mid + 1, hi, std::max(tgt_lo, mid + 1), tgt_hi);
}

void update(Node *&n, int lo, int hi, int tgt_lo, int tgt_hi, const T &d) {
    if (n == nullptr) {
        if (hi < tgt_lo || lo > tgt_hi) {
            return;
        }
        n = new Node(repeat_value(init, hi - lo + 1));
    } else {
        push_delta(n, lo, hi);
    }
    if (hi < tgt_lo || lo > tgt_hi) {
        return;
    }
    if (tgt_lo <= lo && hi <= tgt_hi) {

```

```

    n->delta = d;
    n->pending = true;
    push_delta(n, lo, hi);
    return;
}
int mid = lo + (hi - lo) / 2;
update(n->left, lo, mid, tgt_lo, tgt_hi, d);
update(n->right, mid + 1, hi, tgt_lo, tgt_hi, d);
T left_value = (n->left != nullptr) ? n->left->value : repeat_value(init, mid - lo + 1);
T right_value = (n->right != nullptr) ? n->right->value : repeat_value(init, hi - mid);
n->value = combine(left_value, right_value);
}

template<typename Pred>
int max_right(Node *n, int lo, int hi, int tgt_lo, const Pred &pred, std::optional<T> &acc) {
    if (hi < tgt_lo) {
        return -1;
    }
    if (n != nullptr) {
        push_delta(n, lo, hi);
    }
    T node_value = n != nullptr ? n->value : repeat_value(init, hi - lo + 1);
    if (tgt_lo <= lo) {
        T next = acc ? combine(*acc, node_value) : node_value;
        if (pred(next)) {
            acc = next;
            return -1;
        }
        if (lo == hi) {
            return lo;
        }
    }
    int mid = lo + (hi - lo) / 2;
    int res = max_right(n == nullptr ? nullptr : n->left, lo, mid, tgt_lo, pred, acc);
    return res != -1 ? res
        : max_right(n == nullptr ? nullptr : n->right, mid + 1, hi, tgt_lo, pred, acc);
}

template<typename Pred>
int min_left(Node *n, int lo, int hi, int tgt_hi, const Pred &pred, std::optional<T> &acc) {
    if (tgt_hi <= lo) {
        return -1;
    }
    if (n != nullptr) {
        push_delta(n, lo, hi);
    }
    T node_value = n != nullptr ? n->value : repeat_value(init, hi - lo + 1);
    if (hi < tgt_hi) {
        T next = acc ? combine(node_value, *acc) : node_value;
        if (pred(next)) {
            acc = next;
            return -1;
        }
        if (lo == hi) {
            return lo + 1;
        }
    }
    int mid = lo + (hi - lo) / 2;
    int res = min_left(n == nullptr ? nullptr : n->right, mid + 1, hi, tgt_hi, pred, acc);
    return res != -1 ? res : min_left(n == nullptr ? nullptr : n->left, lo, mid, tgt_hi, pred, acc);
}

void clean_up(Node *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

```

```

}

public:
    explicit SparseSegTree(const T &v = T()) : root(nullptr), init(v) {}

    ~SparseSegTree() { clean_up(root); }
    SparseSegTree(const SparseSegTree &) = delete;
    SparseSegTree &operator=(const SparseSegTree &) = delete;

    T at(int i) {
        assert(0 <= i && i <= N);
        return query(i, i);
    }

    T query(int lo, int hi) {
        assert(0 <= lo && lo <= hi && hi <= N);
        return query(root, 0, N, lo, hi);
    }

    void update(int i, const T &d) {
        assert(0 <= i && i <= N);
        update(i, i, d);
    }

    void update(int lo, int hi, const T &d) {
        assert(0 <= lo && lo <= hi && hi <= N);
        update(root, 0, N, lo, hi, d);
    }

    template<typename Pred>
    int max_right(int lo, const Pred &pred) {
        assert(0 <= lo && lo <= N + 1);
        std::optional<T> acc;
        int res = max_right(root, 0, N, lo, pred, acc);
        return res == -1 ? N + 1 : res;
    }

    template<typename Pred>
    int min_left(int hi, const Pred &pred) {
        assert(0 <= hi && hi <= N + 1);
        std::optional<T> acc;
        int res = min_left(root, 0, N, hi, pred, acc);
        return res == -1 ? 0 : res;
    }
};

```

Example Usage

```

using namespace std;

int main() {
    SparseSegTree<int> t(0);
    t.update(0, 6);
    t.update(1, -2);
    t.update(2, 4);
    t.update(3, 8);
    t.update(4, 10);
    vector<int> expected{6, -2, 4, 8, 10};
    for (int i = 0; i < 5; i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.query(0, 3) == -2);
    t.update(0, 4, 5);
    t.update(3, 2);
    t.update(3, 1);
    expected = {5, 5, 5, 1, 5};
}

```

```

for (int i = 0; i < 5; i++) {
    assert(t.at(i) == expected[i]);
}
assert(t.query(0, 3) == 1);

// Boundary search over the sparse domain; values at indices 0..4 are now {5, 5, 5, 1, 5}.
assert(t.max_right(0, [](int mn) { return mn > 1; }) == 3);
assert(t.min_left(5, [](int mn) { return mn > 1; }) == 4);
return 0;
}

```

2.4.11 Persistent Segment Tree

Maintain immutable versions of a fixed-size array while supporting point updates and range aggregate queries. Each update creates a new root by copying only the $O(\log n)$ nodes on the path from the root to the updated leaf. All unchanged subtrees are shared with previous versions, and updates may branch from any stored version.

The query operation is defined by an associative aggregate function `combine(a, b)`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`.

The point update operation is defined by `apply_delta(v, d)`, which returns the new value at a single updated index. The default definition below supports updates that “set” the chosen array index to a new value. Another possible update operation is “increment”, in which case `apply_delta(v, d)` should return `v + d`.

- `PersistentSegTree<T>(n, v = T())` constructs version 0 of an array of size `n` with indices $[0, n)$, and all values initialized to `v`.
- `PersistentSegTree<T>(lo, hi)` constructs version 0 from the half-open random-access iterator range $[lo, hi)$.
- `size()` returns the size of every version of the array.
- `versions()` returns the number of stored versions. Versions are numbered starting from 0.
- `at(version, i)` returns the value at index `i` in the specified version.
- `query(version, lo, hi)` returns the result of `combine()` applied to all indices in $[lo, hi]$ in the specified version. If `lo == hi`, then the single specified value is returned.
- `update(version, i, d)` creates and returns a new version in which the value at index `i` is assigned to `apply_delta(v, d)`. The specified version remains unchanged.

The boundary-search routines `max_right(lo, pred)` and `min_left(hi, pred)` from the recursive point-update segment tree in 2.4.5 use the same traversal here and may be adapted by starting at `roots[version]` and following the child indices stored in each node.

Time Complexity:

- $O(n)$ per call to either constructor, where n is the size of the array.
- $O(1)$ per call to `size()` and `versions()`.
- $O(\log n)$ per call to `at()`, `query()`, and `update()`.

Space Complexity:

- $O(n + q \log n)$ for storage after q updates.
- $O(\log n)$ auxiliary stack space for construction, `at()`, `query()`, and `update()`.
- $O(1)$ auxiliary for `size()` and `versions()`.

```

#include <algorithm>
#include <cassert>
#include <vector>

template<typename T>
class PersistentSegTree {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &, const T &d) { return d; }

```

```

struct Node {
    T value;
    int left, right;

    explicit Node(const T &value, int left = -1, int right = -1)
        : value(value), left(left), right(right) {}
};

int len;
std::vector<Node> nodes;
std::vector<int> roots;

template<typename Gen>
int build(int lo, int hi, const Gen &gen) {
    if (lo == hi) {
        nodes.emplace_back(gen(lo));
        return static_cast<int>(nodes.size()) - 1;
    }
    int mid = lo + (hi - lo) / 2;
    int left = build(lo, mid, gen);
    int right = build(mid + 1, hi, gen);
    nodes.emplace_back(combine(nodes[left].value, nodes[right].value), left, right);
    return static_cast<int>(nodes.size()) - 1;
}

T query(int i, int lo, int hi, int tgt_lo, int tgt_hi) const {
    if (lo == tgt_lo && hi == tgt_hi) {
        return nodes[i].value;
    }
    int mid = lo + (hi - lo) / 2;
    if (tgt_lo <= mid && mid < tgt_hi) {
        return combine(
            query(nodes[i].left, lo, mid, tgt_lo, mid),
            query(nodes[i].right, mid + 1, hi, mid + 1, tgt_hi)
        );
    }
    if (tgt_lo <= mid) {
        return query(nodes[i].left, lo, mid, tgt_lo, tgt_hi);
    }
    return query(nodes[i].right, mid + 1, hi, tgt_lo, tgt_hi);
}

int update(int cur, int lo, int hi, int target, const T &d) {
    if (lo == hi) {
        nodes.emplace_back(apply_delta(nodes[cur].value, d));
        return static_cast<int>(nodes.size()) - 1;
    }
    int mid = lo + (hi - lo) / 2;
    int left = nodes[cur].left, right = nodes[cur].right;
    if (target <= mid) {
        left = update(left, lo, mid, target, d);
    } else {
        right = update(right, mid + 1, hi, target, d);
    }
    nodes.emplace_back(combine(nodes[left].value, nodes[right].value), left, right);
    return static_cast<int>(nodes.size()) - 1;
}

public:
    explicit PersistentSegTree(int n, const T &v = T()) : len(n) {
        assert(len > 0);
        roots.push_back(build(0, len - 1, [&](int) { return v; }));
    }

    template<typename It>
    PersistentSegTree(It lo, It hi) : len(hi - lo) {
        assert(len > 0);

```

```

    roots.push_back(build(0, len - 1, [&](int i) { return *(lo + i); }));
}

int size() const { return len; }
int versions() const { return static_cast<int>(roots.size()); }

T at(int version, int i) const {
    assert(0 <= version && version < versions());
    assert(0 <= i && i < len);
    return query(version, i, i);
}

T query(int version, int lo, int hi) const {
    assert(0 <= version && version < versions());
    assert(0 <= lo && lo <= hi && hi < len);
    return query(roots[version], 0, len - 1, lo, hi);
}

int update(int version, int i, const T &d) {
    assert(0 <= version && version < versions());
    assert(0 <= i && i < len);
    roots.push_back(update(roots[version], 0, len - 1, i, d));
    return static_cast<int>(roots.size()) - 1;
}
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{6, -2, 1, 8, 10};
    PersistentSegTree<int> t(a.begin(), a.end());
    int v1 = t.update(0, 2, 4);
    int v2 = t.update(0, 1, 7); // Branch from version 0 rather than version 1.

    assert(t.size() == 5);
    assert(t.versions() == 3);
    assert(t.at(0, 2) == 1);
    assert(t.at(v1, 2) == 4);
    assert(t.query(0, 0, 3) == -2);
    assert(t.query(v1, 0, 3) == -2);
    assert(t.query(v2, 0, 3) == 1);
    return 0;
}

```

2.4.12 Merge Sort Tree

A merge sort tree is a static segment tree in which every node stores the sorted multiset of the values in its range. The nodes are exactly the ranges visited by a merge sort, and each node's sorted list is the merge of its two children's lists. This supports queries that count, within a range of indices, how many values fall below a threshold or inside a value interval: decompose the index range into $O(\log n)$ canonical nodes and binary search each node's sorted list.

The structure is built once and not modified afterward. All index ranges are inclusive $[l_0, h_i]$ with 0-based indices, and values may be of any comparable type.

- `MergeSortTree<T>(a)` builds the tree over the array `a`.
- `size()` returns the size of the array.
- `count_leq(l_0, h_i, x)` returns the number of indices `i` in $[l_0, h_i]$ with `a[i] ≤ x`.
- `count_in(l_0, h_i, x, y)` returns the number of indices `i` in $[l_0, h_i]$ such that `a[i] ∈ [x, y]`.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log^2 n)$ per call to `count_leq()` and `count_in()`.

Space Complexity:

- $O(n \log n)$ for storage of the tree.
- $O(\log n)$ auxiliary stack space per query.

```
#include <algorithm>
#include <cassert>
#include <iterator>
#include <vector>

template<typename T>
class MergeSortTree {
    int len;
    std::vector<std::vector<T>> tree;

    void build(int node, int lo, int hi, const std::vector<T> &a) {
        if (lo == hi) {
            tree[node] = {a[lo]};
            return;
        }
        int mid = lo + (hi - lo) / 2;
        build(2 * node, lo, mid, a);
        build(2 * node + 1, mid + 1, hi, a);
        std::merge(
            tree[2 * node].begin(), tree[2 * node].end(), tree[2 * node + 1].begin(),
            tree[2 * node + 1].end(), std::back_inserter(tree[node])
        );
    }

    template<typename Fn>
    int query(int node, int lo, int hi, int tgt_lo, int tgt_hi, const Fn &count_node) const {
        if (tgt_hi < lo || hi < tgt_lo) {
            return 0;
        }
        if (tgt_lo <= lo && hi <= tgt_hi) {
            return count_node(tree[node]);
        }
        int mid = lo + (hi - lo) / 2;
        return query(2 * node, lo, mid, tgt_lo, tgt_hi, count_node) +
            query(2 * node + 1, mid + 1, hi, tgt_lo, tgt_hi, count_node);
    }

public:
    MergeSortTree(const std::vector<T> &a) : len(a.size()), tree(4 * a.size()) {
        assert(len > 0);
        build(1, 0, len - 1, a);
    }

    int size() const { return len; }

    int count_leq(int lo, int hi, const T &x) const {
        assert(0 <= lo && lo <= hi && hi < len);
        return query(1, 0, len - 1, lo, hi, [&](const std::vector<T> &v) {
            return static_cast<int>(std::upper_bound(v.begin(), v.end(), x) - v.begin());
        });
    }

    int count_in(int lo, int hi, const T &x, const T &y) const {
        assert(0 <= lo && lo <= hi && hi < len);
        assert(!(y < x));
        return query(1, 0, len - 1, lo, hi, [&](const std::vector<T> &v) {
```

```

        return static_cast<int>(
            std::upper_bound(v.begin(), v.end(), y) - std::lower_bound(v.begin(), v.end(), x)
        );
    });
}
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{5, 2, 8, 6, 1, 9, 3};
    MergeSortTree<int> t(a);

    assert(t.size() == 7);
    assert(t.count_leq(0, 6, 5) == 4);    // 5, 2, 1, 3
    assert(t.count_leq(2, 4, 6) == 2);    // 6, 1
    assert(t.count_in(0, 6, 3, 8) == 4);   // 5, 8, 6, 3
    assert(t.count_in(1, 5, 2, 6) == 2);   // 2, 6 (8, 1, 9 excluded)
    return 0;
}

```

2.4.13 Wavelet Tree

A wavelet tree is a static structure over an integer array with values in a known range $[\text{min_val}, \text{max_val}]$. The root splits the value range at its midpoint, recording for every prefix of the array how many of its elements belong to the lower value half, then stably partitions the array so that those elements form the left child and the rest form the right child. Recursing on each half builds a balanced tree of height $O(\log s)$, where s is the size of the value range.

The recorded prefix counts let a query at any node translate a range of array positions into the corresponding range of positions in either child, so a single root-to-leaf descent answers order statistics and rank queries. This is strictly more powerful than a merge sort tree: besides counting values below a threshold, it reports the k -th smallest value of a range in $O(\log s)$ time.

All position ranges use 0-based indices and closed intervals $[\text{lo}, \text{hi}]$. If the values are large or sparse, compress them to a small contiguous range first.

- `WaveletTree(a, min_val, max_val)` builds the tree over the array `a`, whose values must all lie in $[\text{min_val}, \text{max_val}]$.
- `size()` returns the size of the array.
- `kth_smallest(lo, hi, k)` returns the k -th smallest value among positions $[\text{lo}, \text{hi}]$, where k is 1-based (so $k == 1$ returns the minimum).
- `count_leq(lo, hi, x)` returns the number of positions $i \in [\text{lo}, \text{hi}]$ such that $a[i] \leq x$.
- `count_in(lo, hi, x, y)` returns the number of positions $i \in [\text{lo}, \text{hi}]$ such that $a[i] \in [x, y]$.

Time Complexity:

- $O(n \log s)$ per call to the constructor, where n is the size of the array and s is the size of the value range.
- $O(1)$ per call to `size()`.
- $O(\log s)$ per call to `kth_smallest()`, `count_leq()`, and `count_in()`.

Space Complexity:

- $O(n \log s)$ for storage of the tree.
- $O(\log s)$ auxiliary stack space per query.

```

#include <algorithm>
#include <cassert>

```

```

#include <cstdint>
#include <limits>
#include <memory>
#include <vector>

class WaveletTree {
    int min_val, max_val; // The value range this node covers.
    std::unique_ptr<WaveletTree> left, right;
    std::vector<int> b; // b[i] = number of the first i elements that go to the left child.

    WaveletTree(std::vector<int>::iterator lo, std::vector<int>::iterator hi, int x, int y)
        : min_val(x), max_val(y) {
        assert(min_val <= max_val);
        b.reserve((hi - lo) + 1);
        b.push_back(0);
        if (lo >= hi) {
            return;
        }
        int mid = min_val + static_cast<int>((static_cast<int64_t>(max_val) - min_val) / 2);
        auto go_left = [mid](int v) { return v <= mid; };
        for (auto it = lo; it != hi; ++it) {
            assert(min_val <= *it && *it <= max_val);
            b.push_back(b.back() + (go_left(*it) ? 1 : 0));
        }
        if (min_val == max_val) {
            return;
        }
        auto pivot = std::stable_partition(lo, hi, go_left);
        left.reset(new WaveletTree(lo, pivot, min_val, mid));
        right.reset(new WaveletTree(pivot, hi, mid + 1, max_val));
    }

    // 1-based positions [lo, hi] for the recursion below.
    int kth(int lo, int hi, int k) const {
        if (min_val == max_val) {
            return min_val;
        }
        int in_left = b[hi] - b[lo - 1];
        if (k <= in_left) {
            return left->kth(b[lo - 1] + 1, b[hi], k);
        }
        return right->kth(lo - b[lo - 1], hi - b[hi], k - in_left);
    }

    int leq(int lo, int hi, int x) const {
        if (lo > hi || x < min_val) {
            return 0;
        }
        if (max_val <= x) {
            return hi - lo + 1;
        }
        return left->leq(b[lo - 1] + 1, b[hi], x) + right->leq(lo - b[lo - 1], hi - b[hi], x);
    }

public:
    WaveletTree(std::vector<int> a, int min_val, int max_val)
        : WaveletTree(a.begin(), a.end(), min_val, max_val) {}

    int size() const { return static_cast<int>(b.size()) - 1; }

    int kth_smallest(int lo, int hi, int k) const {
        assert(0 <= lo && lo <= hi && hi < size());
        assert(1 <= k && k <= hi - lo + 1);
        return kth(lo + 1, hi + 1, k);
    }

    int count_leq(int lo, int hi, int x) const {
        assert(0 <= lo && lo <= hi && hi < size());

```

```

    return leq(lo + 1, hi + 1, x);
}

int count_in(int lo, int hi, int x, int y) const {
    assert(0 <= lo && lo <= hi && hi < size());
    assert(x <= y);
    int below_x = x == std::numeric_limits<int>::lowest() ? 0 : leq(lo + 1, hi + 1, x - 1);
    return leq(lo + 1, hi + 1, y) - below_x;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> a{5, 2, 8, 6, 1, 9, 3};
    WaveletTree t(a, 1, 9);

    assert(t.size() == 7);
    assert(t.kth_smallest(0, 6, 1) == 1); // Minimum of the whole array.
    assert(t.kth_smallest(0, 6, 4) == 5); // Median of the whole array.
    assert(t.kth_smallest(2, 4, 2) == 6); // {8, 6, 1} -> 6 is 2nd smallest.

    assert(t.count_leq(0, 6, 5) == 4); // 5, 2, 1, 3
    assert(t.count_in(0, 6, 3, 8) == 4); // 5, 8, 6, 3
    assert(t.count_in(1, 5, 2, 6) == 2); // 2, 6

    vector<int> extremes{numeric_limits<int>::lowest(), 0, numeric_limits<int>::max()};
    WaveletTree extreme_tree(extremes, numeric_limits<int>::lowest(), numeric_limits<int>::max());
    assert(extreme_tree.count_in(0, 2, numeric_limits<int>::lowest(), 0) == 2);
    return 0;
}

```

2.4.14 Cartesian Tree

Build the min-Cartesian tree of an array in linear time. A Cartesian tree is a binary tree whose inorder traversal is the original array order and whose heap property is determined by array values. The comparator `comp` defines the value ordering and defaults to `std::less<>`, producing a min-Cartesian tree in which each parent precedes or is equivalent to its children. Each subtree therefore represents a contiguous array interval whose first value in comparator order is at the subtree root. Equivalent values are broken by position, so the earlier one becomes an ancestor of the later one.

With the default comparator, this gives a direct connection to range minimum queries: the minimum index in $[lo, hi]$ is the lowest common ancestor of indices `lo` and `hi`. The included search finds that node without extra LCA preprocessing. Starting at the root, it moves right while the current index is left of the range and left while it is right of the range; the first index inside the range is its minimum. This takes $O(h)$ time for tree height h , which can be $O(n)$ for a skewed tree. For many worst-case-fast queries, preprocess the tree for LCA or use the sparse table in 2.4.1 instead. The explicit tree is also useful for interval divide-and-conquer and tree DP; the largest-rectangle application is implemented directly with a monotone stack in 1.2.4.

- `CartesianTree(a, comp = std::less<>())` constructs the tree for array `a`.
- `size()` returns the size of the array.
- `empty()` returns whether the array is empty.
- `root()` returns the root index, or `-1` if the array is empty.
- `parent(i)`, `left(i)`, and `right(i)` return neighboring node indices, or `-1` if absent.
- `range_min_index(lo, hi)` returns the index of the first value in comparator order within $[lo, hi]$.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the array size.
- $O(1)$ per call to `size()`, `empty()`, `root()`, `parent()`, `left()`, and `right()`.
- $O(h)$ per call to `range_min_index()`, where h is the tree height and may be $O(n)$.

Space Complexity: $O(n)$ for tree arrays and the monotone stack.

```
#include <cassert>
#include <functional>
#include <vector>

class CartesianTree {
    int root_;
    std::vector<int> parent_, left_, right_;

public:
    template<typename T, typename Compare = std::less<>>
    explicit CartesianTree(const std::vector<T> &a, Compare comp = Compare())
        : root_(-1), parent_(a.size(), -1), left_(a.size(), -1), right_(a.size(), -1) {
        std::vector<int> st;
        for (int i = 0, n = static_cast<int>(a.size()); i < n; i++) {
            int last = -1;
            while (!st.empty() && comp(a[i], a[st.back()])) {
                last = st.back();
                st.pop_back();
            }
            if (!st.empty()) {
                right_[st.back()] = i;
                parent_[i] = st.back();
            }
            if (last != -1) {
                left_[i] = last;
                parent_[last] = i;
            }
            st.push_back(i);
        }
        root_ = st.empty() ? -1 : st[0];
    }

    int size() const { return static_cast<int>(parent_.size()); }
    bool empty() const { return parent_.empty(); }
    int root() const { return root_; }

    int parent(int i) const {
        assert(0 <= i && i < size());
        return parent_[i];
    }

    int left(int i) const {
        assert(0 <= i && i < size());
        return left_[i];
    }

    int right(int i) const {
        assert(0 <= i && i < size());
        return right_[i];
    }

    int range_min_index(int lo, int hi) const {
        assert(0 <= lo && lo <= hi && hi < size());
        int i = root_;
        while (i < lo || hi < i) {
            i = i < lo ? right_[i] : left_[i];
        }
        return i;
    }
};
```

Example Usage

```
using namespace std;

int main() {
    vector<int> a{3, 2, 6, 1, 9};
    CartesianTree t(a);
    assert(t.size() == 5 && !t.empty());
    assert(t.root() == 3); // The global minimum is a[3] = 1.
    assert(t.parent(1) == 3);
    assert(t.left(3) == 1);
    assert(t.right(3) == 4);
    assert(t.range_min_index(0, 2) == 1); // min({3, 2, 6}) is a[1].
    assert(t.range_min_index(2, 4) == 3); // min({6, 1, 9}) is a[3].

    vector<int> equal{4, 2, 2, 3};
    CartesianTree equal_tree(equal);
    assert(equal_tree.range_min_index(1, 3) == 1); // Earlier equal minima win ties.

    CartesianTree max_tree(a, greater<int>());
    assert(max_tree.root() == 4);
    assert(max_tree.range_min_index(0, 2) == 2); // First under greater<int>: the numeric maximum.
    return 0;
}
```

2.4.15 Implicit Treap

Maintain a dynamically-sized array using a balanced binary search tree while supporting both dynamic queries and updates of contiguous subarrays via the lazy propagation technique. A treap maintains a balanced binary tree structure by preserving the heap property on the randomly generated priority values of nodes, thereby making insertions and deletions run in $O(\log n)$ with high probability.

Array indices are implicit rather than stored as keys: subtree sizes determine each node’s current position, so inserting or erasing a value does not require renumbering later elements. The internal `split(t, left, right, i)` operation separates the first `i` values from the rest, and `merge()` concatenates two such sequences. Range operations isolate their target with two splits, modify or inspect its root, and merge the three pieces back together.

The query operation is defined by an associative aggregate function `combine(a, b)`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`.

Range updates are defined by `apply_delta(v, d, len)`, which applies an update delta `d` to an aggregate summary `v` representing `len` array values, and by `compose_deltas(old, d)`, which combines a pending older delta with a newer delta in that order. These functions do not support arbitrary combinations: applying a delta to a combined segment must be equivalent to applying it to each child segment and then combining the results, and composed deltas must be equivalent to performing their updates sequentially. The default code below defines range assignment. For range increment, `compose_deltas(old, d)` should return `old + d`; `apply_delta(v, d, len)` should return `v + d` for range-min/range-max queries, and `v + d * len` for range-sum queries.

Range reversal is also propagated lazily. Since `combine()` may be order-sensitive, each node stores the aggregate of its subtree in both forward and reverse order. Reversing a subtree then swaps its children and these two aggregates before marking its descendants for later reversal.

- `ImplicitTreap<T>(n = 0, v = T())` constructs an array of size `n` with indices $[0, n)$, with all values initialized to `v`.
- `ImplicitTreap<T>(lo, hi)` constructs an array from the half-open iterator range $[lo, hi)$.
- `size()` returns the size of the array.
- `empty()` returns whether the array is empty.
- `at(i)` returns the value at index `i`.
- `query(lo, hi)` returns the result of `combine()` applied to all indices in $[lo, hi]$.

- `update(lo, hi, d)` applies the delta `d` to every index in `[lo, hi]`.
- `update(i, d)` applies the delta `d` to the single index `i`.
- `reverse(lo, hi)` reverses the order of the values in `[lo, hi]`.
- `insert(i, v)` inserts a new value `v` before index `i`, shifting later elements one position right.
- `erase(i)` removes the element at index `i`, shifting later elements one position left.
- `push_back(v)` appends a new value `v` to the end of the array.
- `pop_back()` removes the last element of the array.

The query and update operations match those of the point- and range-update segment trees in this section;

`insert()`, `erase()`, `push_back()`, and `pop_back()` are additionally analogous to those of `std::vector` (here, `insert()` and `erase()` take an index instead of an iterator).

Time Complexity:

- $O(n \log n)$ on average per call to both constructors, where n is the size of the array.
- $O(1)$ per call to `size()` and `empty()`.
- $O(\log n)$ on average per call to all other operations.

Space Complexity:

- $O(n)$ for storage of the array elements.
- $O(1)$ auxiliary for `size()` and `empty()`.
- $O(\log n)$ auxiliary stack space for all other operations.

```
#include <cassert>
#include <cstdlib>
#include <utility>

template<typename T>
class ImplicitTreap {
    static T combine(const T &a, const T &b) { return a < b ? a : b; }
    static T apply_delta(const T &v, const T &d, int64_t len) { return d; }
    static T compose_deltas(const T &d1, const T &d2) { return d2; }

    struct Node {
        static uint32_t rand32() {
            static uint32_t x = 123456789;
            x ^= x << 13;
            x ^= x >> 17;
            x ^= x << 5;
            return x;
        }

        T value, subtree_value, reverse_value, delta;
        bool pending, reversed;
        int size;
        uint32_t priority;
        Node *left, *right;

        explicit Node(const T &v)
            : value(v),
              subtree_value(v),
              reverse_value(v),
              pending(false),
              reversed(false),
              size(1),
              priority(rand32()),
              left(nullptr),
              right(nullptr) {}
    } *root;

    static int size(Node *n) { return (n == nullptr) ? 0 : n->size; }

    static void update_value(Node *n) {
        if (n == nullptr) {
            return;
        }
    }
}
```

```

    }
    n->subtree_value = n->left != nullptr ? combine(n->left->subtree_value, n->value) : n->value;
    if (n->right != nullptr) {
        n->subtree_value = combine(n->subtree_value, n->right->subtree_value);
    }
    n->reverse_value = n->right != nullptr ? combine(n->right->reverse_value, n->value) : n->value;
    if (n->left != nullptr) {
        n->reverse_value = combine(n->reverse_value, n->left->reverse_value);
    }
    n->size = 1 + size(n->left) + size(n->right);
}

static void update_delta(Node *n, const T &d) {
    if (n != nullptr) {
        n->value = apply_delta(n->value, d, 1);
        n->subtree_value = apply_delta(n->subtree_value, d, n->size);
        n->reverse_value = apply_delta(n->reverse_value, d, n->size);
        n->delta = n->pending ? compose_deltas(n->delta, d) : d;
        n->pending = true;
    }
}

static void reverse_subtree(Node *n) {
    if (n != nullptr) {
        std::swap(n->left, n->right);
        std::swap(n->subtree_value, n->reverse_value);
        n->reversed = !n->reversed;
    }
}

static void push(Node *n) {
    if (n == nullptr) {
        return;
    }
    if (n->pending) {
        update_delta(n->left, n->delta);
        update_delta(n->right, n->delta);
        n->pending = false;
    }
    if (n->reversed) {
        reverse_subtree(n->left);
        reverse_subtree(n->right);
        n->reversed = false;
    }
}

static void merge(Node *&n, Node *left, Node *right) {
    push(left);
    push(right);
    if (left == nullptr) {
        n = right;
    } else if (right == nullptr) {
        n = left;
    } else if (left->priority < right->priority) {
        merge(left->right, left->right, right);
        n = left;
    } else {
        merge(right->left, left, right->left);
        n = right;
    }
    update_value(n);
}

static void split(Node *n, Node *&left, Node *&right, int i) {
    push(n);
    if (n == nullptr) {
        left = right = nullptr;
    } else if (i <= size(n->left)) {

```



```

        split(n->left, left, n->left, i);
        right = n;
    } else {
        split(n->right, n->right, right, i - size(n->left) - 1);
        left = n;
    }
    update_value(n);
}

static void insert(Node *&n, Node *new_node, int i) {
    push(n);
    if (n == nullptr) {
        n = new_node;
    } else if (new_node->priority < n->priority) {
        split(n, new_node->left, new_node->right, i);
        n = new_node;
    } else if (i <= size(n->left)) {
        insert(n->left, new_node, i);
    } else {
        insert(n->right, new_node, i - size(n->left) - 1);
    }
    update_value(n);
}

static void erase(Node *&n, int i) {
    assert(n != nullptr);
    push(n);
    if (i == size(n->left)) {
        Node *left = n->left, *right = n->right;
        delete n;
        merge(n, left, right);
    } else if (i < size(n->left)) {
        erase(n->left, i);
    } else {
        erase(n->right, i - size(n->left) - 1);
    }
    update_value(n);
}

static Node *select(Node *n, int i) {
    assert(n != nullptr);
    push(n);
    if (i < size(n->left)) {
        return select(n->left, i);
    }
    if (i > size(n->left)) {
        return select(n->right, i - size(n->left) - 1);
    }
    return n;
}

void clean_up(Node *&n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
    explicit ImplicitTreap(int n = 0, const T &v = T()) : root(nullptr) {
        assert(n >= 0);
        for (int i = 0; i < n; i++) {
            push_back(v);
        }
    }

    template<typename It>

```

```

ImplicitTreap(It lo, It hi) : root(nullptr) {
    for (; lo != hi; ++lo) {
        push_back(*lo);
    }
}

~ImplicitTreap() { clean_up(root); }
ImplicitTreap(const ImplicitTreap &) = delete;
ImplicitTreap &operator=(const ImplicitTreap &) = delete;
int size() const { return size(root); }
bool empty() const { return root == nullptr; }

void insert(int i, const T &v) {
    assert(0 <= i && i <= size());
    insert(root, new Node(v), i);
}

void erase(int i) {
    assert(0 <= i && i < size());
    erase(root, i);
}

void push_back(const T &v) { insert(size(), v); }

void pop_back() {
    assert(!empty());
    erase(size() - 1);
}

T at(int i) const {
    assert(0 <= i && i < size());
    return select(root, i)->value;
}

T query(int lo, int hi) {
    assert(0 <= lo && lo <= hi && hi < size());
    Node *l1, *r1, *l2, *r2, *t;
    split(root, l1, r1, hi + 1);
    split(l1, l2, r2, lo);
    T res = r2->subtree_value;
    merge(t, l2, r2);
    merge(root, t, r1);
    return res;
}

void update(int lo, int hi, const T &d) {
    assert(0 <= lo && lo <= hi && hi < size());
    Node *l1, *r1, *l2, *r2, *t;
    split(root, l1, r1, hi + 1);
    split(l1, l2, r2, lo);
    update_delta(r2, d);
    merge(t, l2, r2);
    merge(root, t, r1);
}

void update(int i, const T &d) { update(i, i, d); }

void reverse(int lo, int hi) {
    assert(0 <= lo && lo <= hi && hi < size());
    Node *l1, *r1, *l2, *r2, *t;
    split(root, l1, r1, hi + 1);
    split(l1, l2, r2, lo);
    reverse_subtree(r2);
    merge(t, l2, r2);
    merge(root, t, r1);
}
};

```

Example Usage

```
#include <cassert>
#include <vector>
using namespace std;

vector<int> values(ImplicitTreap<int> &t) {
    vector<int> result;
    for (int i = 0; i < t.size(); i++) {
        result.push_back(t.at(i));
    }
    return result;
}

int main() {
    vector<int> a{99, -2, 1, 8, 10};
    ImplicitTreap<int> t(a.begin(), a.end());

    // Append 11, then append and remove 12.
    t.push_back(11);
    t.push_back(12);
    t.pop_back();
    assert(t.size() == 6);
    assert(t.at(5) == 11);
    assert(t.query(0, t.size() - 1) == -2);
    assert((values(t) == vector<int>{99, -2, 1, 8, 10, 11}));

    // Replace the first value by inserting 90 before it and erasing the old 99.
    t.insert(0, 90);
    t.erase(1);
    assert(t.at(0) == 90);
    assert(t.at(1) == -2);
    assert((values(t) == vector<int>{90, -2, 1, 8, 10, 11}));

    // Assign 2 to the first two values.
    t.update(0, 1, 2);
    assert(t.at(0) == 2);
    assert(t.at(1) == 2);
    assert(t.query(0, t.size() - 1) == 1);
    assert((values(t) == vector<int>{2, 2, 1, 8, 10, 11}));

    // Reverse the middle range: {2, 2, 1, 8, 10, 11} becomes {2, 10, 8, 1, 2, 11}.
    t.reverse(1, 4);
    assert(t.at(0) == 2);
    assert(t.at(1) == 10);
    assert(t.at(4) == 2);
    assert(t.at(5) == 11);
    assert(t.query(0, t.size() - 1) == 1);
    assert((values(t) == vector<int>{2, 10, 8, 1, 2, 11}));
    return 0;
}
```

2.5 Range Queries in Two Dimensions

2.5.1 2D Segment Tree

Maintain a dense two-dimensional array while supporting point assignments and rectangle queries. A 2D segment tree applies the ordinary segment-tree decomposition independently to rows and columns: each row-tree node contains a complete segment tree over its columns, so a rectangle decomposes into $O(\log(R) \cdot \log(C))$ canonical blocks.

The query operation is defined by a commutative associative `combine(a, b)`. The default below computes the minimum value in a rectangle. For rectangle maximum queries, use `std::max`; for rectangle sums, use addition.

Point updates are defined by `apply_delta(v, d)`, which applies delta `d` to the current value `v`. The default below assigns `d`; for point increments, return `v + d` instead.

Use this when the grid is small enough for $O(R \cdot C)$ storage and queries need an aggregate other than addition. For additive rectangle sums, a 2D Fenwick tree is simpler and has smaller constants. Use the sparse 2D segment tree in the next section when the coordinate range is too large to allocate dense storage.

- `SegTree2D<T>(rows, cols, value = T())` constructs a `rows` by `cols` array with all entries initialized to `value`.
- `SegTree2D<T>(a)` constructs the tree from the matrix `a`.
- `num_rows()` and `num_cols()` return the dimensions of the array.
- `update(r, c, d)` applies delta `d` to the entry at index `(r, c)`.
- `at(r, c)` returns the value at index `(r, c)`.
- `query(r1, c1, r2, c2)` returns the aggregate over the rectangle with rows in `[r1, r2]` and columns in `[c1, c2]`.

Time Complexity:

- $O(R \cdot C)$ per call to either constructor, where R and C are the matrix dimensions.
- $O(\log(R) \cdot \log(C))$ per call to `update()` and `query()`.
- $O(1)$ per call to `num_rows()`, `num_cols()`, and `at()`.

Space Complexity:

- $O(R \cdot C)$ for storage of the segment tree.
- $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cassert>
#include <optional>
#include <vector>

template<typename T>
class SegTree2D {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d) { return d; }

    int rows, cols;
    std::vector<std::vector<T>> tree;

    template<typename Gen>
    void build(const Gen &gen) {
        tree.assign(2 * rows, std::vector<T>(2 * cols));
        for (int r = 0; r < rows; r++) {
            for (int c = 0; c < cols; c++) {
                tree[r + rows][c + cols] = gen(r, c);
            }
        }
        for (int r = rows; r < 2 * rows; r++) {
            for (int c = cols - 1; c > 0; c--) {
                tree[r][c] = combine(tree[r][2 * c], tree[r][2 * c + 1]);
            }
        }
        for (int r = rows - 1; r > 0; r--) {
            for (int c = 1; c < 2 * cols; c++) {
                tree[r][c] = combine(tree[2 * r][c], tree[2 * r + 1][c]);
            }
        }
    }

    T query_columns(int r, int c1, int c2) const {
        std::optional<T> res;
        for (c1 += cols, c2 += cols + 1; c1 < c2; c1 /= 2, c2 /= 2) {
            if ((c1 & 1) != 0) {
                res = res ? combine(*res, tree[r][c1++]) : tree[r][c1++];
            }
        }
    }
};
```

```

    }
    if ((c2 & 1) != 0) {
        res = res ? combine(*res, tree[r][--c2]) : tree[r][--c2];
    }
}
return *res;
}

public:
SegTree2D(int rows, int cols, const T &value = T()) : rows(rows), cols(cols) {
    assert(rows > 0 && cols > 0);
    build([&](int, int) { return value; });
}

explicit SegTree2D(const std::vector<std::vector<T>> &a)
    : rows(static_cast<int>(a.size())), cols(rows == 0 ? 0 : static_cast<int>(a[0].size())) {
    assert(rows > 0 && cols > 0);
    build([&](int r, int c) { return a[r][c]; });
}

int num_rows() const { return rows; }
int num_cols() const { return cols; }

T at(int r, int c) const {
    assert(0 <= r && r < rows && 0 <= c && c < cols);
    return tree[r + rows][c + cols];
}

void update(int r, int c, const T &d) {
    assert(0 <= r && r < rows && 0 <= c && c < cols);
    int row = r + rows, col = c + cols;
    tree[row][col] = apply_delta(tree[row][col], d);
    for (int j = col; j > 1; j /= 2) {
        tree[row][j / 2] = combine(tree[row][j], tree[row][j ^ 1]);
    }
    for (row /= 2; row > 0; row /= 2) {
        tree[row][col] = combine(tree[2 * row][col], tree[2 * row + 1][col]);
        for (int j = col; j > 1; j /= 2) {
            tree[row][j / 2] = combine(tree[row][j], tree[row][j ^ 1]);
        }
    }
}

T query(int r1, int c1, int r2, int c2) const {
    assert(0 <= r1 && r1 <= r2 && r2 < rows);
    assert(0 <= c1 && c1 <= c2 && c2 < cols);
    std::optional<T> res;
    for (r1 += rows, r2 += rows + 1; r1 < r2; r1 /= 2, r2 /= 2) {
        if ((r1 & 1) != 0) {
            T value = query_columns(r1++, c1, c2);
            res = res ? combine(*res, value) : value;
        }
        if ((r2 & 1) != 0) {
            T value = query_columns(--r2, c1, c2);
            res = res ? combine(*res, value) : value;
        }
    }
    return *res;
}
};

```

Example Usage

```

using namespace std;

int main() {

```

```

SegTree2D<int> tree(3, 4, 100);
tree.update(0, 1, 8);
tree.update(1, 2, 3);
tree.update(2, 0, 5);
assert(tree.num_rows() == 3 && tree.num_cols() == 4);
assert(tree.at(0, 1) == 8);
assert(tree.query(0, 0, 1, 3) == 3);
assert(tree.query(2, 0, 2, 3) == 5);

SegTree2D<int> from_matrix(vector<vector<int>>>{{7, 2, 6}, {4, 9, 1}});
assert(from_matrix.query(0, 0, 1, 1) == 2);
assert(from_matrix.query(0, 1, 1, 2) == 1);
return 0;
}

```

2.5.2 Sparse 2D Segment Tree

Maintain a two-dimensional array over a huge grid while supporting dynamic queries of rectangular subarrays and dynamic updates of individual indices. This is a sparse (a.k.a. dynamic or implicit) 2D segment tree: row and column nodes are allocated lazily as cells are touched, so large coordinate bounds are supported without allocating the full grid.

The query operation is defined by a commutative associative aggregate function `combine(a, b)`. Because untouched regions are implicit, `repeat_value(v, area)` must return the aggregate summary of `area` copies of the initial value `v`. The default code below assumes a numerical array type, defining queries for the “min” of the target range. For rectangle-sum queries, `combine(a, b)` should return `a + b` and `repeat_value(v, area)` should return `v * area`.

The point update operation is defined by `apply_delta(v, d)`, which returns the new value at one updated cell. The default code below defines updates that “set” the chosen cell to a new value. Another possible update operation is “increment”, in which case `apply_delta(v, d)` should return `v + d`.

Rows and columns are split independently, so every rectangle query decomposes into $O(\log(R) \cdot \log(C))$ canonical rectangles regardless of whether it is thin, off-center, or otherwise adversarially placed.

Use the dense 2D segment tree in the previous section when the grid is modest enough for $O(R \cdot C)$ storage; it is simpler and faster. Use this sparse version when the coordinate range is huge but only a small fraction of cells are updated. For dense additive rectangle sums, prefer the 2D Fenwick tree in 2.6.5.

- `SparseSegTree2D<T, R, C>(v = T())` constructs a two-dimensional array with rows $[0, R]$ and columns $[0, C]$. All array values are implicitly initialized to `v`. Nodes are allocated lazily as indices are touched.
- `at(r, c)` returns the value at row `r`, column `c`.
- `query(r1, c1, r2, c2)` returns the result of `combine()` applied to every value in the rectangular region consisting of rows in $[r1, r2]$ and columns in $[c1, c2]$.
- `update(r, c, d)` assigns the value `v` at (r, c) to `apply_delta(v, d)`.

Time Complexity:

- $O(1)$ per call to the constructor.
- $O(\log(R) \cdot \log(C))$ per call to `at()`, `query()`, and `update()`.

Space Complexity:

- $O(n \cdot \log(R) \cdot \log(C))$ for storage after n point updates.
- $O(\log(R) + \log(C))$ auxiliary stack space for `at()`, `query()`, and `update()`.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <optional>

template<typename T, int R = 1000000000, int C = 1000000000>

```

```

class SparseSegTree2D {
    static_assert(R >= 0 && C >= 0);

    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T repeat_value(const T &v, int64_t area) { return v; }
    static T apply_delta(const T &v, const T &d) { return d; }

    struct InnerNode {
        T value;
        int low, high;
        InnerNode *left, *right;

        InnerNode(int lo, int hi, const T &v)
            : value(v), low(lo), high(hi), left(nullptr), right(nullptr) {}
    };

    struct OuterNode {
        InnerNode root;
        int low, high;
        OuterNode *left, *right;

        OuterNode(int lo, int hi, const T &v)
            : root(0, C, v), low(lo), high(hi), left(nullptr), right(nullptr) {}
    } *root;

    T init;

    // Helper variables for query() and update().
    int tgt_c1, tgt_c2;
    int64_t width;

    static int64_t length(int lo, int hi) { return hi - lo + 1LL; }

    void append_result(std::optional<T> &res, const T &v) { res = res ? combine(*res, v) : v; }

    T query(InnerNode *n, int qlo, int qhi, int64_t rows) {
        if (n == nullptr) {
            return repeat_value(init, rows * length(qlo, qhi));
        }
        int lo = n->low, hi = n->high, mid = lo + (hi - lo) / 2;
        std::optional<T> res;
        if (qlo < lo) {
            int missing_hi = std::min(qhi, lo - 1);
            append_result(res, repeat_value(init, rows * length(qlo, missing_hi)));
        }
        int overlap_lo = std::max(qlo, lo), overlap_hi = std::min(qhi, hi);
        if (overlap_lo <= overlap_hi) {
            if (overlap_lo == lo && overlap_hi == hi) {
                append_result(res, n->value);
            } else {
                if (overlap_lo <= mid) {
                    append_result(res, query(n->left, overlap_lo, std::min(overlap_hi, mid), rows));
                }
                if (mid < overlap_hi) {
                    append_result(res, query(n->right, std::max(overlap_lo, mid + 1), overlap_hi, rows));
                }
            }
        }
        if (hi < qhi) {
            int missing_lo = std::max(qlo, hi + 1);
            append_result(res, repeat_value(init, rows * length(missing_lo, qhi)));
        }
        return *res;
    }

    T query(OuterNode *n, int qlo, int qhi) {
        if (n == nullptr) {
            return repeat_value(init, width * length(qlo, qhi));
        }
    }
}

```

```

}
int lo = n->low, hi = n->high, mid = lo + (hi - lo) / 2;
std::optional<T> res;
if (qlo < lo) {
    int missing_hi = std::min(qhi, lo - 1);
    append_result(res, repeat_value(init, width * length(qlo, missing_hi)));
}
int overlap_lo = std::max(qlo, lo), overlap_hi = std::min(qhi, hi);
if (overlap_lo <= overlap_hi) {
    if (overlap_lo == lo && overlap_hi == hi) {
        append_result(res, query(&(n->root), tgt_c1, tgt_c2, length(lo, hi)));
    } else {
        if (overlap_lo <= mid) {
            append_result(res, query(n->left, overlap_lo, std::min(overlap_hi, mid)));
        }
        if (mid < overlap_hi) {
            append_result(res, query(n->right, std::max(overlap_lo, mid + 1), overlap_hi));
        }
    }
}
}
if (hi < qhi) {
    int missing_lo = std::max(qlo, hi + 1);
    append_result(res, repeat_value(init, width * length(missing_lo, qhi)));
}
return *res;
}

void update(InnerNode *n, int c, const T &d, bool leaf_row, int64_t rows) {
    int lo = n->low, hi = n->high, mid = lo + (hi - lo) / 2;
    if (lo == hi) {
        if (leaf_row) {
            n->value = apply_delta(n->value, d);
        } else {
            n->value = d;
        }
        return;
    }
    InnerNode *&target = (c <= mid) ? n->left : n->right;
    if (target == nullptr) {
        target = new InnerNode(c, c, repeat_value(init, rows));
    }
    if (target->low <= c && c <= target->high) {
        update(target, c, d, leaf_row, rows);
    } else {
        int split_lo = lo, split_hi = hi, split_mid = mid;
        do {
            if (c <= split_mid) {
                split_hi = split_mid;
            } else {
                split_lo = split_mid + 1;
            }
            split_mid = split_lo + (split_hi - split_lo) / 2;
        } while ((c <= split_mid) == (target->low <= split_mid));
        InnerNode *tmp =
            new InnerNode(split_lo, split_hi, repeat_value(init, rows * length(split_lo, split_hi)));
        if (target->low <= split_mid) {
            tmp->left = target;
        } else {
            tmp->right = target;
        }
        target = tmp;
        update(tmp, c, d, leaf_row, rows);
    }
    T left_value =
        (n->left != nullptr) ? n->left->value : repeat_value(init, rows * length(lo, mid));
    T right_value =
        (n->right != nullptr) ? n->right->value : repeat_value(init, rows * length(mid + 1, hi));
    n->value = combine(left_value, right_value);
}

```



```

}

void update(OuterNode *n, int r, int c, const T &d) {
    int lo = n->low, hi = n->high, mid = lo + (hi - lo) / 2;
    int64_t rows = length(lo, hi);
    if (lo == hi) {
        update(&(n->root), c, d, true, 1);
        return;
    }
    if (r <= mid) {
        if (n->left == nullptr) {
            n->left = new OuterNode(lo, mid, repeat_value(init, length(lo, mid) * length(0, C)));
        }
        update(n->left, r, c, d);
    } else {
        if (n->right == nullptr) {
            n->right =
                new OuterNode(mid + 1, hi, repeat_value(init, length(mid + 1, hi) * length(0, C)));
        }
        update(n->right, r, c, d);
    }
    T value = repeat_value(init, rows);
    if (n->left != nullptr || n->right != nullptr) {
        T left_value = (n->left != nullptr) ? query(&(n->left->root), c, c, length(lo, mid))
            : repeat_value(init, length(lo, mid));
        T right_value = (n->right != nullptr) ? query(&(n->right->root), c, c, length(mid + 1, hi))
            : repeat_value(init, length(mid + 1, hi));
        value = combine(left_value, right_value);
    }
    update(&(n->root), c, value, false, rows);
}

static void clean_up(InnerNode *n) {
    if (n != nullptr) {
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

static void clean_up(OuterNode *n) {
    if (n != nullptr) {
        clean_up(n->root.left);
        clean_up(n->root.right);
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }
}

public:
    explicit SparseSegTree2D(const T &v = T())
        : root(new OuterNode(0, R, repeat_value(v, length(0, R) * length(0, C))), init(v) {}

    ~SparseSegTree2D() { clean_up(root); }
    SparseSegTree2D(const SparseSegTree2D &) = delete;
    SparseSegTree2D &operator=(const SparseSegTree2D &) = delete;

    T at(int r, int c) {
        assert(0 <= r && r <= R && 0 <= c && c <= C);
        return query(r, c, r, c);
    }

    T query(int r1, int c1, int r2, int c2) {
        assert(0 <= r1 && r1 <= r2 && r2 <= R);
        assert(0 <= c1 && c1 <= c2 && c2 <= C);
        tgt_c1 = c1;
        tgt_c2 = c2;

```

```

        width = length(c1, c2);
        return query(root, r1, r2);
    }

    void update(int r, int c, const T &d) {
        assert(0 <= r && r <= R && 0 <= c && c <= C);
        update(root, r, c, d);
    }
};

```

Example Usage

```

#include <iostream>
using namespace std;

int main() {
    SparseSegTree2D<int> t(0);
    t.update(0, 0, 7);
    t.update(0, 1, 6);
    t.update(1, 0, 5);
    t.update(1, 1, 4);
    t.update(2, 1, 1);
    t.update(2, 2, 9);
    cout << "Values:" << endl;
    for (int i = 0; i < 3; i++) {
        for (int j = 0; j < 3; j++) {
            cout << t.at(i, j) << " ";
        }
        cout << endl;
    }
    assert(t.query(0, 0, 0, 1) == 6);
    assert(t.query(0, 0, 1, 0) == 5);
    assert(t.query(1, 1, 2, 2) == 0);
    assert(t.query(0, 0, 1000000000, 1000000000) == 0);
    t.update(500000000, 500000000, -100);
    assert(t.query(0, 0, 1000000000, 1000000000) == -100);

    SparseSegTree2D<int, 0, 1> rectangular(0);
    rectangular.update(0, 0, 5);
    rectangular.update(0, 1, 6);
    assert(rectangular.query(0, 0, 0, 1) == 5);
    return 0;
}

```

Example Output

```

Values:
7 6 0
5 4 0
0 1 9

```

2.5.3 2D Range Tree

Maintain a set of two-dimensional points while supporting queries for all points that fall inside given rectangular regions. This implementation uses `std::pair` to represent points, requiring operators `<`, `<=`, and `==` to be defined on the numeric template type. A balanced tree over the points sorted by x stores at each node its subrange of points sorted by y , merged from its children's lists as in a merge sort tree. A query decomposes the x -range into $O(\log n)$ nodes and binary searches each node's list for the matching y -range.

Use this for static point-reporting queries when guaranteed worst-case bounds are more important than memory. Compared with a range k-d tree, it uses more space but gives $O(\log^2 n + m)$ query time regardless of point

distribution; the k-d tree is lighter and often faster on typical inputs, but its pruning is more distribution-dependent.

- `RangeTree<T>(lo, hi)` constructs a set of `std::pair` points from the half-open forward-iterator range `[lo, hi)`.
- `query(x1, y1, x2, y2, f)` calls the function `f(i, p)` on each point whose x -coordinate is in `[x1, x2]` and whose y -coordinate is in `[y1, y2]`. The first argument to `f` is the 0-based index of the point in the original range given to the constructor. The second argument is the point itself as an `std::pair`.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the number of points.
- $O(\log^2 n + m)$ per call to `query()`, where m is the number of points that are reported by the query.

Space Complexity:

- $O(n \log n)$ for storage of the points.
- $O(\log n)$ auxiliary stack space for `query()`.

```
#include <algorithm>
#include <cassert>
#include <iterator>
#include <utility>
#include <vector>

template<typename T>
class RangeTree {
    struct IndexedPoint {
        std::pair<T, T> value;
        int original_index;
    };

    std::vector<IndexedPoint> points;
    std::vector<std::vector<std::pair<int, T>>> columns;

    void build(int n, int lo, int hi) {
        if (points[lo].value.first == points[hi].value.first) {
            for (int i = lo; i <= hi; i++) {
                columns[n].emplace_back(i, points[i].value.second);
            }
            return;
        }
        int l = n * 2 + 1, r = n * 2 + 2, mid = lo + (hi - lo) / 2;
        build(l, lo, mid);
        build(r, mid + 1, hi);
        columns[n].resize(columns[l].size() + columns[r].size());
        std::merge(
            columns[l].begin(), columns[l].end(), columns[r].begin(), columns[r].end(),
            columns[n].begin(), [](const auto &a, const auto &b) { return a.second < b.second; }
        );
    }

    // Helper variables for query().
    T x1, y1, x2, y2;

    template<typename Fn>
    void query(int n, int lo, int hi, Fn &f) {
        if (points[hi].value.first < x1 || x2 < points[lo].value.first) {
            return;
        }
        if (!(points[lo].value.first < x1 || x2 < points[hi].value.first)) {
            if (!columns[n].empty() && !(y2 < y1)) {
                auto it = std::lower_bound(
                    columns[n].begin(), columns[n].end(), y1,
                    [](const auto &a, const T &value) { return a.second < value; }
                );
                for (; it != columns[n].end() && it->second <= y2; ++it) {
                    const IndexedPoint &p = points[it->first];
```

```

        f(p.original_index, p.value);
    }
}
} else if (lo != hi) {
    int mid = lo + (hi - lo) / 2;
    query(n * 2 + 1, lo, mid, f);
    query(n * 2 + 2, mid + 1, hi, f);
}
}

public:
template<typename It>
RangeTree(It lo, It hi) {
    int n = std::distance(lo, hi);
    assert(n > 0);
    points.reserve(n);
    int index = 0;
    for (It it = lo; it != hi; ++it) {
        points.push_back({*it, index++});
    }
    columns.resize(4 * n + 1);
    std::sort(points.begin(), points.end(), [](const IndexedPoint &a, const IndexedPoint &b) {
        return a.value < b.value;
    });
    build(0, 0, n - 1);
}

template<typename Fn>
void query(const T &x1, const T &y1, const T &x2, const T &y2, Fn f) {
    assert(!(x2 < x1) && !(y2 < y1));
    this->x1 = x1;
    this->y1 = y1;
    this->x2 = x2;
    this->y2 = y2;
    query(0, 0, static_cast<int>(points.size()) - 1, f);
}
};

```

Example Usage

```

#include <algorithm>
#include <cassert>
using namespace std;

int main() {
    vector<pair<int, int>> v{{1, 4}, {5, 4}, {2, 2}, {3, 1}, {6, -5},
                           {5, -1}, {3, -3}, {-1, -2}, {-1, -1}, {2, -1}};
    RangeTree<int> t(v.begin(), v.end());
    vector<pair<int, int>> got;
    vector<int> indices;
    auto collect = [&](int i, const pair<int, int> &p) {
        indices.push_back(i);
        got.push_back(p);
    };
    t.query(-1, -1, 2, 5, collect);
    sort(got.begin(), got.end());
    assert((got == vector<pair<int, int>>{{-1, -1}, {1, 4}, {2, -1}, {2, 2}}));
    sort(indices.begin(), indices.end());
    assert((indices == vector<int>{0, 2, 8, 9}));
    got.clear();
    indices.clear();
    t.query(1, 1, 4, 8, collect);
    sort(got.begin(), got.end());
    assert((got == vector<pair<int, int>>{{1, 4}, {2, 2}, {3, 1}}));
    return 0;
}

```

2.5.4 K-d Tree for 2D Range Query

Maintain a static set of two-dimensional points while supporting rectangle reporting queries. A k-d tree recursively splits points by alternating coordinates and stores a bounding box for each subtree, allowing whole subtrees to be accepted or pruned during a query.

This implementation uses `std::pair` to represent points, requiring `operator<` to be defined on the numeric template type.

Use this for static point-reporting queries when $O(n)$ space and good average performance are more important than a strict worst-case guarantee. Use the 2D range tree instead when adversarial point sets or query rectangles are expected and the extra $O(n \log n)$ space is acceptable.

- `RangeKdTree<T>(lo, hi)` constructs a set of `std::pair` points from the half-open forward-iterator range `[lo, hi)`.
- `query(x1, y1, x2, y2, f)` calls the function `f(p)` on each point whose x -coordinate is in `[x1, x2]` and whose y -coordinate is in `[y1, y2]`. The argument `p` is the point as an `std::pair`.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the number of points.
- $O(\sqrt{n} + m)$ per call to `query()`, where m is the number of points that are reported by the query.

Space Complexity:

- $O(n)$ for storage of the points.
- $O(\log n)$ auxiliary stack space for `query()`.

```
#include <algorithm>
#include <cassert>
#include <iterator>
#include <utility>
#include <vector>

template<typename T>
class RangeKdTree {
    std::vector<std::pair<T, T>> tree, minp, maxp;
    std::vector<int> l_index, h_index;

    void build(int lo, int hi, bool div_x) {
        if (lo >= hi) {
            return;
        }
        int mid = lo + (hi - lo) / 2;
        std::nth_element(
            tree.begin() + lo, tree.begin() + mid, tree.begin() + hi,
            [div_x](const auto &a, const auto &b) {
                return div_x ? a.first < b.first : a.second < b.second;
            }
        );
        l_index[mid] = lo;
        h_index[mid] = hi;
        minp[mid].first = maxp[mid].first = tree[lo].first;
        minp[mid].second = maxp[mid].second = tree[lo].second;
        for (int i = lo + 1; i < hi; i++) {
            minp[mid].first = std::min(minp[mid].first, tree[i].first);
            minp[mid].second = std::min(minp[mid].second, tree[i].second);
            maxp[mid].first = std::max(maxp[mid].first, tree[i].first);
            maxp[mid].second = std::max(maxp[mid].second, tree[i].second);
        }
        build(lo, mid, !div_x);
        build(mid + 1, hi, !div_x);
    }

    // Helper variables for query().
    T x1, y1, x2, y2;
```

```

template<typename Fn>
void query(int lo, int hi, Fn &f) {
    if (lo >= hi) {
        return;
    }
    int mid = lo + (hi - lo) / 2;
    T ax = minp[mid].first, ay = minp[mid].second;
    T bx = maxp[mid].first, by = maxp[mid].second;
    if (x2 < ax || bx < x1 || y2 < ay || by < y1) {
        return;
    }
    if (!(ax < x1 || x2 < bx || ay < y1 || y2 < by)) {
        for (int i = l_index[mid]; i < h_index[mid]; i++) {
            f(tree[i]);
        }
        return;
    }
    query(lo, mid, f);
    query(mid + 1, hi, f);
    if (tree[mid].first < x1 || x2 < tree[mid].first || tree[mid].second < y1 ||
        y2 < tree[mid].second) {
        return;
    }
    f(tree[mid]);
}

public:
template<typename It>
RangeKdTree(It lo, It hi) : tree(lo, hi) {
    int n = std::distance(lo, hi);
    l_index.resize(n);
    h_index.resize(n);
    minp.resize(n);
    maxp.resize(n);
    build(0, n, true);
}

template<typename Fn>
void query(const T &x1, const T &y1, const T &x2, const T &y2, Fn f) {
    assert(!(x2 < x1) && !(y2 < y1));
    this->x1 = x1;
    this->y1 = y1;
    this->x2 = x2;
    this->y2 = y2;
    query(0, static_cast<int>(tree.size()), f);
}
};

```

Example Usage

```

#include <algorithm>
#include <cassert>
using namespace std;

int main() {
    vector<pair<int, int>> v{{1, 4}, {5, 4}, {2, 2}, {3, 1}, {6, -5},
                          {5, -1}, {3, -3}, {-1, -2}, {-1, -1}, {2, -1}};
    RangeKdTree<int> t(v.begin(), v.end());
    vector<pair<int, int>> got;
    auto collect = [&](const pair<int, int> &p) { got.push_back(p); };
    t.query(-1, -1, 2, 5, collect);
    sort(got.begin(), got.end());
    assert((got == vector<pair<int, int>>{{-1, -1}, {1, 4}, {2, -1}, {2, 2}}));
    got.clear();
    t.query(1, 1, 4, 8, collect);
}

```

```

sort(got.begin(), got.end());
assert((got == vector<pair<int, int>>{{1, 4}, {2, 2}, {3, 1}}));
return 0;
}

```

2.5.5 K-d Tree for Nearest Neighbor

Maintain a static set of two-dimensional points while supporting nearest-neighbor queries. A k-d tree recursively splits points by the coordinate with larger spread, searches the more promising side first, and only explores the other side when its splitting plane can still improve the answer.

This implementation uses `std::pair` to represent points, requiring `operator<` and `long double` casting to be defined on the numeric template type.

Use this for static nearest-neighbor queries on point sets in low dimensions. It is a geometric search tree rather than an aggregate structure: for rectangle reporting use the range k-d tree or 2D range tree, and for grid cell updates or rectangle sums/minima use a Fenwick tree or segment tree.

- `NearestKDTree<T>(lo, hi)` constructs a set of `std::pair` points from the half-open forward-iterator range `[lo, hi)`.
- `nearest(x, y, can_equal = true)` returns a point in the set that is closest to `(x, y)` by Euclidean distance. This may be equal to `(x, y)` only if `can_equal` is `true`; at least one eligible point must exist.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the number of points.
- $O(\log n)$ on average per call to `nearest()`.

Space Complexity:

- $O(n)$ for storage of the points.
- $O(\log n)$ auxiliary stack space for `nearest()`.

```

#include <algorithm>
#include <cassert>
#include <cfloat>
#include <climits>
#include <iterator>
#include <utility>
#include <vector>

template<typename T>
class NearestKDTree {
    std::vector<std::pair<T, T>> tree;
    std::vector<char> div_x;

    void build(int lo, int hi) {
        if (lo >= hi) {
            return;
        }
        int mid = lo + (hi - lo) / 2;
        T minx, maxx, miny, maxy;
        minx = maxx = tree[lo].first;
        miny = maxy = tree[lo].second;
        for (int i = lo + 1; i < hi; i++) {
            minx = std::min(minx, tree[i].first);
            miny = std::min(miny, tree[i].second);
            maxx = std::max(maxx, tree[i].first);
            maxy = std::max(maxy, tree[i].second);
        }
        long double x_span = static_cast<long double>(maxx) - minx;
        long double y_span = static_cast<long double>(maxy) - miny;
        div_x[mid] = !(x_span < y_span);
    }
};

```

```

std::nth_element(
    tree.begin() + lo, tree.begin() + mid, tree.begin() + hi,
    [&](const auto &a, const auto &b) {
        return div_x[mid] ? a.first < b.first : a.second < b.second;
    }
);
if (lo + 1 == hi) {
    return;
}
build(lo, mid);
build(mid + 1, hi);
}

// Helper variables for nearest().
long double min_dist;
int id;

void nearest(int lo, int hi, const T &x, const T &y, bool can_equal) {
    if (lo >= hi) {
        return;
    }
    int mid = lo + (hi - lo) / 2;
    long double dx = static_cast<long double>(x) - tree[mid].first;
    long double dy = static_cast<long double>(y) - tree[mid].second;
    long double d = dx * dx + dy * dy;
    if (d < min_dist && (can_equal || d != 0)) {
        min_dist = d;
        id = mid;
    }
    if (lo + 1 == hi) {
        return;
    }
    d = div_x[mid] ? dx : dy;
    int l1 = lo, r1 = mid, l2 = mid + 1, r2 = hi;
    if (d > 0) {
        std::swap(l1, l2);
        std::swap(r1, r2);
    }
    nearest(l1, r1, x, y, can_equal);
    if (d * static_cast<long double>(d) < min_dist) {
        nearest(l2, r2, x, y, can_equal);
    }
}

public:
template<typename It>
NearestKDTree(It lo, It hi) : tree(lo, hi) {
    int n = std::distance(lo, hi);
    assert(n > 0);
    div_x.resize(n);
    build(0, n);
}

std::pair<T, T> nearest(const T &x, const T &y, bool can_equal = true) {
    min_dist = LDBL_MAX;
    id = -1;
    nearest(0, static_cast<int>(tree.size()), x, y, can_equal);
    assert(id != -1);
    return tree[id];
}
};

```

Example Usage

```
using namespace std;
```



```
int main() {
    vector<pair<int, int>> p{{0, 2}, {0, 3}, {-1, 0}};
    NearestKDTree<int> t(p.begin(), p.end());
    assert(t.nearest(0, 2, true) == (pair<int, int>{0, 2}));
    assert(t.nearest(0, 2, false) == (pair<int, int>{0, 3}));
    assert(t.nearest(0, 0) == (pair<int, int>{-1, 0}));
    assert(t.nearest(-10000, 0) == (pair<int, int>{-1, 0}));

    vector<pair<int, int>> extremes{{INT_MIN, 0}, {INT_MAX, 0}};
    NearestKDTree<int> extreme_tree(extremes.begin(), extremes.end());
    assert(extreme_tree.nearest(0, 0) == (pair<int, int>{INT_MAX, 0}));
    return 0;
}
```

2.6 Fenwick Trees

2.6.1 Fenwick Tree (Point Update, Range Query)

Maintain a numerical array while supporting point increments and range-sum queries. A Fenwick tree (also called a binary indexed tree) stores partial prefix sums: each internal entry `tree[i]` holds the sum of a block of indices whose length is the lowest set bit of `i`. A point update touches $O(\log n)$ covering blocks, and a prefix-sum query combines $O(\log n)$ disjoint blocks.

The value type `T` must represent 0 and support addition and subtraction.

When prefix sums are nondecreasing, `max_prefix(c)` answers prefix-boundary queries by binary lifting: it walks the implicit power-of-two block structure to find the longest prefix with sum at most `c`. This lets the same structure double as a dynamic multiset over a bounded integer domain, useful for coordinate-compressed k -th element lookups and online rank queries.

- `Fenwick<T>(n)` constructs an array with 0-based indices $[0, n)$, with all values initialized to 0.
- `size()` returns the size of the array.
- `add(i, x)` adds `x` to the value at index `i`.
- `set(i, x)` assigns the value at index `i` to `x`.
- `at(i)` returns the value at index `i`.
- `sum(hi)` returns the sum of all values at indices $[0, hi]$.
- `sum(lo, hi)` returns the sum of all values at indices $[lo, hi]$.
- `max_prefix(c)` returns the largest boundary `hi` such that `sum(0, hi - 1) ≤ c`, assuming prefix sums are nondecreasing and `c` is nonnegative. It may return any boundary in $[0, n]$, where `n` is `size()`.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log n)$ per call to `add()`, both `sum()` functions, `set()`, `at()`, and `max_prefix()`.

Space Complexity:

- $O(n)$ for storage of the Fenwick tree.
- $O(1)$ auxiliary for all operations.

```
#include <cassert>
#include <vector>

template<typename T>
class Fenwick {
    int len;
    std::vector<T> tree;
```

```

public:
    explicit Fenwick(int n) : len(n), tree(n + 1) {}

    int size() const { return len; }

    void add(int i, const T &x) {
        assert(0 <= i && i < len);
        for (++i; i <= len; i += i & -i) {
            tree[i] += x;
        }
    }

    T sum(int hi) const {
        assert(-1 <= hi && hi < len);
        T res = 0;
        for (++hi; hi > 0; hi -= hi & -hi) {
            res += tree[hi];
        }
        return res;
    }

    void set(int i, const T &x) { add(i, x - at(i)); }
    T at(int i) const { return sum(i, i); }

    T sum(int lo, int hi) const {
        assert(0 <= lo && lo <= hi && hi < len);
        return sum(hi) - sum(lo - 1);
    }

    int max_prefix(T c) const {
        int pos = 0, pw = 1;
        while (pw * 2 <= len) {
            pw *= 2;
        }
        for (; pw > 0; pw >>= 1) {
            if (pos + pw <= len && tree[pos + pw] <= c) {
                c -= tree[pos + pw];
                pos += pw;
            }
        }
        return pos;
    }
};

```

Example Usage

```

using namespace std;

int main() {
    vector<int> a{10, 1, 2, 3, 4};
    Fenwick<int> t(5);
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        t.set(i, a[i]);
    }
    t.add(0, -5); // a is now {5, 1, 2, 3, 4}.
    assert(t.at(0) == 5);
    assert(t.sum(1, 3) == 6); // 1 + 2 + 3
    assert(t.sum(4) == 15); // 5 + 1 + 2 + 3 + 4

    // With nonnegative values, the tree doubles as a frequency table for order-statistic queries.
    Fenwick<int> freq(8);
    freq.add(1, 1); // One element of value 1.
    freq.add(3, 3); // Three elements of value 3.
    freq.add(6, 1); // One element of value 6.
    assert(freq.max_prefix(0) == 1); // The longest prefix with count at most 0 is [0, 1).
    assert(freq.max_prefix(3) == 3); // Prefix [0, 3) has count 1; [0, 4) has count 4.
}

```

```

    assert(freq.max_prefix(4) == 6); // Prefix [0, 6) has count 4; [0, 7) has count 5.
    return 0;
}

```

2.6.2 Fenwick Tree (Range Update, Point Query)

Maintain a numerical array while supporting range increments and point queries. This is done by storing the difference array in a Fenwick tree: adding x to $[lo, hi]$ adds $+x$ at lo and $-x$ after hi , and the value at index i is the prefix sum of those differences through i .

The value type T must represent 0 and support addition and subtraction.

- `FenwickRUPQ<T>(n)` constructs an array with 0-based indices $[0, n)$, with all values initialized to 0.
- `size()` returns the size of the array.
- `at(i)` returns the value at index i .
- `add(i, x)` adds x to the value at index i .
- `add(lo, hi, x)` adds x to the values at all indices in $[lo, hi]$.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log n)$ per call to `at()` and both `add()` functions.

Space Complexity:

- $O(n)$ for storage of the array elements.
- $O(1)$ auxiliary for all operations.

```

#include <cassert>
#include <vector>

template<typename T>
class FenwickRUPQ {
    int len;
    std::vector<T> tree;

    void add_helper(int i, const T &x) {
        for (i++; i <= len + 1; i += i & -i) {
            tree[i] += x;
        }
    }

public:
    explicit FenwickRUPQ(int n) : len(n), tree(n + 2) {}

    int size() const { return len; }

    T at(int i) const {
        assert(0 <= i && i < len);
        T res = 0;
        for (i++; i > 0; i -= i & -i) {
            res += tree[i];
        }
        return res;
    }

    void add(int i, const T &x) {
        assert(0 <= i && i < len);
        add_helper(i, x);
    }

    void add(int lo, int hi, const T &x) {

```

```

    assert(0 <= lo && lo <= hi && hi < len);
    add_helper(lo, x);
    add_helper(hi + 1, -x);
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    FenwickRUPQ<int> t(5);
    t.add(0, 1, 5);
    t.add(1, 2, 5);
    t.add(2, 4, 10);
    assert(t.size() == 5);
    assert(t.at(0) == 5);
    assert(t.at(1) == 10);
    assert(t.at(2) == 15);
    assert(t.at(4) == 10);
    return 0;
}

```

2.6.3 Fenwick Tree (Range Update, Range Query)

Maintain a numerical array while supporting range increments and range-sum queries. This uses two Fenwick trees to recover prefix sums after difference-array updates: if the difference array stores range additions, then the prefix sum through `hi` can be written as `hi*sum(t1, hi) - sum(t2, hi)`.

The value type `T` must represent 0 and support addition, subtraction, and multiplication by an integer index.

- `FenwickRURQ<T>(n)` constructs an array with 0-based indices $[0, n)$, with all values initialized to 0.
- `size()` returns the size of the array.
- `at(i)` returns the value at index `i`.
- `add(i, x)` increments the value at index `i` by `x`.
- `add(lo, hi, x)` adds `x` to the values at all indices in $[lo, hi]$.
- `set(i, x)` assigns the value at index `i` to `x`.
- `sum(hi)` returns the sum of all values at indices $[0, hi]$.
- `sum(lo, hi)` returns the sum of all values at indices $[lo, hi]$.
- `max_prefix(c)` returns the largest boundary `hi` such that $\text{sum}(0, hi - 1) \leq c$, assuming prefix sums are nondecreasing and `c` is nonnegative. It may return any boundary in $[0, n]$, where n is `size()`.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the size of the array.
- $O(1)$ per call to `size()`.
- $O(\log n)$ per call to all other operations.

Space Complexity:

- $O(n)$ for storage of the array elements.
- $O(1)$ auxiliary for all operations.

```

#include <cassert>
#include <vector>

template<typename T>
class FenwickRURQ {
    int len;

```

```

std::vector<T> t1, t2;

T sum(const std::vector<T> &tree, int i) const {
    assert(0 <= i && i <= len);
    T res = 0;
    for (; i != 0; i -= i & -i) {
        res += tree[i];
    }
    return res;
}

void add(std::vector<T> &tree, int i, const T &x) {
    assert(1 <= i && i <= len + 1);
    for (; i <= len + 1; i += i & -i) {
        tree[i] += x;
    }
}

public:
explicit FenwickRURQ(int n) : len(n), t1(n + 2), t2(n + 2) {}

int size() const { return len; }

void add(int lo, int hi, const T &x) {
    assert(0 <= lo && lo <= hi && hi < len);
    lo++;
    hi++;
    add(t1, lo, x);
    add(t1, hi + 1, -x);
    add(t2, lo, x * (lo - 1)); // Overflow warning.
    add(t2, hi + 1, -x * hi);
}

void add(int i, const T &x) { add(i, i, x); }
void set(int i, const T &x) { add(i, x - at(i)); }

T sum(int hi) const {
    assert(-1 <= hi && hi < len);
    hi++;
    return hi * sum(t1, hi) - sum(t2, hi);
}

T sum(int lo, int hi) const {
    assert(0 <= lo && lo <= hi && hi < len);
    return sum(hi) - sum(lo - 1);
}

T at(int i) const { return sum(i, i); }

int max_prefix(T c) const {
    T s1 = 0, s2 = 0;
    int pos = 0, pw = 1;
    while (pw * 2 <= len) {
        pw *= 2;
    }
    for (; pw > 0; pw >>= 1) {
        int next = pos + pw;
        if (next <= len) {
            T ns1 = s1 + t1[next], ns2 = s2 + t2[next];
            if (next * ns1 - ns2 <= c) {
                s1 = ns1;
                s2 = ns2;
                pos = next;
            }
        }
    }
    return pos;
}

```

```
};
```

Example Usage

```
using namespace std;

int main() {
    vector<int> a{10, 1, 2, 3, 4};
    FenwickRURQ<int> t(5);
    for (int i = 0; i < t.size(); i++) {
        t.set(i, a[i]);
    }
    t.add(0, 2, 5);
    t.set(3, -5);
    vector<int> expected{15, 6, 7, -5, 4};
    for (int i = 0; i < t.size(); i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.sum(0, 4) == 27);
    FenwickRURQ<int> freq(8);
    freq.add(1, 1, 1);
    freq.add(3, 3, 3);
    freq.add(6, 6, 1);
    assert(freq.max_prefix(0) == 1);
    assert(freq.max_prefix(3) == 3);
    assert(freq.max_prefix(4) == 6);
    return 0;
}
```

2.6.4 Sparse Fenwick Tree (Range Update, Range Query)

Maintain a numerical array over a huge index range while supporting range increments and range-sum queries. This is the sparse version of the two-tree Fenwick range-update/range-query trick: only Fenwick nodes reached by previous operations are stored, using `std::unordered_map` instead of dense vectors.

The value type `T` must represent 0 and support addition, subtraction, and multiplication by an integer index.

- `SparseFenwick<T, N>()` constructs an array with 0-based indices $[0, N)$, implicitly initialized to 0.
- `at(i)` returns the value at index `i`.
- `add(i, x)` adds `x` to the value at index `i`.
- `add(lo, hi, x)` adds `x` to the values at all indices in $[lo, hi]$.
- `set(i, x)` assigns the value at index `i` to `x`.
- `sum(hi)` returns the sum of all values at indices $[0, hi]$.
- `sum(lo, hi)` returns the sum of all values at indices $[lo, hi]$.
- `max_prefix(c)` returns the largest boundary `hi` such that $\text{sum}(0, hi - 1) \leq c$, assuming prefix sums are nondecreasing and `c` is nonnegative. It may return any boundary in $[0, N]$.

Time Complexity: $O(\log N)$ expected per call to all member functions.

Space Complexity:

- $O(q \log N)$ expected for storage after q updates.
- $O(1)$ auxiliary for all operations.

```
#include <cassert>
#include <unordered_map>

template<typename T, int N = 1000000001>
class SparseFenwick {
    static_assert(0 < N && N < (1 << 30));
```

```

std::unordered_map<int, T> tmul, tadd;

T get(const std::unordered_map<int, T> &tree, int i) const {
    if (auto it = tree.find(i); it != tree.end()) {
        return it->second;
    }
    return 0;
}

void add_helper(int at, const T &mul, const T &add) {
    assert(1 <= at && at <= N + 1);
    for (int i = at; i <= N; i += i & -i) {
        tmul[i] += mul;
        tadd[i] += add;
    }
}

public:
void add(int lo, int hi, const T &x) {
    assert(0 <= lo && lo <= hi && hi < N);
    lo++;
    hi++;
    add_helper(lo, x, x * (lo - 1)); // Overflow warning.
    add_helper(hi + 1, -x, -x * hi);
}

void add(int i, const T &x) { add(i, i, x); }
void set(int i, const T &x) { add(i, x - at(i)); }

T sum(int hi) const {
    assert(-1 <= hi && hi < N);
    T mul = 0, add = 0;
    hi++;
    for (int i = hi; i > 0; i -= i & -i) {
        mul += get(tmul, i);
        add += get(tadd, i);
    }
    return mul * hi - add;
}

T sum(int lo, int hi) const {
    assert(0 <= lo && lo <= hi && hi < N);
    return sum(hi) - sum(lo - 1);
}

T at(int i) const { return sum(i, i); }

int max_prefix(T c) const {
    T mul = 0, add = 0;
    int pos = 0, pw = 1;
    while (pw * 2 <= N) {
        pw *= 2;
    }
    for (; pw > 0; pw >= 1) {
        int next = pos + pw;
        if (next <= N) {
            T nmul = mul + get(tmul, next), nadd = add + get(tadd, next);
            if (next * nmul - nadd <= c) {
                mul = nmul;
                add = nadd;
                pos = next;
            }
        }
    }
    return pos;
}
};

```

Example Usage

```
#include <vector>
using namespace std;

int main() {
    vector<int> a{10, 1, 2, 3, 4};
    SparseFenwick<long long> t;
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        t.set(i, a[i]);
    }
    t.add(0, 2, 5);
    t.set(3, -5);
    vector<int> expected{15, 6, 7, -5, 4};
    for (int i = 0; i < 5; i++) {
        assert(t.at(i) == expected[i]);
    }
    assert(t.sum(0, 4) == 27);
    t.add(500000001, 500000010, 3);
    t.add(500000011, 500000015, 5);
    t.set(500000000, 10);
    assert(t.sum(0, 1000000000) == 92);

    SparseFenwick<long long> freq;
    freq.add(1, 1, 1);
    freq.add(3, 3, 3);
    freq.add(6, 6, 1);
    assert(freq.max_prefix(0) == 1);
    assert(freq.max_prefix(3) == 3);
    assert(freq.max_prefix(4) == 6);
    return 0;
}
```

2.6.5 2D Fenwick Tree

Maintain a two-dimensional numerical array while supporting point increments and rectangle-sum queries. This is the two-dimensional form of the standard Fenwick tree: each internal entry stores a rectangular block sum, and a prefix query combines $O(\log(R) \cdot \log(C))$ disjoint blocks.

The value type `T` must represent 0 and support addition and subtraction.

Choose among the three 2D Fenwick trees by coordinate range and update style. Use this dense version when the rows and columns are small enough to allocate the full grid: it has the simplest code and the best constants, supporting point updates and rectangle sums. Use the sparse 2D Fenwick tree (next section) when the coordinates are too large to store densely and the updates arrive online; it relies on hashing and is the only one of the three that also supports rectangle updates. Use the offline 2D Fenwick tree when the coordinates are large but every updated cell is known before the queries begin, which lets it coordinate-compress instead of hashing for markedly better constants than the sparse version. For non-additive aggregates such as min/max with custom updates, use a 2D segment tree instead, since Fenwick-tree algebra relies on addition and subtraction.

- `Fenwick2D<T>(rows, cols)` constructs a `rows` by `cols` array with 0-based indices, with all values initialized to 0.
- `num_rows()` and `num_cols()` return the number of rows and columns.
- `add(r, c, x)` adds `x` to the value at index `(r, c)`.
- `set(r, c, x)` assigns `x` to the value at index `(r, c)`.
- `at(r, c)` returns the value at index `(r, c)`.
- `sum(r, c)` returns the sum of the rectangle with rows `[0, r]` and columns `[0, c]`.
- `sum(r1, c1, r2, c2)` returns the sum of the rectangle with rows `[r1, r2]` and columns `[c1, c2]`.

Time Complexity:

- $O(R \cdot C)$ per call to the constructor, where R and C are the number of rows and columns.
- $O(1)$ per call to `num_rows()` and `num_cols()`.
- $O(\log(R) \cdot \log(C))$ per call to all other operations.

Space Complexity:

- $O(R \cdot C)$ for storage of the Fenwick tree.
- $O(1)$ auxiliary for all operations.

```
#include <cassert>
#include <vector>

template<typename T>
class Fenwick2D {
    int rows, cols;
    std::vector<std::vector<T>> tree;

public:
    Fenwick2D(int rows, int cols)
        : rows(rows), cols(cols), tree(rows + 1, std::vector<T>(cols + 1)) {}

    int num_rows() const { return rows; }
    int num_cols() const { return cols; }

    void add(int r, int c, const T &x) {
        assert(0 <= r && r < rows && 0 <= c && c < cols);
        for (int i = r + 1; i <= rows; i += i & -i) {
            for (int j = c + 1; j <= cols; j += j & -j) {
                tree[i][j] += x;
            }
        }
    }

    T sum(int r, int c) const {
        assert(-1 <= r && r < rows && -1 <= c && c < cols);
        T res = 0;
        for (int i = r + 1; i > 0; i -= i & -i) {
            for (int j = c + 1; j > 0; j -= j & -j) {
                res += tree[i][j];
            }
        }
        return res;
    }

    T sum(int r1, int c1, int r2, int c2) const {
        assert(0 <= r1 && r1 <= r2 && r2 < rows);
        assert(0 <= c1 && c1 <= c2 && c2 < cols);
        return sum(r2, c2) - sum(r1 - 1, c2) - sum(r2, c1 - 1) + sum(r1 - 1, c1 - 1);
    }

    T at(int r, int c) const { return sum(r, c, r, c); }
    void set(int r, int c, const T &x) { add(r, c, x - at(r, c)); }
};
```

Example Usage

```
using namespace std;

int main() {
    Fenwick2D<int> t(3, 3);
    t.set(0, 0, 5);
    t.set(0, 1, 6);
    t.set(1, 0, 7);
    t.add(2, 2, 9);
```

```

t.add(1, 0, -4); // Value at (1, 0) is now 3.
assert(t.at(1, 0) == 3);
assert(t.sum(0, 0, 0, 1) == 11); // 5 + 6
assert(t.sum(0, 0, 1, 0) == 8);  // 5 + 3
assert(t.sum(0, 0, 2, 2) == 23); // 5 + 6 + 3 + 9
return 0;
}

```

2.6.6 Sparse 2D Fenwick Tree

Maintain a 2D numerical array over a huge index range while supporting rectangle increments and rectangle-sum queries. This is the sparse two-dimensional analogue of range-update/range-query Fenwick trees: four sparse trees store the inclusion-exclusion coefficients needed to recover any prefix rectangle sum.

The value type `T` must represent 0 and support addition, subtraction, and multiplication by integer coordinates.

Choose this for huge sparse grids when the operation is additive: point or rectangle increments and rectangle sums. It avoids allocating the dense `R*C` table of the dense 2D Fenwick tree, and unlike the dense and offline variants it supports rectangle updates, not just point updates. Prefer it over the offline 2D Fenwick tree only when updates arrive online; if every updated cell is known in advance, that version compresses coordinates instead of hashing and runs with better constants. It is less general than the sparse 2D segment tree because Fenwick-tree algebra relies on addition and subtraction.

- `SparseFenwick2D<T, R, C>()` constructs a 2D array over rows $[0, R)$ and columns $[0, C)$. All values are implicitly initialized to 0; nodes are allocated lazily as indices are touched.
- `add(r, c, x)` adds `x` to the value at index (r, c) .
- `add(r1, c1, r2, c2, x)` adds `x` to all indices in the rectangle with rows $[r1, r2]$ and columns $[c1, c2]$.
- `set(r, c, x)` assigns `x` to the value at index (r, c) .
- `sum(r, c)` returns the sum of the rectangle with rows $[0, r]$ and columns $[0, c]$.
- `sum(r1, c1, r2, c2)` returns the sum of the rectangle with rows in $[r1, r2]$ and columns in $[c1, c2]$.
- `at(r, c)` returns the value at index (r, c) .

Time Complexity: $O(\log(R) \cdot \log(C))$ expected per call to all member functions.

Space Complexity:

- $O(q \cdot \log(R) \cdot \log(C))$ expected for storage after q updates.
- $O(1)$ auxiliary for all operations.

```

#include <cassert>
#include <cstdint>
#include <unordered_map>

template<typename T, int R = 1000000001, int C = 1000000001>
class SparseFenwick2D {
    static_assert(0 < R && R < (1 << 30) && 0 < C && C < (1 << 30));

    std::unordered_map<int64_t, T> t1, t2, t3, t4;

    template<typename Map>
    static T get(const Map &tree, int r, int c) {
        auto it = tree.find(static_cast<int64_t>(r) * (C + 1) + c);
        return it == tree.end() ? T() : it->second;
    }

    template<typename Map>
    static void add(Map &tree, int r, int c, const T &x) {
        assert(0 <= r && r <= R && 0 <= c && c <= C);
        for (int i = r + 1; i <= R; i += i & -i) {
            for (int j = c + 1; j <= C; j += j & -j) {
                tree[static_cast<int64_t>(i) * (C + 1) + j] += x;
            }
        }
    }
};

```

```

    }
}

void add_helper(int r, int c, const T &x) {
    add(t1, 0, 0, x);
    add(t1, 0, c, -x);
    add(t2, 0, c, x * c);
    add(t1, r, 0, -x);
    add(t3, r, 0, x * r);
    add(t1, r, c, x);
    add(t2, r, c, -x * c);
    add(t3, r, c, -x * r);
    add(t4, r, c, x * r * c); // Overflow warning.
}

public:
void add(int r1, int c1, int r2, int c2, const T &x) {
    assert(0 <= r1 && r1 <= r2 && r2 < R);
    assert(0 <= c1 && c1 <= c2 && c2 < C);
    add_helper(r2 + 1, c2 + 1, x);
    add_helper(r1, c2 + 1, -x);
    add_helper(r2 + 1, c1, -x);
    add_helper(r1, c1, x);
}

void add(int r, int c, const T &x) { add(r, c, r, c, x); }
void set(int r, int c, const T &x) { add(r, c, x - at(r, c)); }

T sum(int r, int c) const {
    assert(-1 <= r && r < R && -1 <= c && c < C);
    r++;
    c++;
    T s1 = 0, s2 = 0, s3 = 0, s4 = 0;
    for (int i = r; i > 0; i -= i & -i) {
        for (int j = c; j > 0; j -= j & -j) {
            s1 += get(t1, i, j);
            s2 += get(t2, i, j);
            s3 += get(t3, i, j);
            s4 += get(t4, i, j);
        }
    }
    return s1 * r * c + s2 * r + s3 * c + s4;
}

T sum(int r1, int c1, int r2, int c2) const {
    assert(0 <= r1 && r1 <= r2 && r2 < R);
    assert(0 <= c1 && c1 <= c2 && c2 < C);
    return sum(r2, c2) + sum(r1 - 1, c1 - 1) - sum(r1 - 1, c2) - sum(r2, c1 - 1);
}

T at(int r, int c) const { return sum(r, c, r, c); }
};

```

Example Usage

```

#include <iostream>
using namespace std;

int main() {
    SparseFenwick2D<long long> t;
    t.set(0, 0, 5);
    t.set(0, 1, 6);
    t.set(1, 0, 7);
    t.add(2, 2, 9);
    t.add(1, 0, -4);
    t.add(1, 1, 2, 2, 5);
}

```

```

cout << "Values:" << endl;
for (int i = 0; i < 3; i++) {
    for (int j = 0; j < 3; j++) {
        cout << t.at(i, j) << " ";
    }
    cout << endl;
}
assert(t.sum(0, 0, 0, 1) == 11);
assert(t.sum(0, 0, 1, 0) == 8);
assert(t.sum(1, 1, 2, 2) == 29);
t.set(500000000, 500000000, 1);
assert(t.sum(0, 0, 999999999, 999999999) == 44);
return 0;
}

```

Example Output

```

Values:
5 6 0
3 5 5
0 5 14

```

2.6.7 Offline 2D Fenwick Tree

Maintain a 2D numerical array over an arbitrary, possibly huge index range while supporting point increments and rectangle-sum queries, when every coordinate that will ever be updated is known in advance. This offline restriction is what makes the structure cheap: both coordinates are coordinate-compressed against the update points, so no hashing is required and the storage is exactly the set of cells each Fenwick node can ever touch.

The value type `T` must represent 0 and support addition and subtraction.

The outer dimension is an ordinary Fenwick tree over the sorted distinct update rows. Each outer node owns an inner Fenwick tree over only those columns that reach it along an update path, again stored in a sorted vector. A prefix query maps its row argument to the number of update rows at or below it, walks the outer tree, and within each visited node maps its column argument to a rank by binary search. Compared with the online sparse 2D Fenwick tree backed by hash tables, this is deterministic and has markedly better constants, at the cost of requiring all update locations up front. Prefer the dense 2D Fenwick tree when the coordinates are small enough to allocate directly, and the sparse 2D Fenwick tree when update locations are not known ahead of time or rectangle updates are needed.

Usage is two-phase. First declare every cell that will be updated with `reserve()`; then call `build()` once; then issue `add()` updates and `sum()` / `at()` queries freely. Updates may only target reserved cells, but queries may use any coordinates.

- `OfflineFenwick2D<T>()` constructs an empty structure in the declaration phase.
- `reserve(r, c)` declares that the cell (r, c) may later be updated. Call before `build()`.
- `build()` finalizes the coordinate compression. Call exactly once, after all `reserve()` calls.
- `add(r, c, x)` adds `x` to the value at index (r, c) , which must have been reserved.
- `set(r, c, x)` assigns `x` to the value at index (r, c) , which must have been reserved.
- `sum(r, c)` returns the sum of the rectangle with rows $[0, r]$ and columns $[0, c]$.
- `sum(r1, c1, r2, c2)` returns the sum of the rectangle with rows $[r1, r2]$ and columns $[c1, c2]$.
- `at(r, c)` returns the value at index (r, c) .

Time Complexity:

- $O(n \log^2 n)$ per call to `build()`, where n is the number of reserved cells.
- $O(\log^2 n)$ per call to `add()`, `set()`, `sum()`, and `at()`.

Space Complexity:

- $O(n \log n)$ for storage, where n is the number of reserved cells.
- $O(1)$ auxiliary for all operations after `build()`.

```

#include <algorithm>
#include <cassert>
#include <utility>
#include <vector>

template<typename T>
class OfflineFenwick2D {
    std::vector<int> xs;           // Sorted distinct update rows.
    std::vector<std::vector<int>> ys; // ys[p]: sorted distinct columns at outer node p.
    std::vector<std::vector<T>> tree; // tree[p]: 1-indexed inner Fenwick over ys[p].
    std::vector<std::pair<int, int>> pending; // Reserved (row, column) cells.
    bool built = false;

    // Number of update rows at most r; the outer Fenwick prefix bound for a query at row r.
    int row_rank(int r) const {
        return static_cast<int>(std::upper_bound(xs.begin(), xs.end(), r) - xs.begin());
    }

    // 1-indexed outer position of an update row r, which is guaranteed to be present in xs.
    int row_pos(int r) const {
        return static_cast<int>(std::lower_bound(xs.begin(), xs.end(), r) - xs.begin()) + 1;
    }

    // Number of columns at most c registered at outer node p; the inner prefix bound.
    int col_rank(int p, int c) const {
        return static_cast<int>(std::upper_bound(ys[p].begin(), ys[p].end(), c) - ys[p].begin());
    }

    bool is_reserved(int r, int c) const {
        int p = row_pos(r);
        return p < static_cast<int>(tree.size()) && p > 0 && xs[p - 1] == r &&
            std::binary_search(ys[p].begin(), ys[p].end(), c);
    }

public:
    void reserve(int r, int c) {
        assert(!built);
        pending.emplace_back(r, c);
    }

    void build() {
        assert(!built);
        built = true;
        for (const auto &u : pending) {
            xs.push_back(u.first);
        }
        std::sort(xs.begin(), xs.end());
        xs.erase(std::unique(xs.begin(), xs.end()), xs.end());
        int m = static_cast<int>(xs.size());
        ys.assign(m + 1, {});
        // Register each update's column at every outer node on its update path.
        for (const auto &u : pending) {
            for (int x = row_pos(u.first); x <= m; x += x & -x) {
                ys[x].push_back(u.second);
            }
        }
        tree.assign(m + 1, {});
        for (int x = 1; x <= m; x++) {
            std::sort(ys[x].begin(), ys[x].end());
            ys[x].erase(std::unique(ys[x].begin(), ys[x].end()), ys[x].end());
            tree[x].assign(ys[x].size() + 1, T());
        }
        pending.clear();
    }

    void add(int r, int c, const T &x) {
        assert(built);
        assert(is_reserved(r, c));
    }

```

```

    for (int i = row_pos(r); i < static_cast<int>(tree.size()); i += i & -i) {
        for (int j = col_rank(i, c); j <= static_cast<int>(ys[i].size()); j += j & -j) {
            tree[i][j] += x;
        }
    }
}

T sum(int r, int c) const {
    assert(built);
    T res = T();
    for (int i = row_rank(r); i > 0; i -= i & -i) {
        for (int j = col_rank(i, c); j > 0; j -= j & -j) {
            res += tree[i][j];
        }
    }
    return res;
}

T sum(int r1, int c1, int r2, int c2) const {
    assert(r1 <= r2 && c1 <= c2);
    return sum(r2, c2) - sum(r1 - 1, c2) - sum(r2, c1 - 1) + sum(r1 - 1, c1 - 1);
}

T at(int r, int c) const { return sum(r, c, r, c); }
void set(int r, int c, const T &x) { add(r, c, x - at(r, c)); }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    OfflineFenwick2D<int> t;
    // Declare every cell that will be updated, then finalize.
    t.reserve(0, 0);
    t.reserve(0, 1000000);
    t.reserve(1000000, 0);
    t.reserve(1000000, 1000000);
    t.reserve(500000, 500000);
    t.build();

    t.set(0, 0, 5);
    t.set(0, 1000000, 6);
    t.set(1000000, 0, 7);
    t.add(1000000, 1000000, 9);
    t.add(500000, 500000, 4);

    assert(t.at(1000000, 1000000) == 9);
    assert(t.sum(0, 0, 0, 1000000) == 11);           // 5 + 6.
    assert(t.sum(0, 0, 1000000, 0) == 12);           // 5 + 7.
    assert(t.sum(0, 0, 1000000, 1000000) == 31);      // 5 + 6 + 7 + 9 + 4.
    assert(t.sum(400000, 400000, 600000, 600000) == 4); // Only the center cell.
    return 0;
}

```

2.7 Disjoint Sets and Tree Structures

2.7.1 Disjoint Set Union

Maintain a set of elements partitioned into non-overlapping subsets. Each partition is identified by a unique representative called its root. Each subset is stored as a tree whose nodes point toward its root: finding a representative follows parent pointers, redirecting visited nodes straight to the root along the way (path compression), while unions attach the smaller tree beneath the root of the larger one (union-by-size). Together these keep the trees nearly flat. Union-by-size is interchangeable with the union-by-rank strategy used by Sparse Disjoint Set Union and yields the same complexity bounds, but union-by-size additionally supports `set_size()` queries.

This version is simplified to work only on the fixed set of integer elements $[0, n)$, chosen at construction. For arbitrary element types, see Sparse Disjoint Set Union.

- `DSU(n)` constructs `n` singleton partitions over elements $[0, n)$.
- `find_root(u)` returns the unique representative of the partition containing `u`.
- `sets()` returns the current number of partitions.
- `set_size(u)` returns the number of elements in the partition containing `u`.
- `is_united(u, v)` returns whether elements `u` and `v` belong to the same partition.
- `unite(u, v)` replaces the partitions containing `u` and `v` with a single new partition consisting of the union of elements in the original partitions, returning `true` if a merge occurred or `false` if `u` and `v` already belonged to the same partition.

Time Complexity:

- $O(n)$ per call to the constructor, and $O(1)$ per call to `sets()`.
- $O(\alpha(n))$ per call to `find_root()`, `set_size()`, `is_united()`, and `unite()`, where n is the number of elements, and $\alpha(n)$ is the extremely slow growing inverse of the Ackermann function (effectively a very small constant for all practical values of n).

Space Complexity:

- $O(n)$ for storage of the disjoint set forest elements.
- $O(\log n)$ auxiliary stack space for operations that call `find_root()`.
- $O(1)$ auxiliary for `sets()`.

```
#include <numeric>
#include <utility>
#include <vector>

class DSU {
    std::vector<int> root, size;
    int num_sets;

public:
    explicit DSU(int n) : root(n), size(n, 1), num_sets(n) { std::iota(root.begin(), root.end(), 0); }

    int find_root(int u) {
        if (root[u] != u) {
            root[u] = find_root(root[u]);
        }
        return root[u];
    }

    int sets() const { return num_sets; }
    int set_size(int u) { return size[find_root(u)]; }
    bool is_united(int u, int v) { return find_root(u) == find_root(v); }

    bool unite(int u, int v) {
        int ru = find_root(u), rv = find_root(v);
        if (ru == rv) {
```

```

        return false;
    }
    num_sets--;
    if (size[ru] < size[rv]) {
        std::swap(ru, rv);
    }
    root[rv] = ru;
    size[ru] += size[rv];
    return true;
}
};

```

Example Usage

```

#include <cassert>

int main() {
    DSU dsu(7);
    assert(dsu.unite(0, 1));
    assert(dsu.unite(1, 5));
    assert(dsu.unite(3, 4));
    assert(dsu.unite(3, 6));
    assert(!dsu.unite(0, 5)); // Already united.
    assert(dsu.sets() == 3);
    assert(dsu.set_size(1) == 3);
    assert(dsu.set_size(2) == 1);
    assert(dsu.is_united(0, 1));
    assert(!dsu.is_united(0, 2));
    assert(!dsu.is_united(1, 6));
    assert(dsu.is_united(4, 6));
    return 0;
}

```

2.7.2 Sparse Disjoint Set Union

Maintain a set of elements partitioned into non-overlapping subsets using a collection of trees. Each partition is identified by a unique representative called its root. Each subset is stored as a tree whose nodes point toward its root: finding a representative follows parent pointers, redirecting visited nodes straight to the root along the way (path compression), while unions attach the shallower tree beneath the root of the deeper one (union-by-rank). Union-by-rank is interchangeable with union-by-size as used by the plain Disjoint Set Union, yielding the same complexity bounds; rank stores only the tree-depth estimate, so it cannot answer set-size queries.

This version uses an `std::unordered_map` for storage and coordinate compression (thus, element types must meet the requirements of key types for `std::unordered_map`). The order of sets returned by `get_all_sets()` is unspecified.

- `SparseDSU<T>()` constructs an empty set.
- `make_set(u)` adds `u` as a singleton partition, or does nothing if `u` is already present.
- `size()` returns the number of elements that have been added.
- `sets()` returns the current number of disjoint sets.
- `is_united(u, v)` returns whether elements `u` and `v` belong to the same partition.
- `unite(u, v)` replaces the partitions containing `u` and `v` with a single new partition consisting of the union of elements in the original partitions, returning `true` if a merge occurred or `false` if `u` and `v` already belonged to the same partition.
- `get_all_sets()` returns all current partitions as a vector of vectors.

A precondition to the last three operations is that `make_set()` must have been previously called on their arguments.

Time Complexity:

- $O(1)$ per call to the constructor.

- $O(1)$ per call to `size()` and `sets()`.
- $O(1)$ on average per call to `make_set()`.
- $O(\alpha(n))$ on average per call to `is_united()` and `unite()`, where n is the number of elements that have been added via `make_set()` so far, and $\alpha(n)$ is the extremely slow growing inverse of the Ackermann function (effectively a very small constant for all practical values of n).
- $O(n)$ per call to `get_all_sets()`.

Space Complexity:

- $O(n)$ for storage of the disjoint set forest elements.
- $O(n)$ auxiliary for `get_all_sets()`.
- $O(\log n)$ auxiliary stack space for `is_united()` and `unite()`.
- $O(1)$ auxiliary for all other operations.

```
#include <cassert>
#include <unordered_map>
#include <vector>

template<typename T>
class SparseDSU {
    int num_elements, num_sets;
    std::unordered_map<T, int> id;
    std::vector<int> root, rank;

    int find_root(int u) {
        if (root[u] != u) {
            root[u] = find_root(root[u]);
        }
        return root[u];
    }

    int get_id(const T &u) const {
        auto it = id.find(u);
        assert(it != id.end());
        return it->second;
    }

public:
    SparseDSU() : num_elements(0), num_sets(0) {}

    void make_set(const T &u) {
        if (id.find(u) != id.end()) {
            return;
        }
        id[u] = num_elements;
        root.push_back(num_elements++);
        rank.push_back(0);
        num_sets++;
    }

    int size() const { return num_elements; }
    int sets() const { return num_sets; }
    bool is_united(const T &u, const T &v) { return find_root(get_id(u)) == find_root(get_id(v)); }

    bool unite(const T &u, const T &v) {
        int ru = find_root(get_id(u)), rv = find_root(get_id(v));
        if (ru == rv) {
            return false;
        }
        num_sets--;
        if (rank[ru] < rank[rv]) {
            root[ru] = rv;
        } else {
            root[rv] = ru;
            if (rank[ru] == rank[rv]) {
                rank[ru]++;
            }
        }
    }
};
```

```

    }
}
return true;
}

std::vector<std::vector<T>> get_all_sets() {
    std::unordered_map<int, std::vector<T>> tmp;
    for (auto &[key, val] : id) {
        tmp[find_root(val)].push_back(key);
    }
    std::vector<std::vector<T>> res;
    for (auto &[key, val] : tmp) {
        res.push_back(val);
    }
    return res;
}
};

```

Example Usage

```

#include <algorithm>
#include <cassert>
using namespace std;

int main() {
    SparseDSU<char> dsu;
    for (char c = 'a'; c <= 'g'; c++) {
        dsu.make_set(c);
    }
    assert(dsu.unite('a', 'b'));
    assert(dsu.unite('b', 'f'));
    assert(dsu.unite('d', 'e'));
    assert(dsu.unite('d', 'g'));
    assert(!dsu.unite('a', 'f')); // Already united.
    assert(dsu.size() == 7);
    assert(dsu.sets() == 3);
    auto s = dsu.get_all_sets();
    for (auto &group : s) {
        sort(group.begin(), group.end());
    }
    sort(s.begin(), s.end());
    assert((s == vector<vector<char>>{{'a', 'b', 'f'}, {'c'}, {'d', 'e', 'g'}}));
    return 0;
}

```

2.7.3 Rollback Disjoint Set Union

Maintains disjoint sets while supporting rollback to previous snapshots. Rollback DSU is useful for offline dynamic connectivity, divide-and-conquer over time, and backtracking search where unions must be undone.

Path compression is intentionally not used, because it is hard to undo. Union by size keeps tree height logarithmic, and every successful union records enough information to restore the previous state. Specifically, a union attaches the smaller root below the larger one and records the attached root, its new parent, and that parent's old size. A snapshot is just the current number of history records. Rolling back pops records until reaching that position, detaching each child, restoring the parent size, and increasing the set count. Unions of elements already in the same set change nothing and therefore add no history record.

- `RollbackDSU(n)` constructs `n` singleton sets over elements `[0, n)`.
- `sets()` returns the current number of disjoint sets.
- `find_root(u)` returns the representative of the set containing `u`.
- `is_united(u, v)` returns whether `u` and `v` are in the same set.

- `unite(u, v)` merges two sets and returns whether a merge occurred.
- `snapshot()` returns a token representing the current history size.
- `rollback(snapshot)` undoes all changes made after `snapshot`.

Time Complexity:

- $O(n)$ per call to the constructor.
- $O(1)$ per call to `sets()`.
- $O(\log n)$ worst-case per call to `find_root()`, `is_united()`, and `unite()`.
- $O(1)$ per undone union during `rollback()`.

Space Complexity: $O(n + q)$ for q successful unions stored in the rollback history.

```
#include <cassert>
#include <numeric>
#include <utility>
#include <vector>

class RollbackDSU {
    struct Change {
        int child, parent, parent_size;

        Change(int child = -1, int parent = -1, int parent_size = 0)
            : child(child), parent(parent), parent_size(parent_size) {}
    };

    std::vector<int> root, size;
    std::vector<Change> history;
    int num_sets;

public:
    explicit RollbackDSU(int n) : root(n), size(n, 1), num_sets(n) {
        std::iota(root.begin(), root.end(), 0);
    }

    int find_root(int u) const {
        assert(0 <= u && u < static_cast<int>(root.size()));
        while (root[u] != u) {
            u = root[u];
        }
        return u;
    }

    int sets() const { return num_sets; }
    bool is_united(int u, int v) const { return find_root(u) == find_root(v); }

    bool unite(int u, int v) {
        u = find_root(u);
        v = find_root(v);
        if (u == v) {
            return false;
        }
        if (size[u] > size[v]) {
            std::swap(u, v);
        }
        history.emplace_back(u, v, size[v]);
        root[u] = v;
        size[v] += size[u];
        num_sets--;
        return true;
    }

    int snapshot() const { return static_cast<int>(history.size()); }

    void rollback(int snapshot) {
        assert(0 <= snapshot && snapshot <= static_cast<int>(history.size()));
        while (static_cast<int>(history.size()) > snapshot) {
```

```

        Change c = history.back();
        history.pop_back();
        root[c.child] = c.child;
        size[c.parent] = c.parent_size;
        num_sets++;
    }
}
};

```

Example Usage

```

int main() {
    RollbackDSU dsu(5);
    dsu.unite(0, 1);
    int s = dsu.snapshot();
    dsu.unite(1, 2);
    dsu.unite(3, 4);
    assert(dsu.is_united(0, 2));
    assert(dsu.sets() == 2);
    dsu.rollback(s);
    assert(dsu.is_united(0, 1));
    assert(!dsu.is_united(0, 2));
    assert(!dsu.is_united(3, 4));
    assert(dsu.sets() == 4);
    return 0;
}

```

2.7.4 Lowest Common Ancestor (Sparse Table)

Given a rooted tree or forest, determine the lowest common ancestor of any two nodes in the same tree. The lowest common ancestor of two nodes u and v is the node that has the greatest depth while having both u and v as descendants. A node is considered to be a descendant of itself.

This implementation preprocesses binary ancestor jumps. During depth-first search, it records entry and exit times so ancestry can be tested in $O(1)$, records each node's depth and root, and stores `up[u][i]`, the 2^i -th ancestor of each node u . To answer `lca(u, v)`, it first handles the case where one node is already an ancestor of the other, then jumps u upward by decreasing powers of two until its parent is the lowest common ancestor.

- `SparseTableLCA(adj)` builds the structure over a forest represented by the bidirectional adjacency list `adj`, whose indices represent the nodes.
- `go_up(u, k)` returns the k -th ancestor of node u , stopping at that tree's root if k is larger than u 's depth.
- `lca(u, v)` returns the lowest common ancestor of nodes u and v , or -1 if they are in different trees.
- `is_ancestor(parent, child)` returns whether `parent` is an ancestor of `child`.
- `dist(u, v)` returns the number of edges on the path between nodes u and v , or -1 if they are in different trees.

Time Complexity:

- $O(n \log n)$ for construction, where n is the number of nodes.
- $O(\log n)$ per call to `go_up()`, `lca()`, and `dist()`.
- $O(1)$ per call to `is_ancestor()`.

Space Complexity:

- $O(n \log n)$ to store the binary ancestor table.
- $O(n)$ auxiliary stack space for construction.
- $O(1)$ auxiliary for `go_up()`, `lca()`, `is_ancestor()`, and `dist()`.

```

#include <algorithm>
#include <cassert>

```

```

#include <vector>

class SparseTableLCA {
    std::vector<std::vector<int>>> up;
    std::vector<int> tin, tout, depth, root;
    int len, timer;

    void dfs(const std::vector<std::vector<int>>> &adj, int u, int p, int r, int d) {
        tin[u] = timer++;
        depth[u] = d;
        root[u] = r;
        up[u][0] = p;
        for (int i = 1; i < len; i++) {
            up[u][i] = up[up[u][i - 1]][i - 1];
        }
        for (int v : adj[u]) {
            if (v != p) {
                dfs(adj, v, u, r, d + 1);
            }
        }
        tout[u] = timer++;
    }

public:
    explicit SparseTableLCA(const std::vector<std::vector<int>>> &adj) : timer(0) {
        int n = static_cast<int>(adj.size());
        len = 1;
        while ((1 << len) <= std::max(1, n)) {
            len++;
        }
        up.assign(n, std::vector<int>(len));
        tin.assign(n, 0);
        tout.assign(n, 0);
        depth.assign(n, 0);
        root.assign(n, -1);
        for (int u = 0; u < n; u++) {
            if (root[u] == -1) {
                dfs(adj, u, u, u, 0);
            }
        }
    }

    bool is_ancestor(int parent, int child) const {
        assert(0 <= parent && parent < static_cast<int>(root.size()));
        assert(0 <= child && child < static_cast<int>(root.size()));
        return root[parent] == root[child] && tin[parent] <= tin[child] && tout[child] <= tout[parent];
    }

    int go_up(int u, int k) const {
        assert(0 <= u && u < static_cast<int>(root.size()) && k >= 0);
        k = std::min(k, depth[u]);
        for (int i = 0; i < len; i++) {
            if ((k & (1 << i)) != 0) {
                u = up[u][i];
            }
        }
        return u;
    }

    int lca(int u, int v) const {
        assert(0 <= u && u < static_cast<int>(root.size()));
        assert(0 <= v && v < static_cast<int>(root.size()));
        if (root[u] != root[v]) {
            return -1;
        }
        if (is_ancestor(u, v)) {
            return u;
        }
    }
}

```

```

    if (is_ancestor(v, u)) {
        return v;
    }
    for (int i = len - 1; i >= 0; i--) {
        if (!is_ancestor(up[u][i], v)) {
            u = up[u][i];
        }
    }
    return up[u][0];
}

int dist(int u, int v) const {
    int l = lca(u, v);
    return l == -1 ? -1 : depth[u] + depth[v] - 2 * depth[l];
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

void add_edge(vector<vector<int>> &adj, int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
}

int main() {
    // 0---1---2    5---6
    // |   |
    // 4   3
    vector<vector<int>> adj(7);
    add_edge(adj, 0, 1);
    add_edge(adj, 0, 4);
    add_edge(adj, 1, 2);
    add_edge(adj, 1, 3);
    add_edge(adj, 5, 6);

    SparseTableLCA tree(adj);
    assert(tree.go_up(3, 1) == 1);
    assert(tree.go_up(3, 2) == 0);
    assert(tree.go_up(3, 10) == 0);
    assert(tree.lca(3, 2) == 1);
    assert(tree.lca(2, 4) == 0);
    assert(tree.lca(2, 6) == -1);
    assert(tree.lca(5, 6) == 5);
    assert(tree.dist(3, 2) == 2);    // 3-1-2.
    assert(tree.dist(2, 4) == 3);    // 2-1-0-4.
    assert(tree.dist(5, 6) == 1);    // 5-6.
    assert(tree.dist(2, 6) == -1);   // Different trees.
    return 0;
}

```

2.7.5 Lowest Common Ancestor (Segment Tree)

Given a rooted tree or forest, determine the lowest common ancestor of any two nodes in the same tree. The lowest common ancestor of two nodes u and v is the node that has the greatest depth while having both u and v as descendants. A node is considered to be a descendant of itself.

This reduces LCA to a range-minimum query. An Euler tour of the forest records, at each visit to a node, that node's depth; the lowest common ancestor of u and v in the same tree is the shallowest node visited between their

first occurrences in the tour, found by a range-minimum query over the depth sequence. This version answers those queries with a segment tree.

- `SegTreeLCA(adj)` builds the structure over a forest represented by the bidirectional adjacency list `adj`, whose indices represent the nodes.
- `lca(u, v)` returns the lowest common ancestor of nodes `u` and `v`, or `-1` if they are in different trees.
- `dist(u, v)` returns the number of edges on the path between nodes `u` and `v`, or `-1` if they are in different trees.

Time Complexity:

- $O(n)$ for construction, where n is the number of nodes.
- $O(\log n)$ per call to `lca()` and `dist()`.

Space Complexity:

- $O(n)$ for storage of the segment tree.
- $O(n)$ auxiliary stack space for construction.
- $O(\log n)$ auxiliary stack space for `lca()`.

```
#include <algorithm>
#include <cassert>
#include <vector>

class SegTreeLCA {
    std::vector<int> depth, order, first, root, minpos;
    int len;

    void dfs(const std::vector<std::vector<int>> &adj, int u, int r, int d) {
        root[u] = r;
        depth[u] = d;
        first[u] = static_cast<int>(order.size());
        order.push_back(u);
        for (int v : adj[u]) {
            if (depth[v] == -1) {
                dfs(adj, v, r, d + 1);
                order.push_back(u);
            }
        }
    }

    int better(int a, int b) const { return depth[a] < depth[b] ? a : b; }

    void build(int node, int lo, int hi) {
        if (lo == hi) {
            minpos[node] = order[lo];
            return;
        }
        int mid = lo + (hi - lo) / 2;
        build(2 * node + 1, lo, mid);
        build(2 * node + 2, mid + 1, hi);
        minpos[node] = better(minpos[2 * node + 1], minpos[2 * node + 2]);
    }

    int get_minpos(int a, int b, int node, int lo, int hi) const {
        if (a == lo && b == hi) {
            return minpos[node];
        }
        int mid = lo + (hi - lo) / 2;
        if (b <= mid) {
            return get_minpos(a, b, 2 * node + 1, lo, mid);
        }
        if (mid < a) {
            return get_minpos(a, b, 2 * node + 2, mid + 1, hi);
        }
        return better(
            get_minpos(a, mid, 2 * node + 1, lo, mid), get_minpos(mid + 1, b, 2 * node + 2, mid + 1, hi)
        );
    }
};
```

```

    );
}

public:
explicit SegTreeLCA(const std::vector<std::vector<int>> &adj) {
    int n = static_cast<int>(adj.size());
    depth.assign(n, -1);
    first.assign(n, -1);
    root.assign(n, -1);
    for (int u = 0; u < n; u++) {
        if (depth[u] == -1) {
            dfs(adj, u, u, 0);
        }
    }
    len = static_cast<int>(order.size());
    minpos.assign(std::max(1, 4 * len), 0);
    if (len > 0) {
        build(0, 0, len - 1);
    }
}

int lca(int u, int v) const {
    assert(0 <= u && u < static_cast<int>(root.size()));
    assert(0 <= v && v < static_cast<int>(root.size()));
    if (root[u] != root[v]) {
        return -1;
    }
    return get_minpos(std::min(first[u], first[v]), std::max(first[u], first[v]), 0, 0, len - 1);
}

int dist(int u, int v) const {
    int l = lca(u, v);
    return l == -1 ? -1 : depth[u] + depth[v] - 2 * depth[l];
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

void add_edge(vector<vector<int>> &adj, int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
}

int main() {
    // 0---1---2    5---6
    // |    |
    // 4    3
    vector<vector<int>> adj(7);
    add_edge(adj, 0, 1);
    add_edge(adj, 0, 4);
    add_edge(adj, 1, 2);
    add_edge(adj, 1, 3);
    add_edge(adj, 5, 6);

    SegTreeLCA tree(adj);
    assert(tree.lca(3, 2) == 1);
    assert(tree.lca(2, 4) == 0);
    assert(tree.lca(2, 6) == -1);
    assert(tree.lca(5, 6) == 5);
    assert(tree.dist(3, 2) == 2);    // 3-1-2.
    assert(tree.dist(2, 4) == 3);    // 2-1-0-4.
    assert(tree.dist(5, 6) == 1);    // 5-6.
    assert(tree.dist(2, 6) == -1);   // Different trees.
}

```



```

    return 0;
}

```

2.7.6 Virtual Tree

Given a rooted tree and a subset of its nodes, build the virtual tree (also called the auxiliary or compressed tree): the smallest tree that contains every node of the subset together with the lowest common ancestor of each pair, with chains of irrelevant degree-two ancestors compressed into single edges. Many problems ask repeated queries each restricted to a handful of “important” nodes; the virtual tree has $O(k)$ nodes for a subset of size k , so an algorithm can run over it instead of the whole tree, turning a per-query cost proportional to n into one proportional to the subset size.

The construction sorts the subset by depth-first entry time, which lists the nodes in the order a walk of the tree would meet them. The lowest common ancestor of any contiguous span of the subset is the minimum-depth node among the pairwise LCAs of adjacent entries, so it suffices to add the LCA of each adjacent pair, re-sort, and deduplicate. A single stack pass over the resulting nodes then links each one to its parent in the virtual tree: ancestors that are no longer on the path to the current node are popped, and the node on top becomes the parent. This implementation carries its own binary lifting so it is self-contained; given an existing lowest common ancestor structure, the same sort-and-stack logic applies directly.

- `VirtualTree(adj, root)` builds the structure over a tree rooted at `root`, represented by the bidirectional adjacency list `adj`, whose indices represent the nodes.
- `lca(u, v)` returns the lowest common ancestor of nodes `u` and `v`.
- `depth_of(u)` returns the depth of `u` in the original tree (the root has depth 0). The original path length of a virtual-tree edge `p-c` is `depth_of(c) - depth_of(p)`.
- `build(nodes)` returns the virtual tree induced by `nodes`. The result has `root` set to the topmost node (the common ancestor of `nodes`), `nodes` listing every original node ID in the virtual tree sorted by entry time, and `edges` listing each `(p, c)` pair in original IDs. Both the node set and the edges use the node numbering of the original tree. An empty input yields an empty result with `root` = -1.

Time Complexity:

- $O(n \log n)$ for construction, where n is the number of nodes.
- $O(\log n)$ per call to `lca()`.
- $O(k \log n)$ per call to `build()`, where k is the size of the input subset.

Space Complexity:

- $O(n \log n)$ to store the binary ancestor table.
- $O(k)$ auxiliary plus result size for `build()`.

```

#include <algorithm>
#include <cassert>
#include <utility>
#include <vector>

class VirtualTree {
    std::vector<std::vector<int>>> up;
    std::vector<int> tin, tout, depth;
    int len, timer;

    void dfs(const std::vector<std::vector<int>>> &adj, int u, int p, int d) {
        tin[u] = timer++;
        depth[u] = d;
        up[u][0] = p;
        for (int i = 1; i < len; i++) {
            up[u][i] = up[up[u][i - 1]][i - 1];
        }
        for (int v : adj[u]) {
            if (v != p) {

```

```

        dfs(adj, v, u, d + 1);
    }
}
tout[u] = timer++;
}

public:
struct Tree {
    int root = -1;
    std::vector<int> nodes;
    std::vector<std::pair<int, int>> edges; // (parent, child) in original node ids.
};

VirtualTree(const std::vector<std::vector<int>> &adj, int root) {
    int n = static_cast<int>(adj.size());
    assert(0 <= root && root < n);
    len = 1;
    while ((1 << len) <= std::max(1, n)) {
        len++;
    }
    up.assign(n, std::vector<int>(len));
    tin.assign(n, 0);
    tout.assign(n, 0);
    depth.assign(n, 0);
    timer = 0;
    dfs(adj, root, root, 0);
}

bool is_ancestor(int parent, int child) const {
    assert(0 <= parent && parent < static_cast<int>(tin.size()));
    assert(0 <= child && child < static_cast<int>(tin.size()));
    return tin[parent] <= tin[child] && tout[child] <= tout[parent];
}

// Depth of u in the original tree; depth(child) - depth(parent) is the length of the original
// path that a compressed virtual-tree edge stands for.
int depth_of(int u) const {
    assert(0 <= u && u < static_cast<int>(depth.size()));
    return depth[u];
}

int lca(int u, int v) const {
    assert(0 <= u && u < static_cast<int>(tin.size()));
    assert(0 <= v && v < static_cast<int>(tin.size()));
    if (is_ancestor(u, v)) {
        return u;
    }
    if (is_ancestor(v, u)) {
        return v;
    }
    for (int i = len - 1; i >= 0; i--) {
        if (!is_ancestor(up[u][i], v)) {
            u = up[u][i];
        }
    }
    return up[u][0];
}

Tree build(std::vector<int> nodes) const {
    for (int u : nodes) {
        assert(0 <= u && u < static_cast<int>(tin.size()));
    }
    Tree res;
    if (nodes.empty()) {
        return res;
    }
    auto by_tin = [&](int a, int b) { return tin[a] < tin[b]; };
    std::sort(nodes.begin(), nodes.end(), by_tin);

```

```

nodes.erase(std::unique(nodes.begin(), nodes.end()), nodes.end());
int k = static_cast<int>(nodes.size());
for (int i = 0; i + 1 < k; i++) {
    nodes.push_back(lca(nodes[i], nodes[i + 1]));
}
std::sort(nodes.begin(), nodes.end(), by_tin);
nodes.erase(std::unique(nodes.begin(), nodes.end()), nodes.end());
std::vector<int> ancestors;
for (int v : nodes) {
    while (!ancestors.empty() && !is_ancestor(ancestors.back(), v)) {
        ancestors.pop_back();
    }
    if (!ancestors.empty()) {
        res.edges.emplace_back(ancestors.back(), v);
    }
    ancestors.push_back(v);
}
res.root = nodes.front();
res.nodes = std::move(nodes);
return res;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

void add_edge(vector<vector<int>> &adj, int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
}

int main() {
    //      0
    //     / |
    //    1 2
    //   /| |
    //  6--3 4 5--7
    vector<vector<int>> adj(8);
    add_edge(adj, 0, 1);
    add_edge(adj, 0, 2);
    add_edge(adj, 1, 3);
    add_edge(adj, 1, 4);
    add_edge(adj, 2, 5);
    add_edge(adj, 3, 6);
    add_edge(adj, 5, 7);

    VirtualTree vt(adj, 0);
    assert(vt.lca(6, 4) == 1);
    assert(vt.lca(6, 7) == 0);
    assert(vt.depth_of(0) == 0 && vt.depth_of(6) == 3);
    // The compressed edge 0 -> 7 stands for an original path of length 3 (0-2-5-7).
    assert(vt.depth_of(7) - vt.depth_of(0) == 3);

    // Virtual tree of {6, 7}: their LCA is 0, so the tree is 0 -> {6, 7} with edges compressed.
    //
    //      0
    //     +--+--+
    //    6      7
    VirtualTree::Tree t = vt.build({6, 7});
    assert(t.root == 0);
    assert((t.nodes == vector<int>{0, 6, 7}));
    assert(t.edges.size() == 2);
    for (auto &e : t.edges) {
        assert(e.first == 0 && (e.second == 6 || e.second == 7));
    }
}

```

```

}

// Virtual tree of {6, 4, 7}: node 1 becomes a branch point (LCA of 6 and 4).
//
//      0
//  +--+--+
//  1      7
// +--+
// 6   4
VirtualTree::Tree u = vt.build({6, 4, 7});
assert(u.root == 0);
assert((u.nodes == vector<int>{0, 1, 6, 4, 7}));
// Expected edges: 0->1, 1->6, 1->4, 0->7.
vector<pair<int, int>> want{{0, 1}, {1, 6}, {1, 4}, {0, 7}};
auto got = u.edges;
sort(got.begin(), got.end());
sort(want.begin(), want.end());
assert(got == want);

// A single node is its own virtual tree with no edges.
VirtualTree::Tree s = vt.build({4});
assert(s.root == 4 && s.edges.empty() && s.nodes == vector<int>{4});
return 0;
}

```

2.7.7 Weighted Tree Path Queries (Binary Lifting)

Given a tree or forest with weighted edges, aggregate the edge weights along the path between any two nodes in the same tree. The input is undirected; construction roots each component at the first unvisited node in index order. This builds on the same binary lifting table as LCA: alongside `up[u][i]`, the 2^i -th ancestor of `u`, it stores `agg[u][i]`, the combined weight of the 2^i edges traversed by that jump. A path query splits the path `u-v` at their lowest common ancestor and merges the two upward climbs.

The aggregate is a commutative monoid defined by `combine()` and its identity `identity()`. The two must be edited together; the defaults below compute the maximum edge weight on the path. For the minimum, use `std::min` with `std::numeric_limits<W>::max()`; for the sum (the path length), use `a + b` with `W()`. Commutativity is required because the two half-paths are merged without regard to orientation. For an invertible aggregate such as the sum, an alternative that avoids the `agg` table is to store each node's weighted depth `dw[u]` and return `dw[u] + dw[v] - 2*dw[lca]`.

- `WeightedTree<W>(adj)` builds the structure over a forest given by a weighted, bidirectional adjacency list `adj`, whose indices represent the nodes and where `adj[u]` holds pairs `(v, w)` for each edge `u-v` of weight `w`. Each connected component is rooted at its first node reached by the constructor's outer loop.
- `lca(u, v)` returns the lowest common ancestor of `u` and `v`, or `-1` if they lie in different trees.
- `kth_ancestor(u, k)` returns the k -th ancestor of `u`, stopping at that tree's root if `k` exceeds `u`'s depth.
- `is_ancestor(parent, child)` returns whether `parent` is an ancestor of `child`.
- `path_query(u, v)` returns the combined weight of the edges on the path from `u` to `v`, or `identity()` if they lie in different trees. If `u` equals `v` then the path has no edges, so `identity()` is returned.

Time Complexity:

- $O(n \log n)$ for construction, where n is the number of nodes.
- $O(\log n)$ per call to `lca()`, `kth_ancestor()`, and `path_query()`.
- $O(1)$ per call to `is_ancestor()`.

Space Complexity:

- $O(n \log n)$ to store the binary ancestor and aggregate tables.
- $O(n)$ auxiliary stack space for construction.
- $O(1)$ auxiliary for `lca()`, `kth_ancestor()`, `is_ancestor()`, and `path_query()`.

```

#include <algorithm>
#include <cassert>
#include <limits>
#include <utility>
#include <vector>

template<typename W>
class WeightedTree {
    static W combine(const W &a, const W &b) { return std::max(a, b); }
    static W identity() { return std::numeric_limits<W>::lowest(); }

    std::vector<std::vector<int>>> up;
    std::vector<std::vector<W>>> agg; // agg[u][i] = combined weight of the 2i edges above u.
    std::vector<int> tin, tout, depth, root;
    int len, timer;

    void dfs(int u, int p, int r, int d, const std::vector<std::vector<std::pair<int, W>>> &adj) {
        tin[u] = timer++;
        depth[u] = d;
        root[u] = r;
        up[u][0] = p;
        for (int i = 1; i < len; i++) {
            up[u][i] = up[up[u][i - 1]][i - 1];
            agg[u][i] = combine(agg[u][i - 1], agg[up[u][i - 1]][i - 1]);
        }
        for (const auto &[v, w] : adj[u]) {
            if (v != p) {
                agg[v][0] = w; // The edge into a child is stored on the child.
                dfs(v, u, r, d + 1, adj);
            }
        }
        tout[u] = timer++;
    }

    // Combined weight of the k edges on the upward path of k steps from u.
    W climb(int u, int k) const {
        W res = identity();
        for (int i = 0; i < len; i++) {
            if ((k & (1 << i)) != 0) {
                res = combine(res, agg[u][i]);
                u = up[u][i];
            }
        }
        return res;
    }
};

public:
explicit WeightedTree(const std::vector<std::vector<std::pair<int, W>>> &adj) : timer(0) {
    int n = static_cast<int>(adj.size());
    len = 1;
    while ((1 << len) <= std::max(1, n)) {
        len++;
    }
    up.assign(n, std::vector<int>(len));
    agg.assign(n, std::vector<W>(len, identity()));
    tin.assign(n, 0);
    tout.assign(n, 0);
    depth.assign(n, 0);
    root.assign(n, -1);
    for (int u = 0; u < n; u++) {
        if (root[u] == -1) {
            dfs(u, u, u, 0, adj);
        }
    }
}

bool is_ancestor(int parent, int child) const {
    assert(0 <= parent && parent < static_cast<int>(root.size()));

```

```

    assert(0 <= child && child < static_cast<int>(root.size()));
    return root[parent] == root[child] && tin[parent] <= tin[child] && tout[child] <= tout[parent];
}

int kth_ancestor(int u, int k) const {
    assert(0 <= u && u < static_cast<int>(root.size()) && k >= 0);
    k = std::min(k, depth[u]);
    for (int i = 0; i < len; i++) {
        if ((k & (1 << i)) != 0) {
            u = up[u][i];
        }
    }
    return u;
}

int lca(int u, int v) const {
    assert(0 <= u && u < static_cast<int>(root.size()));
    assert(0 <= v && v < static_cast<int>(root.size()));
    if (root[u] != root[v]) {
        return -1;
    }
    if (is_ancestor(u, v)) {
        return u;
    }
    if (is_ancestor(v, u)) {
        return v;
    }
    for (int i = len - 1; i >= 0; i--) {
        if (!is_ancestor(up[u][i], v)) {
            u = up[u][i];
        }
    }
    return up[u][0];
}

W path_query(int u, int v) const {
    int l = lca(u, v);
    if (l == -1) {
        return identity();
    }
    return combine(climb(u, depth[u] - depth[l]), climb(v, depth[v] - depth[l]));
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

void add_edge(vector<vector<pair<int, int>>> &adj, int u, int v, int w) {
    adj[u].push_back({v, w});
    adj[v].push_back({u, w});
}

int main() {
    //
    //      w=4 /      \ w=2
    //      1         2
    //  w=7 / \ w=1   w=5 / \ w=3
    //    3  4       5  6
    vector<vector<pair<int, int>>> adj(7);
    add_edge(adj, 0, 1, 4);
    add_edge(adj, 0, 2, 2);
    add_edge(adj, 1, 3, 7);
    add_edge(adj, 1, 4, 1);
    add_edge(adj, 2, 5, 5);
}

```

```

add_edge(adj, 2, 6, 3);
WeightedTree<int> tree(adj);

assert(tree.lca(3, 4) == 1);
assert(tree.lca(5, 6) == 2);
assert(tree.lca(3, 6) == 0);

assert(tree.kth_ancestor(3, 1) == 1);
assert(tree.kth_ancestor(3, 2) == 0);
assert(tree.kth_ancestor(3, 5) == 0); // Clamped to the root.

// Default combine() reports the maximum edge weight on the path.
assert(tree.path_query(3, 4) == 7); // 3-1 (7), 1-4 (1).
assert(tree.path_query(5, 6) == 5); // 5-2 (5), 2-6 (3).
assert(tree.path_query(4, 5) == 5); // 4-1 (1), 1-0 (4), 0-2 (2), 2-5 (5).
assert(tree.path_query(0, 3) == 7); // 0-1 (4), 1-3 (7).

vector<vector<pair<int, int>>> forest(2);
WeightedTree<int> disconnected(forest);
assert(disconnected.lca(0, 1) == -1);
assert(disconnected.path_query(0, 1) == numeric_limits<int>::lowest());
return 0;
}

```

2.7.8 Heavy Light Decomposition

Maintain a tree or forest with values associated with either edges or nodes, while supporting both dynamic queries and dynamic updates of all values on a given path between two nodes in the tree. Heavy-light decomposition chooses at most one heavy child for each node and partitions the tree into maximal paths of heavy edges. This implementation keeps a child on its parent’s path when its subtree contains at least half of its parent’s nodes. Every remaining edge is light and at least halves the remaining subtree size, so walking from any node to the root crosses $O(\log n)$ heavy paths. Each path stores its values in its own lazy segment tree, decomposing any path query or update into $O(\log n)$ contiguous range operations.

The query operation is defined by an associative `combine()` function. The default code below assumes a numerical tree type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`. For direction-independent path queries, `combine()` should also be commutative; otherwise, store enough information in each aggregate to combine paths in the required order.

The update operation is defined by `apply_delta()` and `compose_deltas()`. A delta must act on an aggregate summary of a path of length `len`: `apply_delta(v, d, len)` returns the aggregate after applying update `d` to every element represented by aggregate `v`. Pending deltas are combined in chronological order by `compose_deltas(old, d)`, meaning “apply `old`, then apply `d`”. These hooks do not support arbitrary query/update pairings; the delta operation must distribute over `combine()`, and composed deltas must have the same effect as applying their updates sequentially. The default code below defines updates that “set” a path’s edges or nodes to a new value. For range increment updates, `apply_delta(v, d, len)` would return `v + d` for min/max queries, or `v + d * len` for sum queries, and `compose_deltas(old, d)` would return `old + d`.

- `HeavyLight<T, VALUES_ON_EDGES = true>(adj, value)` constructs a new heavy light decomposition on a forest defined by the adjacency list `adj`, with all values initialized to `value`. Its entries must be valid indices into `adj`, and no duplicate edges should exist. Set `VALUES_ON_EDGES` to false to store values on nodes instead.
- `query(u, v)` returns the result of `combine()` applied to all values on the path from node `u` to node `v`. The nodes must be in the same tree and, when values are on edges, must be distinct.
- `update(u, v, d)` modifies all values on the path from node `u` to node `v` by applying the delta `d`. The nodes must be in the same tree.
- `for_each_path(u, v, include_lca, f)` decomposes the path from node `u` to node `v` into heavy path ranges and calls `f(path, lo, hi, up)` on each range, where `up` says the path segment is traversed upward from `u` toward the LCA. If `VALUES_ON_EDGES = true`, pass `include_lca = false` to skip the LCA’s node slot. Returns false if `u` and `v` are in different trees.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the number of nodes.
- $O(\log n)$ calls to `f()` per call to `for_each_path()`.
- $O(\log^2 n)$ per call to `query()` and `update()`.

Space Complexity:

- $O(n)$ for storage of the decomposition.
- $O(n)$ auxiliary stack space for the constructor.
- $O(\log n)$ auxiliary for `for_each_path()`, `query()`, and `update()`.

```
#include <algorithm>
#include <cassert>
#include <tuple>
#include <vector>

template<typename T, bool VALUES_ON_EDGES = true>
class HeavyLight {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d, int len) { return d; }
    static T compose_deltas(const T &old, const T &d) { return d; }

    std::vector<std::vector<int>>> adj;
    std::vector<std::vector<T>> value, delta;
    std::vector<std::vector<char>>> pending;
    std::vector<std::vector<int>>> len;
    std::vector<int> size, parent, root, tin, tout, path, pathlen, pathpos, pathroot;
    int counter, paths;

    void dfs(int u, int p, int r) {
        tin[u] = counter++;
        parent[u] = p;
        root[u] = r;
        size[u] = 1;
        for (int v : adj[u]) {
            if (v != p) {
                dfs(v, u, r);
                size[u] += size[v];
            }
        }
        tout[u] = counter++;
    }

    int new_path(int u) {
        pathroot[paths] = u;
        return paths++;
    }

    void build_paths(int u, int path) {
        this->path[u] = path;
        pathpos[u] = pathlen[path]++;
        for (int v : adj[u]) {
            if (v != parent[u]) {
                build_paths(v, (2 * size[v] >= size[u]) ? path : new_path(v));
            }
        }
    }

    inline T applied_value(int path, int i) {
        return pending[path][i] ? apply_delta(value[path][i], delta[path][i], len[path][i])
            : value[path][i];
    }

    void push_delta(int path, int i) {
        int d = 0;
        while ((i >> d) > 0) {
```



```

    d++;
}
for (d -= 2; d >= 0; d--) {
    int l = (i >> d), r = (1 ^ 1), n = 1 / 2;
    if (pending[path][n]) {
        value[path][n] = applied_value(path, n);
        delta[path][l] =
            pending[path][l] ? compose_deltas(delta[path][l], delta[path][n]) : delta[path][n];
        delta[path][r] =
            pending[path][r] ? compose_deltas(delta[path][r], delta[path][n]) : delta[path][n];
        pending[path][l] = pending[path][r] = true;
        pending[path][n] = false;
    }
}
}

bool query(int path, int u, int v, T *res) {
    push_delta(path, u += static_cast<int>(value[path].size() / 2));
    push_delta(path, v += static_cast<int>(value[path].size() / 2));
    bool found = false;
    for (; u <= v; u = (u + 1) / 2, v = (v - 1) / 2) {
        if ((u & 1) != 0) {
            T value = applied_value(path, u);
            *res = found ? combine(*res, value) : value;
            found = true;
        }
        if ((v & 1) == 0) {
            T value = applied_value(path, v);
            *res = found ? combine(*res, value) : value;
            found = true;
        }
    }
    return found;
}

void update(int path, int u, int v, const T &d) {
    push_delta(path, u += static_cast<int>(value[path].size() / 2));
    push_delta(path, v += static_cast<int>(value[path].size() / 2));
    int tu = -1, tv = -1;
    for (; u <= v; u = (u + 1) / 2, v = (v - 1) / 2) {
        if ((u & 1) != 0) {
            delta[path][u] = pending[path][u] ? compose_deltas(delta[path][u], d) : d;
            pending[path][u] = true;
            if (tu == -1) {
                tu = u;
            }
        }
        if ((v & 1) == 0) {
            delta[path][v] = pending[path][v] ? compose_deltas(delta[path][v], d) : d;
            pending[path][v] = true;
            if (tv == -1) {
                tv = v;
            }
        }
    }
    for (int i = tu; i > 1; i /= 2) {
        value[path][i / 2] = combine(applied_value(path, i), applied_value(path, i ^ 1));
    }
    for (int i = tv; i > 1; i /= 2) {
        value[path][i / 2] = combine(applied_value(path, i), applied_value(path, i ^ 1));
    }
}

inline bool is_ancestor(int parent, int child) {
    return (tin[parent] <= tin[child]) && (tout[child] <= tout[parent]);
}

public:

```

```

explicit HeavyLight(const std::vector<std::vector<int>> &adj, const T &initial_value = T())
: adj(adj),
  size(adj.size()),
  parent(adj.size()),
  root(adj.size(), -1),
  tin(adj.size()),
  tout(adj.size()),
  path(adj.size()),
  pathlen(adj.size()),
  pathpos(adj.size()),
  pathroot(adj.size()),
  counter(0),
  paths(0) {
for (int u = 0; u < static_cast<int>(adj.size()); u++) {
    if (root[u] == -1) {
        dfs(u, -1, u);
        build_paths(u, new_path(u));
    }
}
value.resize(paths);
delta.resize(paths);
pending.resize(paths);
len.resize(paths);
for (int i = 0; i < paths; i++) {
    int m = pathlen[i];
    value[i].assign(2 * m, initial_value);
    delta[i].resize(2 * m);
    pending[i].assign(2 * m, false);
    len[i].assign(2 * m, 1);
    for (int j = 2 * m - 1; j > 1; j -= 2) {
        value[i][j / 2] = combine(value[i][j], value[i][j ^ 1]);
        len[i][j / 2] = len[i][j] + len[i][j ^ 1];
    }
}
}

template<typename Fn>
bool for_each_path(int u, int v, bool include_lca, Fn f) {
    assert(0 <= u && u < static_cast<int>(root.size()));
    assert(0 <= v && v < static_cast<int>(root.size()));
    if (root[u] != root[v]) {
        return false;
    }
    int path_root;
    while (!is_ancestor(path_root = pathroot[path[u]], v)) {
        f(path[u], 0, pathpos[u], true);
        u = parent[path_root];
    }
    std::vector<std::tuple<int, int, int>> down_parts;
    while (!is_ancestor(path_root = pathroot[path[v]], u)) {
        down_parts.emplace_back(path[v], 0, pathpos[v]);
        v = parent[path_root];
    }
    int lo = std::min(pathpos[u], pathpos[v]) + static_cast<int>(!include_lca);
    int hi = std::max(pathpos[u], pathpos[v]);
    if (lo <= hi) {
        f(path[u], lo, hi, pathpos[u] > pathpos[v]);
    }
    for (int i = static_cast<int>(down_parts.size()) - 1; i >= 0; i--) {
        int path, lo, hi;
        std::tie(path, lo, hi) = down_parts[i];
        f(path, lo, hi, false);
    }
    return true;
}

T query(int u, int v) {
    assert(0 <= u && u < static_cast<int>(root.size()));

```

```

    assert(0 <= v && v < static_cast<int>(root.size()));
    assert(root[u] == root[v]);
    assert(!VALUES_ON_EDGES || u != v);
    bool found = false;
    T res = T(), value;
    for_each_path(u, v, !VALUES_ON_EDGES, [&](int path, int lo, int hi, bool) {
        if (query(path, lo, hi, &value)) {
            res = found ? combine(res, value) : value;
            found = true;
        }
    });
    assert(found);
    return res;
}

void update(int u, int v, const T &d) {
    assert(0 <= u && u < static_cast<int>(root.size()));
    assert(0 <= v && v < static_cast<int>(root.size()));
    assert(root[u] == root[v]);
    if (VALUES_ON_EDGES && u == v) {
        return;
    }
    for_each_path(u, v, !VALUES_ON_EDGES, [&](int path, int lo, int hi, bool) {
        update(path, lo, hi, d);
    });
}
};

```

Example Usage

```

using namespace std;

void add_edge(vector<vector<int>> &adj, int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
}

int main() {
    //   v=40   v=20   v=10
    // 0-----1-----2-----3   5-----6
    //           |
    //           +-----4
    //           v=30
    vector<vector<int>> adj(7);
    add_edge(adj, 0, 1);
    add_edge(adj, 1, 2);
    add_edge(adj, 2, 3);
    add_edge(adj, 2, 4);
    add_edge(adj, 5, 6);
    HeavyLight<int> hld(adj, 0);
    hld.update(0, 1, 40);
    hld.update(1, 2, 20);
    hld.update(2, 3, 10);
    hld.update(2, 4, 30);
    assert(hld.query(0, 3) == 10);
    assert(hld.query(2, 4) == 30);
    hld.update(3, 4, 5);

    // Update path 3-2-4 to value 5:
    //   v=40   v=20   v=5
    // 0-----1-----2-----3   5-----6
    //           |
    //           +-----4
    //           v=5
    assert(hld.query(1, 4) == 5);
}

```

```

hld.update(5, 6, 7);

// The disconnected component can be updated independently:
//
// 5-----6
//   v=7
assert(hld.query(5, 6) == 7);
assert(!hld.for_each_path(0, 5, true, [](int, int, int, bool) {}));

// v=100 v=100 v=100 v=100 v=100 v=100
// 0-----1-----2-----3      5-----6
//                               |
//                               +-----4
//                               v=100
HeavyLight<int, false> node_hld(adj, 100);
node_hld.update(1, 4, 8);

// Update nodes on path 1-2-4 to value 8:
//
// v=100 v=8 v=8 v=100 v=100 v=100
// 0-----1-----2-----3      5-----6
//                               |
//                               +-----4
//                               v=8
assert(node_hld.query(0, 3) == 8); // Min over nodes 0, 1, 2, 3.
assert(node_hld.query(5, 5) == 100); // Single-node paths are allowed.
node_hld.update(5, 5, 7);
assert(node_hld.query(5, 5) == 7);
return 0;
}

```

2.7.9 Euler Tour Tree

Maintain an undirected forest dynamically using Euler Tour Trees. Each tree is represented by a sequence of directed edge occurrences in an implicit treap: an undirected edge (u, v) contributes two nodes, $(u \rightarrow v)$ and $(v \rightarrow u)$. Rerooting a represented tree is a cyclic rotation of this sequence, while `link` concatenates two tours with the new edge occurrences and `cut` removes the two occurrences of an existing edge, splitting one tour into two.

This version supports dynamic connectivity, `link`, `cut`, rerooting, and listing the nodes in a connected component. It does not store node aggregates. To add component aggregates, a common extension is to insert one self-loop occurrence per node and maintain aggregate data in the treap.

- `EulerTourTree(n)` constructs a forest on nodes numbered $[0, n)$.
- `connected(u, v)` returns whether `u` and `v` are in the same tree.
- `link(u, v)` adds an edge between different trees and returns false if the edge would create a cycle or already exists.
- `cut(u, v)` removes an existing edge and returns false if it is not present.
- `reroot(u)` rotates the Euler tour of `u`'s tree so an occurrence leaving `u` is first.
- `component_nodes(u)` returns the distinct nodes in `u`'s tree in unspecified order.

Time Complexity:

- $O(\log m)$ expected per call to `connected()`, `link()`, `cut()`, and `reroot()`, where m is the number of edges currently in the forest.
- $O(k)$ expected per call to `component_nodes()`, where k is the number of directed edge occurrences in the component.

Space Complexity:

- $O(n + m)$ for the forest.
- $O(\log m)$ auxiliary recursion stack for treap operations.

```

#include <cassert>
#include <random>
#include <unordered_map>
#include <unordered_set>
#include <utility>
#include <vector>

class EulerTourTree {
    struct Node {
        int u, v, priority, size;
        Node *left, *right, *parent;

        Node(int u, int v, int priority)
            : u(u), v(v), priority(priority), size(1), left(nullptr), right(nullptr), parent(nullptr) {}
    };

    int n;
    std::mt19937 rng;
    std::vector<std::unordered_map<int, Node *>> adj;
    std::unordered_set<Node *> allocated;

    static int size(Node *t) { return t == nullptr ? 0 : t->size; }

    static void set_parent(Node *t, Node *p) {
        if (t != nullptr) {
            t->parent = p;
        }
    }

    static void update(Node *t) {
        if (t != nullptr) {
            t->size = 1 + size(t->left) + size(t->right);
            set_parent(t->left, t);
            set_parent(t->right, t);
        }
    }

    static Node *root(Node *t) {
        if (t == nullptr) {
            return nullptr;
        }
        while (t->parent != nullptr) {
            t = t->parent;
        }
        return t;
    }

    static int index(Node *t) {
        int res = size(t->left);
        while (t->parent != nullptr) {
            if (t == t->parent->right) {
                res += 1 + size(t->parent->left);
            }
            t = t->parent;
        }
        return res;
    }

    static Node *merge(Node *a, Node *b) {
        if (a == nullptr) {
            return b;
        }
        if (b == nullptr) {
            return a;
        }
        if (a->priority > b->priority) {
            a->right = merge(a->right, b);
            update(a);
        }
    }
};

```

```

    a->parent = nullptr;
    return a;
}
b->left = merge(a, b->left);
update(b);
b->parent = nullptr;
return b;
}

static std::pair<Node *, Node *> split(Node *t, int left_size) {
    if (t == nullptr) {
        return {nullptr, nullptr};
    }
    if (size(t->left) >= left_size) {
        auto [a, b] = split(t->left, left_size);
        t->left = b;
        update(t);
        set_parent(a, nullptr);
        t->parent = nullptr;
        return {a, t};
    }
    auto [a, b] = split(t->right, left_size - size(t->left) - 1);
    t->right = a;
    update(t);
    set_parent(b, nullptr);
    t->parent = nullptr;
    return {t, b};
}

static Node *make_first(Node *x) {
    auto [a, b] = split(root(x), index(x));
    return merge(b, a);
}

static Node *remove_first(Node *t) {
    auto [first, rest] = split(t, 1);
    if (first != nullptr) {
        first->left = first->right = first->parent = nullptr;
        first->size = 1;
    }
    return rest;
}

Node *new_node(int u, int v) {
    Node *node = new Node(u, v, std::uniform_int_distribution<int>()(rng));
    allocated.insert(node);
    return node;
}

Node *any_occurrence(int u) const { return adj[u].empty() ? nullptr : adj[u].begin()->second; }

static void collect(Node *t, std::unordered_set<int> &out) {
    if (t == nullptr) {
        return;
    }
    collect(t->left, out);
    out.insert(t->u);
    collect(t->right, out);
}

public:
explicit EulerTourTree(int n) : n(n), rng(std::random_device{}()), adj(n) {}

~EulerTourTree() {
    for (Node *node : allocated) {
        delete node;
    }
}

```

```

EulerTourTree(const EulerTourTree &) = delete;
EulerTourTree &operator=(const EulerTourTree &) = delete;

bool connected(int u, int v) const {
    assert(0 <= u && u < n && 0 <= v && v < n);
    if (u == v) {
        return true;
    }
    Node *a = any_occurrence(u), *b = any_occurrence(v);
    return a != nullptr && b != nullptr && root(a) == root(b);
}

void reroot(int u) {
    assert(0 <= u && u < n);
    Node *a = any_occurrence(u);
    if (a != nullptr) {
        make_first(a);
    }
}

bool link(int u, int v) {
    assert(0 <= u && u < n && 0 <= v && v < n);
    if (u == v || adj[u].count(v) || connected(u, v)) {
        return false;
    }
    Node *tu = any_occurrence(u), *tv = any_occurrence(v);
    if (tu != nullptr) {
        tu = make_first(tu);
    }
    if (tv != nullptr) {
        tv = make_first(tv);
    }
    Node *uv = new_node(u, v), *vu = new_node(v, u);
    adj[u][v] = uv;
    adj[v][u] = vu;
    merge(merge(tu, uv), merge(tv, vu));
    return true;
}

bool cut(int u, int v) {
    assert(0 <= u && u < n && 0 <= v && v < n);
    auto it = adj[u].find(v);
    if (it == adj[u].end()) {
        return false;
    }
    Node *uv = it->second, *vu = adj[v].at(u);
    make_first(uv);
    auto [left, right] = split(root(uv), index(vu));
    remove_first(left); // Remove uv.
    remove_first(right); // Remove vu.
    adj[u].erase(v);
    adj[v].erase(u);
    allocated.erase(uv);
    allocated.erase(vu);
    delete uv;
    delete vu;
    return true;
}

std::vector<int> component_nodes(int u) const {
    assert(0 <= u && u < n);
    Node *a = any_occurrence(u);
    if (a == nullptr) {
        return {u};
    }
    std::unordered_set<int> nodes;
    collect(root(a), nodes);
}

```

```

    return std::vector<int>(nodes.begin(), nodes.end());
}
};

```

Example Usage

```

#include <algorithm>

int main() {
    EulerTourTree ett(6);
    assert(!ett.connected(0, 1));
    assert(ett.connected(0, 0));

    assert(ett.link(0, 1));
    assert(ett.link(1, 2));
    assert(ett.link(3, 4));

    // 0---1---2    3---4    5
    assert(ett.connected(0, 2));
    assert(!ett.connected(0, 4));
    assert(!ett.link(0, 2)); // Would create a cycle.

    auto c = ett.component_nodes(1);
    std::sort(c.begin(), c.end());
    assert((c == std::vector<int>{0, 1, 2}));

    ett.reroot(2);
    assert(ett.connected(0, 2));
    assert(ett.cut(1, 2));

    // 0---1    2    3---4    5
    assert(!ett.connected(0, 2));
    assert((ett.component_nodes(2) == std::vector<int>{2}));

    assert(ett.link(2, 3));

    // 0---1    2---3---4    5
    assert(ett.connected(0, 4) == false);
    assert(ett.connected(2, 4));
    assert(!ett.cut(0, 5));
    return 0;
}

```

2.7.10 Link-Cut Tree

Maintain a forest of trees with values associated with its nodes, while supporting both dynamic queries and dynamic updates of all values on any path between two nodes in a given tree. In addition, support testing of whether two nodes are connected in the forest, as well as the merging and splitting of trees by adding or removing specific edges. Link/cut forests divide each of its trees into node-disjoint paths, each represented by a splay tree.

The query operation is defined by an associative `combine()` function. The default code below assumes a numerical forest type, defining queries for the “min” of the target range. Another possible query operation is “sum”, in which case `combine(a, b)` should return `a + b`. For direction-independent path queries, `combine()` should also be commutative; otherwise, store enough information in each aggregate to combine paths in the required order.

The update operation is defined by `apply_delta()` and `compose_deltas()`. A delta must act on an aggregate summary of a path of length `len`: `apply_delta(v, d, len)` returns the aggregate after applying update `d` to every element represented by aggregate `v`. Pending deltas are combined in chronological order by `compose_deltas(old, d)`, meaning “apply `old`, then apply `d`”. These hooks do not support arbitrary query/update pairings; the delta operation must distribute over `combine()`, and composed deltas must have the same effect as applying their updates sequentially. The default code below defines updates that “set” a path’s nodes to a new

value. For range increment updates, `apply_delta(v, d, len)` would return `v + d` for min/max queries, or `v + d * len` for sum queries, and `compose_deltas(old, d)` would return `old + d`.

- `LinkCut<T>()` constructs an empty forest.
- `size()` returns the number of nodes in the forest.
- `trees()` returns the number of trees in the forest.
- `add_node(i, value = T())` adds a new single-node tree to the forest, labeled with the integer `i` and with value initialized to `value`.
- `is_connected(a, b)` returns whether nodes `a` and `b` are connected.
- `link(a, b)` adds an edge between the nodes `a` and `b`, both of which must exist and not be connected.
- `cut(a, b)` removes the edge between the nodes `a` and `b`, both of which must exist and be connected.
- `query(a, b)` returns the result of `combine()` applied to all values on the path from the node `a` to node `b`.
- `update(a, b, d)` modifies all the values on the path from node `a` to node `b` by applying the delta `d`.
- `reroot(i)` makes node `i` the root of its tree.
- `find_root(i)` returns the label of the root of the tree containing node `i`.
- `lca(a, b)` returns the lowest common ancestor of `a` and `b` relative to the tree's current root.

The forest is unrooted: a tree's root is whichever node was most recently established as such. `reroot(i)` sets it explicitly, while `link`, `query`, and `update` reroot implicitly (the latter two at their first argument). `find_root(i)` and `lca(a, b)` are therefore relative to the current root, so call `reroot(r)` first whenever a specific root `r` is intended.

Time Complexity:

- $O(1)$ per call to the constructor, `size()`, and `trees()`.
- $O(\log n)$ amortized per call to all other operations, where n is the number of nodes.

Space Complexity:

- $O(n)$ for storage of the forest, where n is the number of nodes.
- $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cassert>
#include <unordered_map>
#include <utility>

template<typename T>
class LinkCut {
    static T combine(const T &a, const T &b) { return std::min(a, b); }
    static T apply_delta(const T &v, const T &d, int len) { return d; }
    static T compose_deltas(const T &old, const T &d) { return d; }

    struct Node {
        int id;
        T value, subtree_value, delta;
        int size;
        bool rev, pending;
        Node *left, *right, *parent;

        Node(int id, const T &value)
            : id(id),
              value(value),
              subtree_value(value),
              size(1),
              rev(false),
              pending(false),
              left(nullptr),
              right(nullptr),
              parent(nullptr) {}

        inline bool is_root() const {
            return parent == nullptr || (parent->left != this && parent->right != this);
        }
    };
};
```

```

inline T get_subtree_value() const {
    return pending ? apply_delta(subtree_value, delta, size) : subtree_value;
}

void push() {
    if (rev) {
        rev = false;
        std::swap(left, right);
        if (left != nullptr) {
            left->rev = !left->rev;
        }
        if (right != nullptr) {
            right->rev = !right->rev;
        }
    }
    if (pending) {
        value = apply_delta(value, delta, 1);
        subtree_value = apply_delta(subtree_value, delta, size);
        if (left != nullptr) {
            left->delta = left->pending ? compose_deltas(left->delta, delta) : delta;
            left->pending = true;
        }
        if (right != nullptr) {
            right->delta = right->pending ? compose_deltas(right->delta, delta) : delta;
            right->pending = true;
        }
        pending = false;
    }
}

void update() {
    size = 1;
    subtree_value = value;
    if (left != nullptr) {
        // Combine in in-order (left, value, right) so non-commutative aggregates stay correct.
        subtree_value = combine(left->get_subtree_value(), subtree_value);
        size += left->size;
    }
    if (right != nullptr) {
        subtree_value = combine(subtree_value, right->get_subtree_value());
        size += right->size;
    }
}

};

int num_trees;
std::unordered_map<int, Node *> nodes;

static void connect(Node *child, Node *parent, bool is_left) {
    if (child != nullptr) {
        child->parent = parent;
    }
    if (is_left) {
        parent->left = child;
    } else {
        parent->right = child;
    }
}

static void rotate(Node *n) {
    Node *parent = n->parent, *grandparent = parent->parent;
    bool parent_is_root = parent->is_root(), is_left = (n == parent->left);
    connect(is_left ? n->right : n->left, parent, is_left);
    connect(parent, n, !is_left);
    if (parent_is_root) {
        if (n != nullptr) {
            n->parent = grandparent;
        }
    }
}

```

```

    }
} else {
    connect(n, grandparent, parent == grandparent->left);
}
parent->update();
}

static void splay(Node *n) {
    while (!n->is_root()) {
        Node *parent = n->parent, *grandparent = parent->parent;
        if (!parent->is_root()) {
            grandparent->push();
        }
        parent->push();
        n->push();
        if (!parent->is_root()) {
            if ((n == parent->left) == (parent == grandparent->left)) {
                rotate(parent);
            } else {
                rotate(n);
            }
        }
        rotate(n);
    }
    n->push();
    n->update();
}

// Convention: a node's left child leads toward the tree's root, so an exposed path is ordered by
// depth with the root as its leftmost node (matching the standard link/cut tree presentation).
// cut() and find_root() rely on this orientation.
static Node *expose(Node *n) {
    Node *prev = nullptr;
    for (Node *curr = n; curr != nullptr; curr = curr->parent) {
        splay(curr);
        curr->right = prev;
        curr->update();
        prev = curr;
    }
    splay(n);
    return prev;
}

// Helper variables.
Node *u, *v;

void get_uv(int a, int b) {
    auto it1 = nodes.find(a);
    auto it2 = nodes.find(b);
    assert(it1 != nodes.end() && it2 != nodes.end());
    u = it1->second;
    v = it2->second;
}

public:
    LinkCut() : num_trees(0) {}

    ~LinkCut() {
        for (auto &[key, node] : nodes) {
            delete node;
        }
    }

    LinkCut(const LinkCut &) = delete;
    LinkCut &operator=(const LinkCut &) = delete;
    int size() const { return static_cast<int>(nodes.size()); }
    int trees() const { return num_trees; }

```

```

void add_node(int i, const T &value = T()) {
    assert(nodes.find(i) == nodes.end());
    Node *n = new Node(i, value);
    expose(n);
    n->rev = !n->rev;
    nodes[i] = n;
    num_trees++;
}

bool is_connected(int a, int b) {
    get_uv(a, b);
    if (a == b) {
        return true;
    }
    expose(u);
    expose(v);
    return u->parent != nullptr;
}

void link(int a, int b) {
    assert(!is_connected(a, b));
    get_uv(a, b);
    expose(u);
    u->rev = !u->rev;
    u->parent = v;
    num_trees--;
}

void cut(int a, int b) {
    get_uv(a, b);
    expose(u);
    u->rev = !u->rev;
    expose(v);
    assert(v->left == u && u->right == nullptr);
    v->left->parent = nullptr;
    v->left = nullptr;
    num_trees++;
}

T query(int a, int b) {
    assert(is_connected(a, b));
    get_uv(a, b);
    expose(u);
    u->rev = !u->rev;
    expose(v);
    return v->get_subtree_value();
}

void update(int a, int b, const T &d) {
    assert(is_connected(a, b));
    get_uv(a, b);
    expose(u);
    u->rev = !u->rev;
    expose(v);
    v->delta = v->pending ? compose_deltas(v->delta, d) : d;
    v->pending = true;
}

void reroot(int i) {
    auto it = nodes.find(i);
    assert(it != nodes.end());
    expose(it->second);
    it->second->rev = !it->second->rev;
}

int find_root(int i) {
    auto it = nodes.find(i);
    assert(it != nodes.end());

```

```

Node *n = it->second;
expose(n);
while (n->left != nullptr) { // The leftmost node of the exposed path is the tree's root.
    n = n->left;
    n->push();
}
splay(n);
return n->id;
}

int lca(int a, int b) {
    get_uv(a, b);
    if (a == b) {
        return a;
    }
    expose(u);
    expose(v);
    assert(u->parent != nullptr);
    splay(u); // u->parent is now the node where u's path rejoins v's exposed path, i.e. the LCA.
    return u->parent != nullptr ? u->parent->id : u->id;
}
};

```

Example Usage

```

using namespace std;

int main() {
    // v=10      v=40      v=20      v=10
    // 0-----1-----2-----3
    //           |
    //           +-----4
    //                   v=30
    LinkCut<int> lcf;
    lcf.add_node(0, 10);
    lcf.add_node(1, 40);
    lcf.add_node(2, 20);
    lcf.add_node(3, 10);
    lcf.add_node(4, 30);
    assert(lcf.size() == 5);
    assert(lcf.trees() == 5);
    lcf.link(0, 1);
    lcf.link(1, 2);
    lcf.link(2, 3);
    lcf.link(2, 4);
    assert(lcf.trees() == 1);
    assert(lcf.query(1, 4) == 20);

    // find_root and lca are relative to the current root, set explicitly via reroot.
    lcf.reroot(0);
    assert(lcf.find_root(3) == 0 && lcf.find_root(4) == 0);
    assert(lcf.lca(3, 4) == 2); // Paths 3-2 and 4-2 meet at 2 under root 0.
    assert(lcf.lca(1, 4) == 1); // 1 is an ancestor of 4 under root 0.
    lcf.reroot(4);
    assert(lcf.lca(0, 3) == 2); // Rerooting at 4 changes the LCA of 0 and 3.

    lcf.update(1, 1, 100);
    lcf.update(2, 4, 100);

    // v=10      v=100      v=100      v=10
    // 0-----1-----2-----3
    //           |
    //           +-----4
    //                   v=100
    assert(lcf.query(4, 4) == 100);
    assert(lcf.query(0, 4) == 10);
}

```

```

assert(lcf.query(3, 4) == 10);
lcf.cut(1, 2);

// v=10      v=100      v=100      v=0
// 0-----1      2-----3
//              |
//              +-----4
//                      v=100
assert(lcf.trees() == 2);
assert(!lcf.is_connected(1, 2));
assert(!lcf.is_connected(0, 4));
assert(lcf.is_connected(2, 3));
return 0;
}

```

2.7.11 Top Tree

Maintain a dynamic forest while supporting both path aggregates and rooted-subtree aggregates. Top trees represent each tree as a hierarchy of clusters. A path cluster summarizes a contiguous path, while non-path clusters are raked onto nodes so that changing the preferred path only requires a logarithmic number of local rotations and recomputations.

This is an advanced alternative to link/cut trees. Link/cut trees are usually simpler for dynamic path queries; top trees are useful when the same dynamic forest also needs rooted-subtree queries.

The aggregate operation is defined by an identity element `identity()` and an associative `combine(a, b)`. The default code below computes sums over node and edge values. For minimum queries, use `std::numeric_limits<T>::max()` as the identity and `std::min(a, b)` as the combine function. For non-commutative aggregates, store enough information in `T` to support reversing a path, and update `flip_path()` accordingly.

- `TopTree(n, value = T())` constructs a forest on nodes $[0, n)$ with every node value initialized to `value`.
- `size()` returns the number of nodes in the forest.
- `edges()` returns the number of edge nodes that have been created.
- `connected(u, v)` returns whether nodes `u` and `v` are in the same tree.
- `link(u, v, value = T())` adds an edge with value `value` between different trees, returning its edge id or -1 if the edge would create a cycle.
- `cut(e)` removes edge id `e`. The edge id must currently be present.
- `set_node(u, value)` changes node `u`'s value to `value`.
- `set_edge(e, value)` changes edge `e`'s value to `value`. The edge id must currently be present.
- `path_query(u, v)` returns the aggregate over all node and edge values on the path from `u` to `v`. The nodes must be connected.
- `subtree_query(root, u)` reroots the represented tree at `root` and returns the aggregate over the rooted subtree of `u`.
- `reroot(u)` makes node `u` the root of its represented tree.
- `lca(u, v)` returns the lowest common ancestor of `u` and `v` relative to the current root. The nodes must be connected.
- `maybe_lca(u, v)` returns the lowest common ancestor of `u` and `v`, or -1 if they are disconnected.

Time Complexity:

- $O(n)$ per call to the constructor, and $O(1)$ per call to `size()` and `edges()`.
- $O(\log n)$ amortized per call to all other operations, where n is the number of nodes and edges in the represented tree.

Space Complexity:

- $O(n + m)$ for the internal nodes representing graph nodes and edges, where m is the number of edges ever added.
- $O(\log n)$ auxiliary stack space for all operations other than `size()` and `edges()`.

```

#include <array>
#include <cassert>
#include <utility>
#include <vector>

template<typename T>
class TopTree {
    static T identity() { return T(); }
    static T combine(const T &a, const T &b) { return a + b; }

    struct Node {
        Node *p;
        std::array<Node *, 3> c;
        T value, path_value, subtree_value;
        int graph_node_id, edge_id;
        bool is_path, is_graph_node, lazy_flip_path, alive;

        Node(const T &value, int graph_node_id, int edge_id)
            : p(nullptr),
              c{nullptr, nullptr, nullptr},
              value(value),
              path_value(value),
              subtree_value(value),
              graph_node_id(graph_node_id),
              edge_id(edge_id),
              is_path(graph_node_id != -1),
              is_graph_node(graph_node_id != -1),
              lazy_flip_path(false),
              alive(graph_node_id != -1) {}

        static T path_value_of(Node *n) { return n == nullptr ? identity() : n->path_value; }
        static T subtree_value_of(Node *n) { return n == nullptr ? identity() : n->subtree_value; }

        int dir() const {
            assert(p != nullptr);
            if (this == p->c[0]) {
                return 0;
            } else if (this == p->c[1]) {
                return 1;
            } else if (this == p->c[2]) {
                return 2;
            }
            assert(false);
            return -1;
        }
    };

    Node *parent_child() const { return p->c[dir()]; }
    bool root_of_auxiliary_tree() const { return p == nullptr || p->is_path != is_path; }

    void flip_path() {
        assert(is_path);
        std::swap(c[0], c[1]);
        lazy_flip_path = !lazy_flip_path;
    }

    void push() {
        if (lazy_flip_path) {
            assert(is_path);
            if (!is_graph_node) {
                c[0]->flip_path();
                c[1]->flip_path();
            }
            lazy_flip_path = false;
        }
    }

    void push_all() {
        if (p != nullptr) {

```

```

    p->push_all();
}
push();
}

Node *pull_all() {
    Node *cur = this;
    cur->pull();
    while (cur->p != nullptr) {
        cur = cur->p;
        cur->pull();
    }
    return cur;
}

void pull() {
    if (is_graph_node) {
        path_value = value;
        subtree_value = combine(combine(value, subtree_value_of(c[0])), subtree_value_of(c[1]));
    } else if (is_path) {
        path_value = combine(combine(path_value_of(c[0]), value), path_value_of(c[1]));
        subtree_value = combine(combine(subtree_value_of(c[0]), value), subtree_value_of(c[1]));
    } else {
        path_value = identity();
        subtree_value = combine(
            combine(subtree_value_of(c[0]), value),
            combine(subtree_value_of(c[2]), subtree_value_of(c[1]))
        );
    }
}

void rotate() {
    assert(!is_graph_node);
    assert(!root_of_auxiliary_tree());
    Node *parent = p;
    int x = dir();
    assert(x == 0 || x == 1);
    Node *child = c[!x];
    if (parent->p != nullptr) {
        parent->parent_child() = this;
    }
    p = parent->p;
    parent->c[x] = child;
    if (child != nullptr) {
        child->p = parent;
    }
    c[!x] = parent;
    parent->p = this;
    parent->pull();
}

void rotate_2(int child_dir) {
    assert(!is_graph_node);
    assert(!root_of_auxiliary_tree());
    assert(c[child_dir] != nullptr);
    assert(!c[child_dir]->is_graph_node);
    if (dir() == child_dir) {
        rotate();
        return;
    }
    Node *parent = p;
    int x = dir();
    assert(x == 0 || x == 1);
    assert(child_dir == !x);
    Node *child = c[child_dir]->c[!x];
    if (parent->p != nullptr) {
        parent->parent_child() = this;
    }
}

```



```

    p = parent->p;
    parent->c[x] = child;
    if (child != nullptr) {
        child->p = parent;
    }
    c[child_dir]->c[!x] = parent;
    parent->p = c[child_dir];
    parent->pull();
}

void splay_dir(int x) {
    while (!root_of_auxiliary_tree() && dir() == x) {
        if (!p->root_of_auxiliary_tree() && p->dir() == x) {
            p->rotate();
        }
        rotate();
    }
}

void splay_2(int child_dir) {
    assert(!is_graph_node && is_path);
    assert(c[child_dir] != nullptr && !c[child_dir]->is_graph_node);
    while (!root_of_auxiliary_tree()) {
        if (!p->root_of_auxiliary_tree()) {
            if (p->dir() == dir()) {
                p->rotate();
            } else {
                rotate_2(child_dir);
            }
        }
        rotate_2(child_dir);
    }
}

void splay_2() {
    assert(!is_graph_node && is_path);
    assert(!root_of_auxiliary_tree());
    p->splay_2(dir());
}

void splay_graph_node() {
    assert(is_graph_node);
    if (root_of_auxiliary_tree()) {
        return;
    }
    p->splay_dir(dir());
    if (p->root_of_auxiliary_tree()) {
        return;
    }
    assert(p->dir() != dir());
    if (dir() == 1) {
        p->rotate();
    }
    assert(dir() == 0);
    p->splay_2();
    assert(dir() == 0);
    assert(p->dir() == 1);
    assert(p->p->root_of_auxiliary_tree());
}

void splay() {
    assert(!is_graph_node);
    while (!root_of_auxiliary_tree()) {
        if (!p->root_of_auxiliary_tree()) {
            if (p->dir() == dir()) {
                p->rotate();
            } else {
                rotate();
            }
        }
    }
}

```

```

    }
  }
  rotate();
}
}

Node *cut_right() {
  assert(is_graph_node && is_path);
  splay_graph_node();
  if (root_of_auxiliary_tree() || dir() == 1) {
    assert(root_of_auxiliary_tree() || (dir() == 1 && p->root_of_auxiliary_tree()));
    assert(c[0] == nullptr);
    return nullptr;
  }
  Node *parent = p;
  assert(
    parent->root_of_auxiliary_tree() ||
    (parent->dir() == 1 && parent->p->root_of_auxiliary_tree())
  );
  assert(!parent->is_graph_node);
  assert(parent->is_path);
  assert(parent->c[0] == this);
  assert(parent->c[2] == nullptr);
  if (parent->p != nullptr) {
    parent->parent_child() = this;
  }
  p = parent->p;
  parent->is_path = false;
  parent->c[2] = parent->c[1];
  parent->c[0] = c[0];
  if (c[0] != nullptr) {
    c[0]->p = parent;
  }
  parent->c[1] = c[1];
  if (c[1] != nullptr) {
    c[1]->p = parent;
  }
  c[0] = nullptr;
  c[1] = parent;
  parent->p = this;
  assert(c[2] == nullptr);
  assert(c[0] == nullptr);
  parent->pull();
  return parent;
}

Node *splice_non_path() {
  assert(!is_path);
  assert(!is_graph_node);
  splay();
  assert(p != nullptr && p->is_graph_node && p->is_path);
  p->cut_right();
  if (!p->is_path) {
    rotate();
  }
  assert(p != nullptr && p->is_graph_node && p->is_path);
  assert(p->root_of_auxiliary_tree() || (p->dir() == 1 && p->p->root_of_auxiliary_tree()));
  assert(p->c[dir()] == this && p->c[!dir()] == nullptr);
  Node *parent = p;
  if (parent->p != nullptr) {
    parent->parent_child() = this;
  }
  p = parent->p;
  parent->c[0] = c[0];
  if (c[0] != nullptr) {
    c[0]->p = parent;
  }
  parent->c[1] = c[1];

```

```

    if (c[1] != nullptr) {
        c[1]->p = parent;
    }
    assert(c[2] != nullptr && c[2]->is_path);
    c[1] = c[2];
    c[0] = parent;
    parent->p = this;
    c[2] = nullptr;
    is_path = true;
    parent->pull();
    return parent;
}

Node *splice_all() {
    Node *res = this;
    for (Node *cur = this; cur != nullptr; cur = cur->p) {
        if (!cur->is_path) {
            res = cur->splice_non_path();
        }
        assert(cur->is_path);
    }
    return res;
}

Node *expose() {
    assert(is_graph_node);
    push_all();
    Node *res = splice_all();
    cut_right();
    pull_all();
    return res;
}

Node *expose_edge() {
    assert(!is_graph_node);
    push_all();
    Node *v = is_path ? c[1] : c[2];
    v->push();
    while (!v->is_graph_node) {
        v = v->c[0];
        v->push();
    }
    Node *res = v->splice_all();
    v->cut_right();
    v->pull_all();
    assert(p == nullptr);
    assert(v == c[1]);
    return res == v ? this : res;
}

Node *meld_path_end() {
    assert(p == nullptr);
    Node *rt = this;
    while (true) {
        rt->push();
        if (rt->is_graph_node) {
            break;
        }
        rt = rt->c[1];
    }
    assert(rt->is_graph_node);
    rt->splay_graph_node();
    if (rt->c[0] != nullptr && rt->c[1] != nullptr) {
        Node *child = rt->c[1];
        while (true) {
            child->push();
            if (child->c[0] == nullptr) {
                break;
            }
        }
    }
}

```

```

    }
    child = child->c[0];
  }
  child->splay();
  assert(child->c[0] == nullptr);
  child->c[0] = rt->c[0];
  child->c[0]->p = child;
  rt->c[0] = nullptr;
  child->pull();
} else if (rt->c[0] != nullptr) {
  rt->c[1] = rt->c[0];
  rt->c[0] = nullptr;
}
assert(rt->c[0] == nullptr);
return rt->pull_all();
}

void make_root() {
  expose();
  Node *rt = this;
  while (rt->p != nullptr) {
    assert(rt->dir() == 1);
    rt = rt->p;
  }
  rt->flip_path();
  rt->meld_path_end();
  expose();
  assert(p == nullptr);
}

};

std::vector<Node *> graph_nodes, edge_nodes;

static void link_nodes(Node *edge, Node *u, Node *v) {
  assert(edge != nullptr && u != nullptr && v != nullptr);
  assert(edge->c[0] == nullptr && edge->c[1] == nullptr && edge->c[2] == nullptr);
  u->expose();
  while (u->p != nullptr) {
    u = u->p;
  }
  v->make_root();
  assert(u->p == nullptr);
  assert(v->p == nullptr);
  edge->is_path = true;
  edge->is_graph_node = false;
  edge->alive = true;
  edge->c[0] = u;
  u->p = edge;
  edge->c[1] = v;
  v->p = edge;
  edge->pull();
}

static std::pair<Node *, Node *> cut_node(Node *edge) {
  assert(edge != nullptr && edge->alive && !edge->is_graph_node);
  edge->expose_edge();
  assert(edge->p == nullptr);
  assert(edge->is_path);
  Node *l = edge->c[0], *r = edge->c[1];
  assert(l != nullptr && r != nullptr);
  edge->c[0] = edge->c[1] = edge->c[2] = nullptr;
  l->p = r->p = nullptr;
  l = l->meld_path_end();
  edge->is_path = false;
  edge->alive = false;
  edge->pull();
  return {l, r};
}

```

```

static Node *get_path(Node *u, Node *v) {
    assert(u->is_graph_node && v->is_graph_node);
    u->make_root();
    v->expose();
    if (u == v) {
        assert(v->p == nullptr);
        return v;
    }
    assert(v->p->p == nullptr);
    return v->p;
}

static Node *get_subtree(Node *root, Node *u) {
    root->make_root();
    u->expose();
    return u;
}

static Node *lca_same_component(Node *u, Node *v) {
    u->expose();
    return v->expose();
}

static Node *maybe_lca_node(Node *u, Node *v) {
    u->expose();
    Node *up = u->p;
    assert(up == nullptr || up->p == nullptr);
    Node *res = v->expose();
    assert(v->p == nullptr || v->p->p == nullptr);
    if (u != v && up == u->p && (up == nullptr || up->p == nullptr)) {
        return nullptr;
    }
    return res;
}

public:
explicit TopTree(int n, const T &value = T()) {
    assert(n >= 0);
    graph_nodes.reserve(n);
    for (int i = 0; i < n; i++) {
        graph_nodes.push_back(new Node(value, i, -1));
    }
}

~TopTree() {
    for (Node *v : graph_nodes) {
        delete v;
    }
    for (Node *e : edge_nodes) {
        delete e;
    }
}

TopTree(const TopTree &) = delete;
TopTree &operator=(const TopTree &) = delete;
int size() const { return static_cast<int>(graph_nodes.size()); }
int edges() const { return static_cast<int>(edge_nodes.size()); }

bool connected(int u, int v) {
    assert(0 <= u && u < size());
    assert(0 <= v && v < size());
    return u == v || maybe_lca_node(graph_nodes[u], graph_nodes[v]) != nullptr;
}

int link(int u, int v, const T &value = T()) {
    assert(0 <= u && u < size());
    assert(0 <= v && v < size());

```

```

    if (connected(u, v)) {
        return -1;
    }
    int id = edges();
    edge_nodes.push_back(new Node(value, -1, id));
    link_nodes(edge_nodes.back(), graph_nodes[u], graph_nodes[v]);
    return id;
}

void cut(int e) {
    assert(0 <= e && e < edges());
    cut_node(edge_nodes[e]);
}

void set_node(int u, const T &value) {
    assert(0 <= u && u < size());
    graph_nodes[u]->expose();
    graph_nodes[u]->value = value;
    graph_nodes[u]->pull_all();
}

void set_edge(int e, const T &value) {
    assert(0 <= e && e < edges());
    assert(edge_nodes[e]->alive);
    edge_nodes[e]->expose_edge();
    edge_nodes[e]->value = value;
    edge_nodes[e]->pull_all();
}

T path_query(int u, int v) {
    assert(connected(u, v));
    return get_path(graph_nodes[u], graph_nodes[v])->path_value;
}

T subtree_query(int root, int u) {
    assert(connected(root, u));
    return get_subtree(graph_nodes[root], graph_nodes[u])->subtree_value;
}

void reroot(int u) {
    assert(0 <= u && u < size());
    graph_nodes[u]->make_root();
}

int lca(int u, int v) {
    assert(connected(u, v));
    return lca_same_component(graph_nodes[u], graph_nodes[v])->graph_node_id;
}

int maybe_lca(int u, int v) {
    assert(0 <= u && u < size());
    assert(0 <= v && v < size());
    Node *res = maybe_lca_node(graph_nodes[u], graph_nodes[v]);
    return res == nullptr ? -1 : res->graph_node_id;
}
};

```

Example Usage

```

using namespace std;

int main() {
    TopTree<int> tree(5);
    for (int i = 0; i < 5; i++) {
        tree.set_node(i, i + 1);
    }
}

```

```

int e01 = tree.link(0, 1, 10);
int e12 = tree.link(1, 2, 20);
int e13 = tree.link(1, 3, 30);
int e34 = tree.link(3, 4, 40);
assert(e01 == 0 && e12 == 1 && e13 == 2 && e34 == 3);
assert(tree.link(0, 4, 100) == -1);

//      0
//  v=10 |
//      1
// v=20 / \ v=30
//      2   3
//          | v=40
//          4
assert(tree.path_query(2, 4) == 104);
assert(tree.subtree_query(0, 3) == 49);
assert(tree.lca(2, 4) == 1);

tree.cut(e13);

// Cut edge 1-3:
//
//      0
//  v=10 |
//      1       3
// v=20 /       | v=40
//      2       4
assert(!tree.connected(2, 4));
assert(tree.maybe_lca(2, 4) == -1);
e13 = tree.link(2, 4, 30);
assert(e13 == 4);

// Re-link 2-4 with value 30:
//
//      0
//  v=10 |
//      1       3
// v=20 /       | v=40
//      2-----4
//          v=30
assert(tree.path_query(0, 3) == 115);

tree.set_edge(e34, 4);
tree.set_node(4, 50);

// Change edge 3-4 to value 4 and node 4 to value 50:
//
//      0
//  v=10 |
//      1       3
// v=20 /       | v=4
//      2-----4
//          v=30  v=50
assert(tree.subtree_query(2, 4) == 58);
assert(tree.path_query(0, 3) == 124);
return 0;
}

```

Chapter 3

Strings

3.1 String Utilities

Common string functions, many of which already have standard STL equivalents. Most of the following implementations are presented for educational purposes, and are not heavily optimized. They often depend on certain `std::string` functions that have unspecified complexity.

Time Complexity: $O(n)$ per call to most operations, where n is the length of the input string or total length of all input strings. Exceptions are noted with the individual functions below.

Space Complexity:

- $O(n)$ auxiliary for operations that return a new string or vector of strings.
- $O(1)$ auxiliary for in-place operations.

```
#include <cctype>
#include <cstdint>
#include <regex>
#include <sstream>
#include <string>
#include <vector>
using std::string;
```

Integer Conversion:

- `to_str(i)` returns the string representation of integer `i`, much like `std::to_string()`.
- `to_int(s)` returns the integer representation of string `s`, much like `std::atoi()`, except here we handle special cases of overflow by throwing an exception.
- `itoa(value, &str, base = 10)` implements the non-standard C function which converts `value` into a C string, storing it into pointer `str` in the given `base`. For more generalized base conversion, see the math utilities section.
- `is_integer(s)` returns whether `s` is a valid base-10 integer literal with an optional sign.
- `is_number(s)` returns whether `s` is a decimal number literal with optional sign, decimal point, and exponent.

```
template<typename Int>
string to_str(Int i) {
    std::ostringstream oss;
    oss << i;
    return oss.str();
}
```

```
int to_int(const string &s) {
```



```

std::istringstream iss(s);
int res;
if (!(iss >> res) || !(iss >> std::ws).eof()) {
    throw std::runtime_error("to_int failed");
}
return res;
}

char *itoa(int value, char *str, int base = 10) {
    if (base < 2 || base > 36) {
        *str = '\0';
        return str;
    }
    char *ptr = str, *ptr1 = str, tmp_c;
    int tmp_v;
    do {
        tmp_v = value;
        value /= base;
        *ptr++ =
            "zyxwvutsrqponmlkjihgfedcba9876543210123456789"
            "abcdefghijklmnopqrstuvwxyz"[35 + (tmp_v - value * base)];
    } while (value);
    if (tmp_v < 0) {
        *ptr++ = '-';
    }
    for (*ptr-- = '\0'; ptr1 < ptr; *ptr1++ = tmp_c) {
        tmp_c = *ptr;
        *ptr-- = *ptr1;
    }
    return str;
}

bool is_integer(const string &s) {
    // return std::regex_match(s, std::regex(R"([+-]?\d+)"));
    int i = (s.empty() || (s[0] != '-' && s[0] != '+')) ? 0 : 1;
    if (i == static_cast<int>(s.size())) {
        return false;
    }
    for (; i < static_cast<int>(s.size()); i++) {
        if (!isdigit(static_cast<unsigned char>(s[i]))) {
            return false;
        }
    }
    return true;
}

bool is_number(const string &s) {
    // return std::regex_match(s, std::regex(R"([+-]?(\d+(\.\d*)?|\.\d+)([eE][+-]?\d+)?)"));
    int i = (s.empty() || (s[0] != '-' && s[0] != '+')) ? 0 : 1;
    bool seen_digit = false, seen_dot = false;
    for (; i < static_cast<int>(s.size()); i++) {
        if (isdigit(static_cast<unsigned char>(s[i]))) {
            seen_digit = true;
        } else if (s[i] == '.' && !seen_dot) {
            seen_dot = true;
        } else {
            break;
        }
    }
    if (!seen_digit) {
        return false;
    }
    if (i < static_cast<int>(s.size()) && (s[i] == 'e' || s[i] == 'E')) {
        i++;
        if (i < static_cast<int>(s.size()) && (s[i] == '-' || s[i] == '+')) {
            i++;
        }
        bool seen_exp_digit = false;
    }

```

```

for (; i < static_cast<int>(s.size()); i++) {
    if (!isdigit(static_cast<unsigned char>(s[i]))) {
        return false;
    }
    seen_exp_digit = true;
}
return seen_exp_digit;
}
return i == static_cast<int>(s.size());
}

```

Case Conversion:

- `to_upper(s)` returns `s` with all alphabetical characters converted to uppercase.
- `to_lower(s)` returns `s` with all alphabetical characters converted to lowercase.
- `to_title(s)` returns the title case representation of string `s`, where the first letter of every word (consecutive alphabetical characters) is capitalized.

```

string to_upper(const string &s) {
    string res;
    res.reserve(s.size());
    for (unsigned char c : s) {
        res.push_back(static_cast<char>(toupper(c)));
    }
    return res;
}

```

```

string to_lower(const string &s) {
    string res;
    res.reserve(s.size());
    for (unsigned char c : s) {
        res.push_back(static_cast<char>(tolower(c)));
    }
    return res;
}

```

```

string to_title(const string &s) {
    string res;
    unsigned char prev = '\0';
    for (unsigned char c : s) {
        res.push_back(isalpha(prev) ? static_cast<char>(tolower(c)) : static_cast<char>(toupper(c)));
        prev = res.back();
    }
    return res;
}

```

Stripping:

- `lstrip(s, delim = WHITESPACE)` strips the left side of `s` in-place (that is, the input is modified) using the given delimiters and returns a reference to the stripped string.
- `rstrip(s, delim = WHITESPACE)` strips the right side of `s` in-place using the given delimiters and returns a reference to the stripped string.
- `strip(s, delim = WHITESPACE)` strips both sides of `s` in-place and returns a reference to the stripped string.
- `ltrimmed(s, delim = WHITESPACE)`, `rtrimmed(s, delim = WHITESPACE)`, and `trimmed(s, delim = WHITESPACE)` do not modify `s` and return stripped copies.

```

const string WHITESPACE = " \n\t\v\f\r";

string &lstrip(string &s, const string &delim = WHITESPACE) {
    std::size_t pos = s.find_first_not_of(delim);
    if (pos != string::npos) {
        s.erase(0, pos);
    } else {

```

```

    s.clear();
}
return s;
}

string &rstrip(string &s, const string &delim = WHITESPACE) {
    std::size_t pos = s.find_last_not_of(delim);
    if (pos != string::npos) {
        s.erase(pos + 1);
    } else {
        s.clear();
    }
    return s;
}

string &lstrip(string &s, const string &delim = WHITESPACE) {
    return lstrip(s, delim);
}

string &ltrimmed(string s, const string &delim = WHITESPACE) {
    return lstrip(s, delim);
}

string &rtrimmed(string s, const string &delim = WHITESPACE) {
    return rstrip(s, delim);
}

string &trimmed(string s, const string &delim = WHITESPACE) {
    return strip(s, delim);
}

```

Find and Replace:

- `starts_with(s, prefix)` returns whether `s` begins with `prefix`.
- `ends_with(s, suffix)` returns whether `s` ends with `suffix`.
- `contains(s, needle)` returns whether `needle` occurs in `s`.
- `find_all(haystack, needle)` returns a vector of all positions where the nonempty string `needle` appears in the string `haystack`. Matches may overlap; e.g. `find_all("aaa", "aa")` returns `{0, 1}`. An empty `needle` produces no matches.
- `count(haystack, needle)` returns the number of non-overlapping occurrences of `needle` in `haystack`, like Python's `str.count`. This differs from `find_all().size()`, which counts overlapping matches; e.g. `count("aaa", "aa")` returns 1, but `find_all("aaa", "aa")` finds 2.
- `replace(s, old, replacement)` returns a copy of `s` with all occurrences of the string `old` replaced with the given `replacement`.
- `regex_find_all(s, pattern)` returns every substring of `s` matched by `pattern`.
- `regex_groups(s, pattern)` returns the capture groups of the first match of `pattern` in `s`, excluding the full match. If there is no match, returns an empty vector.
- `regex_replace_all(s, pattern, replacement)` returns a copy of `s` with every regex match replaced by `replacement`.

```

bool starts_with(const string &s, const string &prefix) {
    return s.size() >= prefix.size() && s.compare(0, prefix.size(), prefix) == 0;
}

bool ends_with(const string &s, const string &suffix) {
    return s.size() >= suffix.size() &&
        s.compare(s.size() - suffix.size(), suffix.size(), suffix) == 0;
}

bool contains(const string &s, const string &needle) {
    return s.find(needle) != string::npos;
}

```

```

std::vector<int> find_all(const string &haystack, const string &needle) {
    std::vector<int> res;
    if (needle.empty()) {
        return res;
    }
    std::size_t pos = haystack.find(needle, 0);
    while (pos != string::npos) {
        res.push_back(static_cast<int>(pos));
        pos = haystack.find(needle, pos + 1);
    }
    return res;
}

int count(const string &haystack, const string &needle) {
    if (needle.empty()) {
        return 0;
    }
    int res = 0;
    std::size_t pos = haystack.find(needle, 0);
    while (pos != string::npos) {
        res++;
        pos = haystack.find(needle, pos + needle.size());
    }
    return res;
}

string replace(const string &s, const string &old, const string &replacement) {
    if (old.empty()) {
        return s;
    }
    string res(s);
    std::size_t pos = 0;
    while ((pos = res.find(old, pos)) != string::npos) {
        res.replace(pos, old.length(), replacement);
        pos += replacement.length();
    }
    return res;
}

std::vector<string> regex_find_all(const string &s, const string &pattern) {
    std::regex re(pattern);
    return {std::sregex_token_iterator(s.begin(), s.end(), re, 0), std::sregex_token_iterator()};
}

std::vector<string> regex_groups(const string &s, const string &pattern) {
    std::smatch match;
    if (!std::regex_search(s, match, std::regex(pattern))) {
        return {};
    }
    std::vector<string> res;
    for (int i = 1; i < static_cast<int>(match.size()); i++) {
        res.push_back(match[i].str());
    }
    return res;
}

string regex_replace_all(const string &s, const string &pattern, const string &replacement) {
    return std::regex_replace(s, std::regex(pattern), replacement);
}

```

Joining and Splitting:

- `join(v, delim = " ")` returns the strings in vector `v` concatenated, separated by the given delimiter.
- `repeat(s, n)` returns `s` concatenated with itself `n` times, like Python's `s * n`, or an empty string if `n` is nonpositive. For a single repeated character, prefer `std::string(n, ch)`.

- `ljust(s, width, ch = ' ')` returns `s` left-justified (padded on the right with `ch`) to at least `width` characters, like Python's `str.ljust`.
- `rjust(s, width, ch = ' ')` returns `s` right-justified (padded on the left with `ch`) to at least `width` characters, like Python's `str.rjust`.
- `center(s, width, ch = ' ')` returns `s` centered in a field of at least `width` characters by padding both sides with `ch`, like Python's `str.center`. If the padding is uneven, the extra character goes on the right.
- `split(s, char delim)` returns a vector of tokens of `s`, split on a single character delimiter. Empty tokens are skipped, so `split("a::b", ':')` returns `{"a", "b"}`, not `{"a", "", "b"}`.
- `split(s, string delim = WHITESPACE)` returns a vector of tokens of `s`, split on a set of many possible single character delimiters. All characters of `delim` will be removed from `s`, and the remaining token(s) of `s` will be added sequentially to a vector and returned. As with the first version, empty tokens are skipped. For example, `split("a::b", ":", " ")` returns `{"a", "b"}`, not `{"a", "", "b"}`.
- `explode(s, delim)` returns a vector of tokens of `s`, split on the entire delimiter string `delim`. Unlike the `split()` functions above, `delim` is treated as a contiguous boundary string, not merely a set of possible boundary characters. This will not skip empty tokens. For example, `explode("a:::b", ":", " ")` yields `{"a", "", "b"}`, not `{"a", "b"}`. An empty delimiter returns `{s}`.
- `explode(s, char delim)` is the single-character delimiter version that also preserves empty tokens.

```
string join(const std::vector<string> &v, const string &delim = " ") {
    if (v.empty()) {
        return "";
    }
    string res;
    std::size_t total = delim.size() * (v.size() - 1);
    for (const string &s : v) {
        total += s.size();
    }
    res.reserve(total);
    for (int i = 0; i < static_cast<int>(v.size()); i++) {
        res += (i > 0 ? delim : "") + v[i];
    }
    return res;
}

string repeat(const string &s, int n) {
    if (n <= 0) {
        return "";
    }
    string res;
    res.reserve(s.size() * n);
    for (int i = 0; i < n; i++) {
        res += s;
    }
    return res;
}

string ljust(const string &s, int width, char ch = ' ') {
    int pad = width - static_cast<int>(s.size());
    return pad > 0 ? s + string(pad, ch) : s;
}

string rjust(const string &s, int width, char ch = ' ') {
    int pad = width - static_cast<int>(s.size());
    return pad > 0 ? string(pad, ch) + s : s;
}

string center(const string &s, int width, char ch = ' ') {
    int pad = width - static_cast<int>(s.size());
    if (pad <= 0) {
        return s;
    }
    int left = pad / 2;
    return string(left, ch) + s + string(pad - left, ch);
}
```

```

}

std::vector<string> split(const string &s, char delim) {
    std::vector<string> res;
    std::stringstream ss(s);
    string curr;
    while (std::getline(ss, curr, delim)) {
        if (!curr.empty()) {
            res.push_back(curr);
        }
    }
    return res;
}

std::vector<string> split(const string &s, const string &delim = WHITESPACE) {
    std::vector<string> res;
    string curr;
    for (char c : s) {
        if (delim.find(c) == string::npos) {
            curr += c;
        } else if (!curr.empty()) {
            res.push_back(curr);
            curr = "";
        }
    }
    if (!curr.empty()) {
        res.push_back(curr);
    }
    return res;
}

std::vector<string> explode(const string &s, const string &delim) {
    if (delim.empty()) {
        return {s};
    }
    std::vector<string> res;
    std::size_t last = 0, next = 0;
    while ((next = s.find(delim, last)) != string::npos) {
        res.push_back(s.substr(last, next - last));
        last = next + delim.size();
    }
    res.push_back(s.substr(last));
    return res;
}

std::vector<string> explode(const string &s, char delim) {
    std::vector<string> res;
    std::size_t last = 0, next = 0;
    while ((next = s.find(delim, last)) != string::npos) {
        res.push_back(s.substr(last, next - last));
        last = next + 1;
    }
    res.push_back(s.substr(last));
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    assert(to_str(123) + "4" == "1234");
    assert(to_int("1234") == 1234);
    bool threw = false;
    try {

```

```

    to_int("123x");
} catch (const runtime_error &) {
    threw = true;
}
assert(threw);
vector<char> buffer(50);
assert(string(itoa(1750, buffer.data(), 10)) == "1750");
assert(string(itoa(1750, buffer.data(), 16)) == "6d6");
assert(string(itoa(1750, buffer.data(), 2)) == "11011010110");
assert(is_integer("-123") && !is_integer("12.3"));
assert(is_number("-.5e+2") && is_number("123.") && !is_number("1e"));

assert(to_upper("Hello world") == "HELLO WORLD");
assert(to_lower("Hello World") == "hello world");
assert(to_title("hello world") == "Hello World");

string s("  abc \n");
string t = s;
assert(lstrip(s) == "abc \n");
assert(rstrip(s) == strip(t));
assert(trimmed(" \t abc \n") == "abc");
assert(trimmed("xxx", "x").empty());

vector<int> pos;
pos.push_back(0);
pos.push_back(7);
assert(starts_with("abracadabra", "abra"));
assert(ends_with("abracadabra", "dabra"));
assert(contains("abracadabra", "cad"));
assert(find_all("abracadabra", "ab") == pos);
assert(count("abracadabra", "ab") == 2);
assert(count("aaa", "aa") == 1 && find_all("aaa", "aa").size() == 2);
assert(find_all("abc", "").empty());
assert(replace("abcdabba", "ab", "00") == "00cd00ba");
assert(join(regex_find_all("a12b345", "[0-9]+"), "|") == "12|345");
assert(join(regex_groups("abc-123", "([a-z]+)-([0-9]+)"), "|") == "abc|123");
assert(regex_replace_all("a12b345", "[0-9]+", "#") == "a#b#");

assert(repeat("ab", 3) == "ababab");
assert(repeat("ab", -1).empty());
assert(rjust("42", 5, '0') == "00042");
assert(ljust("hi", 4, '.') == "hi..");
assert(center("ab", 5, '*') == "*ab*");
assert(join(split("a\nb\ncde\nf", '\n'), "|") == "a|b|cde|f"); // split v1
assert(join(split("a::b", ':'), "|") == "a|b"); // split v1 skips empty tokens
assert(join(split("a::b,cde:,f", ":", ","), "|") == "a|b|cde|f"); // split v2
assert(join(explode("a..b.cde...f", "."), "|") == "a|b.cde|f");
assert((explode("abc", "") == vector<string>{"abc"}));
assert(join(explode("a::b:", ':'), "|") == "a|b|");
return 0;
}

```

3.2 Hashing

3.2.1 Hash Functions

Provides a small library of non-cryptographic hash functions and hash functors for common contest keys. These functions are useful for fingerprinting values, combining keys, seeding randomized algorithms, and defining `std::unordered_map` keys for pairs, tuples, and vectors. They are not suitable for passwords, signatures, or adversarial security.

The standalone functions below are deterministic. The `IntHasher` functor adds a per-run random seed before mixing, which is useful for making integer-keyed hash tables harder to hack in open-test contests.

The composite hashers separate the mixing primitive (`mix64`/`hash_combine64`) from each type's traversal of its parts, so swapping in a different mixing function only requires editing one place.

FNV-1a (Fowler-Noll-Vo) is a tiny hash for variable-length byte sequences. Starting from a fixed offset basis, it XORs each byte into the accumulator and then multiplies by a fixed prime; the `1a` variant XORs before multiplying, which mixes slightly better than the original FNV-1. Reach for it when the key is a string or byte buffer, whereas `mix32` and `mix64` are mixers for a single fixed-width integer. Its appeal is simplicity: a few lines, no lookup tables, and no seed, with distribution good enough for the short keys common in contests. For very long inputs a block hash such as MurmurHash or xxHash is faster and mixes more thoroughly, and because FNV-1a is unseeded it offers no protection against adversarial inputs.

- `mix32(x)` mixes a 32-bit unsigned integer `x`, using MurmurHash3's `fmix32` mixer.
- `mix64(x)` mixes a 64-bit unsigned integer `x`, using the `SplitMix64` mixer.
- `hash32(x)` and `hash64(x)` hash any integer type that can be converted to `uint32_t` or `uint64_t`, respectively.
- `hash_combine32(h, x)` and `hash_combine64(h, x)` fold another already-hashed value into an existing hash accumulator.
- `hash_float(x)` and `hash_double(x)` hash floating-point bit patterns, normalizing `-0.0` to `0.0`. Different NaN representations may still hash differently.
- `hash_string_fnv1a32(s)` and `hash_string_fnv1a64(s)` compute FNV-1a hashes of string `s`.
- `hash_range32(lo, hi)` and `hash_range64(lo, hi)` hash a sequence of integer-like values in the half-open iterator range `[lo, hi)`.
- `PairIntHasher` is a self-contained hasher for `std::pair<int, int>` keys, written without any dependency on the templates below so that it can be copy-pasted on its own.
- `IntHasher<Int>`, `PairHasher<A, B>`, `TupleHasher<Ts...>`, and `VectorHasher<T>` are hash functors passed as the third template argument of `std::unordered_map` or `std::unordered_set`.
- `GenericHasher<T>` recursively handles integer, pair, tuple, and vector keys, and falls back to `std::hash<T>` for other hashable types.

Time Complexity:

- $O(1)$ per call to the integer, combiner, and floating-point hash functions.
- $O(n)$ per call to string and range hashing functions, where n is the number of elements processed.
- $O(n)$ per call to `VectorHasher<T>::operator()`, where n is the vector size.

Space Complexity: $O(1)$ auxiliary.

```
#include <chrono>
#include <cstdint>
#include <cstdint>
#include <cstring>
#include <functional>
#include <string>
#include <tuple>
#include <type_traits>
#include <unordered_map>
#include <utility>
#include <vector>

// MurmurHash3's fmix32 mixer.
uint32_t mix32(uint32_t x) {
    x ^= x >> 16;
    x *= 0x85ebca6bU;
    x ^= x >> 13;
    x *= 0xc2b2ae35U;
    return x ^ (x >> 16);
}

// SplitMix64 mixer.
uint64_t mix64(uint64_t x) {
```



```

    x += 0x9e3779b97f4a7c15ULL;
    x = (x ^ (x >> 30)) * 0xbf58476d1ce4e5b9ULL;
    x = (x ^ (x >> 27)) * 0x94d049bb133111ebULL;
    return x ^ (x >> 31);
}

template<typename Int>
uint32_t hash32(Int x) {
    return mix32(static_cast<uint32_t>(x));
}

template<typename Int>
uint64_t hash64(Int x) {
    return mix64(static_cast<uint64_t>(x));
}

uint32_t hash_combine32(uint32_t h, uint32_t x) {
    return mix32(h ^ (x + 0x9e3779b9U + (h << 6) + (h >> 2)));
}

uint64_t hash_combine64(uint64_t h, uint64_t x) {
    return mix64(h ^ (x + 0x9e3779b97f4a7c15ULL + (h << 6) + (h >> 2)));
}

uint32_t hash_float(float x) {
    if (x == 0.0f) {
        x = 0.0f; // Normalize -0.0.
    }
    uint32_t bits = 0;
    std::memcpy(&bits, &x, sizeof(bits));
    return mix32(bits);
}

uint64_t hash_double(double x) {
    if (x == 0.0) {
        x = 0.0; // Normalize -0.0.
    }
    uint64_t bits = 0;
    std::memcpy(&bits, &x, sizeof(bits));
    return mix64(bits);
}

uint32_t hash_string_fnv1a32(const std::string &s) {
    uint32_t h = 2166136261U;
    for (unsigned char c : s) {
        h ^= c;
        h *= 16777619U;
    }
    return h;
}

uint64_t hash_string_fnv1a64(const std::string &s) {
    uint64_t h = 14695981039346656037ULL;
    for (unsigned char c : s) {
        h ^= c;
        h *= 1099511628211ULL;
    }
    return h;
}

template<typename It>
uint32_t hash_range32(It lo, It hi) {
    uint32_t h = 0;
    for (It it = lo; it != hi; ++it) {
        h = hash_combine32(h, hash32(*it));
    }
    return h;
}

```

```

template<typename It>
uint64_t hash_range64(It lo, It hi) {
    uint64_t h = 0;
    for (It it = lo; it != hi; ++it) {
        h = hash_combine64(h, hash64(*it));
    }
    return h;
}

template<typename Int>
struct IntHasher {
    std::size_t operator()(Int x) const {
        static const uint64_t RAND_SEED = std::chrono::steady_clock::now().time_since_epoch().count();
        return static_cast<std::size_t>(mix64(static_cast<uint64_t>(x) + RAND_SEED));
    }
};

template<typename T>
struct GenericHasher;

// Use the seeded hasher for integers, and std::hash by default for everything else.
template<typename T>
struct ScalarHasher {
    std::size_t operator()(const T &x) const {
        if constexpr (std::is_integral_v<T>) {
            return IntHasher<T>{}(x);
        }
        return std::hash<T>()(x);
    }
};

// A self-contained hasher for std::pair<int, int> keys: copy just this struct when you do not need
// the rest of this file. The two 32-bit halves pack losslessly into one 64-bit word, which the
// SplitMix64 mixer (the same one used by mix64 above) then scrambles. Add a random seed to key
// before mixing, as IntHasher does, if you need anti-hacking protection.
struct PairIntHasher {
    std::size_t operator()(const std::pair<int, int> &p) const {
        uint64_t key = (static_cast<uint64_t>(static_cast<uint32_t>(p.first)) << 32) |
            static_cast<uint32_t>(p.second);
        key += 0x9e3779b97f4a7c15ULL;
        key = (key ^ (key >> 30)) * 0xbf58476d1ce4e5b9ULL;
        key = (key ^ (key >> 27)) * 0x94d049bb133111ebULL;
        return static_cast<std::size_t>(key ^ (key >> 31));
    }
};

template<typename A, typename B>
struct PairHasher {
    std::size_t operator()(const std::pair<A, B> &p) const {
        uint64_t h = 0;
        h = hash_combine64(h, static_cast<uint64_t>(GenericHasher<A>()(p.first)));
        h = hash_combine64(h, static_cast<uint64_t>(GenericHasher<B>()(p.second)));
        return static_cast<std::size_t>(h);
    }
};

template<typename T>
struct VectorHasher {
    std::size_t operator()(const std::vector<T> &v) const {
        uint64_t h = 0;
        for (const auto &x : v) {
            h = hash_combine64(h, static_cast<uint64_t>(GenericHasher<T>()(x)));
        }
        h = hash_combine64(h, v.size()); // Mix in the length so different sizes (e.g. empty) separate.
        return static_cast<std::size_t>(h);
    }
};

```

```

template<typename... Ts>
struct TupleHasher {
    std::size_t operator()(const std::tuple<Ts...> &t) const {
        uint64_t h = 0;
        std::apply(
            [&](const Ts &...xs) {
                ((h = hash_combine64(h, static_cast<uint64_t>(GenericHasher<Ts>()(xs)))), ...);
            },
            t
        );
        return static_cast<std::size_t>(h);
    }
};

template<typename T>
struct GenericHasher : ScalarHasher<T> {};

template<typename A, typename B>
struct GenericHasher<std::pair<A, B>> : PairHasher<A, B> {};

template<typename... Ts>
struct GenericHasher<std::tuple<Ts...>> : TupleHasher<Ts...> {};

template<typename T>
struct GenericHasher<std::vector<T>> : VectorHasher<T> {};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    assert(mix32(123U) == mix32(123U));
    assert(mix32(123U) != mix32(124U));
    assert(hash32(-1) == hash32(-1));
    assert(hash64(123) == mix64(123ULL));
    assert(hash_float(-0.0f) == hash_float(0.0f));
    assert(hash_double(-0.0) == hash_double(0.0));
    assert(hash_string_fnv1a32("abc") == hash_string_fnv1a32("abc"));
    assert(hash_string_fnv1a64("abc") != hash_string_fnv1a64("acb"));

    vector<int> v;
    v.push_back(1);
    v.push_back(2);
    v.push_back(3);
    assert(hash_range32(v.begin(), v.end()) == hash_range32(v.begin(), v.end()));
    assert(hash_range64(v.begin(), v.end()) == hash_range64(v.begin(), v.end()));

    unordered_map<int64_t, int, IntHasher<int64_t>> count;
    count[1000000000000LL] = 7;
    assert(count[1000000000000LL] == 7);

    using Point = pair<int, int>;
    unordered_map<Point, int, PairHasher<int, int>> dist;
    dist[Point(3, 4)] = 5;
    assert(dist[Point(3, 4)] == 5);

    // The standalone PairIntHasher is the quick choice for pair<int, int> keys.
    unordered_map<Point, int, PairIntHasher> dist_fast;
    dist_fast[Point(3, 4)] = 5;
    assert(dist_fast[Point(3, 4)] == 5);

    unordered_map<vector<int>, int, VectorHasher<int>> seen;
    seen[v] = 1;
    assert(seen[v] == 1);
}

```

```

using State = pair<vector<int>, double>;
unordered_map<State, int, GenericHasher<State>> dp;
dp[State(v, 4.5)] = 9;
assert(dp[State(v, 4.5)] == 9);

// Tuple keys via the explicit functor.
using Triple = tuple<int, char, double>;
unordered_map<Triple, int, GenericHasher<Triple>> grid;
grid[Tuple(1, 'a', 3.14)] = 6;
assert(grid[Tuple(1, 'a', 3.14)] == 6);
return 0;
}

```

3.2.2 Rolling Hash

Computes polynomial rolling hashes for sequences of an arbitrary hashable type $\mathbb{T}$, supporting $O(1)$ contiguous subsequence hash queries after preprocessing. This is useful for probabilistic string matching, substring equality checks, repeated subarray detection, and binary-searching over candidate lengths.

Given a sequence $a_0, a_1, \dots, a_{n-1}$ and a value hash $h(a_i)$ for each element, the hash is the polynomial $H(a) = \sum_{i=0}^{n-1} h(a_i) B^{n-1-i} \pmod{M}$, where B is `HASH_BASE` and M is `HASH_MOD`. Let $p(i)$ be the hash of the first i values. It is built left-to-right by the prefix hash recurrence $p(0) = 0$ and $p(i+1) = p(i)B + h(a_i) \pmod{M}$. The hash of a half-open subsequence $[l, r)$ is then computed as $p(r) - p(l)B^{r-l} \pmod{M}$.

The implementation works modulo the Mersenne prime $2^{61} - 1$, using `__uint128_t` for the multiplication where available and falling back to a double-and-add modular multiply otherwise. The base `HASH_BASE` should be changed or chosen randomly for open-hacking environments. With a fixed base, hashing remains fast and practical, but it is still probabilistic and should not be used as proof of equality when exact verification is required.

By default, each sequence value is cast to `uint64_t` and mixed. For non-integer element types, instantiate `RollingHash<T, Hash>` and pass a custom value hasher as the final constructor or `hash()` argument. It must map each element to a stable nonzero value in $[1, \text{HASH_MOD})$.

- `RollingHash<T>()` constructs an empty hash sequence.
- `RollingHash<T>(lo, hi)` constructs prefix hashes from the values in the half-open iterator range $[lo, hi)$.
- `RollingHash<T>(v)` constructs prefix hashes for vector `v`.
- `get(lo, hi)` returns the hash of the half-open subsequence $[lo, hi)$.
- `hash(lo, hi)` returns the hash of the half-open iterator range $[lo, hi)$.
- `hash(v)` returns the hash of vector `v`.
- `concat(left, right, right_len)` returns the hash of the concatenation of a sequence with hash `left` and a sequence with hash `right` and length `right_len`.

Time Complexity:

- $O(n)$ per call to the constructor or `hash()`, where n is the sequence length.
- $O(1)$ per call to `get()` and `concat()`.

Space Complexity:

- $O(n)$ for storage of prefix hashes and powers, where n is the maximum sequence length processed so far.
- $O(1)$ auxiliary per query.

```

#include <cassert>
#include <cstdint>
#include <string>
#include <utility>
#include <vector>

const uint64_t HASH_MOD = (1ULL << 61) - 1;
const uint64_t HASH_BASE = 911382323ULL; // Change or randomize.

```

```

uint64_t hash_add(uint64_t a, uint64_t b) {
    uint64_t c = a + b;
    return c >= HASH_MOD ? c - HASH_MOD : c;
}

uint64_t hash_sub(uint64_t a, uint64_t b) {
    return a >= b ? a - b : a + HASH_MOD - b;
}

uint64_t hash_mul(uint64_t a, uint64_t b) {
#if defined(__SIZEOF_INT128__)
    __uint128_t p = static_cast<__uint128_t>(a) * b;
    return hash_add(static_cast<uint64_t>(p) & HASH_MOD, static_cast<uint64_t>(p >> 61));
#else
    // Portable fallback when no 128-bit type exists: double-and-add modular multiply.
    uint64_t res = 0;
    for (a %= HASH_MOD; b > 0; b >>= 1) {
        if (b & 1) {
            res = hash_add(res, a);
        }
        a = hash_add(a, a);
    }
    return res;
#endif
}

// SplitMix64 mixer.
uint64_t mix64(uint64_t x) {
    x += 0x9e3779b97f4a7c15ULL;
    x = (x ^ (x >> 30)) * 0xbf58476d1ce4e5b9ULL;
    x = (x ^ (x >> 27)) * 0x94d049bb133111ebULL;
    return x ^ (x >> 31);
}

template<typename T>
struct ValueHasher {
    // +1 ensures the result is in [1, HASH_MOD) and never zero; a zero element in a polynomial
    // hash is invisible and enables trivial collisions.
    uint64_t operator()(const T &x) const {
        return mix64(static_cast<uint64_t>(x)) % (HASH_MOD - 1) + 1;
    }
};

template<typename T, typename Hash = ValueHasher<T>>
class RollingHash {
    static std::vector<uint64_t> pow_base;
    std::vector<uint64_t> pref;
    Hash hasher;

    static void ensure_powers(int n) {
        while (static_cast<int>(pow_base.size()) <= n) {
            pow_base.push_back(hash_mul(pow_base.back(), HASH_BASE));
        }
    }

    template<typename It>
    void build(It lo, It hi) {
        pref.clear();
        pref.push_back(0);
        for (; lo != hi; ++lo) {
            pref.push_back(hash_add(hash_mul(pref.back(), HASH_BASE), hasher(*lo)));
        }
        ensure_powers(static_cast<int>(pref.size()) - 1);
    }

public:
    explicit RollingHash(const Hash &hasher = Hash()) : hasher(hasher) {

```

```

    ensure_powers(0);
    pref.push_back(0);
}

template<typename It>
RollingHash(It lo, It hi, const Hash &hasher = Hash()) : hasher(hasher) {
    build(lo, hi);
}

explicit RollingHash(const std::vector<T> &v, const Hash &hasher = Hash()) : hasher(hasher) {
    build(v.begin(), v.end());
}

int size() const { return static_cast<int>(pref.size()) - 1; }

uint64_t get(int lo, int hi) const {
    assert(0 <= lo && lo <= hi && hi <= size());
    return hash_sub(pref[hi], hash_mul(pref[lo], pow_base[hi - lo]));
}

template<typename It>
static uint64_t hash(It lo, It hi, const Hash &hasher = Hash()) {
    RollingHash<T, Hash> h(lo, hi, hasher);
    return h.get(0, h.size());
}

static uint64_t hash(const std::vector<T> &v, const Hash &hasher = Hash()) {
    RollingHash<T, Hash> h(v, hasher);
    return h.get(0, h.size());
}

static uint64_t concat(uint64_t left, uint64_t right, int right_len) {
    assert(right_len >= 0);
    ensure_powers(right_len);
    return hash_add(hash_mul(left, pow_base[right_len]), right);
}
};

template<typename T, typename Hash>
std::vector<uint64_t> RollingHash<T, Hash>::pow_base(1, 1);

```

Example Usage

```

using namespace std;

using PointI = pair<int, int>;

struct PointHasher {
    uint64_t operator()(const PointI &p) const {
        // Pack the two 32-bit coordinates losslessly into one 64-bit key before mixing, so distinct
        // points never collide pre-mix (a small-multiplier combine like a * C + b would).
        uint64_t key = (static_cast<uint64_t>(static_cast<uint32_t>(p.first)) << 32) |
            static_cast<uint32_t>(p.second);
        return mix64(key) % (HASH_MOD - 1) + 1;
    }
};

int main() {
    string s = "abracadabra";
    RollingHash<char> hs(s.begin(), s.end());
    assert(hs.get(0, 4) == hs.get(7, 11)); // "abra" == "abra"
    assert(hs.get(0, 4) != hs.get(3, 7));  // "abra" != "acad"

    string a = "abc", b = "def";
    uint64_t ab = RollingHash<char>::concat(
        RollingHash<char>::hash(a.begin(), a.end()), RollingHash<char>::hash(b.begin(), b.end()),

```

```

    static_cast<int>(b.size())
);
string c = a + b;
assert(ab == RollingHash<char>::hash(c.begin(), c.end()));

vector<int> v;
v.push_back(1);
v.push_back(2);
v.push_back(3);
v.push_back(1);
v.push_back(2);
RollingHash<int> hv(v);
assert(hv.get(0, 2) == hv.get(3, 5)); // Both ranges are [1, 2].
assert(hv.get(0, 3) == RollingHash<int>::hash(v.begin(), v.begin() + 3));

vector<PointI> poly{{1, 2}, {3, 4}};
RollingHash<PointI, PointHasher> hp(poly);
assert((hp.get(0, 2) == RollingHash<PointI, PointHasher>::hash(poly)));
return 0;
}

```

3.2.3 Rolling Hash Deque

Maintains a polynomial rolling hash for a sequence of an arbitrary hashable type $\mathbb{T}$, supporting insertion and deletion at either end in $O(1)$ time per operation. Each element is first reduced to a nonnegative integer “digit” by a user-supplied hasher, so characters, integers, whole words, and tuples can all be hashed. This is the deque generalization of the classic Rabin-Karp sliding window: pushing at the back and popping from the front is exactly a moving window, and this structure additionally supports pushing at the front and popping from the back.

The sequence of digits is hashed with the front element as the most significant digit, so a sequence of length n has hash $\sum_{i=0}^{n-1} a_i B^{n-1-i}$ (the front element a_0 carries the highest power). This makes the four end operations pure $O(1)$ updates of a running accumulator:

- `push_back(c)`: $H \rightarrow H \cdot B + c$ (Horner’s step; the existing weights all shift up).
- `pop_front()`: $H \rightarrow H - a_0 \cdot B^{n-1}$ (drop the most significant term).
- `push_front(c)`: $H \rightarrow c \cdot B^n + H$ (the new front becomes the most significant term).
- `pop_back()`: $H \rightarrow (H - a_{n-1}) \cdot B^{-1}$ (drop the constant term, then renormalize).

Only `pop_back()` needs a modular inverse of the base, which is precomputed once. The weight of the front element, B^{n-1} , is maintained incrementally (multiplied by B on a push, by B^{-1} on a pop), so no power table or length-dependent exponentiation is needed. Each element’s digit is stored when pushed, so popping never re-invokes the hasher. Because the hash depends only on the current contents, two deques with equal sequences hash equally regardless of how they were built, and a sliding window can be compared against a fixed target in $O(1)$ per step.

`MOD1` and `MOD2` are two distinct primes, each kept below 2^{31} so that the product of any two residues fits in a 64-bit integer; hashing modulo both and comparing the resulting pair makes a collision astronomically unlikely. `BASE1` and `BASE2` are the corresponding polynomial bases, which must be nonzero modulo their corresponding primes and may be randomized per run to resist adversarial inputs. For a small raw alphabet, choosing bases larger than every digit also prevents carry-style collisions before modular reduction, but arbitrary hasher outputs need not satisfy this bound.

The hasher’s output must be nonnegative. As with any polynomial hash, comparing sequences of different lengths is only safe when digits are positive (for example, map an alphabet to $[1, B)$) so that a leading zero cannot make a shorter sequence collide with a longer one.

The element type defaults to `int`, and integer-like values use `IdentityDigit<T>`. To customize the mapping, instantiate `HashDeque<T, Hash>` and pass the hasher to the constructor.

- `HashDeque<T>()` constructs an empty deque. `HashDeque<char>` and `HashDeque<int>` need no custom hasher.

- `size()` returns the number of elements, and `empty()` whether there are none.
- `push_back(x)` and `push_front(x)` append `x` at the back or front.
- `pop_back()` and `pop_front()` remove the back or front element; the deque must be nonempty.
- `hash()` returns the pair of hashes of the current sequence.
- Operators `==` and `!=` compare whether two deques are equal using the lengths and hash pairs. Equal sequences always compare equal; equality indicates equal sequences with high probability.

Time Complexity: $O(1)$ per call to every operation, plus one hasher evaluation per `push_back()` / `push_front()`.

Space Complexity:

- $O(n)$ for storage, where n is the current number of elements.
- $O(1)$ auxiliary per operation.

```
#include <cassert>
#include <cstdint>
#include <deque>
#include <utility>

template<typename T>
struct IdentityDigit {
    uint64_t operator()(const T &x) const { return static_cast<uint64_t>(x); }
};

template<typename T = int, typename Hash = IdentityDigit<T>>
class HashDeque {
    static const uint64_t MOD1 = 1000000007, MOD2 = 1000000009; // Distinct primes below 2^31.
    static const uint64_t BASE1 = 131, BASE2 = 137;           // Nonzero modulo MOD1 and MOD2.

    Hash hasher; // Maps an element to its nonnegative integer digit.
    std::deque<uint64_t> digits; // Per-element digit values, retained so pops never re-hash.
    uint64_t h1 = 0, h2 = 0; // Hash of the current sequence modulo MOD1, MOD2.
    uint64_t top1, top2; // B^(len-1): weight of the front element (B^-1 while empty).
    uint64_t inv1, inv2; // Modular inverses of BASE1, BASE2, for pop_back().

    // Products and sums of residues below 2^30 stay under 2^60, so 64-bit arithmetic never overflows.
    static uint64_t powmod(uint64_t b, uint64_t e, uint64_t m) {
        uint64_t res = 1;
        for (b %= m; e > 0; e >>= 1) {
            if (e & 1) {
                res = res * b % m;
            }
            b = b * b % m;
        }
        return res;
    }

public:
    explicit HashDeque(Hash hasher = Hash()) : hasher(hasher) {
        inv1 = powmod(BASE1, MOD1 - 2, MOD1);
        inv2 = powmod(BASE2, MOD2 - 2, MOD2);
        top1 = inv1; // B^(0-1) = B^-1, so the first push lifts the front weight to B^0 = 1.
        top2 = inv2;
    }

    int size() const { return static_cast<int>(digits.size()); }
    bool empty() const { return digits.empty(); }
    std::pair<uint64_t, uint64_t> hash() const { return {h1, h2}; }

    void push_back(const T &x) {
        uint64_t d = hasher(x), c1 = d % MOD1, c2 = d % MOD2;
        h1 = (h1 * BASE1 + c1) % MOD1;
        h2 = (h2 * BASE2 + c2) % MOD2;
        top1 = top1 * BASE1 % MOD1;
        top2 = top2 * BASE2 % MOD2;
        digits.push_back(d);
    }
};
```



```

void push_front(const T &x) {
    uint64_t d = hasher(x), c1 = d % MOD1, c2 = d % MOD2;
    top1 = top1 * BASE1 % MOD1; // Now B~len: the weight the new front element receives.
    top2 = top2 * BASE2 % MOD2;
    h1 = (c1 * top1 + h1) % MOD1;
    h2 = (c2 * top2 + h2) % MOD2;
    digits.push_front(d);
}

void pop_front() {
    assert(!digits.empty());
    uint64_t d = digits.front(), c1 = d % MOD1, c2 = d % MOD2;
    h1 = (h1 + MOD1 - c1 * top1 % MOD1) % MOD1;
    h2 = (h2 + MOD2 - c2 * top2 % MOD2) % MOD2;
    top1 = top1 * inv1 % MOD1;
    top2 = top2 * inv2 % MOD2;
    digits.pop_front();
}

void pop_back() {
    assert(!digits.empty());
    uint64_t d = digits.back(), c1 = d % MOD1, c2 = d % MOD2;
    h1 = (h1 + MOD1 - c1) % MOD1 * inv1 % MOD1;
    h2 = (h2 + MOD2 - c2) % MOD2 * inv2 % MOD2;
    top1 = top1 * inv1 % MOD1;
    top2 = top2 * inv2 % MOD2;
    digits.pop_back();
}

// Compares lengths first, so sequences of different lengths never compare equal.
friend bool operator==(const HashDeque &a, const HashDeque &b) {
    return a.digits.size() == b.digits.size() && a.h1 == b.h1 && a.h2 == b.h2;
}

friend bool operator!=(const HashDeque &a, const HashDeque &b) { return !(a == b); }
};

```

Example Usage

```

#include <functional>
#include <string>
using namespace std;

int main() {
    // Default hasher: characters hash by their code point.
    HashDeque<char> hd, hrev;
    for (char c : string("abc")) {
        hd.push_back(c);
    }
    for (char c : string("cba")) {
        hrev.push_front(c); // Prepending c, b, a also yields "abc".
    }
    assert(hd.size() == 3 && hd == hrev);

    // Slide a width-3 window across the text like Rabin-Karp: push_back + pop_front.
    string text = "Xabcab";
    auto target = hd.hash(); // Hash of "abc".
    HashDeque<char> window;
    for (int i = 0; i < 3; i++) {
        window.push_back(text[i]); // "Xab".
    }
    bool found = (window.hash() == target);
    for (int i = 3; i < static_cast<int>(text.size()); i++) {
        window.push_back(text[i]);
        window.pop_front();
    }
}

```

```

    found = found || (window.hash() == target);
}
assert(found); // "abc" occurs in "Xabcbab".

// pop_back undoes a push_back, restoring an earlier state's hash.
HashDeque<char> d;
d.push_back('a');
d.push_back('b');
auto ab = d.hash();
d.push_back('c');
d.pop_back();
assert(d.hash() == ab); // Back to "ab".

// An arbitrary element type via a custom hasher: a deque of whole words.
auto whash = [](const string &w) -> uint64_t { return std::hash<string>{}(w); };
HashDeque<string, decltype(whash)> sentence(whash), reversed(whash);
for (string w : {string("the"), string("quick"), string("fox")}) {
    sentence.push_back(w);
}
for (string w : {string("fox"), string("quick"), string("the")}) {
    reversed.push_front(w); // Same word sequence, assembled from the front.
}
assert(sentence == reversed);

sentence.pop_back(); // Drop "fox", leaving the words "the quick".
HashDeque<string, decltype(whash)> prefix(whash);
prefix.push_back("the");
prefix.push_back("quick");
assert(sentence == prefix);
return 0;
}

```

3.2.4 Dynamic Rolling Hash

Maintains a polynomial rolling hash for a sequence of nonnegative integers, supporting point updates and $O(\log n)$ contiguous subsequence hash queries. Unlike the static rolling hash, which fixes its prefix hashes at construction, this structure allows any element to change between queries, which is useful for problems that compare substrings while characters are edited.

The key idea is to weight each element by a power of the base tied to its absolute position: storing $a_i B^i$ at index i turns the hash of any prefix into an ordinary prefix sum, which a Fenwick tree maintains under point updates. The raw sum over a half-open range $[l, h)$ is $\sum_i a_i B^i$, and multiplying by B^{-l} shifts it to the position-independent value $\sum_i a_i B^{i-l}$, so two ranges with equal contents receive equal hashes (the leftmost element carries the lowest power). A segment tree could replace the Fenwick tree if lazy range updates are needed.

`MOD1` and `MOD2` are two distinct primes, each kept below 2^{31} so that the product of any two residues fits in a 64-bit integer; hashing modulo both and comparing the resulting pair makes a collision astronomically unlikely. `BASE1` and `BASE2` are the corresponding polynomial bases, which must be nonzero modulo their corresponding primes and may be randomized per run to resist adversarial inputs. For a small raw alphabet, choosing bases larger than every value also prevents carry-style collisions before modular reduction, but general integer values need not satisfy this bound.

Element values must be nonnegative. As with any polynomial hash, comparing subsequences of different lengths is only safe when values are positive (for example, map an alphabet to $[1, B)$) so that trailing zeros cannot make a shorter sequence collide with a longer one.

- `DynamicRollingHash(a)` builds prefix hashes over the integer sequence `a`.
- `size()` returns the length of the sequence.
- `set(i, x)` assigns the value at index `i` to `x`.
- `hash(lo, hi)` returns the hash pair of the half-open subsequence $[lo, hi)$. Equal ranges always have equal hash pairs, while equal hash pairs indicate equal ranges with high probability.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the length of the sequence.
- $O(\log n)$ per call to `set()` and `hash()`.
- $O(1)$ per call to `size()`.

Space Complexity:

- $O(n)$ for storage.
- $O(1)$ auxiliary per query.

```
#include <cassert>
#include <cstdint>
#include <utility>
#include <vector>

class DynamicRollingHash {
    static const uint64_t MOD1 = 1000000007, MOD2 = 1000000009; // Distinct primes below 2^31.
    static const uint64_t BASE1 = 131, BASE2 = 137;           // Nonzero modulo MOD1 and MOD2.

    int len;
    std::vector<int> cur;
    std::vector<uint64_t> pw1, pw2, ip1, ip2; // Base powers and inverse base powers.
    std::vector<uint64_t> ft1, ft2;         // Fenwick trees, 1-indexed of size len + 1.

    // Products and sums of residues below 2^30 stay under 2^60, so 64-bit arithmetic never overflows.
    static uint64_t powmod(uint64_t b, uint64_t e, uint64_t m) {
        uint64_t res = 1;
        for (b %= m; e > 0; e >>= 1) {
            if (e & 1) {
                res = res * b % m;
            }
            b = b * b % m;
        }
        return res;
    }

    static void add(std::vector<uint64_t> &ft, int i, uint64_t delta, uint64_t m) {
        for (i++; i < static_cast<int>(ft.size()); i += i & -i) {
            ft[i] = (ft[i] + delta) % m;
        }
    }

    static uint64_t sum(const std::vector<uint64_t> &ft, int i, uint64_t m) {
        uint64_t res = 0;
        for (i++; i > 0; i -= i & -i) {
            res = (res + ft[i]) % m;
        }
        return res;
    }

public:
    explicit DynamicRollingHash(const std::vector<int> &a)
        : len(static_cast<int>(a.size())), cur(len, 0) {
        pw1.resize(len + 1);
        pw2.resize(len + 1);
        ip1.resize(len + 1);
        ip2.resize(len + 1);
        pw1[0] = pw2[0] = ip1[0] = ip2[0] = 1;
        uint64_t invb1 = powmod(BASE1, MOD1 - 2, MOD1), invb2 = powmod(BASE2, MOD2 - 2, MOD2);
        for (int i = 1; i <= len; i++) {
            pw1[i] = pw1[i - 1] * BASE1 % MOD1;
            pw2[i] = pw2[i - 1] * BASE2 % MOD2;
            ip1[i] = ip1[i - 1] * invb1 % MOD1;
            ip2[i] = ip2[i - 1] * invb2 % MOD2;
        }
        ft1.assign(len + 1, 0);
```

```

    ft2.assign(len + 1, 0);
    for (int i = 0; i < len; i++) {
        set(i, a[i]);
    }
}

int size() const { return len; }

void set(int i, int x) {
    assert(0 <= i && i < len && x >= 0);
    uint64_t old1 = static_cast<uint64_t>(cur[i]) % MOD1 * pw1[i] % MOD1;
    uint64_t old2 = static_cast<uint64_t>(cur[i]) % MOD2 * pw2[i] % MOD2;
    uint64_t new1 = static_cast<uint64_t>(x) % MOD1 * pw1[i] % MOD1;
    uint64_t new2 = static_cast<uint64_t>(x) % MOD2 * pw2[i] % MOD2;
    add(ft1, i, (new1 + MOD1 - old1) % MOD1, MOD1);
    add(ft2, i, (new2 + MOD2 - old2) % MOD2, MOD2);
    cur[i] = x;
}

std::pair<uint64_t, uint64_t> hash(int lo, int hi) const {
    assert(0 <= lo && lo <= hi && hi <= len);
    uint64_t raw1 = (sum(ft1, hi - 1, MOD1) + MOD1 - sum(ft1, lo - 1, MOD1)) % MOD1;
    uint64_t raw2 = (sum(ft2, hi - 1, MOD2) + MOD2 - sum(ft2, lo - 1, MOD2)) % MOD2;
    return {raw1 * ip1[lo] % MOD1, raw2 * ip2[lo] % MOD2};
}
};

```

Example Usage

```

#include <string>
using namespace std;

int main() {
    string s = "abcabc";
    vector<int> a(s.begin(), s.end());
    DynamicRollingHash h(a);
    assert(h.size() == 6);
    assert(h.hash(0, 3) == h.hash(3, 6)); // "abc" == "abc".

    h.set(0, 'x'); // "xabcabc".
    assert(!(h.hash(0, 3) == h.hash(3, 6)));
    h.set(3, 'x'); // "xabcxabc".
    assert(h.hash(0, 3) == h.hash(3, 6)); // "xabc" == "xabc".

    // A single element's hash is invariant under where an equal value sits.
    assert(h.hash(1, 2) == h.hash(4, 5)); // both 'b'.
    return 0;
}

```

3.2.5 Zobrist Hashing

Computes randomized XOR fingerprints for keys and sets of keys using Zobrist hashing. Each distinct key is assigned a pseudorandom 64-bit token, and a set is represented by the XOR of its members' tokens. This makes insertion and deletion the same operation: XOR the key's token once.

This technique is useful for probabilistic equality checks on sets, distinct-prefix queries, board-game states, and other unordered collections. It is not cryptographic, and collisions are possible with probability roughly $q^2/2^{64}$ over q compared fingerprints. For multisets, plain XOR only records element parity; use a summed hash or pair it with counts if multiplicity matters.

- `ZobristHash<T>(seed = 0)` constructs a token generator with initial value `seed`.
- `get(x)` returns the stable token assigned to key `x`, creating it if needed.

- `toggle(h, x)` returns the hash obtained by inserting `x` into set hash `h` if absent, or removing `x` from `h` if present.
- `distinct_prefix_hashes(lo, hi)` returns `n + 1` hashes. Element `i` hashes the first `i` values in `[lo, hi)` with each distinct key contributing only on its first occurrence.

Time Complexity:

- $O(1)$ expected per call to `get()` and `toggle()`.
- $O(n)$ expected per call to `distinct_prefix_hashes(lo, hi)`, where n is the distance between `lo` and `hi`.

Space Complexity:

- $O(n)$ for stored key tokens, where n is the number of distinct keys seen so far.
- $O(n)$ auxiliary for `distinct_prefix_hashes()`.

```
#include <cstdint>
#include <unordered_map>
#include <unordered_set>
#include <vector>

template<typename T>
class ZobristHash {
    std::unordered_map<T, uint64_t> token;
    uint64_t state;

    // SplitMix64 mixer.
    static uint64_t mix64(uint64_t x) {
        x += 0x9e3779b97f4a7c15ULL;
        x = (x ^ (x >> 30)) * 0xbf58476d1ce4e5b9ULL;
        x = (x ^ (x >> 27)) * 0x94d049bb133111ebULL;
        return x ^ (x >> 31);
    }

public:
    explicit ZobristHash(uint64_t seed = 0) : state(seed) {}

    uint64_t get(const T &x) {
        if (auto it = token.find(x); it != token.end()) {
            return it->second;
        }
        state = mix64(state);
        token[x] = state;
        return state;
    }

    uint64_t toggle(uint64_t h, const T &x) { return h ^ get(x); }

    template<typename It>
    std::vector<uint64_t> distinct_prefix_hashes(It lo, It hi) {
        std::unordered_set<T> seen;
        std::vector<uint64_t> res(1, 0);
        uint64_t h = 0;
        for (It it = lo; it != hi; ++it) {
            if (seen.insert(*it).second) {
                h ^= get(*it);
            }
            res.push_back(h);
        }
        return res;
    }
};
```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    vector<int> a{1, 2, 1, 3};
    vector<int> b{2, 1, 3};
    ZobristHash<int> zh(1234567); // Fixed seed for reproducibility.
    vector<uint64_t> pa = zh.distinct_prefix_hashes(a.begin(), a.end());
    vector<uint64_t> pb = zh.distinct_prefix_hashes(b.begin(), b.end());

    assert(pa[2] == pb[2]); // Distinct sets {1, 2} and {2, 1}.
    assert(pa[4] == pb[3]); // Distinct sets {1, 2, 3} and {2, 1, 3}.

    uint64_t h = 0;
    h = zh.toggle(h, 5);
    assert(h != 0);
    h = zh.toggle(h, 5);
    assert(h == 0);
    return 0;
}
```

3.3 Pattern Matching

3.3.1 String Searching (KMP)

Given a single string (needle) and subsequent queries of texts (haystacks) to be searched, determine the first positions in which the needle occurs within the given haystacks in linear time using the Knuth-Morris-Pratt algorithm. Unlike `std::string::find`, KMP guarantees a linear worst-case bound. KMP precomputes, for every prefix of the needle, the length of its longest proper border (a prefix that is also a suffix). On a mismatch, the needle falls back to that border instead of restarting, so the position in the haystack never moves backward.

- `KMP(needle)` constructs the partial match table for a string `needle` that is to be searched for subsequently in `haystack` queries.
- `find_in(haystack)` returns the first position that `needle` occurs in `haystack`, or `std::string::npos` if it cannot be found. Note that the function can be modified to return all matches by simply letting the loop run and storing the results instead of returning early.

Time Complexity:

- $O(m)$ per call to the constructor, where m is the length of `needle`.
- $O(n)$ per call to `find_in(haystack)`, where n is the length of `haystack`.

Space Complexity:

- $O(m)$ for storage of the partial match table, where m is the length of `needle`.
- $O(1)$ auxiliary for `find_in()`.

```
#include <cstddef>
#include <string>
#include <vector>
using std::string;

class KMP {
    string needle;
    std::vector<int> table;

public:
    explicit KMP(const string &needle) : needle(needle), table(needle.size()) {
```

```

    for (int i = 1, j = 0; i < static_cast<int>(needle.size()); i++) {
        while (j > 0 && needle[i] != needle[j]) {
            j = table[j - 1];
        }
        if (needle[i] == needle[j]) {
            j++;
        }
        table[i] = j;
    }
}

std::size_t find_in(const string &haystack) const {
    int m = static_cast<int>(needle.size());
    if (m == 0) {
        return 0;
    }
    for (int i = 0, j = 0; i < static_cast<int>(haystack.size()); i++) {
        while (j > 0 && needle[j] != haystack[i]) {
            j = table[j - 1];
        }
        if (needle[j] == haystack[i]) {
            j++;
        }
        if (j == m) {
            return i + 1 - m;
        }
    }
    return string::npos;
}
};

```

Example Usage

```

#include <cassert>

int main() {
    KMP kmp("ABCDABD");
    assert(kmp.find_in("ABC ABCDAB ABCDABCDABDE") == 15);
    assert(kmp.find_in("ABC ABCDAB ABCDABCDAB") == string::npos);
    assert(KMP("aa").find_in("aaa") == 0); // First overlapping occurrence.
    assert(KMP("").find_in("abc") == 0);
    return 0;
}

```

3.3.2 String Searching (Z Algorithm)

Given a single string (needle) and a single text (haystack) to be searched, determine the first position in which the needle occurs within the haystack in linear time using the Z algorithm. In comparison, `std::string::find` does not guarantee a linear worst-case bound.

The Z array stores, for each position of a string, the length of the longest substring starting there that matches a prefix of the whole string. It is computed in linear time by maintaining the rightmost matched window and seeding each position from its mirror earlier in that window.

The `find_substring_z()` function below calls the Z algorithm on the concatenation of `needle` and `haystack`, separated by a sentinel value that is guaranteed not to collide with any byte in either input. Any position whose Z value reaches the needle's length marks an occurrence.

- `z_array(s)` constructs the Z array of sequence `s`. Each `z[i]` is the length of the longest prefix of `s` that also occurs starting at index `i`.

- `find_substring_z(haystack, needle)` returns the first position where `needle` occurs in `haystack`, or `std::string::npos` if it cannot be found. To return all matches, let the loop run and store the results instead of returning early.

Time Complexity:

- $O(n)$ per call to `z_array(s)`, where n is the length of `s`.
- $O(n + m)$ per call to `find_substring_z(haystack, needle)`, where n and m are the lengths of `haystack` and `needle`.

Space Complexity:

- $O(n)$ auxiliary for `z_array(s)`, where n is the length of `s`.
- $O(n + m)$ auxiliary for `find_substring_z(haystack, needle)`, where n is the length of `haystack` and m is the length of `needle`.

```
#include <algorithm>
#include <cstdint>
#include <string>
#include <vector>
using std::string;

template<typename Seq>
std::vector<int> z_array(const Seq &s) {
    std::vector<int> z(s.size());
    for (int i = 1, l = 0, r = 0; i < static_cast<int>(z.size()); i++) {
        if (i <= r) {
            z[i] = std::min(r - i + 1, z[i - 1]);
        }
        while (i + z[i] < static_cast<int>(z.size()) && s[z[i]] == s[i + z[i]]) {
            z[i]++;
        }
        if (r < i + z[i] - 1) {
            l = i;
            r = i + z[i] - 1;
        }
    }
    return z;
}

std::size_t find_substring_z(const string &haystack, const string &needle) {
    if (needle.empty()) {
        return 0;
    }
    std::vector<int> s;
    s.reserve(needle.size() + haystack.size() + 1);
    for (unsigned char c : needle) {
        s.push_back(c + 1);
    }
    s.push_back(0);
    for (unsigned char c : haystack) {
        s.push_back(c + 1);
    }
    auto z = z_array(s);
    int m = static_cast<int>(needle.size());
    for (int i = m + 1; i < static_cast<int>(z.size()); i++) {
        if (z[i] == m) {
            return i - m - 1;
        }
    }
    return string::npos;
}
```


Example Usage

```
#include <cassert>

int main() {
    assert((z_array(string("aaaaa")) == std::vector<int>{0, 4, 3, 2, 1}));
    assert(find_substring_z("ABC ABCDAB ABCDABCDABDE", "ABCDABD") == 15);
    assert(find_substring_z("ABC ABCDAB ABCDABCD", "ABCDABD") == string::npos);
    assert(find_substring_z("aaa", "aa") == 0); // First overlapping occurrence.
    assert(find_substring_z("abc", "") == 0);
    return 0;
}
```

3.3.3 String Searching (Aho-Corasick)

Given a set of strings (needles) and subsequent queries of texts (haystacks) to be searched, determine all positions in which needles occur within the given haystacks in linear time using the Aho-Corasick algorithm. The needles are arranged in a trie whose nodes carry failure links to the longest proper suffix that is also a trie node. Scanning the haystack once while following transitions and failure links then reports every match through each node's output set.

This implementation stores transitions in hash tables, which keeps lookups expected constant time without assuming a fixed alphabet size. The output tables are stored as ordered sets so matches ending at the same position are reported in deterministic needle order.

- `AhoCorasick(needles)` constructs the finite-state automaton for a set of needle strings that are to be searched for subsequently in haystack queries. Empty needles are ignored.
- `find_all_in(haystack, report_match)` calls the function `report_match(s, pos)` once on each occurrence of each needle that occurs in the `haystack`, where `pos` is the starting position in the `haystack` at which string `s` (a matched needle) occurs. The matches will be reported in increasing order of their ending positions within the `haystack`.

Time Complexity:

- $O(L + z \log m)$ expected per call to the constructor, where L is the total length of the m needles and z is the total number of inherited output entries.
- $O(n + z)$ expected per call to `find_all_in(haystack, report_match)`, where n is the length of `haystack` and z is the number of matches.

Space Complexity:

- $O(L + z)$ for storage of the automaton and inherited output entries.
- $O(1)$ auxiliary for `find_all_in()`.

```
#include <queue>
#include <set>
#include <string>
#include <unordered_map>
#include <vector>
using std::string;

class AhoCorasick {
    std::vector<string> needles;
    std::vector<int> fail;
    std::vector<std::unordered_map<char, int>> adj;
    std::vector<std::set<int>> out;

    int next_state(int curr, char c) const {
        int next = curr;
        while (next != 0 && adj[next].find(c) == adj[next].end()) {
            next = fail[next];
        }
    }
};
```

```

    }
    auto it = adj[next].find(c);
    return (it != adj[next].end()) ? it->second : 0;
}

public:
explicit AhoCorasick(const std::vector<string> &needles) : needles(needles) {
    int total_len = 0;
    for (const auto &needle : needles) {
        total_len += static_cast<int>(needle.size());
    }
    // The trie has at most total_len + 1 states: the root plus one per character when no prefixes
    // are shared. Sizing to total_len alone overflows for low-sharing needle sets (e.g. a single
    // needle, which still needs the root plus one node).
    int max_states = total_len + 1;
    fail.resize(max_states, -1);
    adj.resize(max_states);
    out.resize(max_states);
    int states = 1;
    for (int i = 0; i < static_cast<int>(needles.size()); i++) {
        if (needles[i].empty()) {
            continue;
        }
        int curr = 0;
        for (char c : needles[i]) {
            if (auto it = adj[curr].find(c); it != adj[curr].end()) {
                curr = it->second;
            } else {
                curr = adj[curr][c] = states++;
            }
        }
        out[curr].insert(i);
    }
    std::queue<int> q;
    for (auto &[c, v] : adj[0]) {
        if (v != 0) {
            fail[v] = 0;
            q.push(v);
        }
    }
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        for (auto &[c, v] : adj[u]) {
            int f = fail[u];
            while (f != 0 && adj[f].find(c) == adj[f].end()) {
                f = fail[f];
            }
            auto fit = adj[f].find(c);
            f = (fit != adj[f].end()) ? fit->second : 0;
            fail[v] = f;
            out[v].insert(out[f].begin(), out[f].end());
            q.push(v);
        }
    }
}

template<typename Fn>
void find_all_in(const string &haystack, Fn report_match) const {
    int state = 0;
    for (int i = 0; i < static_cast<int>(haystack.size()); i++) {
        state = next_state(state, haystack[i]);
        for (int idx : out[state]) {
            report_match(needles[idx], i - static_cast<int>(needles[idx].size()) + 1);
        }
    }
}
};

```

Example Usage

```
#include <algorithm>
#include <cassert>
#include <utility>
using namespace std;

int main() {
    vector<string> needles;
    needles.push_back("a");
    needles.push_back("ab");
    needles.push_back("bab");
    needles.push_back("bc");
    needles.push_back("bca");
    needles.push_back("c");
    needles.push_back("caa");
    needles.push_back("abccab");

    AhoCorasick ac(needles);
    vector<pair<string, int>> matches;
    ac.find_all_in("abccab", [&](const string &needle, int pos) {
        matches.push_back({needle, pos});
    });
    sort(matches.begin(), matches.end());
    assert(
        (matches ==
         vector<pair<string, int>>{
             {"a", 0}, {"a", 4}, {"ab", 0}, {"ab", 4}, {"abccab", 0}, {"bc", 1}, {"c", 2}, {"c", 3}
         })
    );

    vector<pair<string, int>> none;
    ac.find_all_in("zzzz", [&](const string &needle, int pos) { none.push_back({needle, pos}); });
    assert(none.empty());
    return 0;
}
```

3.3.4 Palindrome Searching (Manacher)

Computes all odd- and even-length palindromic radii of a string in linear time using Manacher's algorithm. These radii can be used to test whether any substring is a palindrome in $O(1)$, find the longest palindromic substring, or count all palindromic substrings. The algorithm tracks the rightmost known palindrome and seeds each new center with the radius of its mirror inside that window, so characters are only ever compared when a palindrome is extended past the window.

For odd palindromes, `odd[i]` is the radius centered at character `s[i]`, including the center character. Thus the palindrome spans `[i - odd[i] + 1, i + odd[i])`. For even palindromes, `even[i]` is the radius centered between `s[i - 1]` and `s[i]`. Thus the palindrome spans `[i - even[i], i + even[i])`.

- `Manacher(s)` constructs the palindromic radii arrays for string `s`.
- `is_palindrome(lo, hi)` returns whether substring `[lo, hi)` is a palindrome.
- `longest_palindrome()` returns one longest palindromic substring of `s`.
- `count_palindromes()` returns the total number of palindromic substrings of `s`, counted with multiplicity by position.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the length of `s`.
- $O(1)$ per call to `is_palindrome()`.
- $O(n)$ per call to `longest_palindrome()` and `count_palindromes()`.

Space Complexity:

- $O(n)$ for storage of the radii arrays.
- $O(1)$ auxiliary per query.

```

#include <algorithm>
#include <cassert>
#include <cstdlib>
#include <string>
#include <vector>
using std::string;

class Manacher {
    string s;
    std::vector<int> odd, even;

public:
    explicit Manacher(const string &s) : s(s), odd(s.size()), even(s.size()) {
        int n = static_cast<int>(s.size());
        for (int i = 0, l = 0, r = -1; i < n; i++) {
            int k = (i > r) ? 1 : std::min(odd[l + r - i], r - i + 1);
            while (0 <= i - k && i + k < n && s[i - k] == s[i + k]) {
                k++;
            }
            odd[i] = k;
            if (i + k - 1 > r) {
                l = i - k + 1;
                r = i + k - 1;
            }
        }
        for (int i = 0, l = 0, r = -1; i < n; i++) {
            int k = (i > r) ? 0 : std::min(even[l + r - i + 1], r - i + 1);
            while (0 <= i - k - 1 && i + k < n && s[i - k - 1] == s[i + k]) {
                k++;
            }
            even[i] = k;
            if (i + k - 1 > r) {
                l = i - k;
                r = i + k - 1;
            }
        }
    }

    bool is_palindrome(int lo, int hi) const {
        assert(0 <= lo && lo <= hi && hi <= static_cast<int>(s.size()));
        int len = hi - lo;
        if (len <= 0) {
            return true;
        }
        int mid = lo + (hi - lo) / 2;
        return (len % 2) ? odd[mid] >= len / 2 + 1 : even[mid] >= len / 2;
    }

    string longest_palindrome() const {
        int best_l = 0, best_len = 0;
        int n = static_cast<int>(s.size());
        for (int i = 0; i < n; i++) {
            int len = 2 * odd[i] - 1;
            if (len > best_len) {
                best_len = len;
                best_l = i - odd[i] + 1;
            }
            len = 2 * even[i];
            if (len > best_len) {
                best_len = len;
                best_l = i - even[i];
            }
        }
    }
}

```

```

    return s.substr(best_l, best_len);
}

int64_t count_palindromes() const {
    int64_t res = 0;
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        res += odd[i] + even[i];
    }
    return res;
}
};

```

Example Usage

```

int main() {
    Manacher m("abacaba");
    assert(m.is_palindrome(0, 7));
    assert(m.is_palindrome(1, 6));
    assert(!m.is_palindrome(0, 6));
    assert(m.longest_palindrome() == "abacaba");
    assert(m.count_palindromes() == 12);

    Manacher e("cabbad");
    assert(e.is_palindrome(1, 5));
    assert(e.longest_palindrome() == "abba");
    return 0;
}

```

3.3.5 Lyndon Factorization (Duval)

A Lyndon word is a nonempty string that is strictly smaller than all of its proper suffixes, or equivalently, strictly smaller than all of its nontrivial cyclic rotations. By the Chen-Fox-Lyndon theorem, every string has a unique factorization into a sequence of Lyndon words whose values are non-increasing: $s = w_1 w_2 \cdots w_k$ with $w_1 \geq w_2 \geq \cdots \geq w_k$. Duval's algorithm computes this factorization in linear time and constant auxiliary space by scanning the string while tracking the current candidate Lyndon prefix and the position it last matched.

The same scan, applied to the string viewed as a cycle, yields the starting position of its lexicographically least rotation. This is the standard way to canonicalize a necklace, for example to test whether two strings are cyclic rotations of one another.

Both functions below operate on any random-access sequence whose elements support `operator<`.

- `lyndon_factorization(s)` returns a vector of the starting indices of the Lyndon factors of `s`, in order. The factor at index `i` of the result spans `[res[i], res[i + 1])`, or `[res[i], s.size())` for the final factor.
- `min_rotation(s)` returns the 0-based position `p` such that the rotation beginning at `p`, that is, `s[p]`, `s[(p + 1) % n]`, ..., `s[(p + n - 1) % n]`, is lexicographically smallest. If several rotations tie, the smallest such `p` is returned.

Time Complexity: $O(n)$ per call to `lyndon_factorization(s)` and `min_rotation(s)`, where n is the length of `s`.

Space Complexity:

- $O(k)$ auxiliary for `lyndon_factorization()`, where k is the number of factors.
- $O(1)$ auxiliary for `min_rotation()`.

```

#include <vector>

template<typename Seq>
std::vector<int> lyndon_factorization(const Seq &s) {
    int n = static_cast<int>(s.size());
    std::vector<int> starts;

```

```

int i = 0;
while (i < n) {
    int j = i + 1, k = i;
    while (j < n && !(s[j] < s[k])) {
        k = (s[k] < s[j]) ? i : k + 1;
        j++;
    }
    while (i <= k) {
        starts.push_back(i);
        i += j - k;
    }
}
return starts;
}

template<typename Seq>
int min_rotation(const Seq &s) {
    int n = static_cast<int>(s.size());
    if (n == 0) {
        return 0;
    }
    int i = 0, res = 0;
    while (i < n) {
        res = i;
        int j = i + 1, k = i;
        while (j < 2 * n && !(s[j % n] < s[k % n])) {
            k = (s[k % n] < s[j % n]) ? i : k + 1;
            j++;
        }
        while (i <= k) {
            i += j - k;
        }
    }
    return res;
}

```

Example Usage

```

#include <cassert>
#include <string>
using std::string;

int main() {
    // "banana" factors as "b" >= "an" >= "an" >= "a".
    assert((lyndon_factorization(string("banana")) == std::vector<int>{0, 1, 3, 5}));

    // Among the rotations of "baca", "abac" (starting at index 3) is smallest.
    assert(min_rotation(string("baca")) == 3);

    // Both routines work on any comparable sequence, not just strings.
    assert(min_rotation(std::vector<int>{3, 1, 2, 1, 3}) == 1);
    return 0;
}

```

3.3.6 De Bruijn Sequence

A de Bruijn sequence $B(k, n)$ is a cyclic string over the alphabet $[0, k)$ of length k^n in which every one of the k^n possible length- n strings appears exactly once as a (cyclic) substring. Reading a window of n symbols and advancing one position at a time therefore cycles through all length- n strings without repetition, which makes the sequence a compact way to probe every short pattern: cracking a rotating lock, addressing a shift register, or building the perfect-hash table that maps an isolated set bit to its index.

The construction here is the Fredricksen-Kessler-Maiorana (FKM) algorithm. It concatenates, in lexicographic order, the Lyndon words (see the Lyndon factorization section) whose length divides n . The result is the lexicographically smallest de Bruijn sequence. The recursion enumerates these necklace representatives directly, appending the first p symbols of each candidate whenever its period p divides n .

- `de_bruijn(k, n)` returns a vector of length k^n holding one period of $B(k, n)$. Each entry lies in $[0, k)$. Both k and n must be at least 1. Treat the result as cyclic: the length- n windows that wrap past the end are exactly the ones still needed to cover all strings.

Time Complexity: $O(k^n)$ per call, which is optimal since the output itself has length k^n .

Space Complexity: $O(k^n)$ for the returned sequence, plus $O(n)$ auxiliary stack space for the recursion.

```
#include <cassert>
#include <vector>

static void de_bruijn_gen(
    int t, int p, int k, int n, std::vector<int> &word, std::vector<int> &seq
) {
    if (t > n) {
        if (n % p == 0) {
            for (int i = 1; i <= p; i++) {
                seq.push_back(word[i]);
            }
        }
        return;
    }
    word[t] = word[t - p];
    de_bruijn_gen(t + 1, p, k, n, word, seq);
    for (int j = word[t - p] + 1; j < k; j++) {
        word[t] = j;
        de_bruijn_gen(t + 1, t, k, n, word, seq);
    }
}

std::vector<int> de_bruijn(int k, int n) {
    assert(k >= 1 && n >= 1);
    std::vector<int> word(n + 1, 0), seq;
    de_bruijn_gen(1, 1, k, n, word, seq);
    return seq;
}
```

Example Usage

```
#include <set>
#include <vector>
using namespace std;

int main() {
    // B(2, 3) has length 8 and every 3-bit string appears once cyclically.
    vector<int> seq = de_bruijn(2, 3);
    assert(seq.size() == 8);

    for (int k = 2; k <= 4; k++) {
        for (int n = 1; n <= 4; n++) {
            vector<int> s = de_bruijn(k, n);
            int len = 1;
            for (int i = 0; i < n; i++) {
                len *= k;
            }
            assert(static_cast<int>(s.size()) == len);

            // Collect every length-n window, wrapping cyclically, and confirm all k^n are distinct.
            set<vector<int>> windows;
            for (int i = 0; i < len; i++) {
```

```

    vector<int> w(n);
    for (int j = 0; j < n; j++) {
        w[j] = s[(i + j) % len];
    }
    windows.insert(w);
}
assert(static_cast<int>(windows.size()) == len);
}
}
return 0;
}

```

3.4 Suffix Arrays and LCP

3.4.1 Suffix Array and LCP (Manber-Myers)

Given a string s , a suffix array stores the starting positions of the suffixes of s in sorted order. That is, the i -th position of the suffix array stores the starting position of the i -th lexicographically smallest suffix of s . For example, $s = \text{"cab"}$ has the suffixes `"cab"`, `"ab"`, and `"b"`. In sorted order they are `"ab"`, `"b"`, and `"cab"`, so their starting indices form the suffix array $[1, 2, 0]$.

For a string s of length n , the longest common prefix (LCP) array of length $n - 1$ stores the lengths of the longest common prefixes between all pairs of lexicographically adjacent suffixes in s . For example, `"baa"` has the sorted suffixes `"a"`, `"aa"`, and `"baa"`, with an LCP array of $[1, 0]$.

The original Manber-Myers doubling algorithm comparison-sorts the suffixes by their first 2^k characters for increasing k : each round orders suffixes by their pair of ranks from the previous round, so a comparison costs $O(1)$ and the full order emerges after $O(\log n)$ rounds.

- `SuffixArrayManberMyers(s)` constructs a suffix array from string `s`.
- `get_sa()` returns the constructed suffix array.
- `get_lcp()` returns the corresponding LCP array for the suffix array.
- `find(needle)` returns one position that `needle` occurs in `s` (not necessarily the first), or `std::string::npos` if it cannot be found. For a `needle` of length m , this implementation uses an $O(m \log n)$ binary search, but can be optimized to $O(m + \log n)$ by first computing the LCP-LR array using the LCP array.

Time Complexity:

- $O(n \log^2 n)$ per call to the constructor, where n is the length of `s`.
- $O(1)$ per call to `get_sa()`.
- $O(n)$ per call to `get_lcp()`, where n is the length of `s`.
- $O(m \log n)$ per call to `find(needle)`, where m is the length of `needle` and n is the length of `s`.

Space Complexity:

- $O(n)$ for storage of the suffix and LCP arrays.
- $O(n)$ auxiliary for the constructor.
- $O(1)$ auxiliary for all other operations.

```

#include <algorithm>
#include <cstdint>
#include <numeric>
#include <string>
#include <utility>
#include <vector>
using std::string;

class SuffixArrayManberMyers {
    string s;

```



```

std::vector<int> sa, rk;

public:
explicit SuffixArrayManberMyers(const string &s) : s(s), sa(s.size()), rk(s.size()) {
    int n = static_cast<int>(s.size());
    std::iota(sa.begin(), sa.end(), 0);
    for (int i = 0; i < n; i++) {
        rk[i] = static_cast<unsigned char>(s[i]);
    }
    std::vector<std::pair<int, int>> rk2(n);
    for (int gap = 1; gap < n; gap *= 2) {
        for (int i = 0; i < n; i++) {
            rk2[i] = {rk[i], i + gap < n ? rk[i + gap] + 1 : 0};
        }
        std::sort(sa.begin(), sa.end(), [&](int i, int j) { return rk2[i] < rk2[j]; });
        for (int i = 0; i < n; i++) {
            rk[sa[i]] = (i > 0 && rk2[sa[i - 1]] == rk2[sa[i]]) ? rk[sa[i - 1]] : i;
        }
    }
}

const std::vector<int> &get_sa() const { return sa; }

std::vector<int> get_lcp() const {
    int n = static_cast<int>(s.size());
    if (n == 0) {
        return {};
    }
    std::vector<int> rank(n), lcp(n - 1);
    for (int i = 0; i < n; i++) {
        rank[sa[i]] = i;
    }
    for (int i = 0, k = 0; i < n; i++) {
        if (rank[i] < n - 1) {
            int j = sa[rank[i] + 1];
            while (std::max(i, j) + k < n && s[i + k] == s[j + k]) {
                k++;
            }
            lcp[rank[i]] = k;
            if (k > 0) {
                k--;
            }
        }
    }
    return lcp;
}

std::size_t find(const string &needle) const {
    if (needle.empty()) {
        return 0;
    }
    int lo = 0, hi = static_cast<int>(s.size()) - 1;
    while (lo <= hi) {
        int mid = lo + (hi - lo) / 2;
        int cmp = s.compare(sa[mid], needle.size(), needle);
        if (cmp < 0) {
            lo = mid + 1;
        } else if (cmp > 0) {
            hi = mid - 1;
        } else {
            return sa[mid];
        }
    }
    return string::npos;
}
};

```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    SuffixArrayManberMyers sa("banana");
    vector<int> sarr = sa.get_sa(), lcp = sa.get_lcp();
    vector<int> sarr_expected{5, 3, 1, 0, 4, 2};
    vector<int> lcp_expected{1, 3, 0, 0, 2};
    assert(sarr == sarr_expected);
    assert(lcp == lcp_expected);
    assert(sa.find("ana") == 1);
    assert(sa.find("x") == string::npos);
    return 0;
}
```

3.4.2 Suffix Array and LCP (Counting Sort)

Given a string s , a suffix array stores the starting positions of the suffixes of s in sorted order. That is, the i -th position of the suffix array stores the starting position of the i -th lexicographically smallest suffix of s . For example, $s = \text{"cab"}$ has the suffixes "cab" , "ab" , and "b" . In sorted order they are "ab" , "b" , and "cab" , so their starting indices form the suffix array $[1, 2, 0]$.

For a string s of length n , the longest common prefix (LCP) array of length $n - 1$ stores the lengths of the longest common prefixes between all pairs of lexicographically adjacent suffixes in s . For example, "baa" has the sorted suffixes "a" , "aa" , and "baa" , with an LCP array of $[1, 0]$.

The original Manber-Myers doubling algorithm comparison-sorts the suffixes by their first 2^k characters for increasing k : each round orders suffixes by their pair of ranks from the previous round, so a comparison costs $O(1)$ and the full order emerges after $O(\log n)$ rounds. Replacing the comparison sort of each round with counting-sort-style rank updates removes a logarithmic factor.

- `SuffixArrayCountingSort(s)` constructs a suffix array from string `s`.
- `get_sa()` returns the constructed suffix array.
- `get_lcp()` returns the corresponding LCP array for the suffix array.
- `find(needle)` returns one position that `needle` occurs in `s` (not necessarily the first), or `std::string::npos` if it cannot be found. For a `needle` of length m , this implementation uses an $O(m \log n)$ binary search, but can be optimized to $O(m + \log n)$ by first computing the LCP-LR array using the LCP array.

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the length of `s`.
- $O(1)$ per call to `get_sa()`.
- $O(n)$ per call to `get_lcp()`, where n is the length of `s`.
- $O(m \log n)$ per call to `find(needle)`, where m is the length of `needle` and n is the length of `s`.

Space Complexity:

- $O(n)$ for storage of the suffix and LCP arrays.
- $O(n)$ auxiliary for the constructor.
- $O(1)$ auxiliary for all other operations.

```
#include <algorithm>
#include <cstdint>
#include <numeric>
#include <string>
#include <vector>
using std::string;

class SuffixArrayCountingSort {
```

```

string s;
std::vector<int> sa, rk;

public:
explicit SuffixArrayCountingSort(const string &s) : s(s), sa(s.size()), rk(s.size()) {
    int n = static_cast<int>(s.size());
    std::iota(sa.rbegin(), sa.rend(), 0);
    for (int i = 0; i < n; i++) {
        rk[i] = static_cast<unsigned char>(s[i]);
    }
    std::stable_sort(sa.begin(), sa.end(), [&](int i, int j) {
        return static_cast<unsigned char>(s[i]) < static_cast<unsigned char>(s[j]);
    });
    for (int gap = 1; gap < n; gap *= 2) {
        std::vector<int> prev_rk(rk), prev_sa(sa), cnt(n);
        std::iota(cnt.begin(), cnt.end(), 0);
        for (int i = 0; i < n; i++) {
            rk[sa[i]] = (i > 0 && prev_rk[sa[i - 1]] == prev_rk[sa[i]] && sa[i - 1] + gap < n &&
                prev_rk[sa[i - 1] + gap / 2] == prev_rk[sa[i] + gap / 2])
                ? rk[sa[i - 1]]
                : i;
        }
        for (int i = 0; i < n; i++) {
            int s1 = prev_sa[i] - gap;
            if (s1 >= 0) {
                sa[cnt[rk[s1]]++] = s1;
            }
        }
    }
}

const std::vector<int> &get_sa() const { return sa; }

std::vector<int> get_lcp() const {
    int n = static_cast<int>(s.size());
    if (n == 0) {
        return {};
    }
    std::vector<int> rank(n), lcp(n - 1);
    for (int i = 0; i < n; i++) {
        rank[sa[i]] = i;
    }
    for (int i = 0, k = 0; i < n; i++) {
        if (rank[i] < n - 1) {
            int j = sa[rank[i] + 1];
            while (std::max(i, j) + k < n && s[i + k] == s[j + k]) {
                k++;
            }
            lcp[rank[i]] = k;
            if (k > 0) {
                k--;
            }
        }
    }
    return lcp;
}

std::size_t find(const string &needle) const {
    if (needle.empty()) {
        return 0;
    }
    int lo = 0, hi = static_cast<int>(s.size()) - 1;
    while (lo <= hi) {
        int mid = lo + (hi - lo) / 2;
        int cmp = s.compare(sa[mid], needle.size(), needle);
        if (cmp < 0) {
            lo = mid + 1;
        } else if (cmp > 0) {

```

```

        hi = mid - 1;
    } else {
        return sa[mid];
    }
}
return string::npos;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    SuffixArrayCountingSort sa("banana");
    vector<int> sarr = sa.get_sa(), lcp = sa.get_lcp();
    vector<int> sarr_expected{5, 3, 1, 0, 4, 2};
    vector<int> lcp_expected{1, 3, 0, 0, 2};
    assert(sarr == sarr_expected);
    assert(lcp == lcp_expected);
    assert(sa.find("ana") == 1);
    assert(sa.find("x") == string::npos);

    SuffixArrayCountingSort repeated("bbba");
    assert((repeated.get_lcp() == vector<int>{0, 1, 2}));
    return 0;
}

```

3.4.3 Suffix Array and LCP (DC3)

Given a string s , a suffix array stores the starting positions of the suffixes of s in sorted order. That is, the i -th position of the suffix array stores the starting position of the i -th lexicographically smallest suffix of s . For example, $s = \text{"cab"}$ has the suffixes "cab" , "ab" , and "b" . In sorted order they are "ab" , "b" , and "cab" , so their starting indices form the suffix array $[1, 2, 0]$.

For a string s of length n , the longest common prefix (LCP) array of length $n - 1$ stores the lengths of the longest common prefixes between all pairs of lexicographically adjacent suffixes in s . For example, "baa" has the sorted suffixes "a" , "aa" , and "baa" , with an LCP array of $[1, 0]$.

The DC3/skew algorithm recursively constructs the suffix array of the suffixes starting at positions not divisible by 3, uses that order to sort the remaining suffixes with radix sort, and merges the two sorted groups in linear time.

- `SuffixArrayDC3(s)` constructs a suffix array from string `s`.
- `get_sa()` returns the constructed suffix array.
- `get_lcp()` returns the corresponding LCP array for the suffix array.
- `find(needle)` returns one position that `needle` occurs in `s` (not necessarily the first), or `std::string::npos` if it cannot be found. For a `needle` of length m , this implementation uses an $O(m \log n)$ binary search, but can be optimized to $O(m + \log n)$ by first computing the LCP-LR array using the LCP array.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the length of `s`.
- $O(1)$ per call to `get_sa()`.
- $O(n)$ per call to `get_lcp()`, where n is the length of `s`.
- $O(m \log n)$ per call to `find(needle)`, where m is the length of `needle` and n is the length of `s`.

Space Complexity:

- $O(n)$ for storage of the suffix and LCP arrays.
- $O(n)$ auxiliary for the constructor.

- $O(1)$ auxiliary for all other operations.

```

#include <algorithm>
#include <cstdint>
#include <string>
#include <vector>
using std::string;

class SuffixArrayDC3 {
    static bool leq(int a1, int a2, int b1, int b2) { return (a1 < b1) || (a1 == b1 && a2 <= b2); }

    static bool leq(int a1, int a2, int a3, int b1, int b2, int b3) {
        return (a1 < b1) || (a1 == b1 && leq(a2, a3, b2, b3));
    }

    template<typename It>
    static void radix_pass(It a, It b, It r, int n, int alphabet) {
        std::vector<int> cnt(alphabet + 1);
        for (int i = 0; i < n; i++) {
            cnt[r[a[i]]]++;
        }
        for (int i = 1; i <= alphabet; i++) {
            cnt[i] += cnt[i - 1];
        }
        for (int i = n - 1; i >= 0; i--) {
            b[--cnt[r[a[i]]]] = a[i];
        }
    }

    template<typename It>
    static void suffix_array_dc3(It s, It sa, int n, int alphabet) {
        int n0 = (n + 2) / 3, n1 = (n + 1) / 3, n2 = n / 3, n02 = n0 + n2;
        std::vector<int> s12(n02 + 3), sa12(n02 + 3), s0(n0), sa0(n0);
        s12[n02] = s12[n02 + 1] = s12[n02 + 2] = 0;
        sa12[n02] = sa12[n02 + 1] = sa12[n02 + 2] = 0;
        for (int i = 0, j = 0; i < n + n0 - n1; i++) {
            if (i % 3 != 0) {
                s12[j++] = i;
            }
        }
        radix_pass(s12.begin(), sa12.begin(), s + 2, n02, alphabet);
        radix_pass(sa12.begin(), s12.begin(), s + 1, n02, alphabet);
        radix_pass(s12.begin(), sa12.begin(), s, n02, alphabet);
        int name = 0, c0 = -1, c1 = -1, c2 = -1;
        for (int i = 0; i < n02; i++) {
            if (s[sa12[i]] != c0 || s[sa12[i] + 1] != c1 || s[sa12[i] + 2] != c2) {
                name++;
                c0 = s[sa12[i]];
                c1 = s[sa12[i] + 1];
                c2 = s[sa12[i] + 2];
            }
            (sa12[i] % 3 == 1 ? s12[sa12[i] / 3] : s12[sa12[i] / 3 + n0]) = name;
        }
        if (name < n02) {
            suffix_array_dc3(s12.begin(), sa12.begin(), n02, name);
            for (int i = 0; i < n02; i++) {
                s12[sa12[i]] = i + 1;
            }
        } else {
            for (int i = 0; i < n02; i++) {
                sa12[s12[i] - 1] = i;
            }
        }
        for (int i = 0, j = 0; i < n02; i++) {
            if (sa12[i] < n0) {
                s0[j++] = 3 * sa12[i];
            }
        }
    }
}

```

```

radix_pass(s0.begin(), sa0.begin(), s, n0, alphabet);
for (int p = 0, t = n0 - n1, k = 0; k < n; k++) {
    int i = (sa12[t] < n0) ? 3 * sa12[t] + 1 : 3 * (sa12[t] - n0) + 2, j = sa0[p];
    if (sa12[t] < n0
        ? leq(s[i], s12[sa12[t] + n0], s[j], s12[j / 3])
        : leq(s[i], s[i + 1], s12[sa12[t] - n0 + 1], s[j], s[j + 1], s12[j / 3 + n0])) {
        sa[k] = i;
        if (++t == n02) {
            for (k++; p < n0; p++, k++) {
                sa[k] = sa0[p];
            }
        }
    } else {
        sa[k] = j;
        if (++p == n0) {
            for (k++; t < n02; t++, k++) {
                sa[k] = (sa12[t] < n0) ? 3 * sa12[t] + 1 : 3 * (sa12[t] - n0) + 2;
            }
        }
    }
}
}

string s;
std::vector<int> sa;

public:
explicit SuffixArrayDC3(const string &s) : s(s), sa(s.size() + 1) {
    int n = static_cast<int>(s.size());
    std::vector<int> scopy(n);
    for (int i = 0; i < n; i++) {
        scopy[i] = static_cast<unsigned char>(s[i]) + 1;
    }
    scopy.resize(n + 4);
    suffix_array_dc3(scopy.begin(), sa.begin(), n + 1, 256);
    sa.erase(sa.begin());
}

const std::vector<int> &get_sa() const { return sa; }

std::vector<int> get_lcp() const {
    int n = static_cast<int>(s.size());
    if (n == 0) {
        return {};
    }
    std::vector<int> rank(n), lcp(n - 1);
    for (int i = 0; i < n; i++) {
        rank[sa[i]] = i;
    }
    for (int i = 0, k = 0; i < n; i++) {
        if (rank[i] < n - 1) {
            int j = sa[rank[i] + 1];
            while (std::max(i, j) + k < n && s[i + k] == s[j + k]) {
                k++;
            }
            lcp[rank[i]] = k;
            if (k > 0) {
                k--;
            }
        }
    }
    return lcp;
}

std::size_t find(const string &needle) const {
    if (needle.empty()) {
        return 0;
    }
}

```

```

int lo = 0, hi = static_cast<int>(s.size()) - 1;
while (lo <= hi) {
    int mid = lo + (hi - lo) / 2;
    int cmp = s.compare(sa[mid], needle.size(), needle);
    if (cmp < 0) {
        lo = mid + 1;
    } else if (cmp > 0) {
        hi = mid - 1;
    } else {
        return sa[mid];
    }
}
return string::npos;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    SuffixArrayDC3 sa("banana");
    vector<int> sarr = sa.get_sa(), lcp = sa.get_lcp();
    vector<int> sarr_expected{5, 3, 1, 0, 4, 2};
    vector<int> lcp_expected{1, 3, 0, 0, 2};
    assert(sarr == sarr_expected);
    assert(lcp == lcp_expected);
    assert(sa.find("ana") == 1);
    assert(sa.find("x") == string::npos);
    return 0;
}

```

3.5 String Data Structures

3.5.1 Trie

Maintain a map of strings to values using a tree data structure. Each node corresponds to a character, and each inserted string corresponds to a path from the root to a node that is flagged as a terminal node. Children are stored in hash tables, and `walk()` sorts child labels as needed to preserve lexicographic traversal.

- `Trie<V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(s, v)` adds an entry with string key `s` and value `v` to the map, returning `true` if a new entry was added or `false` if the string already exists (in which case the map is unchanged and the old value associated with the string key is preserved).
- `erase(s)` removes the entry with string key `s` from the map, returning `true` if the removal was successful or `false` if the string to be removed was not found.
- `find(s)` returns a pointer to a const value associated with string key `s`, or `nullptr` if the key was not found.
- `count_prefix(s)` returns the number of keys currently in the map that have `s` as a prefix (a key equal to `s` counts as well). `count_prefix("")` therefore equals `size()`.
- `walk(f)` calls the function `f(s, v)` on each entry of the map, in lexicographically ascending order of the string keys.

For a small, fixed alphabet, the `std::unordered_map` children can be replaced by a fixed-size array of child pointers indexed by character. This removes the per-step hashing and lets `walk()` traverse in sorted order without sorting each node's labels, at the cost of restricting the alphabet and spending more memory per node.

Time Complexity:

- $O(n)$ expected per call to `insert(s, v)`, `erase(s)`, `find(s)`, and `count_prefix(s)`, where n is the length of `s`.
- $O(l)$ per call to `walk()`, where l is the total length of string keys that are currently in the map, treating the character alphabet as constant.
- $O(1)$ per call to all other operations.

Space Complexity:

- $O(l)$ for storage of the Trie, where l is the total length of string keys that are currently in the map.
- $O(n)$ auxiliary stack space for destruction and `walk()`, where n is the maximum length of any string that has been inserted so far.
- $O(n)$ auxiliary stack space for `erase(s)`, where n is the length of `s`.
- $O(1)$ auxiliary for all other operations.

```
#include <algorithm>
#include <string>
#include <unordered_map>
#include <utility>
#include <vector>
using std::string;

template<typename V>
class Trie {
    struct Node {
        V value;
        bool is_terminal;
        int cnt; // Optional: maintain count of terminal keys in this subtree for count_prefix queries.
        std::unordered_map<char, Node*> children;

        Node() : is_terminal(false), cnt(0) {}
    } *root;

    static bool erase(Node *n, const string &s, int i) {
        if (i == static_cast<int>(s.size())) {
            if (!n->is_terminal) {
                return false;
            }
            n->is_terminal = false;
            n->cnt--;
            return true;
        }
        auto it = n->children.find(s[i]);
        if (it == n->children.end() || !erase(it->second, s, i + 1)) {
            return false;
        }
        n->cnt--;
        // Prune the child only if it is now a non-terminal leaf; a childless terminal is still a key.
        if (it->second->children.empty() && !it->second->is_terminal) {
            delete it->second;
            n->children.erase(it);
        }
        return true;
    }
};

template<typename Fn>
static void walk(Node *n, string &s, Fn f) {
    if (n->is_terminal) {
        f(s, n->value);
    }
    std::vector<std::pair<char, Node*>> children(n->children.begin(), n->children.end());
    std::sort(children.begin(), children.end());
```



```

    for (auto &[c, child] : children) {
        s += c;
        walk(child, s, f);
        s.pop_back();
    }
}

static void clean_up(Node *n) {
    for (auto &[c, child] : n->children) {
        clean_up(child);
    }
    delete n;
}

int num_terminals;

public:
Trie() : root(new Node()), num_terminals(0) {}

~Trie() { clean_up(root); }
Trie(const Trie &) = delete;
Trie &operator=(const Trie &) = delete;
int size() const { return num_terminals; }
bool empty() const { return num_terminals == 0; }

bool insert(const string &s, const V &v) {
    Node *n = root;
    for (char c : s) {
        auto [it, inserted] = n->children.try_emplace(c, nullptr);
        if (inserted) {
            it->second = new Node();
        }
        n = it->second;
    }
    if (n->is_terminal) {
        return false;
    }
    num_terminals++;
    n->is_terminal = true;
    n->value = v;
    // Re-walk from the root to bump the subtree count along the inserted path.
    n = root;
    n->cnt++;
    for (char c : s) {
        n = n->children[c];
        n->cnt++;
    }
    return true;
}

bool erase(const string &s) {
    if (erase(root, s, 0)) {
        num_terminals--;
        return true;
    }
    return false;
}

const V *find(const string &s) const {
    Node *n = root;
    for (char c : s) {
        if (auto it = n->children.find(c); it != n->children.end()) {
            n = it->second;
        } else {
            return nullptr;
        }
    }
    return n->is_terminal ? &(n->value) : nullptr;
}

```

```

}

int count_prefix(const string &s) const {
    Node *n = root;
    for (char c : s) {
        auto it = n->children.find(c);
        if (it == n->children.end()) {
            return 0;
        }
        n = it->second;
    }
    return n->cnt;
}

template<typename Fn>
void walk(Fn f) const {
    string s;
    walk(root, s, f);
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<string> s{"", "a", "to", "tea", "ted", "ten", "i", "in", "inn"};
    Trie<int> t;
    assert(t.empty());
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        assert(t.insert(s[i], i));
    }
    vector<pair<string, int>> entries;
    t.walk([&](string key, int value) { entries.emplace_back(key, value); });
    assert(
        (entries == vector<pair<string, int>>{
            {"", 0},
            {"a", 1},
            {"i", 6},
            {"in", 7},
            {"inn", 8},
            {"tea", 3},
            {"ted", 4},
            {"ten", 5},
            {"to", 2}
        })
    );
    assert(!t.empty());
    assert(t.size() == 9);
    assert(!t.insert(s[0], 2));
    assert(t.size() == 9);
    assert(*t.find("") == 0);
    assert(*t.find("ten") == 5);
    assert(t.count_prefix("") == 9); // Every key has the empty prefix.
    assert(t.count_prefix("te") == 3); // "tea", "ted", "ten".
    assert(t.count_prefix("in") == 2); // "in", "inn".
    assert(t.count_prefix("z") == 0);
    assert(t.erase("tea"));
    assert(t.size() == 8);
    assert(t.find("tea") == nullptr);
    assert(t.count_prefix("te") == 2); // "ted", "ten" remain.
    assert(t.erase(""));
    assert(t.find("") == nullptr);
    assert(t.count_prefix("") == 7);
}

```

```

// Erasing a key must not delete a shorter key sharing its path.
Trie<int> u;
u.insert("bc", 1);
u.insert("bca", 2);
assert(u.erase("bca"));
assert(u.find("bc") != nullptr && *u.find("bc") == 1);
assert(u.size() == 1 && u.count_prefix("bc") == 1);
return 0;
}

```

3.5.2 Radix Tree

Maintain a map of strings to values using a compressed tree data structure. Each node corresponds to a substring of an inserted string, and each inserted string corresponds to a path from the root to a node that is flagged as a terminal node. Contrary to a regular trie, a radix tree is more space efficient as it combines chains of nodes with only a single child. Children are stored in hash tables, and `walk()` sorts child labels as needed to preserve lexicographic traversal.

- `RadixTree<V>()` constructs an empty map.
- `size()` returns the size of the map.
- `empty()` returns whether the map is empty.
- `insert(s, v)` adds an entry with string key `s` and value `v` to the map, returning `true` if a new entry was added or `false` if the string already exists (in which case the map is unchanged and the old value associated with the string key is preserved).
- `erase(s)` removes the entry with string key `s` from the map, returning `true` if the removal was successful or `false` if the string to be removed was not found.
- `find(s)` returns a pointer to a const value associated with string key `s`, or `nullptr` if the key was not found.
- `count_prefix(s)` returns the number of keys currently in the map that have `s` as a prefix (a key equal to `s` counts as well), including the case where `s` ends partway along a compressed edge. `count_prefix("")` therefore equals `size()`.
- `walk(f)` calls the function `f(s, v)` on each entry of the map, in lexicographically ascending order of the string keys.

Time Complexity:

- $O(n)$ expected per call to `insert(s, v)`, `erase(s)`, `find(s)`, and `count_prefix(s)`, where n is the length of `s`, assuming the number of outgoing compressed edges checked at each node is small.
- $O(l \log l)$ per call to `walk()`, where l is the total length of string keys that are currently in the map.
- $O(1)$ per call to all other operations.

Space Complexity:

- $O(l)$ for storage of the radix tree, where l is the total length of string keys that are currently in the map.
- $O(n)$ auxiliary stack space for `insert()`, destruction, and `walk()`, where n is the maximum length of any string that has been inserted so far.
- $O(n)$ auxiliary stack space for `erase(s)`, where n is the length of `s`.
- $O(1)$ auxiliary for all other operations.

```

#include <algorithm>
#include <string>
#include <unordered_map>
#include <utility>
#include <vector>
using std::string;

template<typename V>
class RadixTree {
    struct Node {
        V value;

```

```

bool is_terminal;
int cnt; // Optional: maintain count of terminal keys in this subtree for count_prefix queries.
std::unordered_map<string, Node *> children;

Node(const V &value = V(), bool is_terminal = false)
    : value(value), is_terminal(is_terminal), cnt(0) {}
} *root;

static int lcp_len(const string &s1, const string &s2, int s2start) {
    int i = 0;
    for (int j = s2start; i < static_cast<int>(s1.size()) && j < static_cast<int>(s2.size());
        i++, j++) {
        if (s1[i] != s2[j]) {
            break;
        }
    }
    return i;
}

static bool insert(Node *n, const string &s, int i, const V &v) {
    if (i == static_cast<int>(s.size())) {
        if (n->is_terminal) {
            return false;
        }
        n->value = v;
        n->is_terminal = true;
        n->cnt++;
        return true;
    }
    for (auto it = n->children.begin(); it != n->children.end(); ++it) {
        int len = lcp_len(it->first, s, i);
        if (len == 0) {
            continue;
        }
        if (len == static_cast<int>(it->first.size())) {
            if (!insert(it->second, s, i + len, v)) {
                return false;
            }
            n->cnt++;
            return true;
        }
        string left = it->first.substr(0, len);
        string right = it->first.substr(len);
        Node *tmp = new Node();
        tmp->cnt = it->second->cnt; // The new split node heads the same subtree as the old child.
        tmp->children[right] = it->second;
        n->children.erase(it);
        n->children[left] = tmp;
        if (len == static_cast<int>(s.size()) - i) {
            tmp->value = v;
            tmp->is_terminal = true;
            tmp->cnt++;
        } else {
            insert(tmp, s, i + len, v);
        }
        n->cnt++;
        return true;
    }
    Node *leaf = new Node(v, true);
    leaf->cnt = 1;
    n->children[s.substr(i)] = leaf;
    n->cnt++;
    return true;
}

static bool erase(Node *n, const string &s, int i) {
    if (i == static_cast<int>(s.size())) {
        if (!n->is_terminal) {

```

```

        return false;
    }
    n->is_terminal = false;
    n->cnt--;
    return true;
}
for (auto it = n->children.begin(); it != n->children.end(); ++it) {
    int len = lcp_len(it->first, s, i);
    if (len == 0) {
        continue;
    }
    // The entire edge label must be consumed; a partial match means the key is not present, so
    // erasing must fail rather than descend (otherwise erase("te") would remove "tea").
    if (len < static_cast<int>(it->first.size())) {
        return false;
    }
    Node *child = it->second;
    if (!erase(child, s, i + len)) {
        return false;
    }
    n->cnt--;
    // The merge below leaves counts intact: a non-terminal node with one child already has the
    // same subtree count as that child, so reusing the node keeps its cnt correct.
    if (child->children.empty() && !child->is_terminal) {
        delete child;
        n->children.erase(it);
    } else if (child->children.size() == 1) {
        Node *grandchild = child->children.begin()->second;
        if (!child->is_terminal) {
            string merged_key(it->first + child->children.begin()->first);
            child->value = grandchild->value;
            child->is_terminal = grandchild->is_terminal;
            child->children = grandchild->children;
            delete grandchild;
            n->children.erase(it);
            n->children[merged_key] = child;
        }
    }
    return true;
}
return false;
}

template<typename Fn>
static void walk(Node *n, string &s, Fn f) {
    if (n->is_terminal) {
        f(s, n->value);
    }
    std::vector<std::pair<string, Node *>> children(n->children.begin(), n->children.end());
    std::sort(children.begin(), children.end());
    for (auto &[key, child] : children) {
        s += key;
        walk(child, s, f);
        s.erase(s.size() - key.size());
    }
}

static void clean_up(Node *n) {
    for (auto &[key, child] : n->children) {
        clean_up(child);
    }
    delete n;
}

int num_terminals;

public:
RadixTree() : root(new Node()), num_terminals(0) {}

```

```

~RadixTree() { clean_up(root); }
RadixTree(const RadixTree &) = delete;
RadixTree &operator=(const RadixTree &) = delete;
int size() const { return num_terminals; }
bool empty() const { return num_terminals == 0; }

bool insert(const string &s, const V &v) {
    if (insert(root, s, 0, v)) {
        num_terminals++;
        return true;
    }
    return false;
}

bool erase(const string &s) {
    if (erase(root, s, 0)) {
        num_terminals--;
        return true;
    }
    return false;
}

const V *find(const string &s) const {
    Node *n = root;
    int i = 0;
    while (i < static_cast<int>(s.size())) {
        bool found = false;
        for (auto &[key, child] : n->children) {
            if (key[0] == s[i]) {
                // The entire edge label must be consumed; a partial match (s ends mid-edge or diverges
                // from it) means the key is not present.
                if (lcp_len(key, s, i) < static_cast<int>(key.size())) {
                    return nullptr;
                }
                i += static_cast<int>(key.size());
                n = child;
                found = true;
                break;
            }
        }
        if (!found) {
            return nullptr;
        }
    }
    return n->is_terminal ? &(n->value) : nullptr;
}

int count_prefix(const string &s) const {
    Node *n = root;
    int i = 0;
    while (i < static_cast<int>(s.size())) {
        bool found = false;
        for (auto &[key, child] : n->children) {
            if (key[0] == s[i]) {
                int len = lcp_len(key, s, i);
                if (i + len == static_cast<int>(s.size())) {
                    // s ends within (or exactly at the end of) this edge: every key below qualifies.
                    return child->cnt;
                }
                if (len < static_cast<int>(key.size())) {
                    return 0; // s diverges partway along the edge.
                }
                i += len;
                n = child;
                found = true;
                break;
            }
        }
    }
}

```

```

    }
    if (!found) {
        return 0;
    }
}
return n->cnt;
}

template<typename Fn>
void walk(Fn f) const {
    string s;
    walk(root, s, f);
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<string> s{"", "a", "to", "tea", "ted", "ten", "i", "in", "inn"};
    RadixTree<int> t;
    assert(t.empty());
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        assert(t.insert(s[i], i));
    }
    vector<pair<string, int>> entries;
    t.walk([&](string key, int value) { entries.emplace_back(key, value); });
    assert(
        (entries == vector<pair<string, int>>{
            {"", 0},
            {"a", 1},
            {"i", 6},
            {"in", 7},
            {"inn", 8},
            {"tea", 3},
            {"ted", 4},
            {"ten", 5},
            {"to", 2}
        })
    );
    assert(!t.empty());
    assert(t.size() == 9);
    assert(!t.insert(s[0], 2));
    assert(t.size() == 9);
    assert(t.find("") && *t.find("") == 0);
    assert(*t.find("ten") == 5);
    assert(t.count_prefix("") == 9);    // Every key has the empty prefix.
    assert(t.count_prefix("te") == 3); // "tea", "ted", "ten" (prefix ends mid-edge).
    assert(t.count_prefix("in") == 2); // "in", "inn".
    assert(t.count_prefix("z") == 0);
    assert(t.erase("tea"));
    assert(t.size() == 8);
    assert(t.find("tea") == nullptr);
    assert(t.count_prefix("te") == 2); // "ted", "ten" remain.
    assert(t.erase(""));
    assert(t.find("") == nullptr);
    assert(t.count_prefix("") == 7);

    // Erasing a key must not delete a shorter key sharing its path.
    RadixTree<int> u;
    u.insert("bc", 1);
    u.insert("bca", 2);
    assert(u.erase("bca"));
    assert(u.find("bc") != nullptr && *u.find("bc") == 1);
}

```

```

    assert(u.size() == 1 && u.count_prefix("bc") == 1);
    return 0;
}

```

3.5.3 Suffix Tree (Ukkonen)

A suffix tree is the compressed trie of every suffix of a string: a rooted tree in which each edge is labeled with a substring, no two edges leaving a node start with the same character, and the concatenation of labels from the root to each leaf spells one suffix. Paths from the root therefore spell exactly the distinct substrings of the string, which makes the tree a powerful index for substring queries, repeated-substring problems, and many other tasks. It is the explicit-pointer counterpart of the suffix automaton; the two store the same information and either can be converted into the other.

This implementation is Ukkonen's algorithm, which builds the tree online in linear time by feeding one character at a time. It maintains an "active point" (a node, an edge, and a length along that edge) describing where the next suffix must be inserted, together with suffix links that let it hop between successive insertion points in amortized constant time. Edges to leaves are stored with an open-ended length that grows automatically as the string lengthens, so each phase touches only the suffixes that genuinely change. A unique terminal sentinel smaller than every real character is appended internally so that no suffix is a prefix of another and every suffix ends at its own leaf. The input must therefore not contain the character `'\1'`.

- `SuffixTree(s)` builds the suffix tree of string `s`.
- `contains(t)` returns whether string `t` occurs as a substring of `s`.
- `count_distinct_substrings()` returns the number of distinct nonempty substrings of `s`.
- `longest_repeated_substring()` returns one longest substring of `s` that occurs at least twice, or the empty string if every character is distinct.

Time Complexity:

- $O(n)$ per call to the constructor for a string of length n (using a per-node ordered map, $O(n \log s)$ where s is the alphabet size).
- $O(m \log s)$ per call to `contains(t)`, where m is the length of `t`.
- $O(n)$ per call to `count_distinct_substrings()` and `longest_repeated_substring()`.

Space Complexity:

- $O(n)$ nodes and edges.
- $O(n)$ auxiliary stack space for the traversals.

```

#include <stdint>
#include <map>
#include <string>
#include <vector>

class SuffixTree {
    static const int OPEN = -1; // Marks a leaf edge whose right end follows the growing string.

    struct Node {
        int start, end; // Edge into this node labels str[start..end]; end == OPEN means leaf_end.
        int link;
        std::map<char, int> next;
        Node(int start, int end) : start(start), end(end), link(0) {}
    };

    std::string str;
    std::vector<Node> tree;
    int leaf_end;
    // Active point and bookkeeping for Ukkonen's algorithm.
    int active_node, active_edge, active_length, remaining, last_new;

    int edge_end(int v) const { return tree[v].end == OPEN ? leaf_end : tree[v].end; }
}

```



```

int edge_len(int v) const { return edge_end(v) - tree[v].start + 1; }

int make_node(int start, int end) {
    tree.emplace_back(start, end);
    return static_cast<int>(tree.size()) - 1;
}

// If the active length runs past the edge into v, descend and report that we moved.
bool walk_down(int v) {
    if (active_length >= edge_len(v)) {
        active_edge += edge_len(v);
        active_length -= edge_len(v);
        active_node = v;
        return true;
    }
    return false;
}

void extend(int pos) {
    leaf_end = pos;
    remaining++;
    last_new = -1;
    while (remaining > 0) {
        if (active_length == 0) {
            active_edge = pos;
        }
        char c = str[active_edge];
        auto it = tree[active_node].next.find(c);
        if (it == tree[active_node].next.end()) {
            // Rule 2: no edge starts with c, so grow a new leaf from the active node.
            tree[active_node].next[c] = make_node(pos, OPEN);
            if (last_new != -1) {
                tree[last_new].link = active_node;
                last_new = -1;
            }
        } else {
            int v = it->second;
            if (walk_down(v)) {
                continue; // Re-evaluate the same suffix from the deeper node.
            }
            if (str[tree[v].start + active_length] == str[pos]) {
                // Rule 3: the character is already present; stop and lengthen the active point.
                if (last_new != -1) {
                    tree[last_new].link = active_node;
                }
                active_length++;
                break;
            }
            // Rule 2 with a split: break edge v in two and attach a new leaf at the split point.
            int split = make_node(tree[v].start, tree[v].start + active_length - 1);
            tree[active_node].next[c] = split;
            tree[split].next[str[pos]] = make_node(pos, OPEN);
            tree[v].start += active_length;
            tree[split].next[str[tree[v].start]] = v;
            if (last_new != -1) {
                tree[last_new].link = split;
            }
            last_new = split;
        }
        remaining--;
        if (active_node == 0 && active_length > 0) {
            active_length--;
            active_edge = pos - remaining + 1;
        } else if (active_node != 0) {
            active_node = tree[active_node].link;
        }
    }
}

```

```

// String depth and right end (one past) of the deepest internal node, for repeated substrings.
void deepest_internal(int v, int depth, int &best_depth, int &best_end) const {
    bool internal = false;
    for (const auto &kv : tree[v].next) {
        internal = true;
        int w = kv.second;
        deepest_internal(w, depth + edge_len(w), best_depth, best_end);
    }
    if (internal && v != 0 && depth > best_depth) {
        best_depth = depth;
        best_end = edge_end(v) + 1;
    }
}

public:
explicit SuffixTree(const std::string &s) : str(s) {
    str.push_back('\1'); // Unique terminal sentinel, smaller than every real character.
    make_node(-1, -1); // Root, at index 0; its incoming edge is never read.
    active_node = 0;
    active_edge = 0;
    active_length = 0;
    remaining = 0;
    leaf_end = -1;
    last_new = -1;
    for (int i = 0; i < static_cast<int>(str.size()); i++) {
        extend(i);
    }
}

bool contains(const std::string &t) const {
    int v = 0, i = 0, m = static_cast<int>(t.size());
    while (i < m) {
        auto it = tree[v].next.find(t[i]);
        if (it == tree[v].next.end()) {
            return false;
        }
        v = it->second;
        int len = edge_len(v);
        for (int j = 0; j < len && i < m; j++, i++) {
            if (str[tree[v].start + j] != t[i]) {
                return false;
            }
        }
    }
    return true;
}

int64_t count_distinct_substrings() const {
    int64_t total = 0;
    for (int v = 1; v < static_cast<int>(tree.size()); v++) {
        total += edge_len(v);
    }
    // Every leaf edge ends in the sentinel. The str.size() suffixes of the sentineled string are
    // the only substrings containing it, so removing them leaves the original string's count.
    return total - static_cast<int64_t>(str.size());
}

std::string longest_repeated_substring() const {
    int best_depth = 0, best_end = 0;
    deepest_internal(0, 0, best_depth, best_end);
    return str.substr(best_end - best_depth, best_depth);
}
};

```

Example Usage

```
#include <cassert>
#include <set>
using namespace std;

int main() {
    SuffixTree st("banana");
    assert(st.contains("ana"));
    assert(st.contains("banana"));
    assert(st.contains(""));
    assert(!st.contains("anana_"));
    assert(!st.contains("\nb"));
    assert(st.longest_repeated_substring() == "ana");

    // "banana" has 15 distinct substrings:
    // a, an, ana, anan, anana, b, ba, ban, bana, banan, banana, n, na, nan, nana.
    assert(st.count_distinct_substrings() == 15);

    // Cross-check both queries against brute force on a small string.
    string text = "mississippi";
    SuffixTree t(text);
    set<string> subs;
    int n = static_cast<int>(text.size());
    for (int i = 0; i < n; i++) {
        for (int len = 1; i + len <= n; len++) {
            subs.insert(text.substr(i, len));
        }
    }
    assert(t.count_distinct_substrings() == static_cast<int>(subs.size()));
    for (const string &sub : subs) {
        assert(t.contains(sub));
    }
    assert(!t.contains("issa"));
    assert(t.longest_repeated_substring() == "issi"); // "issi" appears at indices 1 and 4.
    return 0;
}
```

3.5.4 Suffix Automaton

Maintains a compact deterministic automaton representing all substrings of a string. A suffix automaton is useful for substring membership queries, counting distinct substrings, finding longest common substrings, and many other string dynamic programming tasks.

Each state represents an equivalence class of substrings with the same set of ending positions. The value `st[v].len` is the maximum length represented by state `v`, while `st[v].link` points to the state representing the longest proper suffix of that class.

More precisely, state `v` contributes one distinct substring for each length greater than `st[st[v].link].len` and at most `st[v].len`. When a character is appended, `extend()` creates a state for the new full string and walks suffix links from the previous `last`, adding missing transitions to it. If an existing transition leads to a state whose maximum length is too large, a shorter clone of that state is inserted and the affected transitions and suffix links are redirected to it, splitting the ending-position classes without changing the recognized substrings. Summing these length ranges gives the number of distinct substrings. Membership queries follow transitions, while longest common substring queries retreat through suffix links when a match fails.

- `SuffixAutomaton()` constructs an empty automaton.
- `SuffixAutomaton(s)` constructs the automaton for string `s`.
- `extend(c)` appends character `c` to the current string.
- `contains(t)` returns whether string `t` occurs as a substring.
- `first_occurrence(t)` returns the starting index of the first occurrence of `t`, or `-1` if absent.

- `count_distinct_substrings()` returns the number of distinct nonempty substrings.
- `longest_common_substring(t)` returns one longest substring common to the built string and string `t`.

Time Complexity:

- $O(n)$ expected per call to the constructor for a string of length n .
- $O(m)$ expected per call to `contains(t)`, `first_occurrence(t)`, or `longest_common_substring(t)`, where m is the length of `t`.
- $O(n)$ per call to `count_distinct_substrings()`.

Space Complexity:

- $O(n)$ transition storage and $O(n)$ states.
- $O(1)$ auxiliary per character processed.

```
#include <cstdint>
#include <string>
#include <unordered_map>
#include <vector>
using std::string;

class SuffixAutomaton {
    struct State {
        int len, link, first_pos;
        std::unordered_map<char, int> next;

        State() : len(0), link(-1), first_pos(-1) {}
    };

    std::vector<State> st;
    int last;

public:
    SuffixAutomaton() : st(1), last(0) {}

    explicit SuffixAutomaton(const string &s) : st(1), last(0) {
        for (char c : s) {
            extend(c);
        }
    }

    void extend(char c) {
        int cur = static_cast<int>(st.size());
        st.emplace_back();
        st[cur].len = st[last].len + 1;
        st[cur].first_pos = st[cur].len - 1;
        int p = last;
        while (p != -1 && st[p].next.find(c) == st[p].next.end()) {
            st[p].next[c] = cur;
            p = st[p].link;
        }
        if (p == -1) {
            st[cur].link = 0;
        } else {
            int q = st[p].next[c];
            if (st[p].len + 1 == st[q].len) {
                st[cur].link = q;
            } else {
                int clone = static_cast<int>(st.size());
                st.push_back(st[q]);
                st[clone].len = st[p].len + 1;
                while (p != -1 && st[p].next[c] == q) {
                    st[p].next[c] = clone;
                    p = st[p].link;
                }
                st[q].link = st[cur].link = clone;
            }
        }
    }
};
```

```

    }
    last = cur;
}

bool contains(const string &t) const {
    int v = 0;
    for (char c : t) {
        if (auto it = st[v].next.find(c); it != st[v].next.end()) {
            v = it->second;
        } else {
            return false;
        }
    }
    return true;
}

// Returns the starting index of the first (leftmost) occurrence of t as a substring, or -1 if t
// does not occur. first_pos stores the end index of the first occurrence of each state's class,
// so subtracting the length of t recovers its start.
int first_occurrence(const string &t) const {
    int v = 0;
    for (char c : t) {
        auto it = st[v].next.find(c);
        if (it == st[v].next.end()) {
            return -1;
        }
        v = it->second;
    }
    return st[v].first_pos - static_cast<int>(t.size()) + 1;
}

int64_t count_distinct_substrings() const {
    int64_t res = 0;
    for (int v = 1; v < static_cast<int>(st.size()); v++) {
        res += st[v].len - st[st[v].link].len;
    }
    return res;
}

string longest_common_substring(const string &t) const {
    int v = 0, len = 0, best = 0, best_pos = -1;
    for (int i = 0; i < static_cast<int>(t.size()); i++) {
        while (v && !st[v].next.count(t[i])) {
            v = st[v].link;
            len = st[v].len;
        }
        if (auto it = st[v].next.find(t[i]); it != st[v].next.end()) {
            v = it->second;
            len++;
        } else {
            v = 0;
            len = 0;
        }
        if (len > best) {
            best = len;
            best_pos = i;
        }
    }
    return best == 0 ? "" : t.substr(best_pos - best + 1, best);
}
};

```

Example Usage

```
#include <cassert>
```

```

int main() {
    SuffixAutomaton sa("ababa");
    assert(sa.contains("aba"));
    assert(sa.contains("bab"));
    assert(!sa.contains("abba"));
    // "ababa" has 9 distinct nonempty substrings.
    assert(sa.count_distinct_substrings() == 9);
    // The longest substring shared by "ababa" and "zzbabab" is "baba".
    assert(sa.longest_common_substring("zzbabab") == "baba");
    assert(sa.first_occurrence("aba") == 0);
    assert(sa.first_occurrence("bab") == 1);
    assert(sa.first_occurrence("abba") == -1);

    SuffixAutomaton online;
    online.extend('a');
    online.extend('b');
    online.extend('c');
    assert(online.contains("bc"));
    return 0;
}

```

3.5.5 Eertree

Maintains all distinct palindromic substrings of a string online using an Eertree, also known as a palindromic tree. Each non-root node represents one distinct palindrome, and suffix links connect each palindrome to its longest proper palindromic suffix.

The structure has two roots: one of length -1 , which simplifies boundary cases, and one of length 0 , representing the empty palindrome. After each appended character, `last` points to the node for the longest palindromic suffix of the current string.

To append a character, `add()` follows suffix links from `last` until finding the longest suffix palindrome that can be enclosed by the new character. If the corresponding transition already exists, that node becomes `last`; otherwise, a node two characters longer is created and linked to its longest proper palindromic suffix, found by the same search. Exactly one distinct palindrome can be created at each position, so the number of non-root nodes is the distinct count. Each node first counts the times it is the longest palindromic suffix; propagating those counts through suffix links then includes occurrences where it appears as a suffix of a longer palindrome.

- `Eertree()` constructs an empty palindromic tree.
- `Eertree(s)` constructs the tree for string `s`.
- `add(c)` appends character `c` and returns `true` if this creates a new distinct palindrome.
- `count_distinct_palindromes()` returns the number of distinct nonempty palindromic substrings.
- `longest_suffix_length()` returns the length of the current longest palindromic suffix.
- `count_occurrences()` propagates occurrence counts through suffix links and returns `occ[v]` for each node `v`.

Time Complexity:

- $O(n)$ expected per call to the constructor for a string of length n .
- $O(1)$ expected amortized per call to `add()`.
- $O(n)$ per call to `count_occurrences()`.

Space Complexity:

- $O(n)$ transition storage and $O(n)$ nodes.
- $O(1)$ auxiliary per appended character.

```

#include <string>
#include <unordered_map>
#include <vector>
using std::string;

```

```

class Eertree {
    struct Node {
        int len, link, occ;
        std::unordered_map<char, int> next;

        Node(int len = 0) : len(len), link(0), occ(0) {}
    };

    std::vector<Node> tree;
    string str;
    int last;

    int get_suffix(int v, int pos, char c) const {
        while (pos - 1 - tree[v].len < 0 || str[pos - 1 - tree[v].len] != c) {
            v = tree[v].link;
        }
        return v;
    }

public:
    Eertree() : tree(), str(), last(1) {
        tree.emplace_back(-1);
        tree.emplace_back(0);
        tree[0].link = 0;
        tree[1].link = 0;
    }

    explicit Eertree(const string &s) : tree(), str(), last(1) {
        tree.emplace_back(-1);
        tree.emplace_back(0);
        tree[0].link = 0;
        tree[1].link = 0;
        for (char c : s) {
            add(c);
        }
    }

    bool add(char c) {
        str += c;
        int pos = static_cast<int>(str.size()) - 1;
        int cur = get_suffix(last, pos, c);
        auto it = tree[cur].next.find(c);
        if (it != tree[cur].next.end()) {
            last = it->second;
            tree[last].occ++;
            return false;
        }
        last = static_cast<int>(tree.size());
        tree.emplace_back(tree[cur].len + 2);
        tree[last].occ = 1;
        tree[cur].next[c] = last;
        if (tree[last].len == 1) {
            tree[last].link = 1;
            return true;
        }
        int link_parent = get_suffix(tree[cur].link, pos, c);
        tree[last].link = tree[link_parent].next[c];
        return true;
    }

    int count_distinct_palindromes() const { return static_cast<int>(tree.size()) - 2; }
    int longest_suffix_length() const { return tree[last].len; }

    std::vector<int> count_occurrences() const {
        std::vector<int> occ(tree.size());
        for (int i = 0; i < static_cast<int>(tree.size()); i++) {
            occ[i] = tree[i].occ;
        }
    }
}

```

```

    for (int i = static_cast<int>(tree.size()) - 1; i >= 2; i--) {
        occ[tree[i].link] += occ[i];
    }
    return occ;
}
};

```

Example Usage

```

#include <cassert>

int main() {
    Eertree t("abacaba");
    // Distinct palindromes: a, b, c, aba, aca, bacab, abacaba.
    assert(t.count_distinct_palindromes() == 7);
    assert(t.longest_suffix_length() == 7);

    Eertree online;
    assert(online.add('a'));
    assert(online.add('a'));
    assert(online.add('b'));
    assert(online.add('a'));
    assert(online.count_distinct_palindromes() == 4);
    assert(online.longest_suffix_length() == 3);

    Eertree duplicate;
    assert(duplicate.add('a'));
    assert(duplicate.add('b'));
    assert(duplicate.add('c'));
    // The second 'a' creates substring "abca", but palindrome "a" already exists.
    assert(!duplicate.add('a'));

    std::vector<int> occ = t.count_occurrences();
    assert(static_cast<int>(occ.size()) == t.count_distinct_palindromes() + 2);
    return 0;
}

```

3.6 Sequence Dynamic Programming

3.6.1 Longest Common Substring

Given two strings, determine their longest common substring: a contiguous run of characters that appears in both. A dynamic program computes, for each pair of prefixes, the length of their longest common suffix; the largest value over all pairs locates a longest common substring. Only the previous row of the table is retained, so the auxiliary space is linear in the shorter string.

- `longest_common_substring(s1, s2)` returns a longest common substring of `s1` and `s2` (one such substring if several are tied in length), or the empty string if they have none in common.

Time Complexity: $O(n \cdot m)$ per call, where n and m are the lengths of strings `s1` and `s2`, respectively.

Space Complexity: $O(\min(n, m))$ auxiliary, where n and m are the lengths of `s1` and `s2`, respectively.

```

#include <string>
#include <vector>
using std::string;

string longest_common_substring(const string &s1, const string &s2) {
    int n = static_cast<int>(s1.size()), m = static_cast<int>(s2.size());
    if (n == 0 || m == 0) {

```



```

    return "";
}
if (n < m) {
    return longest_common_substring(s2, s1);
}
std::vector<int> curr(m), prev(m);
int pos = 0, len = 0;
for (int i = 0; i < n; i++) {
    for (int j = 0; j < m; j++) {
        if (s1[i] == s2[j]) {
            curr[j] = (i > 0 && j > 0) ? prev[j - 1] + 1 : 1;
            if (len < curr[j]) {
                len = curr[j];
                pos = i - curr[j] + 1;
            }
        } else {
            curr[j] = 0;
        }
    }
    curr.swap(prev);
}
return s1.substr(pos, len);
}

```

Example Usage

```

#include <cassert>

int main() {
    assert(longest_common_substring("bbbabca", "aababcd") == "babc");
    assert(longest_common_substring("abc", "def") == "");
    assert(longest_common_substring("", "abc") == "");
    assert(longest_common_substring("xyabc", "zzabczz") == "abc");
    return 0;
}

```

3.6.2 Longest Common Subsequence

Given two strings, determine their longest common subsequence. A subsequence is a string that can be derived from the original string by deleting some elements without changing the order of the remaining elements (e.g. "ACE" is a subsequence of "ABCDE", but "BAE" is not).

The dynamic program compares prefixes: if the last characters match, they can extend an optimal subsequence of the smaller prefixes; otherwise, one of the two last characters must be skipped. Following the filled table backward reconstructs one optimal subsequence. Hirschberg's algorithm keeps only one DP row at a time, using forward and backward length rows to choose where an optimal answer crosses the midpoint of the longer string.

- `longest_common_subsequence(s1, s2)` returns the longest common subsequence of strings `s1` and `s2` using a classic dynamic programming approach. This implementation computes `dp[i][j]` (the length of the longest common subsequence for the length i prefix of `s1` and the length j prefix of `s2`) before following the path backwards to construct the answer.
- `hirschberg_lcs(s1, s2)` returns the longest common subsequence of strings `s1` and `s2` using the more memory efficient Hirschberg's algorithm.

Time Complexity: $O(n \cdot m)$ per call to `longest_common_subsequence(s1, s2)` as well as `hirschberg_lcs(s1, s2)`, where n and m are the lengths of `s1` and `s2`, respectively.

Space Complexity:

- $O(n \cdot m)$ auxiliary heap space for `longest_common_subsequence(s1, s2)`, where n and m are the lengths of `s1` and `s2`, respectively.

- $O(\log(\max(n, m)))$ auxiliary stack space and $O(\min(n, m))$ auxiliary heap space for `hirschberg_lcs(s1, s2)`, where n and m are the lengths of `s1` and `s2`, respectively.

```

#include <algorithm>
#include <iterator>
#include <string>
#include <vector>
using std::string;

string longest_common_subsequence(const string &s1, const string &s2) {
    int n = static_cast<int>(s1.size()), m = static_cast<int>(s2.size());
    std::vector<std::vector<int>>> dp(n + 1, std::vector<int>(m + 1, 0));
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= m; j++) {
            if (s1[i - 1] == s2[j - 1]) {
                dp[i][j] = dp[i - 1][j - 1] + 1;
            } else {
                dp[i][j] = std::max(dp[i][j - 1], dp[i - 1][j]);
            }
        }
    }
    string res;
    for (int i = n, j = m; i > 0 && j > 0;) {
        if (s1[i - 1] == s2[j - 1]) {
            res += s1[i - 1];
            i--;
            j--;
        } else if (dp[i - 1][j] >= dp[i][j - 1]) {
            i--;
        } else {
            j--;
        }
    }
    std::reverse(res.begin(), res.end());
    return res;
}

template<typename It>
std::vector<int> lcs_len(It lo1, It hi1, It lo2, It hi2) {
    std::vector<int> res(std::distance(lo2, hi2) + 1, prev(res));
    for (It it1 = lo1; it1 != hi1; ++it1) {
        res.swap(prev);
        int i = 0;
        for (It it2 = lo2; it2 != hi2; ++it2) {
            res[i + 1] = (*it1 == *it2) ? prev[i] + 1 : std::max(res[i], prev[i + 1]);
            i++;
        }
    }
    return res;
}

template<typename It>
void hirschberg_rec(It lo1, It hi1, It lo2, It hi2, string *res) {
    if (lo1 == hi1) {
        return;
    }
    if (lo1 + 1 == hi1) {
        if (std::find(lo2, hi2, *lo1) != hi2) {
            *res += *lo1;
        }
        return;
    }
    It mid1 = lo1 + (hi1 - lo1) / 2;
    std::reverse_iterator<It> rlo1(hi1), rmid1(mid1), rlo2(hi2), rhi2(lo2);
    std::vector<int> fwd = lcs_len(lo1, mid1, lo2, hi2);
    std::vector<int> rev = lcs_len(rlo1, rmid1, rlo2, rhi2);
    It mid2 = lo2;

```

```

int maxlen = -1;
for (int i = 0, j = static_cast<int>(rev.size()) - 1; i < static_cast<int>(fwd.size());
    i++, j--) {
    if (fwd[i] + rev[j] > maxlen) {
        maxlen = fwd[i] + rev[j];
        mid2 = lo2 + i;
    }
}
hirschberg_rec(lo1, mid1, lo2, mid2, res);
hirschberg_rec(mid1, hi1, mid2, hi2, res);
}

string hirschberg_lcs(const string &s1, const string &s2) {
    if (s1.size() < s2.size()) {
        return hirschberg_lcs(s2, s1);
    }
    string res;
    hirschberg_rec(s1.begin(), s1.end(), s2.begin(), s2.end(), &res);
    return res;
}

```

Example Usage

```

#include <cassert>

int main() {
    assert(longest_common_subsequence("xmjyauz", "mzjawxu") == "mjau");
    assert(hirschberg_lcs("xmjyauz", "mzjawxu") == "mjau");
    assert(longest_common_subsequence("abc", "def") == "");
    assert(hirschberg_lcs("abc", "def") == "");
    assert(longest_common_subsequence("", "abc") == "");
    assert(hirschberg_lcs("", "abc") == "");
    return 0;
}

```

3.6.3 Edit Distance and Sequence Alignment

Given two strings, determine their edit distance or a minimum-cost alignment. The Levenshtein edit distance is the minimum number of insertions, deletions, and substitutions needed to transform one string into the other. More generally, an alignment inserts gap characters `'_'` into both strings so that the final lengths match; its cost is `gap_cost` times the number of inserted gaps plus `sub_cost` times the number of indices where the aligned characters differ.

Dynamic programming compares prefixes and chooses the cheapest last operation: align the final characters together, leave a gap in the first string, or leave a gap in the second string. Following the filled table backward reconstructs one minimum-cost alignment. This is the Needleman-Wunsch sequence-alignment recurrence. Hirschberg's algorithm uses the same recurrence with only one DP row at a time, splitting the problem where an optimal alignment crosses the longer string's midpoint. Both `gap_cost` and `sub_cost` must be nonnegative.

- `edit_distance(s1, s2)` returns the Levenshtein edit distance between strings `s1` and `s2`.
- `align_sequences(s1, s2, gap_cost, sub_cost)` returns a pair of aligned strings for strings `s1` and `s2`, using a classic dynamic programming approach. This implementation first computes `dp[i][j]` (the cost of aligning the length `i` prefix of `s1` with the length `j` prefix of `s2`) before following the path backwards to construct the answer. For `gap_cost = sub_cost = 1`, `dp[n][m]` will be the Levenshtein edit distance, where `n` and `m` are the lengths of `s1` and `s2`, respectively.
- `hirschberg_align(s1, s2, gap_cost, sub_cost)` returns the sequence alignment of strings `s1` and `s2` using the more memory efficient Hirschberg's algorithm.

Time Complexity: $O(n \cdot m)$ per call to `edit_distance(s1, s2)`, `align_sequences(s1, s2)`, and `hirschberg_align(s1, s2)`, where n and m are the lengths of `s1` and `s2`, respectively.

Space Complexity:

- $O(\min(n, m))$ auxiliary heap space for `edit_distance()`.
- $O(n \cdot m)$ auxiliary heap space for `align_sequences(s1, s2)`, where n and m are the lengths of `s1` and `s2`, respectively.
- $O(\log(\max(n, m)))$ auxiliary stack space and $O(\min(n, m))$ auxiliary heap space for `hirschberg_align(s1, s2)`, where n and m are the lengths of `s1` and `s2`, respectively.

```
#include <algorithm>
#include <iterator>
#include <string>
#include <utility>
#include <vector>
using std::string;

int edit_distance(const string &s1, const string &s2) {
    if (s1.size() < s2.size()) {
        return edit_distance(s2, s1);
    }
    int n = static_cast<int>(s1.size()), m = static_cast<int>(s2.size());
    std::vector<int> dp(m + 1), prev(m + 1);
    for (int j = 0; j <= m; j++) {
        dp[j] = j;
    }
    for (int i = 1; i <= n; i++) {
        dp.swap(prev);
        dp[0] = i;
        for (int j = 1; j <= m; j++) {
            dp[j] =
                (s1[i - 1] == s2[j - 1]) ? prev[j - 1] : std::min({prev[j - 1], prev[j], dp[j - 1]}) + 1;
        }
    }
    return dp[m];
}

std::pair<string, string> align_sequences(
    const string &s1, const string &s2, int gap_cost = 1, int sub_cost = 1
) {
    int n = static_cast<int>(s1.size()), m = static_cast<int>(s2.size());
    std::vector<std::vector<int>> dp(n + 1, std::vector<int>(m + 1, 0));
    for (int i = 0; i <= n; i++) {
        dp[i][0] = i * gap_cost;
    }
    for (int j = 0; j <= m; j++) {
        dp[0][j] = j * gap_cost;
    }
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= m; j++) {
            dp[i][j] = (s1[i - 1] == s2[j - 1]) ? dp[i - 1][j - 1]
                : std::min(
                    dp[i - 1][j - 1] + sub_cost,
                    std::min(dp[i - 1][j], dp[i][j - 1]) + gap_cost
                );
        }
    }
    string res1, res2;
    int i = n, j = m;
    while (i > 0 && j > 0) {
        if (s1[i - 1] == s2[j - 1] || dp[i][j] == dp[i - 1][j - 1] + sub_cost) {
            res1 += s1[--i];
            res2 += s2[--j];
        } else if (dp[i][j] == dp[i - 1][j] + gap_cost) {
            res1 += s1[--i];
            res2 += '_';
        }
    }
}
```

```

    } else if (dp[i][j] == dp[i][j - 1] + gap_cost) {
        res1 += '_';
        res2 += s2[--j];
    }
}
while (i > 0 || j > 0) {
    res1 += (i > 0) ? s1[--i] : '_';
    res2 += (j > 0) ? s2[--j] : '_';
}
std::reverse(res1.begin(), res1.end());
std::reverse(res2.begin(), res2.end());
return {res1, res2};
}

template<typename It>
std::vector<int> row_cost(It lo1, It hi1, It lo2, It hi2, int gap_cost, int sub_cost) {
    std::vector<int> res(std::distance(lo2, hi2) + 1);
    for (int i = 0; i < static_cast<int>(res.size()); i++) {
        res[i] = i * gap_cost;
    }
    std::vector<int> prev(res);
    int row = 0;
    for (It it1 = lo1; it1 != hi1; ++it1) {
        res.swap(prev);
        res[0] = ++row * gap_cost;
        int i = 0;
        for (It it2 = lo2; it2 != hi2; ++it2) {
            res[i + 1] = (*it1 == *it2)
                ? prev[i]
                : std::min({prev[i] + sub_cost, prev[i + 1] + gap_cost, res[i] + gap_cost});
            i++;
        }
    }
    return res;
}

template<typename It>
void hirschberg_rec(
    It lo1, It hi1, It lo2, It hi2, string *res1, string *res2, int gap_cost, int sub_cost
) {
    if (lo1 == hi1) {
        for (It it2 = lo2; it2 != hi2; ++it2) {
            *res1 += '_';
            *res2 += *it2;
        }
        return;
    }
    if (lo1 + 1 == hi1) {
        It pos = std::find(lo2, hi2, *lo1);
        bool use_two_gaps = (pos == hi2) && (2 * gap_cost < sub_cost);
        if (lo2 == hi2 || use_two_gaps) {
            *res1 += *lo1;
            *res2 += '_';
        }
        for (It it2 = lo2; it2 != hi2; ++it2) {
            *res1 += (pos == it2 || (pos == hi2 && !use_two_gaps && it2 == lo2)) ? *lo1 : '_';
            *res2 += *it2;
        }
        return;
    }
    It mid1 = lo1 + (hi1 - lo1) / 2;
    std::reverse_iterator<It> rlo1(hi1), rmid1(mid1), rlo2(hi2), rhi2(lo2);
    std::vector<int> fwd = row_cost(lo1, mid1, lo2, hi2, gap_cost, sub_cost);
    std::vector<int> rev = row_cost(rlo1, rmid1, rlo2, rhi2, gap_cost, sub_cost);
    It mid2 = lo2;
    int mincost = -1;
    for (int i = 0, j = static_cast<int>(rev.size()) - 1; i < static_cast<int>(fwd.size());
        i++, j--) {

```

```

    if (mincost < 0 || fwd[i] + rev[j] < mincost) {
        mincost = fwd[i] + rev[j];
        mid2 = lo2 + i;
    }
}
hirschberg_rec(lo1, mid1, lo2, mid2, res1, res2, gap_cost, sub_cost);
hirschberg_rec(mid1, hi1, mid2, hi2, res1, res2, gap_cost, sub_cost);
}

std::pair<string, string> hirschberg_align(
    const string &s1, const string &s2, int gap_cost = 1, int sub_cost = 1
) {
    if (s1.size() < s2.size()) {
        auto res = hirschberg_align(s2, s1, gap_cost, sub_cost);
        return {res.second, res.first};
    }
    string res1, res2;
    hirschberg_rec(s1.begin(), s1.end(), s2.begin(), s2.end(), &res1, &res2, gap_cost, sub_cost);
    return {res1, res2};
}

```

Example Usage

```

#include <cassert>
#include <iostream>

int main() {
    assert(edit_distance("kitten", "sitting") == 3);
    assert(edit_distance("", "abc") == 3);

    auto alignment = align_sequences("AGGGCT", "AGGCA", 2, 3);
    assert(alignment == (std::pair<string, string>{"AGGGCT", "A_GGCA"}));
    assert(hirschberg_align("AGGGCT", "AGGCA", 2, 3) == alignment);

    assert((hirschberg_align("ab", "zabz") == std::pair<string, string>{"_ab_", "zabz"}));
    assert(
        (hirschberg_align("a", "bbbbbbbbbb", 1, 5) ==
         std::pair<string, string>{"_a", "bbbbbbbbbb_"}
    ));
    assert((hirschberg_align("aa", "ba", 1, 2) == std::pair<string, string>{"a_a", "_ba"}));

    std::cout << "Aligned strings:" << std::endl;
    std::cout << alignment.first << std::endl << alignment.second << std::endl;
    return 0;
}

```

Example Output

```

Aligned strings:
AGGGCT
A_GGCA

```

3.6.4 Palindrome Dynamic Programming

Solves basic palindrome dynamic programming problems on a string. The core interval state considers substring $[l, r]$: if the two endpoints match, they can wrap an optimal answer for the inside substring; otherwise, at least one endpoint must be skipped or inserted around the other side. Since each state depends on shorter substrings, intervals are processed by increasing length.

The longest palindromic subsequence (LPS) is a maximum-length subsequence that reads the same forward and backward. The minimum number of insertions needed to make a string a palindrome is the number of characters outside an LPS, so it equals n minus the LPS length.

- `longest_palindromic_subsequence(s)` returns one longest palindromic subsequence of string `s`.
- `min_insertions_palindrome(s)` returns the minimum number of characters that must be inserted anywhere in string `s` to make it a palindrome.

Time Complexity: $O(n^2)$ per call to `longest_palindromic_subsequence(s)` and `min_insertions_palindrome(s)`, where n is the length of `s`.

Space Complexity:

- $O(n^2)$ auxiliary for `longest_palindromic_subsequence()`.
- $O(n)$ auxiliary for `min_insertions_palindrome()`.

```
#include <algorithm>
#include <cassert>
#include <string>
#include <vector>
using std::string;

string longest_palindromic_subsequence(const string &s) {
    int n = static_cast<int>(s.size());
    if (n == 0) {
        return "";
    }
    std::vector<std::vector<int>> dp(n, std::vector<int>(n, 0));
    for (int len = 1; len <= n; len++) {
        for (int l = 0, r = len - 1; r < n; l++, r++) {
            if (l == r) {
                dp[l][r] = 1;
            } else if (s[l] == s[r]) {
                dp[l][r] = 2 + (l + 1 <= r - 1 ? dp[l + 1][r - 1] : 0);
            } else {
                dp[l][r] = std::max(dp[l + 1][r], dp[l][r - 1]);
            }
        }
    }
    string left, right;
    for (int l = 0, r = n - 1; l <= r;) {
        if (l == r) {
            left += s[l];
            break;
        }
        if (s[l] == s[r] && dp[l][r] == 2 + (l + 1 <= r - 1 ? dp[l + 1][r - 1] : 0)) {
            left += s[l++];
            right += s[r--];
        } else if (dp[l + 1][r] >= dp[l][r - 1]) {
            l++;
        } else {
            r--;
        }
    }
    std::reverse(right.begin(), right.end());
    return left + right;
}

int min_insertions_palindrome(const string &s) {
    int n = static_cast<int>(s.size());
    std::vector<int> dp(n, 0);
    for (int l = n - 1; l >= 0; l--) {
        int inside = 0;
        for (int r = l + 1; r < n; r++) {
            int old = dp[r];
            if (s[l] == s[r]) {
                dp[r] = inside;
            } else {
                dp[r] = 1 + std::min(dp[r], dp[r - 1]);
            }
        }
        inside = old;
    }
}
```

```

    }
}
return n == 0 ? 0 : dp[n - 1];
}

```

Example Usage

```

#include <cassert>

int main() {
    assert(longest_palindromic_subsequence("bbbab") == "bbbb");
    assert(longest_palindromic_subsequence("character") == "carac");
    assert(min_insertions_palindrome("abcda") == 2);
    assert(min_insertions_palindrome("race") == 3);
    assert(min_insertions_palindrome("") == 0);
    return 0;
}

```

3.7 Expression Parsing

3.7.1 Recursive Descent Parsing (Simple)

Evaluate an expression according to the order of operations (parentheses, unary plus and minus signs, multiplication/division, addition/subtraction). The following is a minimalistic recursive descent implementation using iterators. Each precedence level is parsed by its own function that obtains operands by calling the next-tighter level, so the call structure itself enforces the order of operations.

- `eval(s)` returns an evaluation of the arithmetic expression `s`.

Time Complexity: $O(n)$ per call, where n is the length of `s`.

Space Complexity: $O(n)$ auxiliary stack space, where n is the length of `s`.

```

#include <string>

template<typename It>
int eval(It &it, int prec) {
    if (prec == 0) {
        int sign = 1, res = 0;
        for (; *it == '-' || *it == '+'; it++) {
            if (*it == '-') {
                sign *= -1;
            }
        }
        if (*it == '(') {
            res = eval(++it, 2);
            it++;
        } else {
            while (*it >= '0' && *it <= '9') {
                res = 10 * res + (*it++ - '0');
            }
        }
        return sign * res;
    }
}

int num = eval(it, prec - 1);
while (!(prec == 2 && *it != '+' && *it != '-' || (prec == 1 && *it != '*' && *it != '/'))) {
    switch (*it++) {
        case '+':
            num += eval(it, prec - 1);
            break;

```



```

        case '-':
            num -= eval(it, prec - 1);
            break;
        case '*':
            num *= eval(it, prec - 1);
            break;
        case '/':
            num /= eval(it, prec - 1);
            break;
    }
}
return num;
}

int eval(const std::string &s) {
    std::string expr = s + '\0';
    auto it = expr.begin();
    return eval(it, 2);
}

```

Example Usage

```

#include <cassert>

int main() {
    assert(eval("+1") == 1);
    assert(eval("1+-2") == -1);
    assert(eval("1++1") == 2);
    assert(eval("1+2*3*4+3*(2+2)-100") == -63);
    return 0;
}

```

3.7.2 Recursive Descent Parsing

Evaluate an expression with custom operand types, prefix unary operators, binary operators, and precedences. Evaluation is performed by recursive descent: the parser recursively evaluates tighter precedence levels before looser ones. Parentheses are supported, but multiplication by juxtaposition is not.

An arbitrary operand type is supported by changing the `Operand` alias and the static `is_operand()` and `eval_operand()` helpers. For maximum reliability, the string representation of operands should not use characters shared by any operator. For instance, the best practice instead of accepting `-1` as a valid operand (since the `-` sign may conflict with the identical binary operator), is to specify nonnegative numbers as operands alongside the unary operator `-`.

Operators may be non-empty strings of any length, but should not contain any parentheses or shared characters with the string representations of operands. Ideally, operators should not be prefixes or suffixes of one another, else the tokenization process may be ambiguous. For example, if `++` and `+` are both operators, then `++` may be split into either `["+"]`, `["+"]` or `["++"]` depending on the lexicographical ordering of conflicting operators.

- `RecursiveDescentParser(unary_ops, binary_ops)` initializes a parser with operators specified by hash tables `unary_ops` (of operator to unary function object) and `binary_ops` (of operator to pair of binary function object and operator precedence). Operator precedences should be numbered upwards starting at 0 (lowest precedence, evaluated last).
- `split(s)` returns a vector of tokens for the expression `s`, split on the given operators during construction. Each parenthesis, operator, and operand satisfying `is_operand()` will be split into a separate token. The algorithm is naive, matching operators lazily in the case of overlapping operators as mentioned above. Under these circumstances, the parse may not always succeed.
- `eval(lo, hi)` returns the evaluation of a range `[lo, hi)` of already split-up expression tokens, where `lo` and `hi` must be random-access iterators.

- `eval(s)` returns the evaluation of expression `s`, after first calling `split(s)` to obtain the tokens.

Time Complexity:

- $O(m)$ per call to the constructor, where m is the total number of operators.
- $O(n \cdot m \cdot k)$ per call to `split(s)`, where n is the length of `s`, m is the total number of operators defined for the parser instance, and k is the maximum length for any operator representation.
- $O(n)$ expected per call to `eval(lo, hi)`, where n is the distance between `lo` and `hi`.
- $O(n \cdot m \cdot k + n)$ expected per call to `eval(s)`, where n is the length of `s`, and m and k are as defined previously.

Space Complexity:

- $O(m \cdot k)$ for storage of the m operators, of maximum length k .
- $O(n)$ auxiliary stack space for `split(s)`, `eval(lo, hi)`, and `eval(s)`, where n is the length of the argument.

```
#include <algorithm>
#include <cctype>
#include <functional>
#include <set>
#include <sstream>
#include <stdexcept>
#include <string>
#include <unordered_map>
#include <vector>
using std::string;

// Define the custom operand type and representation below.
using Operand = double;
using UnaryOp = std::function<Operand(Operand)>;
using BinaryOp = std::function<Operand(Operand, Operand)>;

struct BinaryRule {
    BinaryOp op;
    int precedence;
};

class RecursiveDescentParser {
    using UnaryOpMap = std::unordered_map<string, UnaryOp>;
    using BinaryOpMap = std::unordered_map<string, BinaryRule>;
    UnaryOpMap unary_ops;
    BinaryOpMap binary_ops;
    std::set<string> op_tokens;
    int max_precedence;

    static bool is_operand(const string &s) {
        int npoints = 0;
        for (char c : s) {
            if (c == '.') {
                if (++npoints > 1) {
                    return false;
                }
            } else if (!isdigit(static_cast<unsigned char>(c))) {
                return false;
            }
        }
        return !s.empty();
    }

    static Operand eval_operand(const string &s) {
        Operand res;
        std::stringstream ss(s);
        ss >> res;
        return res;
    }
};

template<typename StrIt>
```

```

Operand eval_unary(StrIt &lo, StrIt hi) const {
    if (lo == hi) {
        throw std::runtime_error("Unexpected end of expression.");
    }
    if (is_operand(*lo)) {
        return eval_operand(*(lo++));
    }
    if (auto it = unary_ops.find(*lo); it != unary_ops.end()) {
        return (it->second)(eval_unary(++lo, hi));
    }
    if (*lo != "(") {
        throw std::runtime_error("Expected \"(\" during eval.");
    }
    Operand res = eval_binary(++lo, hi, 0);
    if (lo == hi || *lo != ")") {
        throw std::runtime_error("Expected \")\" during eval.");
    }
    ++lo;
    return res;
}

template<typename StrIt>
Operand eval_binary(StrIt &lo, StrIt hi, int precedence) const {
    if (precedence > max_precedence) {
        return eval_unary(lo, hi);
    }
    Operand v = eval_binary(lo, hi, precedence + 1);
    while (lo != hi) {
        auto it = binary_ops.find(*lo);
        if (it == binary_ops.end() || it->second.precedence != precedence) {
            return v;
        }
        v = (it->second.op)(v, eval_binary(++lo, hi, precedence + 1));
    }
    return v;
}

static string strip(string s) {
    auto not_space = [](unsigned char c) { return !std::isspace(c); };
    s.erase(s.begin(), std::find_if(s.begin(), s.end(), not_space));
    s.erase(std::find_if(s.rbegin(), s.rend(), not_space).base(), s.end());
    return s;
}

public:
RecursiveDescentParser(const UnaryOpMap &unary_ops, const BinaryOpMap &binary_ops)
    : unary_ops(unary_ops), binary_ops(binary_ops) {
    for (const auto &[op, _] : unary_ops) {
        op_tokens.insert(op);
    }
    max_precedence = 0;
    for (const auto &[op, fn_prec] : binary_ops) {
        op_tokens.insert(op);
        max_precedence = std::max(max_precedence, fn_prec.precedence);
    }
}

std::vector<string> split(const string &s) const {
    std::vector<string> res;
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        if (std::isspace(static_cast<unsigned char>(s[i]))) {
            continue;
        }
        int next_paren = static_cast<int>(s.size());
        for (int j = i; j < static_cast<int>(s.size()); j++) {
            if (s[j] == '(' || s[j] == ')') {
                next_paren = j;
                break;
            }
        }
        res.push_back(strip(s.substr(i, next_paren - i)));
        i = next_paren;
    }
    return res;
}

```

```

    }
}
while (i < next_paren) {
    int found = next_paren;
    string found_op;
    for (int j = i; j < next_paren && found == next_paren; j++) {
        for (const auto &op : op_tokens) {
            if (s.substr(j, op.size()) == op) {
                found = j;
                found_op = op;
                break;
            }
        }
    }
    string term = strip(s.substr(i, found - i));
    if (!term.empty()) {
        res.push_back(term);
        if (!is_operand(term)) {
            throw std::runtime_error("Failed to split term: \"" + term + "\".");
        }
    }
    if (found < next_paren) {
        res.push_back(found_op);
        i = found + static_cast<int>(found_op.size());
    } else {
        i = next_paren;
    }
}
if (next_paren < static_cast<int>(s.size())) {
    res.emplace_back(1, s[next_paren]);
}
}
return res;
}

template<typename StrIt>
Operand eval(StrIt lo, StrIt hi) const {
    Operand res = eval_binary(lo, hi, 0);
    if (lo != hi) {
        throw std::runtime_error("Eval failed at token " + *lo + ".");
    }
    return res;
}

Operand eval(const string &s) const {
    std::vector<string> tokens = split(s);
    return eval(tokens.begin(), tokens.end());
}
};

```

Example Usage

```

#include <cassert>
#include <cmath>
using namespace std;

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-7;
}

int main() {
    unordered_map<string, UnaryOp> unary_ops;
    unary_ops["+"] = [](double x) { return +x; };
    unary_ops["-"] = [](double x) { return -x; };

    unordered_map<string, BinaryRule> binary_ops;
}

```

```

binary_ops["+"] = {[](double a, double b) { return a + b; }, 0};
binary_ops["-"] = {[](double a, double b) { return a - b; }, 0};
binary_ops["*"] = {[](double a, double b) { return a * b; }, 1};
binary_ops["/"] = {[](double a, double b) { return a / b; }, 1};
binary_ops["^"] = {[](double a, double b) { return std::pow(a, b); }, 2};

RecursiveDescentParser p(unary_ops, binary_ops);
assert(EQ(p.eval("--+((--(-+1)))"), -1));
assert(EQ(p.eval("5*(3+3)-2-2"), 26));
assert(EQ(p.eval("1+2*3*4+3*(+2)-100"), -69));
assert(EQ(p.eval("3*3*3*3*3-2*2*2*2*2*2"), 473));
assert(EQ(p.eval("3.14 + 3 * (7.7/9.8^32.9 )"), 3.14));
assert(EQ(p.eval("5*(3+2)/-1*-2+(-2-2+3)-3-(-2)+15/2/2+(-2)", 45.875));
assert(EQ(p.eval("123456789./3/3/3*2*2+456/6-23/3"), 36579857.6666666667));
assert(EQ(p.eval("10/3+10/4+10/5+10/6+10/7+10/8+10/9+10/10+15*23456"), 351854.28968253968));
assert(
    EQ(p.eval(
        "-(5-(5-(5-(5-(5-2)))))+(3-(3-(3-(3-(3+3)))))*"
        "(7-(7-(7-(7-(7-7+4*5)))))"
    ),
    117)
);
return 0;
}

```

3.7.3 Shunting Yard Parsing

Evaluates an infix expression with custom operand types, prefix unary operators, binary operators, and precedences. In infix notation, binary operators appear between operands; prefix and postfix notation place operators before and after their operands, respectively. Postfix notation is also called reverse Polish notation (RPN).

The shunting yard algorithm commonly converts infix expressions to postfix notation. This implementation instead maintains operator and value stacks and evaluates each operator when it would be emitted to the postfix output, avoiding an intermediate token list. Multiplication by juxtaposition is not supported.

An arbitrary operand type is supported by changing the `Operand` alias and the static `is_operand()` and `eval_operand()` helpers. For maximum reliability, the string representation of operands should not use characters shared by any operator. For instance, the best practice instead of accepting `-1` as a valid operand (since the `-` sign may conflict with the identical binary operator), is to specify nonnegative numbers as operands alongside the unary operator `-`.

Operators may be non-empty strings of any length, but should not contain any parentheses or shared characters with the string representations of operands. Ideally, operators should not be prefixes or suffixes of one another, else the tokenization process may be ambiguous. For example, if `++` and `+` are both operators, then `++` may be split into either `["+"]`, `["+"]` or `["++"]` depending on the lexicographical ordering of conflicting operators.

- `ShuntingYardParser(unary_ops, binary_ops)` initializes a parser with operators specified by hash tables `unary_ops` (of operator to unary function object) and `binary_ops` (of operator to pair of binary function object and operator precedence). Operator precedences should be numbered upwards starting at 0 (lowest precedence, evaluated last).
- `split(s)` returns a vector of tokens for the expression `s`, split on the given operators during construction. Each parenthesis, operator, and operand satisfying `is_operand()` will be split into a separate token. The algorithm is naive, matching operators lazily in the case of overlapping operators as mentioned above. Under these circumstances, the parse may not always succeed.
- `eval(lo, hi)` returns the evaluation of a range `[lo, hi)` of already split-up expression tokens, where `lo` and `hi` must be random-access iterators.
- `eval(s)` returns the evaluation of expression `s`, after first calling `split(s)` to obtain the tokens.

Time Complexity:

- $O(m)$ per call to the constructor, where m is the total number of operators.
- $O(n \cdot m \cdot k)$ per call to `split(s)`, where n is the length of `s`, m is the total number of operators defined for the parser instance, and k is the maximum length for any operator representation.
- $O(n)$ expected per call to `eval(lo, hi)`, where n is the distance between `lo` and `hi`.
- $O(n \cdot m \cdot k + n)$ expected per call to `eval(s)`, where n is the length of `s`, and m and k are as defined previously.

Space Complexity:

- $O(m \cdot k)$ for storage of the m operators, of maximum length k .
- $O(n)$ auxiliary stack space for `split(s)`, `eval(lo, hi)`, and `eval(s)`, where n is the length of the argument.

```
#include <algorithm>
#include <cctype>
#include <functional>
#include <set>
#include <sstream>
#include <stack>
#include <stdexcept>
#include <string>
#include <unordered_map>
#include <utility>
#include <vector>
using std::string;

// Define the custom operand type and representation below.
using Operand = double;
using UnaryOp = std::function<Operand(Operand)>;
using BinaryOp = std::function<Operand(Operand, Operand)>;

struct BinaryRule {
    BinaryOp op;
    int precedence;
};

class ShuntingYardParser {
    using UnaryOpMap = std::unordered_map<string, UnaryOp>;
    using BinaryOpMap = std::unordered_map<string, BinaryRule>;
    UnaryOpMap unary_ops;
    BinaryOpMap binary_ops;
    std::set<string> op_tokens;

    static Operand eval_operand(const string &s) {
        Operand res;
        std::stringstream ss(s);
        ss >> res;
        return res;
    }

    static bool is_operand(const string &s) {
        int npoints = 0;
        for (char c : s) {
            if (c == '.') {
                if (++npoints > 1) {
                    return false;
                }
            } else if (!isdigit(static_cast<unsigned char>(c))) {
                return false;
            }
        }
        return !s.empty();
    }

    static string strip(string s) {
        auto not_space = [](unsigned char c) { return !std::isspace(c); };
    }
};
```

```

    s.erase(s.begin(), std::find_if(s.begin(), s.end(), not_space));
    s.erase(std::find_if(s.rbegin(), s.rend(), not_space).base(), s.end());
    return s;
}

public:
ShuntingYardParser(const UnaryOpMap &unary_ops, const BinaryOpMap &binary_ops)
    : unary_ops(unary_ops), binary_ops(binary_ops) {
    for (const auto &[op, fn] : unary_ops) {
        op_tokens.insert(op);
    }
    for (const auto &[op, fn_prec] : binary_ops) {
        op_tokens.insert(op);
    }
}

std::vector<string> split(const string &s) const {
    std::vector<string> res;
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        if (std::isspace(static_cast<unsigned char>(s[i]))) {
            continue;
        }
        int next_paren = static_cast<int>(s.size());
        for (int j = i; j < static_cast<int>(s.size()); j++) {
            if (s[j] == '(' || s[j] == ')') {
                next_paren = j;
                break;
            }
        }
        while (i < next_paren) {
            int found = next_paren;
            string found_op;
            for (int j = i; j < next_paren && found == next_paren; j++) {
                for (const auto &op : op_tokens) {
                    if (s.substr(j, op.size()) == op) {
                        found = j;
                        found_op = op;
                        break;
                    }
                }
            }
            string term = strip(s.substr(i, found - i));
            if (!term.empty()) {
                res.push_back(term);
                if (!is_operand(term)) {
                    throw std::runtime_error("Failed to split term: \"" + term + "\".");
                }
            }
            if (found < next_paren) {
                res.push_back(found_op);
                i = found + static_cast<int>(found_op.size());
            } else {
                i = next_paren;
            }
        }
        if (next_paren < static_cast<int>(s.size())) {
            res.emplace_back(1, s[next_paren]);
        }
    }
    return res;
}

template<typename StrIt>
Operand eval(StrIt lo, StrIt hi) const {
    if (lo == hi) {
        throw std::runtime_error("Cannot evaluate an empty expression.");
    }
    std::stack<Operand> vals;

```

```

std::stack<std::pair<string, bool>> op_stack;
op_stack.emplace("(", false);
StrIt prev = hi;
do {
    string curr = (lo == hi) ? ")" : *lo;
    if (is_operand(curr)) {
        vals.push(eval_operand(curr));
    } else if (curr == "(") {
        op_stack.emplace(curr, false);
    } else if (
        unary_ops.find(curr) != unary_ops.end() &&
        (prev == hi || *prev == "(" || binary_ops.find(*prev) != binary_ops.end())
    ) {
        op_stack.emplace(curr, true);
    } else {
        while (true) {
            auto [op, is_unary] = op_stack.top();
            auto it1 = binary_ops.find(op);
            auto it2 = binary_ops.find(curr);
            if (!is_unary && (it1 == binary_ops.end() ? -1 : it1->second.precedence) <
                (it2 == binary_ops.end() ? -1 : it2->second.precedence)) {
                break;
            }
            op_stack.pop();
            if (op == "(") {
                break;
            }
            Operand b = vals.top();
            vals.pop();
            if (is_unary) {
                if (auto it = unary_ops.find(op); it != unary_ops.end()) {
                    vals.push((it->second)(b));
                } else {
                    throw std::runtime_error("Failed to eval unary op: " + op);
                }
            } else {
                Operand a = vals.top();
                vals.pop();
                if (it1 == binary_ops.end()) {
                    throw std::runtime_error("Failed to eval binary op: " + op);
                }
                vals.push((it1->second.op)(a, b));
            }
        }
        if (curr != ")") {
            op_stack.emplace(*lo, false);
        }
    }
    prev = lo;
} while (lo++ != hi);
return vals.top();
}

Operand eval(const string &s) const {
    std::vector<string> tokens = split(s);
    return eval(tokens.begin(), tokens.end());
}
};

```

Example Usage

```

#include <cassert>
#include <cmath>
using namespace std;

bool EQ(double a, double b) {

```



```

    return fabs(a - b) < 1e-7;
}

int main() {
    unordered_map<string, UnaryOp> unary_ops;
    unary_ops["+"] = [](double x) { return +x; };
    unary_ops["-"] = [](double x) { return -x; };

    unordered_map<string, BinaryRule> binary_ops;
    binary_ops["+"] = {[](double a, double b) { return a + b; }, 0};
    binary_ops["-"] = {[](double a, double b) { return a - b; }, 0};
    binary_ops["*"] = {[](double a, double b) { return a * b; }, 1};
    binary_ops["/"] = {[](double a, double b) { return a / b; }, 1};
    binary_ops["^"] = {[](double a, double b) { return std::pow(a, b); }, 2};

    ShuntingYardParser p(unary_ops, binary_ops);
    assert(EQ(p.eval("-+-(--(+1))"), -1));
    assert(EQ(p.eval("5*(3+3)-2-2"), 26));
    assert(EQ(p.eval("1+2*3*4+3*(+2)-100"), -69));
    assert(EQ(p.eval("3*3*3*3*3-2*2*2*2*2*2"), 473));
    assert(EQ(p.eval("3.14 + 3 * (7.7/9.8~32.9 )"), 3.14));
    assert(EQ(p.eval("5*(3+2)/-1*-2+(-2-2+3)-3-(-2)+15/2/2+(-2)"), 45.875));
    assert(EQ(p.eval("123456789./3/3/3*2*2*2+456/6-23/3"), 36579857.6666666667));
    assert(EQ(p.eval("10/3+10/4+10/5+10/6+10/7+10/8+10/9+10/10+15*23456"), 351854.28968253968));
    assert(
        EQ(p.eval(
            "-(5-(5-(5-(5-(5-2)))))+(3-(3-(3-(3-(3+3)))))*"
            "(7-(7-(7-(7-(7-7+4*5)))))"
        ),
        117)
    );
    return 0;
}

```

3.8 Encoding and Compression

3.8.1 Entropy and Frequency Tables

Computes byte frequency tables and empirical information-theoretic quantities for strings. These primitives are useful when analyzing compression, building coding schemes such as Huffman coding, or comparing symbol distributions.

For a string with symbol probabilities p_i , empirical entropy is $H = -\sum_i p_i \log_2 p_i$ bits per symbol. This is a lower bound on the average number of bits per symbol achievable by any uniquely decodable code for the same independent symbol model.

- `byte_frequencies(s)` returns a length-256 table of byte counts in string `s`.
- `entropy(freq)` returns empirical entropy in bits per symbol for a frequency table.
- `entropy(s)` returns empirical entropy in bits per symbol for string `s`.
- `expected_code_length(freq, length)` returns the average encoded bits per symbol for code lengths `length[c]`.

The two input vectors must have the same size.

Time Complexity:

- $O(n + m)$ per call to `entropy(s)`, where n is the string length and $m = 256$.
- $O(m)$ per call to `entropy(freq)` and `expected_code_length(freq, length)`.

Space Complexity:

- $O(m)$ for the table returned by `byte_frequencies()` and $O(m)$ auxiliary for `entropy(s)`.
- $O(1)$ auxiliary for the other operations.

```

#include <cassert>
#include <cmath>
#include <cstdint>
#include <string>
#include <vector>
using std::string;

std::vector<int> byte_frequencies(const string &s) {
    std::vector<int> freq(256, 0);
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        freq[static_cast<unsigned char>(s[i])]++;
    }
    return freq;
}

double entropy(const std::vector<int> &freq) {
    int64_t total = 0;
    for (int i = 0; i < static_cast<int>(freq.size()); i++) {
        total += freq[i];
    }
    if (total == 0) {
        return 0;
    }
    double res = 0;
    for (int i = 0; i < static_cast<int>(freq.size()); i++) {
        if (freq[i] == 0) {
            continue;
        }
        double p = static_cast<double>(freq[i]) / total;
        res -= p * std::log2(p);
    }
    return res;
}

double entropy(const string &s) {
    return entropy(byte_frequencies(s));
}

double expected_code_length(const std::vector<int> &freq, const std::vector<int> &length) {
    assert(freq.size() == length.size());
    int64_t total = 0;
    for (int i = 0; i < static_cast<int>(freq.size()); i++) {
        total += freq[i];
    }
    if (total == 0) {
        return 0;
    }
    double res = 0;
    for (int i = 0; i < static_cast<int>(freq.size()); i++) {
        res += static_cast<double>(freq[i]) * length[i] / total;
    }
    return res;
}

```

Example Usage

```

#include <cmath>
using namespace std;

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-9;
}

int main() {
    vector<int> freq = byte_frequencies("aabb");
    assert(freq['a'] == 2);
}

```

```

assert(freq['b'] == 2);
assert(EQ(entropy(freq), 1.0));
assert(EQ(entropy("aaaa"), 0.0));

vector<int> length(256, 0);
length['a'] = 1;
length['b'] = 1;
assert(EQ(expected_code_length(freq, length), 1.0));
return 0;
}

```

3.8.2 Run-Length Encoding

Encodes consecutive equal characters as (`character`, `count`) runs. Run-length encoding is a simple lossless compression technique that works well when strings contain long repeated blocks and poorly when repetitions are rare.

The representation below stores runs in a vector of pairs rather than formatting them as a string. This avoids ambiguity when the original text contains digits or separator characters.

- `run_length_encode(s)` returns the sequence of runs in string `s`.
- `run_length_decode(runs)` reconstructs the original string from a sequence of runs.

Time Complexity: $O(n)$ per call, where n is the input or output length.

Space Complexity:

- $O(r)$ for encoded output, where r is the number of runs.
- $O(n)$ for decoded output.

```

#include <string>
#include <utility>
#include <vector>
using std::string;

std::vector<std::pair<char, int>> run_length_encode(const string &s) {
    std::vector<std::pair<char, int>> res;
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        if (res.empty() || res.back().first != s[i]) {
            res.emplace_back(s[i], 1);
        } else {
            res.back().second++;
        }
    }
    return res;
}

string run_length_decode(const std::vector<std::pair<char, int>> &runs) {
    string res;
    for (int i = 0; i < static_cast<int>(runs.size()); i++) {
        res.append(runs[i].second, runs[i].first);
    }
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    string s = "aaabcccd";

```

```

vector<pair<char, int>> runs = run_length_encode(s);
assert((runs == vector<pair<char, int>>{{'a', 3}, {'b', 1}, {'c', 4}, {'d', 2}}));
assert(run_length_decode(runs) == s);
assert(run_length_encode("").empty());
assert(run_length_decode({}) == "");
assert((run_length_encode("z") == vector<pair<char, int>>{{'z', 1}}));
return 0;
}

```

3.8.3 Base64 Encoding

Encodes binary data as printable ASCII using Base64, and decodes Base64 text back to bytes. Base64 is an encoding, not encryption or compression: it preserves all input data exactly, but usually increases size by about one third. Each group of three input bytes is split into four 6-bit values, each selecting one character of a 64-symbol alphabet, with `=` padding marking a final group of only one or two bytes.

This implementation uses the standard alphabet with `+`, `/`, and `=` padding. Whitespace is not skipped by the decoder; clean the input first if line breaks or spaces may appear.

- `base64_encode(bytes)` returns the Base64 representation of string `bytes`.
- `base64_decode(text)` returns the decoded byte string. It assumes `text` is valid padded Base64.

Time Complexity: $O(n)$ per call, where n is the input length.

Space Complexity: $O(n)$ for the output string.

```

#include <stdint>
#include <string>
#include <vector>
using std::string;

const char *BASE64_ALPHABET = "ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz0123456789+/";

string base64_encode(const string &bytes) {
    string res;
    uint32_t val = 0;
    int bits = -6;
    for (int i = 0; i < static_cast<int>(bytes.size()); i++) {
        val = (val << 8) + static_cast<unsigned char>(bytes[i]);
        bits += 8;
        while (bits >= 0) {
            res.push_back(BASE64_ALPHABET[(val >> bits) & 63]);
            bits -= 6;
        }
    }
    if (bits > -6) {
        res.push_back(BASE64_ALPHABET[((val << 8) >> (bits + 8)) & 63]);
    }
    while (res.size() % 4) {
        res.push_back('=');
    }
    return res;
}

string base64_decode(const string &text) {
    std::vector<int> table(256, -1);
    for (int i = 0; i < 64; i++) {
        table[static_cast<unsigned char>(BASE64_ALPHABET[i])] = i;
    }
    string res;
    uint32_t val = 0;
    int bits = -8;
    for (int i = 0; i < static_cast<int>(text.size()); i++) {
        if (text[i] == '=') {

```

```

        break;
    }
    int x = table[static_cast<unsigned char>(text[i])];
    if (x < 0) {
        break;
    }
    val = (val << 6) + x;
    bits += 6;
    if (bits >= 0) {
        res.push_back(static_cast<char>(((val >> bits) & 255)));
        bits -= 8;
    }
}
return res;
}

```

Example Usage

```

#include <cassert>

int main() {
    assert(base64_encode("") == "");
    assert(base64_decode("") == "");
    assert(base64_encode("f") == "Zg==");
    assert(base64_encode("fo") == "Zm8=");
    assert(base64_encode("hello") == "aGVsbG8=");
    assert(base64_decode("aGVsbG8=") == "hello");
    assert(base64_decode(base64_encode("abc123!?")) == "abc123!?");
    return 0;
}

```

3.8.4 Huffman Coding

Builds a Huffman code for a string and uses it to encode and decode that string. Huffman coding is a lossless prefix-code compression algorithm: more frequent characters receive shorter bit strings, and no code is a prefix of another code. The code tree is built greedily with a priority queue by repeatedly merging the two lowest-frequency subtrees until one tree remains; each character's bit string is then read off its root-to-leaf path.

The implementation below stores encoded bits as a string of `'0'` and `'1'` characters for clarity. For real compression, pack those bits into bytes. The tree is also needed to decode the bit string, so compressed data normally stores enough metadata to reconstruct the same tree.

- `HuffmanTree(text)` constructs a Huffman tree from the character frequencies in `text`.
- `codes()` returns a table mapping each byte value to its bit string.
- `encode(text)` returns the encoded bit string. Every character in `text` must occur in the string used to construct the tree.
- `decode(bits)` returns the decoded text for a bit string produced by this tree.

Time Complexity:

- $O(n + m \log m)$ per call to the constructor, where n is the input length and m is the number of distinct characters.
- $O(n + b)$ per call to `encode()` for n characters producing b bits, and $O(b)$ per call to `decode()`.

Space Complexity:

- $O(m)$ for the tree and code table, where m is the number of distinct characters.
- $O(n)$ for the encoded or decoded output.

```

#include <functional>
#include <queue>

```

```

#include <string>
#include <utility>
#include <vector>
using std::string;

class HuffmanTree {
    struct Node {
        int freq;
        unsigned char ch;
        int left, right;

        Node(int freq = 0, unsigned char ch = 0, int left = -1, int right = -1)
            : freq(freq), ch(ch), left(left), right(right) {}
    };

    int root;
    std::vector<Node> nodes;
    std::vector<string> code;

    void build_codes(int u, const string &path) {
        if (nodes[u].left == -1 && nodes[u].right == -1) {
            code[nodes[u].ch] = path.empty() ? "0" : path;
            return;
        }
        build_codes(nodes[u].left, path + '0');
        build_codes(nodes[u].right, path + '1');
    }

public:
    explicit HuffmanTree(const string &text) : root(-1), code(256) {
        std::vector<int> freq(256, 0);
        for (int i = 0; i < static_cast<int>(text.size()); i++) {
            freq[static_cast<unsigned char>(text[i])]++;
        }
        std::priority_queue<std::pair<int, int>, std::vector<std::pair<int, int>>, std::greater<>> pq;
        for (int c = 0; c < 256; c++) {
            if (freq[c] > 0) {
                nodes.emplace_back(freq[c], static_cast<unsigned char>(c));
                pq.emplace(freq[c], static_cast<int>(nodes.size()) - 1);
            }
        }
        if (pq.empty()) {
            return;
        }
        while (static_cast<int>(pq.size()) > 1) {
            int a = pq.top().second;
            pq.pop();
            int b = pq.top().second;
            pq.pop();
            nodes.emplace_back(nodes[a].freq + nodes[b].freq, 0, a, b);
            pq.emplace(nodes.back().freq, static_cast<int>(nodes.size()) - 1);
        }
        root = pq.top().second;
        build_codes(root, "");
    }

    std::vector<string> codes() const { return code; }

    string encode(const string &text) const {
        string bits;
        for (int i = 0; i < static_cast<int>(text.size()); i++) {
            bits += code[static_cast<unsigned char>(text[i])];
        }
        return bits;
    }

    string decode(const string &bits) const {
        string text;

```

```

    if (root == -1) {
        return text;
    }
    // Degenerate tree: the input had a single distinct character, so the root is itself a leaf and
    // each symbol was encoded as one bit. Walking children would dereference node -1.
    if (nodes[root].left == -1 && nodes[root].right == -1) {
        return string(bits.size(), static_cast<char>(nodes[root].ch));
    }
    int u = root;
    for (int i = 0; i < static_cast<int>(bits.size()); i++) {
        u = bits[i] == '0' ? nodes[u].left : nodes[u].right;
        if (nodes[u].left == -1 && nodes[u].right == -1) {
            text.push_back(static_cast<char>(nodes[u].ch));
            u = root;
        }
    }
    return text;
}
};

```

Example Usage

```

#include <cassert>

int main() {
    string text = "banana bandana";
    HuffmanTree h(text);
    string bits = h.encode(text);
    assert(h.decode(bits) == text);
    assert(bits.size() < text.size() * 8);

    HuffmanTree single("aaaa");
    assert(single.encode("aaaa") == "0000");
    assert(single.decode(single.encode("aaaa")) == "aaaa");

    HuffmanTree empty("");
    assert(empty.encode("") == "");
    assert(empty.decode("") == "");
    return 0;
}

```

3.8.5 Burrows-Wheeler Transform

The Burrows-Wheeler transform (BWT) permutes a string into one that tends to be far more compressible while remaining perfectly invertible. It writes down every cyclic rotation of the string, sorts them lexicographically, and reads off the last character of each sorted rotation. Characters that occur in similar contexts are brought together, producing long runs that run-length and move-to-front coders exploit; it is the heart of the bzip2 compressor and a building block of compressed-text indexes.

To make the permutation invertible, a unique sentinel character smaller than every real character is appended before transforming, so the rotations are all distinct and one of them is distinguished as the original. Sorting the rotations of the sentineled string is the same as sorting its suffixes, so the forward transform is computed with a prefix-doubling rotation sort. The inverse rebuilds the string with the last-to-first (LF) mapping: stably ordering the transformed characters yields, for each position of the sorted first column, the row that supplies the next character, and following that mapping from the sentinel row walks the original string out one character at a time.

The input must not contain the sentinel character `'\0'`.

- `bwt(s)` returns the BWT of `s`. The result has length `s.size() + 1` and contains the sentinel `'\0'` exactly once.

- `inverse_bwt(code)` returns the original string, inverting `bwt`. The argument must be a string produced by that function.

Time Complexity:

- $O(n \log^2 n)$ per call to `bwt(s)`, where n is the length of `s`.
- $O(n \log n)$ per call to `inverse_bwt(code)`.

Space Complexity: $O(n)$ auxiliary for both functions.

```
#include <algorithm>
#include <numeric>
#include <string>
#include <vector>

std::string bwt(std::string s) {
    s.push_back('\0'); // Unique sentinel, smaller than every real character.
    int n = static_cast<int>(s.size());
    // Sort the indices of the cyclic rotations by prefix doubling; with the sentinel this matches
    // sorting the suffixes.
    std::vector<int> order(n), rank(n), tmp(n);
    for (int i = 0; i < n; i++) {
        order[i] = i;
        rank[i] = static_cast<unsigned char>(s[i]);
    }
    for (int k = 1; k < n; k <= 1) {
        auto cmp = [&](int a, int b) {
            if (rank[a] != rank[b]) {
                return rank[a] < rank[b];
            }
            return rank[(a + k) % n] < rank[(b + k) % n];
        };
        std::sort(order.begin(), order.end(), cmp);
        tmp[order[0]] = 0;
        for (int i = 1; i < n; i++) {
            tmp[order[i]] = tmp[order[i - 1]] + (cmp(order[i - 1], order[i]) ? 1 : 0);
        }
        rank = tmp;
        if (rank[order[n - 1]] == n - 1) {
            break;
        }
    }
    std::string code(n, '\0');
    for (int i = 0; i < n; i++) {
        code[i] = s[(order[i] + n - 1) % n]; // Last character of the i-th sorted rotation.
    }
    return code;
}

std::string inverse_bwt(const std::string &code) {
    int n = static_cast<int>(code.size());
    // first[k] is the row supplying the first column's k-th character; following it from row 0 (the
    // sentinel row, sorted first) reconstructs the sentinel followed by the original string.
    std::vector<int> first(n);
    std::iota(first.begin(), first.end(), 0);
    std::stable_sort(first.begin(), first.end(), [&](int a, int b) {
        return static_cast<unsigned char>(code[a]) < static_cast<unsigned char>(code[b]);
    });
    std::string res;
    res.reserve(n);
    for (int i = 0, row = 0; i < n; i++) {
        row = first[row];
        res.push_back(code[row]);
    }
    return res.substr(1); // Drop the leading sentinel.
}
```


Example Usage

```
#include <cassert>
#include <iostream>
using namespace std;

string escaped(const string &s) {
    string res;
    for (char c : s) {
        res += (c == '\\0') ? "\\0" : string(1, c);
    }
    return res;
}

int count_runs(const string &s) {
    int runs = 0;
    for (int i = 0; i < static_cast<int>(s.size()); i++) {
        runs += i == 0 || s[i] != s[i - 1];
    }
    return runs;
}

int main() {
    string code = bwt("banana");
    assert(code.size() == 7); // "banana" plus the sentinel.
    assert(escaped(code) == "annb\\0aa");
    assert(inverse_bwt(code) == "banana");
    cout << "BWT(banana): " << escaped(code) << endl;
    cout << "inverse: " << inverse_bwt(code) << endl;

    // BWT does not compress by itself, but can cluster characters for a subsequent run-length coder.
    string text = "abracadabra abracadabra";
    string transformed = bwt(text);
    assert(inverse_bwt(transformed) == text);
    assert(count_runs(text) == 23);
    assert(count_runs(transformed) == 9);

    // Round-trips on assorted strings, including repeats and a single character.
    for (const string &s :
        {string(""), string("a"), string("aaaa"), string("abracadabra"),
         string("the quick brown fox"), string("\\1\\2abc", 5)}) {
        assert(inverse_bwt(bwt(s)) == s);
    }
    return 0;
}
```

Example Output

```
BWT(banana): annb\0aa
inverse: banana
```

Chapter 4

Graphs

4.1 DFS and Tree Algorithms

4.1.1 Graph Class and DFS

A graph consists of a set of nodes and a set of connections (edges) between pairs of nodes. A graph may be stored as an adjacency list, which is a space-efficient representation that is also efficient for traversals.

The following class implements a simple graph using adjacency lists, along with depth-first search and a few other applications. The constructor takes a flag which specifies whether the instance is a directed or undirected graph. The nodes of the graph are identified by integer indices numbered consecutively starting from 0. The total number of nodes automatically increases based on the maximum node index passed to `add_edge()` so far.

- `Graph(directed = true)` constructs an empty graph, directed if `directed` is true and undirected otherwise.
- `size()` returns the current number of nodes.
- `operator[u]` returns a mutable or const reference to the adjacency list (`std::vector<int>`) of node `u`.
- `add_edge(u, v)` adds an edge from `u` to `v`, plus the reverse edge if the graph is undirected, growing the node count to accommodate the larger index if necessary.
- `dfs(start, f)` runs a depth-first search from node `start`, calling function `f` on each node in the order it is first visited.
- `has_cycle()` returns whether the graph contains a cycle.
- `is_directed()` returns whether the graph is directed.
- `is_forest()` returns whether the graph is undirected and acyclic.
- `is_dag()` returns whether the graph is directed and acyclic.

For undirected graphs, `has_cycle()` assumes a simple graph; parallel edges are not distinguished from the tree edge back to the parent.

Time Complexity:

- $O(1)$ amortized per call to `add_edge()` when no resize is required; $O(r)$ for a call that grows the graph to r nodes.
- $O(\max(n, m))$ per call for `dfs()`, `has_cycle()`, `is_forest()`, or `is_dag()`, where n is the number of nodes and m is the number of edges.
- $O(1)$ per call to all other public member functions.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary stack space for `dfs()`, `has_cycle()`, `is_forest()`, and `is_dag()`.
- $O(1)$ auxiliary for all other public member functions.

```

#include <algorithm>
#include <vector>

class Graph {
    std::vector<std::vector<int>>> adj;
    bool directed;

    template<typename Fn>
    void dfs(int u, std::vector<char> &visit, Fn f) const {
        f(u);
        visit[u] = true;
        for (int v : adj[u]) {
            if (!visit[v]) {
                dfs(v, visit, f);
            }
        }
    }

    bool has_cycle(int u, int prev, std::vector<char> &visit, std::vector<char> &onstack) const {
        visit[u] = true;
        onstack[u] = true;
        for (int v : adj[u]) {
            if (directed && onstack[v]) {
                return true;
            }
            if (!directed && visit[v] && v != prev) {
                return true;
            }
            if (!visit[v] && has_cycle(v, u, visit, onstack)) {
                return true;
            }
        }
        onstack[u] = false;
        return false;
    }

public:
    Graph(bool directed = true) : directed(directed) {}

    int size() const { return static_cast<int>(adj.size()); }
    std::vector<int> &operator[](int u) { return adj[u]; }
    const std::vector<int> &operator[](int u) const { return adj[u]; }

    void add_edge(int u, int v) {
        int n = static_cast<int>(adj.size());
        if (u >= n || v >= n) {
            adj.resize(std::max(u, v) + 1);
        }
        adj[u].push_back(v);
        if (!directed) {
            adj[v].push_back(u);
        }
    }

    bool has_cycle() const {
        int n = static_cast<int>(adj.size());
        std::vector<char> visit(n, false), onstack(n, false);
        for (int i = 0; i < n; i++) {
            if (!visit[i] && has_cycle(i, -1, visit, onstack)) {
                return true;
            }
        }
        return false;
    }

    bool is_directed() const { return directed; }
    bool is_forest() const { return !directed && !has_cycle(); }
    bool is_dag() const { return directed && !has_cycle(); }

```

```

template<typename Fn>
void dfs(int start, Fn f) const {
    std::vector<char> visit(adj.size(), false);
    dfs(start, visit, f);
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    {
        //          0
        //        / / |
        //       v v v
        //      1 6 7
        //     / |   / |
        //    v v   v v
        //   2 5   8 11
        //  / |   / |
        // v v   v v
        // 3 4   9 10
        Graph g;
        g.add_edge(0, 1);
        g.add_edge(0, 6);
        g.add_edge(0, 7);
        g.add_edge(1, 2);
        g.add_edge(1, 5);
        g.add_edge(2, 3);
        g.add_edge(2, 4);
        g.add_edge(7, 8);
        g.add_edge(7, 11);
        g.add_edge(8, 9);
        g.add_edge(8, 10);
        vector<int> order;
        g.dfs(0, [&](int u) { order.push_back(u); });
        assert((order == vector<int>{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11}));
        assert(g.size() == 12 && g[0].size() == 3);
        assert(g.is_dag());
        assert(!g.has_cycle());
    }
    {
        //          0
        //        / |
        //       1 2
        //      / |
        //     3 4
        Graph tree(false);
        tree.add_edge(0, 1);
        tree.add_edge(0, 2);
        tree.add_edge(1, 3);
        tree.add_edge(1, 4);
        assert(tree.is_forest());
        assert(!tree.is_dag());
        tree.add_edge(2, 3);
        assert(!tree.is_forest());
    }
    return 0;
}

```

4.1.2 Topological Sort

Given a directed acyclic graph (DAG), find one of possibly many orderings of the nodes such that for every edge from node u to v , u comes before v in the ordering.

The DFS version uses the fact that every outgoing dependency of a node is finished before the node itself is appended, so reversing DFS post-order places each prerequisite before everything it can reach. Kahn's algorithm instead repeatedly removes nodes with indegree zero: these are exactly the nodes whose remaining prerequisites have all already been placed. If at some point not every node can be removed, the leftover nodes must lie on a directed cycle.

- `toposort_dfs()` returns a valid topological ordering using DFS post-order, or throws if the graph contains a cycle.
- `toposort_kahn()` returns a valid topological ordering using indegrees and a queue, or returns an empty vector if the graph contains a cycle.

Time Complexity: $O(\max(n, m))$ per call to either function, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary stack and heap space for `toposort_dfs()`.
- $O(n)$ auxiliary heap space for the indegree array and queue in `toposort_kahn()`.

```
#include <algorithm>
#include <cassert>
#include <queue>
#include <stdexcept>
#include <vector>

std::vector<std::vector<int>>> adj;

void dfs(int u, std::vector<char> &visit, std::vector<char> &done, std::vector<int> &res) {
    if (visit[u]) {
        throw std::runtime_error("Not a directed acyclic graph.");
    }
    if (done[u]) {
        return;
    }
    visit[u] = true;
    for (int v : adj[u]) {
        dfs(v, visit, done, res);
    }
    visit[u] = false;
    done[u] = true;
    res.push_back(u);
}

std::vector<int> toposort_dfs() {
    int n = static_cast<int>(adj.size());
    std::vector<char> visit(n, false), done(n, false);
    std::vector<int> res;
    for (int i = 0; i < n; i++) {
        if (!done[i]) {
            dfs(i, visit, done, res);
        }
    }
    std::reverse(res.begin(), res.end());
    return res;
}

std::vector<int> toposort_kahn() {
    int n = static_cast<int>(adj.size());
    std::vector<int> indegree(n), res;
```

```

for (int u = 0; u < n; u++) {
    for (int v : adj[u]) {
        indegree[v]++;
    }
}
std::queue<int> q;
for (int i = 0; i < n; i++) {
    if (indegree[i] == 0) {
        q.push(i);
    }
}
while (!q.empty()) {
    int u = q.front();
    q.pop();
    res.push_back(u);
    for (int v : adj[u]) {
        if (--indegree[v] == 0) {
            q.push(v);
        }
    }
}
return static_cast<int>(res.size()) == n ? res : std::vector<int>();
}

```

Example Usage

```

using namespace std;

int main() {
    //      1
    //      |
    //      v
    // 0 ----> 3 ----> 5
    // |      / |
    // v      v v
    // 4 -> 6 7 <--- 2
    // ~-----|
    adj.assign(8, {});
    adj[0].push_back(3);
    adj[0].push_back(4);
    adj[1].push_back(3);
    adj[2].push_back(4);
    adj[2].push_back(7);
    adj[3].push_back(5);
    adj[3].push_back(6);
    adj[3].push_back(7);
    adj[4].push_back(6);
    vector<int> dfs_res = toposort_dfs();
    assert((dfs_res == vector<int>{2, 1, 0, 4, 3, 7, 6, 5}));
    vector<int> kahn = toposort_kahn();
    assert(dfs_res.size() == adj.size());
    assert(kahn.size() == adj.size());
    vector<int> position(kahn.size());
    for (int i = 0; i < static_cast<int>(kahn.size()); i++) {
        position[kahn[i]] = i;
    }
    for (int u = 0; u < static_cast<int>(adj.size()); u++) {
        for (int v : adj[u]) {
            assert(position[u] < position[v]);
        }
    }
    return 0;
}

```

4.1.3 Cycle Finding

Given a directed or undirected graph, find simple cycles encountered by a depth-first search. The free functions return one cycle of a simple graph as a list of nodes. For undirected graphs, an edge to a visited non-parent node closes a cycle. For directed graphs, only an edge to a node still on the recursion stack closes a cycle; edges to already-finished nodes do not.

Both functions remember the DFS parent of each node. When a closing edge `u → cycle_start` is found, the answer is reconstructed by walking parents backward from `u` to `cycle_start`. In the directed version, `state[v]` is 0 for unvisited nodes, 1 for nodes currently on the recursion stack, and 2 for finished nodes.

- `find_cycle_undirected(adj)` returns one node cycle in a simple undirected graph, or an empty vector if no cycle exists.
- `find_cycle_directed(adj)` returns one node cycle in a simple directed graph, or an empty vector if no cycle exists.

Time Complexity: $O(\max(n, m))$ per call to `find_cycle_undirected()` and `find_cycle_directed()`, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph.
- $O(n)$ auxiliary stack and heap space for the DFS, plus $O(n)$ for the returned cycle.

```
#include <algorithm>
#include <cassert>
#include <utility>
#include <vector>

std::vector<int> find_cycle_undirected(const std::vector<std::vector<int>> &adj) {
    int n = static_cast<int>(adj.size());
    if (n == 0) {
        return {};
    }
    std::vector<char> visit(n, false);
    std::vector<int> parent(n, -1);
    int cycle_start = -1, cycle_end = -1;
    auto dfs = [&](auto &&dfs, int u, int p) -> bool {
        visit[u] = true;
        for (int v : adj[u]) {
            if (v == p) {
                continue;
            }
            if (visit[v]) {
                cycle_start = v;
                cycle_end = u;
                return true;
            }
            parent[v] = u;
            if (dfs(dfs, v, u)) {
                return true;
            }
        }
        return false;
    };
    for (int u = 0; u < n && cycle_start == -1; u++) {
        if (!visit[u]) {
            dfs(dfs, u, -1);
        }
    }
    if (cycle_start == -1) {
        return {};
    }
    std::vector<int> cycle{cycle_start};
    for (int u = cycle_end; u != cycle_start; u = parent[u]) {
        cycle.push_back(u);
    }
}
```

```

    }
    std::reverse(cycle.begin() + 1, cycle.end());
    return cycle;
}

std::vector<int> find_cycle_directed(const std::vector<std::vector<int>> &adj) {
    int n = static_cast<int>(adj.size());
    if (n == 0) {
        return {};
    }
    std::vector<char> state(n, 0); // 0 = unvisited, 1 = currently on DFS stack, 2 = finished.
    std::vector<int> parent(n, -1);
    int cycle_start = -1, cycle_end = -1;
    auto dfs = [&](auto &&dfs, int u) -> bool {
        state[u] = 1;
        for (int v : adj[u]) {
            if (state[v] == 0) {
                parent[v] = u;
                if (dfs(dfs, v)) {
                    return true;
                }
            } else if (state[v] == 1) {
                cycle_start = v;
                cycle_end = u;
                return true;
            }
        }
        state[u] = 2;
        return false;
    };
    for (int u = 0; u < n && cycle_start == -1; u++) {
        if (state[u] == 0) {
            dfs(dfs, u);
        }
    }
    if (cycle_start == -1) {
        return {};
    }
    std::vector<int> cycle{cycle_start};
    for (int u = cycle_end; u != cycle_start; u = parent[u]) {
        cycle.push_back(u);
    }
    std::reverse(cycle.begin() + 1, cycle.end());
    return cycle;
}

```

The `CycleFinder` class is the more general multigraph variant. Its cycles are returned as edge IDs rather than node lists, which makes the representation robust for parallel edges. In an undirected graph, the cycles found this way form a DFS cycle basis: one cycle for each non-tree edge. In a directed graph, this finds cycles created by back edges to nodes still on the recursion stack. Here `state[v]` is `-1` before discovery, `-2` after finishing, and otherwise the index where the node first appeared in the current edge stack; that index marks which stack suffix forms a cycle.

- `CycleFinder(n, directed)` constructs a graph of `n` nodes numbered `[0, n)`. The graph is directed if `directed` is true, or undirected otherwise.
- `add_edge(u, v)` adds an edge and returns its edge ID.
- `find_cycles(max_cycles, max_total_size)` returns simple cycles as vectors of edge IDs, stopping once either optional limit is reached.
- `cycle_nodes(edge_cycle)` converts one edge-ID cycle into a corresponding cyclic list of nodes.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the number of nodes.
- $O(1)$ per call to `add_edge()`.

- $O(\max(n, m + s))$ per call to `find_cycles()`, where n is the number of nodes, m is the number of edges, and s is the total size of the returned cycles.
- $O(k)$ per call to `cycle_nodes()`, where k is the number of edges in the cycle.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph.
- $O(n)$ auxiliary stack and heap space for `find_cycles()`, plus $O(s)$ for the returned cycles, where s is their total size.
- $O(k)$ for the vector returned by `cycle_nodes()`, where k is the number of edges in the cycle.

```
class CycleFinder {
    std::vector<std::pair<int, int>> edges;
    std::vector<std::vector<int>> adj;
    bool directed;

    void dfs_cycles(
        int u, int parent_edge, int max_cycles, int max_total_size, std::vector<int> &state,
        std::vector<int> &edge_stack, std::vector<std::vector<int>> &cycles, int &total_size
    ) const {
        if (static_cast<int>(cycles.size()) >= max_cycles || total_size >= max_total_size) {
            return;
        }
        // -1 = unvisited, -2 = finished, otherwise the node's index in the current edge stack.
        state[u] = static_cast<int>(edge_stack.size());
        for (int id : adj[u]) {
            if (!directed && id == parent_edge) {
                continue;
            }
            const auto &[eu, ev] = edges[id];
            int v = directed ? ev : eu ^ ev ^ u;
            if (state[v] >= 0) {
                std::vector<int> cycle{id};
                for (int i = state[v]; i < static_cast<int>(edge_stack.size()); i++) {
                    cycle.push_back(edge_stack[i]);
                }
                cycles.push_back(cycle);
                total_size += static_cast<int>(cycle.size());
                if (static_cast<int>(cycles.size()) >= max_cycles || total_size >= max_total_size) {
                    return;
                }
            } else if (state[v] == -1) {
                edge_stack.push_back(id);
                dfs_cycles(v, id, max_cycles, max_total_size, state, edge_stack, cycles, total_size);
                edge_stack.pop_back();
            }
        }
        state[u] = -2;
    }

public:
    CycleFinder(int n, bool directed) : adj(n), directed(directed) {}

    int add_edge(int u, int v) {
        int id = static_cast<int>(edges.size());
        edges.emplace_back(u, v);
        adj[u].push_back(id);
        if (!directed) {
            adj[v].push_back(id);
        }
        return id;
    }

    std::vector<std::vector<int>> find_cycles(
        int max_cycles = 1 << 30, int max_total_size = 1 << 30
    ) const {
        int n = static_cast<int>(adj.size());
```

```

std::vector<int> state(n, -1), edge_stack;
std::vector<std::vector<int>> cycles;
int total_size = 0;
for (int u = 0; u < n; u++) {
    if (state[u] == -1) {
        dfs_cycles(u, -1, max_cycles, max_total_size, state, edge_stack, cycles, total_size);
    }
}
return cycles;
}

std::vector<int> cycle_nodes(const std::vector<int> &edge_cycle) const {
    int size = static_cast<int>(edge_cycle.size());
    std::vector<int> nodes;
    if (size == 0) {
        return nodes;
    }
    if (size <= 2) {
        nodes.push_back(edges[edge_cycle[0]].first);
        if (size == 2) {
            nodes.push_back(edges[edge_cycle[0]].second);
        }
        return nodes;
    }
    for (int i = 0; i < size; i++) {
        const auto &[au, av] = edges[edge_cycle[i]];
        const auto &[bu, bv] = edges[edge_cycle[(i + 1) % size]];
        nodes.push_back((bu == au || bv == au) ? av : au);
    }
    return nodes;
}
};

```

Example Usage

```

#include <algorithm>
#include <cassert>
using namespace std;

int main() {
    {
        // 0---1    3---4
        // |  /
        // | /
        // 2
        vector<vector<int>> adj(5);
        auto add_edge = [&](int u, int v) {
            adj[u].push_back(v);
            adj[v].push_back(u);
        };
        add_edge(0, 1);
        add_edge(1, 2);
        add_edge(2, 0);
        add_edge(3, 4);
        // DFS encounters the triangle component before the acyclic 3-4 component.
        assert((find_cycle_undirected(adj) == vector<int>{0, 1, 2}));
    }
    {
        // 0 --> 1 --> 2    3 --> 4 --> 5
        //                  ^         |
        //                  |-----|
        vector<vector<int>> adj(6);
        adj[0].push_back(1);
        adj[1].push_back(2);
        adj[3].push_back(4);
        adj[4].push_back(5);
    }
}

```

```

adj[5].push_back(3);
assert((find_cycle_directed(adj) == vector<int>{3, 4, 5}));
assert(find_cycle_directed({{1}, {2}, {}}).empty());
}
{
    // 0---a---1
    // |      /
    // c    b
    // |    /
    // 2-----3
    CycleFinder g(4, false);
    int a = g.add_edge(0, 1);
    int b = g.add_edge(1, 2);
    int c = g.add_edge(2, 0);
    g.add_edge(2, 3);
    vector<vector<int>> cycles = g.find_cycles();
    assert(cycles.size() == 1);
    sort(cycles[0].begin(), cycles[0].end());
    assert((cycles[0] == vector<int>{a, b, c}));
    vector<int> nodes = g.cycle_nodes(cycles[0]);
    assert(nodes.size() == 3);
}
{
    // 0---a---1
    // 0---b---1
    CycleFinder g(2, false);
    int a = g.add_edge(0, 1);
    int b = g.add_edge(0, 1);
    // Parallel undirected edges form a 2-edge cycle; edge IDs make this representable.
    vector<vector<int>> cycles = g.find_cycles();
    assert(cycles.size() == 1);
    sort(cycles[0].begin(), cycles[0].end());
    assert((cycles[0] == vector<int>{a, b}));
}
{
    // 0 ---> 1
    // ^      /
    // |      /
    // 2 <-+
    CycleFinder g(3, true);
    g.add_edge(0, 1);
    g.add_edge(1, 2);
    g.add_edge(2, 0);
    assert(g.find_cycles(1).size() == 1);
}
return 0;
}

```

4.1.4 Tree Centers and Diameter

An unweighted tree possesses a center, centroid, and diameter. The centers are found by repeatedly peeling away the leaves of the tree until one or two nodes remain. The diameter relies on the fact that the farthest node from any start is always one endpoint of a diameter, so a second traversal from that endpoint measures the full diameter. The following functions apply to a global, bidirectionally pre-populated adjacency list `adj` which must form a valid tree with nodes numbered $[0, n)$, where n is `adj.size()`.

- `find_centers()` returns a vector of either one or two tree Jordan centers. The Jordan center of a tree is the set of all nodes with minimum eccentricity, that is, the set of all nodes where the maximum distance to all other nodes in the tree is minimal.
- `find_centroid(u = 0, p = -1)` returns the node where all of its subtrees have a size less than or equal to $n/2$, where n is the number of nodes in the tree.

- `diameter()` returns the maximum distance between any two nodes in the tree, using a well-known double depth-first search technique.

Time Complexity: $O(\max(n, m))$ per call to `find_centers()`, `find_centroid()`, and `diameter()`, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(n)$ auxiliary heap space for `find_centers()`, where n is the number of nodes.
- $O(n)$ auxiliary stack space for `find_centroid()` and `diameter()`.

```
#include <algorithm>
#include <utility>
#include <vector>

std::vector<std::vector<int>> adj;

std::vector<int> find_centers() {
    int n = static_cast<int>(adj.size());
    std::vector<int> leaves, degree(n);
    for (int i = 0; i < n; i++) {
        degree[i] = static_cast<int>(adj[i].size());
        if (degree[i] <= 1) {
            leaves.push_back(i);
        }
    }
    int removed = static_cast<int>(leaves.size());
    while (removed < n) {
        std::vector<int> nleaves;
        for (int u : leaves) {
            for (int v : adj[u]) {
                if (--degree[v] == 1) {
                    nleaves.push_back(v);
                }
            }
        }
        leaves = nleaves;
        removed += static_cast<int>(leaves.size());
    }
    return leaves;
}

// Returns the centroid index if found in this subtree, or -(subtree size) to propagate
// the size up to the parent so it can check the complementary component's size.
int find_centroid(int u = 0, int p = -1) {
    int n = static_cast<int>(adj.size());
    int count = 1;
    bool good_center = true;
    for (int v : adj[u]) {
        if (v == p) {
            continue;
        }
        int res = find_centroid(v, u);
        if (res >= 0) {
            return res;
        }
        int size = -res;
        good_center &= (size <= n / 2);
        count += size;
    }
    good_center &= (n - count <= n / 2);
    return good_center ? u : -count;
}

std::pair<int, int> dfs(int u, int p, int depth) {
    std::pair<int, int> res{depth, u};
    for (int v : adj[u]) {
```

```

    if (v != p) {
        res = std::max(res, dfs(v, u, depth + 1));
    }
}
return res;
}

int diameter() {
    int furthest_node = dfs(0, -1, 0).second;
    return dfs(furthest_node, -1, 0).first;
}

```

Example Usage

```

#include <cassert>
using namespace std;

void add_edge(int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
}

int main() {
    // 0---1---4---3
    //   |   |
    //   2   5
    adj.assign(6, {});
    add_edge(0, 1);
    add_edge(1, 2);
    add_edge(1, 4);
    add_edge(3, 4);
    add_edge(4, 5);
    vector<int> centers = find_centers();
    assert(centers.size() == 2 && centers[0] == 1 && centers[1] == 4);
    assert(find_centroid() == 4);
    assert(diameter() == 3);
    return 0;
}

```

4.1.5 Tree Rerooting DP

Given a tree DP that can answer one rooted query with a recursive `dfs(u, p)`, compute that same answer for every possible root. Rerooting DP first stores parent pointers, depths, preorder, subtree sizes, and entry/exit times. A bottom-up pass computes the usual child-to-parent DFS values, then a top-down pass uses prefix/suffix aggregation to compute the missing parent-to-child values. The prefix/suffix aggregation requires `combine` to be associative for a running time of $O(n)$, which is usually preferred, but a segment-tree/exclusive-combine variant can support a left-fold style interface in $O(n \log n)$.

- `RerootingTree(n)` constructs a tree with `n` nodes numbered $[0, n)$.
- `add_edge(u, v)` adds an undirected edge between nodes `u` and `v`.
- `build_rooted_tree(root = 0)` populates `parent`, `depth`, `subtree_size`, `order`, `tin`, and `tout` for the tree rooted at `root`. This must be called only after the added edges define a connected tree of size `n`.
- `is_ancestor(u, v)` returns whether `u` is an ancestor of `v` in the last rooted tree.
- `rerooting_dp(identity, combine, finalize, lift)` returns the value that `dfs(root, p)` would produce for every possible `root` with sentinel parent $p = -1$. This must be called only after the added edges define a connected tree of size `n`.

The aggregation scheme for `rerooting_dp()` is defined by the following pieces:

- `Data` is the return type of a normal fixed-root DFS call `dfs(u, p)`. It summarizes the component containing `u` after removing edge `u-p`, with all values measured from node `u`.
- `identity` is the neutral value for an empty component, which for all aggregate values `a` must satisfy `combine(a, identity) = combine(identity, a) = a`.
- `combine(a, b)` merges two independent neighbor contributions into the same node `u`, like two different child DFS returns.
- `finalize(acc, u)` finalizes node `u` after all neighbor contributions have been merged. This is the last line of a normal `dfs(u, p)` before returning.
- `lift(acc, u, p)` converts a finished `dfs(u, p)` value into the contribution that node `u` gives to adjacent node `p`. Here `p` is the neighbor treated as the parent/excluded side of the directed DFS state, not necessarily `u`'s parent in the initial root-0 traversal.

This covers problems such as subtree sizes, sum/max of distances from every root, and rerooted polynomial/hash-like aggregates, as long as sibling contributions can be combined associatively. The usage example shows how to compute the sum of distances from every root: `Data` stores `sum`, the total distance from `u` to all nodes in `dfs(u, p)`, and `nodes`, the node count in that component. The node count is needed because `lift(acc, u, p)` crosses one edge, so every represented node becomes 1 farther from `p` and `sum` increases by `nodes`.

Time Complexity:

- $O(1)$ per call to `add_edge()`.
- $O(n)$ per call to `build_rooted_tree()`, where n is the number of nodes.
- $O(1)$ per call to `is_ancestor()`.
- $O(n)$ per call to `rerooting_dp()`, assuming $O(1)$ callback functions.

Space Complexity:

- $O(n)$ for the tree, metadata, and DP arrays.
- $O(n)$ auxiliary stack space for the recursive search.

```
#include <stdint>
#include <vector>

class RerootingTree {
    int timer;

    void dfs_rooted(int u, int p) {
        parent[u] = p;
        tin[u] = timer++;
        order.push_back(u);
        subtree_size[u] = 1;
        for (int v : adj[u]) {
            if (v == p) {
                continue;
            }
            depth[v] = depth[u] + 1;
            dfs_rooted(v, u);
            subtree_size[u] += subtree_size[v];
        }
        tout[u] = timer;
    }

public:
    std::vector<std::vector<int>>> adj;
    std::vector<int> parent, depth, subtree_size, order, tin, tout;

    explicit RerootingTree(int n) : timer(0), adj(n) {}

    void add_edge(int u, int v) {
        adj[u].push_back(v);
        adj[v].push_back(u);
    }

    void build_rooted_tree(int root = 0) {
```

```

    int n = static_cast<int>(adj.size());
    parent.assign(n, -1);
    depth.assign(n, 0);
    subtree_size.assign(n, 0);
    order.clear();
    tin.assign(n, 0);
    tout.assign(n, 0);
    timer = 0;
    dfs_rooted(root, -1);
}

bool is_ancestor(int u, int v) const { return tin[u] <= tin[v] && tout[v] <= tout[u]; }

template<typename Data, typename Combine, typename Finalize, typename Lift>
std::vector<Data> rerooting_dp(
    const Data &identity, Combine combine, Finalize finalize, Lift lift
) {
    int n = static_cast<int>(adj.size());
    std::vector<Data> down(n, identity), outside(n, identity), answer(n, identity);
    build_rooted_tree();
    for (int i = n - 1; i >= 0; i--) {
        int u = order[i];
        Data acc = identity;
        for (int v : adj[u]) {
            if (parent[v] == u) {
                acc = combine(acc, lift(down[v], v, u));
            }
        }
        down[u] = finalize(acc, u);
    }
    for (int u : order) {
        std::vector<Data> vals;
        for (int v : adj[u]) {
            if (parent[v] == u) {
                vals.push_back(lift(down[v], v, u));
            } else {
                vals.push_back(outside[u]);
            }
        }
        std::vector<Data> suffix(vals.size() + 1, identity);
        for (int i = static_cast<int>(vals.size()) - 1; i >= 0; i--) {
            suffix[i] = combine(vals[i], suffix[i + 1]);
        }
        Data prefix = identity;
        for (int i = 0; i < static_cast<int>(adj[u].size()); i++) {
            int v = adj[u][i];
            if (parent[v] == u) {
                Data without_child = finalize(combine(prefix, suffix[i + 1]), u);
                outside[v] = lift(without_child, u, v);
            }
            prefix = combine(prefix, vals[i]);
        }
        answer[u] = finalize(prefix, u);
    }
    return answer;
}
};

```

Example Usage

```

#include <cassert>
#include <string>
using namespace std;

vector<int64_t> sum_distances_all_roots(RerootingTree &tree) {
    struct DistanceData {

```

```

    int64_t sum = 0;
    int nodes = 0;
};
auto combine = [](DistanceData a, DistanceData b) {
    return DistanceData{a.sum + b.sum, a.nodes + b.nodes};
};
auto finalize = [](DistanceData a, int) {
    a.nodes++;
    return a;
};
auto lift = [](DistanceData a, int, int) {
    a.sum += a.nodes;
    return a;
};
vector<DistanceData> dp = tree.rerooting_dp(DistanceData{}, combine, finalize, lift);
vector<int64_t> answer(dp.size());
for (int i = 0; i < static_cast<int>(dp.size()); i++) {
    answer[i] = dp[i].sum;
}
return answer;
}

int main() {
    // 0---1---2
    //   |
    //   3---4
    RerootingTree tree(5);
    tree.add_edge(0, 1);
    tree.add_edge(1, 2);
    tree.add_edge(1, 3);
    tree.add_edge(3, 4);
    tree.build_rooted_tree();
    assert(tree.parent[4] == 3);
    assert(tree.depth[4] == 3);
    assert(tree.subtree_size[1] == 4);
    assert(tree.is_ancestor(1, 4));
    assert(!tree.is_ancestor(3, 2));
    vector<int64_t> distances = sum_distances_all_roots(tree);
    assert((distances == vector<int64_t>{8, 5, 8, 6, 9}));

    // Associative contributions retain adjacency-list order even when combine is not commutative.
    RerootingTree ordered(3);
    ordered.add_edge(1, 2);
    ordered.add_edge(0, 1); // adj[1] is [2, 0], while 0 is its parent in the initial rooting.
    auto combine = [](string a, const string &b) { return a + b; };
    auto finalize = [](string a, int u) { return a + to_string(u); };
    auto lift = [](const string &a, int, int) { return a; };
    vector<string> traversal = ordered.rerooting_dp(string{}, combine, finalize, lift);
    assert(traversal[1] == "201");
    return 0;
}

```

4.1.6 Subtree Small-to-Large Merging

Given a rooted tree, compute a mutable summary for every subtree by merging child summaries into one kept summary. Small-to-large merging always keeps the largest child summary and absorbs smaller summaries into it. Each stored item can move only into summaries at least twice as large as before, so the total number of moved items is logarithmic per item.

This pattern is often called DSU on tree, although it is unrelated to the disjoint-set union data structure used by Kruskal's algorithm. It is useful for subtree queries whose answer can be maintained by inserting all items from child summaries into one mutable summary, such as distinct colors, frequency modes, or other multiset-style statistics.

- `subtree_small_to_large_merge(adj, root, init, merge, answer)` visits the tree rooted at `root`, given by the bidirectional adjacency list `adj`, which must form a valid tree. The callback `init(u)` returns the initial summary for node `u`; `merge(big, small)` absorbs one child summary into another; and `answer(u, summary)` records the final subtree summary for `u`.

Summary objects must support `.size()`, and `merge(big, small)` should iterate over `small` and insert into `big`.

Time Complexity: $O(n \log n)$ expected per call when `merge()` does expected $O(1)$ work per moved item, where n is the number of nodes.

Space Complexity: $O(n)$ auxiliary stack space and $O(n \log n)$ auxiliary heap space in the worst case, since absorbed summaries remain stored.

```
#include <algorithm>
#include <cassert>
#include <unordered_map>
#include <vector>

template<typename Summary, typename Init, typename Merge, typename Answer>
int dfs_subtree_small_to_large(
    int u, int p, const std::vector<std::vector<int>>> &adj, std::vector<Summary> &summary,
    Init &init, Merge &merge, Answer &answer
) {
    int keep = u;
    summary[u] = init(u);
    for (int v : adj[u]) {
        if (v == p) {
            continue;
        }
        int child = dfs_subtree_small_to_large(v, u, adj, summary, init, merge, answer);
        if (summary[keep].size() < summary[child].size()) {
            std::swap(keep, child);
        }
        merge(summary[keep], summary[child]);
    }
    answer(u, summary[keep]);
    return keep;
}

template<typename Init, typename Merge, typename Answer>
void subtree_small_to_large_merge(
    const std::vector<std::vector<int>>> &adj, int root, Init init, Merge merge, Answer answer
) {
    int n = static_cast<int>(adj.size());
    assert(0 <= root && root < n);
    using Summary = decltype(init(0));
    std::vector<Summary> summary(n);
    dfs_subtree_small_to_large(root, -1, adj, summary, init, merge, answer);
}
```

Example Usage

```
int main() {
    //      0:A
    //      /  |
    //     1:B  2:A
    //    / |  |
    //   3:C 4:B 5:C
    std::vector<std::vector<int>>> adj(6);
    auto add_edge = [&](int u, int v) {
        adj[u].push_back(v);
        adj[v].push_back(u);
    };
    add_edge(0, 1);
    add_edge(0, 2);
```

```

add_edge(1, 3);
add_edge(1, 4);
add_edge(2, 5);

std::vector<int> color{0, 1, 0, 2, 1, 2};
std::vector<int> ans(6);
auto init = [&](int u) {
    std::unordered_map<int, int> freq;
    freq[color[u]] = 1;
    return freq;
};
auto merge = [](std::unordered_map<int, int> &big, const std::unordered_map<int, int> &small) {
    for (const auto &[c, cnt] : small) {
        big[c] += cnt;
    }
};
auto answer = [&](int u, const std::unordered_map<int, int> &freq) {
    ans[u] = static_cast<int>(freq.size());
};
subtree_small_to_large_merge(adj, 0, init, merge, answer);
assert((ans == std::vector<int>{3, 2, 2, 1, 1, 1}));
return 0;
}

```

4.1.7 Centroid Decomposition

Given a tree, build its centroid decomposition. A centroid is a node whose removal splits the current subtree into connected components of size at most half of that subtree. Recursively choosing centroids creates a decomposition tree of height $O(\log n)$, which is useful for queries based on distances to marked nodes, nearest special node, and other “path through a centroid” problems.

- `CentroidDecomposition(adj)` builds the centroid decomposition of the tree given by a bidirectional adjacency list `adj`, which must form a valid tree and whose indices represent the nodes.
- `parent(u)` returns the parent of node `u` in the centroid tree, or -1 if `u` is the root centroid.

Time Complexity: $O(n \log n)$ for construction, where n is the number of nodes.

Space Complexity:

- $O(n)$ for storage of the centroid tree.
- $O(n)$ auxiliary stack space for the recursive searches.

```

#include <vector>

class CentroidDecomposition {
    std::vector<std::vector<int>>> adj;
    std::vector<int> subtree_size, par;
    std::vector<char> removed;

    int get_size(int u, int p) {
        subtree_size[u] = 1;
        for (int v : adj[u]) {
            if (v != p && !removed[v]) {
                subtree_size[u] += get_size(v, u);
            }
        }
        return subtree_size[u];
    }

    int get_centroid(int u, int p, int total) {
        for (int v : adj[u]) {
            if (v != p && !removed[v] && subtree_size[v] > total / 2) {
                return get_centroid(v, u, total);
            }
        }
    }
}

```

```

    }
    return u;
}

void decompose(int entry, int parent) {
    int total = get_size(entry, -1);
    int centroid = get_centroid(entry, -1, total);
    par[centroid] = parent;
    removed[centroid] = true;
    for (int v : adj[centroid]) {
        if (!removed[v]) {
            decompose(v, centroid);
        }
    }
}

public:
    explicit CentroidDecomposition(const std::vector<std::vector<int>> &adj)
        : adj(adj), subtree_size(adj.size()), par(adj.size(), -1), removed(adj.size()) {
        if (adj.size() > 0) {
            decompose(0, -1);
        }
    }

    int parent(int u) const { return par[u]; }
};

```

Example Usage

```

#include <algorithm>
#include <cassert>
#include <climits>
using namespace std;

// Example-only distance helper; use LCA preprocessing for O(1) distances in a full solution.
int tree_dist(const vector<vector<int>> &adj, int u, int target, int p = -1) {
    if (u == target) {
        return 0;
    }
    for (int v : adj[u]) {
        if (v != p) {
            int d = tree_dist(adj, v, target, u);
            if (d != -1) {
                return d + 1;
            }
        }
    }
    return -1;
}

int main() {
    // 0---1---3---5---6
    //   |   |
    //   2   4
    vector<vector<int>> adj(7);
    auto add_edge = [&](int u, int v) {
        adj[u].push_back(v);
        adj[v].push_back(u);
    };
    add_edge(0, 1);
    add_edge(1, 2);
    add_edge(1, 3);
    add_edge(3, 4);
    add_edge(3, 5);
    add_edge(5, 6);
    CentroidDecomposition tree(adj);
}

```

```

vector<int> parent(7);
for (int u = 0; u < 7; u++) {
    parent[u] = tree.parent(u);
}
assert((parent == vector<int>{1, 3, 1, -1, 3, 3, 5}));
assert(tree.parent(3) == -1);
assert(tree.parent(1) == 3);
assert(tree.parent(5) == 3);

const int INF = INT_MAX / 2;
vector<int> best(7, INF);
auto mark = [&](int u) {
    for (int c = u; c != -1; c = tree.parent(c)) {
        best[c] = min(best[c], tree_dist(adj, u, c));
    }
};
auto nearest_marked = [&](int u) {
    int answer = INF;
    for (int c = u; c != -1; c = tree.parent(c)) {
        answer = min(answer, best[c] + tree_dist(adj, u, c));
    }
    return answer;
};
mark(0);
mark(6);
assert(nearest_marked(2) == 2);
assert(nearest_marked(4) == 3);
assert(nearest_marked(5) == 1);
return 0;
}

```

4.1.8 Prufer Code

Encode and decode labeled trees using Prufer codes. A Prufer code is a sequence of length $n - 2$ that uniquely represents a labeled tree on nodes $0, 1, \dots, n - 1$. It is useful for counting labeled trees, generating test cases, and converting between trees and compact sequences. Encoding repeatedly removes a leaf and records its neighbor until two nodes remain; decoding reverses this, deducing each removed leaf from the node degrees implied by the code. Both functions choose the smallest available leaf at each step. This makes the implementation deterministic and matches the usual textbook convention.

- `encode_prufer()` returns the Prufer code for the global, bidirectionally pre-populated adjacency list `adj`, which must form a valid tree and whose indices represent the nodes.
- `decode_prufer()` takes a Prufer code and returns the corresponding tree edges.

The encoded tree must contain at least two nodes, since the empty code represents the unique tree on two nodes.

Time Complexity: $O(n \log n)$ per call to `encode_prufer()` or `decode_prufer()`, where n is the number of nodes in the tree.

Space Complexity: $O(n)$ auxiliary for `encode_prufer()` and `decode_prufer()`.

```

#include <set>
#include <utility>
#include <vector>

std::vector<int> encode_prufer(const std::vector<std::vector<int>>& adj) {
    int n = static_cast<int>(adj.size());
    std::vector<int> degree(n), parent(n, -1), code;
    std::set<int> leaves;
    if (n <= 2) {
        return code;
    }
    // Root at node n - 1, which is guaranteed to survive smallest-leaf removal (it is never the

```

```

// smallest leaf while other nodes remain). Rooting elsewhere risks recording parent[root] == -2
// if the chosen root itself gets stripped to a leaf (e.g. when n - 1 happens to be a leaf).
int root = n - 1;
for (int u = 0; u < n; u++) {
    degree[u] = static_cast<int>(adj[u].size());
    if (degree[u] == 1) {
        leaves.insert(u);
    }
}
std::vector<int> st(1, root);
parent[root] = -2;
while (!st.empty()) {
    int u = st.back();
    st.pop_back();
    for (int v : adj[u]) {
        if (parent[v] == -1) {
            parent[v] = u;
            st.push_back(v);
        }
    }
}
for (int i = 0; i < n - 2; i++) {
    int leaf = *leaves.begin();
    leaves.erase(leaves.begin());
    int p = parent[leaf];
    code.push_back(p);
    if (--degree[p] == 1) {
        leaves.insert(p);
    }
}
return code;
}

std::vector<std::pair<int, int>> decode_prufer(const std::vector<int> &code) {
    int n = static_cast<int>(code.size()) + 2;
    std::vector<int> degree(n, 1);
    std::set<int> leaves;
    std::vector<std::pair<int, int>> edges;
    for (int c : code) {
        degree[c]++;
    }
    for (int u = 0; u < n; u++) {
        if (degree[u] == 1) {
            leaves.insert(u);
        }
    }
    for (int u : code) {
        int leaf = *leaves.begin();
        leaves.erase(leaves.begin());
        edges.emplace_back(leaf, u);
        if (--degree[u] == 1) {
            leaves.insert(u);
        }
    }
    edges.emplace_back(*leaves.begin(), *leaves.rbegin());
    return edges;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // 0
    // 1

```

```

// 1---3---2
//   |
//   4
vector<vector<int>>> adj(5);
auto add_edge = [&](int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
};
add_edge(0, 3);
add_edge(1, 3);
add_edge(2, 3);
add_edge(3, 4);
vector<int> code = encode_prufer(adj);
assert(code.size() == 3 && code[0] == 3 && code[1] == 3 && code[2] == 3);
vector<pair<int, int>>> edges = decode_prufer(code);
assert(edges.size() == 4);
return 0;
}

```

4.1.9 Eulerian Trails (Hierholzer)

An Eulerian trail is a path in a graph which contains every edge exactly once. An Eulerian cycle or circuit is an Eulerian trail which begins and ends on the same node. A directed graph has an Eulerian trail when all nonzero-degree nodes belong to one connected part of the underlying graph, and either every node has equal in-degree and out-degree or exactly one node has one extra outgoing edge and exactly one node has one extra incoming edge.

Hierholzer's algorithm walks unused edges until stuck, then backtracks to splice each closed detour into the final trail. For a directed graph known to have a trail from `start`, the core algorithm can simply consume outgoing edges by popping them from a local copy of the adjacency list.

- `known_eulerian_path_directed(adj, start)` returns a node trail using every directed edge in `adj` exactly once, assuming such a trail exists and begins at `start`.
- `eulerian_path_directed(adj, start = -1)` returns a node trail using every directed edge in `adj` exactly once, or an empty vector if no such trail exists. If `start = -1`, a valid start is chosen automatically; otherwise the trail must begin at `start`.

Parallel directed edges are supported by storing duplicate neighbors in `adj`. Since these functions return only nodes, use `EulerianGraph` below when edge IDs are needed.

Time Complexity: $O(\max(n, m))$ per call to `known_eulerian_path_directed()` and `eulerian_path_directed()`, where n is the number of nodes and m is the number of edges.

Space Complexity: $O(\max(n, m))$ auxiliary for `known_eulerian_path_directed()` and `eulerian_path_directed()`.

```

#include <algorithm>
#include <cassert>
#include <utility>
#include <vector>

std::vector<int> known_eulerian_path_directed(std::vector<std::vector<int>>> adj, int start) {
    assert(0 <= start && start < static_cast<int>(adj.size()));
    std::vector<int> st{start}, path;
    while (!st.empty()) {
        int u = st.back();
        if (adj[u].empty()) {
            path.push_back(u);
            st.pop_back();
        } else {
            st.push_back(adj[u].back());
            adj[u].pop_back();
        }
    }
}

```

```

    }
    std::reverse(path.begin(), path.end());
    return path;
}

std::vector<int> eulerian_path_directed(const std::vector<std::vector<int>>> &adj, int start = -1) {
    int n = static_cast<int>(adj.size()), m = 0;
    if (n == 0) {
        return {};
    }
    assert(-1 <= start && start < n);
    std::vector<int> indeg(n), outdeg(n);
    for (int u = 0; u < n; u++) {
        outdeg[u] = static_cast<int>(adj[u].size());
        m += outdeg[u];
        for (int v : adj[u]) {
            indeg[v]++;
        }
    }
    if (m == 0) {
        return {start == -1 ? 0 : start};
    }
    int source = -1, sources = 0, sinks = 0;
    for (int u = 0; u < n; u++) {
        int diff = outdeg[u] - indeg[u];
        if (diff == 1) {
            source = u;
            sources++;
        } else if (diff == -1) {
            sinks++;
        } else if (diff != 0) {
            return {};
        }
    }
    if (!(sources == 0 && sinks == 0) || (sources == 1 && sinks == 1)) {
        return {};
    }
    if (start == -1) {
        if (source != -1) {
            start = source;
        } else {
            for (int u = 0; u < n && start == -1; u++) {
                if (outdeg[u] > 0) {
                    start = u;
                }
            }
        }
    }
    if ((source != -1 && start != source) || (source == -1 && outdeg[start] == 0 && m > 0)) {
        return {};
    }
    std::vector<int> path = known_eulerian_path_directed(adj, start);
    return static_cast<int>(path.size()) == m + 1 ? path : std::vector<int>();
}

```

The `EulerianGraph` class is the more general edge-ID variant. An undirected graph has an Eulerian trail when all nonzero-degree nodes are connected and either zero or two nodes have odd degree. This class stores and returns edge IDs, which supports directed graphs, undirected graphs, and multigraphs: parallel edges are distinct because each edge receives its own ID.

- `EulerianGraph(n, directed)` constructs a graph of `n` nodes numbered $[0, n)$. The graph is directed if `directed` is true, or undirected otherwise.
- `add_edge(u, v)` adds an edge and returns its edge ID.
- `eulerian_path(start = -1)` returns an `EulerianTrail` containing every edge exactly once, or an invalid result with `start = -1` if no such trail exists. In a valid result, `start` is the first node, `edges` lists edge IDs in traversal order, and `nodes` lists the visited nodes in order. If the input `start = -1`, a valid start is chosen

automatically. For a valid result, `is_cycle()` returns whether the valid trail uses at least one edge and begins and ends on the same node.

Time Complexity: $O(\max(n, m))$ per call to `eulerian_path()`, where n is the number of nodes and m is the number of edges.

Space Complexity: $O(\max(n, m))$ for graph storage and auxiliary arrays.

```
class EulerianGraph {
    std::vector<std::pair<int, int>> edges;
    std::vector<std::vector<int>> adj;
    bool directed;

    bool valid_degrees(int start) const {
        int n = static_cast<int>(adj.size());
        if (directed) {
            std::vector<int> indeg(n), outdeg(n);
            for (const auto &[eu, ev] : edges) {
                outdeg[eu]++;
                indeg[ev]++;
            }
            int source = -1, sink = -1;
            for (int u = 0; u < n; u++) {
                int diff = outdeg[u] - indeg[u];
                if (diff == 1) {
                    if (source != -1) {
                        return false;
                    }
                    source = u;
                } else if (diff == -1) {
                    if (sink != -1) {
                        return false;
                    }
                    sink = u;
                } else if (diff != 0) {
                    return false;
                }
            }
            if ((source == -1) != (sink == -1)) {
                return false;
            }
            if (start != -1) {
                return source == -1 ? outdeg[start] > 0 : start == source;
            }
            return true;
        }
        std::vector<int> degree(n);
        for (const auto &[eu, ev] : edges) {
            degree[eu]++;
            degree[ev]++;
        }
        std::vector<int> odd;
        for (int u = 0; u < n; u++) {
            if (degree[u] % 2 == 1) {
                odd.push_back(u);
            }
        }
        if (!odd.empty() && odd.size() != 2) {
            return false;
        }
        if (start != -1) {
            return odd.empty() ? degree[start] > 0 : (start == odd[0] || start == odd[1]);
        }
        return true;
    }

    int choose_start() const {
        int n = static_cast<int>(adj.size());
```



```

    if (directed) {
        std::vector<int> indeg(n), outdeg(n);
        for (const auto &[eu, ev] : edges) {
            outdeg[eu]++;
            indeg[ev]++;
        }
        for (int u = 0; u < n; u++) {
            if (outdeg[u] - indeg[u] == 1) {
                return u;
            }
        }
        for (int u = 0; u < n; u++) {
            if (outdeg[u] > 0) {
                return u;
            }
        }
    } else {
        std::vector<int> degree(n);
        for (const auto &[eu, ev] : edges) {
            degree[eu]++;
            degree[ev]++;
        }
        for (int u = 0; u < n; u++) {
            if (degree[u] % 2 == 1) {
                return u;
            }
        }
        for (int u = 0; u < n; u++) {
            if (degree[u] > 0) {
                return u;
            }
        }
    }
    return 0;
}

std::vector<int> build_nodes(int start, const std::vector<int> &trail_edges) const {
    std::vector<int> nodes{start};
    int u = start;
    for (int id : trail_edges) {
        const auto &[eu, ev] = edges[id];
        u = directed ? ev : eu ^ ev ^ u;
        nodes.push_back(u);
    }
    return nodes;
}

public:
    struct EulerianTrail {
        int start = -1;
        std::vector<int> edges, nodes;

        bool is_cycle() const {
            return start != -1 && nodes.size() > 1 && nodes.front() == nodes.back();
        }
    };

    EulerianGraph(int n, bool directed) : adj(n), directed(directed) {}

    int add_edge(int u, int v) {
        int id = static_cast<int>(edges.size());
        edges.emplace_back(u, v);
        adj[u].push_back(id);
        if (!directed) {
            adj[v].push_back(id);
        }
        return id;
    }
}

```

```

EulerianTrail eulerian_path(int start = -1) const {
    int n = static_cast<int>(adj.size());
    int m = static_cast<int>(edges.size());
    if (n == 0) {
        return EulerianTrail{};
    }
    assert(-1 <= start && start < n);
    if (m == 0) {
        int s = (start == -1 ? 0 : start);
        return EulerianTrail{s, {}, {s}};
    }
    if (!valid_degrees(start)) {
        return EulerianTrail{};
    }
    if (start == -1) {
        start = choose_start();
    }
    std::vector<char> used(m, false);
    std::vector<int> ptr(n), node_stack{start}, edge_stack{-1}, trail_edges;
    while (!node_stack.empty()) {
        int u = node_stack.back();
        while (ptr[u] < static_cast<int>(adj[u].size()) && used[adj[u][ptr[u]]]) {
            ptr[u]++;
        }
        if (ptr[u] == static_cast<int>(adj[u].size())) {
            if (edge_stack.back() != -1) {
                trail_edges.push_back(edge_stack.back());
            }
            node_stack.pop_back();
            edge_stack.pop_back();
        } else {
            int id = adj[u][ptr[u]++];
            used[id] = true;
            const auto &[eu, ev] = edges[id];
            int v = directed ? ev : eu ^ ev ^ u;
            node_stack.push_back(v);
            edge_stack.push_back(id);
        }
    }
    if (static_cast<int>(trail_edges.size()) != m) {
        return EulerianTrail{};
    }
    std::reverse(trail_edges.begin(), trail_edges.end());
    return EulerianTrail{start, trail_edges, build_nodes(start, trail_edges)};
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    {
        // 3 <-- 0 --> 1
        //      ^   |
        //      \   v
        //      +--2
        vector<vector<int>> adj(4);
        adj[0].push_back(1);
        adj[1].push_back(2);
        adj[2].push_back(0);
        adj[0].push_back(3);
        assert((known_eulerian_path_directed(adj, 0) == vector<int>{0, 1, 2, 0, 3}));
        vector<int> path = eulerian_path_directed(adj);
    }
}

```

```

    assert((path == vector<int>{0, 1, 2, 0, 3}));
    assert(eulerian_path_directed(adj, 1).empty());
}
{
    // 0 <--- 2
    // |     ^
    // |    /
    // v /
    // 1 ----> 3 ----> 4
    // ^-----|
    EulerianGraph g(5, true);
    g.add_edge(0, 1);
    g.add_edge(1, 2);
    g.add_edge(2, 0);
    g.add_edge(1, 3);
    g.add_edge(3, 4);
    g.add_edge(4, 1);
    auto trail = g.eulerian_path(0);
    assert(trail.edges.size() == 6);
    assert((trail.nodes == vector<int>{0, 1, 3, 4, 1, 2, 0}));
    assert(trail.is_cycle());
}
{
    // 0 ---- 2
    // |     /
    // |    /
    // | /
    // 1 ---- 3 ---- 4
    // |-----|
    EulerianGraph g(5, false);
    g.add_edge(0, 1);
    g.add_edge(1, 2);
    g.add_edge(2, 0);
    g.add_edge(1, 3);
    g.add_edge(3, 4);
    g.add_edge(4, 1);
    auto trail = g.eulerian_path();
    vector<int> used_edges = trail.edges;
    sort(used_edges.begin(), used_edges.end());
    assert((used_edges == vector<int>{0, 1, 2, 3, 4, 5}));
    assert(trail.nodes.size() == 7);
    assert(trail.is_cycle());
}
{
    EulerianGraph g(2, false);
    int a = g.add_edge(0, 1);
    int b = g.add_edge(0, 1);
    auto trail = g.eulerian_path(0);
    assert(trail.is_cycle());
    assert((trail.nodes == vector<int>{0, 1, 0}));
    assert((trail.edges == vector<int>{a, b} || trail.edges == vector<int>{b, a}));
}
{
    EulerianGraph g(0, false);
    assert(g.eulerian_path().start == -1);
}
{
    EulerianGraph g(3, true);
    g.add_edge(0, 1);
    g.add_edge(1, 2);
    auto trail = g.eulerian_path();
    assert(trail.start == 0);
    assert((trail.nodes == vector<int>{0, 1, 2}));
    assert(!trail.is_cycle());
}
return 0;
}

```

4.2 Connectivity

4.2.1 Connected Components

Given an undirected graph, label every node with its connected component and collect the nodes in each component. This is the standard graph reachability primitive behind many larger routines: run one depth-first search from every unseen node, and all nodes visited by that search belong to the same component. For directed graphs, use a strongly connected components algorithm instead.

- `connected_components()` populates the global arrays `comp_id[]` (component ID for each node) and `components[]` (nodes for each component ID) given a global, bidirectionally pre-populated adjacency list `adj` whose indices represent the nodes.

Time Complexity: $O(\max(n, m))$ per call, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary stack space.

```
#include <vector>

std::vector<std::vector<int>> adj;
std::vector<std::vector<int>> components;
std::vector<int> comp_id;

void dfs(int u, int id) {
    comp_id[u] = id;
    components[id].push_back(u);
    for (int v : adj[u]) {
        if (comp_id[v] == -1) {
            dfs(v, id);
        }
    }
}

void connected_components() {
    int n = static_cast<int>(adj.size());
    components.clear();
    comp_id.assign(n, -1);
    for (int u = 0; u < n; u++) {
        if (comp_id[u] == -1) {
            components.push_back({});
            dfs(u, static_cast<int>(components.size()) - 1);
        }
    }
}
```

Example Usage

```
#include <cassert>
using namespace std;

void add_edge(int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
}

int main() {
    // 0---1---2    3---4
    adj.assign(6, {});
    add_edge(0, 1);
    add_edge(1, 2);
}
```

```

add_edge(3, 4);
connected_components();
assert((components == vector<vector<int>>{{0, 1, 2}, {3, 4}, {5}}));
assert(components.size() == 3);
assert(comp_id[0] == comp_id[2]);
assert(comp_id[0] != comp_id[3]);
assert(comp_id[5] == 2);
return 0;
}

```

4.2.2 Strongly Connected Components (Kosaraju)

Given a directed graph, determine the strongly connected components (SCCs) using the Kosaraju-Sharir algorithm. A strongly connected component is a maximal set of nodes where every node can reach every other node. Condensing each SCC into one node produces a directed acyclic graph. The algorithm runs two passes of depth-first search: the first records the order in which nodes finish, and the second explores the transposed graph in reverse finish order, with each search collecting exactly one component.

- `KosarajuSCC(n = 0)` constructs a directed graph of `n` nodes numbered $[0, n)$.
- `add_edge(u, v)` adds the directed edge from `u` to `v`.
- `build_scc()` computes the strongly connected components.
- `components()` returns the strongly connected components from the last `build_scc()` call.
- `component_id(v)` returns the component ID containing node `v`. Component IDs are in topological order: for every edge from component a to a different component b , $a < b$.

Time Complexity: $O(\max(n, m))$ per call to `build_scc()`, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, reverse graph, SCCs, and component IDs.
- $O(n)$ auxiliary stack space.

```

#include <algorithm>
#include <vector>

class KosarajuSCC {
    std::vector<std::vector<int>> adj, rev, scc;
    std::vector<int> component;
    std::vector<char> visited;

    void dfs(const std::vector<std::vector<int>> &g, std::vector<int> &res, int u) {
        visited[u] = true;
        for (int v : g[u]) {
            if (!visited[v]) {
                dfs(g, res, v);
            }
        }
        res.push_back(u);
    }

    void dfs_component(int u, int id, std::vector<int> &res) {
        visited[u] = true;
        component[u] = id;
        res.push_back(u);
        for (int v : rev[u]) {
            if (!visited[v]) {
                dfs_component(v, id, res);
            }
        }
    }
}

```

```

public:
KosarajuSCC(int n = 0) : adj(n), rev(n) {}

void add_edge(int u, int v) {
    adj[u].push_back(v);
    rev[v].push_back(u);
}

void build_scc() {
    int n = static_cast<int>(adj.size());
    visited.assign(n, false);
    std::vector<int> order;
    for (int i = 0; i < n; i++) {
        if (!visited[i]) {
            dfs(adj, order, i);
        }
    }
    std::reverse(order.begin(), order.end());
    visited.assign(n, false);
    component.assign(n, -1);
    scc.clear();
    for (int u : order) {
        if (visited[u]) {
            continue;
        }
        std::vector<int> comp_nodes;
        dfs_component(u, static_cast<int>(scc.size()), comp_nodes);
        scc.push_back(comp_nodes);
    }
}

const std::vector<std::vector<int>> &components() const { return scc; }
int component_id(int v) const { return component[v]; }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // 0 ----> 1 ----> 2 <---> 3
    // ^      / |      |      ^
    // |      / |      |      |
    // |      / |      |      v
    // | v    v      v      7
    // 4 ----> 5 <---> 6 <---> /
    KosarajuSCC g(8);
    g.add_edge(0, 1);
    g.add_edge(1, 2);
    g.add_edge(1, 4);
    g.add_edge(1, 5);
    g.add_edge(2, 3);
    g.add_edge(2, 6);
    g.add_edge(3, 2);
    g.add_edge(3, 7);
    g.add_edge(4, 0);
    g.add_edge(4, 5);
    g.add_edge(5, 6);
    g.add_edge(6, 5);
    g.add_edge(7, 3);
    g.add_edge(7, 6);
    g.build_scc();
    // SCC condensation DAG:
    // {0,1,4} -> {2,3,7} -> {5,6}
    //      \-----^

```

```

assert(g.components().size() == 3);
vector<vector<int>> components = g.components();
for (auto &component : components) {
    sort(component.begin(), component.end());
}
sort(components.begin(), components.end());
assert((components == vector<vector<int>>{{0, 1, 4}, {2, 3, 7}, {5, 6}}));
assert(g.component_id(0) == g.component_id(1) && g.component_id(1) == g.component_id(4));
assert(g.component_id(2) == g.component_id(3) && g.component_id(3) == g.component_id(7));
assert(g.component_id(5) == g.component_id(6));
assert(g.component_id(0) != g.component_id(2) && g.component_id(2) != g.component_id(5));
return 0;
}

```

4.2.3 Strongly Connected Components (Tarjan)

Given a directed graph, determine the strongly connected components (SCCs) using Tarjan's algorithm. A strongly connected component is a maximal set of nodes where every node can reach every other node. Condensing each SCC into one node produces a directed acyclic graph. A single depth-first search keeps visited nodes on a stack and tracks each node's low-link, the smallest entry time reachable from its subtree; a node whose low-link equals its own entry time roots a component, which is popped off the stack in one piece.

- `TarjanSCC(n = 0)` constructs a directed graph of `n` nodes numbered $[0, n)$.
- `add_edge(u, v)` adds the directed edge from `u` to `v`.
- `build_scc()` computes the strongly connected components.
- `components()` returns the strongly connected components from the last `build_scc()` call.
- `component_id(v)` returns the component ID containing node `v`. Component IDs are in reverse topological order: for every edge from component a to a different component b , $a > b$.

Time Complexity: $O(\max(n, m))$ per call to `build_scc()`, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, SCCs, and component IDs.
- $O(n)$ auxiliary stack space.

```

#include <algorithm>
#include <climits>
#include <vector>

class TarjanSCC {
    static const int INF = INT_MAX / 2;
    std::vector<std::vector<int>> adj, scc;
    std::vector<int> component, active, lowlink;
    std::vector<char> visited;
    int timer;

    void dfs(int u) {
        lowlink[u] = timer++;
        visited[u] = true;
        active.push_back(u);
        bool is_component_root = true;
        for (int v : adj[u]) {
            if (!visited[v]) {
                dfs(v);
            }
            if (lowlink[u] > lowlink[v]) {
                lowlink[u] = lowlink[v];
                is_component_root = false;
            }
        }
        if (is_component_root) {
            scc.push_back(active);
            component.push_back(scc.size() - 1);
            active.clear();
        }
    }
}

```

```

    if (!is_component_root) {
        return;
    }
    std::vector<int> comp_nodes;
    int id = static_cast<int>(scc.size());
    int v;
    do {
        v = active.back();
        active.pop_back();
        lowlink[v] = INF; // marks v as removed from the active stack
        component[v] = id;
        comp_nodes.push_back(v);
    } while (u != v);
    scc.push_back(comp_nodes);
}

public:
TarjanSCC(int n = 0) : adj(n) {}

void add_edge(int u, int v) { adj[u].push_back(v); }

void build_scc() {
    int n = static_cast<int>(adj.size());
    scc.clear();
    component.assign(n, -1);
    active.clear();
    lowlink.assign(n, 0);
    visited.assign(n, false);
    timer = 0;
    for (int i = 0; i < n; i++) {
        if (!visited[i]) {
            dfs(i);
        }
    }
}

const std::vector<std::vector<int>>> &components() const { return scc; }
int component_id(int v) const { return component[v]; }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // 0 ----> 1 ----> 2 <----> 3
    // ~    / |      |      ^
    // |    / |      |      |
    // |    / |      |      v
    // | v    v      v      7
    // 4 ----> 5 <----> 6 <----/
    TarjanSCC g(8);
    g.add_edge(0, 1);
    g.add_edge(1, 2);
    g.add_edge(1, 4);
    g.add_edge(1, 5);
    g.add_edge(2, 3);
    g.add_edge(2, 6);
    g.add_edge(3, 2);
    g.add_edge(3, 7);
    g.add_edge(4, 0);
    g.add_edge(4, 5);
    g.add_edge(5, 6);
    g.add_edge(6, 5);
    g.add_edge(7, 3);
}

```



```

g.add_edge(7, 6);
g.build_scc();
// SCC condensation DAG:
// {0,1,4} -> {2,3,7} -> {5,6}
//      \-----^
assert(g.components().size() == 3);
vector<vector<int>> components = g.components();
for (auto &component : components) {
    sort(component.begin(), component.end());
}
sort(components.begin(), components.end());
assert((components == vector<vector<int>>{{0, 1, 4}, {2, 3, 7}, {5, 6}}));
assert(g.component_id(0) == g.component_id(1) && g.component_id(1) == g.component_id(4));
assert(g.component_id(2) == g.component_id(3) && g.component_id(3) == g.component_id(7));
assert(g.component_id(5) == g.component_id(6));
assert(g.component_id(0) != g.component_id(2) && g.component_id(2) != g.component_id(5));
return 0;
}

```

4.2.4 Dominator Tree (Lengauer-Tarjan)

Given a directed graph and a start node, find the immediate dominator of every node reachable from the start. A node u dominates node v if every path from the start to v passes through u . The immediate dominator $\text{idom}[v]$ is the closest strict dominator of v ; these parent links form the dominator tree rooted at the start. Dominator trees are useful in control-flow graphs, program analysis, and reachability problems with unavoidable checkpoints.

Lengauer-Tarjan's algorithm numbers nodes by DFS order, computes each reachable node's semidominator using a disjoint-set structure with path compression, and then resolves immediate dominators from semidominator buckets.

- `Dominators(n)` constructs a directed graph of n nodes numbered $[0, n)$.
- `add_edge(u, v)` adds a directed edge from u to v .
- `find_dominators(start)` returns a vector `idom` where `idom[start] = start`, `idom[v]` is the immediate dominator of reachable node v , and `idom[v] = -1` if v is unreachable.

Time Complexity: $O(\max(n, m) \log n)$ per call to `find_dominators()` in this path-compressed implementation, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for the graph and auxiliary heap arrays.
- $O(n)$ auxiliary stack space for the DFS.

```

#include <algorithm>
#include <vector>

class Dominators {
    std::vector<std::vector<int>> adj, pred, bucket;
    std::vector<int> time, node_at_time, parent, sdom, idom_index, dsu, best;
    int timer;

    void dfs(int u) {
        time[u] = ++timer;
        node_at_time[time] = u;
        sdom[time] = dsu[time] = best[time] = timer;
        for (int v : adj[u]) {
            if (time[v] == 0) {
                dfs(v);
                parent[time[v]] = time[u];
            }
            if (time[v] != 0) {
                pred[time[v]].push_back(time[u]);
            }
        }
    }
}

```

```

    }
}

int find_best(int u, bool compress) {
    if (dsu[u] != u) {
        int v = find_best(dsu[u], true);
        if (sdom[v] < sdom[best[u]]) {
            best[u] = v;
        }
        if (compress) {
            dsu[u] = dsu[dsu[u]];
        }
    }
    return best[u];
}

public:
    explicit Dominators(int n) : adj(n), timer(0) {}

    void add_edge(int u, int v) { adj[u].push_back(v); }

    std::vector<int> find_dominators(int start) {
        int n = static_cast<int>(adj.size());
        timer = 0;
        time.assign(n, 0);
        node_at_time.assign(n + 1, 0);
        parent.assign(n + 1, 0);
        sdom.assign(n + 1, 0);
        idom_index.assign(n + 1, 0);
        dsu.assign(n + 1, 0);
        best.assign(n + 1, 0);
        pred.assign(n + 1, {});
        bucket.assign(n + 1, {});
        dfs(start);
        for (int i = timer; i >= 1; i--) {
            for (int v : pred[i]) {
                sdom[i] = std::min(sdom[i], sdom[find_best(v, false)]);
            }
            if (i > 1) {
                bucket[sdom[i]].push_back(i);
            }
            for (int v : bucket[i]) {
                int u = find_best(v, false);
                idom_index[v] = (sdom[u] == sdom[v]) ? sdom[v] : u;
            }
            if (i > 1) {
                dsu[i] = parent[i];
            }
        }
        std::vector<int> idom(adj.size(), -1);
        idom[start] = start;
        for (int i = 2; i <= timer; i++) {
            if (idom_index[i] != sdom[i]) {
                idom_index[i] = idom_index[idom_index[i]];
            }
            idom[node_at_time[i]] = node_at_time[idom_index[i]];
        }
        return idom;
    }
};

```

Example Usage

```

#include <cassert>
using namespace std;

```

```

int main() {
    // 0 ---> 1 ---+ +-----> 4
    // |           | /         |
    // v           v /         v
    // 2 -----> 3 ----> 5 ----> 6
    Dominators g(7);
    g.add_edge(0, 1);
    g.add_edge(0, 2);
    g.add_edge(1, 3);
    g.add_edge(2, 3);
    g.add_edge(3, 4);
    g.add_edge(3, 5);
    g.add_edge(4, 6);
    g.add_edge(5, 6);
    vector<int> idom = g.find_dominators(0);
    assert((idom == vector<int>{0, 0, 0, 0, 3, 3, 3}));
    return 0;
}

```

4.2.5 Articulation Points, Biconnected Components, Block-Cut Forest

Given an undirected graph, compute articulation points, vertex-biconnected components (BCCs), and the block-cut forest using Tarjan's algorithm. An articulation point (a.k.a. cut vertex) is a node whose removal increases the number of connected components in the graph. A vertex-biconnected component is a maximal subgraph that cannot be disconnected by removing one node, with single-edge and isolated-node components handled as degenerate cases. A single depth-first search tracks each node's low-link, the earliest entry time reachable from its subtree: a node is an articulation point exactly when some child's subtree cannot reach above it, and the edges of each BCC are collected on a stack popped at that moment.

The block-cut forest is a bipartite forest with one node for each BCC and one node for each articulation point, with an edge whenever an articulation point belongs to a BCC. Note that this differs from a bridge forest: the block-cut forest describes vertex connectivity, while a bridge forest describes edge connectivity after compressing 2-edge-connected components.

- `BiconnectedComponents(n = 0)` constructs an undirected graph of `n` nodes numbered `[0, n)`.
- `add_edge(u, v)` adds the undirected edge `u-v`. Parallel edges are supported.
- `build_bcc()` computes the articulation points and vertex-biconnected components.
- `articulation_points()` and `components()` return those results.
- `build_block_cut_forest()` computes the block-cut forest using the results of the previous `build_bcc()` call.
- `block_cut_forest()` returns that forest. Its component nodes are numbered in the range `[0, components().size())`.
- `block_cut_id(v)` returns the block-cut forest node representing articulation point `v`, or `-1` if `v` is not an articulation point.

Time Complexity: $O(\max(n, m))$ per call to `build_bcc()` and `build_block_cut_forest()`, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, edge stack, BCCs, and block-cut forest.
- $O(n)$ auxiliary stack space for `build_bcc()`.

```

#include <algorithm>
#include <utility>
#include <vector>

class BiconnectedComponents {
    std::vector<std::vector<int>>> adj, bccs, forest;
    std::vector<int> lowlink, tin, forest_id, articulation;
    std::vector<char> visited, is_articulation;
    std::vector<std::pair<int, int>>> edges;
    std::vector<int> edge_stack;

```

```

int timer;

int other(int id, int u) const { return edges[id].first ^ edges[id].second ^ u; }

void add_component_until(int stop_id) {
    std::vector<int> component;
    while (true) {
        int id = edge_stack.back();
        edge_stack.pop_back();
        component.push_back(edges[id].first);
        component.push_back(edges[id].second);
        if (id == stop_id) {
            break;
        }
    }
    std::sort(component.begin(), component.end());
    component.erase(std::unique(component.begin(), component.end()), component.end());
    bccs.push_back(component);
}

void dfs(int u, int p) {
    visited[u] = true;
    lowlink[u] = tin[u] = timer++;
    int children = 0;
    for (int id : adj[u]) {
        if (id == p) {
            continue;
        }
        int v = other(id, u);
        if (visited[v]) {
            if (tin[v] < tin[u]) {
                edge_stack.push_back(id);
                lowlink[u] = std::min(lowlink[u], tin[v]);
            }
        } else {
            edge_stack.push_back(id);
            dfs(v, id);
            lowlink[u] = std::min(lowlink[u], lowlink[v]);
            children++;
            if (lowlink[v] >= tin[u]) {
                if (p != -1 || children > 1) {
                    is_articulation[u] = true;
                }
                add_component_until(id);
            }
        }
    }
    if (p == -1 && children == 0) {
        bccs.push_back({u});
    }
}

public:
BiconnectedComponents(int n = 0) : adj(n) {}

void add_edge(int u, int v) {
    int id = static_cast<int>(edges.size());
    adj[u].push_back(id);
    adj[v].push_back(id);
    edges.emplace_back(u, v);
}

void build_bcc() {
    int n = static_cast<int>(adj.size());
    articulation.clear();
    bccs.clear();
    edge_stack.clear();
    lowlink.assign(n, 0);
}

```

```

tin.assign(n, 0);
visited.assign(n, false);
is_articulation.assign(n, false);
timer = 0;
for (int i = 0; i < n; i++) {
    if (!visited[i]) {
        dfs(i, -1);
    }
}
for (int i = 0; i < n; i++) {
    if (is_articulation[i]) {
        articulation.push_back(i);
    }
}
}

void build_block_cut_forest() {
    int n = static_cast<int>(adj.size());
    int blocks = static_cast<int>(bccs.size());
    forest_id.assign(n, -1);
    int total = blocks;
    for (int v : articulation) {
        forest_id[v] = total++;
    }
    forest.assign(total, {});
    for (int b = 0; b < blocks; b++) {
        for (int v : bccs[b]) {
            if (forest_id[v] != -1) {
                forest[b].push_back(forest_id[v]);
                forest[forest_id[v]].push_back(b);
            }
        }
    }
}

const std::vector<int> &articulation_points() const { return articulation; }
const std::vector<std::vector<int>> &components() const { return bccs; }
const std::vector<std::vector<int>> &block_cut_forest() const { return forest; }
int block_cut_id(int v) const { return forest_id[v]; }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // 0---1---2    3---7
    //   \ |
    //     5---4
    BiconnectedComponents g(8);
    g.add_edge(0, 1);
    g.add_edge(0, 5);
    g.add_edge(1, 2);
    g.add_edge(1, 5);
    g.add_edge(3, 7);
    g.add_edge(4, 5);
    g.build_bcc();
    assert((g.articulation_points() == vector<int>{1, 5}));
    assert((g.components() == vector<vector<int>>{{1, 2}, {4, 5}, {0, 1, 5}, {3, 7}, {6}}));
    g.build_block_cut_forest();
    assert((g.block_cut_forest() == vector<vector<int>>{{5}, {6}, {5, 6}, {}, {}, {0, 2}, {1, 2}}));
    return 0;
}

```

4.2.6 Bridges, 2-Edge-Connected Components, Bridge Forest

Given an undirected graph, compute bridges, 2-edge-connected components, and the bridge forest using Tarjan's algorithm. A bridge is an edge whose removal increases the number of connected components in the graph. A 2-edge-connected component is a maximal set of nodes connected without crossing a bridge. A single depth-first search tracks each node's low-link, the earliest entry time reachable from its subtree: a tree edge is a bridge exactly when the child's subtree cannot reach back to the parent's side by any other route.

After each 2-edge-connected component is condensed into one node, the original bridges connect those nodes into a bridge tree, or a bridge forest when the original graph is disconnected. This differs from a block-cut forest: the bridge forest describes edge connectivity, while the block-cut forest describes vertex connectivity using articulation points and vertex-biconnected components.

- `BridgeDecomposition(n = 0)` constructs an undirected graph of `n` nodes numbered `[0, n)`.
- `add_edge(u, v)` adds the undirected edge `u-v`. Parallel edges are supported.
- `build_bridges()` computes the bridges, returned by `bridges()`.
- `build_bridge_forest()` computes the 2-edge-connected components and bridge forest using the results of the previous `build_bridges()` call.
- `components()` and `bridge_forest()` return those results.
- `component_id(u)` returns the 2-edge-connected component ID containing node `u`.

Time Complexity: $O(\max(n, m))$ per call to `build_bridges()` followed by `build_bridge_forest()`, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, bridges, 2-edge-connected components, and bridge forest.
- $O(n)$ auxiliary stack space for the DFS calls.

```
#include <algorithm>
#include <utility>
#include <vector>

class BridgeDecomposition {
    std::vector<std::vector<int>>> adj, two_edge_components, forest;
    std::vector<int> lowlink, tin, component;
    std::vector<char> visited, is_bridge_edge;
    std::vector<std::pair<int, int>>> edges, bridge_edges;
    int timer;

    int other(int id, int u) const { return edges[id].first ^ edges[id].second ^ u; }

    void dfs_bridges(int u, int p) {
        visited[u] = true;
        lowlink[u] = tin[u] = timer++;
        for (int id : adj[u]) {
            if (id == p) {
                continue;
            }
            int v = other(id, u);
            if (visited[v]) {
                lowlink[u] = std::min(lowlink[u], tin[v]);
            } else {
                dfs_bridges(v, id);
                lowlink[u] = std::min(lowlink[u], lowlink[v]);
                if (lowlink[v] > tin[u]) {
                    bridge_edges.emplace_back(u, v);
                    is_bridge_edge[id] = true;
                }
            }
        }
    }

    void dfs_component(int u, int id) {
```

```

    component[u] = id;
    two_edge_components[id].push_back(u);
    for (int edge_id : adj[u]) {
        int v = other(edge_id, u);
        if (component[v] == -1 && !is_bridge_edge[edge_id]) {
            dfs_component(v, id);
        }
    }
}

public:
    BridgeDecomposition(int n = 0) : adj(n) {}

    void add_edge(int u, int v) {
        int id = static_cast<int>(edges.size());
        adj[u].push_back(id);
        adj[v].push_back(id);
        edges.emplace_back(u, v);
    }

    void build_bridges() {
        int n = static_cast<int>(adj.size());
        bridge_edges.clear();
        lowlink.assign(n, 0);
        tin.assign(n, 0);
        visited.assign(n, false);
        is_bridge_edge.assign(edges.size(), false);
        timer = 0;
        for (int i = 0; i < n; i++) {
            if (!visited[i]) {
                dfs_bridges(i, -1);
            }
        }
    }

    void build_bridge_forest() {
        int n = static_cast<int>(adj.size());
        component.assign(n, -1);
        two_edge_components.clear();
        for (int i = 0; i < n; i++) {
            if (component[i] == -1) {
                two_edge_components.push_back({});
                dfs_component(i, static_cast<int>(two_edge_components.size() - 1));
            }
        }
        forest.assign(two_edge_components.size(), {});
        for (auto &[u, v] : bridge_edges) {
            int cu = component[u], cv = component[v];
            forest[cu].push_back(cv);
            forest[cv].push_back(cu);
        }
    }

    const std::vector<std::pair<int, int>> &bridges() const { return bridge_edges; }
    const std::vector<std::vector<int>> &components() const { return two_edge_components; }
    const std::vector<std::vector<int>> &bridge_forest() const { return forest; }
    int component_id(int u) const { return component[u]; }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // 0---1---2    3---7

```

```

// \ |
// 5---4
BridgeDecomposition g(8);
g.add_edge(0, 1);
g.add_edge(0, 5);
g.add_edge(1, 2);
g.add_edge(1, 5);
g.add_edge(3, 7);
g.add_edge(4, 5);
g.build_bridges();
assert((g.bridges() == vector<pair<int, int>>{{1, 2}, {5, 4}, {3, 7}}));
g.build_bridge_forest();
vector<vector<int>> components = g.components();
for (auto &component : components) {
    sort(component.begin(), component.end());
}
sort(components.begin(), components.end());
assert((components == vector<vector<int>>{{0, 1, 5}, {2}, {3}, {4}, {6}, {7}}));
assert((g.bridge_forest() == vector<vector<int>>{{1, 3}, {0}, {5}, {0}, {}, {2}}));
return 0;
}

```

4.2.7 Online Bridge Finding

Maintain the number of bridges in an undirected graph while edges are inserted one at a time, along with the 2-edge-connected components and ordinary connected components. When an inserted edge joins two different connected components, it creates a new bridge. When it joins two nodes already in the same connected component, it creates a cycle and every bridge on the path between the two endpoints stops being a bridge.

- `OnlineBridges(n = 0)` constructs a graph with `n` isolated nodes numbered $[0, n)$.
- `add_edge(u, v)` adds the undirected edge `u-v` and updates the 2-edge-connected and ordinary connected components.
- `bridge_count()` returns the current number of bridges.
- `find_2ecc(u)` returns node `u`'s representative in the 2-edge-connected components partition.
- `find_cc(u)` returns node `u`'s representative in the ordinary connected components partition.

Parallel edges are supported. Adding a second edge between two nodes that are already in the same connected component dissolves every bridge on the path between them, merging those nodes into the same 2-edge-connected component.

Time Complexity: $O(\log n)$ amortized per call to `add_edge()` in this implementation, where n is the number of nodes.

Space Complexity: $O(n)$ for the disjoint-set arrays and auxiliary parent links.

```

#include <utility>
#include <vector>

class OnlineBridges {
    std::vector<int> dsu_2ecc, dsu_cc, dsu_cc_size, parent, last_visit;
    int lca_iteration, num_bridges;

    void make_root(int u) {
        u = find_2ecc(u);
        int root = u, child = -1;
        while (u != -1) {
            int p = find_2ecc(parent[u]);
            parent[u] = child;
            dsu_cc[u] = root;
            child = u;
            u = p;
        }
    }
}

```



```

    dsu_cc_size[root] = dsu_cc_size[child];
}

void merge_path(int u, int v) {
    lca_iteration++;
    std::vector<int> path_a, path_b;
    int lca = -1;
    while (lca == -1) {
        if (u != -1) {
            u = find_2ecc(u);
            path_a.push_back(u);
            if (last_visit[u] == lca_iteration) {
                lca = u;
                break;
            }
            last_visit[u] = lca_iteration;
            u = parent[u];
        }
        if (v != -1) {
            v = find_2ecc(v);
            path_b.push_back(v);
            if (last_visit[v] == lca_iteration) {
                lca = v;
                break;
            }
            last_visit[v] = lca_iteration;
            v = parent[v];
        }
    }
    for (int a : path_a) {
        dsu_2ecc[a] = lca;
        if (a == lca) {
            break;
        }
    }
    num_bridges--;
    for (int b : path_b) {
        dsu_2ecc[b] = lca;
        if (b == lca) {
            break;
        }
    }
    num_bridges--;
}

public:
OnlineBridges(int n = 0) {
    dsu_2ecc.resize(n);
    dsu_cc.resize(n);
    dsu_cc_size.assign(n, 1);
    parent.assign(n, -1);
    last_visit.assign(n, 0);
    lca_iteration = 0;
    num_bridges = 0;
    for (int i = 0; i < n; i++) {
        dsu_2ecc[i] = dsu_cc[i] = i;
    }
}

int find_2ecc(int u) {
    if (u == -1) {
        return -1;
    }
    if (dsu_2ecc[u] == u) {
        return u;
    }
    return dsu_2ecc[u] = find_2ecc(dsu_2ecc[u]);
}

```

```

int find_cc(int u) {
    u = find_2ecc(u);
    if (dsu_cc[u] == u) {
        return u;
    }
    return dsu_cc[u] = find_cc(dsu_cc[u]);
}

void add_edge(int u, int v) {
    u = find_2ecc(u);
    v = find_2ecc(v);
    if (u == v) {
        return;
    }
    int ca = find_cc(u), cb = find_cc(v);
    if (ca != cb) {
        num_bridges++;
        if (dsu_cc_size[ca] > dsu_cc_size[cb]) {
            std::swap(u, v);
            std::swap(ca, cb);
        }
        make_root(u);
        parent[u] = v;
        dsu_cc[u] = v;
        dsu_cc_size[cb] += dsu_cc_size[u];
    } else {
        merge_path(u, v);
    }
}

int bridge_count() const { return num_bridges; }
};

```

Example Usage

```

#include <cassert>

int main() {
    // 0---1
    //   \ |
    //     2---3
    OnlineBridges g(4);
    g.add_edge(0, 1);
    assert(g.bridge_count() == 1);
    g.add_edge(1, 2);
    assert(g.bridge_count() == 2);
    g.add_edge(2, 0);
    assert(g.bridge_count() == 0);
    g.add_edge(2, 3);
    assert(g.bridge_count() == 1);
    return 0;
}

```

4.2.8 2-SAT

Solve a Boolean formula in 2-CNF, where each clause contains at most two literals. A clause like $(a \vee b)$ is represented by the implications $\neg a \rightarrow b$ and $\neg b \rightarrow a$. The formula is satisfiable if and only if no variable and its negation belong to the same strongly connected component of the implication graph.

Variables are numbered $[0, n)$. A literal is represented by `literal(variable, value)`, where `value == true` means the variable itself and `value == false` means its negation.

- `TwoSAT(n = 0)` constructs an empty formula over `n` variables.
- `literal(variable, value)` returns the integer ID for a variable or its negation.
- `add_implication(a, b)` adds the implication $a \rightarrow b$.
- `add_or(a, b)` adds the clause $(a \vee b)$.
- `add_true(a)` forces literal `a` to be true.
- `add_false(a)` forces literal `a` to be false.
- `add(variable, value)` forces a variable to equal `value`.
- `add(x, xval, y, yval)` adds the clause $(x = xval \vee y = yval)$.
- `satisfiable()` returns whether all added clauses can be satisfied.
- `assignment()` returns the valid assignment found by the last successful `satisfiable()` call.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the number of variables.
- $O(1)$ per call to `literal()`, `add_implication()`, `add_or()`, `add_true()`, `add_false()`, and `add()`.
- $O(\max(n, m))$ per call to `satisfiable()`, where n is the number of variables and m is the number of edges in the implication graph.

Space Complexity: $O(\max(n, m))$ for the implication graph and DFS stacks.

```
#include <algorithm>
#include <vector>

class TwoSAT {
    int variables;
    std::vector<std::vector<int>>> adj, rev;
    std::vector<int> order, component, solution;

    static int neg(int x) { return x ^ 1; }

    void dfs_order(int u, std::vector<char> &visited) {
        visited[u] = true;
        for (int v : adj[u]) {
            if (!visited[v]) {
                dfs_order(v, visited);
            }
        }
        order.push_back(u);
    }

    void dfs_component(int u, int id) {
        component[u] = id;
        for (int v : rev[u]) {
            if (component[v] == -1) {
                dfs_component(v, id);
            }
        }
    }

public:
    TwoSAT(int n = 0) : variables(n), adj(2 * n), rev(2 * n), solution(n, false) {}

    int literal(int variable, bool value) const { return 2 * variable + (value ? 0 : 1); }

    void add_implication(int a, int b) {
        adj[a].push_back(b);
        rev[b].push_back(a);
    }

    void add_or(int a, int b) {
        add_implication(neg(a), b);
        add_implication(neg(b), a);
    }

    void add_true(int a) { add_implication(neg(a), a); }
```

```

void add_false(int a) { add_implication(a, neg(a)); }
void add(int variable, bool value) { add_true(literal(variable, value)); }
void add(int x, bool xval, int y, bool yval) { add_or(literal(x, xval), literal(y, yval)); }

bool satisfiable() {
    order.clear();
    component.assign(2 * variables, -1);
    std::vector<char> visited(2 * variables, false);
    for (int i = 0; i < 2 * variables; i++) {
        if (!visited[i]) {
            dfs_order(i, visited);
        }
    }
    std::reverse(order.begin(), order.end());
    int id = 0;
    for (int u : order) {
        if (component[u] == -1) {
            dfs_component(u, id++);
        }
    }
    for (int i = 0; i < variables; i++) {
        if (component[2 * i] == component[2 * i + 1]) {
            return false;
        }
        solution[i] = component[2 * i] > component[2 * i + 1];
    }
    return true;
}

const std::vector<int> &assignment() const { return solution; }
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    TwoSAT solver(3);
    int x0 = solver.literal(0, true);
    int not_x0 = solver.literal(0, false);
    int x1 = solver.literal(1, true);
    int x2 = solver.literal(2, true);
    solver.add_or(x0, x1);
    solver.add_or(not_x0, x2);
    solver.add_true(x0);
    solver.add(1, true, 2, false);
    assert(solver.satisfiable());
    assert((solver.assignment() == vector<int>{true, true, true}));
    return 0;
}

```

4.2.9 Functional Graph

A functional graph is a directed graph in which every node has exactly one outgoing edge, given by a successor function f , where $f(i)$ is the node reached in one step from node i . Following the edges from any node eventually enters a cycle, so each weakly connected component consists of one directed cycle with trees of nodes feeding into it. This structure underlies problems about iterating a function: where a sequence $i, f(i), f(f(i)), \dots$ ends up, and how far it is from repeating.

Building the structure peels away nodes of in-degree zero to expose the cycles (the survivors), then labels each cycle and walks the remaining tail nodes to record their distance to the cycle. A binary lifting table stores the

function applied 2^k times, so any number of steps can be taken in logarithmic time. Together these answer, for any node, which cycle it eventually reaches, how many steps until it gets there, and exactly where it lands after a huge number of steps.

- `FunctionalGraph(f)` builds the structure for the successor function `f`, whose indices represent the nodes.
- `on_cycle(i)` returns whether node `i` lies on a cycle.
- `dist_to_cycle(i)` returns the number of steps from `i` to the first cycle node it reaches (0 if `i` is itself on a cycle).
- `cycle_id(i)` returns an identifier for the cycle that `i` eventually reaches; two nodes share an identifier exactly when they funnel into the same cycle.
- `cycle_length(id)` returns the length of the cycle with the given identifier.
- `num_cycles()` returns the number of distinct cycles.
- `successor(i, steps)` returns the node reached from `i` after applying `f` exactly `steps` times, for any nonnegative `steps` (which may be astronomically large).

Time Complexity:

- $O(n \log n)$ per call to the constructor, where n is the number of nodes.
- $O(\log n)$ per call to `successor()`.
- $O(1)$ per call to all other member functions.

Space Complexity: $O(n \log n)$ for storage of the binary lifting table.

```
#include <cstdint>
#include <vector>

class FunctionalGraph {
    std::vector<std::vector<int>>> up; // up[k][i] = f applied 2^k times to i.
    std::vector<char> on_cyc;
    std::vector<int> to_cyc; // Steps from i to the first cycle node.
    std::vector<int> comp; // Identifier of the cycle that i reaches.
    std::vector<int> clen; // clen[c] = length of cycle c.

    int jump(int i, int64_t steps) const {
        for (int k = 0; steps > 0; k++, steps >>= 1) {
            if (steps & 1) {
                i = up[k][i];
            }
        }
        return i;
    }

public:
    FunctionalGraph(const std::vector<int> &f) {
        int n = static_cast<int>(f.size()), levels = 1;
        while ((1LL << levels) < n) {
            levels++;
        }
        up.assign(levels + 1, std::vector<int>(n));
        up[0] = f;
        for (int k = 1; k <= levels; k++) {
            for (int i = 0; i < n; i++) {
                up[k][i] = up[k - 1][up[k - 1][i]];
            }
        }
        // Peel nodes of in-degree zero; the survivors are exactly the cycle nodes.
        std::vector<int> indeg(n, 0);
        for (int i = 0; i < n; i++) {
            indeg[f[i]]++;
        }
        std::vector<char> removed(n, false);
        std::vector<int> st;
        for (int i = 0; i < n; i++) {
            if (indeg[i] == 0) {
                st.push_back(i);
            }
        }
    }
};
```

```

    }
}
while (!st.empty()) {
    int u = st.back();
    st.pop_back();
    removed[u] = true;
    if (--indeg[f[u]] == 0) {
        st.push_back(f[u]);
    }
}
on_cyc.assign(n, false);
for (int i = 0; i < n; i++) {
    on_cyc[i] = !removed[i];
}
// Label each cycle and record its length.
comp.assign(n, -1);
to_cyc.assign(n, 0);
for (int i = 0; i < n; i++) {
    if (on_cyc[i] && comp[i] == -1) {
        int len = 0;
        for (int u = i; comp[u] == -1; u = f[u]) {
            comp[u] = static_cast<int>(clen.size());
            len++;
        }
        clen.push_back(len);
    }
}
// Walk each tail to the cycle, assigning its component and distance.
for (int i = 0; i < n; i++) {
    if (comp[i] != -1) {
        continue;
    }
    std::vector<int> path;
    int u = i;
    while (comp[u] == -1) {
        path.push_back(u);
        u = f[u];
    }
    for (int j = static_cast<int>(path.size()) - 1, step = 1; j >= 0; j--, step++) {
        comp[path[j]] = comp[u];
        to_cyc[path[j]] = to_cyc[u] + step;
    }
}
}

bool on_cycle(int i) const { return on_cyc[i]; }
int dist_to_cycle(int i) const { return to_cyc[i]; }
int cycle_id(int i) const { return comp[i]; }
int cycle_length(int id) const { return clen[id]; }
int num_cycles() const { return static_cast<int>(clen.size()); }

int successor(int i, int64_t steps) const {
    if (steps <= to_cyc[i]) {
        return jump(i, steps);
    }
    steps -= to_cyc[i];
    i = jump(i, to_cyc[i]); // Advance to the cycle entry node.
    steps %= clen[comp[i]];
    return jump(i, steps);
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

```

```

int main() {
    // 0 --> 1 <-- 3 <-- 4 <-- 5
    //   ^   /
    //   \ v
    //   2 <-- 6
    FunctionalGraph g({1, 2, 0, 1, 3, 4, 2});
    assert(g.on_cycle(0) && g.on_cycle(1) && g.on_cycle(2));
    assert(!g.on_cycle(3) && !g.on_cycle(5));
    assert(g.num_cycles() == 1);
    assert(g.cycle_length(0) == 3);

    assert(g.dist_to_cycle(2) == 0);
    assert(g.dist_to_cycle(5) == 3);
    assert(g.dist_to_cycle(6) == 1);
    assert(g.cycle_id(5) == g.cycle_id(0));

    assert(g.successor(5, 1) == 4);
    assert(g.successor(5, 3) == 1);    // 5 -> 4 -> 3 -> 1.
    assert(g.successor(5, 4) == 2);    // One more step around into the cycle.
    assert(g.successor(0, 100) == 1);  // 100 mod 3 == 1 step around the cycle.
    return 0;
}

```

4.2.10 Offline Dynamic Connectivity

Maintains the number of connected components of an undirected graph as edges are added and removed over time, answering queries offline once the entire sequence of operations is known. Each edge is alive over a contiguous interval of time, from when it is added until it is removed. These intervals are placed into a segment tree over the time axis, so that for T operations, each edge appears in $O(\log T)$ nodes whose ranges it spans. A depth-first traversal of this segment tree unites the edges stored at each node on the way down and undoes them on the way back up, so at each leaf the disjoint set forest reflects exactly the edges alive at that moment.

This offline method supports arbitrary undirected graphs. When updates must be processed online and the graph remains a forest, see the Euler Tour Tree in Section 2.7.9.

Undoing unions requires a disjoint set forest with rollback (see the disjoint sets section): it joins by size or rank without path compression, recording each change on a stack so it can be reversed. Path compression is incompatible with rollback, so finding a representative costs $O(\log n)$ rather than near-constant time.

- `OfflineDynamicConnectivity(n)` creates a structure of `n` nodes numbered $[0, n)$, initially with no edges.
- `add_edge(u, v)` and `remove_edge(u, v)` record adding or removing the undirected edge `u-v` at the current time, then advance the time by one step. Each edge may be present at most once at a time: `add_edge` must not name an edge that is already present, and `remove_edge` must name one that is. Both preconditions are checked with `assert` during `solve()`.
- `count_components()` records a query at the current time and advances the time by one step.
- `solve()` processes all recorded operations and returns a vector holding, for each `count_components()` query in the order issued, the number of connected components at that time.

Time Complexity:

- $O(T \log T + m \log T \log n)$ per call to `solve()`, where T is the total number of operations (`add_edge()`, `remove_edge()`, and `count_components()` combined), m is the number of `add_edge()` operations, and n is the number of nodes.
- $O(1)$ amortized per call to `add_edge()`, `remove_edge()`, and `count_components()`.

Space Complexity: $O(n + T + m \log T)$ auxiliary.

```

#include <algorithm>
#include <cassert>
#include <map>

```

```

#include <numeric>
#include <utility>
#include <vector>

class OfflineDynamicConnectivity {
    static const int ADD = 0, REMOVE = 1, QUERY = 2;

    struct Op {
        int type, u, v;
    };

    int n, num_sets;
    std::vector<Op> ops;
    std::vector<int> root, rank;
    std::vector<std::pair<int, int>> history; // (child root, whether its parent's rank increased)
    std::vector<std::vector<std::pair<int, int>>> seg;
    std::vector<int> answers;

    int find_root(int u) const {
        while (root[u] != u) {
            u = root[u];
        }
        return u;
    }

    void unite(int u, int v) {
        u = find_root(u);
        v = find_root(v);
        if (u == v) {
            history.emplace_back(-1, 0);
            return;
        }
        if (rank[u] < rank[v]) {
            std::swap(u, v);
        }
        history.emplace_back(v, rank[u] == rank[v] ? 1 : 0);
        root[v] = u;
        if (rank[u] == rank[v]) {
            rank[u]++;
        }
        num_sets--;
    }

    void rollback() {
        auto [child, raised] = history.back();
        history.pop_back();
        if (child == -1) {
            return;
        }
        rank[root[child]] -= raised;
        root[child] = child;
        num_sets++;
    }

    void add_interval(
        int node, int lo, int hi, int tgt_lo, int tgt_hi, const std::pair<int, int> &e
    ) {
        if (tgt_hi < lo || hi < tgt_lo) {
            return;
        }
        if (tgt_lo <= lo && hi <= tgt_hi) {
            seg[node].push_back(e);
            return;
        }
        int mid = lo + (hi - lo) / 2;
        add_interval(2 * node, lo, mid, tgt_lo, tgt_hi, e);
        add_interval(2 * node + 1, mid + 1, hi, tgt_lo, tgt_hi, e);
    }
}

```



```

void dfs(int node, int lo, int hi) {
    int applied = 0;
    for (const auto &[eu, ev] : seg[node]) {
        unite(eu, ev);
        applied++;
    }
    if (lo == hi) {
        if (ops[lo].type == QUERY) {
            answers.push_back(num_sets);
        }
    } else {
        int mid = lo + (hi - lo) / 2;
        dfs(2 * node, lo, mid);
        dfs(2 * node + 1, mid + 1, hi);
    }
    while (applied-- > 0) {
        rollback();
    }
}

public:
OfflineDynamicConnectivity(int n) : n(n) {}

void add_edge(int u, int v) { ops.push_back({ADD, std::min(u, v), std::max(u, v)}); }
void remove_edge(int u, int v) { ops.push_back({REMOVE, std::min(u, v), std::max(u, v)}); }
void count_components() { ops.push_back({QUERY, 0, 0}); }

std::vector<int> solve() {
    int t = static_cast<int>(ops.size());
    answers.clear();
    if (t == 0) {
        return answers;
    }
    seg.assign(4 * t, {});
    std::map<std::pair<int, int>, int> active;
    for (int i = 0; i < t; i++) {
        std::pair<int, int> key(ops[i].u, ops[i].v);
        if (ops[i].type == ADD) {
            assert(active.find(key) == active.end()); // The edge must not already be present.
            active[key] = i;
        } else if (ops[i].type == REMOVE) {
            auto it = active.find(key);
            assert(it != active.end()); // The edge must currently be present.
            add_interval(1, 0, t - 1, it->second, i - 1, {ops[i].u, ops[i].v});
            active.erase(it);
        }
    }
    for (const auto &[key, start] : active) {
        add_interval(1, 0, t - 1, start, t - 1, {key.first, key.second});
    }
    root.resize(n);
    std::iota(root.begin(), root.end(), 0);
    rank.assign(n, 0);
    num_sets = n;
    history.clear();
    dfs(1, 0, t - 1);
    return answers;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

```

```

int main() {
    OfflineDynamicConnectivity dc(4);
    // 0 1 2 3
    dc.count_components(); // 4 isolated nodes.
    dc.add_edge(0, 1);
    dc.add_edge(2, 3);
    // 0---1 2---3
    dc.count_components(); // 2 components.
    dc.add_edge(1, 2);
    // 0---1---2---3
    dc.count_components(); // All connected => 1 component.
    dc.remove_edge(1, 2);
    // 0---1 2---3
    dc.count_components(); // Back to 2 components.
    assert((dc.solve() == vector<int>{4, 2, 1, 2}));
    return 0;
}

```

4.3 Shortest Paths

4.3.1 Unweighted Shortest Path (BFS)

Given one or more starting nodes in an unweighted, directed graph, visit every reachable node and determine its minimum distance from any start. Breadth-first search explores nodes in order of increasing distance using a queue, so the first time a node is reached is via a shortest path. Initializing the queue with several distance-0 sources gives multi-source BFS without changing the traversal. Optionally, reconstruct a shortest path using the predecessor vector `pred`.

- `bfs(starts)` populates `dist` and `pred` from every node in `starts`, assigning each node its minimum distance from any start. The global, pre-populated adjacency list `adj` uses its indices as nodes and must contain only valid node indices. Pass a singleton vector for single-source BFS.
- `get_path(dest)` returns the path from a nearest starting node to `dest`, or an empty vector if `dest` is unreachable, after the latest call to `bfs()`.

For path reconstruction, `pred[v]` stores the node immediately before `v` on the shortest path from a nearest start to `v`, or `-1` if `v` is a starting node or unreachable. Follow `pred` backward from the destination to its starting node, then reverse that sequence to recover the path.

Time Complexity:

- $O(\max(n, m, s))$ per call, where n is the number of nodes, m is the number of edges, and s is the number of supplied starting nodes. For single-source BFS, $s = 1$.
- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary.
- $O(p)$ for the path returned by `get_path()`.

```

#include <algorithm>
#include <climits>
#include <queue>
#include <vector>

const int INF = INT_MAX / 2;
std::vector<std::vector<int>>> adj;
std::vector<int> dist, pred;

void bfs(const std::vector<int> &starts) {

```

```

int n = static_cast<int>(adj.size());
dist.assign(n, INF);
pred.assign(n, -1);
std::queue<int> q;
for (int start : starts) {
    if (dist[start] == 0) {
        continue;
    }
    dist[start] = 0;
    q.push(start);
}
while (!q.empty()) {
    int u = q.front();
    q.pop();
    for (int v : adj[u]) {
        if (dist[v] != INF) {
            continue;
        }
        dist[v] = dist[u] + 1;
        pred[v] = u;
        q.push(v);
    }
}

std::vector<int> get_path(int dest) {
    if (dist[dest] == INF) {
        return {};
    }
    std::vector<int> path;
    for (int v = dest; v != -1; v = pred[v]) {
        path.push_back(v);
    }
    std::reverse(path.begin(), path.end());
    return path;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // 0 --> 1 --> 2
    //      |   |
    //      v   |
    //      3 <---+
    adj.assign(4, {});
    adj[0].push_back(1);
    adj[1].push_back(2);
    adj[1].push_back(3);
    adj[2].push_back(3);
    int start = 0, dest = 3;
    bfs(vector<int>{start});
    assert(dist[dest] == 2);
    assert((get_path(dest) == vector<int>{0, 1, 3}));

    bfs(vector<int>{0, 2});
    assert(dist[2] == 0);
    assert(dist[3] == 1);
    assert((get_path(3) == vector<int>{2, 3}));
    return 0;
}

```

4.3.2 0-1 Shortest Path (Deque BFS)

Given one or more starting nodes in a weighted graph whose edge weights are only 0 or 1, compute the shortest distance from any start to every reachable node. Optionally, reconstruct the shortest path to a specific destination node using the shortest-path tree from the predecessor array `pred`.

0-1 BFS is a specialized version of Dijkstra's algorithm. Because every relaxation changes the distance by either 0 or 1, a deque maintains nodes in nondecreasing distance order: push weight-0 relaxations to the front and weight-1 relaxations to the back.

- `bfs_zero_one(starts)` populates `dist` and `pred` for shortest paths from the nodes in `starts`, assigning each node its minimum distance from any start. The global, pre-populated adjacency list `adj` uses its indices as nodes. Each edge is stored as `(neighbor, weight)`, where `weight` is either 0 or 1. Pass a singleton vector for a single source.
- `get_path(dest)` returns the path from a nearest starting node to `dest`, or an empty vector if `dest` is unreachable, after the latest call to `bfs_zero_one()`.

For path reconstruction, `pred[v]` stores the node immediately before `v` on the shortest path from a nearest start to `v`, or `-1` if `v` is a starting node or unreachable. Follow `pred` backward from the destination to its starting node, then reverse that sequence to recover the path.

Time Complexity:

- $O(\max(n, m, s))$ per call, where n is the number of nodes, m is the number of edges, and s is the number of supplied starting nodes. For a single source, $s = 1$.
- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary deque space.
- $O(p)$ for the path returned by `get_path()`.

```
#include <algorithm>
#include <climits>
#include <deque>
#include <utility>
#include <vector>

const int INF = INT_MAX / 2;
std::vector<std::vector<std::pair<int, int>>> adj;
std::vector<int> dist, pred;

void bfs_zero_one(const std::vector<int> &starts) {
    int n = static_cast<int>(adj.size());
    dist.assign(n, INF);
    pred.assign(n, -1);
    std::deque<int> dq;
    for (int start : starts) {
        if (dist[start] == 0) {
            continue;
        }
        dist[start] = 0;
        dq.push_back(start);
    }
    while (!dq.empty()) {
        int u = dq.front();
        dq.pop_front();
        for (auto &[v, w] : adj[u]) {
            if (dist[v] > dist[u] + w) {
                dist[v] = dist[u] + w;
                pred[v] = u;
                if (w == 0) {
                    dq.push_front(v);
                } else {
                    dq.push_back(v);
                }
            }
        }
    }
}
```

```

        dq.push_back(v);
    }
}
}
}

std::vector<int> get_path(int dest) {
    if (dist[dest] == INF) {
        return {};
    }
    std::vector<int> path;
    for (int v = dest; v != -1; v = pred[v]) {
        path.push_back(v);
    }
    std::reverse(path.begin(), path.end());
    return path;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //      w=1
    //      0 -----> 1
    //      |         ^ |
    // w=0 |         w=0 | w=1
    //      v /         v
    //      2 -----> 3
    //      w=1
    adj.assign(4, {});
    adj[0].emplace_back(1, 1);
    adj[0].emplace_back(2, 0);
    adj[2].emplace_back(1, 0);
    adj[1].emplace_back(3, 1);
    adj[2].emplace_back(3, 1);
    int start = 0, dest = 3;
    bfs_zero_one(vector<int>{start});
    assert(dist[dest] == 1);
    assert((get_path(dest) == vector<int>{0, 2, 3}));

    bfs_zero_one(vector<int>{1, 0});
    assert(dist[3] == 1);
    assert((get_path(3) == vector<int>{1, 3}));
    return 0;
}

```

4.3.3 Shortest Path (Dijkstra)

Given a starting node in a weighted, directed graph with nonnegative weights only, visit every reachable node and determine the minimum distance to each such node. Optionally, reconstruct the shortest path to a destination node using the shortest-path tree from the predecessor array `pred`.

Dijkstra's algorithm repeatedly selects the unvisited node of smallest tentative distance using a priority queue and relaxes its outgoing edges. Dijkstra's algorithm requires nonnegative edge weights. Use Bellman-Ford or SPFA instead when negative edges are present. Because the weights are nonnegative, a node's distance is final the first time it is removed from the queue.

- `dijkstra(start)` populates `dist` and `pred` for a global, pre-populated adjacency list `adj` whose indices represent the nodes. Each edge is stored as `(neighbor, weight)`, where `weight` is nonnegative.

- `get_path(dest)` returns the path from `start` to `dest`, or an empty vector if `dest` is unreachable, after the latest call to `dijkstra()`.

For path reconstruction, `pred[v]` stores the node immediately before `v` on the shortest path from `start` to `v`, or `-1` if `v` is `start` or unreachable. Follow `pred` backward from the destination to `start`, then reverse that sequence to recover the path.

Time Complexity:

- $O(n + m \log n)$ per call, where n is the number of nodes and m is the number of edges.
- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(\max(n, m))$ auxiliary.
- $O(p)$ for the path returned by `get_path()`.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <functional>
#include <queue>
#include <utility>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::vector<std::pair<int, int>>> adj;
std::vector<int64_t> dist;
std::vector<int> pred;

void dijkstra(int start) {
    int n = static_cast<int>(adj.size());
    dist.assign(n, INF);
    pred.assign(n, -1);
    dist[start] = 0;
    using qnode = std::pair<int64_t, int>; // (distance, node)
    std::priority_queue<qnode, std::vector<qnode>, std::greater<>> pq;
    pq.emplace(0, start);
    while (!pq.empty()) {
        auto [du, u] = pq.top();
        pq.pop();
        if (du != dist[u]) {
            continue;
        }
        for (auto &[v, w] : adj[u]) {
            if (dist[v] > dist[u] + w) {
                dist[v] = dist[u] + w;
                pred[v] = u;
                pq.emplace(dist[v], v);
            }
        }
    }
}

std::vector<int> get_path(int dest) {
    if (dist[dest] == INF) {
        return {};
    }
    std::vector<int> path;
    for (int v = dest; v != -1; v = pred[v]) {
        path.push_back(v);
    }
    std::reverse(path.begin(), path.end());
    return path;
}
```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    //      w=2      w=2
    //  0 ----> 1 ----> 2
    //   \      |      /
    // w=8\  w=4|  /w=1
    //   \      v  /
    //      --> 3 <--
    adj.assign(4, {});
    adj[0].emplace_back(1, 2);
    adj[0].emplace_back(3, 8);
    adj[1].emplace_back(2, 2);
    adj[1].emplace_back(3, 4);
    adj[2].emplace_back(3, 1);
    int start = 0, dest = 3;
    dijkstra(start);
    assert(dist[dest] == 5);
    assert((get_path(dest) == vector<int>{0, 1, 2, 3}));
    return 0;
}
```

4.3.4 Shortest Path (Bellman-Ford)

Given a starting node in a weighted, directed graph with possibly negative weights, visit every reachable node and determine the minimum distance to each such node. Optionally, reconstruct the shortest path to a destination node using the shortest-path tree from the predecessor array `pred`.

Bellman-Ford relaxes every edge in the graph up to $n - 1$ times, stopping early if a pass makes no changes. Since any shortest path uses at most $n - 1$ edges, all distances are correct after these passes unless a negative-weight cycle keeps reducing them, which a further pass detects. Bellman-Ford will also detect whether the graph contains a negative-weight cycle reachable from the start node, in which case the affected shortest paths are undefined and an error will be thrown. (To detect a negative cycle anywhere in the graph, add a virtual source with zero-weight edges to every node and start from it.)

- `bellman_ford(n, start)` populates `dist` and `pred` for a global, pre-populated edge list `edges` whose endpoints must be numbered $[0, n)$.
- `get_path(dest)` returns the path from `start` to `dest`, or an empty vector if `dest` is unreachable, after the latest call to `bellman_ford()`.

For path reconstruction, `pred[v]` stores the node immediately before `v` on the shortest path from `start` to `v`, or -1 if `v` is `start` or unreachable. Follow `pred` backward from the destination to `start`, then reverse that sequence to recover the path.

Time Complexity:

- $O(n \cdot m)$ per call, where n is the number of nodes and m is the number of edges.
- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary.
- $O(p)$ for the path returned by `get_path()`.

```
#include <algorithm>
#include <cstdint>
#include <stdexcept>
```

```

#include <tuple>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::tuple<int, int, int>> edges; // (u, v, w)
std::vector<int64_t> dist;
std::vector<int> pred;

void bellman_ford(int n, int start) {
    dist.assign(n, INF);
    pred.assign(n, -1);
    dist[start] = 0;
    for (int i = 0; i < n - 1; i++) {
        bool changed = false;
        for (auto &[u, v, w] : edges) {
            // The dist[u] != INF guard avoids relaxing out of unreachable nodes: a negative edge from an
            // unreachable u would otherwise give v a bogus finite distance (INF + w < INF).
            if (dist[u] != INF && dist[v] > dist[u] + w) {
                dist[v] = dist[u] + w;
                pred[v] = u;
                changed = true;
            }
        }
        if (!changed) break;
    }
    // Optional: report a negative-weight cycle reachable from the start node.
    for (auto &[u, v, w] : edges) {
        if (dist[u] != INF && dist[v] > dist[u] + w) {
            throw std::runtime_error("Negative-weight cycle found.");
        }
    }
}

std::vector<int> get_path(int dest) {
    if (dist[dest] == INF) {
        return {};
    }
    std::vector<int> path;
    for (int v = dest; v != -1; v = pred[v]) {
        path.push_back(v);
    }
    std::reverse(path.begin(), path.end());
    return path;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //      w=1      w=2
    // 0 -----> 1 -----> 2
    // |               ^
    // +-----+-----+
    //           w=5
    edges.emplace_back(0, 1, 1);
    edges.emplace_back(1, 2, 2);
    edges.emplace_back(0, 2, 5);
    int start = 0, dest = 2;
    bellman_ford(3, start);
    assert(dist[dest] == 3);
    assert((get_path(dest) == vector<int>{0, 1, 2}));
    return 0;
}

```


4.3.5 Shortest Path (SPFA)

Given a starting node in a weighted directed graph, compute shortest paths even when some edge weights are negative. The Shortest Path Faster Algorithm (SPFA) is a queue-based optimization of Bellman-Ford: instead of relaxing every edge in every round, it keeps a queue of nodes whose distances have improved and relaxes only their outgoing edges. It is often fast on benign inputs, but it still has Bellman-Ford's worst-case behavior and can be forced to run in $O(n \cdot m)$. Prefer Dijkstra for nonnegative weights, and use SPFA mainly when negative edges are present and the input is not adversarial.

- `spfa(start)` populates `dist` and `pred` for a global, pre-populated adjacency list `adj` which uses its indices as nodes. Each edge is stored as `(neighbor, weight)`. The function returns `false` if it detects a reachable negative cycle, and returns `true` otherwise.
- `get_path(dest)` returns the path from `start` to `dest` after a successful call to `spfa()`, or an empty vector if `dest` is unreachable.

For path reconstruction, `pred[v]` stores the node immediately before `v` on the shortest path from `start` to `v`, or `-1` if `v` is `start` or unreachable. Follow `pred` backward from the destination to `start`, then reverse that sequence to recover the path.

Time Complexity:

- $O(n \cdot m)$ per call in the worst case, where n is the number of nodes and m is the number of edges.
- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary queue space.
- $O(p)$ for the path returned by `get_path()`.

```
#include <algorithm>
#include <cstdint>
#include <queue>
#include <utility>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::vector<std::pair<int, int>>>> adj;
std::vector<int64_t> dist;
std::vector<int> pred;

bool spfa(int start) {
    int n = static_cast<int>(adj.size());
    dist.assign(n, INF);
    pred.assign(n, -1);
    std::vector<int> path_edges(n, 0);
    std::vector<char> in_queue(n, false);
    std::queue<int> q;
    dist[start] = 0;
    q.push(start);
    in_queue[start] = true;
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        in_queue[u] = false;
        for (auto &[v, w] : adj[u]) {
            if (dist[v] > dist[u] + w) {
                dist[v] = dist[u] + w;
                pred[v] = u;
                path_edges[v] = path_edges[u] + 1;
                if (path_edges[v] >= n) {
                    return false;
                }
            }
            if (!in_queue[v]) {
                q.push(v);
            }
        }
    }
    return true;
}
```

```

        in_queue[v] = true;
    }
}
}
return true;
}

std::vector<int> get_path(int dest) {
    if (dist[dest] == INF) {
        return {};
    }
    std::vector<int> path;
    for (int v = dest; v != -1; v = pred[v]) {
        path.push_back(v);
    }
    std::reverse(path.begin(), path.end());
    return path;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //      w=4
    //    0 ----> 1
    //    |      /
    // w=5 |    / w=-2
    //    |  /
    //    v v  w=3
    //    2 -----> 3
    adj.assign(4, {});
    adj[0].emplace_back(1, 4);
    adj[0].emplace_back(2, 5);
    adj[1].emplace_back(2, -2);
    adj[2].emplace_back(3, 3);
    int start = 0, dest = 3;
    spfa(start);
    assert(dist[dest] == 5);
    assert((get_path(dest) == vector<int>{0, 1, 2, 3}));

    adj[3].emplace_back(1, -2); // The cycle 1 -> 2 -> 3 -> 1 now has total weight -1.
    assert(!spfa(start));
    return 0;
}

```

4.3.6 DAG Shortest Path (Topological DFS)

Given a starting node in a weighted, directed acyclic graph (DAG), determine the minimum distance to every reachable node. Because the graph is acyclic, its nodes admit a topological ordering in which every edge points forward. Relaxing edges in this order settles each node's distance in a single pass, with no priority queue and no restriction on edge signs. This makes it both faster than Dijkstra's algorithm and able to handle negative edge weights, as long as the graph has no cycles. The same pass computes longest paths if the relaxation comparison is reversed, which is the standard way to find the critical path in a schedule of dependent tasks.

- `dag_shortest_path(start)` populates `dist` and `pred` for a global, pre-populated adjacency list `adj` whose indices represent the nodes. Each edge is stored as `(neighbor, weight)` and may have any sign. `dist[v]` is set to `INF` for nodes not reachable from `start`, and `pred` stores the shortest-path tree for path reconstruction.

- `get_path(dest)` returns the path from `start` to `dest`, or an empty vector if `dest` is unreachable, after the latest call to `dag_shortest_path()`.

For path reconstruction, `pred[v]` stores the node immediately before `v` on the shortest path from `start` to `v`, or `-1` if `v` is `start` or unreachable. Follow `pred` backward from the destination to `start`, then reverse that sequence to recover the path.

Time Complexity:

- $O(n + m)$ per call, where n is the number of nodes and m is the number of edges.
- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph.
- $O(n)$ auxiliary.
- $O(p)$ for the path returned by `get_path()`.

```
#include <algorithm>
#include <cstdint>
#include <utility>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::vector<std::pair<int, int>>> adj;
std::vector<int64_t> dist;
std::vector<int> pred;

void topological_dfs(int u, std::vector<char> &visit, std::vector<int> &order) {
    visit[u] = true;
    for (auto &[v, w] : adj[u]) {
        if (!visit[v]) {
            topological_dfs(v, visit, order);
        }
    }
    order.push_back(u);
}

void dag_shortest_path(int start) {
    int n = static_cast<int>(adj.size());
    std::vector<char> visit(n, false);
    std::vector<int> order;
    for (int i = 0; i < n; i++) {
        if (!visit[i]) {
            topological_dfs(i, visit, order);
        }
    }
    std::reverse(order.begin(), order.end());
    dist.assign(n, INF);
    pred.assign(n, -1);
    dist[start] = 0;
    for (int u : order) {
        if (dist[u] == INF) {
            continue;
        }
        for (auto &[v, w] : adj[u]) {
            if (dist[u] + w < dist[v]) {
                dist[v] = dist[u] + w;
                pred[v] = u;
            }
        }
    }
}

std::vector<int> get_path(int dest) {
    if (dist[dest] == INF) {
        return {};
    }
}
```

```

}
std::vector<int> path;
for (int v = dest; v != -1; v = pred[v]) {
    path.push_back(v);
}
std::reverse(path.begin(), path.end());
return path;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //           w=1    w=-2    w=3
    //      0 ----> 1 ----> 3 ----> 4
    //      |      /      ^
    // w=5 |    / w=1    /
    //      v v    /w=1
    //      2 -----+
    adj.assign(5, {});
    adj[0].emplace_back(1, 1);
    adj[0].emplace_back(2, 5);
    adj[1].emplace_back(2, 1);
    adj[1].emplace_back(3, -2); // A negative edge, which Dijkstra could not handle.
    adj[2].emplace_back(3, 1);
    adj[3].emplace_back(4, 3);
    int start = 0, dest = 4;
    dag_shortest_path(start);
    assert(dist[dest] == 2);
    assert((get_path(dest) == vector<int>{0, 1, 3, 4}));
    return 0;
}

```

4.3.7 All-Pairs Shortest Path (Floyd-Warshall)

Given a weighted, directed graph with possibly negative weights, determine the minimum distance between all pairs of start and destination nodes in the graph. Optionally, reconstruct a shortest path between two nodes using the precomputed next-hop matrix `next_node`.

Floyd-Warshall's algorithm is a dynamic program over intermediate nodes: for each node k in turn, every pair (i, j) is relaxed by considering a path through k , so once all k have been processed the matrix holds the all-pairs shortest distances.

- `init_floyd(n)` initializes `dist` and `next_node` for a graph of n nodes numbered $[0, n)$.
- `floyd_warshall()` updates the global adjacency matrix `dist` so `dist[u][v]` stores the shortest-path distance from u to v , and updates `next_node` for path reconstruction. If the graph contains negative-weight cycles, there is no shortest path and an error will be thrown.
- `get_path(start, dest)` returns the shortest path from `start` to `dest`, or an empty vector if `dest` is unreachable from `start`, after the latest call to `floyd_warshall()`.

For path reconstruction, `next_node[i][j]` stores the next node to visit after i on a current shortest path from i to j . It is initialized to j for every pair and, when a shorter route $i \rightarrow k \rightarrow j$ is found, becomes `next_node[i][k]`. Repeatedly setting i to `next_node[i][j]` therefore walks the path from source to destination. This value is meaningful only when `dist[i][j] != INF`.

Time Complexity:

- $O(n^2)$ per call to `init_floyd()`, where n is the number of nodes.
- $O(n^3)$ per call to `floyd_warshall()`.

- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(n^2)$ for storage of the graph, where n is the number of nodes.
- $O(1)$ auxiliary.
- $O(p)$ for the path returned by `get_path()`.

```
#include <stdint>
#include <stdexcept>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::vector<int64_t>> dist;
std::vector<std::vector<int>> next_node;

void init_floyd(int n) {
    dist.assign(n, std::vector<int64_t>(n));
    next_node.assign(n, std::vector<int>(n));
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            dist[i][j] = (i == j) ? 0 : INF;
            next_node[i][j] = j;
        }
    }
}

void floyd_warshall() {
    int n = static_cast<int>(dist.size());
    for (int k = 0; k < n; k++) {
        for (int i = 0; i < n; i++) {
            for (int j = 0; j < n; j++) {
                // The INF guards avoid relaxing through an unreachable intermediate: with a negative edge,
                // INF + w < INF would otherwise give an unreachable pair a bogus finite distance.
                if (dist[i][k] != INF && dist[k][j] != INF && dist[i][j] > dist[i][k] + dist[k][j]) {
                    dist[i][j] = dist[i][k] + dist[k][j];
                    next_node[i][j] = next_node[i][k];
                }
            }
        }
    }
    // Optional: Report negative-weight cycles.
    for (int i = 0; i < n; i++) {
        if (dist[i][i] < 0) {
            throw std::runtime_error("Negative-weight cycle found.");
        }
    }
}

std::vector<int> get_path(int start, int dest) {
    if (dist[start][dest] == INF) {
        return {};
    }
    std::vector<int> path{start};
    while (start != dest) {
        start = next_node[start][dest];
        path.push_back(start);
    }
    return path;
}
```

Example Usage

```
#include <cassert>
using namespace std;
```

```

int main() {
    //      w=1      w=2
    // 0 -----> 1 -----> 2
    // |           ^
    // +-----+
    //      w=5
    init_floyd(3);
    dist[0][1] = 1;
    dist[1][2] = 2;
    dist[0][2] = 5;
    floyd_warshall();
    int start = 0, dest = 2;
    assert(dist[start][dest] == 3);
    assert((get_path(start, dest) == vector<int>{0, 1, 2}));
    return 0;
}

```

4.3.8 All-Pairs Shortest Path (Johnson)

Given a sparse, weighted directed graph with possibly negative edge weights, determine the minimum distance between every pair of nodes. Johnson's algorithm first computes a potential $h(v)$ using Bellman-Ford, then replaces each edge weight $w(u, v)$ with $w'(u, v) = w(u, v) + h(u) - h(v)$. Every reduced weight is nonnegative, so Dijkstra's algorithm can run from every node. Restoring the original distance afterward preserves all shortest paths because the potential terms telescope along each path.

Initializing every potential to 0 is equivalent to adding a virtual source with zero-weight edges to all nodes. Bellman-Ford therefore detects a negative cycle anywhere in the graph, not only one reachable from a particular original node. If such a cycle exists, shortest paths are undefined.

- `johnson_all_pairs(n)` populates `dist` and `next_node` for a global, pre-populated edge list `edges` whose nodes are numbered $[0, n)$. Each edge is stored as $(u, v, weight)$. Afterward, `dist[u][v]` is the shortest distance from u to v , or `INF` if v is unreachable from u . An error is thrown if the graph contains a negative-weight cycle.
- `get_path(start, dest)` returns the shortest path from `start` to `dest`, or an empty vector if `dest` is unreachable from `start`, after the latest call to `johnson_all_pairs()`.

For path reconstruction, `next_node[u][v]` stores the first node after u on a shortest path to v , or -1 if v is unreachable from u . Repeatedly advance u to `next_node[u][v]` until it equals v .

All potentials, reduced weights, and path distances must fit in `int64_t`.

Time Complexity:

- $O(n \cdot m + n \cdot (n + m) \cdot \log n)$ per call, where n is the number of nodes and m is the number of edges. For a sparse connected graph, this is $O(n \cdot m \cdot \log n)$.
- $O(p)$ per call to `get_path()`, where p is the number of nodes in the returned path.

Space Complexity:

- $O(n^2 + \max(n, m))$ for the graph, distance matrix, and auxiliary storage.
- $O(p)$ for the path returned by `get_path()`.

```

#include <stdint>
#include <functional>
#include <queue>
#include <stdexcept>
#include <tuple>
#include <utility>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::tuple<int, int, int>> edges; // (u, v, weight)

```

```

std::vector<std::vector<int64_t>> dist;
std::vector<std::vector<int>> next_node;

void johnson_all_pairs(int n) {
    std::vector<int64_t> potential(n, 0);
    for (int i = 0; i < n; i++) {
        bool changed = false;
        for (auto [u, v, w] : edges) {
            int64_t pv = potential[u] + w; // Overflow warning.
            if (pv < potential[v]) {
                if (i == n - 1) {
                    throw std::runtime_error("Negative-weight cycle found.");
                }
                potential[v] = pv;
                changed = true;
            }
        }
        if (!changed) {
            break;
        }
    }
    std::vector<std::vector<std::pair<int, int64_t>>> adj(n);
    for (auto [u, v, w] : edges) {
        int64_t reduced = static_cast<int64_t>(w) + potential[u] - potential[v]; // Overflow warning.
        adj[u].emplace_back(v, reduced);
    }
    dist.assign(n, std::vector<int64_t>(n, INF));
    next_node.assign(n, std::vector<int>(n, -1));
    using qnode = std::pair<int64_t, int>; // (distance, node)
    for (int start = 0; start < n; start++) {
        std::priority_queue<qnode, std::vector<qnode>, std::greater<>> pq;
        dist[start][start] = 0;
        next_node[start][start] = start;
        pq.emplace(0, start);
        while (!pq.empty()) {
            auto [du, u] = pq.top();
            pq.pop();
            if (du != dist[start][u]) {
                continue;
            }
            for (auto [v, w] : adj[u]) {
                int64_t dv = du + w; // Overflow warning.
                if (dv < dist[start][v]) {
                    dist[start][v] = dv;
                    next_node[start][v] = (u == start) ? v : next_node[start][u];
                    pq.emplace(dv, v);
                }
            }
        }
        for (int v = 0; v < n; v++) {
            if (dist[start][v] != INF) {
                dist[start][v] += potential[v] - potential[start]; // Overflow warning.
            }
        }
    }
}

std::vector<int> get_path(int start, int dest) {
    if (next_node[start][dest] == -1) {
        return {};
    }
    std::vector<int> path{start};
    while (start != dest) {
        start = next_node[start][dest];
        path.push_back(start);
    }
    return path;
}

```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    //
    //           w=3
    //           +-----+
    //           |         |
    //       w=1  v  w=-2   w=2  |
    // 0 -----> 1 -----> 2 -----> 3
    // |               ^
    // +-----+
    //       w=4
    edges = {{0, 1, 1}, {0, 2, 4}, {1, 2, -2}, {2, 3, 2}, {3, 1, 3}};
    johnson_all_pairs(5); // Node 4 is isolated.
    assert((dist[0] == vector<int64_t>{0, 1, -1, 1, INF}));
    assert((dist[1] == vector<int64_t>{INF, 0, -2, 0, INF}));
    assert((dist[2] == vector<int64_t>{INF, 5, 0, 2, INF}));
    assert((dist[3] == vector<int64_t>{INF, 3, 1, 0, INF}));
    assert((dist[4] == vector<int64_t>{INF, INF, INF, INF, 0}));
    assert((get_path(0, 3) == vector<int>{0, 1, 2, 3}));
    assert(get_path(0, 4).empty());

    edges = {{0, 1, -1}, {1, 0, -1}};
    bool threw = false;
    try {
        johnson_all_pairs(2);
    } catch (const runtime_error &) {
        threw = true;
    }
    assert(threw);
    return 0;
}
```

4.4 Spanning Trees

4.4.1 Minimum Spanning Tree (Prim)

For a connected, undirected, weighted graph with possibly negative weights, a minimum spanning tree (MST) connects all nodes using a subset of edges with minimum possible total weight. On disconnected input, this implementation instead finds the minimum spanning forest.

Prim's algorithm grows the tree from an arbitrary start node, repeatedly adding the minimum-weight edge that joins a new node to the current tree, with a priority queue supplying the cheapest such edge at each step.

- `prim_mst()` populates `mst` with the edges in the minimum spanning forest and returns the total MST weight for a global, bidirectionally pre-populated adjacency list `adj`, whose indices represent the nodes. Adjacency entries are stored as `(neighbor, weight)`, while each MST edge is stored as `(from, to, weight)`.

The priority queue stores candidate edges as `(weight, from, to)` and uses `std::greater` to make it a min-heap. To find a maximum spanning tree instead, use the default max-heap ordering. Multigraphs are supported; parallel edges are stored as separate adjacency entries and the algorithm automatically selects the minimum-weight one to each unvisited node.

Time Complexity: $O(n + m \log m)$ per call, where n is the number of nodes and m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(\max(n, m))$ auxiliary.


```

#include <cstdint>
#include <functional>
#include <queue>
#include <tuple>
#include <utility>
#include <vector>

std::vector<std::vector<std::pair<int, int>>>> adj; // adj[u] = {(v, weight), ...}
std::vector<std::tuple<int, int, int>> mst; // (u, v, weight)

int64_t prim_mst() {
    int n = static_cast<int>(adj.size());
    mst.clear();
    std::vector<char> visit(n);
    int64_t total_dist = 0;
    for (int i = 0; i < n; i++) {
        if (visit[i]) {
            continue;
        }
        visit[i] = true;
        using qnode = std::tuple<int, int, int>; // (weight, u, v)
        std::priority_queue<qnode, std::vector<qnode>, std::greater<>> pq;
        for (auto [v, w] : adj[i]) {
            pq.emplace(w, i, v);
        }
        while (!pq.empty()) {
            auto [w, u, v] = pq.top();
            pq.pop();
            if (visit[v]) {
                continue;
            }
            visit[v] = true;
            mst.emplace_back(u, v, w);
            total_dist += w;
            for (auto [to, ew] : adj[v]) {
                pq.emplace(ew, v, to);
            }
        }
    }
    return total_dist;
}

```

Example Usage

```

#include <cassert>
using namespace std;

void add_edge(int u, int v, int w) {
    adj[u].emplace_back(v, w);
    adj[v].emplace_back(u, w);
}

int main() {
    // Two connected components; MST routine returns a minimum spanning forest.
    //      w=4          w=1
    //  0 ---- 1      3 ---- 4
    //   \    /      / |
    //  w=3 \ / w=6   w=2 / | w=4
    //      \ /      / w=3 |
    //       2      5 ---- 6
    adj.assign(7, {});
    add_edge(0, 1, 4);
    add_edge(1, 2, 6);
    add_edge(2, 0, 3);
    add_edge(3, 4, 1);
    add_edge(4, 5, 2);
}

```

```

add_edge(5, 6, 3);
add_edge(6, 4, 4);
int64_t total = prim_mst();
assert(total == 13);
assert(
    (mst == vector<tuple<int, int, int>>{{0, 2, 3}, {0, 1, 4}, {3, 4, 1}, {4, 5, 2}, {5, 6, 3}})
);
return 0;
}

```

4.4.2 Minimum Spanning Tree (Kruskal)

For a connected, undirected, weighted graph with possibly negative weights, a minimum spanning tree (MST) connects all nodes using a subset of edges with minimum possible total weight. On disconnected input, this implementation instead finds the minimum spanning forest.

Kruskal's algorithm scans the edges in nondecreasing weight order, adding each edge whose endpoints lie in different components (tracked with a disjoint-set structure) and skipping any edge that would form a cycle.

- `kruskal_mst(n)` populates `mst` with the edge IDs in the minimum spanning forest and returns the total MST weight for a global, pre-populated edge list `edges` whose endpoints must be numbered $[0, n)$. Each edge is stored as a tuple `(weight, u, v)`.

Multigraphs are supported; parallel edges appear as separate entries in `edges` and any redundant ones are skipped once the cheaper edge has already connected the two components.

Time Complexity: $O(n + m \log m)$ per call, where n is the number of nodes and m is the number of edges. The internal DSU uses path compression and union by size, making each DSU operation $O(\alpha(n))$ amortized.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(\max(n, m))$ auxiliary heap space and $O(\log m)$ auxiliary stack space.

```

#include <algorithm>
#include <cstdint>
#include <numeric>
#include <tuple>
#include <utility>
#include <vector>

std::vector<std::tuple<int, int, int>> edges; // (weight, u, v)
std::vector<int> dsu_root, dsu_size, mst;

int find_root(int u) {
    if (dsu_root[u] != u) {
        dsu_root[u] = find_root(dsu_root[u]);
    }
    return dsu_root[u];
}

int64_t kruskal_mst(int n) {
    mst.clear();
    std::vector<int> order(edges.size());
    std::iota(order.begin(), order.end(), 0);
    std::sort(order.begin(), order.end(), [](int a, int b) {
        return std::get<0>(edges[a]) < std::get<0>(edges[b]);
    });
    int64_t total_dist = 0;
    dsu_root.assign(n, 0);
    dsu_size.assign(n, 1);
    std::iota(dsu_root.begin(), dsu_root.end(), 0);
    for (int id : order) {

```

```

    auto &[w, a, b] = edges[id];
    int u = find_root(a), v = find_root(b);
    if (u != v) {
        if (dsu_size[u] < dsu_size[v]) {
            std::swap(u, v);
        }
        dsu_root[v] = u;
        dsu_size[u] += dsu_size[v];
        mst.push_back(id);
        total_dist += w;
    }
}
return total_dist;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // Two connected components; MST routine returns a minimum spanning forest.
    //      w=4          w=1
    //  0 ---- 1      3 ---- 4
    //   \      /      / |
    // w=3 \    / w=6   w=2 / | w=4
    //      \  /      / w=3 |
    //       2      5 ---- 6
    edges.emplace_back(4, 0, 1); // (weight, u, v)
    edges.emplace_back(6, 1, 2);
    edges.emplace_back(3, 2, 0);
    edges.emplace_back(1, 3, 4);
    edges.emplace_back(2, 4, 5);
    edges.emplace_back(3, 5, 6);
    edges.emplace_back(4, 6, 4);
    int64_t total = kruskal_mst(7);
    assert(total == 13);
    assert((mst == vector<int>{3, 4, 2, 5, 0}));
    return 0;
}

```

4.4.3 Kruskal Reconstruction Tree

Given a connected, undirected, weighted graph, build the Kruskal reconstruction tree of its minimum spanning tree. Kruskal's algorithm adds MST edges in nondecreasing weight order; whenever it joins two components, create a new internal node whose children are the two component roots and whose value is the joining edge weight. The original graph nodes become leaves.

The lowest common ancestor of two leaves is exactly the merge step that first connects their MST components, so its value is the maximum edge weight on the MST path between those two nodes. This turns bottleneck path queries into ordinary LCA queries on the reconstruction tree.

- `KruskalReconstructionTree(n, edges)` builds the tree from weighted edges stored as `(weight, u, v)`.
- `root()` returns the reconstruction-tree root.
- `max_edge_on_path(u, v)` returns the maximum edge weight on the MST path between original nodes `u` and `v`.

The input graph must be connected. Parallel edges are supported and are handled naturally by Kruskal's algorithm.

Time Complexity:

- $O(m \log m + n \log n)$ for construction, where n is the number of nodes and m is the number of edges.
- $O(\log n)$ per call to `max_edge_on_path()`.

Space Complexity: $O(n \log n)$ for the reconstruction tree and LCA table, plus $O(m)$ auxiliary during construction.

```
#include <algorithm>
#include <cassert>
#include <stdint>
#include <numeric>
#include <tuple>
#include <vector>

class KruskalReconstructionTree {
    std::vector<int> dsu_root, dsu_tree_root, depth, tin, tout;
    std::vector<std::vector<int>>> up;
    std::vector<std::vector<int>>> tree;
    std::vector<int64_t> value;
    int timer = 0, root_node = -1;

    int find_root(int u) {
        if (dsu_root[u] != u) {
            dsu_root[u] = find_root(dsu_root[u]);
        }
        return dsu_root[u];
    }

    void dfs_lca(int u, int p) {
        tin[u] = timer++;
        up[0][u] = p;
        for (int k = 1; k < static_cast<int>(up.size()); k++) {
            up[k][u] = up[k - 1][up[k - 1][u]];
        }
        for (int v : tree[u]) {
            depth[v] = depth[u] + 1;
            dfs_lca(v, u);
        }
        tout[u] = timer;
    }

    bool is_ancestor(int u, int v) const { return tin[u] <= tin[v] && tout[v] <= tout[u]; }

    int lca(int u, int v) const {
        if (is_ancestor(u, v)) {
            return u;
        }
        if (is_ancestor(v, u)) {
            return v;
        }
        for (int k = static_cast<int>(up.size()) - 1; k >= 0; k--) {
            if (!is_ancestor(up[k][u], v)) {
                u = up[k][u];
            }
        }
        return up[0][u];
    }

public:
    KruskalReconstructionTree(int n, std::vector<std::tuple<int64_t, int, int>> edges)
        : dsu_root(2 * n - 1),
          dsu_tree_root(2 * n - 1),
          depth(2 * n - 1),
          tin(2 * n - 1),
          tout(2 * n - 1),
          tree(2 * n - 1),
          value(2 * n - 1) {
```

```

assert(n > 0);
std::iota(dsu_root.begin(), dsu_root.begin() + n, 0);
std::iota(dsu_tree_root.begin(), dsu_tree_root.begin() + n, 0);
std::sort(edges.begin(), edges.end());
int nodes = n;
for (const auto &[w, a, b] : edges) {
    int u = find_root(a), v = find_root(b);
    if (u == v) {
        continue;
    }
    value[nodes] = w;
    tree[nodes].push_back(dsu_tree_root[u]);
    tree[nodes].push_back(dsu_tree_root[v]);
    dsu_root[u] = nodes;
    dsu_root[v] = nodes;
    dsu_root[nodes] = nodes;
    dsu_tree_root[nodes] = nodes;
    nodes++;
}
assert(nodes == 2 * n - 1);
root_node = nodes - 1;
int lg = 1;
while ((1 << lg) <= nodes) {
    lg++;
}
up.assign(lg, std::vector<int>(nodes));
dfs_lca(root_node, root_node);
}

int root() const { return root_node; }
int64_t max_edge_on_path(int u, int v) const { return value[lca(u, v)]; }
};

```

Example Usage

```

int main() {
    //          0
    //      w=2 / | w=8
    //          / |
    //      1-----2
    //      | w=4 /
    // w=5 | /
    //      | / w=9
    //      3
    std::vector<std::tuple<int64_t, int, int>> edges{
        {2, 0, 1}, {4, 1, 2}, {5, 1, 3}, {8, 0, 2}, {9, 2, 3}
    };
    KruskalReconstructionTree tree(4, edges);

    // Reconstruction tree; internal nodes are labeled with their joining edge weight.
    //          6 (5)
    //          / |
    //      5 (4) 3
    //      / |
    //  4 (2) 2
    //  / |
    // 0 1
    assert(tree.root() == 6);
    assert(tree.max_edge_on_path(0, 2) == 4);
    assert(tree.max_edge_on_path(0, 3) == 5);
    assert(tree.max_edge_on_path(2, 3) == 5);
    return 0;
}

```

4.4.4 Minimum Spanning Arborescence (Edmonds)

Given a weighted, directed graph and a root node, a spanning arborescence is a set of edges forming a tree in which every other node is reachable from the root along directed edges. Equivalently, every node except the root has exactly one incoming edge. The minimum spanning arborescence minimizes the total edge weight and is the directed analog of the minimum spanning tree.

Edmonds' (a.k.a. the Chu-Liu/Edmonds) algorithm computes the arborescence by first selecting, for each non-root node, its cheapest incoming edge. If these choices form no cycle, they already constitute the answer. Otherwise each cycle is contracted into a single supernode, and every edge entering the cycle has its weight reduced by the cost of the edge it would replace inside the cycle. Solving the problem on the contracted graph and expanding the cycles back yields the optimum. The implementation below repeats this contraction in rounds, accumulating the selected weights, until no cycle remains.

- `directed_mst(n, edges, root)` returns the total weight of the minimum spanning arborescence rooted at `root` over `n` nodes numbered `[0, n)`, or `std::nullopt` if no arborescence exists. Edges are given as `(from, to, weight)` triples; parallel edges and self-loops are allowed, with self-loops simply ignored.

Time Complexity: $O(n \cdot m)$ per call, where n is the number of nodes and m is the number of edges.

Space Complexity: $O(n + m)$ auxiliary.

```
#include <cstdint>
#include <optional>
#include <tuple>
#include <vector>

using Edge = std::tuple<int, int, int64_t>; // (u, v, weight)

std::optional<int64_t> directed_mst(int n, std::vector<Edge> edges, int root) {
    const int64_t INF = INT64_MAX / 4;
    int64_t total_dist = 0;
    while (true) {
        std::vector<int64_t> min_in(n, INF);
        std::vector<int> pre(n, -1);
        for (const auto &[u, v, w] : edges) {
            if (u != v && w < min_in[v]) {
                min_in[v] = w;
                pre[v] = u;
            }
        }
        for (int v = 0; v < n; v++) {
            if (v != root && min_in[v] == INF) {
                return std::nullopt;
            }
        }
        // Identify cycles formed by the chosen incoming edges, labeling each with an ID in comp.
        int cycles = 0;
        std::vector<int> comp(n, -1), seen(n, -1);
        min_in[root] = 0;
        for (int v = 0; v < n; v++) {
            total_dist += min_in[v];
            int u = v;
            while (seen[u] != v && comp[u] == -1 && u != root) {
                seen[u] = v;
                u = pre[u];
            }
            if (u != root && comp[u] == -1) {
                for (int w = pre[u]; w != u; w = pre[w]) {
                    comp[w] = cycles;
                }
                comp[u] = cycles++;
            }
        }
        if (cycles == 0) {
```

```

    break; // No cycle remains; the selected edges form the arborescence.
}
// Contract each cycle into a supernode and reduce the weights of edges entering it.
for (int v = 0; v < n; v++) {
    if (comp[v] == -1) {
        comp[v] = cycles++;
    }
}
for (auto &[u, v, w] : edges) {
    int cu = comp[u], cv = comp[v];
    if (cu != cv) {
        w -= min_in[v];
    }
    u = cu;
    v = cv;
}
n = cycles;
root = comp[root];
}
return total_dist;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //      w=1
    // root=0 -----> 1 <-----+
    //      |      / |      |
    //      |      w=2 |      |
    // w=5| /      |w=2 |w=4
    //      v v      v
    //      2 -----> 3 -----+
    //      w=3
    vector<Edge> edges{
        {0, 1, 1}, {0, 2, 5}, {1, 2, 2}, {2, 3, 3}, {3, 1, 4},
    };
    auto weight = directed_mst(4, edges, 0);
    assert(weight == 6); // Cheapest edges 0->1->2->3 form a tree of weight 6.

    //      w=1      w=2
    // root=0 ----> 1 ----> 2      3
    vector<Edge> disconnected{{0, 1, 1}, {1, 2, 1}};
    auto disconnected_weight = directed_mst(4, disconnected, 0);
    assert(!disconnected_weight); // Node 3 is unreachable from root 0.

    vector<Edge> negative{{0, 1, -1}};
    assert(directed_mst(2, negative, 0) == -1); // A weight of -1 is a valid result.
    return 0;
}

```

4.5 Flows and Cuts

4.5.1 Maximum Flow (Ford-Fulkerson)

Given a flow network with integer capacities, find the maximum flow from a given source node to a given sink node. The flow along each edge may not exceed its capacity, and flow is conserved at every node other than the source and sink.

The Ford-Fulkerson method repeatedly finds any augmenting path from source to sink in the residual graph (here by depth-first search) and pushes as much flow as the path allows, until no augmenting path remains. This algorithm should only be used on graphs with integer capacities, because some real-valued inputs can make it fail to terminate. The Edmonds-Karp algorithm instead uses breadth-first search and always terminates in polynomial time.

- `max_flow(source, sink)` modifies the global residual capacity matrix `cap` and returns the maximum flow. Nodes are numbered $[0, n)$, where n is `cap.size()`.

Time Complexity: $O(n^2 \cdot f)$ per call, where n is the number of nodes and f is the maximum flow.

Space Complexity:

- $O(n^2)$ for storage of the flow network, where n is the number of nodes.
- $O(n)$ auxiliary stack space.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::vector<int64_t>> cap;
std::vector<char> visit;

int64_t dfs(int u, int64_t f, int sink) {
    int n = static_cast<int>(cap.size());
    if (u == sink) {
        return f;
    }
    visit[u] = true;
    for (int v = 0; v < n; v++) {
        if (!visit[v] && cap[u][v] > 0) {
            int64_t flow = dfs(v, std::min(f, cap[u][v]), sink);
            if (flow > 0) {
                cap[u][v] -= flow;
                cap[v][u] += flow;
                return flow;
            }
        }
    }
    return 0;
}

int64_t max_flow(int source, int sink) {
    int n = static_cast<int>(cap.size());
    assert(source != sink);
    int64_t total = 0;
    while (true) {
        visit.assign(n, false);
        int64_t flow = dfs(source, INF, sink);
        if (flow == 0) {
            break;
        }
        total += flow;
    }
    return total;
}
```

Example Usage

```
#include <cassert>

int main() {
    // Example graph after max flow, with each edge labeled flow/capacity:
```



```

//          2/2
//      1 -----> 3
//      / \      |
// 3/4 /   \ 1/1 | 2/2
//    /     v     v
//  0       4 -----> 5
//    \     ^     3/3
// 2/3 \   / 2/2
//      v /
//      2
int nodes = 6;
cap.assign(nodes, std::vector<int64_t>(nodes));
cap[0][1] = 4;
cap[0][2] = 3;
cap[1][3] = 2;
cap[1][4] = 1;
cap[2][4] = 2;
cap[3][5] = 2;
cap[4][5] = 3;
assert(max_flow(0, 5) == 5);
return 0;
}

```

4.5.2 Maximum Flow (Edmonds-Karp)

Given a flow network with integer or floating-point capacities, find the maximum flow from a given source node to a given sink node. The flow along each edge may not exceed its capacity, and flow is conserved at every node other than the source and sink.

Edmonds-Karp is the Ford-Fulkerson method with the augmenting path always chosen as a shortest one (fewest edges) via breadth-first search. This bounds the number of augmentations at $O(n \cdot m)$, independent of the capacity magnitudes.

- `EdmondsKarp<T>(n)` constructs an empty flow network with nodes numbered $[0, n)$.
- `add_edge(u, v, cap)` adds a directed residual-network edge from `u` to `v`.
- `max_flow(source, sink)` modifies the residual network and returns the maximum flow.
- `clear_flow()` resets all edge flows to zero, allowing a fresh recomputation on the same graph.
- `min_cut(source)` returns the source side of a minimum cut after `max_flow()` has been called.

Repeated calls to `max_flow()` continue augmenting from the current flow. This is useful after adding new edges or increasing capacities; call `clear_flow()` first to recompute from zero. The capacity type `T` should be signed, since reverse residual edges store negative flow.

Time Complexity:

- $O(n \cdot m^2)$ per call to `max_flow()`, where n is the number of nodes and m is the number of edges. For integer capacities, this is also $O(m \cdot f)$, where f is the maximum flow.
- $O(n + m)$ per call to `min_cut()`.

Space Complexity: $O(\max(n, m))$ for storage of the flow network, where n is the number of nodes and m is the number of edges.

```

#include <algorithm>
#include <cassert>
#include <limits>
#include <queue>
#include <vector>

template<typename T>
class EdmondsKarp {
    struct Edge {
        int u, v, rev;

```

```

    T cap, flow;
};

static constexpr T EPS = static_cast<T>(1e-9);
static bool residual(const Edge &e) { return e.cap - e.flow > EPS; }

int nodes;
T flow;
std::vector<std::vector<Edge>> adj;

public:
explicit EdmondsKarp(int n) : nodes(n), flow(0), adj(n) {}

void add_edge(int u, int v, T cap) {
    adj[u].push_back(Edge{u, v, static_cast<int>(adj[v].size()), cap, 0});
    adj[v].push_back(Edge{v, u, static_cast<int>(adj[u].size()) - 1, 0, 0});
}

void clear_flow() {
    for (auto &edges : adj) {
        for (Edge &e : edges) {
            e.flow = 0;
        }
    }
    flow = 0;
}

T max_flow(int source, int sink) {
    assert(source != sink);
    while (true) {
        std::vector<Edge*> pred(nodes, nullptr);
        std::queue<int> q;
        q.push(source);
        while (!q.empty() && pred[sink] == nullptr) {
            int u = q.front();
            q.pop();
            for (Edge &e : adj[u]) {
                if (e.v != source && pred[e.v] == nullptr && residual(e)) {
                    pred[e.v] = &e;
                    q.push(e.v);
                }
            }
        }
        if (pred[sink] == nullptr) {
            break;
        }
        T aug = std::numeric_limits<T>::max();
        for (int u = sink; u != source; u = pred[u]->u) {
            aug = std::min(aug, pred[u]->cap - pred[u]->flow);
        }
        for (int u = sink; u != source; u = pred[u]->u) {
            pred[u]->flow += aug;
            adj[pred[u]->v][pred[u]->rev].flow -= aug;
        }
        flow += aug;
    }
    return flow;
}

std::vector<char> min_cut(int source) const {
    std::vector<char> reachable(nodes);
    std::queue<int> q;
    reachable[source] = true;
    q.push(source);
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        for (const Edge &e : adj[u]) {

```

```

        if (!reachable[e.v] && residual(e)) {
            reachable[e.v] = true;
            q.push(e.v);
        }
    }
}
return reachable;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // Example graph after max flow, with each edge labeled flow/capacity:
    //           2/2
    //           1 -----> 3
    //          / \      |
    // 3/4 /   \ 1/1   | 2/2
    //   /       v       v
    //  0       4 -----> 5
    //   \      ^   3/3
    // 2/3 \   / 2/2
    //       v /
    //       2
    EdmondsKarp<int> g(6);
    g.add_edge(0, 1, 4);
    g.add_edge(0, 2, 3);
    g.add_edge(1, 3, 2);
    g.add_edge(1, 4, 1);
    g.add_edge(2, 4, 2);
    g.add_edge(3, 5, 2);
    g.add_edge(4, 5, 3);
    int flow = g.max_flow(0, 5);
    assert(flow == 5);
    assert(g.max_flow(0, 5) == 5);
    vector<char> cut = g.min_cut(0);
    assert((cut == vector<char>{true, true, true, false, false, false}));
    g.clear_flow();
    assert(g.max_flow(0, 5) == 5);
    return 0;
}

```

4.5.3 Maximum Flow and Decomposition (Dinic)

Given a flow network with integer or floating-point capacities, find the maximum flow from a given source node to a given sink node. The flow along each edge may not exceed its capacity, and flow is conserved at every node other than the source and sink.

Dinic's algorithm proceeds in phases. Each phase runs a breadth-first search to build a level graph of shortest residual distances from the source, then finds a blocking flow that saturates that level graph using depth-first search before repeating on the new residual graph. Residual edges are stored in a flat array so callers can inspect edge flows and decompose the final flow into source-to-sink paths.

- `Dinic<T>(n)` constructs an empty flow network with nodes numbered $[0, n)$.
- `add_edge(u, v, cap)` adds a directed residual-network edge from `u` to `v` and returns its edge ID.
- `edge_flow(id)` returns the flow through a previously added edge.
- `max_flow(source, sink)` modifies the residual network and returns the maximum flow.
- `clear_flow()` resets all edge flows to zero, allowing a fresh recomputation on the same graph.

- `min_cut(source)` returns the source side of a minimum cut after `max_flow()` has been called.
- `decompose(source, sink)` decomposes the current flow into paths from `source` to `sink`.

Repeated calls to `max_flow()` continue augmenting from the current flow. This is useful after adding new edges or increasing capacities; call `clear_flow()` first to recompute from zero. The capacity type `T` should be signed, since reverse residual edges store negative flow.

Time Complexity:

- $O(n^2 \cdot m)$ per call to `max_flow()`, where n is the number of nodes and m is the number of edges.
- $O(n + m)$ per call to `min_cut()`.
- $O(p \cdot m)$ per call to `decompose()`, where p is the number of paths returned.

Space Complexity:

- $O(\max(n, m))$ for storage of the flow network.
- $O(n)$ auxiliary stack space and $O(n)$ auxiliary heap space for `max_flow()`.
- $O(n + m)$ auxiliary heap space for `decompose()`, not counting the returned paths.

```
#include <algorithm>
#include <cassert>
#include <functional>
#include <limits>
#include <queue>
#include <vector>

template<typename T>
class Dinic {
    struct Edge {
        int u, v;
        T cap, flow;
    };

    static constexpr T EPS = static_cast<T>(1e-9);
    static bool residual(const Edge &e) { return e.cap - e.flow > EPS; }

    int nodes;
    T flow;
    std::vector<Edge> edges;
    std::vector<std::vector<int>>> adj;
    std::vector<int> dist, ptr;

    bool bfs(int source, int sink) {
        dist.assign(nodes, -1);
        dist[source] = 0;
        std::queue<int> q;
        q.push(source);
        while (!q.empty()) {
            int u = q.front();
            q.pop();
            for (int id : adj[u]) {
                const Edge &e = edges[id];
                if (dist[e.v] < 0 && residual(e)) {
                    dist[e.v] = dist[u] + 1;
                    q.push(e.v);
                }
            }
        }
        return dist[sink] >= 0;
    }

    T dfs(int u, T pushed, int sink) {
        if (u == sink || pushed <= EPS) {
            return pushed;
        }
        for (; ptr[u] < static_cast<int>(adj[u].size()); ptr[u]++) {
            int id = adj[u][ptr[u]];
            Edge &e = edges[id];
            T res = 0;
            if (dist[e.v] == dist[u] + 1) {
                T cur = min(e.cap - e.flow, pushed);
                e.flow += cur;
                res += cur;
                pushed -= cur;
            }
        }
        return res;
    }

    T max_flow(int source, int sink) {
        T res = 0;
        while (bfs(source, sink)) {
            res += dfs(source, INF, sink);
        }
        return res;
    }

    std::vector<int> min_cut(int source) {
        max_flow(source, sink);
        std::vector<int> sset(nodes);
        for (int v = 0; v < nodes; v++) {
            if (dist[v] < 0) sset[v] = 1;
        }
        return sset;
    }

    std::vector<std::vector<int>>> decompose(int source, int sink) {
        max_flow(source, sink);
        std::vector<std::vector<int>>> paths;
        std::queue<int> q;
        q.push(source);
        while (!q.empty()) {
            int u = q.front();
            q.pop();
            for (int id : adj[u]) {
                const Edge &e = edges[id];
                if (dist[e.v] == dist[u] + 1 && e.cap - e.flow > 0) {
                    paths.push_back({u, e.v});
                    q.push(e.v);
                }
            }
        }
        return paths;
    }

    void clear_flow() {
        for (Edge &e : edges) e.flow = 0;
    }
};
```

```

    Edge &e = edges[id];
    if (dist[e.v] == dist[u] + 1 && residual(e)) {
        T next = dfs(e.v, std::min(pushed, e.cap - e.flow), sink);
        if (next > EPS) {
            e.flow += next;
            edges[id ^ 1].flow -= next;
            return next;
        }
    }
}
return 0;
}

public:
    struct FlowPath {
        T flow;
        std::vector<int> nodes, edges;
    };

    explicit Dinic(int n) : nodes(n), flow(0), adj(n), dist(n), ptr(n) {}

    int add_edge(int u, int v, T cap) {
        int id = static_cast<int>(edges.size());
        adj[u].push_back(id);
        edges.push_back(Edge{u, v, cap, 0});
        adj[v].push_back(id + 1);
        edges.push_back(Edge{v, u, 0, 0});
        return id;
    }

    T edge_flow(int id) const { return edges[id].flow; }

    void clear_flow() {
        for (Edge &e : edges) {
            e.flow = 0;
        }
        flow = 0;
    }

    T max_flow(int source, int sink) {
        assert(source != sink);
        while (bfs(source, sink)) {
            ptr.assign(nodes, 0);
            while (true) {
                T pushed = dfs(source, std::numeric_limits<T>::max(), sink);
                if (pushed <= EPS) {
                    break;
                }
                flow += pushed;
            }
        }
        return flow;
    }

    std::vector<char> min_cut(int source) const {
        std::vector<char> reachable(nodes);
        std::queue<int> q;
        reachable[source] = true;
        q.push(source);
        while (!q.empty()) {
            int u = q.front();
            q.pop();
            for (int id : adj[u]) {
                const Edge &e = edges[id];
                if (!reachable[e.v] && residual(e)) {
                    reachable[e.v] = true;
                    q.push(e.v);
                }
            }
        }
    }

```

```

    }
}
return reachable;
}

std::vector<FlowPath> decompose(int source, int sink) const {
    std::vector<T> rem(edges.size());
    for (int id = 0; id < static_cast<int>(edges.size()); id++) {
        rem[id] = edges[id].flow;
    }
    std::vector<FlowPath> res;
    while (true) {
        FlowPath path{std::numeric_limits<T>::max(), {source}, {}};
        std::vector<char> seen(nodes);
        std::function<bool(int)> find_path = [&](int u) {
            if (u == sink) {
                return true;
            }
            seen[u] = true;
            for (int id : adj[u]) {
                const Edge &e = edges[id];
                if (rem[id] > EPS && !seen[e.v]) {
                    path.edges.push_back(id);
                    path.nodes.push_back(e.v);
                    if (find_path(e.v)) {
                        return true;
                    }
                    path.nodes.pop_back();
                    path.edges.pop_back();
                }
            }
            return false;
        };
        if (!find_path(source)) {
            break;
        }
        for (int id : path.edges) {
            path.flow = std::min(path.flow, rem[id]);
        }
        for (int id : path.edges) {
            rem[id] -= path.flow;
        }
        res.push_back(path);
    }
    return res;
}
};

```

Example Usage

```

#include <cassert>
#include <iostream>
using namespace std;

int main() {
    // Example graph after max flow, with each edge labeled flow/capacity:
    //
    //           1 -----> 3
    //          / \         |
    // 3/4 /   \ 1/1   | 2/2
    //   /     v       v
    // 0      4 -----> 5
    //   \     ^   3/3
    // 2/3 \   / 2/2
    //      v /
    //      2

```

```

Dinic<int> g(6);
int id01 = g.add_edge(0, 1, 4);
int id02 = g.add_edge(0, 2, 3);
int id13 = g.add_edge(1, 3, 2);
int id14 = g.add_edge(1, 4, 1);
int id24 = g.add_edge(2, 4, 2);
int id35 = g.add_edge(3, 5, 2);
int id45 = g.add_edge(4, 5, 3);
int flow = g.max_flow(0, 5);
cout << "Max flow from 0 to 5: " << flow << endl;
assert(flow == 5);
assert(g.max_flow(0, 5) == 5);
assert(g.edge_flow(id01) == 3);
assert(g.edge_flow(id02) == 2);
assert(g.edge_flow(id13) == 2);
assert(g.edge_flow(id14) == 1);
assert(g.edge_flow(id24) == 2);
assert(g.edge_flow(id35) == 2);
assert(g.edge_flow(id45) == 3);
vector<char> cut = g.min_cut(0);
assert(cut[0] && !cut[5]);
auto paths = g.decompose(0, 5);
cout << "Flow decomposition:" << endl;
int total = 0;
for (const auto &path : paths) {
    cout << path.flow << ": ";
    for (int i = 0; i < static_cast<int>(path.nodes.size()); i++) {
        if (i > 0) {
            cout << " -> ";
        }
        cout << path.nodes[i];
    }
    cout << endl;
    assert(path.nodes.front() == 0 && path.nodes.back() == 5);
    assert(path.edges.size() + 1 == path.nodes.size());
    total += path.flow;
}
assert(total == 5);
g.clear_flow();
assert(g.max_flow(0, 5) == 5);
return 0;
}

```

Example Output

```

Max flow from 0 to 5: 5
Flow decomposition:
2: 0 -> 1 -> 3 -> 5
1: 0 -> 1 -> 4 -> 5
2: 0 -> 2 -> 4 -> 5

```

4.5.4 Maximum Flow (Push-Relabel)

Given a flow network with integer capacities, find the maximum flow from a given source node to a given sink node. The flow along each edge may not exceed its capacity, and flow is conserved at every node other than the source and sink.

Rather than augmenting along whole paths, the push-relabel algorithm maintains a preflow and a height label on each node, repeatedly pushing excess flow to lower-labeled neighbors and relabeling (raising) a node whose excess cannot yet be pushed, until no node other than the source and sink holds excess.

- `push_relabel(source, sink)` returns the maximum flow for a global capacity matrix `cap`, whose row and column indices represent the nodes.
- `min_cut(source)` returns the source side of a minimum cut after `push_relabel()` has been called.

Unlike Ford-Fulkerson, the $O(n^3)$ bound does not improve when the integral maximum flow f is small; the matrix-based Ford-Fulkerson implementation runs in $O(n^2 \cdot f)$.

Time Complexity:

- $O(n^3)$ per call to `push_relabel()`, where n is the number of nodes.
- $O(n^2)$ per call to `min_cut()`.

Space Complexity:

- $O(n^2)$ for storage of the flow network, where n is the number of nodes.
- $O(n)$ auxiliary for `push_relabel()`.

```
#include <algorithm>
#include <cassert>
#include <climits>
#include <cstdint>
#include <queue>
#include <vector>

std::vector<std::vector<int64_t>> cap;
std::vector<std::vector<int64_t>> f;

int64_t push_relabel(int source, int sink) {
    int n = static_cast<int>(cap.size());
    assert(source != sink);
    f.assign(n, std::vector<int64_t>(n, 0));
    std::vector<int64_t> excess(n, 0);
    std::vector<int> height(n, 0), max_height(n, 0);
    height[source] = n - 1;
    for (int i = 0; i < n; i++) {
        f[source][i] = cap[source][i];
        f[i][source] = -f[source][i];
        excess[i] = cap[source][i];
    }
    int size = 0;
    while (true) {
        if (size == 0) {
            for (int i = 0; i < n; i++) {
                if (i != source && i != sink && excess[i] > 0) {
                    if (size != 0 && height[i] > height[max_height[0]]) {
                        size = 0;
                    }
                    max_height[size++] = i;
                }
            }
        }
        if (size == 0) {
            break;
        }
        while (size != 0) {
            int i = max_height[size - 1];
            bool pushed = false;
            for (int j = 0; j < n && excess[i] != 0; j++) {
                if (height[i] == height[j] + 1 && cap[i][j] - f[i][j] > 0) {
                    int64_t df = std::min(cap[i][j] - f[i][j], excess[i]);
                    f[i][j] += df;
                    f[j][i] -= df;
                    excess[i] -= df;
                    excess[j] += df;
                    if (excess[i] == 0) {
                        size--;
                    }
                    pushed = true;
                }
            }
            if (pushed) {
                continue;
            }
        }
    }
}
```



```

    }
    height[i] = INT_MAX / 2;
    for (int j = 0; j < n; j++) {
        if (height[i] > height[j] + 1 && cap[i][j] - f[i][j] > 0) {
            height[i] = height[j] + 1;
        }
    }
    if (height[i] > height[max_height[0]]) {
        size = 0;
        break;
    }
}
}
int64_t max_flow = 0;
for (int i = 0; i < n; i++) {
    max_flow += f[source][i];
}
return max_flow;
}

std::vector<char> min_cut(int source) {
    int n = static_cast<int>(cap.size());
    std::vector<char> reachable(n);
    std::queue<int> q;
    reachable[source] = true;
    q.push(source);
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        for (int v = 0; v < n; v++) {
            if (!reachable[v] && cap[u][v] - f[u][v] > 0) {
                reachable[v] = true;
                q.push(v);
            }
        }
    }
    return reachable;
}

```

Example Usage

```

#include <cassert>

int main() {
    // Example graph after max flow, with each edge labeled flow/capacity:
    //           2/2
    //           1 -----> 3
    //          / \      |
    // 3/4 /   \ 1/1   | 2/2
    //   /     v     v
    //  0      4 ----> 5
    //   \     ^   3/3
    // 2/3 \   / 2/2
    //      v /
    //      2
    int nodes = 6;
    cap.assign(nodes, std::vector<int64_t>(nodes));
    cap[0][1] = 4;
    cap[0][2] = 3;
    cap[1][3] = 2;
    cap[1][4] = 1;
    cap[2][4] = 2;
    cap[3][5] = 2;
    cap[4][5] = 3;
    assert(push_relabel(0, 5) == 5);
    std::vector<char> cut = min_cut(0);
}

```

```

    assert(cut[0] && !cut[5]);
    return 0;
}

```

4.5.5 Flow with Lower Bounds

Given a directed network where each edge has lower and upper capacity bounds, find whether the lower bounds can be satisfied, and optionally optimize an s - t flow value. Reserve each edge's lower bound, leaving the difference between its bounds as residual capacity, then record the forced imbalance: for an edge from u to v , its lower-bound flow must leave u and enter v . A super-source supplies every node with net demand, and every node with net surplus sends that surplus to a super-sink. The bounds are feasible exactly when all such auxiliary edges can be saturated.

For s - t flow, add an infinite edge from t back to s before checking feasibility. The flow on that artificial edge is one feasible value. After the feasibility check, augmenting from s to t maximizes the value; augmenting from t to s cancels as much flow as possible and minimizes it.

- `BoundedFlow(n)` constructs a directed lower-bound flow network with nodes in $[0, n)$.
- `add_edge(u, v, lo, hi)` adds an edge with lower capacity `lo` and upper capacity `hi`.
- `feasible_circulation()` returns whether all edge bounds can be satisfied with flow conserved at every node.
- `max_flow(source, sink)` returns the maximum feasible flow from `source` to `sink`, or `std::nullopt` if no feasible flow exists.
- `min_flow(source, sink)` returns the minimum feasible flow from `source` to `sink`, or `std::nullopt` if no feasible flow exists.
- `edge_flows()` returns one feasible flow value for each original edge after a successful call.

All capacities should be nonnegative integers and `lo` $\leq$ `hi`. Choose `INF` larger than any possible finite flow value. Each feasibility or optimization call starts from the original network, so different variants may be solved successively on the same instance.

Time Complexity: $O(n^2 \cdot (n + m))$ per feasibility check or optimized flow call, where n is the number of nodes and m is the number of edges.

Space Complexity: $O(\max(n, m))$ for network storage and auxiliary arrays.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <optional>
#include <queue>
#include <vector>

class BoundedFlow {
    struct Edge {
        int v, rev;
        int64_t cap;
    };

    struct OriginalEdge {
        int u, v, index;
        int64_t lo, hi;
    };

    static constexpr int64_t INF = (1LL << 60);
    int nodes;
    std::vector<std::vector<Edge>> adj;
    std::vector<OriginalEdge> original;
    std::vector<int> level, ptr;

    void add_residual_edge(int u, int v, int64_t cap) {
        adj[u].push_back(Edge{v, static_cast<int>(adj[v].size()), cap});
    }
};

```

```

    adj[v].push_back(Edge{u, static_cast<int>(adj[u].size() - 1, 0)});
}

bool bfs(int source, int sink) {
    level.assign(static_cast<int>(adj.size()), -1);
    level[source] = 0;
    std::queue<int> q;
    q.push(source);
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        for (const Edge &e : adj[u]) {
            if (e.cap > 0 && level[e.v] == -1) {
                level[e.v] = level[u] + 1;
                q.push(e.v);
            }
        }
    }
    return level[sink] != -1;
}

int64_t dfs(int u, int sink, int64_t pushed) {
    if (u == sink || pushed == 0) {
        return pushed;
    }
    for (; ptr[u] < static_cast<int>(adj[u].size()); ptr[u]++) {
        Edge &e = adj[u][ptr[u]];
        if (e.cap > 0 && level[e.v] == level[u] + 1) {
            int64_t next = dfs(e.v, sink, std::min(pushed, e.cap));
            if (next > 0) {
                e.cap -= next;
                adj[e.v][e.rev].cap += next;
                return next;
            }
        }
    }
    return 0;
}

int64_t dinic(int source, int sink) {
    int64_t flow = 0;
    while (bfs(source, sink)) {
        ptr.assign(static_cast<int>(adj.size()), 0);
        while (int64_t pushed = dfs(source, sink, INF)) {
            flow += pushed;
        }
    }
    return flow;
}

bool satisfy_demands(int source, int sink, int &artificial_index, int64_t &forced_flow) {
    int super_source = static_cast<int>(adj.size());
    int super_sink = super_source + 1;
    adj.resize(super_sink + 1);
    level.resize(super_sink + 1);
    ptr.resize(super_sink + 1);
    std::vector<int64_t> balance(nodes);
    for (const OriginalEdge &e : original) {
        balance[e.u] -= e.lo;
        balance[e.v] += e.lo;
    }
    if (source != -1) {
        artificial_index = static_cast<int>(adj[sink].size());
        add_residual_edge(sink, source, INF);
    }
    int64_t need = 0;
    for (int u = 0; u < nodes; u++) {
        if (balance[u] > 0) {

```

```

        add_residual_edge(super_source, u, balance[u]);
        need += balance[u];
    } else if (balance[u] < 0) {
        add_residual_edge(u, super_sink, -balance[u]);
    }
}
bool ok = dinic(super_source, super_sink) == need;
forced_flow = source == -1 ? 0 : adj[source][adj[sink][artificial_index].rev].cap;
return ok;
}

std::optional<int64_t> optimize_flow(int source, int sink, bool maximize) {
    int artificial_index = -1;
    int64_t value = 0;
    if (!satisfy_demands(source, sink, artificial_index, value)) {
        return std::nullopt;
    }
    Edge &artificial = adj[sink][artificial_index];
    Edge &reverse = adj[source][artificial.rev];
    artificial.cap = reverse.cap = 0;
    return maximize ? value + dinic(source, sink) : value - dinic(sink, source);
}

void reset() {
    adj.assign(nodes, {});
    level.assign(nodes, 0);
    ptr.assign(nodes, 0);
    for (OriginalEdge &e : original) {
        e.index = static_cast<int>(adj[e.u].size());
        add_residual_edge(e.u, e.v, e.hi - e.lo);
    }
}

public:
explicit BoundedFlow(int n) : nodes(n), adj(n), level(n), ptr(n) {}

void add_edge(int u, int v, int64_t lo, int64_t hi) {
    assert(0 <= u && u < nodes && 0 <= v && v < nodes);
    assert(0 <= lo && lo <= hi);
    original.push_back(OriginalEdge{u, v, static_cast<int>(adj[u].size()), lo, hi});
    add_residual_edge(u, v, hi - lo);
}

bool feasible_circulation() {
    reset();
    int artificial_index = -1;
    int64_t value = 0;
    return satisfy_demands(-1, -1, artificial_index, value);
}

std::optional<int64_t> max_flow(int source, int sink) {
    assert(0 <= source && source < nodes && 0 <= sink && sink < nodes && source != sink);
    reset();
    return optimize_flow(source, sink, true);
}

std::optional<int64_t> min_flow(int source, int sink) {
    assert(0 <= source && source < nodes && 0 <= sink && sink < nodes && source != sink);
    reset();
    return optimize_flow(source, sink, false);
}

std::vector<int64_t> edge_flows() const {
    std::vector<int64_t> flow;
    flow.reserve(original.size());
    for (const OriginalEdge &e : original) {
        const Edge &forward = adj[e.u][e.index];
        flow.push_back(e.hi - forward.cap);
    }
}

```

```

    }
    return flow;
}
};

```

Example Usage

```

int main() {
    //           [1,4]       [1,3]
    //      0 -----> 1 -----> 3
    //      |               ^
    // [0,2] |               | [0,2]
    //      v               |
    //      2 -----+-----+
    BoundedFlow network(4);
    network.add_edge(0, 1, 1, 4);
    network.add_edge(0, 2, 0, 2);
    network.add_edge(1, 3, 1, 3);
    network.add_edge(2, 3, 0, 2);
    assert(network.max_flow(0, 3) == 5);
    std::vector<int64_t> max_flow = network.edge_flows();
    assert(max_flow[0] + max_flow[1] == 5);

    assert(network.min_flow(0, 3) == 1);

    //           [2,3]
    //      0 -----> 1
    //      ^               |
    // [2,3] |               | [2,3]
    //      |               |
    //      2 <-----+-----+
    BoundedFlow circulation(3);
    circulation.add_edge(0, 1, 2, 3);
    circulation.add_edge(1, 2, 2, 3);
    circulation.add_edge(2, 0, 2, 3);
    assert(circulation.feasible_circulation());
    assert((circulation.edge_flows() == std::vector<int64_t>{2, 2, 2}));
    return 0;
}

```

4.5.6 Minimum-Cost Maximum-Flow (SPFA)

Given a flow network with capacities and edge costs, find a minimum-cost flow from a source node to a sink node. This is the successive shortest path method: repeatedly find a minimum-cost augmenting path from source to sink in the residual graph and push as much flow as it allows, until the flow target is met or no path remains. The path search uses SPFA (queue-based Bellman-Ford), since residual edges carry negated costs.

- `MinCostMaxFlow<T, C>(n)` constructs an empty residual network with nodes numbered $[0, n)$.
- `add_edge(u, v, cap, cost, rev_cap = 0)` adds a directed residual-network edge and returns its edge ID.
- `edge_flow(id)` returns the flow through a previously added edge.
- `clear_flow()` resets all edge flows to zero.
- `min_cost_flow(source, sink, target_flow)` sends up to `target_flow` additional units of flow and returns a pair $(\text{flow}, \text{cost})$ for the flow sent by that call. If the returned flow is smaller than `target_flow`, the residual network cannot carry the requested amount.

This implementation uses shortest augmenting paths with SPFA, so it supports negative edge costs as long as the residual graph has no reachable negative-cost cycle. For dense or adversarial inputs, a potential-based Dijkstra version is usually faster when reduced costs are nonnegative. The capacity type `T` should be signed, since reverse residual edges store negative flow.

Overflow warning: the returned cost accumulates products of sent flow and edge cost. Ensure each product and the total cost fit in the cost type `C`.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the number of nodes.
- $O(1)$ per call to `add_edge()`.
- $O(a \cdot n \cdot m)$ per call to `min_cost_flow()`, where a is the number of augmenting paths, n is the number of nodes, and m is the number of edges. For integer capacities, a is at most the flow sent.

Space Complexity: $O(\max(n, m))$ for storage of the residual network and auxiliary arrays.

```
#include <algorithm>
#include <cassert>
#include <cstdlib>
#include <limits>
#include <queue>
#include <utility>
#include <vector>

template<typename T, typename C = int64_t>
class MinCostMaxFlow {
    struct Edge {
        int u, v;
        T cap, flow;
        C cost;
    };

    static constexpr T EPS = static_cast<T>(1e-9);

    int nodes;
    std::vector<Edge> edges;
    std::vector<std::vector<int>>> adj;

public:
    explicit MinCostMaxFlow(int n) : nodes(n), adj(n) {}

    int add_edge(int u, int v, T cap, C cost, T rev_cap = 0) {
        int id = static_cast<int>(edges.size());
        adj[u].push_back(id);
        edges.push_back(Edge{u, v, cap, 0, cost});
        adj[v].push_back(id + 1);
        edges.push_back(Edge{v, u, rev_cap, 0, -cost});
        return id;
    }

    T edge_flow(int id) const { return edges[id].flow; }

    void clear_flow() {
        for (Edge &e : edges) {
            e.flow = 0;
        }
    }

    std::pair<T, C> min_cost_flow(int source, int sink, T target_flow) {
        assert(source != sink);
        static const C INF_COST = std::numeric_limits<C>::max() / 4;
        T flow = 0;
        C cost = 0;
        std::vector<C> dist(nodes);
        std::vector<int> parent_edge(nodes);
        std::vector<char> in_queue(nodes);
        while (flow < target_flow) {
            dist.assign(nodes, INF_COST);
            in_queue.assign(nodes, false);
            std::queue<int> q;
            dist[source] = 0;
```

```

q.push(source);
in_queue[source] = true;
while (!q.empty()) {
    int u = q.front();
    q.pop();
    in_queue[u] = false;
    for (int id : adj[u]) {
        const Edge &e = edges[id];
        if (e.cap - e.flow > EPS && dist[u] + e.cost < dist[e.v]) {
            dist[e.v] = dist[u] + e.cost;
            parent_edge[e.v] = id;
            if (!in_queue[e.v]) {
                q.push(e.v);
                in_queue[e.v] = true;
            }
        }
    }
}
if (dist[sink] == INF_COST) {
    break;
}
T aug = target_flow - flow;
for (int v = sink; v != source; v = edges[parent_edge[v]].u) {
    const Edge &e = edges[parent_edge[v]];
    aug = std::min(aug, e.cap - e.flow);
}
for (int v = sink; v != source; v = edges[parent_edge[v]].u) {
    int id = parent_edge[v];
    Edge &e = edges[id];
    e.flow += aug;
    edges[id ^ 1].flow -= aug;
    cost += aug * e.cost; // Overflow warning!
}
flow += aug;
}
return {flow, cost};
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // Example graph after max flow, with each edge labeled flow/capacity:
    //           2/2
    //           1 -----> 3
    //          / \      |
    // 3/4 /   \ 1/1   | 2/2
    //   /     v     v
    //  0      4 ----> 5
    //   \     ^     3/3
    // 2/3 \   / 2/2
    //      v /
    //      2
    MinCostMaxFlow<int> g(6);
    int id01 = g.add_edge(0, 1, 4, 1);
    g.add_edge(0, 2, 3, 1);
    g.add_edge(1, 3, 2, 1);
    g.add_edge(1, 4, 1, 2);
    g.add_edge(2, 4, 2, 1);
    g.add_edge(3, 5, 2, 1);
    g.add_edge(4, 5, 3, 1);
    auto [flow, cost] = g.min_cost_flow(0, 5, 5);
    assert(flow == 5);
}

```

```

assert(cost == 16);
assert(g.edge_flow(id01) == 3);
g.clear_flow();
auto [flow2, cost2] = g.min_cost_flow(0, 5, 5);
assert(flow2 == 5);
assert(cost2 == 16);
return 0;
}

```

4.5.7 Minimum-Cost Maximum-Flow (Potentials)

Given a flow network with capacities and edge costs, find a minimum-cost flow from a source node to a sink node. This is the successive shortest path method with Johnson potentials. A node potential $h[v]$ transforms each edge cost $c(u, v)$ into a reduced cost $c(u, v) + h[u] - h[v]$; when all reduced costs are nonnegative, Dijkstra's algorithm can find the shortest augmenting path safely. After each augmentation, h is updated with the new shortest-path distances so that reduced costs remain nonnegative on the updated residual graph.

- `MinCostMaxFlow<T, C>(n)` constructs an empty residual network with nodes numbered $[0, n)$.
- `add_edge(u, v, cap, cost, rev_cap = 0)` adds a directed residual-network edge and returns its edge ID.
- `edge_flow(id)` returns the flow through a previously added edge.
- `clear_flow()` resets all edge flows to zero.
- `min_cost_flow(source, sink, target_flow)` sends up to `target_flow` additional units of flow and returns a pair `(flow, cost)` for the flow sent by that call. If the returned flow is smaller than `target_flow`, the residual network cannot carry the requested amount.

If all initial edge costs are nonnegative, the initial potentials can be zero and the first shortest path search is already Dijkstra-safe. If negative costs are present, this implementation first runs SPFA to compute valid initial potentials; it assumes there is no reachable negative-cost cycle. The capacity type `T` should be signed, since reverse residual edges store negative flow.

Overflow warning: the returned cost accumulates products of sent flow and edge cost. Ensure each product and the total cost fit in the cost type `C`.

Time Complexity:

- $O(n)$ per call to the constructor, where n is the number of nodes.
- $O(1)$ per call to `add_edge()`.
- $O(m)$ to check whether initial potentials are needed, plus $O(n \cdot m)$ to compute them when negative residual costs are reachable.
- $O(a \cdot m \cdot \log n)$ for the augmenting paths, where a is the number of augmenting paths, n is the number of nodes, and m is the number of edges. For integer capacities, a is at most the flow sent.

Space Complexity: $O(\max(n, m))$ for storage of the residual network and auxiliary arrays.

```

#include <algorithm>
#include <cassert>
#include <stdint>
#include <functional>
#include <limits>
#include <queue>
#include <utility>
#include <vector>

template<typename T, typename C = int64_t>
class MinCostMaxFlow {
    struct Edge {
        int u, v;
        T cap, flow;
        C cost;
    };
};

```



```

static constexpr T EPS = static_cast<T>(1e-9);
static bool residual(const Edge &e) { return e.cap - e.flow > EPS; }

int nodes;
std::vector<Edge> edges;
std::vector<std::vector<int>>> adj;
std::vector<C> potential;

void init_potentials(int source) {
    static const C INF_COST = std::numeric_limits<C>::max() / 4;
    potential.assign(nodes, 0);
    bool has_negative_cost = false;
    for (const Edge &e : edges) {
        has_negative_cost |= residual(e) && e.cost < 0;
    }
    if (!has_negative_cost) {
        return;
    }
    potential.assign(nodes, INF_COST);
    std::vector<char> in_queue(nodes);
    std::queue<int> q;
    potential[source] = 0;
    q.push(source);
    in_queue[source] = true;
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        in_queue[u] = false;
        for (int id : adj[u]) {
            const Edge &e = edges[id];
            if (residual(e) && potential[u] + e.cost < potential[e.v]) {
                potential[e.v] = potential[u] + e.cost;
                if (!in_queue[e.v]) {
                    q.push(e.v);
                    in_queue[e.v] = true;
                }
            }
        }
    }
    for (C &p : potential) {
        if (p == INF_COST) {
            p = 0;
        }
    }
}

public:
explicit MinCostMaxFlow(int n) : nodes(n), adj(n), potential(n, 0) {}

int add_edge(int u, int v, T cap, C cost, T rev_cap = 0) {
    int id = static_cast<int>(edges.size());
    adj[u].push_back(id);
    edges.push_back(Edge{u, v, cap, 0, cost});
    adj[v].push_back(id + 1);
    edges.push_back(Edge{v, u, rev_cap, 0, -cost});
    return id;
}

T edge_flow(int id) const { return edges[id].flow; }

void clear_flow() {
    for (Edge &e : edges) {
        e.flow = 0;
    }
    potential.assign(nodes, 0);
}

std::pair<T, C> min_cost_flow(int source, int sink, T target_flow) {

```

```

assert(source != sink);
static const C INF_COST = std::numeric_limits<C>::max() / 4;
init_potentials(source);
T flow = 0;
C cost = 0;
std::vector<C> dist(nodes);
std::vector<int> parent_edge(nodes);
using qnode = std::pair<C, int>; // (distance, node)
while (flow < target_flow) {
    std::priority_queue<qnode, std::vector<qnode>, std::greater<>> pq;
    dist.assign(nodes, INF_COST);
    dist[source] = 0;
    pq.push({0, source});
    while (!pq.empty()) {
        auto [du, u] = pq.top();
        pq.pop();
        if (du != dist[u]) {
            continue;
        }
        for (int id : adj[u]) {
            const Edge &e = edges[id];
            C reduced_cost = e.cost + potential[u] - potential[e.v];
            if (residual(e) && dist[u] + reduced_cost < dist[e.v]) {
                dist[e.v] = dist[u] + reduced_cost;
                parent_edge[e.v] = id;
                pq.push({dist[e.v], e.v});
            }
        }
    }
    if (dist[sink] == INF_COST) {
        break;
    }
    for (int i = 0; i < nodes; i++) {
        if (dist[i] < INF_COST) {
            potential[i] += dist[i];
        }
    }
    T aug = target_flow - flow;
    for (int v = sink; v != source; v = edges[parent_edge[v]].u) {
        const Edge &e = edges[parent_edge[v]];
        aug = std::min(aug, e.cap - e.flow);
    }
    for (int v = sink; v != source; v = edges[parent_edge[v]].u) {
        int id = parent_edge[v];
        Edge &e = edges[id];
        e.flow += aug;
        edges[id ^ 1].flow -= aug;
        cost += aug * e.cost; // Overflow warning!
    }
    flow += aug;
}
return {flow, cost};
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // Example graph after max flow, with each edge labeled flow/capacity:
    //      2/2
    //      1 -----> 3
    //      / \      |
    // 3/4 /   \ 1/1 | 2/2

```

```

//      /      v      v
//      0      4 ----> 5
//      \      ^      3/3
// 2/3 \      / 2/2
//      v /
//      2
MinCostMaxFlow<int> g(6);
int id01 = g.add_edge(0, 1, 4, 1);
g.add_edge(0, 2, 3, 1);
g.add_edge(1, 3, 2, 1);
g.add_edge(1, 4, 1, 2);
g.add_edge(2, 4, 2, 1);
g.add_edge(3, 5, 2, 1);
g.add_edge(4, 5, 3, 1);
auto [flow, cost] = g.min_cost_flow(0, 5, 5);
assert(flow == 5);
assert(cost == 16);
assert(g.edge_flow(id01) == 3);
g.clear_flow();
auto [flow2, cost2] = g.min_cost_flow(0, 5, 5);
assert(flow2 == 5);
assert(cost2 == 16);
return 0;
}

```

4.5.8 Gomory-Hu Tree

Given an undirected graph with nonnegative edge weights, build a Gomory-Hu tree: a weighted tree on the same nodes that represents every pairwise minimum s - t cut value. For any two nodes u and v , the minimum cut value between them in the original graph is the minimum edge weight on the path from u to v in the Gomory-Hu tree.

The algorithm starts with every node attached to node 0, then performs $n - 1$ max-flow/min-cut computations. After computing the minimum cut between node s and its current parent, nodes on the s side of that cut are reparented under s , gradually refining the cut-equivalent tree.

- `gomory_hu(n, edges)` returns the $n - 1$ edges of a Gomory-Hu tree as tuples (u, v, weight) for an undirected, weighted graph with n nodes and `edges` of the same shape.
- `min_cut_value(n, tree, source, sink)` returns the minimum cut value between two nodes using a Gomory-Hu tree. Note that for many pairwise cut queries on the same tree, it's more efficient to prebuild the tree adjacency once and answer minimum edge-on-path queries with LCA/binary lifting instead of calling `min_cut_value()` each time.

Time Complexity:

- $O(n)$ calls to maximum flow. With the included Dinic implementation, this is $O(n^3 \cdot m)$ in the worst case, where n is the number of nodes and m is the number of undirected edges.
- $O(n)$ per call to `min_cut_value()`.

Space Complexity: $O(\max(n, m))$ for the residual network and auxiliary arrays.

```

#include <algorithm>
#include <cassert>
#include <functional>
#include <limits>
#include <queue>
#include <tuple>
#include <utility>
#include <vector>

template<typename T>
class Dinic {
    int nodes;

```

```

std::vector<std::vector<std::tuple<int, int, T>>> adj; // (v, rev, cap)
std::vector<int> dist, ptr;

bool bfs(int source, int sink) {
    dist.assign(nodes, -1);
    dist[source] = 0;
    std::queue<int> q;
    q.push(source);
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        for (const auto &[v, rev, cap] : adj[u]) {
            if (dist[v] < 0 && cap > 0) {
                dist[v] = dist[u] + 1;
                q.push(v);
            }
        }
    }
    return dist[sink] >= 0;
}

T dfs(int u, int sink, T pushed) {
    if (u == sink || pushed == 0) {
        return pushed;
    }
    for (; ptr[u] < static_cast<int>(adj[u].size()); ptr[u]++) {
        auto &[v, rev, cap] = adj[u][ptr[u]];
        if (dist[v] == dist[u] + 1 && cap > 0) {
            T pushed_next = dfs(v, sink, std::min(pushed, cap));
            if (pushed_next > 0) {
                cap -= pushed_next;
                std::get<2>(adj[v][rev]) += pushed_next;
                return pushed_next;
            }
        }
    }
    return 0;
}

public:
explicit Dinic(int n) : nodes(n), adj(n), dist(n), ptr(n) {}

void add_edge(int u, int v, T cap) {
    adj[u].emplace_back(v, static_cast<int>(adj[v].size()), cap);
    adj[v].emplace_back(u, static_cast<int>(adj[u].size()) - 1, 0);
}

T max_flow(int source, int sink) {
    T flow = 0;
    while (bfs(source, sink)) {
        ptr.assign(nodes, 0);
        while (true) {
            T pushed = dfs(source, sink, std::numeric_limits<T>::max());
            if (pushed == 0) {
                break;
            }
            flow += pushed;
        }
    }
    return flow;
}

std::vector<char> min_cut(int source) const {
    std::vector<char> reachable(nodes);
    std::queue<int> q;
    reachable[source] = true;
    q.push(source);
    while (!q.empty()) {

```

```

    int u = q.front();
    q.pop();
    for (const auto &[v, rev, cap] : adj[u]) {
        if (!reachable[v] && cap > 0) {
            reachable[v] = true;
            q.push(v);
        }
    }
}
return reachable;
}
};

template<typename T>
using Edge = std::tuple<int, int, T>;

template<typename T>
std::vector<Edge<T>> gomory_hu(int n, const std::vector<Edge<T>> &edges) {
    assert(n >= 1);
    std::vector<int> parent(n, 0);
    std::vector<T> cut_value(n);
    for (int s = 1; s < n; s++) {
        int t = parent[s];
        Dinic<T> mf(n);
        for (const auto &[u, v, w] : edges) {
            mf.add_edge(u, v, w);
            mf.add_edge(v, u, w);
        }
        T flow = mf.max_flow(s, t);
        std::vector<char> cut = mf.min_cut(s);
        for (int v = s + 1; v < n; v++) {
            if (parent[v] == t && cut[v] == cut[s]) {
                parent[v] = s;
            }
        }
        if (cut[parent[t]]) {
            parent[s] = parent[t];
            parent[t] = s;
            cut_value[s] = cut_value[t];
            cut_value[t] = flow;
        } else {
            cut_value[s] = flow;
        }
    }
    std::vector<Edge<T>> tree;
    for (int v = 1; v < n; v++) {
        tree.emplace_back(v, parent[v], cut_value[v]);
    }
    return tree;
}

template<typename T>
T min_cut_value(int n, const std::vector<Edge<T>> &tree, int source, int sink) {
    assert(0 <= source && source < n && 0 <= sink && sink < n && source != sink);
    // For many queries, prebuild this adjacency and add LCA/binary lifting for min edge on path.
    std::vector<std::vector<std::pair<int, T>>> adj(n);
    for (const auto &[u, v, w] : tree) {
        adj[u].emplace_back(v, w);
        adj[v].emplace_back(u, w);
    }
    std::vector<char> seen(n);
    std::function<T(int, T)> dfs = [&](int u, T best) {
        if (u == sink) {
            return best;
        }
        seen[u] = true;
        for (auto [v, w] : adj[u]) {
            if (!seen[v]) {

```

```

    T res = dfs(v, std::min(best, w));
    if (res != -1) {
        return res;
    }
}
}
return static_cast<T>(-1);
};
return dfs(source, std::numeric_limits<T>::max());
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //      w=3      w=2
    //  0 -----> 1 -----> 3
    //      \      w=1|      ^
    // w=2 \      v      / w=4
    //      +----> 2 ----+
    vector<Edge<int>> edges{
        {0, 1, 3}, {0, 2, 2}, {1, 2, 1}, {1, 3, 2}, {2, 3, 4},
    };
    auto tree = gomory_hu(4, edges);
    assert((tree == vector<Edge<int>>{{1, 0, 5}, {2, 1, 5}, {3, 2, 6}}));
    assert(min_cut_value(4, tree, 0, 1) == 5);
    assert(min_cut_value(4, tree, 0, 2) == 5);
    assert(min_cut_value(4, tree, 0, 3) == 5);
    assert(min_cut_value(4, tree, 2, 3) == 6);
    return 0;
}

```

4.5.9 Global Minimum Cut (Stoer-Wagner)

Given an undirected graph with nonnegative edge weights, the global minimum cut is the partition of the nodes into two nonempty sets that minimizes the total weight of edges crossing between them. Unlike a minimum s - t cut, no source or sink is fixed; the goal is the cheapest way to split the graph in two. The Stoer-Wagner algorithm finds it without any maximum flow computation by repeatedly running a “minimum cut phase”.

A phase grows a set, starting from an arbitrary node, by repeatedly adding the most tightly connected outside node (the one with the greatest total weight to the set so far), exactly as in Prim’s algorithm. The last node added, together with its connection weight, yields one candidate cut, the “cut of the phase”. The last two nodes added are then merged into a single supernode and the next phase begins on the smaller graph. After $n - 1$ phases every candidate has been considered, and the smallest is the global minimum cut.

- `global_min_cut(cap)` returns a pair (`weight`, `side`) for a graph as a symmetric capacity matrix `cap` (with 0 for the diagonal and absent edges), where `weight` is the total weight of the global minimum cut and `side` lists the nodes on one side of that cut. The graph must have at least two nodes.

Parallel edges should be pre-summed into the capacity matrix before calling: `cap[u][v]` holds a single combined weight, so the sum of all parallel edge weights between `u` and `v` should be stored there.

Time Complexity: $O(n^3)$ per call, where n is the number of nodes.

Space Complexity: $O(n^2)$ auxiliary.

```

#include <cassert>
#include <cstdint>
#include <utility>

```

```

#include <vector>

std::pair<int64_t, std::vector<int>> global_min_cut(std::vector<std::vector<int64_t>> cap) {
    int n = static_cast<int>(cap.size());
    assert(n >= 2);
    for (int i = 0; i < n; i++) {
        for (int j = i + 1; j < n; j++) {
            assert(cap[i][j] >= 0);
            assert(cap[i][j] == cap[j][i]);
        }
    }
    std::pair<int64_t, std::vector<int>> best = {INT64_MAX, {}};
    std::vector<std::vector<int>> groups(n);
    std::vector<int64_t> weight(n);
    std::vector<char> exists(n, true), added(n);
    for (int i = 0; i < n; i++) {
        groups[i] = {i};
    }
    for (int phase = 1; phase < n; phase++) {
        weight.assign(n, 0);
        added.assign(n, false);
        int s = -1;
        for (int it = 0; it <= n - phase; it++) {
            int sel = -1;
            for (int i = 0; i < n; i++) {
                if (exists[i] && !added[i] && (sel == -1 || weight[i] > weight[sel])) {
                    sel = i;
                }
            }
            if (it == n - phase) {
                if (weight[sel] < best.first) {
                    best = {weight[sel], groups[sel]};
                }
                groups[s].insert(groups[s].end(), groups[sel].begin(), groups[sel].end());
                for (int i = 0; i < n; i++) {
                    cap[s][i] += cap[sel][i]; // Overflow warning!
                    cap[i][s] = cap[s][i];
                }
                exists[sel] = false;
                break;
            }
            added[sel] = true;
            for (int i = 0; i < n; i++) {
                weight[i] += cap[sel][i];
            }
            s = sel;
        }
    }
    return best;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //      w=2
    //  0 ----- 1
    //   \       /
    // w=1 \   / w=3
    //      2
    vector<vector<int64_t>> cap{
        {0, 2, 1},
        {2, 0, 3},
        {1, 3, 0},
    }
}

```

```

};
// Cuts: {0}|{1,2} = 3, {1}|{0,2} = 5, {2}|{0,1} = 4. Minimum is 3.
auto [weight, side] = global_min_cut(cap);
assert(weight == 3);
assert((side == vector<int>{1, 2})); // The shore {1, 2}, opposite the single node 0.

vector<vector<int64_t>> disconnected{{0, 0}, {0, 0}};
assert(global_min_cut(disconnected).first == 0);
return 0;
}

```

4.6 Matching and Assignment

4.6.1 Maximum Bipartite Matching (Kuhn)

Given two sets of nodes $A = \{0, 1, \dots, n_1 - 1\}$ and $B = \{0, 1, \dots, n_2 - 1\}$, as well as a set of edges E mapping nodes from set A to set B , find the largest possible subset of E containing no edges that share the same node.

Kuhn's algorithm builds the matching one left node at a time. For each, a depth-first search looks for an augmenting path: an alternating sequence of unmatched and matched edges that reaches a free right node. Flipping the edges along such a path enlarges the matching by one.

- `bipartite_matching(n2)` populates `match_left` and `match_right`, then returns maximum matching size for a global, pre-populated adjacency list `adj`, whose indices represent left-side nodes and whose entries contain right-side neighbors numbered $[0, n_2)$.

Time Complexity: $O(n_1 \cdot m)$ per call, where m is the number of edges.

Space Complexity:

- $O(\max(n_1 + n_2, m))$ for storage of the graph, where m is the number of edges.
- $O(n_1 + n_2)$ auxiliary heap space and $O(n_1)$ auxiliary stack space.

```

#include <vector>

std::vector<int> match_left, match_right;
std::vector<int> visit;
std::vector<std::vector<int>> adj;
int timer;

bool dfs(int u) {
    visit[u] = timer;
    for (int nb : adj[u]) {
        if (match_right[nb] == -1) {
            match_left[u] = nb;
            match_right[nb] = u;
            return true;
        }
    }
    for (int nb : adj[u]) {
        int v = match_right[nb];
        if (visit[v] != timer && dfs(v)) {
            match_left[u] = nb;
            match_right[nb] = u;
            return true;
        }
    }
    return false;
}

int bipartite_matching(int n2) {
    int n1 = static_cast<int>(adj.size());

```



```

match_left.assign(n1, -1);
match_right.assign(n2, -1);
visit.assign(n1, 0);
timer = 0;
int matches = 0;
for (int i = 0; i < n1; i++) {
    timer++;
    if (dfs(i)) {
        matches++;
    }
}
return matches;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // Left nodes a0..a2, right nodes b0..b3:
    //   a0 -- b1
    //   a1 -- b0, b1, b2
    //   a2 -- b2, b3
    int n1 = 3, n2 = 4;
    adj.assign(n1, {});
    adj[0].push_back(1);
    adj[1].push_back(0);
    adj[1].push_back(1);
    adj[1].push_back(2);
    adj[2].push_back(2);
    adj[2].push_back(3);
    int pairs = bipartite_matching(n2);
    assert(pairs == 3);
    assert((match_right == vector<int>{1, 0, 2, -1}));
    return 0;
}

```

4.6.2 Maximum Bipartite Matching (Hopcroft-Karp)

Given two sets of nodes $A = \{0, 1, \dots, n_1 - 1\}$ and $B = \{0, 1, \dots, n_2 - 1\}$, as well as a set of edges E mapping nodes from set A to set B , find the largest possible subset of E containing no edges that share the same node.

Hopcroft-Karp augments along many shortest augmenting paths per phase. Each phase first runs a breadth-first search to layer the graph by distance, then a depth-first search to find a maximal set of node-disjoint shortest augmenting paths to flip at once, for $O(m \cdot \sqrt{n_1 + n_2})$ running time overall.

- `bipartite_matching_hk(n2)` populates `match_left` and `match_right`, then returns maximum matching size for a global, pre-populated adjacency list `adj`, whose indices represent left-side nodes and whose entries contain right-side neighbors numbered $[0, n_2)$.

Time Complexity: $O(m \cdot \sqrt{n_1 + n_2})$ per call, where m is the number of edges.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n_1 + n_2)$ auxiliary stack space and $O(n_1 + n_2)$ auxiliary heap space.

```

#include <queue>
#include <vector>

std::vector<std::vector<int>> adj;

```

```

std::vector<char> used, visit;
std::vector<int> match_left, match_right, dist;

void bfs() {
    int n1 = static_cast<int>(adj.size());
    dist.assign(n1, -1);
    std::queue<int> q;
    for (int u = 0; u < n1; u++) {
        if (!used[u]) {
            q.push(u);
            dist[u] = 0;
        }
    }
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        for (int nb : adj[u]) {
            int v = match_right[nb];
            if (v >= 0 && dist[v] < 0) {
                dist[v] = dist[u] + 1;
                q.push(v);
            }
        }
    }
}

bool dfs(int u) {
    visit[u] = true;
    for (int nb : adj[u]) {
        int v = match_right[nb];
        if (v < 0 || (!visit[v] && dist[v] == dist[u] + 1 && dfs(v))) {
            match_left[u] = nb;
            match_right[nb] = u;
            used[u] = true;
            return true;
        }
    }
    return false;
}

int bipartite_matching_hk(int n2) {
    int n1 = static_cast<int>(adj.size());
    match_left.assign(n1, -1);
    match_right.assign(n2, -1);
    used.assign(n1, false);
    int res = 0;
    while (true) {
        bfs();
        visit.assign(n1, false);
        int f = 0;
        for (int u = 0; u < n1; u++) {
            if (!used[u] && dfs(u)) {
                f++;
            }
        }
        if (f == 0) {
            return res;
        }
        res += f;
    }
    return res;
}

```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    // Left nodes a0..a2, right nodes b0..b3:
    //   a0 -- b1
    //   a1 -- b0, b1, b2
    //   a2 -- b2, b3
    int n1 = 3, n2 = 4;
    adj.assign(n1, {});
    adj[0].push_back(1);
    adj[1].push_back(0);
    adj[1].push_back(1);
    adj[1].push_back(2);
    adj[2].push_back(2);
    adj[2].push_back(3);
    int pairs = bipartite_matching_hk(n2);
    assert(pairs == 3);
    assert((match_right == vector<int>{1, 0, 2, -1}));
    return 0;
}
```

4.6.3 Minimum Vertex Cover (Konig)

Given a bipartite graph with left nodes $A = \{0, \dots, n_1 - 1\}$ and right nodes $B = \{0, \dots, n_2 - 1\}$, find a minimum vertex cover: a smallest set of nodes such that every edge has at least one endpoint in the set. By Konig's theorem, in a bipartite graph the size of a minimum vertex cover equals the size of a maximum matching, and the nodes left uncovered form a maximum independent set.

The construction first computes a maximum matching with Kuhn's algorithm. It then marks every left node reachable from an unmatched left node by an alternating path (following unmatched edges from left to right and matched edges from right to left). Konig's theorem says the minimum vertex cover is exactly the left nodes that are not reachable together with the right nodes that are reachable; each matched edge contributes exactly one endpoint, so the cover size equals the matching size.

- `min_vertex_cover(n2)` returns the nodes of a minimum vertex cover for the global adjacency list `adj` (left nodes $[0, n_1)$, right neighbors $[0, n_2)$), where `n1` is `adj.size()`. Each returned value is a left node `i` in $[0, n_1)$, or a right node `j` encoded as `n1 + j`.

The adjacency list `adj` is shared with the matching routine: `adj[u]` lists the right-side neighbors of left node `u`. Cover nodes are reported with right node `j` encoded as `n1 + j` so the two sides share one index space.

Time Complexity: $O(n_1 \cdot m)$ per call, dominated by the matching, where m is the number of edges.

Space Complexity:

- $O(\max(n_1 + n_2, m))$ for storage of the graph.
- $O(n_1 + n_2)$ auxiliary.

```
#include <vector>

std::vector<std::vector<int>> adj;
std::vector<int> match_left, match_right;
std::vector<int> visit;
int timer;

bool dfs(int u) {
    visit[u] = timer;
    for (int nb : adj[u]) {
        if (match_right[nb] == -1) {
```

```

        match_left[u] = nb;
        match_right[nb] = u;
        return true;
    }
}
for (int nb : adj[u]) {
    int v = match_right[nb];
    if (visit[v] != timer && dfs(v)) {
        match_left[u] = nb;
        match_right[nb] = u;
        return true;
    }
}
return false;
}

int bipartite_matching(int n2) {
    int n1 = static_cast<int>(adj.size());
    match_left.assign(n1, -1);
    match_right.assign(n2, -1);
    visit.assign(n1, 0);
    timer = 0;
    int matches = 0;
    for (int i = 0; i < n1; i++) {
        timer++;
        if (dfs(i)) {
            matches++;
        }
    }
    return matches;
}

std::vector<int> min_vertex_cover(int n2) {
    int n1 = static_cast<int>(adj.size());
    bipartite_matching(n2);
    std::vector<char> reachable_left(n1), reachable_right(n2);
    std::vector<int> st;
    for (int u = 0; u < n1; u++) {
        if (match_left[u] == -1) {
            reachable_left[u] = true;
            st.push_back(u);
        }
    }
    while (!st.empty()) {
        int u = st.back();
        st.pop_back();
        for (int v : adj[u]) {
            if (v != match_left[u] && !reachable_right[v]) {
                reachable_right[v] = true;
                int next = match_right[v];
                if (next != -1 && !reachable_left[next]) {
                    reachable_left[next] = true;
                    st.push_back(next);
                }
            }
        }
    }
    std::vector<int> cover;
    for (int u = 0; u < n1; u++) {
        if (!reachable_left[u]) {
            cover.push_back(u);
        }
    }
    for (int v = 0; v < n2; v++) {
        if (reachable_right[v]) {
            cover.push_back(n1 + v);
        }
    }
}

```

```
    return cover;
}
```

Example Usage

```
#include <algorithm>
#include <cassert>
using namespace std;

int main() {
    // Left nodes a0..a2, right nodes b0..b3:
    //  a0 -- b1
    //  a1 -- b0, b1, b2
    //  a2 -- b2, b3
    int n1 = 3, n2 = 4;
    adj.assign(n1, {});
    adj[0].push_back(1);
    adj[1].push_back(0);
    adj[1].push_back(1);
    adj[1].push_back(2);
    adj[2].push_back(2);
    adj[2].push_back(3);
    vector<int> cover = min_vertex_cover(n2);
    assert(cover.size() == 3); // Equals the maximum matching size, by Konig's theorem.
    // The cover is valid: every edge has at least one endpoint in it.
    for (int u = 0; u < n1; u++) {
        for (int v : adj[u]) {
            bool left_in = find(cover.begin(), cover.end(), u) != cover.end();
            bool right_in = find(cover.begin(), cover.end(), n1 + v) != cover.end();
            assert(left_in || right_in);
        }
    }
    return 0;
}
```

4.6.4 Maximum Graph Matching (Blossom)

Given an undirected, unweighted simple graph, determine a maximum matching: a maximum subset of its edges such that no node is shared between different edges in the resulting subset.

Edmonds' blossom algorithm extends augmenting-path search to general (non-bipartite) graphs, where odd-length cycles called blossoms can conceal augmenting paths. Each blossom is contracted into a single node so the search can proceed, and is expanded afterward to recover the matching. For maximum-weight matching in bipartite graphs, use Hungarian or min-cost max-flow; for maximum-weight matching in general graphs, use the weighted blossom algorithm in the next section.

The graph must be simple: self-loops and parallel edges are not supported.

- `maximum_matching(adj)` returns a matching for a bidirectionally pre-populated adjacency list `adj`, whose indices represent the nodes. The returned vector `match` has `match[u] == v` and `match[v] == u` when `u` and `v` are matched, or `match[u]` is `-1` when `u` is unmatched.

Time Complexity: $O(n^3)$ per call, where n is the number of nodes.

Space Complexity:

- $O(\max(n, m))$ for storage of the graph, where n is the number of nodes and m is the number of edges.
- $O(n)$ auxiliary.

```
#include <algorithm>
#include <numeric>
#include <queue>
```

```

#include <utility>
#include <vector>

std::vector<int> maximum_matching(const std::vector<std::vector<int>>> &adj) {
    int n = static_cast<int>(adj.size());
    std::vector<int> match(n, -1), label(n), parent(n), base(n), aux(n, -1);
    std::queue<int> q;
    int aux_time = -1;
    auto lca = [&](int u, int v) {
        aux_time++;
        while (true) {
            if (u != -1) {
                u = base[u];
                if (aux[u] == aux_time) {
                    return u;
                }
            }
            aux[u] = aux_time;
            u = (match[u] == -1) ? -1 : parent[match[u]];
        }
        std::swap(u, v);
    };
};
auto mark_blossom = [&](int u, int v, int b) {
    while (base[u] != b) {
        parent[u] = v;
        v = match[u];
        if (label[v] == 1) {
            label[v] = 0;
            q.push(v);
        }
        base[u] = base[v] = b;
        u = parent[v];
    }
};
auto augment = [&](int u) {
    while (u != -1) {
        int p = parent[u];
        int next = match[p];
        match[u] = p;
        match[p] = u;
        u = next;
    }
};
auto find_path = [&](int root) {
    label.assign(n, -1);
    parent.assign(n, -1);
    std::iota(base.begin(), base.end(), 0);
    while (!q.empty()) {
        q.pop();
    }
    label[root] = 0;
    q.push(root);
    while (!q.empty()) {
        int u = q.front();
        q.pop();
        for (int v : adj[u]) {
            if (base[u] == base[v] || match[u] == v) {
                continue;
            }
            if (label[v] == -1) {
                label[v] = 1;
                parent[v] = u;
                if (match[v] == -1) {
                    augment(v);
                    return true;
                }
            }
            label[match[v]] = 0;
            q.push(match[v]);
        }
    }
};

```

```

        } else if (label[v] == 0) {
            int b = lca(u, v);
            mark_blossom(u, v, b);
            mark_blossom(v, u, b);
        }
    }
}
return false;
};
for (int i = 0; i < n; i++) {
    if (match[i] == -1) {
        for (int v : adj[i]) {
            if (i != v && match[v] == -1) {
                match[i] = v;
                match[v] = i;
                break;
            }
        }
    }
}
for (int i = 0; i < n; i++) {
    if (match[i] == -1) {
        find_path(i);
    }
}
return match;
}

```

Example Usage

```

#include <cassert>
using namespace std;

vector<vector<int>>> adj;

void add_edge(int u, int v) {
    adj[u].push_back(v);
    adj[v].push_back(u);
}

int matching_size(const vector<int> &match) {
    return std::count_if(match.begin(), match.end(), [](int v) { return v != -1; }) / 2;
}

int main() {
    {
        // 0---1---2
        //  \   |
        //   +-----3
        int nodes = 4;
        adj.assign(nodes, {});
        add_edge(0, 1);
        add_edge(1, 2);
        add_edge(2, 3);
        add_edge(3, 0);
        vector<int> match = maximum_matching(adj);
        assert(matching_size(match) == 2);
        assert((match == vector<int>{1, 0, 3, 2}));
    }
    {
        // 3---0---1---4
        //   |   |
        //   +---2
        int nodes = 5;
        adj.assign(nodes, {});
        add_edge(0, 1);
    }
}

```

```

    add_edge(1, 2);
    add_edge(2, 0);
    add_edge(0, 3);
    add_edge(1, 4);
    assert(matching_size(maximum_matching(adj)) == 2);
}
return 0;
}

```

4.6.5 Maximum Weight Graph Matching (Weighted Blossom)

Given an undirected graph with positive edge weights, determine a maximum-weight matching: a subset of edges, no two sharing a node, whose total weight is as large as possible. Unlike the bipartite case (which Hungarian or min-cost max-flow solves), general graphs may contain odd cycles, so this is the weighted generalization of Edmonds' blossom algorithm.

The algorithm runs a primal-dual method. Each node carries a dual label, and an edge is “tight” when its label sum equals twice its weight; the search uses only tight edges, exactly as unweighted blossom matching does, contracting odd cycles into blossoms. When no tight augmenting structure remains, the labels are adjusted by the largest step that keeps every dual constraint satisfied, which makes a new edge tight or lets a blossom expand, and the search resumes. Because the matching is over the implicit complete graph with absent edges weighted 0, a node may be left effectively unmatched whenever staying unmatched is at least as good as any incident edge.

All edge weights must be positive; an absent edge is treated as weight 0. Self-loops are not supported. Parallel edges are supported by retaining only the maximum weight added between each pair of nodes. The node count should be kept modest, as the algorithm uses dense storage.

Overflow warning: labels and the returned matching weight are stored in `int64_t`; all doubled weights, label sums, and the total matching weight must fit.

- `WeightedGeneralMatching(n)` creates a graph of `n` nodes numbered $[0, n)$.
- `add_edge(u, v, w)` adds an undirected edge of positive weight `w` between `u` and `v`. If several edges join the same pair, only the maximum weight is retained.
- `solve()` computes a maximum-weight matching and returns its total weight.
- `partner(u)` returns the node matched with `u`, or `-1` if `u` is unmatched. Valid after `solve()`.

Time Complexity: $O(n^3)$ per call to `solve()`, where n is the number of nodes.

Space Complexity: $O(n^2)$ for the dense edge and blossom storage.

```

#include <algorithm>
#include <cstdint>
#include <queue>
#include <utility>
#include <vector>

class WeightedGeneralMatching {
    struct Edge {
        int u, v;
        int64_t w;
    };

    int n, nx; // n original nodes (1-based internally); nx includes blossom nodes.
    std::vector<std::vector<Edge>> g;
    std::vector<std::vector<int>> flo, flo_from;
    std::vector<int64_t> lab;
    std::vector<int> match_, slack, st, par, s_, aux_;
    std::queue<int> q;
    int aux_clock;

    int64_t e_delta(const Edge &e) const {
        return lab[e.u] + lab[e.v] - 2 * e.w; // Overflow warning!
    }
};

```



```

}

void update_slack(int u, int x) {
    if (slack[x] == 0 || e_delta(g[u][x]) < e_delta(g[slack[x]][x])) {
        slack[x] = u;
    }
}

void set_slack(int x) {
    slack[x] = 0;
    for (int u = 1; u <= n; u++) {
        if (g[u][x].w > 0 && st[u] != x && s_[st[u]] == 0) {
            update_slack(u, x);
        }
    }
}

void queue_push(int x) {
    if (x <= n) {
        q.push(x);
    } else {
        for (int t : flo[x]) {
            queue_push(t);
        }
    }
}

void set_st(int x, int b) {
    st[x] = b;
    if (x > n) {
        for (int t : flo[x]) {
            set_st(t, b);
        }
    }
}

int get_pr(int b, int xr) {
    int pr = static_cast<int>(std::find(flo[b].begin(), flo[b].end(), xr) - flo[b].begin());
    if (pr & 1) {
        std::reverse(flo[b].begin() + 1, flo[b].end());
        return static_cast<int>(flo[b].size()) - pr;
    }
    return pr;
}

void set_match(int u, int v) {
    match_[u] = g[u][v].v;
    if (u <= n) {
        return;
    }
    Edge e = g[u][v];
    int xr = flo_from[u][e.u], pr = get_pr(u, xr);
    for (int i = 0; i < pr; i++) {
        set_match(flo[u][i], flo[u][i ^ 1]);
    }
    set_match(xr, v);
    std::rotate(flo[u].begin(), flo[u].begin() + pr, flo[u].end());
}

void augment(int u, int v) {
    while (true) {
        int xnv = st[match_[u]];
        set_match(u, v);
        if (xnv == 0) {
            return;
        }
        set_match(xnv, st[par[xnv]]);
        u = st[par[xnv]];
    }
}

```

```

    v = xnv;
}
}

int lca(int u, int v) {
    aux_clock++;
    while (u != 0 || v != 0) {
        if (u == 0) {
            std::swap(u, v);
            continue;
        }
        if (aux_[u] == aux_clock) {
            return u;
        }
        aux_[u] = aux_clock;
        u = st[match_[u]];
        if (u != 0) {
            u = st[par[u]];
        }
        std::swap(u, v);
    }
    return 0;
}

void add_blossom(int u, int anc, int v) {
    int b = n + 1;
    while (b <= nx && st[b] != 0) {
        b++;
    }
    if (b > nx) {
        nx++;
    }
    lab[b] = 0;
    s_[b] = 0;
    match_[b] = match_[anc];
    flo[b] = {anc};
    auto build = [&](int x) {
        for (int y; x != anc; x = st[par[y]]) {
            flo[b].push_back(x);
            y = st[match_[x]];
            flo[b].push_back(y);
            queue_push(y);
        }
    };
    build(u);
    std::reverse(flo[b].begin() + 1, flo[b].end());
    build(v);
    set_st(b, b);
    for (int x = 1; x <= nx; x++) {
        g[b][x].w = g[x][b].w = 0;
    }
    for (int x = 1; x <= n; x++) {
        flo_from[b][x] = 0;
    }
    for (int xs : flo[b]) {
        for (int x = 1; x <= nx; x++) {
            if (g[b][x].w == 0 || e_delta(g[xs][x]) < e_delta(g[b][x])) {
                g[b][x] = g[xs][x];
                g[x][b] = g[x][xs];
            }
        }
    }
    for (int x = 1; x <= n; x++) {
        if (flo_from[xs][x] != 0) {
            flo_from[b][x] = xs;
        }
    }
    set_slack(b);
}

```

```

}

void expand_blossom(int b) {
    for (int t : flo[b]) {
        set_st(t, t);
    }
    int xr = flo_from[b][g[b][par[b]].u], pr = get_pr(b, xr);
    for (int i = 0; i < pr; i += 2) {
        int xs = flo[b][i], xns = flo[b][i + 1];
        par[xs] = g[xns][xs].u;
        s_[xs] = 1;
        s_[xns] = 0;
        slack[xs] = 0;
        set_slack(xns);
        queue_push(xns);
    }
    s_[xr] = 1;
    par[xr] = par[b];
    for (int i = pr + 1; i < static_cast<int>(flo[b].size()); i++) {
        int xs = flo[b][i];
        s_[xs] = -1;
        set_slack(xs);
    }
    st[b] = 0;
}

bool on_found_edge(const Edge &e) {
    int u = st[e.u], v = st[e.v];
    if (s_[v] == -1) {
        par[v] = e.u;
        s_[v] = 1;
        int nu = st[match_[v]];
        slack[v] = slack[nu] = 0;
        s_[nu] = 0;
        queue_push(nu);
    } else if (s_[v] == 0) {
        int anc = lca(u, v);
        if (anc == 0) {
            augment(u, v);
            augment(v, u);
            return true;
        }
        add_blossom(u, anc, v);
    }
    return false;
}

bool matching() {
    std::fill(s_.begin(), s_.begin() + nx + 1, -1);
    std::fill(slack.begin(), slack.begin() + nx + 1, 0);
    q = std::queue<int>();
    for (int x = 1; x <= nx; x++) {
        if (st[x] == x && match_[x] == 0) {
            par[x] = s_[x] = 0;
            queue_push(x);
        }
    }
    if (q.empty()) {
        return false;
    }
    while (true) {
        while (!q.empty()) {
            int u = q.front();
            q.pop();
            if (s_[st[u]] == 1) {
                continue;
            }
            for (int v = 1; v <= n; v++) {

```

```

        if (g[u][v].w > 0 && st[u] != st[v]) {
            if (e_delta(g[u][v]) == 0) {
                if (on_found_edge(g[u][v])) {
                    return true;
                }
            } else {
                update_slack(u, st[v]);
            }
        }
    }
}
}
int64_t d = INT64_MAX;
for (int b = n + 1; b <= nx; b++) {
    if (st[b] == b && s_[b] == 1) {
        d = std::min(d, lab[b] / 2);
    }
}
for (int x = 1; x <= nx; x++) {
    if (st[x] == x && slack[x] != 0) {
        if (s_[x] == -1) {
            d = std::min(d, e_delta(g[slack[x]][x]));
        } else if (s_[x] == 0) {
            d = std::min(d, e_delta(g[slack[x]][x]) / 2);
        }
    }
}
for (int u = 1; u <= n; u++) {
    if (s_[st[u]] == 0) {
        if (lab[u] <= d) {
            return false;
        }
        lab[u] -= d;
    } else if (s_[st[u]] == 1) {
        lab[u] += d;
    }
}
for (int b = n + 1; b <= nx; b++) {
    if (st[b] == b && s_[b] != -1) {
        lab[b] += (s_[b] == 0 ? 1 : -1) * d * 2;
    }
}
q = std::queue<int>();
for (int x = 1; x <= nx; x++) {
    if (st[x] == x && slack[x] != 0 && st[slack[x]] != x && e_delta(g[slack[x]][x]) == 0) {
        if (on_found_edge(g[slack[x]][x])) {
            return true;
        }
    }
}
for (int b = n + 1; b <= nx; b++) {
    if (st[b] == b && s_[b] == 1 && lab[b] == 0) {
        expand_blossom(b);
    }
}
}
}

public:
explicit WeightedGeneralMatching(int n)
: n(n),
  nx(n),
  g(2 * n + 1, std::vector<Edge>(2 * n + 1)),
  flo(2 * n + 1),
  flo_from(2 * n + 1, std::vector<int>(n + 1, 0)),
  lab(2 * n + 1, 0),
  match_(2 * n + 1, 0),
  slack(2 * n + 1, 0),
  st(2 * n + 1, 0),

```

```

        par(2 * n + 1, 0),
        s_(2 * n + 1, 0),
        aux_(2 * n + 1, 0),
        aux_clock(0) {
    for (int u = 1; u <= n; u++) {
        for (int v = 1; v <= n; v++) {
            g[u][v] = {u, v, 0};
        }
    }
}

void add_edge(int u, int v, int64_t w) {
    w = std::max(w, g[u + 1][v + 1].w);
    g[u + 1][v + 1].w = g[v + 1][u + 1].w = w;
}

int64_t solve() {
    nx = n;
    for (int u = 1; u <= n; u++) {
        st[u] = u;
        match_[u] = 0;
        flo[u].clear();
    }
    int64_t w_max = 0;
    for (int u = 1; u <= n; u++) {
        for (int v = 1; v <= n; v++) {
            flo_from[u][v] = (u == v ? u : 0);
            w_max = std::max(w_max, g[u][v].w);
        }
    }
    for (int u = 1; u <= n; u++) {
        lab[u] = w_max;
    }
    while (matching()) {
    }
    int64_t weight = 0;
    for (int u = 1; u <= n; u++) {
        if (match_[u] != 0 && match_[u] < u) {
            weight += g[u][match_[u]].w; // Overflow warning!
        }
    }
    return weight;
}

int partner(int u) const {
    int v = match_[u + 1];
    // A node "matched" through an absent (weight 0) edge is really unmatched.
    if (v == 0 || g[u + 1][v].w == 0) {
        return -1;
    }
    return v - 1;
}
};

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    //   w=5   w=6   w=5   w=6
    // 0-----1-----2-----3-----4
    //  \-----/
    //           w=5
    WeightedGeneralMatching g(5);
    g.add_edge(0, 1, 5);
}

```

```

g.add_edge(1, 2, 6);
g.add_edge(2, 3, 5);
g.add_edge(3, 4, 6);
g.add_edge(4, 0, 5);
int64_t weight = g.solve();
assert(weight == 12); // Edges 1-2 and 3-4 with weights 6 + 6.
assert(g.partner(1) == 2 && g.partner(2) == 1);
assert(g.partner(3) == 4 && g.partner(4) == 3);
assert(g.partner(0) == -1); // Node 0 is left unmatched.

// w=3 w=3 w=5
// 0-----1-----2-----3
// \-----/
//      w=3
WeightedGeneralMatching t(4);
t.add_edge(0, 1, 3);
t.add_edge(1, 2, 3);
t.add_edge(0, 2, 3);
t.add_edge(2, 3, 5);
t.add_edge(2, 3, 1); // A lighter parallel edge does not replace the heavier one.
assert(t.solve() == 8); // Edges 0-1 and 2-3 with weights 3 + 5.
return 0;
}

```

4.6.6 Assignment (Hungarian Algorithm)

Solves the minimum-cost or maximum-value assignment problem for a rectangular matrix. Given n workers and m jobs (with $n \leq m$), choose a distinct job for each worker so that the total cost is minimized or the total value is maximized.

The classical Hungarian algorithm using potentials transforms a cost matrix using the principle that adding or subtracting constants from rows or columns does not change the optimal assignment. Workers are assigned one row at a time: each new row triggers a search for a least-cost augmenting path over zero-slack edges, with the potentials adjusted along the way to expose new zero-slack edges until the path reaches a free job. Note that while the input matrix is 0-based here, the internal calculations are 1-based.

Maximum-value assignment reduces to minimum-cost assignment by replacing each value x with $M - x$, where M is the largest matrix entry. Every assignment contains exactly n entries, so minimizing the transformed total $nM - \sum x$ maximizes the original total.

- `min_assignment(cost)` returns a pair (`min_cost`, `assignment`) for a 0-based `cost` matrix, where `assignment` contains n values and `assignment[i]` is the chosen job for worker `i`, using 0-based job numbers.
- `max_assignment(value)` returns the analogous pair (`max_value`, `assignment`) maximizing a 0-based `value` matrix.

Time Complexity: $O(n^2 \cdot m)$ per call, where the input matrix has n rows and m columns.

Space Complexity:

- $O(n + m)$ auxiliary, not counting the input matrix and returned assignment.
- $O(n \cdot m)$ auxiliary for `max_assignment()` due to the transformed matrix.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <limits>
#include <utility>
#include <vector>

template<typename T>
std::pair<T, std::vector<int>> min_assignment(const std::vector<std::vector<T>> &cost) {
    static const T INF = std::numeric_limits<T>::max() / 4;

```

```

int n = static_cast<int>(cost.size());
int m = cost.empty() ? 0 : static_cast<int>(cost[0].size());
assert(n <= m);
std::vector<T> u(n + 1), v(m + 1), minv(m + 1);
std::vector<int> p(m + 1), way(m + 1);
std::vector<char> used(m + 1);
for (int i = 1; i <= n; i++) {
    minv.assign(m + 1, INF);
    used.assign(m + 1, false);
    p[0] = i;
    int j0 = 0;
    do {
        used[j0] = true;
        int i0 = p[j0], j1 = 0;
        T delta = INF;
        for (int j = 1; j <= m; j++) {
            if (used[j]) {
                continue;
            }
            T cur = cost[i0 - 1][j - 1] - u[i0] - v[j]; // Overflow warning!
            if (cur < minv[j]) {
                minv[j] = cur;
                way[j] = j0;
            }
            if (minv[j] < delta) {
                delta = minv[j];
                j1 = j;
            }
        }
        for (int j = 0; j <= m; j++) {
            if (used[j]) {
                u[p[j]] += delta;
                v[j] -= delta;
            } else {
                minv[j] -= delta;
            }
        }
        j0 = j1;
    } while (p[j0] != 0);
    do {
        int j1 = way[j0];
        p[j0] = p[j1];
        j0 = j1;
    } while (j0 != 0);
}
std::vector<int> assignment(n, -1);
for (int j = 1; j <= m; j++) {
    if (p[j] != 0) {
        assignment[p[j] - 1] = j - 1;
    }
}
return {-v[0], assignment}; // Overflow warning!
}

template<typename T>
std::pair<T, std::vector<int>> max_assignment(std::vector<std::vector<T>> value) {
    if (value.empty()) {
        return {0, {}};
    }
    T max_value = value[0][0];
    for (const auto &row : value) {
        max_value = std::max(max_value, *std::max_element(row.begin(), row.end()));
    }
    for (auto &row : value) {
        for (T &x : row) {
            x = max_value - x;
        }
    }
}

```

```

    auto [min_cost, assignment] = min_assignment(value);
    return {max_value * static_cast<T>(value.size()) - min_cost, assignment}; // Overflow warning!
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<vector<int64_t>> cost{
        {9, 2, 7, 8},
        {6, 4, 3, 7},
        {5, 8, 1, 8},
    };
    auto [min_cost, assignment] = min_assignment(cost);
    assert(min_cost == 9);
    assert((assignment == vector<int>{1, 0, 2}));

    auto [max_value, max_jobs] = max_assignment(cost);
    assert(max_value == 24);
    assert((max_jobs == vector<int>{0, 3, 1}));
    return 0;
}

```

4.6.7 Stable Marriage (Gale-Shapley)

The stable marriage problem takes n men and n women, with every person ranking all members of the opposite group in strict order of preference, and seeks a bijection between the two groups. An unmatched pair (m, w) is a blocking pair if man m prefers woman w to his assigned partner and woman w prefers man m to hers. A matching is stable if it has no blocking pair. Unlike maximum matching, the goal is therefore not to maximize the number of pairs, but to find a complete pairing that satisfies this stability condition.

The Gale-Shapley algorithm produces a stable matching, and a stable matching always exists. Each free man proposes to the most preferred woman who has not yet rejected him. A woman provisionally accepts her best proposal so far and rejects the rest, possibly breaking a previous engagement when a man she prefers proposes later. Because every man descends his list at most once and each proposal is constant work after precomputing the women's rankings, the algorithm runs in quadratic time. The result is man-optimal: every man receives the best partner he could have in any stable matching.

- `stable_marriage(men_pref, women_pref)` returns a vector `wife` where `wife[m]` is the woman matched to man `m` in the man-optimal stable matching, given the input preference lists each as n by n matrices. `men_pref[m]` lists the women in `m`'s order of preference, most preferred first, and `women_pref[w]` lists the men in `w`'s order of preference.

Time Complexity: $O(n^2)$ per call, where n is the number of men (equivalently, women).

Space Complexity: $O(n^2)$ auxiliary for the precomputed rankings.

```

#include <cassert>
#include <queue>
#include <vector>

std::vector<int> stable_marriage(
    const std::vector<std::vector<int>>& men_pref, const std::vector<std::vector<int>>& women_pref
) {
    int n = static_cast<int>(men_pref.size());
    assert(static_cast<int>(women_pref.size()) == n);
    std::vector<std::vector<int>> rank(n, std::vector<int>(n));
    for (int w = 0; w < n; w++) {

```



```

    for (int i = 0; i < n; i++) {
        rank[w][women_pref[w][i]] = i; // How woman w ranks each man (smaller is better).
    }
}
std::vector<int> next_proposal(n, 0), husband(n, -1), wife(n, -1);
std::queue<int> free_men;
for (int m = 0; m < n; m++) {
    free_men.push(m);
}
while (!free_men.empty()) {
    int m = free_men.front();
    free_men.pop();
    int w = men_pref[m][next_proposal[m]++];
    if (husband[w] == -1) {
        husband[w] = m;
        wife[m] = w;
    } else if (rank[w][m] < rank[w][husband[w]]) {
        wife[husband[w]] = -1;
        free_men.push(husband[w]);
        husband[w] = m;
        wife[m] = w;
    } else {
        free_men.push(m); // Rejected; m remains free and will propose to the next woman.
    }
}
return wife;
}

```

Example Usage

```

using namespace std;

int main() {
    // All three men prefer w0 > w1 > w2; all three women prefer m2 > m1 > m0.
    vector<vector<int>> men_pref{{0, 1, 2}, {0, 1, 2}, {0, 1, 2}};
    vector<vector<int>> women_pref{{2, 1, 0}, {2, 1, 0}, {2, 1, 0}};

    // Cascading proposals leave m2 with w0, m1 with w1, m0 with w2.
    assert((stable_marriage(men_pref, women_pref) == vector<int>{2, 1, 0}));
    return 0;
}

```

4.7 Exponential Graph Problems

4.7.1 Maximum Clique (Bron-Kerbosch)

Given an undirected graph, determine a maximum clique: a largest subset of nodes in which every pair is connected by an edge. A weighted variant instead seeks the clique of maximum total node weight, given a weight for each node.

The Bron-Kerbosch algorithm recursively extends a growing clique, tracking the set of candidate nodes that may still be added and the set of nodes already excluded. Choosing a pivot node to avoid branching on its neighbors prunes large parts of the search, keeping it efficient on most graphs.

- `max_clique()` returns the maximum clique size for a global, bidirectionally pre-populated adjacency matrix `adj`, whose row and column indices represent the nodes.
- `max_clique_weighted()` additionally uses global `w` and returns the maximum clique weight.

These implementations use `Mask`, which is `uint64_t` by default, so the number of nodes must be less than `MASK_BITS`. Node weights must be nonnegative.

Time Complexity: $O(3^{n/3})$ per call to `max_clique()` and `max_clique_weighted()`, where n is the number of nodes.

Space Complexity:

- $O(n^2)$ for storage of the graph, where n is the number of nodes.
- $O(n)$ auxiliary stack space for `max_clique()` and `max_clique_weighted()`.

```
#include <algorithm>
#include <cassert>
#include <climits>
#include <cstdint>
#include <vector>

using Mask = uint64_t;
const int MASK_BITS = sizeof(Mask) * CHAR_BIT;

std::vector<std::vector<char>> adj;
std::vector<int64_t> w;

std::vector<Mask> build_mask_graph() {
    int n = static_cast<int>(adj.size());
    assert(n < MASK_BITS);
    std::vector<Mask> g(n);
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            if (adj[i][j]) {
                g[i] |= Mask{1} << j;
            }
        }
    }
    return g;
}

int max_clique_rec(const std::vector<Mask> &g, Mask curr, Mask pool, Mask excl) {
    if (pool == 0) {
        return __builtin_popcountll(curr);
    }
    int res = 0;
    int pivot = __builtin_ctzll(pool | excl);
    Mask candidates = pool & ~g[pivot];
    while (candidates != 0) {
        int u = __builtin_ctzll(candidates);
        res = std::max(res, max_clique_rec(g, curr | (Mask{1} << u), pool & g[u], excl & g[u]));
        pool ^= Mask{1} << u;
        excl |= Mask{1} << u;
        candidates &= candidates - 1;
    }
    return res;
}

int max_clique() {
    int n = static_cast<int>(adj.size());
    std::vector<Mask> g = build_mask_graph();
    return max_clique_rec(g, 0, (Mask{1} << n) - 1, 0);
}

int64_t max_clique_weighted_rec(const std::vector<Mask> &g, int64_t curr, Mask pool, Mask excl) {
    if (pool == 0) {
        return curr;
    }
    int64_t res = -1;
    int pivot = __builtin_ctzll(pool | excl);
    Mask z = pool & ~g[pivot];
    while (z != 0) {
        int u = __builtin_ctzll(z);
        int64_t next_weight = curr + w[u]; // Overflow warning!
        int64_t next = max_clique_weighted_rec(g, next_weight, pool & g[u], excl & g[u]);
        res = std::max(res, next);
    }
}
```

```

    pool ^= Mask{1} << u;
    excl |= Mask{1} << u;
    z &= z - 1;
}
return res;
}

int64_t max_clique_weighted() {
    int n = static_cast<int>(adj.size());
    assert(static_cast<int>(w.size()) == n);
    std::vector<Mask> g = build_mask_graph();
    return max_clique_weighted_rec(g, 0, (Mask{1} << n) - 1, 0);
}

```

Example Usage

```

void add_edge(int u, int v) {
    adj[u][v] = true;
    adj[v][u] = true;
}

int main() {
    //           0(10)
    //         // |
    //        // |
    //       // |
    //      // |
    // 1(20)---+---3(40)---4(50)
    //        \ | /
    //         \ | /
    //        \ | /
    //       // |
    //      // |
    //     2(30)----+
    int nodes = 5;
    adj.assign(nodes, std::vector<char>(nodes));
    w.assign(nodes, 0);
    add_edge(0, 1);
    add_edge(0, 2);
    add_edge(0, 3);
    add_edge(1, 2);
    add_edge(1, 3);
    add_edge(2, 3);
    add_edge(3, 4);
    add_edge(4, 2);
    w[0] = 10;
    w[1] = 20;
    w[2] = 30;
    w[3] = 40;
    w[4] = 50;
    assert(max_clique() == 4);
    assert(max_clique_weighted() == 120);
    return 0;
}

```

4.7.2 Graph Coloring (Backtracking)

Given an undirected graph, assign a color to every node such that no pair of adjacent nodes have the same color, and that the total number of colors used is minimized.

This implementation finds the chromatic number by backtracking: nodes (ordered by degree to prune sooner) are colored one at a time, each tried with every color already in use plus one new color, while the fewest colors seen in a complete coloring so far bounds the search and cuts off any branch at least as costly.

- `color_graph()` populates `color` and returns minimum color count for a global, bidirectionally pre-populated adjacency matrix `adj`, whose row and column indices represent the nodes.

Time Complexity: $O(n^n)$ per call in the worst case, where n is the number of nodes.

Space Complexity:

- $O(n^2)$ for storage of the graph, where n is the number of nodes.
- $O(n)$ auxiliary stack space and $O(n)$ auxiliary heap space.

```
#include <algorithm>
#include <numeric>
#include <vector>

std::vector<std::vector<char>>> adj;
int min_colors;
std::vector<int> color, curr, order;

void color_rec(int pos, int used_colors) {
    if (used_colors >= min_colors) {
        return;
    }
    if (pos == static_cast<int>(order.size())) {
        color = curr;
        min_colors = used_colors;
        return;
    }
    int u = order[pos];
    for (int c = 0; c < used_colors; c++) {
        bool valid = true;
        for (int i = 0; i < pos; i++) {
            int v = order[i];
            if (adj[u][v] && curr[v] == c) {
                valid = false;
                break;
            }
        }
        if (valid) {
            curr[u] = c;
            color_rec(pos + 1, used_colors);
            curr[u] = -1;
        }
    }
    curr[u] = used_colors;
    color_rec(pos + 1, used_colors + 1);
    curr[u] = -1;
}

int color_graph() {
    int n = static_cast<int>(adj.size());
    if (n == 0) {
        color.clear();
        return 0;
    }
    std::vector<int> degree(n);
    for (int u = 0; u < n; u++) {
        for (int v = 0; v < n; v++) {
            degree[u] += adj[u][v];
        }
    }
    order.resize(n);
    std::iota(order.begin(), order.end(), 0);
    std::sort(order.begin(), order.end(), [&](int u, int v) {
        if (degree[u] != degree[v]) {
            return degree[u] > degree[v];
        }
        return u < v;
    });
};
```

```

min_colors = n + 1;
color.assign(n, -1);
curr.assign(n, -1);
color_rec(0, 0);
std::vector<int> remap(min_colors, -1);
int colors = 0;
for (int &c : color) {
    if (remap[c] == -1) {
        remap[c] = colors++;
    }
    c = remap[c];
}
return min_colors;
}

```

Example Usage

```

#include <cassert>
using namespace std;

void add_edge(int u, int v) {
    adj[u][v] = true;
    adj[v][u] = true;
}

int main() {
    //      0
    //     / |
    //    1---4---2
    //     \ / /
    //      3---+
    int nodes = 5;
    adj.assign(nodes, std::vector<char>(nodes));
    add_edge(0, 1);
    add_edge(0, 4);
    add_edge(1, 3);
    add_edge(1, 4);
    add_edge(2, 3);
    add_edge(2, 4);
    add_edge(3, 4);
    int colors = color_graph();
    assert(colors == 3);
    for (int u = 0; u < nodes; u++) {
        for (int v = u + 1; v < nodes; v++) {
            if (adj[u][v]) {
                assert(color[u] != color[v]);
            }
        }
    }
    assert((color == vector<int>{0, 1, 1, 0, 2}));
    return 0;
}

```

4.7.3 Edge Coloring (Misra-Gries)

Given a simple undirected graph with maximum degree D , assign a color to every edge so that no two edges sharing an endpoint receive the same color. Vizing's theorem guarantees that $D + 1$ colors always suffice for a simple graph (while D colors are clearly necessary), and the Misra-Gries algorithm constructs such a coloring. Deciding between D and $D + 1$ colors is NP-hard in general, though bipartite graphs can always be edge-colored with exactly D colors.

The algorithm colors edges one at a time. To color an edge (u, v) it grows a maximal “fan” of u : a sequence of neighbors starting at v in which each successive edge’s color is free at the previous fan node. It then picks a color c free at u and a color d free at the fan’s last node, and flips the colors along the maximal alternating c/d path leaving u . This frees d at u , after which rotating the fan shifts colors over by one and assigns d to the new edge. Each step keeps the partial coloring proper and never introduces a $(D + 2)$ -th color.

The graph must be simple: the implementation stores one color per ordered node pair in a dense matrix, so self-loops and parallel edges between the same pair cannot be represented.

- `edge_coloring(n, edges)` returns a vector of colors, one per entry of `edges`, using colors in the range $[0, D]$ where D is the maximum degree. Nodes are numbered $[0, n)$, and `edges[i]` is an undirected edge $u-v$.

Time Complexity: $O(n^3 \cdot m)$ per call in the worst case, where n is the number of nodes and m is the number of edges; substantially faster in practice.

Space Complexity:

- $O(n^2)$ to store the per-edge colors.
- $O(n)$ auxiliary heap and stack space.

```
#include <utility>
#include <vector>

class EdgeColoring {
    int n;
    std::vector<std::vector<int>>> color; // color[u][v] is the color of edge (u, v), or 0 if none.

    bool is_free(int x, int c) const {
        for (int w = 0; w < n; w++) {
            if (color[x][w] == c) {
                return false;
            }
        }
        return true;
    }

    int free_color(int x) const {
        int c = 1;
        while (!is_free(x, c)) {
            c++;
        }
        return c;
    }

    void set_color(int u, int v, int c) {
        color[u][v] = c;
        color[v][u] = c;
    }

    // Flip colors c and d along the maximal alternating path of those colors starting at x.
    void invert_path(int x, int d, int c) {
        for (int w = 0; w < n; w++) {
            if (color[x][w] == d) {
                set_color(x, w, 0);
                invert_path(w, c, d);
                set_color(x, w, c);
                return;
            }
        }
    }

public:
    EdgeColoring(int n, const std::vector<std::pair<int, int>> &edges)
        : n(n), color(n, std::vector<int>(n)) {
        for (const auto &e : edges) {
            add_edge(e.first, e.second);
        }
    }
};
```

```

}

void add_edge(int u, int v) {
    // Build a maximal fan of u starting at v: fan[0] = v, and color(u, fan[i+1]) is free at fan[i].
    std::vector<int> fan{v};
    std::vector<char> used(n, false);
    used[v] = true;
    while (true) {
        int last = fan.back(), next = -1;
        for (int w = 0; w < n; w++) {
            if (!used[w] && color[u][w] != 0 && is_free(last, color[u][w])) {
                next = w;
                break;
            }
        }
        if (next == -1) {
            break;
        }
        fan.push_back(next);
        used[next] = true;
    }
    int c = free_color(u), d = free_color(fan.back());
    invert_path(u, d, c);
    // Find the prefix of the fan to rotate: the first node on which d is now free.
    int k = 0;
    while (k < static_cast<int>(fan.size()) && color[u][fan[k]] != d && !is_free(fan[k], d)) {
        k++;
    }
    for (int j = 0; j < k; j++) {
        set_color(u, fan[j], color[u][fan[j + 1]]);
    }
    set_color(u, fan[k], d);
}

std::vector<int> coloring(const std::vector<std::pair<int, int>> &edges) const {
    std::vector<int> result;
    result.reserve(edges.size());
    for (const auto &e : edges) {
        result.push_back(color[e.first][e.second] - 1); // Shift internal 1..D+1 colors to 0..D.
    }
    return result;
}

};

std::vector<int> edge_coloring(int n, const std::vector<std::pair<int, int>> &edges) {
    return EdgeColoring(n, edges).coloring(edges);
}

```

Example Usage

```

#include <cassert>
#include <set>
using namespace std;

int main() {
    // A 5-cycle has maximum degree 2 but is odd, so it needs 3 colors (D + 1).
    vector<pair<int, int>> cycle{{0, 1}, {1, 2}, {2, 3}, {3, 4}, {4, 0}};
    vector<int> col = edge_coloring(5, cycle);
    set<int> palette(col.begin(), col.end());
    assert(palette.size() == 3);

    // Verify the coloring is proper: edges sharing an endpoint differ in color.
    int m = static_cast<int>(cycle.size());
    for (int i = 0; i < m; i++) {
        for (int j = i + 1; j < m; j++) {
            bool share = cycle[i].first == cycle[j].first || cycle[i].first == cycle[j].second ||

```

```

        cycle[i].second == cycle[j].first || cycle[i].second == cycle[j].second;
    if (share) {
        assert(col[i] != col[j]);
    }
}

// A bipartite graph (here K_{2,3}) is edge-colored with exactly D colors.
vector<pair<int, int>> bip{{0, 2}, {0, 3}, {0, 4}, {1, 2}, {1, 3}, {1, 4}};
vector<int> bcol = edge_coloring(5, bip);
set<int> bpal(bcol.begin(), bcol.end());
assert(static_cast<int>(bpal.size()) == 3); // Max degree is 3.
return 0;
}

```

4.7.4 Shortest Hamiltonian Path and Cycle (Held-Karp)

Given a complete, weighted, directed graph, find a minimum-length Hamiltonian path or cycle. A Hamiltonian path visits each node exactly once. A Hamiltonian cycle additionally returns to its starting node; this cycle version is the traveling salesman problem (TSP).

Both functions use the Held-Karp subset dynamic program. Let $dp(S, i)$ be the shortest path that visits exactly the nodes in S and ends at i . Its final edge must come from some $j \in S \setminus \{i\}$, giving the recurrence $dp(S, i) = \min_j (dp(S \setminus \{i\}, j) + w(j, i))$. Paths initialize every singleton set; cycles initialize only node 0 and consider only subsets containing it. After processing the full set, the algorithm chooses the best final node, including its return edge to node 0 for a cycle. Reconstruction repeats the same minimization backward instead of storing a separate parent table.

- `shortest_hamiltonian_path()` populates `path` and returns the minimum Hamiltonian path length for a global, pre-populated adjacency matrix `adj`.
- `shortest_hamiltonian_cycle()` fixes the start at node 0, populates `path`, and returns the minimum Hamiltonian cycle length for a global, pre-populated adjacency matrix `adj`.

Since this implementation uses bitmasks with signed 32-bit integers, the maximum number of nodes must be less than 31.

Time Complexity: $O(2^n \cdot n^2)$ per call, where n is the number of nodes.

Space Complexity:

- $O(n^2)$ for storage of the graph, where n is the number of nodes.
- $O(2^n \cdot n)$ auxiliary.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <vector>

const int64_t INF = INT64_MAX / 4;
std::vector<std::vector<int64_t>> adj;
std::vector<int> path;

int64_t shortest_hamiltonian_path() {
    int n = static_cast<int>(adj.size());
    assert(1 <= n && n < 31);
    int max_mask = (1 << n) - 1;
    std::vector<std::vector<int64_t>> dp(max_mask + 1, std::vector<int64_t>(n, INF));
    path.assign(n, 0);
    for (int i = 0; i < n; i++) {
        dp[1 << i][i] = 0;
    }
    for (int mask = 1; mask <= max_mask; mask++) {

```



```

    for (int i = 0; i < n; i++) {
        if ((mask & (1 << i)) != 0) {
            for (int j = 0; j < n; j++) {
                if ((mask & (1 << j)) != 0 && dp[mask ^ (1 << i)][j] != INF) {
                    int64_t candidate = dp[mask ^ (1 << i)][j] + adj[j][i]; // Overflow warning!
                    dp[mask][i] = std::min(dp[mask][i], candidate);
                }
            }
        }
    }
}

int64_t res = INF;
for (int i = 0; i < n; i++) {
    res = std::min(res, dp[max_mask][i]);
}

int mask = max_mask, old = -1;
for (int i = n - 1; i >= 0; i--) {
    int best = -1;
    for (int j = 0; j < n; j++) {
        if ((mask & (1 << j)) != 0 &&
            (best == -1 || dp[mask][best] + (old == -1 ? 0 : adj[best][old]) >
             dp[mask][j] + (old == -1 ? 0 : adj[j][old]))) {
            best = j;
        }
    }
    path[i] = best;
    mask ^= 1 << best;
    old = best;
}

return res;
}

int64_t shortest_hamiltonian_cycle() {
    int n = static_cast<int>(adj.size());
    assert(1 <= n && n < 31);
    if (n == 1) {
        path = {0};
        return 0;
    }

    int max_mask = (1 << n) - 1;
    std::vector<std::vector<int64_t>> dp(max_mask + 1, std::vector<int64_t>(n, INF));
    path.assign(n, 0);
    dp[1][0] = 0;
    for (int mask = 1; mask <= max_mask; mask += 2) {
        for (int i = 1; i < n; i++) {
            if ((mask & (1 << i)) != 0) {
                for (int j = 0; j < n; j++) {
                    if ((mask & (1 << j)) != 0 && dp[mask ^ (1 << i)][j] != INF) {
                        int64_t candidate = dp[mask ^ (1 << i)][j] + adj[j][i]; // Overflow warning!
                        dp[mask][i] = std::min(dp[mask][i], candidate);
                    }
                }
            }
        }
    }

    int64_t res = INF;
    for (int i = 1; i < n; i++) {
        res = std::min(res, dp[max_mask][i] + adj[i][0]);
    }

    int mask = max_mask, old = 0;
    for (int i = n - 1; i >= 1; i--) {
        int best = -1;
        for (int j = 1; j < n; j++) {
            if ((mask & (1 << j)) != 0 &&
                (best == -1 || dp[mask][best] + adj[best][old] > dp[mask][j] + adj[j][old])) {
                best = j;
            }
        }
    }
}

```

```

    path[i] = best;
    mask ^= 1 << best;
    old = best;
}
return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // Complete directed graph where 0->1->2->0 costs 1 + 2 + 3.
    // The reverse cycle is intentionally more expensive.
    //
    //      0 <-----+
    //      ^ | \
    //      / | \
    //      w=7 | \ w=1 | w=3
    //      | | \
    //      | | \
    //      | v w=2 v |
    //      +--- 1 -----> 2
    //      ^-----/
    //      w=5
    int nodes = 3;
    adj.assign(nodes, std::vector<int64_t>(nodes));
    adj[0][1] = 1;
    adj[0][2] = 1;
    adj[1][0] = 7;
    adj[1][2] = 2;
    adj[2][0] = 3;
    adj[2][1] = 5;
    int64_t path_len = shortest_hamiltonian_path();
    assert(path_len == 3);
    assert((path == vector<int>{0, 1, 2}));
    int64_t cycle_len = shortest_hamiltonian_cycle();
    assert(cycle_len == 6);
    assert((path == vector<int>{0, 1, 2}));
    return 0;
}

```

Chapter 5

Optimization

5.1 Monotone Predicate Search

5.1.1 Binary Search

Binary search locates a target in sorted data, or more generally the transition point of a monotonic Boolean predicate that changes from false to true or true to false at most once. Each step evaluates the midpoint of the current interval and keeps the half that still contains the transition, halving the search space until it closes on the answer. The same technique applies to non-increasing or non-decreasing numeric functions. Unlike searching through an array, discrete binary search is not restricted by available memory, making it useful for implicit domains that are too large to materialize, including broad integer or real intervals.

- `binary_search_first_true(lo, hi, pred)` takes signed integer boundaries for the search space $[lo, hi)$ (i.e. including `lo`, but excluding `hi`) and returns the smallest integer `k` in $[lo, hi)$ for which the predicate `pred(k)` tests true. If `pred(k)` tests false for the entire input range, then the original `hi` is returned. The caller must ensure `pred` is monotonically ascending on the input range, i.e. returning all false for some (possibly empty) prefix, followed by all true in some (possibly empty) suffix. E.g., patterns `01`, `00`, and `11` are allowed, but `10` is disallowed.
- `binary_search_last_true(lo, hi, pred)` takes signed integer boundaries for the search space $[lo, hi)$ (i.e. including `lo`, but excluding `hi`) and returns the largest integer `k` in $[lo, hi)$ for which the predicate `pred(k)` tests true. If `pred(k)` tests false for the entire input range, then the original `lo - 1` is returned, so `lo` must exceed the minimum representable value. The caller must ensure `pred` is monotonically descending on the input range, i.e. returning all true for some (possibly empty) prefix, followed by all false in some (possibly empty) suffix. E.g., patterns `10`, `00`, and `11` are allowed, but `01` is disallowed.
- `fbinary_search(lo, hi, pred)` is the equivalent of `binary_search_first_true()` on floating point predicates. Since any interval of real numbers is dense, the exact target cannot be found due to floating point error. Instead, a value that is very close to the border between false and true is returned. The precision of the answer depends on the number of repetitions the function performs. Since each repetition bisects the search space, the absolute error of the answer is $1/(2^n)$ times the distance between `lo` and `hi` after n repetitions. Although the error can be controlled by looping until the distance shrinks to an arbitrary epsilon, it is simpler to let the loop run for a desired number of iterations until floating point arithmetic breaks down. 100 iterations is usually sufficient, since the search space will be reduced to 2^{-100} (roughly 10^{-30}) times its original size. This implementation can be modified to find the “last true” point by simply interchanging the assignments of `lo` and `hi` in the if-else statements.

Time Complexity:

- $O(\log n)$ calls to `pred()` per call to `binary_search_first_true()` and `binary_search_last_true()`, where n is the distance between `lo` and `hi`.
- $O(n)$ calls to `pred()` per call to `fbinary_search()`, where n is the number of iterations.

Space Complexity: $O(1)$ auxiliary for all operations.

```

template<typename Int, typename Pred>
Int binary_search_first_true(Int lo, Int hi, Pred pred) { // 000[1]11
    while (lo < hi) {
        Int mid = lo + (hi - lo) / 2;
        if (pred(mid)) {
            hi = mid;
        } else {
            lo = mid + 1;
        }
    }
    return lo; // hi if all false
}

template<typename Int, typename Pred>
Int binary_search_last_true(Int lo, Int hi, Pred pred) { // 11[1]000
    while (lo < hi) {
        Int mid = lo + (hi - lo) / 2;
        if (pred(mid)) {
            lo = mid + 1;
        } else {
            hi = mid;
        }
    }
    return lo - 1;
}

template<typename Pred>
double fbinary_search(double lo, double hi, Pred pred) { // 000[1]11
    double mid;
    for (int i = 0; i < 100; i++) {
        mid = (lo + hi) / 2.0;
        if (pred(mid)) {
            hi = mid;
        } else {
            lo = mid;
        }
    }
    return lo;
}

```

Example Usage

```

#include <cassert>
#include <cmath>

int main() {
    assert(binary_search_first_true(0, 7, [](int x) { return x >= 3; }) == 3);
    assert(binary_search_first_true(0, 7, [](int x) { return true; }) == 0);
    assert(binary_search_first_true(0, 7, [](int x) { return false; }) == 7);
    assert(binary_search_first_true(4, 4, [](int x) { return x >= 4; }) == 4);

    assert(binary_search_last_true(0, 7, [](int x) { return x <= 5; }) == 5);
    assert(binary_search_last_true(0, 7, [](int x) { return true; }) == 6);
    assert(binary_search_last_true(0, 7, [](int x) { return false; }) == -1);
    assert(binary_search_last_true(4, 4, [](int x) { return x <= 4; }) == 3);

    double res = fbinary_search(-10.0, 10.0, [](double x) { return x >= 1.2345; });
    assert(fabs(res - 1.2345) < 1e-15);
    return 0;
}

```

5.1.2 Exponential Search

Finds the first value where a monotone Boolean predicate becomes true when an upper bound is not known in advance. Exponential search first grows the search range by powers of two until it brackets the transition point, then finishes with ordinary binary search.

This is useful for answer-search problems over unbounded or very large integer domains. The predicate must eventually become true, or the search may overflow or run forever unless an explicit limit is added.

- `exponential_search_first_true(lo, pred)` returns the smallest integer `x` greater than or equal to `lo` such that `pred(x)` is true.

Time Complexity: $O(\log n)$ calls to `pred()` per call, where n is the distance from `lo` to the first true value.

Space Complexity: $O(1)$ auxiliary.

```
#include <cstdint>

template<typename Int, typename Pred>
Int exponential_search_first_true(Int lo, Pred pred) { // 000[1]11
    if (pred(lo)) {
        return lo;
    }
    Int step = 1;
    while (!pred(lo + step)) {
        step *= 2;
    }
    Int hi = lo + step;
    lo += step / 2;
    while (lo < hi) {
        Int mid = lo + (hi - lo) / 2;
        if (pred(mid)) {
            hi = mid;
        } else {
            lo = mid + 1;
        }
    }
    return lo;
}
```

Example Usage

```
#include <cassert>

int main() {
    auto at_least_1000 = [](int x) { return x >= 1000; };
    assert(exponential_search_first_true(1000, at_least_1000) == 1000);
    assert(exponential_search_first_true(0, at_least_1000) == 1000);
    assert(exponential_search_first_true(5, at_least_1000) == 1000);
    auto square_large_enough = [](int64_t x) { return x * x >= 123456789LL; };
    assert(exponential_search_first_true(0LL, square_large_enough) == 11112);
    return 0;
}
```

5.1.3 Local Minimum (Binary Search)

Finds a non-strict local minimum in an array, that is, an index i whose value is no greater than the values at its existing neighbors. Unlike ternary search, this does not require the array to be unimodal, but the result is only a local minimum and need not be the global minimum.

Compare the middle element with its right neighbor. If the sequence slopes downward, the right half contains a local minimum; otherwise, the middle element or the left half contains one. Each comparison therefore discards half of the remaining indices.

- `local_min_index(a)` returns the index of a non-strict local minimum in the nonempty array `a`. Either endpoint may be returned when it is no greater than its sole neighbor.

Time Complexity: $O(\log n)$ comparisons per call, where n is the size of `a`.

Space Complexity: $O(1)$ auxiliary.

```
#include <cassert>
#include <vector>

template<typename T>
int local_min_index(const std::vector<T> &a) {
    assert(!a.empty());
    int lo = 0, hi = static_cast<int>(a.size()) - 1;
    while (lo < hi) {
        int mid = lo + (hi - lo) / 2;
        if (a[mid] > a[mid + 1]) {
            lo = mid + 1;
        } else {
            hi = mid;
        }
    }
    return lo;
}
```

Example Usage

```
int main() {
    assert(local_min_index(std::vector<int>{9, 7, 3, 4, 8}) == 2);
    assert(local_min_index(std::vector<int>{1, 2, 3, 4}) == 0);

    std::vector<int> duplicates{4, 2, 2, 3};
    int i = local_min_index(duplicates);
    assert((i == 1 || i == 2) && duplicates[i] == 2);
    return 0;
}
```

5.2 Unimodal and Continuous Search

5.2.1 Ternary and Golden Section Search

Given a unimodal function f , find the position of its global minimum or maximum: a value x such that the function strictly decreases up to x and strictly increases after it (for a minimum), or the reverse (for a maximum). All three searches below exploit the same idea: evaluate two interior points, and the comparison reveals one side of the interval that cannot contain the optimum, so it can be discarded. They differ only in how the probes are placed and whether the domain is continuous or discrete. Each maximum routine is the minimum routine applied to the negated function, since the location of a maximum of f is the location of a minimum of $-f$.

Ternary search splits the interval into thirds. It is the simplest to reason about, but it recomputes both probe points each step, spending two function evaluations per iteration while the interval shrinks to two-thirds.

Golden-section search places the probes at the golden ratio so that one probe is reused as an interior point of the next, smaller interval. It therefore performs only one new evaluation per iteration while shrinking the interval by a

factor of $1/\varphi$ (about 0.618), and reaches a given accuracy in roughly 2.4 times fewer evaluations than ternary search. Prefer it when the function is expensive to evaluate.

The discrete version operates on an integer domain, narrowing by thirds until at most three candidates remain and then scanning them. It returns the index of the optimum rather than a real coordinate. As with ternary search on reals, the function must be strictly unimodal.

- `ternary_search_min(lo, hi, f, eps = 1e-12)` and `ternary_search_max(lo, hi, f, eps = 1e-12)` return a `double x` in `[lo, hi]` optimizing the continuous unimodal function `f`, to within the optional absolute error `eps`.
- `golden_section_min(lo, hi, f, eps = 1e-12)` and `golden_section_max(lo, hi, f, eps = 1e-12)` do the same with one function evaluation per iteration.
- `discrete_ternary_min(lo, hi, f)` and `discrete_ternary_max(lo, hi, f)` return the integer index in range `[lo, hi]` optimizing the unimodal function `f`, breaking ties toward the smaller index.

Time Complexity:

- $O(\log(n/\text{eps}))$ calls to `f()` per continuous-search call, where n is the distance between `lo` and `hi`. Golden-section makes about half as many per iteration as ternary search.
- $O(\log n)$ calls to `f()` per discrete-search call.

Space Complexity: $O(1)$ auxiliary for all operations.

```
template<typename Fn>
double ternary_search_min(double lo, double hi, Fn f, double eps = 1e-12) {
    while (hi - lo > eps) {
        double m1 = lo + (hi - lo) / 3, m2 = hi - (hi - lo) / 3;
        if (f(m1) < f(m2)) {
            hi = m2;
        } else {
            lo = m1;
        }
    }
    return (lo + hi) / 2;
}

template<typename Fn>
double ternary_search_max(double lo, double hi, Fn f, double eps = 1e-12) {
    return ternary_search_min(lo, hi, [&](double x) { return -f(x); }, eps);
}

template<typename Fn>
double golden_section_min(double lo, double hi, Fn f, double eps = 1e-12) {
    const double r = 0.6180339887498949; // 1 / phi = (sqrt(5) - 1) / 2.
    double m1 = hi - r * (hi - lo), m2 = lo + r * (hi - lo);
    double f1 = f(m1), f2 = f(m2);
    while (hi - lo > eps) {
        if (f1 < f2) {
            hi = m2;
            m2 = m1;
            f2 = f1;
            m1 = hi - r * (hi - lo);
            f1 = f(m1);
        } else {
            lo = m1;
            m1 = m2;
            f1 = f2;
            m2 = lo + r * (hi - lo);
            f2 = f(m2);
        }
    }
    return (lo + hi) / 2;
}

template<typename Fn>
double golden_section_max(double lo, double hi, Fn f, double eps = 1e-12) {
    return golden_section_min(lo, hi, [&](double x) { return -f(x); }, eps);
}
```

```

}

template<typename Int, typename Fn>
Int discrete_ternary_min(Int lo, Int hi, Fn f) {
    while (hi - lo > 2) {
        Int m1 = lo + (hi - lo) / 3, m2 = hi - (hi - lo) / 3;
        if (f(m1) < f(m2)) {
            hi = m2 - 1;
        } else {
            lo = m1 + 1;
        }
    }
    Int best = lo;
    auto best_value = f(best);
    for (Int k = lo + 1; k <= hi; k++) {
        auto value = f(k);
        if (value < best_value) {
            best = k;
            best_value = value;
        }
    }
    return best;
}

template<typename Int, typename Fn>
Int discrete_ternary_max(Int lo, Int hi, Fn f) {
    return discrete_ternary_min(lo, hi, [&](Int x) { return -f(x); });
}

```

Example Usage

```

#include <cassert>
#include <cmath>
#include <vector>
using namespace std;

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-6;
}

int main() {
    // Parabola opening up with vertex at x = -2.
    auto up = [](double x) { return 3 * x * x + 12 * x - 12; };
    // Parabola opening down with vertex at x = 2/19.
    auto down = [](double x) { return -5.7 * x * x + 1.2 * x + 88; };

    assert(EQ(ternary_search_min(-1000, 1000, up), -2));
    assert(EQ(golden_section_min(-1000, 1000, up), -2));
    assert(EQ(ternary_search_max(-1000, 1000, down), 2.0 / 19));
    assert(EQ(golden_section_max(-1000, 1000, down), 2.0 / 19));

    // Discrete unimodal arrays: a valley and a peak.
    vector<int> valley{5, 3, 1, 2, 4, 6};
    vector<int> peak{1, 3, 7, 6, 2};
    auto v = [&](int i) { return valley[i]; };
    auto p = [&](int i) { return peak[i]; };
    assert(discrete_ternary_min(0, static_cast<int>(valley.size()) - 1, v) == 2); // Value 1.
    assert(discrete_ternary_max(0, static_cast<int>(peak.size()) - 1, p) == 2);    // Value 7.
    return 0;
}

```


5.2.2 2D Hill Climbing

Given a continuous geometric objective $f(x, y)$ and a (possibly arbitrary) starting guess (x_0, y_0) , search for a small value using two-dimensional hill climbing.

The heuristic starts at the guess and evaluates one step in each of eight evenly spaced directions. It moves to the best improving neighbor, checks all eight directions again, and reduces the step size when no neighbor improves the answer. This repeats until the step size falls below the chosen threshold. Success depends heavily on the behavior of f and the initial guess, so the result is not guaranteed to be the global minimum.

- `hill_climb_min(f, x0, y0, &best_x, &best_y, step_min = 1e-9, step_max = 1e6)` returns an approximate minimum value of function `f` reached by hill climbing an initial guess `(x0, y0)`. If either optional pointer `best_x` or `best_y` is supplied, the corresponding coordinate of the point attaining the returned value is stored through it. The search starts with step size `step_max` and stops below `step_min`.

Time Complexity: $O(k + s)$ calls to `f()` per call, where k is the number of times the step size is reduced and s is the total number of improving moves accepted.

Space Complexity: $O(1)$ auxiliary.

```
template<typename Fn>
double hill_climb_min(
    Fn f, double x0, double y0, double *best_x = nullptr, double *best_y = nullptr,
    const double step_min = 1e-9, const double step_max = 1e6
) {
    static const double inv_sqrt2 = 0.7071067811865476;
    static const double dx[] = {1, inv_sqrt2, 0, -inv_sqrt2, -1, -inv_sqrt2, 0, inv_sqrt2};
    static const double dy[] = {0, inv_sqrt2, 1, inv_sqrt2, 0, -inv_sqrt2, -1, -inv_sqrt2};
    double x = x0, y = y0, res = f(x0, y0);
    for (double step = step_max; step > step_min;) {
        double next_value = res, next_x = x, next_y = y;
        bool found = false;
        for (int i = 0; i < 8; i++) {
            double x2 = x + step * dx[i], y2 = y + step * dy[i];
            double value = f(x2, y2);
            if (value < next_value) {
                next_x = x2;
                next_y = y2;
                next_value = value;
                found = true;
            }
        }
        if (!found) {
            step /= 2.0;
        } else {
            x = next_x;
            y = next_y;
            res = next_value;
        }
    }
    if (best_x != nullptr) {
        *best_x = x;
    }
    if (best_y != nullptr) {
        *best_y = y;
    }
    return res;
}
```

Example Usage

```
#include <cassert>
#include <cmath>
```

```

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-9;
}

// Paraboloid with global minimum at f(2, 3) = 0.
double f(double x, double y) {
    return (x - 2) * (x - 2) + (y - 3) * (y - 3);
}

int main() {
    double x, y;
    assert(EQ(hill_climb_min(f, 0, 0, &x, &y), 0));
    assert(EQ(x, 2) && EQ(y, 3));
    return 0;
}

```

5.2.3 Simulated Annealing

Approximately minimizes an energy function $E(s)$ over an arbitrary state space using simulated annealing. This randomized heuristic is useful when the states may be discrete, the energy has many local minima, and gradients or stronger structural properties are unavailable.

Starting from an initial state s , each iteration generates a random neighboring state s' . An improving move is always accepted, while a worsening move is accepted with probability $\exp(-(E(s') - E(s))/T)$ at temperature T . High temperatures encourage exploration across energy barriers; as the temperature decreases, the search increasingly favors improvements. The best state is tracked separately because the current state may later move to a worse one.

- `anneal_min(initial, energy, rand_neighbor, rng, ...)` returns `(best_energy, best_state)` given an initial state `initial`, a callable `energy(state)` that returns the state's energy, and a callable `rand_neighbor(state, temp, rng)` that returns a randomly chosen nearby state. Optional parameters default to `temp_start = 1000`, `temp_end = 1e-6`, and `cooling_rate = 0.995`. `temp_end` must be positive and less than `temp_start`, while `cooling_rate` must be strictly between 0 and 1.

The geometric schedule repeatedly multiplies the temperature by `cooling_rate`, performing $k = \lceil \log(T_{end}/T_{start}) / \log(r) \rceil$ iterations. Temperature has the same scale as energy: multiplying every energy by a constant requires multiplying both temperature bounds by the same constant to preserve acceptance probabilities. For continuous states, temperature can also control the neighbor step size. Slower cooling, independent restarts, and a time-based stopping condition often improve results, but simulated annealing never proves optimality.

This generic interface copies each proposed state and recomputes its energy. For large permutation states, adapt the loop to mutate and revert moves in place and update the energy from the affected terms instead.

Time Complexity: $O(k(C_E + C_N))$ per call, where k is the number of temperature levels and C_E and C_N are the costs of evaluating the energy and generating a neighboring state.

Space Complexity: $O(S)$ for the current, next, and best states, where S is the size of one state.

```

#include <cassert>
#include <cmath>
#include <random>
#include <utility>

template<typename State, typename Energy, typename Neighbor>
std::pair<double, State> anneal_min(
    State initial, Energy energy, Neighbor rand_neighbor, std::mt19937 &rng,
    const double temp_start = 1000, const double temp_end = 1e-6, const double cooling_rate = 0.995
) {
    assert(0 < temp_end && temp_end < temp_start);
    assert(0 < cooling_rate && cooling_rate < 1);
    std::uniform_real_distribution<double> unit(0.0, 1.0);

```

```

State current = initial, best = initial;
double current_energy = energy(current), best_energy = current_energy;
for (double temp = temp_start; temp > temp_end; temp *= cooling_rate) {
    State next = rand_neighbor(current, temp, rng);
    double next_energy = energy(next);
    if (next_energy <= current_energy ||
        unit(rng) < std::exp((current_energy - next_energy) / temp)) {
        current = next;
        current_energy = next_energy;
        if (current_energy < best_energy) {
            best = current;
            best_energy = current_energy;
        }
    }
}
return {best_energy, best};
}

```

Example Usage

```

#include <algorithm>
#include <vector>
using namespace std;

int main() {
    mt19937 rng(1234567); // Fixed seed for reproducibility.

    // Escape the local minimum near x = 1 and find the lower minimum near x = -1.
    auto energy_1d = [](double x) { return (x * x - 1) * (x * x - 1) + 0.2 * x; };
    auto neighbor_1d = [](double x, double temp, mt19937 &gen) {
        uniform_real_distribution<double> move(-1.0, 1.0);
        return x + move(gen) * sqrt(temp);
    };
    auto [best_energy, best_x] = anneal_min(1.0, energy_1d, neighbor_1d, rng, 2.0, 1e-8);
    assert(best_energy < -0.2 && best_x < -1);

    // Approximate TSP with a permutation state and random 2-opt reversals.
    vector<pair<double, double>> points{{0, 0}, {1, 0}, {2, 0}, {2, 1},
                                         {2, 2}, {1, 2}, {0, 2}, {0, 1}};
    auto tour_length = [&](const vector<int> &tour) {
        double length = 0;
        for (int i = 0; i < static_cast<int>(tour.size()); i++) {
            auto [x1, y1] = points[tour[i]];
            auto [x2, y2] = points[tour[(i + 1) % tour.size()]];
            length += hypot(x1 - x2, y1 - y2);
        }
        return length;
    };
    auto reverse_segment = [](const vector<int> &tour, double, mt19937 &gen) {
        vector<int> next = tour;
        uniform_int_distribution<int> pick(0, static_cast<int>(tour.size()) - 1);
        int lo = pick(gen), hi = pick(gen);
        if (lo > hi) {
            swap(lo, hi);
        }
        reverse(next.begin() + lo, next.begin() + hi + 1);
        return next;
    };
    vector<int> initial_tour{0, 4, 2, 6, 1, 5, 3, 7};
    auto [best_length, best_tour] =
        anneal_min(initial_tour, tour_length, reverse_segment, rng, 5.0, 1e-6, 0.995);
    assert(fabs(best_length - 8) < 1e-9);
    assert(fabs(tour_length(best_tour) - best_length) < 1e-9);
    return 0;
}

```

5.3 Dynamic Programming Tricks

5.3.1 Divide-and-Conquer DP Optimization

Computes one layer of a dynamic program of the form $dp_cur[i] = \min(dp_prev[k] + cost(k, i))$, where $0 \leq k \leq i$. Let $opt[i]$ be the smallest index k attaining this minimum. Divide-and-conquer optimization applies when $opt[i] \leq opt[i + 1]$, so the best transition index moves only to the right as i increases.

To compute $dp_cur[lo..hi]$, evaluate its midpoint and record its best transition as $best_k$. Monotonicity restricts every optimum in the left half to indices at most $best_k$, and every optimum in the right half to indices at least $best_k$. Recursively applying these bounds avoids checking every transition for every state. The caller must verify the required monotonicity property.

- `compute_dp_layer(dp_prev, dp_cur, lo, hi, opt_lo, opt_hi, cost)` fills $dp_cur[lo..hi]$ using candidate transition indices in $[opt_lo, opt_hi]$. The template parameter `cost` must be callable such that `cost(k, i)` returns the transition cost from previous state k to current state i .

Time Complexity: $O(n \log n)$ calls to `cost()` per call, when computing one layer whose states each have $O(n)$ possible transitions.

Space Complexity: $O(\log n)$ auxiliary stack space.

```
#include <algorithm>
#include <stdint>
#include <vector>

const int64_t INF = (1LL << 62);

template<typename Cost>
void compute_dp_layer(
    const std::vector<int64_t> &dp_prev, std::vector<int64_t> &dp_cur, int lo, int hi, int opt_lo,
    int opt_hi, Cost cost
) {
    if (lo > hi) {
        return;
    }
    int mid = lo + (hi - lo) / 2;
    int64_t best = INF;
    int best_k = opt_lo;
    int upper = std::min(mid, opt_hi);
    for (int k = opt_lo; k <= upper; k++) {
        int64_t candidate = dp_prev[k] + cost(k, mid); // Overflow warning!
        if (candidate < best) {
            best = candidate;
            best_k = k;
        }
    }
    dp_cur[mid] = best;
    compute_dp_layer(dp_prev, dp_cur, lo, mid - 1, opt_lo, best_k, cost);
    compute_dp_layer(dp_prev, dp_cur, mid + 1, hi, best_k, opt_hi, cost);
}
```

Example Usage

```
#include <cassert>
using namespace std;

struct SquareSegmentCost {
    vector<int64_t> prefix;

    explicit SquareSegmentCost(const vector<int> &a) : prefix(a.size() + 1) {
        for (int i = 0; i < static_cast<int>(a.size()); i++) {
            prefix[i + 1] = prefix[i] + a[i];
        }
    }
};
```

```

    }
}

int64_t operator()(int k, int i) const {
    // Cost of placing the half-open segment a[k..i) in one group.
    int64_t sum = prefix[i] - prefix[k];
    return sum * sum;
}
};

int main() {
    vector<int> a{1, 2, 3, 4};
    int n = static_cast<int>(a.size());
    SquareSegmentCost cost(a);

    // Base layer for zero groups: only the empty prefix is reachable.
    vector<int64_t> dp_prev(n + 1, INF), dp_cur(n + 1, INF);
    dp_prev[0] = 0;

    // The first layer places each prefix a[0..j) into one group.
    compute_dp_layer(dp_prev, dp_cur, 1, n, 0, n - 1, cost);
    assert(dp_cur[4] == 100); // One group: (1 + 2 + 3 + 4)^2.

    // The second layer chooses the best split between two groups.
    vector<int64_t> dp_two(n + 1, INF);
    compute_dp_layer(dp_cur, dp_two, 1, n, 0, n - 1, cost);
    assert(dp_two[4] == 52); // (1 + 2 + 3)^2 + 4^2.
    return 0;
}

```

5.3.2 SMAWK Row Minima

Finds the position of the minimum in every row of a totally monotone matrix in time linear in its dimensions, without ever materializing the matrix. A matrix is totally monotone (for row minima) when the column index of each row's leftmost minimum never decreases as the row index increases, and this holds not just for the whole matrix but for every submatrix. Monge matrices, where $M_{i,j} + M_{i+1,j+1} \leq M_{i,j+1} + M_{i+1,j}$, are the most common source of such matrices and arise throughout dynamic programming on intervals and partitions.

The SMAWK algorithm interleaves two steps. REDUCE discards columns that cannot hold the minimum of any remaining row, leaving at most as many candidate columns as active rows. INTERPOLATE recurses on the odd-indexed active rows, then recovers each even-indexed active row's minimum by scanning only the candidate-column range bounded by its already-solved neighbors; total monotonicity guarantees those scans advance monotonically and cost $O(R + C)$ overall, where R is the number of rows and C is the number of columns in the original matrix. This removes the logarithmic factor of divide-and-conquer DP optimization, but unlike that technique it is purely offline: every matrix entry must be available on demand from the oracle, so it cannot be used when an entry depends on a row minimum computed earlier in the same pass.

- `smawk_row_minima(rows, cols, get)` returns a vector `arg` of length `rows`, where `arg[r]` is the column index minimizing `get(r, c)` over $0 \leq c < \text{cols}$ (on ties, the smallest qualifying column is chosen). The input matrix is supplied as a totally monotone function `get(r, c)` with bounds `rows` ≥ 0 and `cols` ≥ 1 .

Time Complexity: $O(R + C)$ evaluations of `get` per call, where R is `rows` and C is `cols`.

Space Complexity: $O(R + C)$ auxiliary.

```

#include <cassert>
#include <numeric>
#include <vector>

template<typename Get>
std::vector<int> smawk_rec(
    const Get &get, const std::vector<int> &active_rows, const std::vector<int> &candidate_cols

```

```

) {
    int num_rows = static_cast<int>(active_rows.size());
    std::vector<int> arg(num_rows, -1);
    if (num_rows == 0) {
        return arg;
    }
    // Base case: a single row or column is cheap to scan directly.
    if (num_rows <= 2 || candidate_cols.size() <= 2) {
        for (int row_pos = 0; row_pos < num_rows; row_pos++) {
            for (int c : candidate_cols) {
                if (arg[row_pos] == -1 ||
                    get(active_rows[row_pos], c) < get(active_rows[row_pos], arg[row_pos])) {
                    arg[row_pos] = c;
                }
            }
        }
        return arg;
    }
    // REDUCE: keep only columns that can be the minimum of some row.
    std::vector<int> kept;
    if (static_cast<int>(candidate_cols.size()) > num_rows) {
        for (int c : candidate_cols) {
            while (!kept.empty()) {
                int r = active_rows[static_cast<int>(kept.size()) - 1];
                if (get(r, kept.back()) <= get(r, c)) {
                    break;
                }
                kept.pop_back();
            }
            if (static_cast<int>(kept.size()) < num_rows) {
                kept.push_back(c);
            }
        }
    } else {
        kept = candidate_cols;
    }
    // INTERPOLATE: solve the odd rows recursively, then fill the even rows between known bounds.
    std::vector<int> odd_rows;
    for (int row_pos = 1; row_pos < num_rows; row_pos += 2) {
        odd_rows.push_back(active_rows[row_pos]);
    }
    std::vector<int> odd_arg = smawk_rec(get, odd_rows, kept);
    for (int odd_pos = 0; odd_pos < static_cast<int>(odd_arg.size()); odd_pos++) {
        arg[2 * odd_pos + 1] = odd_arg[odd_pos];
    }
    int col_pos = 0;
    for (int row_pos = 0; row_pos < num_rows; row_pos += 2) {
        int hi = (row_pos + 1 < num_rows) ? arg[row_pos + 1] : kept.back();
        int r = active_rows[row_pos];
        while (true) {
            int c = kept[col_pos];
            if (arg[row_pos] == -1 || get(r, c) < get(r, arg[row_pos])) {
                arg[row_pos] = c;
            }
            if (c == hi) {
                break; // Stop at the upper bound; the next even row starts scanning from here.
            }
            col_pos++;
        }
    }
    return arg;
}

template<typename Get>
std::vector<int> smawk_row_minima(int rows, int cols, const Get &get) {
    assert(rows >= 0 && cols >= 1);
    std::vector<int> row_ids(rows), col_ids(cols);
    std::iota(row_ids.begin(), row_ids.end(), 0);

```

```
std::iota(col_ids.begin(), col_ids.end(), 0);
return smawk_rec(get, row_ids, col_ids);
}
```

Example Usage

```
#include <cassert>
#include <vector>
using namespace std;

int main() {
    // An explicit totally monotone matrix; row minima sit at non-decreasing columns.
    vector<vector<int>> m{
        {10, 17, 13, 28, 23},
        {17, 22, 16, 29, 23},
        {24, 28, 22, 34, 24},
        {11, 13, 6, 17, 7},
        {45, 44, 32, 37, 23},
    };
    auto get = [&](int r, int c) { return m[r][c]; };
    vector<int> arg = smawk_row_minima(5, 5, get);
    assert((arg == vector<int>{0, 2, 2, 2, 4}));
    return 0;
}
```

5.3.3 Knuth DP Optimization

Computes minimum costs for every half-open interval $[l, r)$ using the recurrence

$dp[l][r] = \min(dp[l][k] + dp[k][r] + cost(l, r))$. Each state chooses a split k strictly inside the interval and pays $cost(l, r)$ after solving the two resulting subintervals.

A direct implementation checks every split and takes $O(n^3)$. Knuth optimization applies when the smallest optimal splits satisfy $opt[l][r - 1] \leq opt[l][r] \leq opt[l + 1][r]$. The two neighboring states therefore bound the splits that must be checked for $dp[l][r]$, reducing the total number of candidate checks to $O(n^2)$. Common applications include optimal binary search trees and some range merging problems. The caller must verify the required monotonicity property for the chosen $cost(l, r)$; the quadrangle inequality and interval monotonicity of the cost are common sufficient conditions.

- `knuth_interval_dp(n, cost, &opt)` computes minimum costs for all half-open intervals $[l, r)$ over n items. The template parameter `cost` must be callable such that `cost(l, r)` returns the interval cost added after choosing the best split. If the optional pointer `opt` is supplied, it is filled with the chosen split points.

Time Complexity: $O(n^2)$ calls to `cost()` and $O(n^2)$ candidate split checks per call.

Space Complexity: $O(n^2)$ for the `dp` and `opt` tables.

```
#include <algorithm>
#include <cstdint>
#include <vector>

const int64_t INF = (1LL << 62);

template<typename Cost>
std::vector<std::vector<int64_t>> knuth_interval_dp(
    int n, Cost cost, std::vector<std::vector<int>> *out = nullptr
) {
    std::vector<std::vector<int64_t>> dp(n + 1, std::vector<int64_t>(n + 1, 0));
    std::vector<std::vector<int>> opt(n + 1, std::vector<int>(n + 1, 0));
    // Empty and one-item intervals have cost 0.
    for (int i = 0; i <= n; i++) {
        opt[i][i] = i;
```

```

    if (i < n) {
        opt[i][i + 1] = i + 1;
    }
}
for (int len = 2; len <= n; len++) {
    for (int l = 0; l + len <= n; l++) {
        int r = l + len;
        dp[l][r] = INF;
        int split_lo = std::max(opt[l][r - 1], l + 1);
        int split_hi = std::min(opt[l + 1][r], r - 1);
        int64_t interval_cost = cost(l, r);
        for (int k = split_lo; k <= split_hi; k++) {
            int64_t candidate = dp[l][k] + dp[k][r] + interval_cost; // Overflow warning!
            if (candidate < dp[l][r]) {
                dp[l][r] = candidate;
                opt[l][r] = k;
            }
        }
    }
}
}
if (out != nullptr) {
    *out = opt;
}
return dp;
}

```

Example Usage

```

#include <cassert>
using namespace std;

struct MergeCost {
    vector<int64_t> prefix;

    explicit MergeCost(const vector<int> &a) : prefix(a.size() + 1) {
        for (int i = 0; i < static_cast<int>(a.size()); i++) {
            prefix[i + 1] = prefix[i] + a[i];
        }
    }

    int64_t operator()(int l, int r) const {
        // Cost of merging the half-open interval a[l..r) into one group.
        return prefix[r] - prefix[l];
    }
};

int main() {
    vector<int> a{1, 2, 3, 4};
    vector<vector<int>> opt;
    vector<vector<int64_t>> dp = knuth_interval_dp(a.size(), MergeCost(a), &opt);
    assert(dp[0][4] == 19); // Merge 1 + 2, then 3 + 3, then 6 + 4.
    assert(opt[0][4] == 3); // The final merge splits [1, 2, 3] from [4].
    return 0;
}

```

5.3.4 Lagrangian Relaxation (Aliens Trick)

Given an optimization problem with a constraint on the number of chosen objects, Lagrangian relaxation moves that count into the objective by charging a penalty per object. If the relaxed problem can be solved quickly and the number chosen is monotone as the penalty changes, binary search recovers the best value for the desired count.

For maximization, solve the relaxed problem as `score - penalty*count`, breaking ties toward larger `count`. As `penalty` increases, the optimal count never increases. The largest penalty whose relaxed optimum still chooses at least `target_count` objects gives the exact constrained score after adding back `penalty*target_count`.

- `lagrangian_maximize(target_count, lo, hi, solve)` returns the maximum original score using exactly `target_count` objects. The callable `solve(penalty)` must return `relaxed_score` and `count` as a pair and must tie-break toward larger `count`. Penalties are searched over the half-open range `[lo, hi)`; `lo` must choose at least `target_count` objects, while `hi` must be greater than the last penalty that does so.

This technique is often called the Aliens trick. Widen the penalty range when in doubt.

Time Complexity: $O(\log n)$ calls to `solve()` per call, where n is the distance between `lo` and `hi`.

Space Complexity: $O(1)$ auxiliary.

```
#include <cassert>
#include <cstdint>
#include <utility>
#include <vector>

template<typename Solve>
int64_t lagrangian_maximize(int target_count, int64_t lo, int64_t hi, Solve solve) {
    // The chosen count is nonincreasing, so find the last penalty that still chooses enough objects.
    while (lo < hi) {
        int64_t mid = lo + (hi - lo) / 2;
        if (solve(mid).second >= target_count) {
            lo = mid + 1;
        } else {
            hi = mid;
        }
    }
    int64_t penalty = lo - 1;
    auto [relaxed_score, count] = solve(penalty);
    assert(count >= target_count);
    return relaxed_score + penalty * target_count; // Overflow warning!
}
```

Example Usage

```
int main() {
    std::vector<int64_t> value{10, 7, 3, 2};
    auto choose_profitable = [&](int64_t penalty) {
        int64_t score = 0;
        int count = 0;
        for (int64_t x : value) {
            int64_t gain = x - penalty;
            if (gain >= 0) { // Include zero gains to break ties toward larger counts.
                score += gain;
                count++;
            }
        }
        return std::pair<int64_t, int>{score, count};
    };

    assert(lagrangian_maximize(2, 0, 20, choose_profitable) == 17);
    assert(lagrangian_maximize(3, 0, 20, choose_profitable) == 20);
    return 0;
}
```

5.3.5 Convex Hull Trick (Semi-Dynamic)

Given a set of pairs (m, b) specifying lines of the form $y = mx + b$, answer queries at specified x -coordinates, each asking for the minimum y -value over all given lines. This is useful for dynamic programming recurrences of the form `dp[i] = min(m[j] * x[i] + b[j])`. Only the lower envelope of the lines can ever answer a query, so each added line pops previously stored lines that it renders useless, and ascending queries advance a pointer along the envelope so every line is visited at most twice.

The following implementation is a concise, semi-dynamic version of the convex hull optimization technique. It supports an interlaced sequence of `add_line()` and `query()` calls, as long as the preconditions of descending `m` and ascending `x` are satisfied. As a result, it may be necessary to sort the lines and queries before calling the functions. In that case, the overall time complexity will be dominated by the sorting step.

- `SemiDynamicCHT()` constructs an empty hull.
- `add_line(m, b)` inserts the line $y = mx + b$. The caller must ensure that slope `m` is less than or equal to the slope of every line added so far.
- `query(x)` returns the minimum y -value among all inserted lines at coordinate `x`. The caller must ensure that query coordinates are nondecreasing across calls. At least one line must have been inserted.

Overflow warning: `add_line()` compares intersections by cross-multiplying slope and intercept differences, a product on the order of the squared coefficient magnitude. For large `m/b` (roughly beyond 10^9 with `int64_t`), cast that comparison to `__int128`. `query()` evaluates `m * x + b`, which also needs to fit in `int64_t`.

Time Complexity: $O(n + q)$ for any interlaced sequence containing n calls to `add_line()` and q calls to `query()`. Each line is removed or passed by the query pointer at most once, so both operations take amortized $O(1)$ time.

Space Complexity:

- $O(n)$ for storage of the lines.
- $O(1)$ auxiliary for `add_line()` and `query()`.

```
#include <cassert>
#include <cstdint>
#include <vector>

class SemiDynamicCHT {
    std::vector<int64_t> slopes, intercepts;
    int ptr = 0;

public:
    void add_line(int64_t m, int64_t b) {
        int len = static_cast<int>(slopes.size());
        if (len > 0 && slopes.back() == m) {
            if (intercepts.back() <= b) {
                return;
            }
            slopes.pop_back();
            intercepts.pop_back();
            len--;
        }
        if (ptr > len) {
            ptr = len;
        }
        // Overflow warning!
        while (len > 1 && (intercepts[len - 2] - intercepts[len - 1]) * (m - slopes[len - 1]) >=
                    (intercepts[len - 1] - b) * (slopes[len - 1] - slopes[len - 2])) {
            len--;
        }
        slopes.resize(len);
        intercepts.resize(len);
        slopes.push_back(m);
        intercepts.push_back(b);
    }

    int64_t query(int64_t x) {
```

```

assert(!slopes.empty());
if (ptr >= static_cast<int>(slopes.size())) {
    ptr = static_cast<int>(slopes.size()) - 1;
}
// Overflow warning!
while (ptr + 1 < static_cast<int>(slopes.size()) &&
        slopes[ptr + 1] * x + intercepts[ptr + 1] <= slopes[ptr] * x + intercepts[ptr]) {
    ptr++;
}
return slopes[ptr] * x + intercepts[ptr];
}
};

```

Example Usage

```

#include <cassert>

int main() {
    SemiDynamicCHT cht;
    cht.add_line(3, 0);
    cht.add_line(2, 1);
    cht.add_line(1, 2);
    cht.add_line(0, 6);
    assert(cht.query(0) == 0);
    assert(cht.query(1) == 3);
    assert(cht.query(2) == 4);
    assert(cht.query(3) == 5);
    cht.add_line(0, 7); // Same slope, worse intercept; ignored.
    assert(cht.query(4) == 6);
    cht.add_line(0, 5); // Same slope, better intercept; replaces the old constant line.
    assert(cht.query(5) == 5);
    return 0;
}

```

5.3.6 Convex Hull Trick (Fully-Dynamic)

Maintains a dynamic set of lines $y = mx + b$ and answers minimum value queries at integer points. Lines and queries may arrive in arbitrary order, making this useful for dynamic programming recurrences of the form `dp[i] = min(m[j] * x[i] + b[j])` without monotone slopes or query coordinates. The same interface can answer maximum queries when constructed with `query_max = true`.

This fully dynamic convex hull stores the envelope in a self-balancing binary search tree (`std::set`). Inserting a line removes any neighbors it dominates, and each remaining line records the last coordinate where it is optimal, so a query is one tree lookup. Unlike a Li Chao tree, this structure requires no fixed coordinate domain and its complexity depends on the number of lines rather than the domain width. Prefer it when the domain is unknown or very large, or when maximum queries are needed; a Li Chao tree is often simpler when a manageable domain is known.

- `HullOptimizer(query_max = false)` constructs an empty hull. By default, `query(x)` minimizes; if `query_max` is true, `query(x)` maximizes.
- `add_line(m, b)` inserts line $y = mx + b$. Lines may be added in any order.
- `query(x)` returns the best y -value among all inserted lines at coordinate x . At least one line must have been inserted, and query coordinates may be supplied in any order.

Overflow warning: minimization negates the coefficients, border updates subtract slopes/intercepts, and `query()` evaluates `m * x + b`, so those intermediate values must fit in `int64_t`.

Time Complexity: $O(\log n)$ amortized per call to `add_line()` and $O(\log n)$ per call to `query()`, where n is the number of lines added.

Space Complexity:

- $O(n)$ for storage of the lines.
- $O(1)$ auxiliary for `add_line()` and `query()`.

```

#include <cassert>
#include <cstdint>
#include <set>

class HullOptimizer {
    struct Line {
        int64_t m, b;
        mutable int64_t xhi;
        bool is_query;

        Line(int64_t m, int64_t b, int64_t xhi = 0, bool is_query = false)
            : m(m), b(b), xhi(xhi), is_query(is_query) {}

        bool operator< (const Line &l) const { return l.is_query ? xhi < l.xhi : m < l.m; }
    };

    std::multiset<Line> hull;
    bool query_max;

    using HullIter = std::multiset<Line>::iterator;

    bool update_border(HullIter x, HullIter y) {
        if (y == hull.end()) {
            x->xhi = INT64_MAX;
            return false;
        }
        if (x->m == y->m) {
            x->xhi = (x->b > y->b) ? INT64_MAX : INT64_MIN;
        } else {
            int64_t a = y->b - x->b, b = x->m - y->m; // Overflow warning!
            x->xhi = a / b - ((a ^ b) < 0 && a % b);
        }
        return x->xhi >= y->xhi;
    }

public:
    explicit HullOptimizer(bool query_max = false) : query_max(query_max) {}

    void add_line(int64_t m, int64_t b) {
        if (!query_max) {
            m = -m;
            b = -b;
        }
        auto z = hull.insert(Line(m, b));
        auto y = z++;
        auto x = y;
        while (update_border(y, z)) {
            z = hull.erase(z);
        }
        if (x != hull.begin() && update_border(--x, y)) {
            update_border(x, y = hull.erase(y));
        }
        while ((y = x) != hull.begin() && (--x)->xhi >= y->xhi) {
            update_border(x, hull.erase(y));
        }
    }

    int64_t query(int64_t x) const {
        assert(!hull.empty());
        Line q(0, 0, x, true);
        auto it = hull.lower_bound(q);
        int64_t res = it->m * x + it->b; // Overflow warning!
        return query_max ? res : -res;
    }
};

```

```

    }
};

```

Example Usage

```

#include <cassert>

int main() {
    HullOptimizer h;
    h.add_line(3, 0);
    h.add_line(0, 6);
    h.add_line(1, 2);
    h.add_line(2, 1);
    // Minimize among y = 3x, 6, x + 2, and 2x + 1.
    assert(h.query(0) == 0);
    assert(h.query(2) == 4);
    assert(h.query(1) == 3);
    assert(h.query(3) == 5);

    HullOptimizer mx(true);
    mx.add_line(3, 0);
    mx.add_line(0, 6);
    mx.add_line(1, 2);
    // Same interface can maximize when constructed with query_max = true.
    assert(mx.query(0) == 6);
    assert(mx.query(3) == 9);
    return 0;
}

```

5.3.7 Li Chao Tree

Maintains a dynamic set of lines $y = mx + b$ and answers minimum value queries at integer points. Lines and queries may arrive in arbitrary order, making this useful for dynamic programming recurrences of the form `dp[i] = min(m[j] * x[i] + b[j])` without monotone slopes or query coordinates.

A Li Chao tree fixes an inclusive coordinate domain $[l_0, h_1]$ and stores one line at each node of a dynamic segment tree. Inserting a line keeps the line that wins at the midpoint and recurses only into the half where the other line may still win; a query takes the best value along one root-to-leaf path. Prefer this structure when the domain is known and manageable. The fully dynamic convex hull in the previous section needs no fixed domain and also supports maximum queries, while a Li Chao tree has a simpler invariant and is easier to extend to lines active only on subintervals.

- `LiChaoTree(l0, hi)` constructs an empty tree over integer domain $[l_0, h_1]$.
- `add_line(m, b)` inserts line $y = mx + b$. Lines may be added in any order.
- `query(x)` returns the minimum y -value among all inserted lines at coordinate x . At least one line must have been inserted, and query coordinates may be supplied in any order.

Overflow warning: each comparison and query evaluates `m * x + b`, which must fit in `int64_t`.

Time Complexity: $O(\log d)$ per call to `add_line()` and `query()`, where d is the distance between `l0` and `hi`.

Space Complexity:

- $O(n)$ node storage for n lines inserted, since each insertion creates at most one node.
- $O(\log d)$ auxiliary stack space per operation.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <utility>

```

```

const int64_t INF = INT64_MAX / 4;

class LiChaoTree {
    struct Line {
        int64_t m, b;
        Line(int64_t m = 0, int64_t b = INF) : m(m), b(b) {}
        int64_t eval(int64_t x) const { return m * x + b; } // Overflow warning!
    };

    struct Node {
        Line f;
        Node *left, *right;
        explicit Node(const Line &f) : f(f), left(nullptr), right(nullptr) {}
    };

    Node *root;
    int64_t lo, hi;

    static void clean_up(Node *n) {
        if (n == nullptr) {
            return;
        }
        clean_up(n->left);
        clean_up(n->right);
        delete n;
    }

    static void add_line(Node *&n, int64_t lo, int64_t hi, Line f) {
        if (n == nullptr) {
            n = new Node(f);
            return;
        }
        int64_t mid = lo + (hi - lo) / 2;
        bool left_better = f.eval(lo) < n->f.eval(lo);
        bool mid_better = f.eval(mid) < n->f.eval(mid);
        if (mid_better) {
            std::swap(f, n->f);
        }
        if (lo == hi) {
            return;
        }
        if (left_better != mid_better) {
            add_line(n->left, lo, mid, f);
        } else {
            add_line(n->right, mid + 1, hi, f);
        }
    }

    static int64_t query(Node *n, int64_t lo, int64_t hi, int64_t x) {
        if (n == nullptr) {
            return INF;
        }
        int64_t res = n->f.eval(x);
        if (lo == hi) {
            return res;
        }
        int64_t mid = lo + (hi - lo) / 2;
        if (x <= mid) {
            return std::min(res, query(n->left, lo, mid, x));
        }
        return std::min(res, query(n->right, mid + 1, hi, x));
    }

public:
    LiChaoTree(int64_t lo, int64_t hi) : root(nullptr), lo(lo), hi(hi) { assert(lo <= hi); }

    ~LiChaoTree() { clean_up(root); }
    LiChaoTree(const LiChaoTree &) = delete;

```

```

LiChaoTree &operator=(const LiChaoTree &) = delete;
void add_line(int64_t m, int64_t b) { add_line(root, lo, hi, Line(m, b)); }

int64_t query(int64_t x) const {
    assert(root != nullptr);
    assert(lo <= x && x <= hi);
    return query(root, lo, hi, x);
}
};

```

Example Usage

```

#include <cassert>

int main() {
    LiChaoTree h(-10, 10);
    h.add_line(3, 0);
    h.add_line(0, 6);
    h.add_line(1, 2);
    h.add_line(2, 1);
    // Minimize among y = 3x, 6, x + 2, and 2x + 1.
    assert(h.query(0) == 0);
    assert(h.query(2) == 4);
    assert(h.query(1) == 3);
    assert(h.query(3) == 5);
    return 0;
}

```

5.4 Numerical Methods

5.4.1 Root Finding (Bracketing)

Finds an x in an interval $[a, b]$ for a continuous function f such that $f(x) = 0$. By the intermediate value theorem, a root must exist in $[a, b]$ if either endpoint is a root or the signs of $f(a)$ and $f(b)$ differ. Each bisection step evaluates the midpoint of the interval and keeps the half whose endpoints still differ in sign, so a root always remains bracketed. After n iterations, the interval width is $(b - a)/2^n$. Although it is possible to control the error by looping while $b - a$ is greater than an arbitrary epsilon, it is simpler to let the loop run for a desired number of iterations until floating point arithmetic breaks down. 100 iterations is usually sufficient, since the search space will be reduced to 2^{-100} (roughly 10^{-30}) times its original size.

- `bisection_root(f, a, b, iterations = 100)` returns a root in an interval $[a, b]$ for a continuous function f where either endpoint is a root or the endpoint values have opposite signs, using the bisection method.
- `falsi_illinois_root(f, a, b, iterations = 100)` returns a root in an interval $[a, b]$ for a continuous function f under the same endpoint conditions, using the Illinois algorithm variant of the false position (a.k.a. regula falsi) method.

Time Complexity: $O(n)$ calls to `f()` per call to `bisection_root()` and `falsi_illinois_root()`, where n is the number of iterations.

Space Complexity: $O(1)$ auxiliary for both operations.

```

#include <cassert>

template<typename Fn>
double bisection_root(Fn f, double a, double b, const int iterations = 100) {
    double fa = f(a), fb = f(b);
    assert(a <= b && ((fa <= 0 && fb >= 0) || (fa >= 0 && fb <= 0)));
    if (fa == 0) {

```

```

    return a;
}
if (fb == 0) {
    return b;
}
double m = a;
for (int i = 0; i < iterations; i++) {
    m = a + (b - a) / 2;
    double fm = f(m);
    if (fm == 0) {
        return m;
    }
    if ((fa < 0) == (fm < 0)) {
        a = m;
        fa = fm;
    } else {
        b = m;
    }
}
return m;
}

template<typename Fn>
double falsi_illinois_root(Fn f, double a, double b, const int iterations = 100) {
    double fa = f(a), fb = f(b);
    assert(a <= b && ((fa <= 0 && fb >= 0) || (fa >= 0 && fb <= 0)));
    if (fa == 0) {
        return a;
    }
    if (fb == 0) {
        return b;
    }
    double m = a;
    int side = 0;
    for (int i = 0; i < iterations; i++) {
        m = (fa * b - fb * a) / (fa - fb);
        double fm = f(m);
        if (fm == 0) {
            break;
        }
        if ((fb < 0) == (fm < 0)) {
            b = m;
            fb = fm;
            if (side < 0) {
                fa /= 2;
            }
            side = -1;
        } else {
            a = m;
            fa = fm;
            if (side > 0) {
                fb /= 2;
            }
            side = 1;
        }
    }
    return m;
}

```

Example Usage

```

#include <cmath>

double f(double x) {
    return x * x - 4 * sin(x);
}

```



```
int main() {
    assert(fabs(f(bisection_root(f, 1, 3))) < 1e-10);
    assert(fabs(f(falsi_illinois_root(f, 1, 3))) < 1e-10);
    assert(bisection_root(f, 0, 1) == 0);
    assert(falsi_illinois_root(f, 0, 1) == 0);
    return 0;
}
```

5.4.2 Root Finding (Iteration)

Finds an x for a continuous function f such that $f(x) = 0$ using iterative approximation by an initial guess that is close to the answer. Each step follows the tangent line at the current guess down to its x -intercept, that is, $x \leftarrow x - f(x)/f'(x)$. Newton's method requires an explicit definition of the function's derivative while the secant method starts with two initial guesses and approximates the derivative using the secant slope from the previous iteration. With smooth functions and good initial guesses, Newton's method has quadratic convergence, while the secant method has convergence order approximately 1.618.

- `newton_root(f, fprime, x0, eps = 1e-15, iterations = 100)` returns a root x for a function `f` with derivative `fprime` using an initial guess `x0` which should be relatively close to x .
- `secant_root(f, x0, x1, eps = 1e-15, iterations = 100)` returns a root x for a function `f` using two initial guesses `x0` and `x1` which should be relatively close to x .

Time Complexity: $O(n)$ calls to `f()` and `fprime()` per call to `newton_root()`, and $O(n)$ calls to `f()` per call to `secant_root()`, where n is the number of iterations.

Space Complexity: $O(1)$ auxiliary for both operations.

```
#include <cmath>
#include <stdexcept>

template<typename Fn, typename Deriv>
double newton_root(
    Fn f, Deriv fprime, double x0, const double eps = 1e-15, const int iterations = 100
) {
    double x = x0, error = eps + 1;
    for (int i = 0; std::isfinite(error) && error > eps && i < iterations; i++) {
        double fx = f(x);
        if (fx == 0) {
            return x;
        }
        double xnew = x - fx / fprime(x);
        error = fabs(xnew - x);
        x = xnew;
    }
    if (!std::isfinite(error) || error > eps) {
        throw std::runtime_error("Newton's method failed to converge.");
    }
    return x;
}

template<typename Fn>
double secant_root(
    Fn f, double x0, double x1, const double eps = 1e-15, const int iterations = 100
) {
    double xold = x0, fxold = f(x0), x = x1, error = eps + 1;
    for (int i = 0; std::isfinite(error) && error > eps && i < iterations; i++) {
        double fx = f(x);
        if (fx == 0) {
            return x;
        }
        double xnew = x - fx * ((x - xold) / (fx - fxold));
        error = fabs(xnew - x);
    }
}
```

```

    xold = x;
    fxold = fx;
    x = xnew;
}
if (!std::isfinite(error) || error > eps) {
    throw std::runtime_error("Secant method failed to converge.");
}
return x;
}

```

Example Usage

```

#include <cassert>

double f(double x) {
    return x * x - 4 * sin(x);
}

double fprime(double x) {
    return 2 * x - 4 * cos(x);
}

int main() {
    assert(fabs(f(newton_root(f, fprime, 3))) < 1e-10);
    assert(fabs(f(secant_root(f, 3, 2))) < 1e-10);
    auto g = [](double x) { return x * x - 2; };
    auto gprime = [](double x) { return 2 * x; };
    assert(fabs(g(newton_root(g, gprime, 1))) < 1e-10);
    return 0;
}

```

5.4.3 Polynomial Root Finding (Differentiation)

Finds every real root x of a polynomial p satisfying $p(x) = 0$ by recursively finding the roots of its derivative. Each adjacent pair of local extrema is searched using the bisection method.

The key observation is that derivative roots split the real line into intervals where the polynomial is monotone. A monotone interval contains at most one root, and if an endpoint is a root or the endpoint values have opposite signs, bisection isolates that root. Recursively solving the derivative therefore provides all interval boundaries needed to find every real root in the requested range, including roots of even multiplicity.

- `horner_eval(p, x)` evaluates the polynomial `p` of degree d (represented as a vector of size $d + 1$ where `p[i]` stores the coefficient for the x^i term) at `x`, using Horner's method.
- `find_one_root(p, a, b, eps = 1e-15)` returns a root in the interval `[a, b]` for a polynomial `p` where either endpoint is a root or the endpoint values have opposite signs, using the bisection method. If this precondition is not satisfied, then `NaN` is returned. The root is found to a tolerance of `eps` in absolute or relative error (whichever is reached first).
- `find_all_roots(p, a = -1e20, b = 1e20, eps = 1e-15)` returns a vector of all roots in the interval `[a, b]` for a polynomial `p` using the bisection method. The roots are found to a tolerance of `eps` in absolute or relative error (whichever is reached first).

Time Complexity:

- $O(n)$ per call to `horner_eval()`, where n is the degree of the polynomial.
- $O(nt)$ per call to `find_one_root()`, where n is the degree of the polynomial and t is the number of bisection iterations required to reach the requested tolerance.
- $O(n^3t)$ per call to `find_all_roots()`, where n is the degree of the polynomial and t is the number of bisection iterations required to reach the requested tolerance.

Space Complexity:

- $O(1)$ auxiliary for `horner_eval()` and `find_one_root()`.
- $O(n)$ auxiliary stack space and $O(n^2)$ auxiliary heap space for `find_all_roots()`, where n is the degree of the polynomial.

```

#include <cmath>
#include <limits>
#include <vector>

double horner_eval(const std::vector<double> &p, double x) {
    double res = p.back();
    for (int i = static_cast<int>(p.size()) - 2; i >= 0; i--) {
        res = res * x + p[i];
    }
    return res;
}

double find_one_root(const std::vector<double> &p, double a, double b, const double eps = 1e-15) {
    double pa = horner_eval(p, a), pb = horner_eval(p, b);
    if (fabs(pa) <= eps) {
        return a;
    }
    if (fabs(pb) <= eps) {
        return b;
    }
    bool paneg = pa < 0, pbneg = pb < 0;
    if (paneg == pbneg) {
        return std::numeric_limits<double>::quiet_NaN();
    }
    while (b - a > eps && a * (1 + eps) < b && a < b * (1 + eps)) {
        double m = a + (b - a) / 2;
        if ((horner_eval(p, m) < 0) == paneg) {
            a = m;
        } else {
            b = m;
        }
    }
    return a;
}

std::vector<double> find_all_roots(
    const std::vector<double> &p, double a = -1e20, double b = 1e20, const double eps = 1e-15
) {
    std::vector<double> pprime;
    pprime.reserve(p.size() > 0 ? p.size() - 1 : 0);
    for (int i = 1; i < static_cast<int>(p.size()); i++) {
        pprime.push_back(p[i] * i);
    }
    if (pprime.empty()) {
        return {};
    }
    std::vector<double> res, r = find_all_roots(pprime, a, b, eps);
    r.push_back(b);
    for (int i = 0; i < static_cast<int>(r.size()); i++) {
        double root = find_one_root(p, i == 0 ? a : r[i - 1], r[i], eps);
        if (!std::isnan(root) && (res.empty() || root != res.back())) {
            res.push_back(root);
        }
    }
    return res;
}

```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    { // -1 + 2x - 6x^2 + 2x^3
        vector<double> p{-1.0, 2.0, -6.0, 2.0};
        vector<double> roots = find_all_roots(p);
        assert(roots.size() == 1 && fabs(horner_eval(p, roots[0])) < 1e-10);
    }
    { // -20 + 4x + 3x^2
        vector<double> p{-20.0, 4.0, 3.0};
        vector<double> roots = find_all_roots(p);
        assert(roots.size() == 2);
        assert(fabs(horner_eval(p, roots[0])) < 1e-10);
        assert(fabs(horner_eval(p, roots[1])) < 1e-10);
    }
    { // (x - 1)^2
        vector<double> p{1.0, -2.0, 1.0};
        vector<double> roots = find_all_roots(p);
        assert(roots.size() == 1 && fabs(roots[0] - 1) < 1e-10);
    }
    return 0;
}
```

5.4.4 Polynomial Root Finding (Laguerre)

Finds every complex root x for a polynomial p with complex coefficients such that $p(x) = 0$. Laguerre's method is a root-finding iteration with reliable, near-cubic convergence from almost any starting point. Once a root is found it is removed by polynomial deflation, and the iteration repeats on the lower-degree quotient until every root has been extracted.

- `horner_eval(p, x)` evaluates a complex polynomial p of degree d (represented as a vector of size $d + 1$ where `p[i]` stores the complex coefficient for the x^i term) at `x`, using Horner's method, returning a pair where the first value is the final result $p(x)$ and the second value is the quotient polynomial q from the identity $p(t) = (t - x)q(t) + p(x)$.
- `find_one_root(p, x0, eps = 1e-15, iterations = 10000)` returns a complex root x for polynomial `p` (represented as a vector of size $d + 1$ where `p[i]` stores the complex coefficient for the x^i term) using an initial guess `x0` which should be relatively close to x . The root is found to a tolerance of `eps` in absolute or relative error (whichever is reached first).
- `find_all_roots(p, eps = 1e-15, iterations = 10000)` returns a vector of all complex roots for a complex polynomial `p`. The roots are found to a tolerance of `eps` in absolute or relative error (whichever is reached first).

Time Complexity:

- $O(n)$ per call to `horner_eval()`, where n is the degree of the polynomial.
- $O(nt)$ per call to `find_one_root()`, where n is the degree of the polynomial and t is the number of iterations performed.
- $O(n^2t)$ per call to `find_all_roots()`, where n is the degree of the polynomial and t is the number of iterations performed per root.

Space Complexity:

- $O(n)$ auxiliary for `horner_eval()` and `find_one_root()`, where n is the degree of the polynomial.
- $O(n)$ auxiliary for `find_all_roots()`, where n is the degree of the polynomial.

```
#include <algorithm>
#include <cmath>
```

```

#include <complex>
#include <random>
#include <utility>
#include <vector>

using cdouble = std::complex<double>;
using cpoly = std::vector<cdouble>;

std::pair<cdouble, cpoly> horner_eval(const cpoly &p, const cdouble &x) {
    int n = static_cast<int>(p.size());
    cpoly b(std::max(1, n - 1));
    for (int i = n - 1; i > 0; i--) {
        b[i - 1] = p[i] + (i < n - 1 ? b[i] * x : 0);
    }
    return {p[0] + b[0] * x, b};
}

cpoly derivative(const cpoly &p) {
    int n = static_cast<int>(p.size());
    cpoly res(std::max(1, n - 1));
    for (int i = 1; i < n; i++) {
        res[i - 1] = p[i] * cdouble(i);
    }
    return res;
}

int comp(const cdouble &a, const cdouble &b, const double eps = 1e-15) {
    double diff = std::abs(a) - std::abs(b);
    return (diff < -eps) ? -1 : (diff > eps ? 1 : 0);
}

cdouble find_one_root(
    const cpoly &p, const cdouble &x0, const double eps = 1e-15, const int iterations = 10000
) {
    cdouble x = x0;
    int n = static_cast<int>(p.size()) - 1;
    cpoly p1 = derivative(p), p2 = derivative(p1);
    for (int i = 0; i < iterations; i++) {
        cdouble y0 = horner_eval(p, x).first;
        if (comp(y0, 0, eps) == 0) {
            break;
        }
        cdouble g = horner_eval(p1, x).first / y0;
        cdouble h = g * g - horner_eval(p2, x).first / y0;
        cdouble r = std::sqrt(cdouble(n - 1) * (h * cdouble(n) - g * g));
        cdouble d1 = g + r, d2 = g - r;
        cdouble a = cdouble(n) / (comp(d1, d2, eps) > 0 ? d1 : d2);
        x -= a;
        if (comp(a, 0, eps) == 0) {
            break;
        }
    }
    return x;
}

std::vector<cdouble> find_all_roots(
    const cpoly &p, const double eps = 1e-15, const int iterations = 10000
) {
    std::vector<cdouble> res;
    cpoly q = p;
    if (q.size() <= 1) {
        return res;
    }
    static std::mt19937 rng(std::random_device{}());
    std::uniform_real_distribution<double> unit(0.0, 1.0);
    while (q.size() > 2) {
        cdouble z(unit(rng), unit(rng));
        z = find_one_root(p, find_one_root(q, z, eps, iterations), eps, iterations);
    }
}

```

```

    q = horner_eval(q, z).second;
    res.push_back(z);
}
res.push_back(-q[0] / q[1]);
return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

void assert_roots(vector<cdouble> actual, vector<cdouble> expected) {
    assert(actual.size() == expected.size());
    vector<char> matched(actual.size());
    for (const cdouble &root : expected) {
        int found = -1;
        for (int i = 0; i < static_cast<int>(actual.size()); i++) {
            if (!matched[i] && abs(actual[i] - root) < 1e-8) {
                found = i;
                break;
            }
        }
        assert(found != -1);
        matched[found] = true;
    }
}

int main() {
    { // 140 - 13x - 8x^2 + x^3 = (x + 4)(x - 5)(x - 7)
        cpoly p;
        p.push_back(140);
        p.push_back(-13);
        p.push_back(-8);
        p.push_back(1);
        assert_roots(find_all_roots(p), {-4, 5, 7});
    }
    { // (-24+36i) + (-26+12i)x + (-30+40i)x^2 + (-26+12i)x^3 + (-6+4i)x^4
      // = ((2 + 3i)x + 6)(x + i)(2x + (6 + 4i))(xi + 1):
        cpoly p;
        p.emplace_back(-24, 36);
        p.emplace_back(-26, 12);
        p.emplace_back(-30, 40);
        p.emplace_back(-26, 12);
        p.emplace_back(-6, 4);
        assert_roots(find_all_roots(p), {{-3, -2}, {-12.0 / 13, 18.0 / 13}, {0, -1}, {0, 1}});
    }
    return 0;
}

```

5.4.5 Polynomial Root Finding (Ehrlich-Aberth)

Finds every complex root x for a polynomial p such that $p(x) = 0$. The Ehrlich-Aberth method is a simultaneous iteration: every root estimate is updated using both the Newton correction and a repulsion term from the other estimates.

This routine is intended for well-scaled polynomials when all complex roots are wanted. If only real roots are needed, a real-root isolator such as derivative recursion with bisection is usually more reliable. Multiple or tightly clustered roots may converge slowly, and very large or very small root scales may require rescaling the input or using multiprecision arithmetic.

- `eval_with_derivative(p, x)` returns a pair $(p(x), p'(x))$ for a polynomial p given as a vector `p` where `p[i]` stores the coefficient for the x^i term.
- `find_all_roots(p, eps = ROOT_EPS, iterations = 2000)` returns a vector of all complex roots for a complex polynomial given by the vector of coefficients `p`. A `vector<dbl>` overload is provided for polynomials with real coefficients. The roots are found to a tolerance of `eps` in absolute or relative error (whichever is reached first), and zero roots are removed exactly before the simultaneous iteration starts.

Time Complexity:

- $O(n)$ per call to `eval_with_derivative()`, where n is the degree of the polynomial.
- $O(n^2t)$ per call to `find_all_roots()`, where n is the degree of the polynomial and t is the number of iterations required to reach the desired precision.

Space Complexity:

- $O(1)$ auxiliary for `eval_with_derivative()`.
- $O(n)$ auxiliary for `find_all_roots()`, where n is the degree of the polynomial.

```
#include <algorithm>
#include <cmath>
#include <complex>
#include <utility>
#include <vector>

using dbl = long double;
using cdbl = std::complex<dbl>;
using cpoly = std::vector<cdbl>;

const dbl PI = acosl(-1.0L);
const dbl ZERO_EPS = 1e-30L;    // Treat coefficients and denominators this small as zero.
const dbl ROOT_EPS = 1e-18L;    // Stop once root updates are this small relative to the root.
const dbl CHECK_EPS = 1e-12L;   // Residual tolerance used by the example assertions.

bool is_zero(const cdbl &z, const dbl eps = ZERO_EPS) {
    return std::abs(z) <= eps;
}

bool is_finite(const cdbl &z) {
    return std::isfinite(z.real()) && std::isfinite(z.imag());
}

std::pair<cdbl, cdbl> eval_with_derivative(const cpoly &p, const cdbl &x) {
    cdbl value = p.back(), derivative = 0;
    for (int i = static_cast<int>(p.size()) - 2; i >= 0; i--) {
        derivative = derivative * x + value;
        value = value * x + p[i];
    }
    return {value, derivative};
}

dbl root_bound(const cpoly &p) {
    dbl res = 0;
    int n = static_cast<int>(p.size()) - 1;
    for (int i = 0; i < n; i++) {
        dbl ratio = std::abs(p[i] / p.back());
        if (ratio > 0) {
            res = std::max(res, powl(ratio, 1.0L / (n - i)));
        }
    }
    return 2 * std::max(static_cast<dbl>(1), res);
}

bool root_less(const cdbl &a, const cdbl &b) {
    if (a.real() != b.real()) {
        return a.real() < b.real();
    }
}
```

```

    return a.imag() < b.imag();
}

cpoly find_all_roots(cpoly p, const dbl eps = ROOT_EPS, const int iterations = 2000) {
    while (!p.empty() && is_zero(p.back())) {
        p.pop_back();
    }
    cpoly roots;
    while (p.size() > 1 && is_zero(p[0])) {
        roots.push_back(0);
        p.erase(p.begin());
    }
    if (p.size() <= 1) {
        return roots;
    }
    dbl scale = 0;
    for (int i = 0; i < static_cast<int>(p.size()); i++) {
        scale = std::max(scale, std::abs(p[i]));
    }
    for (int i = 0; i < static_cast<int>(p.size()); i++) {
        p[i] /= scale;
    }
    int n = static_cast<int>(p.size()) - 1;
    if (n == 1) {
        roots.push_back(-p[0] / p[1]);
        return roots;
    }
    cpoly z(n);
    dbl radius = root_bound(p), offset = PI / (2 * n);
    dbl max_radius = 2 * radius;
    for (int i = 0; i < n; i++) {
        dbl angle = offset + 2 * PI * i / n;
        z[i] = radius * cdbl(cosl(angle), sinl(angle));
    }
    for (int it = 0; it < iterations; it++) {
        bool done = true;
        cpoly next = z;
        for (int i = 0; i < n; i++) {
            auto [fx, dfx] = eval_with_derivative(p, z[i]);
            if (std::abs(fx) <= eps) {
                continue;
            }
        }
        cdbl repulsion = 0;
        for (int j = 0; j < n; j++) {
            if (i != j) {
                cdbl diff = z[i] - z[j];
                if (std::abs(diff) <= eps * eps) {
                    dbl angle = 2 * PI * (i + 1) / (n + 1);
                    diff = eps * (1 + std::abs(z[i])) * cdbl(cosl(angle), sinl(angle));
                }
                repulsion += cdbl(1) / diff;
            }
        }
        cdbl denom = dfx - fx * repulsion;
        if (is_zero(denom)) {
            done = false;
            continue;
        }
        cdbl step = fx / denom;
        if (!is_finite(step) && !is_zero(dfx)) {
            step = fx / dfx;
        }
        if (!is_finite(step)) {
            done = false;
            continue;
        }
        dbl limit = 2 * max_radius;
        if (std::abs(step) > limit) {

```



```

        step *= limit / std::abs(step);
    }
    next[i] = z[i] - step;
    if (!is_finite(next[i])) {
        next[i] = z[i];
        done = false;
        continue;
    }
    if (std::abs(next[i]) > max_radius) {
        next[i] *= max_radius / std::abs(next[i]);
    }
    if (std::abs(step) > eps * (1 + std::abs(next[i]))) {
        done = false;
    }
}
z = next;
if (done) {
    break;
}
}
for (int i = 0; i < n; i++) {
    roots.emplace_back(z[i]);
}
std::sort(roots.begin(), roots.end(), root_less);
return roots;
}

std::vector<cdbl> find_all_roots(
    const std::vector<dbl> &p, const dbl eps = ROOT_EPS, const int iterations = 2000
) {
    cpoly q(p.begin(), p.end());
    return find_all_roots(q, eps, iterations);
}

```

Example Usage

```

#include <cassert>
using namespace std;

cdbl eval(const cpoly &p, const cdbl &x) {
    return eval_with_derivative(p, x).first;
}

void assert_roots(const vector<cdbl> &actual, const vector<cdbl> &expected) {
    assert(actual.size() == expected.size());
    vector<char> matched(actual.size());
    for (const cdbl &root : expected) {
        int found = -1;
        for (int i = 0; i < static_cast<int>(actual.size()); i++) {
            if (!matched[i] && abs(actual[i] - root) < CHECK_EPS) {
                found = i;
                break;
            }
        }
        assert(found != -1);
        matched[found] = true;
    }
}

int main() {
    { // -1 + 2x - 6x^2 + 2x^3
        vector<dbl> p{-1.0, 2.0, -6.0, 2.0};
        vector<cdbl> roots = find_all_roots(p);
        assert(roots.size() == 3);
        for (const auto &root : roots) {
            assert(abs(eval(cpoly(p.begin(), p.end()), root)) < CHECK_EPS);
        }
    }
}

```

```

    }
}
{ // (-24+36i) + (-26+12i)x + (-30+40i)x^2 + (-26+12i)x^3 + (-6+4i)x^4
  // = ((2 + 3i)x + 6)(x + i)(2x + (6 + 4i))(xi + 1):
  cpoly p;
  p.emplace_back(-24, 36);
  p.emplace_back(-26, 12);
  p.emplace_back(-30, 40);
  p.emplace_back(-26, 12);
  p.emplace_back(-6, 4);
  vector<cdbl> roots = find_all_roots(p);
  assert(roots.size() == 4);
  for (const auto &root : roots) {
    assert(abs(eval(p, root)) < CHECK_EPS);
  }
  assert_roots(roots, {{-3, -2}, {-12.0L / 13, 18.0L / 13}, {0, -1}, {0, 1}});
}
return 0;
}

```

5.4.6 Integration (Simpson)

Computes the definite integral from a to b for a continuous function f using the composite Simpson's rule: the interval is split into n equal subintervals and each consecutive pair is approximated by a parabola, giving the approximation $\int_a^b f(x) dx \approx h/3 \cdot [f(x_0) + 4f(x_1) + 2f(x_2) + \dots + 4f(x_{n-1}) + f(x_n)]$ with step $h = (b - a)/n$.

Unlike adaptive quadrature, the step size here is uniform, so a feature narrower than h may be under-resolved. The benefit is a simple, non-recursive routine with predictable cost that is not fooled by the premature-termination heuristic that adaptive Simpson's rule can suffer on functions whose coarse sample points happen to fit a parabola.

- `integrate(f, a, b, n = 1000000)` returns the definite integral of a function `f` from `a` to `b` using `n` subintervals, where `n` must be even. Larger `n` trades runtime for accuracy; the leading error scales with `(b - a) * h^4` and the maximum magnitude of the fourth derivative of f over the interval.

Time Complexity: $O(n)$ per call, requiring $n + 1$ evaluations of f .

Space Complexity: $O(1)$ auxiliary.

```

#include <cassert>

template<typename Fn>
double integrate(Fn f, double a, double b, int n = 1000000) {
  assert(n > 0 && n % 2 == 0);
  double h = (b - a) / n, s = f(a) + f(b);
  for (int i = 1; i < n; i++) {
    s += f(a + h * i) * (i & 1 ? 4 : 2);
  }
  return s * h / 3;
}

```

Example Usage

```

#include <cmath>
using namespace std;

double f(double x) {
  return sin(x);
}

int main() {
  const double PI = acos(-1.0);
  assert(fabs(integrate(f, 0.0, PI / 2) - 1) < 1e-10);
}

```

```

assert(fabs(integrate(f, PI / 2, 0.0) + 1) < 1e-10);
assert(fabs(integrate(f, 1.0, 1.0)) < 1e-12);
return 0;
}

```

5.4.7 Linear Programming (Simplex)

Solves a linear programming problem using Dantzig's simplex algorithm. The canonical form of a linear programming problem is to maximize (or minimize) the dot product cx , subject to $ax \leq b$ and $x \geq 0$, where x is a vector of unknowns to be solved, c is a vector of coefficients, a is a matrix of constraint coefficients, and b is a vector of bounds.

The simplex algorithm walks between vertices of the feasible polytope, at each step pivoting to an adjacent vertex that improves the objective, until no improving move remains (an optimum is reached) or an unbounded improving direction is found.

The implementation uses a two-phase simplex tableau. Phase 1 introduces an artificial variable to find a feasible starting basis, which is essential when some constraint bounds are negative. Phase 2 then optimizes the original objective. The `basis` and `nonbasis` arrays track which variable each tableau row or column currently represents, making solution recovery independent of pivot order.

- `simplex_solve(a, b, c, &x, maximize = true, eps = 1e-10)` solves the linear programming problem for an m by n matrix `a` of real values, a length m vector `b`, and a length n vector `c`, returning 0 if an optimum was found, -1 if there are no feasible solutions, or 1 if the objective is unbounded. Set `maximize = false` to minimize instead, and use `eps` as the pivot and feasibility tolerance. If an optimum is found, `x` is assigned a vector of length `n`, where `x[j]` is the value of the j -th variable.

Time Complexity: $O(pmn)$ per call, where p is the total number of pivots across both phases. Each pivot takes $O(mn)$, and p can be exponential in the worst case, although it is usually much smaller in practice.

Space Complexity: $O(m \cdot n)$ auxiliary.

```

#include <cassert>
#include <cmath>
#include <utility>
#include <vector>

int simplex_solve(
    const std::vector<std::vector<double>>& &a, const std::vector<double> &b,
    const std::vector<double> &c, std::vector<double> *x, const bool maximize = true,
    const double eps = 1e-10
) {
    int m = static_cast<int>(a.size()), n = static_cast<int>(c.size());
    assert(x != nullptr && b.size() == a.size());
    std::vector<int> basis(m), nonbasis(n + 1);
    std::vector<std::vector<double>> tab(m + 2, std::vector<double>(n + 2));
    for (int i = 0; i < m; i++) {
        for (int j = 0; j < n; j++) {
            tab[i][j] = a[i][j];
        }
        basis[i] = n + i;
        tab[i][n] = -1;
        tab[i][n + 1] = b[i];
    }
    for (int j = 0; j < n; j++) {
        nonbasis[j] = j;
        tab[m][j] = maximize ? -c[j] : c[j];
    }
    nonbasis[n] = -1;
    tab[m + 1][n] = 1;
    auto pivot = [&](int row, int col) {
        double inv = 1.0 / tab[row][col];

```

```

for (int i = 0; i < m + 2; i++) {
    if (i == row || fabs(tab[i][col]) <= eps) {
        continue;
    }
    double factor = tab[i][col] * inv;
    for (int j = 0; j < n + 2; j++) {
        if (j != col) {
            tab[i][j] -= factor * tab[row][j];
        }
    }
    tab[i][col] = -factor;
}
for (int j = 0; j < n + 2; j++) {
    if (j != col) {
        tab[row][j] *= inv;
    }
}
tab[row][col] = inv;
std::swap(basis[row], nonbasis[col]);
};
auto simplex = [&](int phase) {
    int objective = (phase == 1 ? m + 1 : m);
    while (true) {
        int col = -1;
        for (int j = 0; j <= n; j++) {
            if (phase == 2 && nonbasis[j] == -1) {
                continue;
            }
            if (col == -1 || tab[objective][j] < tab[objective][col] - eps ||
                (fabs(tab[objective][j] - tab[objective][col]) <= eps && nonbasis[j] < nonbasis[col])) {
                col = j;
            }
        }
        if (tab[objective][col] >= -eps) {
            return true;
        }
        int row = -1;
        for (int i = 0; i < m; i++) {
            if (tab[i][col] <= eps) {
                continue;
            }
            double cur = tab[i][n + 1] / tab[i][col];
            double best = row == -1 ? 0 : tab[row][n + 1] / tab[row][col];
            if (row == -1 || cur < best - eps || (fabs(cur - best) <= eps && basis[i] < basis[row])) {
                row = i;
            }
        }
        if (row == -1) {
            return false;
        }
        pivot(row, col);
    }
};
int row = 0;
for (int i = 1; i < m; i++) {
    if (tab[i][n + 1] < tab[row][n + 1]) {
        row = i;
    }
}
if (m > 0 && tab[row][n + 1] < -eps) {
    pivot(row, n);
    if (!simplex(1) || fabs(tab[m + 1][n + 1]) > eps) {
        return -1;
    }
}
for (int i = 0; i < m; i++) {
    if (basis[i] != -1) {
        continue;
    }
}

```

```

    int col = -1;
    for (int j = 0; j <= n; j++) {
        if (fabs(tab[i][j]) > eps && (col == -1 || nonbasis[j] < nonbasis[col])) {
            col = j;
        }
    }
    if (col != -1) {
        pivot(i, col);
    }
}
}
if (!simplex(2)) {
    return 1;
}
x->assign(n, 0);
for (int i = 0; i < m; i++) {
    if (basis[i] < n) {
        (*x)[basis[i]] = tab[i][n + 1];
    }
}
return 0;
}

```

Example Usage

```

#include <iostream>
using namespace std;

int main() {
    // Solve [x, y] that maximizes 3x + 4y, subject to x, y >= 0 and:
    // -2x + 1y <= 0
    // 1x + 0.85y <= 9
    // 1x + 2y <= 14
    vector<vector<double>> va{{-2, 1}, {1, 0.85}, {1, 2}};
    vector<double> vb{0, 9, 14};
    vector<double> vc{3, 4};
    vector<double> x;
    assert(simplex_solve(va, vb, vc, &x) == 0);
    double maxval = 0;
    for (int i = 0; i < static_cast<int>(x.size()); i++) {
        maxval += vc[i] * x[i];
    }
    cout << "Solution = " << maxval << " at (" << x[0];
    for (int i = 1; i < static_cast<int>(x.size()); i++) {
        cout << ", " << x[i];
    }
    cout << ")." << endl;

    // Feasible even though the first RHS is negative: x >= 1 and x <= 3.
    vector<vector<double>> negative_rhs{{-1}, {1}};
    vector<double> negative_bounds{-1, 3};
    vector<double> one_var{1};
    assert(simplex_solve(negative_rhs, negative_bounds, one_var, &x) == 0);
    assert(fabs(x[0] - 3) < 1e-9);

    // Infeasible: x <= -1 contradicts x >= 0.
    vector<vector<double>> infeasible{{1}};
    vector<double> bad_bounds{-1};
    assert(simplex_solve(infeasible, bad_bounds, one_var, &x) == -1);

    // Unbounded: maximize x subject only to x >= 0.
    vector<vector<double>> unbounded{{-1}};
    vector<double> lower_bound{0};
    assert(simplex_solve(unbounded, lower_bound, one_var, &x) == 1);
    return 0;
}

```

Example Output

Solution = 33.3043 at (5.30435, 4.34783).

Chapter 6

Mathematics

6.1 Math Utilities

Common mathematic constants and functions, most of which already have standard STL equivalents. The implementations below are for educational purposes only and may not be as heavily optimized as their standard library counterparts.

Time Complexity:

- $O(1)$ per call to most operations.
- $O(1)$ per call to `mulmod()` with `__uint128_t`, or $O(\log n)$ with the portable fallback, where n is the second argument.
- $O(\log n)$ calls to `mulmod()` per call to `powmod(x, n, m)`.
- $O(d \cdot e)$ per call to `convert_base(d, a, b)`, where d is the number of input digits and e is the number of output digits.
- $O(\log_b(x + 1) + 1)$ per call to `to_base()`.
- $O(x/1000 + 1)$ per call to `to_roman()` due to the repeated `M` prefix.

Space Complexity:

- $O(1)$ auxiliary for most operations.
- $O(d + e)$ auxiliary for `convert_base(d, a, b)`, where d is the number of input digits and e is the number of output digits.
- $O(e)$ auxiliary for `to_base(x, b)`, where e is the number of output digits.
- $O(x/1000 + 1)$ auxiliary for `to_roman(x)`.

```
#include <algorithm>
#include <cassert>
#include <climits>
#include <cmath>
#include <cstdint>
#include <limits>
#include <random>
#include <string>
#include <type_traits>
#include <utility>
#include <vector>

#ifdef M_PI
const double M_PI = acos(-1.0); // or std::numbers::pi and std::numbers::pi_v<> in C++20 and later
#endif
#ifdef M_E
const double M_E = exp(1.0); // or std::numbers::e and std::numbers::e_v<> in C++20 and later
#endif
```

```
const double M_PHI = (1.0 + sqrt(5.0)) / 2.0;
const double M_INF = std::numeric_limits<double>::infinity();
const double M_NAN = std::numeric_limits<double>::quiet_NaN();
```

Epsilon Comparisons:

- `EQ()`, `NE()`, `LT()`, `GT()`, `LE()`, and `GE()` relationally compare two values x and y . Arguments may be of different types; the common type governs behavior. If the common type is a floating-point type, exactly equal values (including same-signed infinities) compare equal; otherwise absolute-error epsilon comparison is used. Values within `EPS` of each other are considered equal, and `LT`/`GT`/`LE`/`GE` shift the boundary by `EPS` accordingly. Otherwise exact comparison is used (`==`, `<`, etc.). The branch is selected with `if constexpr`, so exact types without floating-point arithmetic (e.g. integers, `Modular`, `Rational`) compose as well.
- `rEQ(ref, val)` returns whether `val` equals reference `ref` within relative error `EPS`. The tolerance scales with `|ref|`, so `rEQ(ref, val)` is NOT the same as `rEQ(val, ref)`. Use this when one argument is a known exact value and the other is a computed approximation. Degenerates to exact comparison when `ref` is 0, since the tolerance collapses to 0; use `EQ` near zero.
- `rEQ_sym(x, y)` is the symmetric (commutative) variant: tolerance scales with $\max(|x|, |y|)$, so the result is the same regardless of argument order. Still degenerates near zero when both arguments are close to 0.

```
const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U> bool NE(T a, U b) { return !EQ(a, b); }
template<typename T, typename U> bool GT(T a, U b) { return LT(b, a); }
template<typename T, typename U> bool LE(T a, U b) { return !LT(b, a); }
template<typename T, typename U> bool GE(T a, U b) { return !LT(a, b); }

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool rEQ(T ref, U val) {
    return C(ref) == C(val) || fabs(C(ref) - C(val)) <= EPS * fabs(C(ref));
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool rEQ_sym(T x, U y) {
    return C(x) == C(y) || fabs(C(x) - C(y)) <= EPS * std::max(fabs(C(x)), fabs(C(y)));
}
```

Sign Functions:

- `sgn(x)` returns -1 (if $x < 0$), 0 (if $x = 0$), or 1 (if $x > 0$). Unlike `std::signbit()` or `std::copysign()`, this does not handle the sign of `NaN`.
- `signbit_x` is analogous to `std::signbit()`, returning whether the sign bit of the floating point number is set to true. If so, then `x` is considered “negative.” Note that this works as expected on `+0.0`, `-0.0`, `Inf`, `-Inf`, `NaN`, as well as `-NaN`. Warning: This assumes that the sign bit is the leading (most significant) bit in the internal IEEE representation and that bytes are stored in little-endian order.
- `copysign(x, y)` is analogous to `std::copysign()`, returning a number with the magnitude of `x` but the sign of `y`.

```
template<typename T>
```



```

int sgn(const T &x) {
    return (T(0) < x) - (x < T(0));
}

template<typename Double>
bool signbit_(Double x) {
    return (((unsigned char *)&x)[sizeof(x) - 1] >> (CHAR_BIT - 1)) & 1;
}

template<typename Double>
Double copysign_(Double x, Double y) {
    return signbit_(y) ? -fabs(x) : fabs(x);
}

```

Rounding Functions:

- `floor0(x)` returns `x` rounded down, symmetrically towards zero. This function is analogous to `std::trunc()`.
- `ceil0(x)` returns `x` rounded up, symmetrically away from zero. This function is analogous to `std::round()`.
- `round_half_up(x)` returns `x` rounded half up, towards positive infinity.
- `round_half_down(x)` returns `x` rounded half down, towards negative infinity.
- `round_half_to0(x)` returns `x` rounded half down, symmetrically towards zero.
- `round_half_from0(x)` returns `x` rounded half up, symmetrically away from zero.
- `round_half_even(x, eps = 1e-9)` returns `x` rounded half to even, using `eps` to detect ties for banker's rounding.
- `round_half_alterate(x)` returns `x` rounded, where ties are broken by alternating rounds towards positive and negative infinity.
- `round_half_alterate0(x)` returns `x` rounded, where ties are broken by alternating symmetric rounds towards and away from zero.
- `round_half_random(x)` returns `x` rounded, where ties are broken randomly.
- `round_n_places(x, n, f)` returns `x` rounded to `n` digits after the decimal, using the specified rounding function `f(x)`.

```

template<typename Double>
Double floor0(const Double &x) {
    Double res = floor(fabs(x));
    return (x < 0.0) ? -res : res;
}

template<typename Double>
Double ceil0(const Double &x) {
    Double res = ceil(fabs(x));
    return (x < 0.0) ? -res : res;
}

template<typename Double>
Double round_half_up(const Double &x) {
    return floor(x + 0.5);
}

template<typename Double>
Double round_half_down(const Double &x) {
    return ceil(x - 0.5);
}

template<typename Double>
Double round_half_to0(const Double &x) {
    Double res = round_half_down(fabs(x));
    return (x < 0.0) ? -res : res;
}

template<typename Double>
Double round_half_from0(const Double &x) {
    Double res = round_half_up(fabs(x));
}

```

```

    return (x < 0.0) ? -res : res;
}

template<typename Double>
Double round_half_even(const Double &x, const Double &eps = 1e-9) {
    if (x < 0.0) {
        return -round_half_even(-x, eps);
    }
    Double ipart;
    modf(x, &ipart);
    if (fabs(x - (ipart + 0.5)) < eps) { // exactly halfway: break the tie towards even
        return (fmod(ipart, 2.0) < eps) ? ipart : ceil0(ipart + 0.5);
    }
    return round_half_from0(x);
}

template<typename Double>
Double round_half_alterate(const Double &x) {
    static bool up = true;
    return (up = !up) ? round_half_up(x) : round_half_down(x);
}

template<typename Double>
Double round_half_alterate0(const Double &x) {
    static bool up = true;
    return (up = !up) ? round_half_from0(x) : round_half_to0(x);
}

template<typename Double>
Double round_half_random(const Double &x) {
    static std::mt19937 rng(std::random_device{}());
    return (rng() % 2 == 0) ? round_half_from0(x) : round_half_to0(x);
}

template<typename Double, typename RoundFn>
Double round_n_places(const Double &x, unsigned int n, RoundFn f) {
    Double scale = pow(10, n);
    return f(x * scale) / scale;
}

```

Modular Arithmetic:

- `addmod(a, b, m)` and `submod(a, b, m)` respectively return addition and subtraction modulo `m`, each result in $[0, m)$. Both operands may be negative or unreduced; they are normalized before arithmetic to avoid signed overflow. The modulus `m` must be positive.
- `mulmod(x, n, m)` returns `x` multiplied by `n`, modulo `m`. This is done in a way to avoid overflow: on compilers with `__uint128_t` it uses one wide product, while the portable fallback uses double-and-add multiplication. The fallback is slower by a logarithmic factor, but avoids relying on nonstandard 128-bit integers. Unlike `addmod`/`submod`, this takes unsigned operands and supports a full 64-bit modulus.
- `powmod(x, n, m)` returns `x` raised to the power `n`, modulo `m`.

These are lightweight standalone helpers for the occasional modular computation. For a full value type that wraps a fixed modulus with overloaded operators, modular inverses, and combinatorics, see the `Modular` (“Mint”) class in section 6.3.2.

```

int64_t addmod(int64_t a, int64_t b, int64_t m) {
    assert(m > 0);
    a %= m;
    b %= m;
    if (a < 0) a += m;
    if (b < 0) b += m;
    return a >= m - b ? a - (m - b) : a + b;
}

```

```

int64_t submod(int64_t a, int64_t b, int64_t m) {
    assert(m > 0);
    a %= m;
    b %= m;
    if (a < 0) a += m;
    if (b < 0) b += m;
    return a >= b ? a - b : m - (b - a);
}

uint64_t mulmod(uint64_t x, uint64_t n, uint64_t m) {
    assert(m > 0);
    #if defined(__SIZEOF_INT128__)
    return static_cast<uint64_t>(static_cast<__uint128_t>(x) * n % m);
    #else
    uint64_t a = 0, b = x % m;
    for (; n > 0; n >>= 1) {
        if (n & 1) {
            a = a >= m - b ? a - (m - b) : a + b;
        }
        b = b >= m - b ? b - (m - b) : b + b;
    }
    return a;
    #endif
}

uint64_t powmod(uint64_t x, uint64_t n, uint64_t m) {
    assert(m > 0);
    uint64_t a = 1, b = x;
    for (; n > 0; n >>= 1) {
        if (n & 1) {
            a = mulmod(a, b, m);
        }
        b = mulmod(b, b, m);
    }
    return a % m;
}

```

Base Conversion:

- `to_base(x, b = 10)` returns the digits of the unsigned integer `x` in base `b`, where index 0 of the result stores the least significant digit.
- `to_roman(x)` returns the Roman numeral representation of the unsigned integer `x` as a string.
- `convert_base(d, a, b)` converts an integer in base `a` as a vector `d` of digits (where `d[0]` is the least significant digit) to base `b` as a vector of digits (again with index 0 holding the least significant digit). This uses repeated long division, so the value itself does not need to fit in a machine integer. Each intermediate `rem*a + digit` must fit in `uint64_t`.

```

std::vector<int> to_base(unsigned int x, int b = 10) {
    assert(b >= 2);
    std::vector<int> res;
    do { // do-while so that a value of 0 yields the single digit {0}
        res.push_back(x % b);
        x /= b;
    } while (x != 0);
    return res;
}

std::string to_roman(unsigned int x) {
    static const std::string h[] = {"", "C", "CC", "CCC", "CD", "D", "DC", "DCC", "DCCC", "CM"};
    static const std::string t[] = {"", "X", "XX", "XXX", "XL", "L", "LX", "LXX", "LXXX", "XC"};
    static const std::string o[] = {"", "I", "II", "III", "IV", "V", "VI", "VII", "VIII", "IX"};
    std::string prefix(x / 1000, 'M');
    x %= 1000;
    return prefix + h[x / 100] + t[x / 10 % 10] + o[x % 10];
}

```

```

std::vector<int> convert_base(const std::vector<int> &d, int a, int b) {
    assert(a >= 2 && b >= 2);
    std::vector<int> cur = d, res;
    auto trim = [](std::vector<int> &v) {
        while (v.size() > 1 && v.back() == 0) {
            v.pop_back();
        }
    };
    trim(cur);
    if (cur.empty() || (cur.size() == 1 && cur[0] == 0)) {
        return {0};
    }
    while (!(cur.size() == 1 && cur[0] == 0)) {
        std::vector<int> q(cur.size());
        uint64_t rem = 0;
        for (int i = static_cast<int>(cur.size()) - 1; i >= 0; i--) {
            uint64_t x = rem * static_cast<uint64_t>(a) + static_cast<uint64_t>(cur[i]);
            q[i] = static_cast<int>(x / static_cast<uint64_t>(b));
            rem = x % static_cast<uint64_t>(b);
        }
        res.push_back(static_cast<int>(rem));
        trim(q);
        cur = std::move(q);
    }
    return res;
}

```

Example Usage

```

using namespace std;

int main() {
    assert(EQ(M_PI, 3.14159265359));
    assert(EQ(M_INF, M_INF) && rEQ(M_INF, M_INF) && rEQ_sym(M_INF, M_INF));
    assert(EQ(M_E, 2.718281828459));
    assert(EQ(M_PHI, 1.61803398875));

    double x = -12345.6789;
    assert((-M_INF < x) && (x < M_INF));
    assert((M_INF + x == M_INF) && (M_INF - x == M_INF));
    assert((M_INF + M_INF == M_INF) && (-M_INF - M_INF == -M_INF));
    assert((M_NAN != x) && (M_NAN != M_INF) && (M_NAN != M_NAN));
    assert(!(M_NAN < x) && !(M_NAN > x) && !(M_NAN <= x) && !(M_NAN >= x));
    assert(isnan(0.0 * M_INF) && isnan(0.0 * -M_INF) && isnan(M_INF / -M_INF));
    assert(isnan(M_NAN) && isnan(-M_NAN) && isnan(M_INF - M_INF));

    assert(sgn(x) == -1 && sgn(0.0) == 0 && sgn(5678) == 1);
    assert(signbit(x) && !signbit_(0.0) && signbit_(-0.0));
    assert(!signbit_(M_INF) && signbit_(-M_INF));
    assert(!signbit_(M_NAN) && signbit_(-M_NAN));
    assert(copysign(1.0, +2.0) == +1.0 && copysign(M_INF, -2.0) == -M_INF);
    assert(copysign(1.0, -2.0) == -1.0 && std::signbit(copysign(M_NAN, -2.0)));

    assert(EQ(floor0(1.5), 1.0) && EQ(ceil0(1.5), 2.0));
    assert(EQ(floor0(-1.5), -1.0) && EQ(ceil0(-1.5), -2.0));
    assert(EQ(round_half_up(+1.5), +2) && EQ(round_half_down(+1.5), +1));
    assert(EQ(round_half_up(-1.5), -1) && EQ(round_half_down(-1.5), -2));
    assert(EQ(round_half_to0(+1.5), +1) && EQ(round_half_from0(+1.5), +2));
    assert(EQ(round_half_to0(-1.5), -1) && EQ(round_half_from0(-1.5), -2));
    assert(EQ(round_half_even(+1.5), +2) && EQ(round_half_even(-1.5), -2));
    assert(EQ(round_half_even(3.1), 3) && EQ(round_half_even(3.4), 3)); // non-ties round normally
    assert(NE(round_half_alterate(+1.5), round_half_alterate(+1.5)));
    assert(NE(round_half_alterate0(-1.5), round_half_alterate0(-1.5)));
    assert(EQ(round_n_places(-1.23456, 3, round_half_to0<double>), -1.235));
}

```

```

assert(addmod(7, 8, 10) == 5 && submod(2, 5, 10) == 7);
// Negative and unreduced operands are normalized into [0, m).
assert(addmod(-3, -4, 10) == 3 && submod(-3, 4, 10) == 3);
assert(addmod(25, -7, 10) == 8);
assert(addmod(INT64_MAX - 1, INT64_MAX - 1, INT64_MAX) == INT64_MAX - 2);

assert(powmod(2, 10, 1000000007) == 1024);
assert(powmod(2, 62, 1000000) == 387904);
assert(powmod(10001, 10001, 100000) == 10001);

assert(to_roman(1234) == "MCCXXXIV");
assert(to_roman(5678) == "MMMMDCLXXVIII");
std::vector<int> digits{6, 5, 4, 3, 2, 1};
std::vector<int> base20 = to_base(123456, 20);
assert(convert_base(base20, 20, 10) == digits);
assert(convert_base(std::vector<int>{0, 0, 0}, 10, 2) == std::vector<int>{0});

std::vector<int> big_decimal(30, 9); // 10^30 - 1, larger than uint64_t.
std::vector<int> big_binary = convert_base(big_decimal, 10, 2);
assert(convert_base(big_binary, 2, 10) == big_decimal);
return 0;
}

```

6.2 Combinatorics

6.2.1 Combinatorial Calculations

The following functions implement common operations in combinatorics. All size and count inputs must be nonnegative, and all moduli must be positive. All return values and table entries are computed as 64-bit integers modulo an input argument m or p . Modular products use ordinary `int64_t` multiplication, so the square of the chosen modulus must fit in `int64_t`.

- `factorial(n, m = MOD)` returns $n! \bmod m$.
- `factorial_without_p(n, p = MOD)` returns $n!$ modulo the prime p after removing every factor of p from the factorial.
- `binomial_table(n, m = MOD)` returns the first $n + 1$ rows of Pascal's triangle as a two-dimensional vector t such that $t[i][j] = \binom{i}{j} \bmod m$.
- `permute(n, k, m = MOD)` returns $(n \text{ permute } k) \bmod m$.
- `choose(n, k, p = MOD)` returns $\binom{n}{k} \bmod p$, where p is prime and $n < p$.
- `multichoose(n, k, p = MOD)` returns $(n \text{ multichoose } k) \bmod p$, where p is prime and $n + k - 1 < p$ when $n > 0$.
- `catalan(n, p = MOD)` returns the n -th Catalan number mod p , where p is prime and $2n < p$.
- `partitions(n, m = MOD)` returns the number of partitions of n , mod m .
- `partitions(n, k, m = MOD)` returns the number of partitions of n into k parts, mod m .
- `stirling1(n, k, m = MOD)` returns the (n, k) unsigned Stirling number of the 1st kind mod m .
- `stirling2(n, k, m = MOD)` returns the (n, k) Stirling number of the 2nd kind mod m .
- `eulerian1(n, k, m = MOD)` returns the (n, k) Eulerian number of the 1st kind mod m , where $n > k$.
- `eulerian2(n, k, m = MOD)` returns the (n, k) Eulerian number of the 2nd kind mod m , where $n > k$.

Time Complexity:

- $O(n)$ per call to `factorial()`.
- $O(p \log_p(n))$ per call to `factorial_without_p()`.
- $O(n^2)$ per call to `binomial_table()`.
- $O(k)$ per call to `permute()`.
- $O(\min(k, n - k))$ per call to `choose()`.
- $O(k)$ per call to `multichoose()`.
- $O(n)$ per call to `catalan()`.

- $O(n^2)$ per call to `partitions(n, m)`.
- $O(n \cdot k)$ per call to `partitions(n, k, m)`, `stirling1()`, `stirling2()`, `eulerian1()`, and `eulerian2()`.

Space Complexity:

- $O(n^2)$ auxiliary for `binomial_table(n, m)`.
- $O(n \cdot k)$ auxiliary for `partitions(n, k, m)`, `stirling1(n, k, m)`, `stirling2(n, k, m)`, `eulerian1(n, k, m)`, and `eulerian2(n, k, m)`.
- $O(1)$ auxiliary for all other operations.

```
#include <stdint>
#include <vector>

const int64_t MOD = 1000000007;

int64_t factorial(int n, int64_t m = MOD) {
    int64_t res = 1;
    for (int i = 2; i <= n; i++) {
        res = (res * i) % m;
    }
    return res % m;
}

int64_t factorial_without_p(int64_t n, int64_t p = MOD) {
    int64_t res = 1;
    while (n > 1) {
        if (n / p % 2 == 1) {
            res = res * (p - 1) % p;
        }
        int max = n % p;
        for (int i = 2; i <= max; i++) {
            res = (res * i) % p;
        }
        n /= p;
    }
    return res % p;
}

std::vector<std::vector<int64_t>> binomial_table(int n, int64_t m = MOD) {
    std::vector<std::vector<int64_t>> t(n + 1);
    for (int i = 0; i <= n; i++) {
        for (int j = 0; j <= i; j++) {
            if (i < 2 || j == 0 || i == j) {
                t[i].push_back(1);
            } else {
                t[i].push_back((t[i - 1][j - 1] + t[i - 1][j]) % m);
            }
        }
    }
    return t;
}

int64_t permute(int n, int k, int64_t m = MOD) {
    if (n < k) {
        return 0;
    }
    int64_t res = 1;
    for (int i = 0; i < k; i++) {
        res = res * (n - i) % m;
    }
    return res % m;
}

int64_t powmod(int64_t x, int64_t n, int64_t m) {
    int64_t a = 1, b = x % m;
    for (; n > 0; n >>= 1) {
        if (n & 1) {
```

```

    a = a * b % m;
}
b = b * b % m;
}
return a % m;
}

int64_t choose(int n, int k, int64_t p = MOD) {
    if (n < k) {
        return 0;
    }
    if (k > n - k) {
        k = n - k;
    }
    int64_t num = 1, den = 1;
    for (int i = 0; i < k; i++) {
        num = num * (n - i) % p;
    }
    for (int i = 1; i <= k; i++) {
        den = den * i % p;
    }
    return num * powmod(den, p - 2, p) % p;
}

int64_t multichoose(int n, int k, int64_t p = MOD) {
    if (n == 0) {
        return k == 0;
    }
    return choose(n + k - 1, k, p);
}

int64_t catalan(int n, int64_t p = MOD) {
    return choose(2 * n, n, p) * powmod(n + 1, p - 2, p) % p;
}

int64_t partitions(int n, int64_t m = MOD) {
    std::vector<int64_t> t(n + 1, 0);
    t[0] = 1;
    for (int i = 1; i <= n; i++) {
        for (int j = i; j <= n; j++) {
            t[j] = (t[j] + t[j - i]) % m;
        }
    }
    return t[n] % m;
}

int64_t partitions(int n, int k, int64_t m = MOD) {
    std::vector<std::vector<int64_t>> t(n + 1, std::vector<int64_t>(k + 1, 0));
    t[0][0] = 1;
    for (int i = 1; i <= n; i++) {
        for (int j = 1, h = k < i ? k : i; j <= h; j++) {
            t[i][j] = (t[i - 1][j - 1] + t[i - j][j]) % m;
        }
    }
    return t[n][k] % m;
}

int64_t stirling1(int n, int k, int64_t m = MOD) {
    std::vector<std::vector<int64_t>> t(n + 1, std::vector<int64_t>(k + 1, 0));
    t[0][0] = 1;
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= k; j++) {
            t[i][j] = (i - 1) * t[i - 1][j] % m;
            t[i][j] = (t[i][j] + t[i - 1][j - 1]) % m;
        }
    }
    return t[n][k] % m;
}

```

```

int64_t stirling2(int n, int k, int64_t m = MOD) {
    std::vector<std::vector<int64_t>> t(n + 1, std::vector<int64_t>(k + 1, 0));
    t[0][0] = 1;
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= k; j++) {
            t[i][j] = j * t[i - 1][j] % m;
            t[i][j] = (t[i][j] + t[i - 1][j - 1]) % m;
        }
    }
    return t[n][k] % m;
}

int64_t eulerian1(int n, int k, int64_t m = MOD) {
    if (k > n - 1 - k) {
        k = n - 1 - k;
    }
    std::vector<std::vector<int64_t>> t(n + 1, std::vector<int64_t>(k + 1, 1));
    for (int j = 1; j <= k; j++) {
        t[0][j] = 0;
    }
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= k; j++) {
            t[i][j] = (i - j) * t[i - 1][j - 1] % m;
            t[i][j] = (t[i][j] + ((j + 1) * t[i - 1][j] % m)) % m;
        }
    }
    return t[n][k] % m;
}

int64_t eulerian2(int n, int k, int64_t m = MOD) {
    std::vector<std::vector<int64_t>> t(n + 1, std::vector<int64_t>(k + 1, 1));
    for (int i = 1; i <= n; i++) {
        for (int j = 1; j <= k; j++) {
            if (i == j) {
                t[i][j] = 0;
            } else {
                t[i][j] = (j + 1) * t[i - 1][j] % m;
                t[i][j] = ((2 * i - 1 - j) * t[i - 1][j - 1] % m + t[i][j]) % m;
            }
        }
    }
    return t[n][k] % m;
}

```

Example Usage

```

#include <cassert>

int main() {
    auto t = binomial_table(10);
    for (int i = 0; i < static_cast<int>(t.size()); i++) {
        for (int j = 0; j < static_cast<int>(t[i].size()); j++) {
            assert(t[i][j] == choose(i, j));
        }
    }
    assert(factorial(10) == 3628800);
    assert(factorial_without_p(123456) == 639390503);
    assert(factorial_without_p(7, 5) == 3); // 7! / 5 = 1008 = 3 (mod 5).
    assert(permute(10, 4) == 5040);
    assert(choose(20, 7) == 77520);
    assert(multichoose(20, 7) == 657800);
    assert(catalan(10) == 16796);
    assert(partitions(4) == 5);
    assert(partitions(0, 0) == 1);
    assert(partitions(100, 5) == 38225);
}

```



```

assert(stirling1(4, 2) == 11);
assert(stirling2(4, 3) == 6);
assert(eulerian1(9, 5) == 88234);
assert(eulerian2(8, 3) == 195800);
return 0;
}

```

6.2.2 Enumerating Arrangements

For the purposes of this section, we define a “size k arrangement of n ” to be a permutation of a size k subset of the integers in $[0, n)$, for $0 \leq k \leq n$. There are n permute k possible arrangements, but n^k possible arrangements if repeated values are allowed.

For arrangements without repeats, each prefix fixes some used values and leaves a shrinking set of choices for the remaining positions. Ranking counts how many unused smaller choices could have been placed at the current position, multiplied by the number of possible suffix arrangements; unranking performs the inverse search by subtracting those block sizes. With repeats allowed, the sequence is just a base- n counter.

- `next_arrangement(n, a)` tries to rearrange `a` to the next lexicographically greater arrangement, returning true if such an arrangement exists or false if the array is already in descending order (in which case `a` is unchanged). The input `a` must consist of distinct integers in the range $[0, n)$.
- `arrangement_by_rank(n, k, r)` returns the size k arrangement of n which is lexicographically ranked r out of all size k arrangements of n , where r is a 0-based rank in the range $[0, n \text{ permute } k)$.
- `rank_by_arrangement(n, a)` returns an integer representing the 0-based rank of arrangement `a`, which must consist of distinct integers in the range $[0, n)$.
- `next_arrangement_with_repeats(n, a)` tries to rearrange `a` to the next lexicographically greater arrangement with repeats, returning true if such an arrangement exists or false if the array is already in descending order (in which case `a` is unchanged). The input `a` must consist of integers in the range $[0, n)$. If `a` were interpreted as a k digit integer in base n , this function could be thought of as incrementing the integer.

Time Complexity:

- $O(n \cdot k)$ per call to `next_arrangement()`, `arrangement_by_rank()`, and `rank_by_arrangement(n, a)`, where k is the size of `a`.
- $O(k)$ per call to `next_arrangement_with_repeats()`.

Space Complexity:

- $O(n)$ auxiliary for `next_arrangement()`, `arrangement_by_rank()`, and `rank_by_arrangement()`.
- $O(1)$ auxiliary for `next_arrangement_with_repeats()`.

```

#include <algorithm>
#include <cstdint>
#include <numeric>
#include <vector>

bool next_arrangement(int n, std::vector<int> &a) {
    int k = static_cast<int>(a.size());
    std::vector<char> used(n);
    for (int i = 0; i < k; i++) {
        used[a[i]] = true;
    }
    for (int i = k - 1; i >= 0; i--) {
        used[a[i]] = false;
        for (int j = a[i] + 1; j < n; j++) {
            if (!used[j]) {
                a[i++] = j;
                used[j] = true;
                for (int x = 0; x < k; x++) {
                    if (!used[x]) {
                        a[i++] = x;
                    }
                }
                return true;
            }
        }
    }
    return false;
}

```

```

    }
    }
    return true;
}
}
return false;
}

int64_t n_permute_k(int n, int k) {
    int64_t res = 1;
    for (int i = 0; i < k; i++) {
        res *= n - i;
    }
    return res;
}

std::vector<int> arrangement_by_rank(int n, int k, int64_t r) {
    std::vector<int> values(n), res(k);
    std::iota(values.begin(), values.end(), 0);
    for (int i = 0; i < k; i++) {
        int64_t count = n_permute_k(n - 1 - i, k - 1 - i);
        int pos = r / count;
        res[i] = values[pos];
        values.erase(values.begin() + pos);
        r %= count;
    }
    return res;
}

int64_t rank_by_arrangement(int n, const std::vector<int> &a) {
    int k = static_cast<int>(a.size());
    int64_t res = 0;
    std::vector<char> used(n);
    for (int i = 0; i < k; i++) {
        int count = 0;
        for (int j = 0; j < a[i]; j++) {
            if (!used[j]) {
                count++;
            }
        }
        res += count * n_permute_k(n - i - 1, k - i - 1);
        used[a[i]] = true;
    }
    return res;
}

bool next_arrangement_with_repeats(int n, std::vector<int> &a) {
    int k = static_cast<int>(a.size());
    for (int i = k - 1; i >= 0; i--) {
        if (a[i] < n - 1) {
            a[i]++;
            std::fill(a.begin() + i + 1, a.end(), 0);
            return true;
        }
    }
    return false;
}
}

```

Example Usage

```

#include <cassert>
#include <iostream>
using namespace std;

template<typename It>

```

```

void print_range(It lo, It hi) {
    cout << "{";
    for (; lo != hi; ++lo) {
        cout << *lo << (lo == hi - 1 ? "" : ",");
    }
    cout << "} ";
}

int main() {
    {
        int n = 4, k = 3;
        vector<int> a{0, 1, 2};
        cout << n << " permute " << k << " arrangements:" << endl;
        int count = 0;
        do {
            print_range(a.begin(), a.end());
            auto b = arrangement_by_rank(n, k, count);
            assert(a == b);
            assert(rank_by_arrangement(n, a) == count);
            count++;
            if (count == 10 || count == 20) {
                cout << endl;
            }
        } while (next_arrangement(n, a));
        assert(count == 24);
        cout << endl;
    }
    {
        int n = 4, k = 2;
        vector<int> a{0, 0};
        cout << endl << n << "^" << k << " arrangements with repeats:" << endl;
        int count = 0;
        do {
            print_range(a.begin(), a.end());
            count++;
            if (count == 13) {
                cout << endl;
            }
        } while (next_arrangement_with_repeats(n, a));
        assert(count == 16);
        cout << endl;
    }
    return 0;
}

```

Example Output

```

4 permute 3 arrangements:
{0,1,2} {0,1,3} {0,2,1} {0,2,3} {0,3,1} {0,3,2} {1,0,2} {1,0,3} {1,2,0} {1,2,3}
{1,3,0} {1,3,2} {2,0,1} {2,0,3} {2,1,0} {2,1,3} {2,3,0} {2,3,1} {3,0,1} {3,0,2}
{3,1,0} {3,1,2} {3,2,0} {3,2,1}

4^2 arrangements with repeats:
{0,0} {0,1} {0,2} {0,3} {1,0} {1,1} {1,2} {1,3} {2,0} {2,1} {2,2} {2,3} {3,0}
{3,1} {3,2} {3,3}

```

6.2.3 Enumerating Permutations

A permutation is an ordered list consisting of n (not necessarily distinct) elements.

The lexicographic successor is found by locating the rightmost ascent, increasing that position by the smallest possible amount, then reversing the suffix into its minimum order. Ranking and unranking permutations use the factorial number system: each position contributes the number of unused smaller values before the chosen one, scaled by the number of possible suffix permutations. The cycle decomposition views a permutation as a function on indices and follows each unvisited orbit until it returns to its start.

- `next_permutation2(lo, hi, comp = std::less<>())` is analogous to `std::next_permutation()`, taking two BidirectionalIterators as a range `[lo, hi)` for which the function tries to rearrange to the next lexicographically greater permutation according to `comp`. The function returns true if such a permutation exists, or false if the range is already in reverse comparator order, in which case it is rearranged into comparator order.
- `next_permutation(a, comp = std::less<>())` is analogous to `next_permutation2()`, except that it takes a vector instead of a range.
- `next_permutation_mask(x)` returns the next integer having the same number of 1-bits. Treating each 1-bit as whether to take its corresponding item generates combinations of a set of n items.
- `permutation_by_rank(n, r)` returns the permutation of the integers in the range `[0, n)` which is lexicographically ranked r , where r is a 0-based rank in the range `[0, n!)`.
- `rank_by_permutation(a)` returns an integer representing the 0-based rank of permutation `a`, which must be a permutation of the integers `[0, n)`.
- `permutation_cycles(a)` returns the orbits of the index mapping $i \mapsto a[i]$. For example, `{3, 1, 0, 2}` decomposes into cycles `{0, 3, 2}` and `{1}` because following the mapping from index 0 gives $0 \rightarrow 3 \rightarrow 2 \rightarrow 0$, while index 1 maps to itself.

Time Complexity:

- $O(n)$ per call to `next_permutation2()` and `next_permutation()`, where n is the input size.
- $O(n^2)$ per call to `permutation_by_rank()` and `rank_by_permutation()`.
- $O(1)$ per call to `next_permutation_mask()`.
- $O(n)$ per call to `permutation_cycles()`.

Space Complexity:

- $O(1)$ auxiliary for `next_permutation2()` and `next_permutation()`.
- $O(n)$ auxiliary for `permutation_by_rank()`, `rank_by_permutation()`, and `permutation_cycles()`.

```
#include <algorithm>
#include <stdint>
#include <functional>
#include <numeric>
#include <utility>
#include <vector>

template<typename It, typename Compare = std::less<>>
bool next_permutation2(It lo, It hi, Compare comp = Compare()) {
    if (lo == hi) {
        return false;
    }
    It i = lo;
    if (++i == hi) {
        return false;
    }
    i = hi;
    --i;
    while (true) {
        It j = i;
        if (comp(*--i, *j)) {
            It k = hi;
            while (!comp(*i, *--k)) {
            }
            std::iter_swap(i, k);
            std::reverse(j, hi);
            return true;
        }
    }
    if (i == lo) {
        std::reverse(lo, hi);
        return false;
    }
}
```

```

template<typename T, typename Compare = std::less<>>
bool next_permutation(std::vector<T> &a, Compare comp = Compare()) {
    int n = static_cast<int>(a.size());
    for (int i = n - 2; i >= 0; i--) {
        if (comp(a[i], a[i + 1])) {
            for (int j = n - 1; j--;) {
                if (comp(a[i], a[j])) {
                    std::swap(a[i++], a[j]);
                    for (j = n - 1; i < j; i++, j--) {
                        std::swap(a[i], a[j]);
                    }
                    return true;
                }
            }
        }
    }
    std::reverse(a.begin(), a.end());
    return false;
}

```

```

int64_t next_permutation_mask(int64_t x) {
    if (x == 0) {
        return 0;
    }
    int64_t s = x & -x, r = x + s;
    return r | ((x ^ r) >> 2) / s;
}

```

```

std::vector<int> permutation_by_rank(int n, int64_t r) {
    if (n == 0) {
        return {};
    }
    std::vector<int64_t> factorial(n);
    std::vector<int> values(n), res(n);
    factorial[0] = 1;
    for (int i = 1; i < n; i++) {
        factorial[i] = i * factorial[i - 1];
    }
    std::iota(values.begin(), values.end(), 0);
    for (int i = 0; i < n; i++) {
        int pos = r / factorial[n - 1 - i];
        res[i] = values[pos];
        values.erase(values.begin() + pos);
        r %= factorial[n - 1 - i];
    }
    return res;
}

```

```

int64_t rank_by_permutation(const std::vector<int> &a) {
    int n = static_cast<int>(a.size());
    if (n == 0) {
        return 0;
    }
    std::vector<int64_t> factorial(n);
    factorial[0] = 1;
    for (int i = 1; i < n; i++) {
        factorial[i] = i * factorial[i - 1];
    }
    int64_t res = 0;
    for (int i = 0; i < n; i++) {
        int v = a[i];
        for (int j = 0; j < i; j++) {
            if (a[j] < a[i]) {
                v--;
            }
        }
        res += v * factorial[n - 1 - i];
    }
}

```

```

    }
    return res;
}

using Cycles = std::vector<std::vector<int>>>;

Cycles permutation_cycles(const std::vector<int> &a) {
    int n = static_cast<int>(a.size());
    std::vector<char> visit(n);
    Cycles res;
    for (int i = 0; i < n; i++) {
        if (!visit[i]) {
            int j = i;
            std::vector<int> curr;
            do {
                curr.push_back(j);
                visit[j] = true;
                j = a[j];
            } while (j != i);
            res.push_back(std::move(curr));
        }
    }
    return res;
}

```

Example Usage

```

#include <bitset>
#include <cassert>
#include <iostream>
using namespace std;

template<typename It>
void print_range(It lo, It hi) {
    cout << "{";
    for (; lo != hi; ++lo) {
        cout << *lo << (lo == hi - 1 ? " " : ",");
    }
    cout << "} ";
}

int main() {
    {
        const int n = 4;
        vector<int> a{0, 1, 2, 3}, b = a, c = a;
        cout << "Permutations of [0, " << n << "):" << endl;
        int count = 0;
        do {
            print_range(a.begin(), a.end());
            assert(b == a);
            assert(c == a);
            vector<int> d = permutation_by_rank(n, count);
            assert(d == a);
            assert(rank_by_permutation(a) == count);
            count++;
            if (count == 8 || count == 16) {
                cout << endl;
            }
            std::next_permutation(b.begin(), b.end());
            next_permutation(c);
        } while (next_permutation(a));
        assert(count == 24);
        assert((a == vector<int>{0, 1, 2, 3}));
        cout << endl;
    }
    { // Permutations of binary digits.

```

```

const int n = 5;
cout << "\nPermutations of 2 zeros and 3 ones:" << endl;
int lo = bitset<5>(string("00111")).to_ulong();
int hi = bitset<6>(string("100011")).to_ulong();
int count = 0;
do {
    cout << bitset<n>(lo).to_string() << " ";
    count++;
} while ((lo = next_permutation_mask(lo)) != hi);
assert(count == 10);
cout << endl;
}
{
    vector<int> a{3, 2, 1}, b = a;
    assert(next_permutation(a, greater<int>()));
    assert(next_permutation2(b.begin(), b.end(), greater<int>()));
    assert((a == vector<int>{3, 1, 2} && b == a));
}
{ // Decomposition into cycles.
    vector<int> a{3, 1, 0, 2};
    cout << "\nDecomposition of {3,1,0,2} into cycles:" << endl;
    Cycles c = permutation_cycles(a);
    assert(c.size() == 2);
    assert((c[0] == vector<int>{0, 3, 2}));
    assert((c[1] == vector<int>{1}));
    for (const auto &cycle : c) {
        print_range(cycle.begin(), cycle.end());
    }
    cout << endl;
}
return 0;
}

```

Example Output

```

Permutations of [0, 4):
{0,1,2,3} {0,1,3,2} {0,2,1,3} {0,2,3,1} {0,3,1,2} {0,3,2,1} {1,0,2,3} {1,0,3,2}
{1,2,0,3} {1,2,3,0} {1,3,0,2} {1,3,2,0} {2,0,1,3} {2,0,3,1} {2,1,0,3} {2,1,3,0}
{2,3,0,1} {2,3,1,0} {3,0,1,2} {3,0,2,1} {3,1,0,2} {3,1,2,0} {3,2,0,1} {3,2,1,0}

Permutations of 2 zeros and 3 ones:
00111 01011 01101 01110 10011 10101 10110 11001 11010 11100

Decomposition of {3,1,0,2} into cycles:
{0,3,2} {1}

```

6.2.4 Enumerating Combinations

A combination is a subset of size k chosen from a total of n (not necessarily distinct) elements, where order does not matter.

The lexicographic successor advances the rightmost chosen position that can still move right, then packs every later position as far left as possible. Ranking and unranking use the combinatorial number system: at each position, count how many combinations would be skipped by choosing a smaller next value, then either add that count to the rank or subtract it while searching for the requested rank. Bitmask successors use the same order as increasing integers with a fixed popcount.

- `next_combination(lo, mid, hi, comp = std::less<>())` takes random-access iterators `lo`, `mid`, and `hi` as a range `[lo, hi)` of n elements for which the function will rearrange such that the k elements in `[lo, mid)` become the next lexicographically greater combination. The function returns true if such a combination exists, or false if `[lo, mid)` already consists of the lexicographically greatest combination of the elements in `[lo, hi)`, in which case the range is reset to its first combination. The comparator `comp` defines the element ordering.

- `next_combination(n, a)` rearranges `a` to become the next lexicographically greater combination of distinct integers in the range $[0, n)$. The vector `a` must be sorted and contain distinct integers in the range $[0, n)$.
- `next_combination_mask(x)` interprets the bits of an integer `x` as a mask with 1-bits specifying the chosen items for a combination and returns the mask of the next lexicographically greater combination (that is, the lowest integer greater than `x` with the same number of 1 bits). Note that this does not generate combinations in the same order as `next_combination()`, nor does it work if the corresponding n items are not distinct (in that case, duplicate combinations will be generated).
- `combination_by_rank(n, k, r)` returns the combination of k distinct integers in the range $[0, n)$ that is lexicographically ranked r , where r is a 0-based rank in the range $[0, \binom{n}{k})$.
- `rank_by_combination(n, a)` returns an integer representing the 0-based rank of combination `a`, which must contain sorted distinct integers in $[0, n)$.
- `next_combination_with_repeats(n, a)` rearranges `a` to become the next lexicographically greater combination of not necessarily distinct integers in the range $[0, n)$. The vector `a` must be sorted. Note that there is a total of n multichoose k combinations if repetition is allowed, where n multichoose $k = \binom{n+k-1}{k}$.

Time Complexity:

- $O(n)$ per call to `next_combination(lo, mid, hi)`, where n is the distance between `lo` and `hi`.
- $O(k)$ per call to `next_combination(n, a)` and `next_combination_with_repeats(n, a)`, where k is the size of `a`.
- $O(1)$ per call to `next_combination_mask()`.
- $O(n \cdot k)$ per call to `combination_by_rank()` and `rank_by_combination()`.

Space Complexity:

- $O(k)$ auxiliary for `combination_by_rank(n, k, r)`.
- $O(1)$ auxiliary for all other operations.

```
#include <algorithm>
#include <cstdint>
#include <functional>
#include <utility>
#include <vector>

template<typename It, typename Compare = std::less<>>
bool next_combination(It lo, It mid, It hi, Compare comp = Compare()) {
    if (lo == mid || mid == hi) {
        return false;
    }
    It l = mid - 1, h = hi - 1;
    int len1 = 1, len2 = 1;
    while (l != lo && !comp(*l, *h)) {
        --l;
        ++len1;
    }
    if (l == lo && !comp(*l, *h)) {
        std::rotate(lo, mid, hi);
        return false;
    }
    for (; mid < h; ++len2) {
        if (!comp(*l, *--h)) {
            ++h;
            break;
        }
    }
    if (len1 == 1 || len2 == 1) {
        std::iter_swap(l, h);
    } else if (len1 == len2) {
        std::swap_ranges(l, mid, h);
    } else {
        std::iter_swap(l++, h++);
        int total = (--len1) + (--len2), gcd = total;
        for (int i = len1; i != 0;) {
            std::swap(gcd %= i, i);
        }
    }
}
```



```

    int skip = total / gcd - 1;
    for (int i = 0; i < gcd; i++) {
        It curr = (i < len1) ? (1 + i) : (h + (i - len1));
        int k = i;
        auto prev = *curr;
        for (int j = 0; j < skip; j++) {
            k = (k + len1) % total;
            It next = (k < len1) ? (1 + k) : (h + (k - len1));
            *curr = *next;
            curr = next;
        }
        *curr = prev;
    }
}
return true;
}

bool next_combination(int n, std::vector<int> &a) {
    int k = static_cast<int>(a.size());
    for (int i = k - 1; i >= 0; i--) {
        if (a[i] < n - k + i) {
            a[i]++;
            while (++i < k) {
                a[i] = a[i - 1] + 1;
            }
            return true;
        }
    }
    return false;
}

int64_t next_combination_mask(int64_t x) {
    if (x == 0) {
        return 0;
    }
    int64_t s = x & -x, r = x + s;
    return r | (((x ^ r) >> 2) / s);
}

int64_t n_choose_k(int64_t n, int64_t k) {
    if (k > n - k) {
        k = n - k;
    }
    int64_t res = 1;
    for (int i = 0; i < k; i++) {
        res = res * (n - i) / (i + 1);
    }
    return res;
}

std::vector<int> combination_by_rank(int n, int k, int64_t r) {
    std::vector<int> res(k);
    int count = n;
    for (int i = 0; i < k; i++) {
        int j = 1;
        for (;;) {
            int64_t am = n_choose_k(count - j, k - 1 - i);
            if (r < am) {
                break;
            }
            r -= am;
        }
        res[i] = (i > 0) ? (res[i - 1] + j) : (j - 1);
        count -= j;
    }
    return res;
}

```

```

int64_t rank_by_combination(int n, const std::vector<int> &a) {
    int k = static_cast<int>(a.size());
    int64_t res = 0;
    int prev = -1;
    for (int i = 0; i < k; i++) {
        for (int j = prev + 1; j < a[i]; j++) {
            res += n_choose_k(n - 1 - j, k - 1 - i);
        }
        prev = a[i];
    }
    return res;
}

bool next_combination_with_repeats(int n, std::vector<int> &a) {
    int k = static_cast<int>(a.size());
    for (int i = k - 1; i >= 0; i--) {
        if (a[i] < n - 1) {
            for (++a[i]; ++i < k;) {
                a[i] = a[i - 1];
            }
            return true;
        }
    }
    return false;
}

```

Example Usage

```

#include <cassert>
#include <iostream>
using namespace std;

template<typename It>
void print_range(It lo, It hi) {
    cout << "{";
    for (; lo != hi; ++lo) {
        cout << *lo << (lo == hi - 1 ? "" : ",");
    }
    cout << "} ";
}

int main() {
    {
        int k = 3;
        string s = "11234";
        cout << "\n" << s << "\n choose " << k << ":" << endl;
        int count = 0;
        do {
            cout << s.substr(0, k) << " ";
            count++;
        } while (next_combination(s.begin(), s.begin() + k, s.end()));
        assert(count == 7); // Duplicate '1' values collapse equal output strings.
        cout << endl;
    }
    { // Unordered combinations using masks.
        int n = 5, k = 3;
        string char_set = "abcde"; // Must be distinct.
        cout << "\n" << char_set << "\n choose " << k << " with masks:" << endl;
        int64_t mask = (1LL << k) - 1, limit = 1LL << n;
        int count = 0;
        do {
            for (int i = 0; i < n; i++) {
                if ((mask >> i) & 1) {
                    cout << char_set[i];
                }
            }
        }
    }
}

```

```

    cout << " ";
    count++;
    mask = next_combination_mask(mask);
} while (mask < limit);
assert(count == 10);
cout << endl;
}
{
    string s = "4321";
    assert(next_combination(s.begin(), s.begin() + 2, s.end(), greater<char>()));
    assert(s.substr(0, 2) == "42");
}
{ // Combinations of distinct integers in [0, n).
    int n = 5, k = 3;
    vector<int> a{0, 1, 2};
    cout << endl << n << " choose " << k << ":" << endl;
    int count = 0;
    do {
        print_range(a.begin(), a.end());
        vector<int> b = combination_by_rank(n, k, count);
        assert(a == b);
        assert(rank_by_combination(n, a) == count);
        count++;
    } while (next_combination(n, a));
    assert(count == 10);
    cout << endl;
}
{ // Combinations with repeats.
    int n = 3, k = 2;
    vector<int> a{0, 0};
    cout << endl << n << " multichoose " << k << ":" << endl;
    int count = 0;
    do {
        print_range(a.begin(), a.end());
        count++;
    } while (next_combination_with_repeats(n, a));
    assert(count == 6);
    cout << endl;
}
return 0;
}

```

Example Output

```

"11234" choose 3:
112 113 114 123 124 134 234

"abcde" choose 3 with masks:
abc abd acd bcd abe ace bce ade bde cde

5 choose 3:
{0,1,2} {0,1,3} {0,1,4} {0,2,3} {0,2,4} {0,3,4} {1,2,3} {1,2,4} {1,3,4} {2,3,4}

3 multichoose 2:
{0,0} {0,1} {0,2} {1,1} {1,2} {2,2}

```

6.2.5 Enumerating Partitions

A partition of a natural number n is a way to write n as a sum of positive integers where the order of the addends does not matter.

The lexicographic successor edits the suffix of a non-increasing partition: increase the rightmost part that can grow without breaking the order, then refill the remaining sum with ones. Ranking and unranking use a dynamic-programming table that counts partitions of a remaining sum by their next part. This lets the algorithms skip whole blocks of partitions sharing a common prefix.

- `next_partition(p)` takes a reference to a vector `p` of positive integers as a partition of n for which the function will re-assign to become the next lexicographically greater partition. The function returns true if such a partition exists, or false if `p` already consists of the lexicographically greatest partition (i.e. the single integer n).
- `partition_by_rank(n, r)` returns the partition of n that is lexicographically ranked r if addends in each partition were sorted in non-increasing order, where r is a 0-based rank in the range $[0, p(n))$ where $p(n)$ is the number of partitions of n .
- `rank_by_partition(p)` returns an integer representing the 0-based rank of the partition specified by vector `p`, which must consist of positive integers sorted in non-increasing order.
- `generate_increasing_partitions(n, f)` calls the function `f(lo, hi)` on strictly increasing partitions of n in lexicographically increasing order of partition, where `lo` and `hi` are random-access iterators to a range $[lo, hi)$ of integers. Note that non-strictly increasing partitions like $\{1, 1, 1, 1\}$ are skipped.

Time Complexity:

- $O(n)$ per call to `next_partition()`.
- $O(n^2)$ per call to `partition_by_rank()` and `rank_by_partition()`.
- $O(p(n))$ per call to `generate_increasing_partitions()`, where $p(n)$ is the number of partitions of n .

Space Complexity:

- $O(1)$ auxiliary for `next_partition()`.
- $O(n^2)$ auxiliary heap space for `partition_function()`, `partition_by_rank()`, and `rank_by_partition()`.
- $O(n)$ auxiliary stack space for `generate_increasing_partitions()`.

```
#include <cstdint>
#include <vector>

bool next_partition(std::vector<int> &p) {
    int n = static_cast<int>(p.size());
    if (n <= 1) {
        return false;
    }
    int s = p[n - 1] - 1, i = n - 2;
    p.pop_back();
    for (; i > 0 && p[i] == p[i - 1]; i--) {
        s += p[i];
        p.pop_back();
    }
    for (p[i]++; s > 0; s--) {
        p.push_back(1);
    }
    return true;
}

int64_t partition_function(int a, int b) {
    static std::vector<std::vector<int64_t>> p(1, std::vector<int64_t>(1, 1));
    if (a >= static_cast<int>(p.size())) {
        int old = static_cast<int>(p.size());
        p.resize(a + 1);
        for (int i = 0; i < old; i++) {
            p[i].resize(a + 1);
        }
        for (int i = old; i <= a; i++) {
            p[i].resize(a + 1);
            for (int j = 1; j <= i; j++) {
                p[i][j] = p[i - 1][j - 1] + p[i - j][j];
            }
        }
    }
    return p[a][b];
}

std::vector<int> partition_by_rank(int n, int64_t r) {
    std::vector<int> res;
```

```

    for (int i = n, j; i > 0; i -= j) {
        for (j = 1;; j++) {
            int64_t count = partition_function(i, j);
            if (r < count) {
                break;
            }
            r -= count;
        }
        res.push_back(j);
    }
    return res;
}

int64_t rank_by_partition(const std::vector<int> &p) {
    int64_t res = 0;
    int sum = 0;
    for (int x : p) {
        sum += x;
    }
    for (int x : p) {
        for (int j = 0; j < x; j++) {
            res += partition_function(sum, j);
        }
        sum -= x;
    }
    return res;
}

template<typename Fn>
void generate_increasing_partitions(int left, int prev, int i, std::vector<int> &p, Fn f) {
    if (left == 0) {
        f(p.begin(), p.begin() + i);
        return;
    }
    for (p[i] = prev + 1; p[i] <= left; p[i]++) {
        generate_increasing_partitions(left - p[i], p[i], i + 1, p, f);
    }
}

template<typename Fn>
void generate_increasing_partitions(int n, Fn f) {
    std::vector<int> p(n, 0);
    generate_increasing_partitions(n, 0, 0, p, f);
}

```

Example Usage

```

#include <algorithm>
#include <cassert>
#include <iostream>
using namespace std;

template<typename It>
void print_range(It lo, It hi) {
    cout << "{";
    for (; lo != hi; ++lo) {
        cout << *lo << (lo == hi - 1 ? "" : "," );
    }
    cout << "} ";
}

int main() {
    // Exercise incremental growth of the cached dynamic-programming table.
    assert(partition_function(3, 1) == 1);
    assert(partition_function(8, 3) == 5);
    {

```

```

int n = 4;
vector<int> a(n, 1);
cout << "Partitions of " << n << ":" << endl;
int count = 0;
do {
    print_range(a.begin(), a.end());
    vector<int> b = partition_by_rank(n, count);
    assert(equal(a.begin(), a.end(), b.begin()));
    assert(rank_by_partition(a) == count);
    count++;
} while (next_partition(a));
assert(count == 5);
cout << endl;
}
{
    int n = 8;
    cout << "\nIncreasing partitions of " << n << ":" << endl;
    int count = 0;
    generate_increasing_partitions(n, [&](auto lo, auto hi) {
        print_range(lo, hi);
        count++;
    });
    assert(count == 6);
    cout << endl;
}
return 0;
}

```

Example Output

```

Partitions of 4:
{1,1,1,1} {2,1,1} {2,2} {3,1} {4}

Increasing partitions of 8:
{1,2,5} {1,3,4} {1,7} {2,6} {3,5} {8}

```

6.2.6 Enumerating Generic Combinatorial Sequences

Rank, unrank, and enumerate combinatorial sequences in lexicographic order using only a prefix counting oracle. This is useful when the objects are too numerous to precompute, but there is a simple dynamic-programming formula for the number of valid completions after a fixed prefix. Derived classes of `AbstractEnumerator` implement `count(prefix)`, which returns how many valid full sequences begin with that prefix.

Let A be the number of candidate values considered at each position (the base class `range`), L be the output sequence length (the base class `length`), and T be `total_count()`, the number of valid sequences. The base class tries candidate values in $[0, A)$ at each of the L positions, asking `count()` how many completions each prefix admits. The concrete classes may still use their own conventional constructor names, such as `n` and `k` for combinations.

- `to_rank(a)` returns an integer representing the 0-based rank of the combinatorial sequence `a`.
- `from_rank(r)` returns a combinatorial sequence of integers that is lexicographically ranked `r`, where `r` is a 0-based rank in the range $[0, \text{total_count}())$.
- `total_count()` returns the number of valid sequences.
- `enumerate(f)` calls the function `f(lo, hi)` on every specified combinatorial sequence in lexicographically increasing order, where `lo` and `hi` are two random-access iterators to a range $[\text{lo}, \text{hi})$ of integers.

Time Complexity:

- $O(A \cdot L)$ calls to `count()` per call to `to_rank()` and `from_rank()`.
- $O(T \cdot A \cdot L)$ calls to `count()` per call to `enumerate()`, excluding the cost of the callback `f`.
- $O(A^2)$ constructor time for `CombinationEnumerator` and $O(L^2)$ constructor time for `PartitionEnumerator` to build their tables.

- The actual running time multiplies these bounds by the subclass cost of `count()`. In the implementations below, `CombinationEnumerator::count()` is $O(1)$, while `ArrangementEnumerator::count()` and `PartitionEnumerator::count()` are $O(L)$.

Space Complexity:

- $O(L)$ auxiliary for `to_rank()`, `from_rank()`, and `enumerate()`, excluding output storage.
- $O(1)$ object storage for `ArrangementEnumerator` and `PermutationEnumerator`.
- $O(A^2)$ object storage for `CombinationEnumerator` and $O(L^2)$ object storage for `PartitionEnumerator` tables.

```
#include <stdint>
#include <vector>

class AbstractEnumerator {
protected:
    int range, length;

    AbstractEnumerator(int r, int l) : range(r), length(l) {}
    virtual int64_t count(const std::vector<int> &prefix) = 0;

public:
    virtual ~AbstractEnumerator() = default;
    int64_t total_count() { return count({}); }

    int64_t to_rank(const std::vector<int> &a) {
        int64_t res = 0;
        std::vector<int> prefix;
        for (int i = 0; i < static_cast<int>(a.size()); i++) {
            prefix.push_back(0);
            for (; prefix.back() < a[i]; prefix.back()++) {
                res += count(prefix);
            }
        }
        return res;
    }

    std::vector<int> from_rank(int64_t r) {
        std::vector<int> a;
        for (int i = 0; i < length; i++) {
            a.push_back(0);
            for (; a.back() < range; a.back()++) {
                int64_t curr = count(a);
                if (r < curr) {
                    break;
                }
                r -= curr;
            }
        }
        return a;
    }

    // Accepts any callable f(lo, hi), including capturing lambdas and functors.
    template<typename Fn>
    void enumerate(Fn f) {
        int64_t total = total_count();
        for (int64_t i = 0; i < total; i++) {
            std::vector<int> curr = from_rank(i);
            f(curr.begin(), curr.end());
        }
    }
};

class ArrangementEnumerator : public AbstractEnumerator {
public:
    ArrangementEnumerator(int n, int k) : AbstractEnumerator(n, k) {}

    int64_t count(const std::vector<int> &prefix) override {
```

```

    int n = static_cast<int>(prefix.size());
    for (int i = 0; i < n - 1; i++) {
        if (prefix[i] == prefix[n - 1]) {
            return 0;
        }
    }
    int64_t res = 1;
    for (int i = 0; i < length - n; i++) {
        res *= range - n - i;
    }
    return res;
}
};

class PermutationEnumerator : public ArrangementEnumerator {
public:
    PermutationEnumerator(int n) : ArrangementEnumerator(n, n) {}
};

class CombinationEnumerator : public AbstractEnumerator {
    std::vector<std::vector<int64_t>> table;

public:
    CombinationEnumerator(int n, int k)
        : AbstractEnumerator(n, k), table(n + 1, std::vector<int64_t>(n + 1)) {
        for (int i = 0; i <= n; i++) {
            for (int j = 0; j <= i; j++) {
                table[i][j] = (j == 0) ? 1 : table[i - 1][j - 1] + table[i - 1][j];
            }
        }
    }

    int64_t count(const std::vector<int> &prefix) override {
        int n = static_cast<int>(prefix.size());
        if (n >= 2 && prefix[n - 1] <= prefix[n - 2]) {
            return 0;
        }
        if (n == 0) {
            return table[range][length - n];
        }
        return table[range - prefix[n - 1] - 1][length - n];
    }
};

class PartitionEnumerator : public AbstractEnumerator {
    std::vector<std::vector<int64_t>> table;

public:
    PartitionEnumerator(int n)
        : AbstractEnumerator(n + 1, n), table(n + 1, std::vector<int64_t>(n + 1)) {
        std::vector<std::vector<int64_t>> tmp(table);
        tmp[0][0] = 1;
        for (int i = 1; i <= n; i++) {
            for (int j = 1; j <= i; j++) {
                tmp[i][j] = tmp[i - 1][j - 1] + tmp[i - j][j];
            }
        }
        for (int i = 1; i <= n; i++) {
            for (int j = 1; j <= n; j++) {
                table[i][j] = tmp[i][j] + table[i][j - 1];
            }
        }
    }

    int64_t count(const std::vector<int> &prefix) override {
        int n = static_cast<int>(prefix.size()), sum = 0;
        for (int x : prefix) {
            sum += x;
        }
    }
};

```



```

    }
    if (sum == range - 1) {
        return 1;
    }
    if (sum > range - 1 || (n > 0 && prefix[n - 1] == 0) ||
        (n >= 2 && prefix[n - 1] > prefix[n - 2])) {
        return 0;
    }
    if (n == 0) {
        return table[range - sum - 1][range - 1];
    }
    return table[range - sum - 1][prefix[n - 1]];
}
};

```

Example Usage

```

#include <cassert>
#include <iostream>
using namespace std;

template<typename It>
void print_range(It lo, It hi) {
    cout << "{";
    for (; lo != hi; ++lo) {
        cout << *lo << (lo == hi - 1 ? " " : ",");
    }
    cout << "} ";
}

int main() {
    {
        cout << "3 permute 2 arrangements:" << endl;
        ArrangementEnumerator arr(3, 2);
        int count = 0;
        arr.enumerate([&](auto lo, auto hi) {
            print_range(lo, hi);
            count++;
        });
        assert(count == 6);
        assert((arr.from_rank(5) == vector<int>{2, 1}));
        assert(arr.to_rank(vector<int>{2, 1}) == 5);
        cout << endl;
    }
    {
        cout << "\nPermutations of [0, 3):" << endl;
        PermutationEnumerator perm(3);
        int count = 0;
        perm.enumerate([&](auto lo, auto hi) {
            print_range(lo, hi);
            count++;
        });
        assert(count == 6);
        cout << endl;
    }
    {
        cout << "\n4 choose 3 combinations:" << endl;
        CombinationEnumerator comb(4, 3);
        int count = 0;
        comb.enumerate([&](auto lo, auto hi) {
            print_range(lo, hi);
            count++;
        });
        assert(count == 4);
        assert((comb.from_rank(3) == vector<int>{1, 2, 3}));
        assert(comb.to_rank(vector<int>{1, 2, 3}) == 3);
    }
}

```

```

    cout << endl;
}
{
    cout << "\nPartition of 4:" << endl;
    PartitionEnumerator part(4);
    int count = 0;
    part.enumerate([&](auto lo, auto hi) {
        print_range(lo, hi);
        count++;
    });
    assert(count == 5);
    cout << endl;
}
return 0;
}

```

Example Output

```

3 permute 2 arrangements:
{0,1} {0,2} {1,0} {1,2} {2,0} {2,1}

Permutations of [0, 3):
{0,1,2} {0,2,1} {1,0,2} {1,2,0} {2,0,1} {2,1,0}

4 choose 3 combinations:
{0,1,2} {0,1,3} {0,2,3} {1,2,3}

Partition of 4:
{1,1,1,1} {2,1,1,0} {2,2,0,0} {3,1,0,0} {4,0,0,0}

```

6.3 Number Theory

6.3.1 GCD, LCM, Mod Inverse, Chinese Remainder

Common number theory operations relating to modular arithmetic.

- `gcd(a, b)` returns the greatest common divisor of `a` and `b` using the Euclidean algorithm. This is mainly for educational purposes, as `std::gcd(a, b)` from `<numeric>` is available as of C++17 (`__gcd(a, b)` from `<algorithm>` in C++14 and earlier).
- `lcm(a, b)` returns the least common multiple of `a` and `b`. This implementation is mainly for educational purposes, as `std::lcm(a, b)` from `<numeric>` is available as of C++17.
- `extended_euclid(a, b)` returns a pair (x, y) of integers such that $\text{gcd}(a, b) = ax + by$.
- `diophantine(a, b, c, &x, &y, &g)` solves the linear Diophantine equation $ax + by = c$, returning whether a solution exists (one does if and only if $\text{gcd}(a, b)$ divides c), and on success sets `g` to $\text{gcd}(a, b)$ and `(x, y)` to a particular solution bounded by $\max(|a|, |b|, |c|)$ in magnitude. A 128-bit intermediate is used where available to keep the scaling step overflow-free.
- `mod(a, m)` returns the value of $a \bmod m$ under the true Euclidean definition of modulo, that is, the smallest nonnegative integer r satisfying $a + mn = r$ for some integer n . Note that this is identical to the remainder operator `%` in C++ for nonnegative operands `a` and `m`, but the result will differ when an operand is negative.
- `mod_inverse(a, m)` returns an integer x such that $ax \equiv 1 \pmod{m}$, where the arguments must satisfy $m > 0$ and $\text{gcd}(a, m) = 1$.
- `generate_inverse(p)` returns a vector `v` of integers where for each index i in the vector, $i \cdot v[i] \equiv 1 \pmod{p}$, where the argument p is prime.
- `crt(r1, m1, r2, m2, r, m)` merges the two congruences $x \equiv r_1 \pmod{m_1}$ and $x \equiv r_2 \pmod{m_2}$ for arbitrary moduli (not necessarily coprime). It returns whether the system is consistent, and on success sets `m` to `lcm(m1, m2)` and `r` to the unique solution in $[0, m)$. Both moduli must be positive, and their least common multiple must fit in `int64_t`. Fold it pairwise to merge more than two congruences.

- `garner_restore(a, p)` returns the smallest nonnegative solution x for the system of simultaneous congruences $x \equiv a[i] \pmod{p[i]}$ for all indices i in $[0, n)$, where n is `a.size()` and `p` consists of pairwise coprime integers (unlike `crt`, which allows shared factors). The exact solution is unique modulo the product of all moduli, so that product and the final answer must fit in `int64_t`.
- `garner_restore_mod(a, p, m)` returns that same CRT solution modulo `m`. This is the right variant when the product of the moduli is too large to fit in `int64_t`.

Time Complexity:

- $O(\log(a + b))$ per call to `gcd(a, b)`, `lcm(a, b)`, `extended_euclid(a, b)`, `diophantine(a, b, c, ...)`, `mod_inverse(a, m)`, and `crt(...)`.
- $O(1)$ per call to `mod()`.
- $O(p)$ per call to `generate_inverse()`.
- $O(n^2)$ per call to `garner_restore()` and `garner_restore_mod()`.

Space Complexity:

- $O(p)$ auxiliary for `generate_inverse()`.
- $O(n)$ auxiliary for `garner_restore()` and `garner_restore_mod()`.
- $O(1)$ auxiliary for all other operations.

```
#include <cstdint>
#include <type_traits>
#include <utility>
#include <vector>

template<typename Int>
Int gcd(Int a, Int b) {
    while (b != 0) {
        Int t = b;
        b = a % b;
        a = t;
    }
    return (a < 0 ? -a : a);
}

template<typename Int>
Int lcm(Int a, Int b) {
    if (a == 0 || b == 0) {
        return 0;
    }
    Int res = a / gcd(a, b) * b;
    return (res < 0 ? -res : res);
}

template<typename Int>
std::pair<Int, Int> extended_euclid(Int a, Int b) {
    Int x = 1, y = 0, x1 = 0, y1 = 1;
    while (b != 0) {
        Int q = a / b, prev_x1 = x1, prev_y1 = y1, prev_b = b;
        x1 = x - q * x1;
        y1 = y - q * y1;
        b = a - q * b;
        x = prev_x1;
        y = prev_y1;
        a = prev_b;
    }
    return (a > 0) ? std::pair<Int, Int>{x, y} : std::pair<Int, Int>{-x, -y};
}

template<typename Int>
bool diophantine(Int a, Int b, Int c, Int &x, Int &y, Int &g) {
    if (a == 0 && b == 0) {
        x = y = g = 0;
        return c == 0;
    }
}
```

```

if (a == 0) {
    g = (b < 0) ? -b : b;
    if (c % b != 0) {
        return false;
    }
    x = 0;
    y = c / b;
    return true;
}
if (b == 0) {
    g = (a < 0) ? -a : a;
    if (c % a != 0) {
        return false;
    }
    x = c / a;
    y = 0;
    return true;
}
g = gcd(a, b);
if (c % g != 0) {
    return false;
}
// Absorb the bulk of c into a*dx + b*dy, then scale the base solution by the small remainder so
// the reported (x, y) stay bounded by max(|a|, |b|, |c|). The scaled product is taken modulo b
// (resp. a) through a 128-bit intermediate where available, to avoid overflow.
std::pair<Int, Int> base = extended_euclid(a, b);
Int dx = c / a;
c -= dx * a;
Int dy = c / b;
c -= dy * b;
Int f = c / g;
#ifdef __SIZEOF_INT128__
    __extension__ typedef std::conditional_t<sizeof(Int) <= 4, int64_t, __int128> wide_t;
#else
    using wide_t = int64_t;
#endif
x = dx + static_cast<Int>(static_cast<wide_t>(base.first) * f % b);
y = dy + static_cast<Int>(static_cast<wide_t>(base.second) * f % a);
return true;
}

template<typename Int>
Int mod(Int a, Int m) {
    Int r = a % m;
    return (r >= 0) ? r : (r + m);
}

template<typename Int>
Int mod_inverse(Int a, Int m) {
    return mod(extended_euclid(a, m).first, m);
}

std::vector<int> generate_inverse(int p) {
    std::vector<int> res(p);
    res[1] = 1;
    for (int i = 2; i < p; i++) {
        res[i] = (p - (p / i) * res[p % i] % p) % p;
    }
    return res;
}

bool crt(int64_t r1, int64_t m1, int64_t r2, int64_t m2, int64_t &r, int64_t &m) {
    r1 = mod(r1, m1);
    r2 = mod(r2, m2);
    int64_t x, y, g;
    if (!diophantine(m1, -m2, r2 - r1, x, y, g)) {
        return false;
    }
}

```

```

    m = m1 / g * m2;
#if defined(__SIZEOF_INT128__)
    __extension__ typedef __int128 int128_t;
    r = static_cast<int64_t>((
        (static_cast<int128_t>(r1) + static_cast<int128_t>(m1) * mod(x, m2 / g)) % m
    ));
#else
    r = mod(r1 + m1 * mod(x, m2 / g), m); // Overflow warning without int128.
#endif
    return true;
}

std::vector<int64_t> garner_digits(const std::vector<int> &a, const std::vector<int> &p) {
    int n = static_cast<int>(a.size());
    std::vector<int64_t> x(a.begin(), a.end());
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < i; j++) {
            // Reduce mod p[i] each step; otherwise x[i] compounds to ~p^i and overflows int64_t.
            x[i] = mod_inverse(static_cast<int64_t>(p[j]), static_cast<int64_t>(p[i])) * (x[i] - x[j]);
            x[i] = (x[i] % p[i] + p[i]) % p[i];
        }
    }
    return x;
}

int64_t garner_restore(const std::vector<int> &a, const std::vector<int> &p) {
    int n = static_cast<int>(a.size());
    if (n == 0) {
        return 0;
    }
    std::vector<int64_t> x = garner_digits(a, p);
    int64_t res = x[0], m = 1;
    for (int i = 1; i < n; i++) {
        m *= p[i - 1];
        res += x[i] * m;
    }
    return res;
}

int64_t garner_restore_mod(const std::vector<int> &a, const std::vector<int> &p, int64_t m) {
    int n = static_cast<int>(a.size());
    if (n == 0) {
        return 0;
    }
    std::vector<int64_t> x = garner_digits(a, p);
    int64_t res = 0;
    for (int i = n - 1; i >= 0; i--) {
#if defined(__SIZEOF_INT128__)
        __extension__ typedef __int128 int128_t;
        res = static_cast<int64_t>((static_cast<int128_t>(res) * p[i] + x[i]) % m);
#else
        res = (res * p[i] + x[i]) % m; // Requires the intermediate product to fit without int128.
#endif
    }
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    {
        for (int a = -20; a <= 20; a++) {
            for (int b = -20; b <= 20; b++) {

```

```

    int g = gcd(a, b);
    auto [x, y] = extended_euclid(a, b);
    assert(g == a * x + b * y);
    if (g == 1 && b > 1) {
        int inv = mod_inverse(a, b);
        assert(mod(a * inv, b) == 1);
    }
}
}
}
{
    int p = 17;
    auto res = generate_inverse(p);
    for (int i = 0; i < p; i++) {
        if (i > 0) {
            assert(mod(i * res[i], p) == 1);
        }
    }
}
{
    vector<int> a{2, 3, 1}, m{3, 4, 5};
    int n = static_cast<int>(a.size());
    int x = garner_restore(a, m);
    assert(x == garner_restore_mod(a, m, 1000000007));
    for (int i = 0; i < n; i++) {
        assert(mod(x, m[i]) == a[i]);
    }
    assert(x == 11);
}
#ifdef __SIZEOF_INT128__
{ // Garner modulo another number works even when the product of CRT moduli is too large.
    vector<int> p{1000000007, 1000000009, 10000000033};
    __extension__ typedef __int128 int128_t;
    int128_t x = (static_cast<int128_t>(1) << 80) + 123456789;
    vector<int> a;
    for (int m : p) {
        a.push_back(static_cast<int>(x % m));
    }
    int64_t m = 998244353;
    assert(garner_restore_mod(a, p, m) == static_cast<int64_t>(x % m));
}
#endif
{ // Diophantine: exhaustively verify solvability and that solutions satisfy a*x + b*y == c.
    for (int a = -8; a <= 8; a++) {
        for (int b = -8; b <= 8; b++) {
            for (int c = -8; c <= 8; c++) {
                int x, y, g, gg = gcd(a, b);
                bool ok = diophantine(a, b, c, x, y, g);
                assert(ok == (gg == 0 ? c == 0 : c % gg == 0));
                if (ok) {
                    assert(a * x + b * y == c);
                }
            }
        }
    }
}
{ // CRT: merge two congruences, including non-coprime and inconsistent moduli.
    for (int m1 = 1; m1 <= 12; m1++) {
        for (int m2 = 1; m2 <= 12; m2++) {
            for (int r1 = 0; r1 < m1; r1++) {
                for (int r2 = 0; r2 < m2; r2++) {
                    int64_t r, m;
                    bool ok = crt(r1, m1, r2, m2, r, m);
                    int limit = lcm(m1, m2), want = -1;
                    for (int v = 0; v < limit && want < 0; v++) {
                        if (v % m1 == r1 && v % m2 == r2) {
                            want = v;
                        }
                    }
                }
            }
        }
    }
}

```

```

    }
    assert(ok == (want >= 0));
    if (ok) {
        assert(m == limit && r == want);
    }
}
}
}
}
}
return 0;
}

```

6.3.2 Modular Arithmetic

Wraps arithmetic modulo a compile-time constant in a small value type. This is a common contest helper for dynamic programming, combinatorics, polynomial operations, and any calculation where all answers are taken modulo some number p such as $10^9 + 7$.

The implementation is intentionally close to common contest “Mint” templates: normalization happens at construction, arithmetic operators are overloaded, mixed integer operations are supported through implicit construction, and combinations can be computed from lazy factorial tables. The operators use the hidden friend idiom: each `friend` defined inside the class is a non-template function, so mixed expressions like `2 + Mint(3)` convert the integer operand through the implicit constructor. (Free function template operators cannot do this - template argument deduction ignores implicit conversions - which is why some contest templates carry per-operator `Modular op U` and `U op Modular` overloads instead.)

The modulus is a positive signed `auto` template argument. For example, `Modular<1000000007>` and `Modular<998244353>` need no other annotation. Its literal type determines the storage type, so a 64-bit modulus (e.g. `Modular<(1LL << 61) - 1>`) is stored in 64 bits. The modulus must be a compile-time constant; a runtime-chosen modulus is not supported by this design.

Two integer types are at play: the type of the modulus stores the representative, while construction, multiplication, and inverses widen through an intermediate type chosen automatically from the storage width - `int64_t` for 32-bit storage, or `__int128` (a GCC/Clang extension, in line with this book’s use of `__builtin` functions) for 64-bit storage; on compilers without `__int128`, only 32-bit moduli are supported. The requirements are that $2p$ fits the storage type (for addition) and $(p-1)^2$ fits the intermediate type (for the widening multiply), so any modulus up to half the storage type’s range works out of the box; 64-bit moduli simply pay the cost of slower 128-bit multiplies. (Some contest templates instead estimate the multiply’s quotient in `long double` to avoid `__int128`; beware that this requires x86’s 80-bit `long double` and silently breaks for large moduli on ARM, where `long double` is only 64-bit.)

Division uses the extended Euclidean algorithm, so the divisor only needs to be coprime to the modulus. For the factorial-table combination helper, the usual contest assumption is that p is prime and the requested factorials are invertible modulo p .

- `Modular<MOD>(x = 0)` constructs the residue class of integer `x` modulo `MOD`.
- `value()` and `operator()()` return the stored representative in $[0, \text{MOD})$.
- Explicit casts to `int`, `long long`, `double`, and `long double` convert that stored representative to the corresponding primitive.
- `pow(n)` returns this value raised to nonnegative integer exponent `n`.
- `inv()` returns the multiplicative inverse, asserting it exists.
- Operators `+`, `-`, `*`, `/`, comparison, increment, decrement, and stream I/O are overloaded.
- `ModCombinatorics<Mint>::factorial(n)` returns $n!$ using a lazy factorial table.
- `ModCombinatorics<Mint>::choose(n, k)` returns $\binom{n}{k}$ using lazy factorial and inverse-factorial tables.
- `ModCombinatorics<Mint>::permute(n, k)` returns the number of ordered selections of `k` distinct elements from `n`.

- `ModCombinatorics<ModT>::multichoose(n, k)` returns the number of size- k multisets drawn from n types.

Time Complexity:

- $O(1)$ per addition, subtraction, multiplication, comparison, and stream output.
- $O(\log n)$ per call to `pow()`.
- $O(\log m)$ per call to `inv()` and division, where $m = \text{MOD}$.
- $O(n)$ total table growth to answer factorials and combinations up to size n .

Space Complexity:

- $O(1)$ auxiliary for Modular arithmetic.
- $O(n)$ for factorial and inverse-factorial tables grown through `choose(n, k)`.

```
#include <cassert>
#include <cstdint>
#include <iostream>
#include <limits>
#include <type_traits>
#include <utility>
#include <vector>

template<typename T>
T inverse(T a, T m) {
    T original_m = m;
    T u = 0, v = 1;
    while (a != 0) {
        T q = m / a;
        m -= q * a;
        std::swap(a, m);
        u -= q * v;
        std::swap(u, v);
    }
    assert(m == 1);
    u %= original_m;
    return u < 0 ? u + original_m : u;
}

template<auto MOD>
class Modular {
    using value_t = decltype(MOD);

    static_assert(std::is_integral_v<value_t> && std::is_signed_v<value_t>);
    static_assert(MOD > 0 && MOD <= std::numeric_limits<value_t>::max() / 2);

    // Widen through int64_t for 32-bit storage, or __int128 for 64-bit storage where it exists
    // (__extension__ silences -pedantic). Without __int128, only 32-bit moduli are supported.
#if defined(__SIZEOF_INT128__)
    __extension__ typedef std::conditional_t<sizeof(value_t) <= 4, int64_t, __int128> wide_t;
#else
    using wide_t = int64_t;
    static_assert(sizeof(value_t) <= 4, "64-bit moduli require __int128, which is unavailable");
#endif

    value_t v;

    static value_t normalize(wide_t x) {
        if (-static_cast<wide_t>(mod()) <= x && x < static_cast<wide_t>(mod())) {
            value_t y = static_cast<value_t>(x);
            return y < 0 ? y + mod() : y;
        }
        x %= mod();
        if (x < 0) {
            x += mod();
        }
        return static_cast<value_t>(x);
    }
};
```



```

public:
    Modular(wide_t x = 0) { v = normalize(x); }

    static value_t mod() { return MOD; }
    value_t value() const { return v; }
    value_t operator()() const { return v; }
    explicit operator int() const { return static_cast<int>(v); }
    explicit operator long long() const { return static_cast<long long>(v); }
    explicit operator double() const { return static_cast<double>(v); }
    explicit operator long double() const { return static_cast<long double>(v); }

    Modular pow(int64_t n) const {
        assert(n >= 0);
        Modular base = *this, res = 1;
        while (n > 0) {
            if (n & 1) {
                res *= base;
            }
            base *= base;
            n >>= 1;
        }
        return res;
    }

    Modular inv() const {
        assert(v != 0);
        return Modular(inverse(static_cast<wide_t>(v), static_cast<wide_t>(mod())));
    }

    Modular &operator+=(const Modular &other) {
        v += other.v;
        if (v >= mod()) {
            v -= mod();
        }
        return *this;
    }

    Modular &operator--(const Modular &other) {
        v -= other.v;
        if (v < 0) {
            v += mod();
        }
        return *this;
    }

    Modular &operator*=(const Modular &other) {
        v = normalize(static_cast<wide_t>(v) * other.v);
        return *this;
    }

    Modular operator-() const { return Modular(-v); }
    Modular &operator++() { return *this += 1; }
    Modular &operator--() { return *this -= 1; }
    Modular &operator/=(const Modular &other) { return *this *= other.inv(); }
    Modular operator++(int) { Modular res = *this; ++*this; return res; }
    Modular operator--(int) { Modular res = *this; --*this; return res; }
    friend Modular operator+(Modular a, const Modular &b) { return a += b; }
    friend Modular operator-(Modular a, const Modular &b) { return a -= b; }
    friend Modular operator*(Modular a, const Modular &b) { return a *= b; }
    friend Modular operator/(Modular a, const Modular &b) { return a /= b; }
    friend bool operator==(const Modular &a, const Modular &b) { return a.v == b.v; }
    friend bool operator!=(const Modular &a, const Modular &b) { return !(a == b); }
    friend bool operator<(const Modular &a, const Modular &b) { return a.v < b.v; }
    friend std::ostream &operator<<(std::ostream &os, const Modular &x) { return os << x.v; }

    friend std::istream &operator>>(std::istream &is, Modular &x) {
        int64_t value;

```

```

    is >> value;
    x = Modular(value);
    return is;
}
};

template<typename Mint>
class ModCombinatorics {
    std::vector<Mint> fact, inv_fact;

    void ensure(int n) {
        auto old = static_cast<int>(fact.size());
        if (old > n) {
            return;
        }
        fact.resize(n + 1);
        inv_fact.resize(n + 1);
        for (int i = old; i <= n; i++) {
            fact[i] = fact[i - 1] * i;
        }
        inv_fact[n] = fact[n].inv();
        for (int i = n; i > old; i--) {
            inv_fact[i - 1] = inv_fact[i] * i;
        }
    }

public:
    ModCombinatorics() : fact(1, Mint(1)), inv_fact(1, Mint(1)) {}

    Mint factorial(int n) {
        ensure(n);
        return fact[n];
    }

    Mint choose(int n, int k) {
        if (k < 0 || k > n) {
            return 0;
        }
        ensure(n);
        return fact[n] * inv_fact[k] * inv_fact[n - k];
    }

    Mint permute(int n, int k) {
        if (k < 0 || k > n) {
            return 0;
        }
        ensure(n);
        return fact[n] * inv_fact[n - k];
    }

    Mint multichoose(int n, int k) { return n == 0 ? Mint(k == 0) : choose(n + k - 1, k); }
};

using Mint = Modular<1000000007>;

```

Example Usage

```

#include <cassert>
#include <sstream>
using namespace std;

int main() {
    Mint a = 1000000008LL;
    Mint b = -2;
    assert(a.value() == 1);
    assert(b() == Mint::mod() - 2);
}

```

```

assert(static_cast<int>(a) == 1);
assert(static_cast<long long>(a) == 1LL);
assert(static_cast<double>(a) == 1.0);
assert(a + b == -1);
assert(2 + Mint(3) == 5);
assert(Mint(2) * 3 == 6);
assert(Mint(2).pow(10) == 1024);
assert(Mint(5) / 5 == 1);

Mint x = 0;
x++;
++x;
assert(x == 2);
stringstream ss("1000000008");
ss >> x;
assert(x == 1);

ModCombinatorics<Mint> comb;
assert(comb.factorial(5) == 120);
assert(comb.choose(5, 2) == 10);
assert(comb.permute(5, 2) == 20);
assert(comb.multichoose(3, 2) == 6);
assert(comb.multichoose(0, 0) == 1);

// Each distinct modulus is just another instantiation.
using Mint2 = Modular<998244353>;
assert(Mint2(2).pow(10) == 1024);
assert(Mint2(-1) == Mint2::mod() - 1);

#if defined(__SIZEOF_INT128__)
// A 64-bit modulus (deduced as long long) automatically widens through __int128.
using Mint3 = Modular<(1LL << 61) - 1>;
assert(Mint3(2).pow(62) == 2); // 2^62 mod (2^61 - 1)
assert(Mint3(3).inv() * 3 == 1);
#endif
return 0;
}

```

6.3.3 Primitive Roots

A primitive root modulo m is a number g whose powers generate every invertible residue modulo m . Equivalently, the multiplicative order of g modulo m is exactly $\phi(m)$. Primitive roots are useful for constructing generators of finite multiplicative groups, choosing roots for number theoretic transforms, discrete logarithm setups, and randomized modular hashing or testing schemes that need a full-period multiplicative generator.

Primitive roots do not exist for every m . They exist exactly for $m \in \{1, 2, 4, p^k, 2p^k\}$, where p is an odd prime and $k \geq 1$. This implementation first checks that classification by factoring m , then factors $\phi(m)$. A candidate g is primitive exactly when $g^{\phi(m)} \equiv 1 \pmod{m}$ and $g^{\phi(m)/q} \not\equiv 1 \pmod{m}$ for every prime factor q of $\phi(m)$.

For a prime modulus p , this reduces to the common contest test: factor $p - 1$, then scan for the first g such that $g^{(p-1)/q} \not\equiv 1 \pmod{p}$ for every prime factor q of $p - 1$.

- `primitive_root(m)` returns the smallest primitive root modulo m , or -1 if no primitive root exists.

Modular multiplication uses `__uint128_t` when available for speed, with a slower portable double-and-add fallback to avoid overflow on compilers without 128-bit integers.

Time Complexity: $O(\sqrt{m} + g \cdot k \cdot \log m)$ per call, where g is the answer candidate tested and k is the number of distinct prime factors of $\phi(m)$.

Space Complexity: $O(k)$ auxiliary.

```
#include <cstdlib>
```

```

#include <numeric>
#include <set>
#include <utility>
#include <vector>

int64_t mulmod(int64_t a, int64_t b, int64_t m) {
    #if defined(__SIZEOF_INT128__)
        return static_cast<int64_t>(static_cast<__uint128_t>(a) * b % m);
    #else
        int64_t res = 0;
        for (a %= m; b > 0; b >>= 1) {
            if (b & 1) {
                res = res >= m - a ? res - (m - a) : res + a;
            }
            a = a >= m - a ? a - (m - a) : a + a;
        }
        return res;
    #endif
}

int64_t powmod(int64_t a, int64_t n, int64_t m) {
    int64_t res = 1 % m;
    for (a %= m; n > 0; n >>= 1) {
        if (n & 1) {
            res = mulmod(res, a, m);
        }
        a = mulmod(a, a, m);
    }
    return res;
}

std::vector<std::pair<int64_t, int>> factorize(int64_t n) {
    std::vector<std::pair<int64_t, int>> res;
    for (int64_t p = 2; p <= n / p; p += (p == 2 ? 1 : 2)) {
        if (n % p != 0) {
            continue;
        }
        int count = 0;
        while (n % p == 0) {
            n /= p;
            count++;
        }
        res.emplace_back(p, count);
    }
    if (n > 1) {
        res.emplace_back(n, 1);
    }
    return res;
}

int64_t primitive_root(int64_t m) {
    if (m <= 0) {
        return -1;
    }
    if (m == 1 || m == 2 || m == 4) {
        return m - 1;
    }
    std::vector<std::pair<int64_t, int>> factors = factorize(m);
    bool has_root = (factors.size() == 1 && factors[0].first != 2) ||
        (factors.size() == 2 && factors[0].first == 2 && factors[0].second == 1);
    if (!has_root) {
        return -1;
    }
    int64_t phi = m;
    std::set<int64_t> phi_factors;
    for (auto [p, e] : factors) {
        phi = phi / p * (p - 1);
        if (e > 1) {

```

```

    phi_factors.insert(p);
}
for (auto [q, count] : factorize(p - 1)) {
    phi_factors.insert(q);
}
}
for (int64_t g = 1; g < m; g++) {
    if (std::gcd(g, m) != 1 || powmod(g, phi, m) != 1) {
        continue;
    }
    bool ok = true;
    for (int64_t q : phi_factors) {
        if (powmod(g, phi / q, m) == 1) {
            ok = false;
            break;
        }
    }
    if (ok) {
        return g;
    }
}
return -1;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    assert(primitive_root(1) == 0);
    assert(primitive_root(2) == 1);
    assert(primitive_root(4) == 3);

    assert(primitive_root(17) == 3);           // 3 has order phi(17) = 16.
    assert(primitive_root(998244353) == 3);    // Common NTT prime.
    assert(primitive_root(9) == 2);            // Odd prime power.
    assert(primitive_root(18) == 5);           // Twice an odd prime power.

    assert(primitive_root(8) == -1);
    assert(primitive_root(12) == -1);
    return 0;
}

```

6.3.4 Discrete Logarithm (Baby-Step Giant-Step)

Given integers a , b , and a modulus m , the discrete logarithm problem asks for the smallest nonnegative integer x with a^x congruent to b modulo m . This is the modular analogue of a logarithm and underlies Diffie-Hellman key exchange, order computations in the multiplicative group, and many number-theory reductions.

The baby-step giant-step method writes $x = np - q$ for a block size n near the square root of the modulus. It tabulates the “baby steps” $b \cdot a^q$ for q in $[0, n]$, then walks the “giant steps” a^{np} for increasing p until one lands in the table, yielding $x = np - q$. This trades the linear scan of all exponents for $O(\sqrt{m})$ work and memory. A short reduction loop first strips common factors between a and m , so the routine works for any modulus, not only when a and m are coprime.

- `discrete_log(a, b, m)` returns the smallest nonnegative `x` with `ax` congruent to `b` modulo `m`, or `-1` if no such `x` exists. Requires `m` ≥ 1 ; `a` and `b` are reduced modulo `m` internally.

Time Complexity: $O(\sqrt{m})$ expected per call with constant-time modular multiplication, and $O(\sqrt{m} \cdot \log(m))$ with the portable double-and-add fallback.

Space Complexity: $O(\sqrt{m})$ auxiliary for the table.

```

#include <cmath>
#include <stdint>
#include <numeric>
#include <unordered_map>

int64_t mulmod(int64_t a, int64_t b, int64_t m) {
    #if defined(__SIZEOF_INT128__)
        return static_cast<int64_t>(static_cast<__uint128_t>(a) * b % m);
    #else
        int64_t res = 0;
        for (a %= m; b > 0; b >>= 1) {
            if (b & 1) {
                res = res >= m - a ? res - (m - a) : res + a;
            }
            a = a >= m - a ? a - (m - a) : a + a;
        }
        return res;
    #endif
}

int64_t discrete_log(int64_t a, int64_t b, int64_t m) {
    a %= m;
    b %= m;
    if (a < 0) {
        a += m;
    }
    if (b < 0) {
        b += m;
    }
    // Strip common factors of a and m so that a becomes invertible; add counts the exponents removed.
    int64_t k = 1 % m, add = 0, g;
    while ((g = std::gcd(a, m)) > 1) {
        if (b == k) {
            return add;
        }
        if (b % g != 0) {
            return -1;
        }
        b /= g;
        m /= g;
        add++;
        k = mulmod(k, a / g, m);
    }
    int64_t n = static_cast<int64_t>(std::sqrt(static_cast<double>(m))) + 1;
    int64_t an = 1;
    for (int64_t i = 0; i < n; i++) {
        an = mulmod(an, a, m); // an = a^n, the giant step.
    }
    std::unordered_map<int64_t, int64_t> baby;
    int64_t cur = b;
    for (int64_t q = 0; q <= n; q++) {
        baby[cur] = q; // Keep the largest q for each value to make the answer minimal.
        cur = mulmod(cur, a, m);
    }
    cur = k;
    for (int64_t p = 1; p <= n; p++) {
        cur = mulmod(cur, an, m);
        auto it = baby.find(cur);
        if (it != baby.end()) {
            return n * p - it->second + add;
        }
    }
    return -1;
}

```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    // 2^x = 9 (mod 11): 2^6 = 64 = 9, and 6 is the smallest such exponent.
    assert(discrete_log(2, 9, 11) == 6);

    assert(discrete_log(3, 1, 7) == 0);    // a^0 = 1 always.
    assert(discrete_log(2, 3, 5) == 3);    // 2^3 = 8 = 3 (mod 5).
    assert(discrete_log(2, 1, 6) == 0);    // Works for composite, non-coprime moduli.
    assert(discrete_log(4, 6, 8) == -1);    // No power of 4 is 6 (mod 8).
    assert(discrete_log(5, 33, 58) == 9);    // 5^9 = 33 (mod 58).
    return 0;
}
```

6.3.5 Modular Square Root (Tonelli-Shanks)

Given an integer a and an odd prime p , find a modular square root: an x with x^2 congruent to a modulo p . Exactly half of the nonzero residues modulo a prime are squares (quadratic residues), so a solution exists only for those; when it does, the two roots are x and $p - x$.

The Tonelli-Shanks algorithm handles the general case. Euler's criterion, $a^{(p-1)/2}$ modulo p , first tests whether a is a residue at all. When p is congruent to 3 (mod 4), the root is simply $a^{(p+1)/4}$. Otherwise the algorithm factors $p - 1$ as $s \cdot 2^e$, picks any quadratic non-residue to generate the 2-power part of the group, and repeatedly squares a running value to find the order of the current error term, correcting the candidate root until the error becomes 1.

- `mod_sqrt(a, p)` returns the smaller of the two square roots of `a` modulo the odd prime `p`, or `-1` if `a` is not a quadratic residue. The value 0 maps to 0. The argument `a` is reduced modulo `p` internally. The modulus `p` must be an odd prime.

Time Complexity: $O(\log^2 p + z \log p)$ per call, where z is the smallest integer at least 2 that is a quadratic non-residue modulo p ; thus the unconditional worst-case bound is $O(p \log p)$.

Space Complexity: $O(1)$ auxiliary.

```
#include <cstdint>

int64_t mulmod(int64_t a, int64_t b, int64_t m) {
    #if defined(__SIZEOF_INT128__)
        return static_cast<int64_t>(static_cast<__uint128_t>(a) * b % m);
    #else
        int64_t res = 0;
        for (a %= m; b > 0; b >>= 1) {
            if (b & 1) {
                res = res >= m - a ? res - (m - a) : res + a;
            }
            a = a >= m - a ? a - (m - a) : a + a;
        }
        return res;
    #endif
}

int64_t powmod(int64_t a, int64_t n, int64_t m) {
    int64_t res = 1 % m;
    for (a %= m; n > 0; n >>= 1) {
        if (n & 1) {
            res = mulmod(res, a, m);
        }
        a = mulmod(a, a, m);
    }
    return res;
}
```

```

}

int64_t mod_sqrt(int64_t a, int64_t p) {
    a %= p;
    if (a < 0) {
        a += p;
    }
    if (a == 0) {
        return 0;
    }
    if (powmod(a, (p - 1) / 2, p) != 1) {
        return -1; // a is a quadratic non-residue.
    }
    if (p % 4 == 3) {
        int64_t x = powmod(a, p / 4 + 1, p);
        return x < p - x ? x : p - x;
    }
    // Write p - 1 = s * 2^e with s odd.
    int64_t s = p - 1, e = 0;
    while (s % 2 == 0) {
        s /= 2;
        e++;
    }
    // Any quadratic non-residue generates the 2-power part of the group.
    int64_t z = 2;
    while (powmod(z, (p - 1) / 2, p) != p - 1) {
        z++;
    }
    int64_t x = powmod(a, (s + 1) / 2, p); // Candidate root, off by a 2-power factor.
    int64_t b = powmod(a, s, p);         // Error term, a power of two in order.
    int64_t g = powmod(z, s, p);         // Generator of the relevant 2-group.
    int64_t r = e;
    while (true) {
        int64_t t = b, order = 0;
        while (order < r && t != 1) {
            t = mulmod(t, t, p);
            order++;
        }
        if (order == 0) {
            return x < p - x ? x : p - x;
        }
        int64_t gs = powmod(g, static_cast<int64_t>(1) << (r - order - 1), p);
        g = mulmod(gs, gs, p);
        x = mulmod(x, gs, p);
        b = mulmod(b, g, p);
        r = order;
    }
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    int64_t r = mod_sqrt(2, 113);
    assert(r != -1 && mulmod(r, r, 113) == 2);
    assert(r == 51); // 51^2 = 2601 = 2 (mod 113); the other root is 62.

    assert(mod_sqrt(0, 7) == 0);
    assert(mod_sqrt(5, 41) != -1); // p = 1 (mod 4) exercises the full algorithm.
    assert(mod_sqrt(4, 7) == 2);    // p = 3 (mod 4) fast path.
    assert(mod_sqrt(3, 7) == -1);   // 3 is a non-residue modulo 7.

    // Every residue's reported root squares back to it.
    for (int64_t a = 0; a < 41; a++) {

```



```

    int64_t x = mod_sqrt(a, 41);
    if (x != -1) {
        assert(mulmod(x, x, 41) == a);
    }
}
return 0;
}

```

6.3.6 Jacobi Symbol

The Jacobi symbol (a/n) , defined for an odd positive integer n , generalizes the Legendre symbol. For an odd prime p the Legendre symbol (a/p) is 0 when p divides a , 1 when a is a nonzero quadratic residue modulo p , and -1 when a is a non-residue. The Jacobi symbol extends this to composite n by multiplying the Legendre symbols of the prime factors of n , counted with multiplicity. It is computed without factoring n , using the law of quadratic reciprocity together with the supplementary rules for -1 and 2 , in the style of the binary GCD.

The symbol is fully multiplicative in a and in n . For prime n it exactly decides quadratic residuosity, so it is the fast way to test whether a modular square root exists. For composite n , beware that $(a/n) = 1$ does not guarantee that a is a quadratic residue; this asymmetry is what the Solovay-Strassen primality test exploits.

- `jacobi(a, n)` returns the Jacobi symbol (a / n) as one of -1 , 0 , or 1 . The modulus `n` must be a positive odd integer; `a` is reduced modulo `n` internally and may be negative. The result is 0 exactly when `a` and `n` share a common factor.

Time Complexity: $O(\log n)$ arithmetic operations per call.

Space Complexity: $O(1)$ auxiliary.

```

#include <cassert>
#include <cstdint>
#include <utility>

int jacobi(int64_t a, int64_t n) {
    assert(n > 0 && n % 2 == 1);
    a %= n;
    if (a < 0) {
        a += n;
    }
    int result = 1;
    while (a != 0) {
        while (a % 2 == 0) {
            a /= 2;
            // (2 / n) is -1 when n is 3 or 5 modulo 8.
            if (n % 8 == 3 || n % 8 == 5) {
                result = -result;
            }
        }
        std::swap(a, n);
        // Quadratic reciprocity flips the sign when both are 3 modulo 4.
        if (a % 4 == 3 && n % 4 == 3) {
            result = -result;
        }
        a %= n;
    }
    return n == 1 ? result : 0;
}

```

Example Usage

```
using namespace std;
```

```

int64_t legendre_via_euler(int64_t a, int64_t p) {
    int64_t r = 1 % p;
    for (int64_t b = ((a % p) + p) % p, e = (p - 1) / 2; e > 0; e >= 1) {
        if (e & 1) {
            r = r * b % p;
        }
        b = b * b % p;
    }
    return r == p - 1 ? -1 : r; // r is 0, 1, or p - 1.
}

int main() {
    assert(jacobi(2, 15) == 1); // (2/15) = (2/3)(2/5) = (-1)(-1).
    assert(jacobi(7, 15) == -1);
    assert(jacobi(3, 15) == 0); // gcd(3, 15) = 3.
    assert(jacobi(1001, 9907) == -1);

    // For odd primes the Jacobi symbol equals the Legendre symbol from Euler's criterion.
    for (int64_t p : {3, 5, 7, 11, 13, 101, 9973}) {
        for (int64_t a = 0; a < 50; a++) {
            assert(jacobi(a, p) == legendre_via_euler(a, p));
        }
    }
    return 0;
}

```

6.3.7 Prime Generation

Generate prime numbers using the Sieve of Eratosthenes, which repeatedly marks the multiples of each prime as composite.

The range overload `sieve(lo, hi)` uses a segmented sieve, letting primes in a high but narrow window be found without sieving everything below it. Since every composite at most `hi` has a prime factor at most $\sqrt{hi}$, it first sieves the primes up to $\sqrt{hi}$, then uses only those primes to mark composites within $[lo, hi]$ (each prime p starting at the first of its multiples that is at least both p^2 and `lo`). This needs only $O(w + \sqrt{h})$ space, where $w = hi - lo + 1$ and $h = hi$, instead of the $O(h)$ space that sieving all of $[2, hi]$ would require.

The linear sieve finds the least prime factor of each integer. When processing `i`, it marks `i * p` for each known prime `p` no greater than `least[i]`. Every composite `x` is therefore generated exactly once: writing `x = i * p`, where `p` is its least prime factor, gives `p ≤ least[i]`, while any representation using a larger prime is skipped. This single visit per composite gives the $O(n)$ running time.

- `sieve(n)` returns a vector of all the primes less than or equal to `n`.
- `sieve(lo, hi)` returns a vector of all the primes in the range $[lo, hi]$.
- `linear_sieve(n, least_out = nullptr)` returns a vector of all primes less than or equal to `n` in linear time. If `least_out` is not null, it is filled so that `(*least_out)[x]` is the least prime factor of `x` for every `x` ≥ 2 .

Time Complexity:

- $O(n \log \log n)$ per call to `sieve(n)`.
- $O(\sqrt{hi} \cdot \log(\log(hi)) + w \cdot \log(\log(hi)))$ per call to `sieve(lo, hi)`, where $w = hi - lo + 1$.
- $O(n)$ per call to `linear_sieve()`.

Space Complexity:

- $O(n)$ auxiliary for `sieve(n)`.
- $O(w + \sqrt{hi})$ auxiliary for `sieve(lo, hi)`.
- $O(n)$ auxiliary for `linear_sieve()`.

```

#include <algorithm>
#include <cmath>
#include <cstdint>

```

```

#include <vector>

std::vector<int> sieve(int n) {
    if (n < 2) {
        return {};
    }
    std::vector<char> prime(n + 1, true);
    for (int i = 2; i <= n / i; i++) {
        if (prime[i]) {
            for (int64_t j = 1LL * i * i; j <= n; j += i) {
                prime[j] = false;
            }
        }
    }
    std::vector<int> res;
    for (int i = 2; i <= n; i++) {
        if (prime[i]) {
            res.push_back(i);
        }
    }
    return res;
}

std::vector<int> sieve(int lo, int hi) {
    if (hi < 2 || lo > hi) {
        return {};
    }
    int sqrt_hi = sqrt(hi);
    while (1LL * (sqrt_hi + 1) * (sqrt_hi + 1) <= hi) {
        sqrt_hi++;
    }
    int fourth_root_hi = sqrt(sqrt_hi);
    while (1LL * (fourth_root_hi + 1) * (fourth_root_hi + 1) <= sqrt_hi) {
        fourth_root_hi++;
    }
    std::vector<char> prime1(sqrt_hi + 1, true), prime2(hi - lo + 1, true);
    for (int i = 2; i <= fourth_root_hi; i++) {
        if (prime1[i]) {
            for (int64_t j = 1LL * i * i; j <= sqrt_hi; j += i) {
                prime1[j] = false;
            }
        }
    }
    for (int i = 2; i <= sqrt_hi; i++) {
        if (prime1[i]) {
            int64_t first_multiple = ((static_cast<int64_t>(lo) + i - 1) / i) * i;
            int64_t start = std::max<int64_t>(i) * i, first_multiple);
            for (int64_t j = start; j <= hi; j += i) {
                prime2[j - lo] = false;
            }
        }
    }
    std::vector<int> res;
    for (int i = (lo > 1) ? lo : 2; i <= hi; i++) {
        if (prime2[i - lo]) {
            res.push_back(i);
        }
    }
    return res;
}

std::vector<int> linear_sieve(int n, std::vector<int> *least_out = nullptr) {
    if (n < 2) {
        if (least_out != nullptr) {
            least_out->assign(std::max(n + 1, 0), 0);
        }
        return {};
    }

```

```

std::vector<int> least(n + 1, 0), primes;
for (int i = 2; i <= n; i++) {
    if (least[i] == 0) {
        least[i] = i;
        primes.push_back(i);
    }
    for (int p : primes) {
        if (p > least[i] || i > n / p) {
            break;
        }
        least[i * p] = p;
    }
}
if (least_out != nullptr) {
    *least_out = least;
}
return primes;
}

```

Example Usage

```

#include <cassert>
#include <ctime>
#include <iostream>
using namespace std;

int main() {
    int pmax = 10000000;
    time_t start;
    double delta;

    start = clock();
    auto p = sieve(pmax);
    delta = static_cast<double>((clock() - start)) / CLOCKS_PER_SEC;
    cout << "sieve(n=" << pmax << "): " << delta << "s" << endl;
    assert((sieve(1) == vector<int>{}));
    assert((sieve(10) == vector<int>{2, 3, 5, 7}));

    vector<int> least;
    assert(linear_sieve(pmax, &least) == p);
    assert(least[999983] == 999983);
    assert(least[999984] == 2);

    int lo = 1000000000, hi = 1005000000;
    start = clock();
    p = sieve(lo, hi);
    delta = static_cast<double>((clock() - start)) / CLOCKS_PER_SEC;
    cout << "sieve([" << lo << ", " << hi << "]): " << delta << "s" << endl;
    assert((sieve(14, 16) == vector<int>{}));
    assert((sieve(17, 19) == vector<int>{17, 19}));
    assert(!p.empty() && p.front() >= lo && p.back() <= hi);
    return 0;
}

```

Example Output

```

sieve(n=10000000): 0.042641s
sieve([1000000000, 1005000000]): 0.01969s

```

6.3.8 Prime Tests and Factorization

Test whether an integer is prime and compute the prime factorization of an integer n . The Miller-Rabin test writes $n - 1 = d \cdot 2^r$ and squares a^d repeatedly, checking that the sequence behaves the only way it can for a

prime; a base a that breaks the pattern is a witness that n is composite. For large composite inputs, Pollard's rho iterates a pseudorandom map modulo n ; two iterates that collide modulo a hidden prime factor reveal that factor through a GCD with their difference. Prime factorizations are represented as sorted vectors of (`prime`, `exponent`) pairs. For 0 and 1, the prime factorization is empty.

- `is_prime_slow(n)` returns whether the integer `n` is prime using trial division.
- `factorize_slow(n)` returns the prime factorization of `n` using trial division.
- `is_probable_prime(n, k = 20)` returns whether `n` is probably prime using `k` random Miller-Rabin bases. The result is guaranteed correct when `n` is prime. For composite `n`, it returns false with probability at least $1 - (1/4)^k$. Thus a true result has one-sided error probability at most $(1/4)^k$. This is mostly useful when adapting Miller-Rabin outside the known 64-bit deterministic setting.
- `rho_factor(n)` returns a factor of `n` that is not necessarily prime using Pollard's rho with Brent's optimization. If `n` is prime, then `n` itself is returned. While this algorithm is non-deterministic and may fail to detect factors on certain runs of the same input, it can be retried until a nontrivial factor is found, as done in `factorize()`.
- `is_prime(n)` returns whether the signed 64-bit integer `n` is prime using deterministic Miller-Rabin bases sufficient up to and including $2^{63} - 1$.
- `cached_sieve(n)` returns a cached linear sieve up to and including `n`, storing both primes and least prime factors.
- `merge_factors(a, b)` merges two sorted compressed factorizations.
- `factorize_rho(n)` returns the prime factorization of a 64-bit integer using Miller-Rabin and Pollard's rho without initial trial division.
- `factorize(n, small_prime_limit = 1000000)` returns the prime factorization of a 64-bit integer `n` using a combination of trial division, the Miller-Rabin primality test, and Pollard's rho algorithm. `small_prime_limit` specifies the largest prime to test with trial division before falling back to the rho algorithm. This supports 64-bit integers up to and including $2^{63} - 1$.
- `divisors_from_factors(factors)` returns all divisors from a compressed factorization.
- `get_divisors(n)` returns all divisors of a 64-bit integer using `factorize()` followed by `divisors_from_factors()`.

Modular multiplication inside Miller-Rabin and Pollard's rho uses `_uint128_t` when available for speed, with a slower portable double-and-add fallback to avoid overflow on compilers without 128-bit integers.

Time Complexity:

- $O(\sqrt{n})$ per call to `is_prime_slow()` and `factorize_slow()`.
- $O(k \cdot \log(n))$ modular multiplications per call to `is_probable_prime()` and $O(\log(n))$ per call to `is_prime()`. The portable multiplication fallback adds another $O(\log(n))$ factor.
- $O(n)$ per call to `cached_sieve()` when rebuilding the cache, or $O(1)$ when the existing cache is large enough.
- $O(\sqrt{p})$ expected modular multiplications per call to `rho_factor(n)`, where p is the smallest prime factor of `n`. For composite `n`, this is at most $O(n^{1/4})$.
- $O(n^{1/4})$ expected per call to `factorize_rho()`. `factorize(n, L)` additionally performs $O(L/\log(L))$ trial divisions and takes $O(L)$ time if the prime cache must be rebuilt to `L`.
- $O(|a| + |b|)$ per call to `merge_factors()`.
- $O(d \log(d))$ per call to `divisors_from_factors()`, where d is the number of divisors generated, due to sorting the result.
- $O(L/\log(L) + n^{1/4} + d \log(d))$ expected per call to `get_divisors()`, plus $O(L)$ if the prime cache must be rebuilt.

Space Complexity:

- $O(c)$ cached space for the largest `cached_sieve(c)` call.
- $O(f + \log(n))$ auxiliary for recursive factorization, where f is the number of compressed prime factors returned.
- $O(d)$ output space and $O(\log(d))$ auxiliary stack space for `divisors_from_factors()` and `get_divisors()`.
- $O(1)$ auxiliary for the primality tests, `rho_factor()`, `factorize_slow()`, and `merge_factors()`, excluding returned vectors.

```

#include <algorithm>
#include <cstdlib>
#include <numeric>
#include <random>
#include <utility>
#include <vector>

template<typename Int>
bool is_prime_slow(Int n) {
    if (n == 2 || n == 3) {
        return true;
    }
    if (n < 2 || n % 2 == 0 || n % 3 == 0) {
        return false;
    }
    for (Int i = 5, w = 4; i <= n / i; i += w) {
        if (n % i == 0) {
            return false;
        }
        w = 6 - w;
    }
    return true;
}

template<typename Int>
std::vector<std::pair<Int, int>> factorize_slow(Int n) {
    if (n <= 1) {
        return {};
    }
    std::vector<std::pair<Int, int>> res;
    for (Int i = 2; i <= n / i; i++) {
        if (n % i == 0) {
            int exponent = 0;
            do {
                n /= i;
                exponent++;
            } while (n % i == 0);
            res.emplace_back(i, exponent);
        }
    }
    if (n > 1) {
        res.emplace_back(n, 1);
    }
    return res;
}

uint64_t mulmod(uint64_t x, uint64_t n, uint64_t m) {
#ifdef __SIZEOF_INT128__
    return static_cast<uint64_t>((static_cast<__uint128_t>(x) * n % m));
#else
    uint64_t a = 0, b = x % m;
    for (; n > 0; n >>= 1) {
        if (n & 1) {
            a = a >= m - b ? a - (m - b) : a + b;
        }
        b = b >= m - b ? b - (m - b) : b + b;
    }
    return a;
#endif
}

uint64_t powmod(uint64_t x, uint64_t n, uint64_t m) {
    uint64_t a = 1, b = x;
    for (; n > 0; n >>= 1) {
        if (n & 1) {
            a = mulmod(a, b, m);
        }
        b = mulmod(b, b, m);
    }

```

```

    }
    return a % m;
}

uint64_t rand64u() {
    static std::mt19937_64 rng(std::random_device{}());
    return rng();
}

bool is_probable_prime(int64_t n, int k = 20) {
    if (n == 2 || n == 3) {
        return true;
    }
    if (n < 2 || n % 2 == 0 || n % 3 == 0) {
        return false;
    }
    uint64_t s = n - 1, p = n - 1;
    while (!(s & 1)) {
        s >>= 1;
    }
    for (int i = 0; i < k; i++) {
        uint64_t x, r = powmod(rand64u() % (n - 3) + 2, s, n); // random witness base in [2, n - 2]
        for (x = s; x != p && r != 1 && r != p; x <= 1) {
            r = mulmod(r, r, n);
        }
        if (r != p && !(x & 1)) {
            return false;
        }
    }
    return true;
}

int64_t rho_factor(int64_t n) {
    if (n % 2 == 0) {
        return 2;
    }
    uint64_t y = rand64u() % (n - 1) + 1;
    uint64_t c = rand64u() % (n - 1) + 1;
    uint64_t m = rand64u() % (n - 1) + 1;
    uint64_t g = 1, r = 1, q = 1, ys = 0, x = 0;
    for (r = 1; g == 1; r <= 1) {
        x = y;
        for (uint64_t i = 0; i < r; i++) {
            y = (mulmod(y, y, n) + c) % n;
        }
        for (uint64_t k = 0; k < r && g == 1; k += m) {
            ys = y;
            int64_t lim = std::min(m, r - k);
            for (int j = 0; j < lim; j++) {
                y = (mulmod(y, y, n) + c) % n;
                q = mulmod(q, (x > y) ? (x - y) : (y - x), n);
            }
            g = std::gcd(q, n);
        }
    }
    if (g == static_cast<uint64_t>(n)) {
        do {
            ys = (mulmod(ys, ys, n) + c) % n;
            g = std::gcd((x > ys) ? (x - ys) : (ys - x), n);
        } while (g <= 1);
    }
    return g;
}

bool is_prime(int64_t n) {
    if (n < 2) {
        return false;
    }

```

```

static const int small_primes[] = {2, 3, 5, 7, 11, 13, 17, 19, 23, 29};
for (int p : small_primes) {
    if (n % p == 0) {
        return n == p;
    }
}
if (n < 31 * 31) {
    return true;
}
uint64_t t;
int s = 0;
for (t = n - 1; !(t & 1); t >>= 1) {
    s++;
}
static const uint64_t bases[] = {2, 325, 9375, 28178, 450775, 9780504, 1795265022};
for (uint64_t a : bases) {
    if (a % n == 0) {
        continue;
    }
    uint64_t r = powmod(a, t, n);
    if (r == 1) {
        continue;
    }
    bool ok = false;
    for (int j = 0; j < s && !ok; j++) {
        ok |= (r == static_cast<uint64_t>(n) - 1);
        r = mulmod(r, r, n);
    }
    if (!ok) {
        return false;
    }
}
return true;
}

struct SieveCache {
    int limit = 1;
    std::vector<int> least{0, 1}, primes;
};

SieveCache &cached_sieve(int n) {
    static SieveCache cache;
    if (n <= cache.limit) {
        return cache;
    }
    cache.limit = std::max(1, n);
    cache.least.assign(cache.limit + 1, 0);
    cache.primes.clear();
    for (int i = 2; i <= cache.limit; i++) {
        if (cache.least[i] == 0) {
            cache.least[i] = i;
            cache.primes.push_back(i);
        }
        for (int p : cache.primes) {
            if (p > cache.least[i] || i > cache.limit / p) {
                break;
            }
            cache.least[i * p] = p;
        }
    }
    return cache;
}

using Factors = std::vector<std::pair<int64_t, int>>;

Factors merge_factors(const Factors &a, const Factors &b) {
    Factors res;
    int i = 0, j = 0;

```



```

while (i < static_cast<int>(a.size()) || j < static_cast<int>(b.size())) {
    if (i < static_cast<int>(a.size()) && j < static_cast<int>(b.size()) &&
        a[i].first == b[j].first) {
        res.emplace_back(a[i].first, a[i].second + b[j].second);
        i++;
        j++;
    } else if (
        j == static_cast<int>(b.size()) ||
        (i < static_cast<int>(a.size()) && a[i].first < b[j].first)
    ) {
        res.push_back(a[i++]);
    } else {
        res.push_back(b[j++]);
    }
}
return res;
}

```

```

Factors factorize_rho(int64_t n) {
    if (n <= 1) {
        return {};
    }
    if (n % 2 == 0) {
        return merge_factors(Factors{{2, 1}}, factorize_rho(n / 2));
    }
    if (is_prime(n)) {
        return {{n, 1}};
    }
    int64_t p;
    do {
        p = rho_factor(n);
    } while (p == n);
    return merge_factors(factorize_rho(p), factorize_rho(n / p));
}

```

```

Factors factorize(int64_t n, int small_prime_limit = 1000000) {
    if (n <= 1) {
        return {};
    }
    const SieveCache &sieve = cached_sieve(small_prime_limit);
    if (n <= sieve.limit) {
        Factors res;
        while (n > 1) {
            int p = sieve.least[n], cnt = 0;
            do {
                n /= p;
                cnt++;
            } while (n > 1 && sieve.least[n] == p);
            res.emplace_back(p, cnt);
        }
        return res;
    }
    Factors res;
    const std::vector<int> &primes = sieve.primes;
    for (int p : primes) {
        if (p > small_prime_limit || static_cast<int64_t>(p) > n / p) {
            break;
        }
        if (n % p == 0) {
            int cnt = 0;
            do {
                n /= p;
                cnt++;
            } while (n % p == 0);
            res.emplace_back(p, cnt);
        }
    }
    return merge_factors(res, factorize_rho(n));
}

```

```

}

std::vector<int64_t> divisors_from_factors(const Factors &factors) {
    std::vector<int64_t> res{1};
    for (const auto &factor : factors) {
        int old_size = static_cast<int>(res.size());
        int64_t power = 1;
        for (int e = 1; e <= factor.second; e++) {
            power *= factor.first;
            for (int i = 0; i < old_size; i++) {
                res.push_back(res[i] * power);
            }
        }
    }
    std::sort(res.begin(), res.end());
    return res;
}

std::vector<int64_t> get_divisors(int64_t n) {
    if (n <= 1) {
        return (n < 1) ? std::vector<int64_t>() : std::vector<int64_t>(1, 1);
    }
    return divisors_from_factors(factorize(n));
}

```

Example Usage

```

#include <cassert>
#include <ctime>
#include <iomanip>
#include <iostream>
#include <set>
using namespace std;

void validate(int64_t n, const Factors &factors) {
    if (n <= 1) {
        assert(factors.empty());
        return;
    }
    if (is_prime(n)) {
        assert((factors == Factors{{n, 1}}));
        return;
    }
    int64_t prod = 1;
    for (const auto &factor : factors) {
        assert(is_prime(factor.first));
        for (int i = 0; i < factor.second; i++) {
            prod *= factor.first;
        }
    }
    assert(prod == n);
}

int main() {
    { // Small primality tests.
        vector<pair<int, bool>> cases{
            {-10, false}, {-1, false}, {0, false}, {1, false}, // non-primes
            {2, true}, {3, true}, {4, false}, {5, true},
            {9, false}, // divisible by 3
            {25, false}, // square of small prime
            {29, true}, // last small-prime table entry
            {31, true}, // just past small-prime table
            {49, false}, // square, not divisible by 2 or 3
            {121, false}, // 11^2
            {169, false}, // 13^2
            {961, false}, // 31^2, boundary for n < 31*31 shortcut
        };
    }
}

```

```

    {997, true},    // prime near 1000
    {1001, false}, // 7*11*13
    {2047, false}, // 23*89, pseudoprime to base 2
    {341, false},  // 11*31, Fermat pseudoprime to base 2
    {561, false},  // Carmichael number
};
for (const auto &[n, expected] : cases) {
    assert(is_prime_slow(n) == expected);
    assert(is_probable_prime(n) == expected);
    assert(is_prime(n) == expected);
}
}
{ // Primality test benchmark.
    vector<int64_t> nums{
        772023803LL, 2147483647LL, 5705234089LL, 6339503641LL, 999966000289LL,
    };
    cout << "Primality test:" << endl;
    cout << setw(15) << left << "num:";
    cout << setw(18) << left << "is_prime_slow():";
    cout << setw(20) << left << "is_probable_prime():";
    cout << setw(13) << right << "is_prime():" << endl;
    cout << fixed << setprecision(3);
    for (int64_t n : nums) {
        clock_t start = clock();
        bool p1 = is_prime_slow(n);
        double t1 = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;
        start = clock();
        bool p2 = is_probable_prime(n);
        double t2 = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;
        start = clock();
        bool p3 = is_prime(n);
        double t3 = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;
        assert(p1 == p2 && p1 == p3);
        cout << setw(12) << left << n << right;
        cout << setw(14) << 1000 * t1 << "ms";
        cout << setw(16) << 1000 * t2 << "ms";
        cout << setw(16) << 1000 * t3 << "ms" << endl;
    }
    cout << endl;
}
{ // Small factorization tests.
    for (int64_t i = 1; i <= 10000; i++) {
        auto v1 = factorize_slow(i);
        auto v2 = factorize(i);
        validate(i, v1);
        assert(v1 == v2);
        auto d = get_divisors(i);
        set<int> s(d.begin(), d.end());
        assert(d.size() == s.size());
        for (int j = 1; j <= i; j++) {
            if (i % j == 0) {
                assert(s.count(j));
            }
        }
    }
}
{ // Compressed factors are convenient for building divisors.
    Factors factors{{2, 3}, {3, 2}, {5, 1}};
    vector<int64_t> divisors = divisors_from_factors(factors);
    assert(divisors.size() == 24);
    assert(divisors.front() == 1 && divisors.back() == 360);
    assert(factorize(360) == factors);
    assert(get_divisors(360) == divisors);
}
{ // Large factorization tests.
    const vector<int64_t> nums{
        (1LL << 62), // high power of two
        9223372036854775807LL, // 2^63 - 1, many small factors
    };
}

```

```

    999983LL * 999983,           // square of a large prime
    1900009LL * 1910009 * 2541547, // three medium-size prime factors
    50000017LL * 50001037,       // 16-digit semiprime
};
cout << "Factorization test:" << endl;
cout << setw(24) << left << "num:";
cout << setw(20) << left << "factorize_slow():";
cout << "factorize():" << endl;
cout << fixed << setprecision(3);
for (int64_t n : nums) {
    clock_t start = clock();
    Factors factors1 = factorize_slow(n);
    double t1 = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;
    start = clock();
    Factors factors2 = factorize(n);
    double t2 = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;
    validate(n, factors1);
    validate(n, factors2);
    assert(factors1 == factors2);
    cout << setw(26) << left << n << right << setw(8) << 1000 * t1 << "ms";
    cout << setw(16) << 1000 * t2 << "ms" << endl;
    cout << left;
}
}
return 0;
}

```

Example Output

```

Primality test:
num:      is_prime_slow():  is_probable_prime():  is_prime():
772023803      0.006ms      0.005ms      0.001ms
2147483647      0.009ms      0.006ms      0.002ms
5705234089      0.014ms      0.000ms      0.001ms
6339503641      0.013ms      0.001ms      0.001ms
999966000289      0.171ms      0.001ms      0.001ms

Factorization test:
num:      factorize_slow():  factorize():
4611686018427387904      0.002ms      0.002ms
9223372036854775807      0.049ms      0.007ms
999966000289      0.503ms      0.043ms
9223361212852495307      0.952ms      0.096ms
2500052700017629      24.587ms      0.194ms

```

6.3.9 Euler's Totient Function

Euler's totient function $\phi(n)$ returns the number of positive integers less than or equal to n that are relatively prime to n ; that is, the number of integers k in the range $[1, n]$ for which $\gcd(n, k) = 1$.

The function is multiplicative, with $\phi(p^k) = p^k - p^{k-1}$ for a prime power, which yields the product formula $\phi(n) = n \prod_{p|n} (1 - 1/p)$ taken over the distinct primes p dividing n . `phi(n)` applies this formula directly, factoring n by trial division. `phi_table(n)` computes the whole range at once with a sieve of Eratosthenes: starting from `v[i] = i`, it sweeps each prime p across its multiples `j` and folds in the factor $(1 - 1/p)$ by replacing `v[j]` with `v[j] - v[j]/p`. This is faster than the alternative that uses the divisor-sum identity $\sum_{d|n} \phi(d) = n$ to subtract each `v[i]` from its multiples, which costs $O(n \log n)$.

The overload `phi_table(lo, hi)` computes the totients of a high but narrow range with a segmented sieve, analogous to the segmented prime sieve. It sieves the primes up to $\sqrt{hi}$, applies each one's $(1 - 1/p)$ factor to its multiples within $[lo, hi]$ while dividing that prime out of a parallel copy of every value, and finally folds in any single leftover prime factor greater than $\sqrt{hi}$ (a number at most `hi` can have at most one such factor).

- `phi(n)` returns Euler's totient of `n`.

- `phi_table(n)` returns a vector `v` of length `n + 1` such that `v[i]` stores `phi(i)` for every `i` in the range `[0, n]`.
- `phi_table(lo, hi)` returns a vector `v` of length `hi - lo + 1` such that `v[i - lo]` stores `phi(i)` for every `i` in the range `[lo, hi]`. This overload assumes `lo ≥ 0`; if `hi < lo`, it returns an empty vector.

Time Complexity:

- $O(\sqrt{n})$ per call to `phi(n)`.
- $O(n \log \log n)$ per call to `phi_table(n)`.
- $O(\sqrt{hi} \cdot \log(\log(hi)) + w \cdot \log(\log(hi)))$ per call to `phi_table(lo, hi)`, where $w = hi - lo + 1$.

Space Complexity:

- $O(1)$ auxiliary for `phi(n)`.
- $O(n)$ auxiliary for `phi_table(n)`.
- $O(w + \sqrt{hi})$ auxiliary for `phi_table(lo, hi)`.

```
#include <algorithm>
#include <cmath>
#include <cstdint>
#include <numeric>
#include <vector>

int phi(int n) {
    int res = n;
    for (int i = 2; i <= n / i; i++) {
        if (n % i == 0) {
            while (n % i == 0) {
                n /= i;
            }
            res -= res / i;
        }
    }
    if (n > 1) {
        res -= res / n;
    }
    return res;
}

std::vector<int> phi_table(int n) {
    std::vector<int> res(n + 1);
    std::iota(res.begin(), res.end(), 0);
    for (int i = 2; i <= n; i++) {
        if (res[i] == i) { // i is prime, so apply its (1 - 1/i) factor to every multiple.
            for (int j = i; j <= n; j += i) {
                res[j] -= res[j] / i;
            }
        }
    }
    return res;
}

std::vector<int> phi_table(int lo, int hi) {
    if (hi < lo) {
        return {};
    }
    int root = static_cast<int>(sqrt(hi));
    while (1LL * (root + 1) * (root + 1) <= hi) {
        root++;
    }
    // Sieve the primes up to sqrt(hi); any larger prime divisor is the one leftover factor.
    std::vector<char> composite(root + 1, false);
    std::vector<int> primes;
    for (int i = 2; i <= root; i++) {
        if (!composite[i]) {
            primes.push_back(i);
            for (int64_t j = 1LL * i * i; j <= root; j += i) {
                composite[j] = true;
            }
        }
    }
    std::vector<int> res(hi - lo + 1);
    for (int i = lo; i <= hi; i++) {
        int val = i;
        for (int p : primes) {
            if (p > val) break;
            if (val % p == 0) {
                val -= val / p;
            }
        }
        res[i - lo] = val;
    }
    return res;
}
```

```

    }
}
std::vector<int> res(hi - lo + 1);
std::vector<int64_t> remaining(hi - lo + 1);
for (int i = lo; i <= hi; i++) {
    res[i - lo] = i;
    remaining[i - lo] = i;
}
for (int p : primes) {
    // Start at the first multiple of p at least lo, but never below p itself (this skips 0).
    int64_t start = std::max(1LL * p, ((1LL * lo + p - 1) / p) * p);
    for (int64_t j = start; j <= hi; j += p) {
        int idx = static_cast<int>(j - lo);
        res[idx] -= res[idx] / p;
        while (remaining[idx] % p == 0) {
            remaining[idx] /= p;
        }
    }
}
for (int i = lo; i <= hi; i++) {
    int idx = i - lo;
    if (remaining[idx] > 1) { // a single prime factor greater than sqrt(hi) is left over
        res[idx] -= static_cast<int>(res[idx] / remaining[idx]);
    }
}
return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    assert(phi(1) == 1);
    assert(phi(9) == 6);
    assert(phi(1234567) == 1224720);
    const int n = 1000;
    vector<int> pt = phi_table(n);
    for (int i = 0; i <= n; i++) {
        assert(pt[i] == phi(i));
    }
    // The segmented overload agrees with phi(), including over a low range (with 0 and 1) and a
    // high, narrow window where many values have a prime factor exceeding sqrt(hi).
    vector<int> lphi = phi_table(0, 50);
    for (int i = 0; i <= 50; i++) {
        assert(lphi[i] == phi(i));
    }
    int lo = 1000000000, hi = 1000000100;
    vector<int> hphi = phi_table(lo, hi);
    for (int i = lo; i <= hi; i++) {
        assert(hphi[i - lo] == phi(i));
    }
    return 0;
}

```

6.3.10 Mobius Function and Inclusion-Exclusion

The Mobius function $\mu(n)$ is the multiplicative function defined as 1 if $n = 1$, 0 if n is divisible by the square of a prime, and $(-1)^k$ if n is a product of k distinct primes. It is the coefficient that appears in Mobius inversion: if

$g(n) = \sum_{d|n} f(d)$ for all n , then $f(n) = \sum_{d|n} \mu(d) g(n/d)$. This formalizes the principle of inclusion-exclusion in the divisor lattice, letting one recover a summand function from its divisor sums.

Inclusion-exclusion itself counts the size of a union by alternately adding and subtracting the sizes of intersections. A canonical number-theoretic instance is counting integers in $[1, n]$ that are coprime to a modulus m : subtract the multiples of each prime factor of m , add back those counted twice, and so on. The signed terms are precisely the values of μ over the squarefree divisors of m , so `count_coprime(n, m)` equals $\sum_{d|m} \mu(d) \lfloor n/d \rfloor$.

- `mobius(n)` returns $\mu(n)$ for a single `n` by trial division.
- `mobius_sieve(n)` returns a vector of length $n + 1$ where index i holds $\mu(i)$, computed together with a linear sieve.
- `count_coprime(n, m)` returns the number of integers in $[1, n]$ that are coprime to m , via inclusion-exclusion over the distinct prime factors of m .

Time Complexity:

- $O(\sqrt{n})$ per call to `mobius()`.
- $O(n)$ per call to `mobius_sieve()`.
- $O(\sqrt{m} + 2^k \cdot k)$ per call to `count_coprime()`, where k is the number of distinct prime factors of m (at most 15 for m within 64-bit range).

Space Complexity:

- $O(n)$ auxiliary for `mobius_sieve()`.
- $O(k)$ auxiliary for `count_coprime()`.
- $O(1)$ auxiliary for `mobius()`.

```
#include <stdint>
#include <vector>

int mobius(int n) {
    if (n == 1) {
        return 1;
    }
    int res = 1;
    for (int p = 2; p <= n / p; p++) {
        if (n % p == 0) {
            n /= p;
            if (n % p == 0) {
                return 0; // p^2 divides the original n.
            }
            res = -res;
        }
    }
    if (n > 1) {
        res = -res; // One remaining prime factor larger than its square root.
    }
    return res;
}

std::vector<int> mobius_sieve(int n) {
    std::vector<int> mu(n + 1, 0), primes;
    std::vector<char> composite(n + 1, false);
    if (n >= 1) {
        mu[1] = 1;
    }
    for (int i = 2; i <= n; i++) {
        if (!composite[i]) {
            primes.push_back(i);
            mu[i] = -1;
        }
        for (int p : primes) {
            if (p > n / i) {
                break;
            }

```

```

    composite[i * p] = true;
    if (i % p == 0) {
        mu[i * p] = 0; // p^2 divides i * p.
        break;
    }
    mu[i * p] = -mu[i];
}
}
return mu;
}

int64_t count_coprime(int64_t n, int64_t m) {
    std::vector<int64_t> primes;
    for (int64_t p = 2; p <= m / p; p++) {
        if (m % p == 0) {
            primes.push_back(p);
            while (m % p == 0) {
                m /= p;
            }
        }
    }
    if (m > 1) {
        primes.push_back(m);
    }
    int k = static_cast<int>(primes.size());
    int64_t total = 0;
    for (int mask = 0; mask < (1 << k); mask++) {
        int64_t product = 1;
        int bits = 0;
        for (int i = 0; i < k; i++) {
            if ((mask >> i) & 1) {
                product *= primes[i];
                bits++;
            }
        }
        total += (bits & 1 ? -1 : 1) * (n / product);
    }
    return total;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    assert(mobius(1) == 1);
    assert(mobius(2) == -1);
    assert(mobius(6) == 1); // 2 * 3, two distinct primes.
    assert(mobius(12) == 0); // Divisible by 2^2.
    assert(mobius(30) == -1); // 2 * 3 * 5, three distinct primes.

    vector<int> mu = mobius_sieve(12);
    assert((mu == vector<int>{0, 1, -1, -1, 0, -1, 1, -1, 0, 0, 1, -1, 0}));

    assert(count_coprime(10, 6) == 3); // 1, 5, 7
    assert(count_coprime(100, 1) == 100);
    return 0;
}

```


6.4 Arbitrary Precision Arithmetic

6.4.1 Big Integers (Simple)

Perform simple arithmetic operations on arbitrary precision big integers whose digits are internally represented as an `std::string` in little-endian order.

This version is intentionally small and decimal-oriented. Addition and subtraction are the usual schoolbook digit scans with carry or borrow. Multiplication keeps a shifted copy of the left operand for each digit of the right operand, adds that row once per digit value, then shifts by one decimal place. Division is long division: it scans the dividend from most significant digit to least, maintains a running remainder, and repeatedly subtracts the divisor to discover each quotient digit.

- `BigInt(n = 0)` constructs a big integer from an integer `n`.
- `BigInt(s)` constructs a big integer from a string `s`, which must strictly consist of a sequence of numeric digits, optionally preceded by a minus sign.
- `to_string()` returns the string representation of the big integer.
- `comp(a, b)` returns `-1`, `0`, or `1` depending on whether the big integers `a` and `b` compare less, equal, or greater, respectively.
- `add(a, b)` returns the sum of big integers `a` and `b`.
- `sub(a, b)` returns the difference of big integers `a` and `b`.
- `mul(a, b)` returns the product of big integers `a` and `b`.
- `div(a, b)` returns the quotient of big integers `a` and `b`.

Time Complexity:

- $O(n)$ per call to the constructor, `to_string()`, `comp()`, `add()`, and `sub()`, where n is the total number of digits in the arguments and result for each operation.
- $O(n \cdot m)$ per call to `mul(a, b)` and `div(a, b)` where n is the number of digits in `a` and m is the number of digits in `b`.

Space Complexity:

- $O(n)$ for storage of the big integer, where n is the number of digits.
- $O(n)$ auxiliary for `to_string()`, `add()`, `sub()`, `mul()`, and `div()`, where n is the total number of digits in the arguments and result for each operation.

```
#include <algorithm>
#include <cassert>
#include <cctype>
#include <cstdint>
#include <cstdint>
#include <stdexcept>
#include <string>

class BigInt {
    std::string digits;
    int sign;

    void normalize() {
        std::size_t pos = digits.find_last_not_of('0');
        if (pos != std::string::npos) {
            digits.erase(pos + 1);
        } else {
            digits = "0"; // An all-zero (or empty) magnitude collapses to a single "0".
        }
        if (digits.size() == 1 && digits[0] == '0') {
            sign = 1;
        }
    }
}

static int comp(const std::string &a, const std::string &b, int asgn, int bsign) {
```

```

    if (assign != bsign) {
        return assign < bsign ? -1 : 1;
    }
    if (a.size() != b.size()) {
        return a.size() < b.size() ? -assign : assign;
    }
    for (int i = static_cast<int>(a.size()) - 1; i >= 0; i--) {
        if (a[i] != b[i]) {
            return a[i] < b[i] ? -assign : assign;
        }
    }
    return 0;
}

static BigInt add(const std::string &a, const std::string &b, int assign, int bsign) {
    if (assign != bsign) {
        return (assign == 1) ? sub(a, b, assign, 1) : sub(b, a, bsign, 1);
    }
    BigInt res;
    res.sign = assign;
    res.digits.resize(std::max(a.size(), b.size()) + 1, '0');
    for (int i = 0, carry = 0; i < static_cast<int>(res.digits.size()); i++) {
        int d = carry;
        if (i < static_cast<int>(a.size())) {
            d += a[i] - '0';
        }
        if (i < static_cast<int>(b.size())) {
            d += b[i] - '0';
        }
        res.digits[i] = '0' + (d % 10);
        carry = d / 10;
    }
    res.normalize();
    return res;
}

static BigInt sub(const std::string &a, const std::string &b, int assign, int bsign) {
    if (assign == -1 || bsign == -1) {
        return add(a, b, assign, -bsign);
    }
    BigInt res;
    if (comp(a, b, assign, bsign) < 0) {
        res = sub(b, a, bsign, assign);
        res.sign = -1;
        return res;
    }
    res.digits.assign(a.size(), '0');
    for (int i = 0, borrow = 0; i < static_cast<int>(res.digits.size()); i++) {
        int d = (i < static_cast<int>(b.size())) ? a[i] - b[i] : a[i] - '0' - borrow;
        if (a[i] > '0') {
            borrow = 0;
        }
        if (d < 0) {
            d += 10;
            borrow = 1;
        }
        res.digits[i] = '0' + (d % 10);
    }
    res.normalize();
    return res;
}

public:
    BigInt(int64_t n = 0) {
        sign = (n < 0) ? -1 : 1;
        if (n == 0) {
            digits = "0";
            return;
        }

```

```

    }
    for (; n != 0; n /= 10) {
        int digit = static_cast<int>(n % 10);
        digits += '0' + (digit < 0 ? -digit : digit);
    }
    normalize();
}

BigInt(const std::string &s) {
    if (s.empty() || (s[0] == '-' && s.size() == 1)) {
        throw std::runtime_error("Invalid string format to construct BigInt.");
    }
    digits.assign(s.rbegin(), s.rend());
    if (s[0] == '-') {
        sign = -1;
        digits.erase(digits.size() - 1);
    } else {
        sign = 1;
    }
    if (digits.find_first_not_of("0123456789") != std::string::npos) {
        throw std::runtime_error("Invalid string format to construct BigInt.");
    }
    normalize();
}

std::string to_string() const {
    return (sign < 0 ? "-" : "") + std::string(digits.rbegin(), digits.rend());
}

friend int comp(const BigInt &a, const BigInt &b) {
    return comp(a.digits, b.digits, a.sign, b.sign);
}

friend BigInt add(const BigInt &a, const BigInt &b) {
    return add(a.digits, b.digits, a.sign, b.sign);
}

friend BigInt sub(const BigInt &a, const BigInt &b) {
    return sub(a.digits, b.digits, a.sign, b.sign);
}

friend BigInt mul(const BigInt &a, const BigInt &b) {
    BigInt res, row(a);
    for (char dc : b.digits) {
        for (int j = 0; j < (dc - '0'); j++) {
            res = add(res.digits, row.digits, res.sign, row.sign);
        }
        if (row.digits.size() > 1 || row.digits[0] != '0') {
            row.digits.insert(0, "0");
        }
    }
    res.sign = a.sign * b.sign;
    res.normalize();
    return res;
}

friend BigInt div(const BigInt &a, const BigInt &b) {
    assert(b.digits.size() != 1 || b.digits[0] != '0');
    BigInt res, row;
    res.digits.assign(a.digits.size(), '0');
    row.digits.clear(); // Running remainder; kept normalized and nonnegative below.
    for (int i = static_cast<int>(a.digits.size()) - 1; i >= 0; i--) {
        row.digits.insert(row.digits.begin(), a.digits[i]);
        row.normalize(); // Strip the spurious high-order zero left by the previous remainder.
        // Subtract while the remainder is >= b. A strict ">" undercounts the quotient digit when the
        // remainder equals b, and leaves the remainder too large, overflowing the next digit past 9.
        while (comp(row.digits, b.digits, 1, 1) >= 0) {
            res.digits[i]++;
        }
    }
}

```

```

        row = sub(row.digits, b.digits, 1, 1);
    }
}
res.sign = a.sign * b.sign;
res.normalize();
return res;
}
};

```

Example Usage

```

int main() {
    BigInt a("-9899819294989142124"), b("12398124981294214");
    assert(add(a, b).to_string() == "-9887421170007847910");
    assert(sub(a, b).to_string() == "-9912217419970436338");
    assert(mul(a, b).to_string() == "-122739196911503356525379735104870536");
    assert(div(a, b).to_string() == "-798");
    assert(comp(a, b) == -1 && comp(a, a) == 0 && comp(b, a) == 1);
    assert(BigInt(INT64_MIN).to_string() == "-9223372036854775808");
    return 0;
}

```

6.4.2 Big Integers

Perform operations on arbitrary precision big integers internally represented as a vector of base-`BASE` digits in little-endian order. `BASE` defaults to 10^9 and must be a power of 10 no larger than 10^9 ; `USE_FFT_MULT` selects FFT multiplication when true and Karatsuba multiplication when false. Typical arithmetic operations involving mixed numeric primitives and strings are supported through implicit construction and hidden friend operators, as long as at least one operand is a `BigInt` at any given level of evaluation.

- `BigInt()` constructs zero, and `BigInt(n)` constructs a big integer from an integer `n`.
- `BigInt(s)` constructs a big integer from a C string or an `std::string s`.
- `operator=` is defined to copy from another big integer or to assign from a 64-bit integer primitive.
- `size()` returns the number of digits in the base-10 representation.
- Operators `>>` and `<<` are defined to support stream-based input and output.
- `v.to_string()`, `v.to_llong()`, `v.to_double()`, and `v.to_ldouble()` return the big integer `v` converted to an `std::string`, `int64_t`, `double`, and `long double` respectively. For the latter three data types, overflow behavior is based on that of inputting from `std::istream`. Equivalent conversions are also available through explicit casts to `int`, `long long`, `double`, and `long double`.
- `v.abs()` returns the absolute value of big integer `v`.
- `a.comp(v)` returns `-1`, `0`, or `1` depending on whether the big integers `a` and `v` compare less, equal, or greater, respectively.
- Operators `<`, `>`, `<=`, `>=`, `==`, `!=`, `+`, `-`, `*`, `/`, `%`, `++`, `-`, `+=`, `-=`, `*=`, `/=`, and `%=` are defined analogous to those on integer primitives. Addition, subtraction, and comparisons are performed using the standard linear algorithms. Multiplication first converts limbs from base 10^9 to base 10^4 so coefficient products fit safely, then uses either Karatsuba multiplication (if `USE_FFT_MULT` is false) or complex FFT convolution (if `USE_FFT_MULT` is true), converting the result back to base 10^9 . Division and modulo are computed together by normalized long division: scale the operands so the divisor's leading limb is large, estimate each quotient limb from the top one or two limbs of the running remainder, correct the estimate by adding the divisor back if necessary, then unscale the remainder.
- `a.div(v)` returns a pair consisting of the quotient and remainder when `a` is divided by `v`.
- `v.pow(n)` returns `v` raised to the nonnegative power `n`.
- `v.sqrt()` returns the integral part of the square root of big integer `v`.
- `v.nth_root(n)` returns the integral part of the n -th root of big integer `v`. Negative inputs require odd n .
- `rand(n)` returns a random, positive big integer with n digits.

`pow(n)` uses binary exponentiation. `sqrt()` uses a digit-by-digit square-root algorithm in the internal base, while `nth_root(n)` uses binary search over the answer range.

Time Complexity:

- $O(n)$ per call to the constructors, `size()`, `to_string()`, `to_llong()`, `to_double()`, `to_ldouble()`, `abs()`, `comp()`, `rand()`, and all comparison and arithmetic operators except multiplication, division, and modulo, where n is total number of digits in the arguments and result for each operation.
- $O(n \cdot \log(n) \cdot \log(\log(n)))$ or $O(n^{1.585})$ per call to multiplication operations, depending on whether `USE_FFT_MULT` is set to true or false.
- $O(n \cdot m)$ per call to division and modulo operations, where n and m are the number of digits in the dividend and divisor, respectively.
- $O(M(d) \log n)$ per call to `pow()`, where d is the maximum digit length reached during exponentiation and $M(d)$ denotes the cost of one multiplication of two d -digit big integers using the selected multiplication method above.
- $O(d^2)$ per call to `sqrt()`, where d is the number of digits.
- $O((d/k) \cdot M(d) \cdot \log k)$ per call to `nth_root()`, from binary search over a d/k -digit answer.

Space Complexity:

- $O(n)$ for storage of the big integer.
- $O(n)$ auxiliary for negation, addition, subtraction, multiplication, division, `abs()`, `sqrt()`, `pow()`, and `nth_root()`.
- $O(1)$ auxiliary for all other operations.

```
#include <algorithm>
#include <cassert>
#include <climits>
#include <cmath>
#include <complex>
#include <cstdint>
#include <cstring>
#include <iomanip>
#include <istream>
#include <iterator>
#include <ostream>
#include <random>
#include <sstream>
#include <string>
#include <utility>
#include <vector>

class BigInt {
    static const int BASE = 1000000000, BASE_DIGITS = 9;
    static const bool USE_FFT_MULT = true;
    static_assert(10 <= BASE && BASE <= 1000000000, "BASE must be between 10 and 10^9");
    static_assert(BASE_DIGITS > 0, "BASE must be a power of 10");

    using vint = std::vector<int>;
    using vint64 = std::vector<int64_t>;
    using vcd = std::vector<std::complex<double>>;

    vint digits;
    int sign;

    void normalize() {
        while (!digits.empty() && digits.back() == 0) {
            digits.pop_back();
        }
        if (digits.empty()) {
            sign = 1;
        }
    }

    void read(int n, const char *s) {
```

```

    sign = 1;
    digits.clear();
    int pos = 0;
    while (pos < n && (s[pos] == '-' || s[pos] == '+')) {
        if (s[pos] == '-') {
            sign = -sign;
        }
        pos++;
    }
    for (int i = n - 1; i >= pos; i -= BASE_DIGITS) {
        int x = 0;
        for (int j = std::max(pos, i - BASE_DIGITS + 1); j <= i; j++) {
            x = x * 10 + s[j] - '0';
        }
        digits.push_back(x);
    }
    normalize();
}

static int comp(const vint &a, const vint &b, int asign, int bsign) {
    if (asign != bsign) {
        return asign < bsign ? -1 : 1;
    }
    if (a.size() != b.size()) {
        return a.size() < b.size() ? -asign : asign;
    }
    for (int i = static_cast<int>(a.size()) - 1; i >= 0; i--) {
        if (a[i] != b[i]) {
            return a[i] < b[i] ? -asign : asign;
        }
    }
    return 0;
}

static BigInt add(const vint &a, const vint &b, int asign, int bsign) {
    if (asign != bsign) {
        return (asign == 1) ? sub(a, b, asign, 1) : sub(b, a, bsign, 1);
    }
    BigInt res;
    res.digits = a;
    res.sign = asign;
    int carry = 0, size = static_cast<int>(std::max(a.size(), b.size()));
    for (int i = 0; i < size || carry; i++) {
        if (i == static_cast<int>(res.digits.size())) {
            res.digits.push_back(0);
        }
        res.digits[i] += carry + (i < static_cast<int>(b.size()) ? b[i] : 0);
        carry = (res.digits[i] >= BASE) ? 1 : 0;
        if (carry) {
            res.digits[i] -= BASE;
        }
    }
    return res;
}

static BigInt sub(const vint &a, const vint &b, int asign, int bsign) {
    if (asign == -1 || bsign == -1) {
        return add(a, b, asign, -bsign);
    }
    BigInt res;
    if (comp(a, b, asign, bsign) < 0) {
        res = sub(b, a, bsign, asign);
        res.sign = -1;
        return res;
    }
    res.digits = a;
    res.sign = asign;
    for (int i = 0, borrow = 0; i < static_cast<int>(a.size()) || borrow; i++) {

```

```

    res.digits[i] -= borrow + (i < static_cast<int>(b.size()) ? b[i] : 0);
    borrow = res.digits[i] < 0;
    if (borrow) {
        res.digits[i] += BASE;
    }
}
res.normalize();
return res;
}

static vint convert_base(const vint &digits, int l1, int l2) {
    vint64 p(std::max(l1, l2) + 1);
    p[0] = 1;
    for (int i = 1; i < static_cast<int>(p.size()); i++) {
        p[i] = p[i - 1] * 10;
    }
    vint res;
    int64_t curr = 0;
    for (int i = 0, curr_digits = 0; i < static_cast<int>(digits.size()); i++) {
        curr += digits[i] * p[curr_digits];
        curr_digits += l1;
        while (curr_digits >= l2) {
            res.push_back(static_cast<int>((curr % p[l2])));
            curr /= p[l2];
            curr_digits -= l2;
        }
    }
    res.push_back(static_cast<int>(curr));
    while (!res.empty() && res.back() == 0) {
        res.pop_back();
    }
    return res;
}

template<typename It>
static vint64 karatsuba(It alo, It ahi, It blo, It bhi) {
    int n = std::distance(alo, ahi), k = n / 2;
    vint64 res(n * 2);
    if (n <= 32) {
        for (int i = 0; i < n; i++) {
            for (int j = 0; j < n; j++) {
                res[i + j] += alo[i] * blo[j];
            }
        }
        return res;
    }
    auto a1b1 = karatsuba(alo, alo + k, blo, blo + k);
    auto a2b2 = karatsuba(alo + k, ahi, blo + k, bhi);
    vint64 a2(alo + k, ahi), b2(blo + k, bhi);
    for (int i = 0; i < k; i++) {
        a2[i] += alo[i];
        b2[i] += blo[i];
    }
    auto r = karatsuba(a2.begin(), a2.end(), b2.begin(), b2.end());
    for (int i = 0; i < static_cast<int>(a1b1.size()); i++) {
        r[i] -= a1b1[i];
        res[i] += a1b1[i];
    }
    for (int i = 0; i < static_cast<int>(a2b2.size()); i++) {
        r[i] -= a2b2[i];
        res[i + n] += a2b2[i];
    }
    for (int i = 0; i < static_cast<int>(r.size()); i++) {
        res[i + k] += r[i];
    }
    return res;
}

```

```

template<typename It>
static vcd fft(It lo, It hi, bool invert = false) {
    int n = std::distance(lo, hi), k = 0, high1 = -1;
    while ((1 << k) < n) {
        k++;
    }
    std::vector<int> rev(n, 0);
    for (int i = 1; i < n; i++) {
        if (!(i & (i - 1))) {
            high1++;
        }
        rev[i] = rev[i ^ (1 << high1)];
        rev[i] |= (1 << (k - high1 - 1));
    }
    vcd roots(n), res(n);
    for (int i = 0; i < n; i++) {
        double alpha = 2 * 3.14159265358979323846 * i / n;
        roots[i] = std::complex<double>(cos(alpha), sin(alpha));
        res[i] = *(lo + rev[i]);
    }
    for (int len = 1; len < n; len <= 1) {
        vcd tmp(n);
        int rstep = roots.size() / (len << 1);
        for (int pdest = 0; pdest < n; pdest += len) {
            int p = pdest;
            for (int i = 0; i < len; i++) {
                std::complex<double> c = roots[i * rstep] * res[p + len];
                tmp[pdest] = res[p] + c;
                tmp[pdest + len] = res[p] - c;
                pdest++;
                p++;
            }
        }
        res.swap(tmp);
    }
    if (invert) {
        for (int i = 0; i < static_cast<int>(res.size()); i++) {
            res[i] /= n;
        }
        std::reverse(res.begin() + 1, res.end());
    }
    return res;
}

public:
BigInt() : sign(1) {}
BigInt(int v) { *this = static_cast<int64_t>(v); }
BigInt(int64_t v) { *this = v; }
BigInt(const char *s) { read(strlen(s), s); }
BigInt(const std::string &s) { read(s.size(), s.c_str()); }

BigInt &operator=(int64_t v) {
    sign = v < 0 ? -1 : 1;
    digits.clear();
    for (; v != 0; v /= BASE) {
        int limb = static_cast<int>(v % BASE);
        digits.push_back(limb < 0 ? -limb : limb);
    }
    return *this;
}

int size() const {
    if (digits.empty()) {
        return 1;
    }
    std::ostringstream oss;
    oss << digits.back();
    return oss.str().length() + BASE_DIGITS * (digits.size() - 1);
}

```



```

}

friend std::istream &operator>>(std::istream &in, BigInt &v) {
    std::string s;
    in >> s;
    v.read(s.size(), s.c_str());
    return in;
}

friend std::ostream &operator<<(std::ostream &out, const BigInt &v) {
    if (v.sign == -1) {
        out << '-';
    }
    out << (v.digits.empty() ? 0 : v.digits.back());
    for (int i = static_cast<int>(v.digits.size()) - 2; i >= 0; i--) {
        out << std::setw(BASE_DIGITS) << std::setfill('0') << v.digits[i];
    }
    return out;
}

std::string to_string() const {
    std::ostringstream oss;
    if (sign == -1) {
        oss << '-';
    }
    oss << (digits.empty() ? 0 : digits.back());
    for (int i = static_cast<int>(digits.size()) - 2; i >= 0; i--) {
        oss << std::setw(BASE_DIGITS) << std::setfill('0') << digits[i];
    }
    return oss.str();
}

int64_t to_llong() const {
    std::stringstream ss(to_string());
    int64_t res;
    ss >> res;
    return res;
}

double to_double() const {
    std::stringstream ss(to_string());
    double res;
    ss >> res;
    return res;
}

long double to_ldouble() const {
    std::stringstream ss(to_string());
    long double res;
    ss >> res;
    return res;
}

explicit operator int() const { return static_cast<int>(to_llong()); }
explicit operator long long() const { return static_cast<long long>(to_llong()); }
explicit operator double() const { return to_double(); }
explicit operator long double() const { return to_ldouble(); }

int comp(const BigInt &v) const { return comp(digits, v.digits, sign, v.sign); }

// The comparison and binary arithmetic operators are hidden friends, so a raw integer operand on
// either side converts through the implicit constructor.
friend bool operator<(const BigInt &a, const BigInt &b) { return a.comp(b) < 0; }
friend bool operator>(const BigInt &a, const BigInt &b) { return a.comp(b) > 0; }
friend bool operator<=(const BigInt &a, const BigInt &b) { return a.comp(b) <= 0; }
friend bool operator>=(const BigInt &a, const BigInt &b) { return a.comp(b) >= 0; }
friend bool operator==(const BigInt &a, const BigInt &b) { return a.comp(b) == 0; }
friend bool operator!=(const BigInt &a, const BigInt &b) { return a.comp(b) != 0; }

```

```

BigInt abs() const {
    BigInt res(*this);
    res.sign = 1;
    return res;
}

BigInt operator-() const {
    BigInt res(*this);
    if (!digits.empty()) res.sign = -sign;
    return res;
}

friend BigInt operator+(const BigInt &a, const BigInt &b) {
    return add(a.digits, b.digits, a.sign, b.sign);
}

friend BigInt operator-(const BigInt &a, const BigInt &b) {
    return sub(a.digits, b.digits, a.sign, b.sign);
}

BigInt &operator*=(int v) {
    int64_t factor = v;
    if (factor < 0) {
        sign = -sign;
        factor = -factor;
    }
    int64_t carry = 0;
    for (int i = 0; i < static_cast<int>(digits.size()) || carry; i++) {
        if (i == static_cast<int>(digits.size())) {
            digits.push_back(0);
        }
        int64_t curr = digits[i] * factor + carry;
        carry = curr / BASE;
        digits[i] = static_cast<int>((curr % BASE));
    }
    normalize();
    return *this;
}

BigInt operator*(int v) const {
    BigInt res(*this);
    res *= v;
    return res;
}

friend BigInt operator*(const BigInt &u, const BigInt &v) {
    static const int TEMP_BASE = 10000, TEMP_BASE_DIGITS = 4;
    vint a = convert_base(u.digits, BASE_DIGITS, TEMP_BASE_DIGITS);
    vint b = convert_base(v.digits, BASE_DIGITS, TEMP_BASE_DIGITS);
    int n = 1;
    while (n < 2 * static_cast<int>(std::max(a.size(), b.size()))) {
        n <= 1;
    }
    a.resize(n, 0);
    b.resize(n, 0);
    vint64 c;
    if (USE_FFT_MULT) {
        auto at = fft(a.begin(), a.end()), bt = fft(b.begin(), b.end());
        for (int i = 0; i < n; i++) {
            at[i] *= bt[i];
        }
        at = fft(at.begin(), at.end(), true);
        c.resize(n);
        for (int i = 0; i < n; i++) {
            c[i] = at[i].real() + 0.5;
        }
    } else {

```

```

    c = karatsuba(a.begin(), a.end(), b.begin(), b.end());
}
BigInt res;
res.sign = u.sign * v.sign;
for (int i = 0, carry = 0; i < static_cast<int>(c.size()); i++) {
    int64_t d = c[i] + carry;
    res.digits.push_back(d % TEMP_BASE);
    carry = d / TEMP_BASE;
}
res.digits = convert_base(res.digits, TEMP_BASE_DIGITS, BASE_DIGITS);
res.normalize();
return res;
}

BigInt &operator/=(int v) {
    assert(v != 0);
    int64_t divisor = v;
    if (divisor < 0) {
        sign = -sign;
        divisor = -divisor;
    }
    int64_t rem = 0;
    for (int i = static_cast<int>(digits.size()) - 1; i >= 0; i--) {
        int64_t curr = digits[i] + rem * static_cast<int64_t>(BASE);
        digits[i] = static_cast<int>((curr / divisor));
        rem = curr % divisor;
    }
    normalize();
    return *this;
}

BigInt operator/(int v) const {
    BigInt res(*this);
    res /= v;
    return res;
}

int operator%(int v) const {
    assert(v != 0);
    int64_t divisor = v;
    if (divisor < 0) {
        divisor = -divisor;
    }
    int64_t m = 0;
    for (int i = static_cast<int>(digits.size()) - 1; i >= 0; i--) {
        m = (digits[i] + m * static_cast<int64_t>(BASE)) % divisor;
    }
    return static_cast<int>(m * sign);
}

std::pair<BigInt, BigInt> div(const BigInt &v) const {
    assert(v != 0);
    if (comp(digits, v.digits, 1, 1) < 0) {
        return {BigInt(0), *this};
    }
    int norm = BASE / (v.digits.back() + 1);
    BigInt an = abs() * norm, bn = v.abs() * norm, q, r;
    q.digits.resize(an.digits.size());
    for (int i = static_cast<int>(an.digits.size()) - 1; i >= 0; i--) {
        r *= BASE;
        r += an.digits[i];
        int s1 = (r.digits.size() <= bn.digits.size()) ? 0 : r.digits[bn.digits.size()];
        int s2 = (r.digits.size() <= bn.digits.size() - 1) ? 0 : r.digits[bn.digits.size() - 1];
        int d = (static_cast<int64_t>(s1) * BASE + s2) / bn.digits.back();
        for (r -= bn * d; r < 0; r += bn) {
            d--;
        }
        q.digits[i] = d;
    }
}

```

```

    }
    q.sign = sign * v.sign;
    r.sign = sign;
    q.normalize();
    r.normalize();
    return {q, r / norm};
}

friend BigInt operator/(const BigInt &a, const BigInt &b) { return a.div(b).first; }
friend BigInt operator%(const BigInt &a, const BigInt &b) { return a.div(b).second; }
BigInt operator++(int){ BigInt t(*this); operator++(); return t; }
BigInt operator--(int){ BigInt t(*this); operator--(); return t; }
BigInt &operator++() { *this = *this + BigInt(1); return *this; }
BigInt &operator--() { *this = *this - BigInt(1); return *this; }
BigInt &operator+=(const BigInt &v) { *this = *this + v; return *this; }
BigInt &operator-=(const BigInt &v) { *this = *this - v; return *this; }
BigInt &operator*=(const BigInt &v) { *this = *this * v; return *this; }
BigInt &operator/=(const BigInt &v) { *this = *this / v; return *this; }
BigInt &operator%=(const BigInt &v) { *this = *this % v; return *this; }

BigInt pow(int n) const {
    assert(n >= 0);
    if (n == 0) {
        return BigInt(1);
    }
    if (*this == 0) {
        return BigInt(0);
    }
    BigInt x(*this), res(1);
    for (; n != 0; n >>= 1) {
        if (n & 1) {
            res *= x;
        }
        x *= x;
    }
    return res;
}

BigInt sqrt() const {
    assert(sign != -1);
    BigInt v(*this);
    while (v.digits.empty() || v.digits.size() % 2 == 1) {
        v.digits.push_back(0);
    }
    int n = static_cast<int>(v.digits.size());
    int ldig =
        static_cast<int> (::sqrt(static_cast<double>(v.digits[n - 1]) * BASE + v.digits[n - 2]));
    int norm = BASE / (ldig + 1);
    v *= norm;
    v *= norm;
    while (v.digits.empty() || v.digits.size() % 2 == 1) {
        v.digits.push_back(0);
    }
    BigInt r(static_cast<int64_t>(v.digits[n - 1]) * BASE + v.digits[n - 2]);
    int q = ldig =
        static_cast<int> (::sqrt(static_cast<double>(v.digits[n - 1]) * BASE + v.digits[n - 2]));
    BigInt res;
    for (int j = n / 2 - 1; j >= 0; j--) {
        for (; q-- > 0) {
            BigInt r1 =
                (r - (res * 2 * BASE + q) * q) * BASE * BASE +
                (j > 0 ? static_cast<int64_t>(v.digits[2 * j - 1]) * BASE + v.digits[2 * j - 2] : 0);
            if (r1 >= 0) {
                r = r1;
                break;
            }
        }
    }
    res = res * BASE + q;
}

```

```

    if (j > 0) {
        int sz1 = static_cast<int>(res.digits.size());
        int sz2 = static_cast<int>(r.digits.size());
        int d1 = (sz1 + 2 < sz2) ? r.digits[sz1 + 2] : 0;
        int d2 = (sz1 + 1 < sz2) ? r.digits[sz1 + 1] : 0;
        int d3 = (sz1 < sz2) ? r.digits[sz1] : 0;
        q = (static_cast<int64_t>(d1) * BASE * BASE + static_cast<int64_t>(d2) * BASE + d3) /
            (ldig * 2);
    }
}
res.normalize();
return res / norm;
}

BigInt nth_root(int n) const {
    assert(n > 0 && (sign != -1 || n % 2 == 1));
    if (*this == 0) {
        return BigInt(0);
    }
    BigInt magnitude = abs();
    if (n >= size()) {
        int p = 1;
        while (magnitude.comp(BigInt(p).pow(n)) > 0) {
            p++;
        }
        BigInt root = magnitude.comp(BigInt(p).pow(n)) < 0 ? p - 1 : p;
        return sign == -1 ? -root : root;
    }
    BigInt lo(BigInt(10).pow(static_cast<int>(ceil(static_cast<double>(size()) / n) - 1)),
        hi(lo * 10), mid;
    while (lo < hi) {
        mid = (lo + hi) / 2;
        int cmp = comp(digits, mid.pow(n).digits, 1, 1);
        if (lo < mid && cmp > 0) {
            lo = mid;
        } else if (mid < hi && cmp < 0) {
            hi = mid;
        } else {
            return (sign == -1) ? -mid : mid;
        }
    }
    return (sign == -1) ? -(mid + 1) : (mid + 1);
}

static BigInt rand(int n) {
    if (n == 0) {
        return BigInt(0);
    }
    static std::mt19937 rng(std::random_device{}());
    std::uniform_int_distribution<int> first_digit(1, 9), digit(0, 9);
    std::string s(1, static_cast<char>('0' + first_digit(rng)));
    for (int i = 1; i < n; i++) {
        s += static_cast<char>('0' + digit(rng));
    }
    return BigInt(s);
}

friend int comp(const BigInt &a, const BigInt &b) { return a.comp(b); }
friend BigInt abs(const BigInt &v) { return v.abs(); }
friend BigInt pow(const BigInt &v, int n) { return v.pow(n); }
friend BigInt sqrt(const BigInt &v) { return v.sqrt(); }
friend BigInt nth_root(const BigInt &v, int n) { return v.nth_root(n); }
};

```

Example Usage

```

int main() {
    BigInt a("-9899819294989142124"), b("12398124981294214");
    assert(a + b == "-9887421170007847910");
    assert(a - b == "-9912217419970436338");
    assert(a * b == "-122739196911503356525379735104870536");
    assert(a / b == "-798");

    // Raw integers and strings work on either side of comparisons and arithmetic.
    assert("12398124981294214" == b);
    assert(5 + BigInt(7) == 12);
    assert(-100 < b && 5 % BigInt(3) == 2);

    assert(BigInt(20).pow(12345).size() == 16062);
    assert(BigInt("9812985918924981892491829").nth_root(4) == 1769906);
    assert(BigInt(-8).nth_root(3) == -2);
    for (int i = -100; i <= 100; i++) {
        if (i >= 0) {
            assert(BigInt(i).sqrt() == static_cast<int>(sqrt(i)));
        }
        for (int j = -100; j <= 100; j++) {
            assert(BigInt(i) + BigInt(j) == i + j);
            assert(BigInt(i) - BigInt(j) == i - j);
            assert(BigInt(i) * BigInt(j) == i * j);
            if (j != 0) {
                assert(BigInt(i) / BigInt(j) == i / j);
            }
            if (0 < i && i <= 10 && 0 < j && j <= 10) {
                assert(BigInt(i).nth_root(j) == static_cast<int64_t>((pow(i, 1.0 / j) + 1E-5)));
                int64_t p = 1;
                for (int k = 0; k < j; k++) {
                    p *= i;
                }
                assert(BigInt(i).pow(j) == p);
            }
        }
    }
    std::mt19937 rng(1234567); // Fixed seed for reproducibility.
    std::uniform_int_distribution<int> length_dist(1, 100);
    for (int i = 0; i < 20; i++) {
        int n = length_dist(rng);
        BigInt a(BigInt::rand(n)), s(a.sqrt()), xx(s * s), yy(s + 1);
        yy *= yy;
        assert(xx <= a && a < yy);
        std::uniform_int_distribution<int> divisor_length_dist(1, n);
        BigInt b(BigInt::rand(divisor_length_dist(rng)) + 1), q(a / b);
        xx = q * b;
        yy = b * (q + 1);
        assert(a >= xx && a < yy);
    }
    BigInt x(-6);
    assert(x.to_string() == "-6");
    assert(x.to_llong() == -6LL);
    assert(x.to_double() == -6.0);
    assert(x.to_ldouble() == -6.0);
    assert(static_cast<int>(x) == -6);
    assert(static_cast<long long>(x) == -6LL);
    assert(static_cast<double>(x) == -6.0);
    assert(BigInt(INT64_MIN).to_string() == "-9223372036854775808");
    assert(BigInt(INT64_MIN).to_llong() == INT64_MIN);
    assert(-BigInt(0) == 0);
    assert(BigInt(3) * INT_MIN == "-6442450944");
    assert(BigInt("6442450944") / INT_MIN == -3);
    assert(BigInt("6442450945") % INT_MIN == 1);
    return 0;
}

```

6.4.3 Rational Numbers

Perform operations on rational numbers internally represented as two integers: a numerator and a denominator. The template integer type must support streamed input/output, comparisons, and arithmetic operations. Overflow is not checked for in internal operations: comparisons and arithmetic cross-multiply numerators and denominators, so instantiate with a wider integer type (such as `__int128`, or even `BigInt`) if the values may grow large.

- `Rational<Int>(n)` constructs a rational number with numerator `n` and denominator 1.
- `Rational<Int>(n, d)` constructs a rational number with numerator `n` and denominator `d`.
- `operator>>` inputs a rational number using the next integer from the stream as the numerator and 1 as the denominator.
- `operator<<` outputs the rational number as a string consisting of possibly a minus sign followed by the numerator, followed by a slash, followed by the denominator.
- `v.to_string()`, `v.to_llong()`, `v.to_double()`, and `v.to_ldouble()` return the rational `v` converted to an `std::string`, `int64_t`, `double`, and `long double` respectively with potential truncation or approximation for the primitive types. Equivalent conversions are also available through explicit casts to `int`, `long long`, `double`, and `long double`.
- `v.abs()`, `abs(v)`, `v.floor()`, and `v.ceil()` return the absolute value, floor, and ceiling of rational `v`.
- Operators `<`, `>`, `<=`, `>=`, `==`, `!=`, `+`, `-`, `*`, `/`, `%`, `++`, `-`, `+=`, `-=`, `*=`, `/=`, and `%=` are defined analogous to those on numerical primitives. The comparisons and binary arithmetic are hidden friends, so a raw integer operand works on either side.

Time Complexity:

- $O(\log s)$ per call to constructor `Rational(n, d)`, where $s = |n| + |d|$.
- $O(1)$ per call to all other operations, assuming that corresponding operations on the template integer type are $O(1)$ as well.

Space Complexity:

- $O(1)$ for storage of the rational number.
- $O(1)$ auxiliary for all operations.

```
#include <cassert>
#include <cstdint>
#include <istream>
#include <ostream>
#include <sstream>
#include <string>

template<typename Int>
class Rational {
    Int num, den;

public:
    Rational() : num(0), den(1) {}
    Rational(const Int &n) : num(n), den(1) {}

    template<typename T1, typename T2>
    Rational(const T1 &n, const T2 &d) : num(n), den(d) {
        assert(den != 0);
        if (den < 0) {
            num = -num;
            den = -den;
        }
        Int a(num < 0 ? -num : num), b(den), tmp;
        while (a != 0 && b != 0) {
            tmp = a % b;
            a = b;
            b = tmp;
        }
        Int gcd = (b == 0) ? a : b;
```

```

    num /= gcd;
    den /= gcd;
}

friend std::istream &operator>>(std::istream &in, Rational &r) {
    in >> r.num;
    r.den = 1;
    return in;
}

friend std::ostream &operator<<(std::ostream &out, const Rational &r) {
    out << r.num << "/" << r.den;
    return out;
}

std::string to_string() const {
    std::stringstream ss;
    ss << num << "/" << den;
    return ss.str();
}

// to_llong/to_double/to_ldouble round-trip through a stringstream so that any Int supporting
// streamed I/O (e.g. a big-integer type) converts, even without a direct cast to the target type.
int64_t to_llong() const {
    std::stringstream ss;
    ss << num << " " << den;
    int64_t n, d;
    ss >> n >> d;
    return n / d;
}

double to_double() const {
    std::stringstream ss;
    ss << num << " " << den;
    double n, d;
    ss >> n >> d;
    return n / d;
}

long double to_ldouble() const {
    std::stringstream ss;
    ss << num << " " << den;
    long double n, d;
    ss >> n >> d;
    return n / d;
}

explicit operator int() const { return static_cast<int>(to_llong()); }
explicit operator long long() const { return static_cast<long long>(to_llong()); }
explicit operator double() const { return to_double(); }
explicit operator long double() const { return to_ldouble(); }

// The comparison and binary arithmetic operators are hidden friends, so a raw integer operand on
// either side converts through the implicit constructor.
friend bool operator<(const Rational &a, const Rational &b) {
    return a.num * b.den < b.num * a.den;
}

friend bool operator==(const Rational &a, const Rational &b) {
    return a.num == b.num && a.den == b.den;
}

friend bool operator>(const Rational &a, const Rational &b) { return b < a; }
friend bool operator<=(const Rational &a, const Rational &b) { return !(b < a); }
friend bool operator>=(const Rational &a, const Rational &b) { return !(a < b); }
friend bool operator!=(const Rational &a, const Rational &b) { return !(a == b); }

Rational abs() const { return Rational(num < 0 ? -num : num, den); }

```



```

friend Rational abs(const Rational &r) { return r.abs(); }
Int floor() const { return num < 0 ? -((-num + den - 1) / den) : num / den; }
Int ceil() const { return num < 0 ? -(-num / den) : (num + den - 1) / den; }

friend Rational operator+(const Rational &a, const Rational &b) {
    return Rational(a.num * b.den + b.num * a.den, a.den * b.den);
}

friend Rational operator-(const Rational &a, const Rational &b) {
    return Rational(a.num * b.den - b.num * a.den, a.den * b.den);
}

friend Rational operator*(const Rational &a, const Rational &b) {
    return Rational(a.num * b.num, a.den * b.den);
}

friend Rational operator/(const Rational &a, const Rational &b) {
    return Rational(a.num * b.den, a.den * b.num);
}

friend Rational operator%(const Rational &a, const Rational &b) {
    return a - b * Rational(a.num * b.den / (b.num * a.den), 1);
}

Rational operator-() const { return Rational(-num, den); }
Rational operator++(int) { Rational t(*this); operator++(); return t; }
Rational operator--(int) { Rational t(*this); operator--(); return t; }
Rational &operator++() { *this = *this + 1; return *this; }
Rational &operator--() { *this = *this - 1; return *this; }
Rational &operator+=(const Rational &r) { *this = *this + r; return *this; }
Rational &operator-=(const Rational &r) { *this = *this - r; return *this; }
Rational &operator*=(const Rational &r) { *this = *this * r; return *this; }
Rational &operator/=(const Rational &r) { *this = *this / r; return *this; }
Rational &operator%=(const Rational &r) { *this = *this % r; return *this; }
};

```

Example Usage

```

#include <cmath>

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-9;
}

int main() {
    using Rational = Rational<int64_t>;
    assert(Rational(-21, 1) % 2 == -1);
    Rational r(Rational(-53, 10) % Rational(-17, 10));
    assert(EQ(r.to_ldouble(), fmod(-5.3, -1.7)));
    assert(r.to_string() == "-1/5");
    assert(static_cast<int>(Rational(7, 2)) == 3);
    assert(static_cast<long long>(Rational(7, 2)) == 3LL);
    assert(EQ(static_cast<double>(Rational(1, 2)), 0.5));

    // Raw integers work on either side of comparisons and arithmetic.
    assert(2 + Rational(1, 2) == Rational(5, 2));
    assert(1 < Rational(3, 2) && Rational(3, 2) < 2);
    assert(2 % Rational(3, 4) == Rational(1, 2));
    return 0;
}

```

6.5 Linear Algebra

6.5.1 Matrix Utilities

Basic matrix operations defined on a two-dimensional vector of numeric values.

- `make_matrix<T>(m, n, v)` constructs and returns an m by n matrix with 0-based indices (row indices $[0, m)$ and column indices $[0, n)$), where every value is initialized to `v`.
- `identity_matrix<T>(n)` returns the n by n identity matrix, that is, a matrix where each `a[i][j]` equals 1 if $i = j$, or 0 otherwise.
- `rows(a)` returns the number of rows in matrix `a`.
- `columns(a)` returns the number of columns in matrix `a`.
- `a[i][j]` may be used to access or modify the specified entry of `a`.
- Operators `<`, `>`, `<=`, `>=`, `==`, and `!=` define lexicographical comparison based on that of `std::vector`.
- Operators `+`, `-`, `*`, `/`, `+=`, `-=`, `*=`, and `/=` define scalar addition, subtraction, multiplication, and division involving a matrix and a numeric scalar value.
- Operators `*` and `*=` define vector and matrix multiplication.
- Operators `^` and `^=` define iterative binary matrix exponentiation of a square matrix `a` by an integer power `p`.
- `power_sum(a, p)` returns the power sum of a square matrix `a` up to an integer power `p`, that is, $a + a^2 + \dots + a^p$.
- `transpose(a)` returns the transpose of an m by n matrix `a`, that is, a new n by m matrix `b` such that `a[i][j] = b[j][i]` for every $i \in [0, m)$ and $j \in [0, n)$.
- `transpose_in_place(a)` assigns the square matrix `a` to its transpose, returning a reference to the modified argument itself.
- `rotate(a, d = 90)` returns the matrix `a` rotated `d` degrees clockwise. A negative `d` specifies a counter-clockwise rotation, and `d` must be a multiple of 90.
- `rotate_in_place(a, d = 90)` assigns the square matrix `a` to its rotation by `d` degrees clockwise, returning a reference to the modified argument itself. A negative `d` specifies a counter-clockwise rotation, and `d` must be a multiple of 90.

Exponentiation uses the usual iterative binary exponentiation pattern: keep an accumulated result, square the base each round, and multiply the result when the current exponent bit is set. The power sum uses a block matrix identity that $\begin{bmatrix} A & A \\ 0 & I \end{bmatrix}^p$ has $A + A^2 + \dots + A^p$ in its upper-right block.

Time Complexity:

- $O(m \cdot n)$ per construction, output, comparison, or scalar-arithmetic operation on m by n matrices.
- $O(1)$ per call to `rows()` and `columns()`.
- $O(m \cdot n)$ per matrix-matrix addition or subtraction of m by n matrices.
- $O(n^3 \log p)$ per exponentiation of an n by n matrix to power p .
- $O(n^3 \log p)$ per call to `power_sum()` for an n by n matrix and power p .
- $O(m \cdot n \cdot k)$ per multiplication of an m by n matrix by an n by k matrix.
- $O(m \cdot n)$ per call to `transpose()`, `transpose_in_place()`, `rotate()`, and `rotate_in_place()` on an m by n matrix.

Space Complexity:

- $O(1)$ auxiliary for `rows()`, `columns()`, `a[i][j]` access, comparison and scalar compound operators, `transpose_in_place()`, and `rotate_in_place()`.
- $O(n^2)$ auxiliary for exponentiation of an n by n matrix to power p .
- $O(n^2)$ auxiliary for `power_sum()` of an n by n matrix up to power p .
- $O(m \cdot n)$ auxiliary for all non-in-place operations returning an m by n matrix, `transpose()`, and `rotate()`.

```
#include <algorithm>
#include <cassert>
#include <iomanip>
```

```

#include <ostream>
#include <utility>
#include <vector>

template<typename T>
using Matrix = std::vector<std::vector<T>>>;

template<typename T = int>
Matrix<T> make_matrix(int m, int n, const T &v = T{}) {
    return Matrix<T>(m, std::vector<T>(n, v));
}

template<typename T = int>
Matrix<T> identity_matrix(int n) {
    Matrix<T> res = make_matrix<T>(n, n);
    for (int i = 0; i < n; i++) {
        res[i][i] = 1;
    }
    return res;
}

template<typename T>
int rows(const Matrix<T> &a) {
    return static_cast<int>(a.size());
}

template<typename T>
int columns(const Matrix<T> &a) {
    return a.empty() ? 0 : static_cast<int>(a[0].size());
}

template<typename T>
std::ostream &operator<<(std::ostream &out, const Matrix<T> &a) {
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(a); j++) {
            out << std::setw(10) << std::fixed << std::setprecision(5) << a[i][j];
        }
        out << std::endl;
    }
    return out;
}

template<typename T, typename U>
Matrix<T> &operator+=(Matrix<T> &a, const U &v) {
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(a); j++) {
            a[i][j] += v;
        }
    }
    return a;
}

template<typename T, typename U>
Matrix<T> &operator-=(Matrix<T> &a, const U &v) {
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(a); j++) {
            a[i][j] -= v;
        }
    }
    return a;
}

template<typename T, typename U>
Matrix<T> &operator*=(Matrix<T> &a, const U &v) {
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(a); j++) {
            a[i][j] *= v;
        }
    }
}

```

```

    }
    return a;
}

template<typename T, typename U>
Matrix<T> &operator/=(Matrix<T> &a, const U &v) {
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(a); j++) {
            a[i][j] /= v;
        }
    }
    return a;
}

template<typename T>
Matrix<T> &operator+=(Matrix<T> &a, const Matrix<T> &b) {
    assert(rows(a) == rows(b) && columns(a) == columns(b));
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(a); j++) {
            a[i][j] += b[i][j];
        }
    }
    return a;
}

template<typename T>
Matrix<T> &operator-=(Matrix<T> &a, const Matrix<T> &b) {
    assert(rows(a) == rows(b) && columns(a) == columns(b));
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(a); j++) {
            a[i][j] -= b[i][j];
        }
    }
    return a;
}

template<typename T>
Matrix<T> operator+(const Matrix<T> &a, const Matrix<T> &b) {
    Matrix<T> c(a);
    return c += b;
}

template<typename T>
Matrix<T> operator-(const Matrix<T> &a, const Matrix<T> &b) {
    Matrix<T> c(a);
    return c -= b;
}

template<typename T>
Matrix<T> &operator*=(Matrix<T> &a, const std::vector<T> &v) {
    assert(columns(a) == static_cast<int>(v.size()) && !v.empty());
    int cols = columns(a);
    for (int i = 0; i < rows(a); i++) {
        T sum = 0;
        for (int j = 0; j < cols; j++) {
            sum += a[i][j] * v[j];
        }
        a[i].assign(1, sum);
    }
    return a;
}

template<typename T>
Matrix<T> operator*(const Matrix<T> &a, const Matrix<T> &b) {
    assert(columns(a) == rows(b));
    Matrix<T> res = make_matrix<T>(rows(a), columns(b));
    for (int i = 0; i < rows(a); i++) {
        for (int j = 0; j < columns(b); j++) {

```

```

        for (int k = 0; k < rows(b); k++) {
            res[i][j] += a[i][k] * b[k][j];
        }
    }
}
return res;
}

template<typename T>
Matrix<T> &operator+=(Matrix<T> &a, const Matrix<T> &b) { return a = a * b; }

template<typename T, typename U>
Matrix<T> operator+(const Matrix<T> &a, const U &v) { Matrix<T> m(a); return m += v; }

template<typename T, typename U>
Matrix<T> operator-(const Matrix<T> &a, const U &v) { Matrix<T> m(a); return m -= v; }

template<typename T, typename U>
Matrix<T> operator*(const Matrix<T> &a, const U &v) { Matrix<T> m(a); return m *= v; }

template<typename T, typename U>
Matrix<T> operator/(const Matrix<T> &a, const U &v) { Matrix<T> m(a); return m /= v; }

template<typename T, typename U>
Matrix<T> operator+(const U &v, const Matrix<T> &a) { return a + v; }

template<typename T, typename U>
Matrix<T> operator*(const U &v, const Matrix<T> &a) { return a * v; }

template<typename T, typename U>
Matrix<T> operator-(const U &v, const Matrix<T> &a) {
    Matrix<T> m = make_matrix<T>(rows(a), columns(a), v);
    return m -= a;
}

template<typename T>
Matrix<T> operator^(Matrix<T> a, unsigned int p) {
    assert(rows(a) == columns(a));
    Matrix<T> res = identity_matrix<T>(rows(a));
    while (p > 0) {
        if (p & 1) {
            res *= a;
        }
        a *= a;
        p >>= 1;
    }
    return res;
}

template<typename T>
Matrix<T> &operator^=(Matrix<T> &a, unsigned int p) {
    return a = a ^ p;
}

template<typename T>
Matrix<T> power_sum(const Matrix<T> &a, unsigned int p) {
    assert(rows(a) == columns(a));
    int n = rows(a);
    if (p == 0) {
        return make_matrix<T>(n, n);
    }
    // The upper-right block of  $[[A, A], [0, I]]^p$  is  $A + A^2 + \dots + A^p$ .
    Matrix<T> block = make_matrix<T>(2 * n, 2 * n);
    for (int i = 0; i < n; i++) {
        block[i + n][i + n] = 1;
        for (int j = 0; j < n; j++) {
            block[i][j] = a[i][j];
            block[i][j + n] = a[i][j];
        }
    }
}

```

```

    }
}
block = block ^ p;
Matrix<T> res = make_matrix<T>(n, n);
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        res[i][j] = block[i][j + n];
    }
}
return res;
}

template<typename T>
Matrix<T> transpose(const Matrix<T> &a) {
    Matrix<T> res = make_matrix<T>(columns(a), rows(a));
    for (int i = 0; i < rows(res); i++) {
        for (int j = 0; j < columns(res); j++) {
            res[i][j] = a[j][i];
        }
    }
    return res;
}

template<typename T>
Matrix<T> &transpose_in_place(Matrix<T> &a) {
    assert(rows(a) == columns(a));
    for (int i = 0; i < rows(a); i++) {
        for (int j = i + 1; j < columns(a); j++) {
            std::swap(a[i][j], a[j][i]);
        }
    }
    return a;
}

template<typename T>
Matrix<T> rotate(const Matrix<T> &a, int degrees = 90) {
    assert(degrees % 90 == 0);
    if (degrees < 0) {
        degrees = 360 - ((-degrees) % 360);
    }
    Matrix<T> res;
    switch (degrees % 360) {
        case 90: {
            res = make_matrix<T>(columns(a), rows(a));
            for (int i = 0; i < columns(a); i++) {
                for (int j = 0; j < rows(a); j++) {
                    res[i][j] = a[rows(a) - j - 1][i];
                }
            }
            break;
        }
        case 180: {
            res = make_matrix<T>(rows(a), columns(a));
            for (int i = 0; i < rows(a); i++) {
                for (int j = 0; j < columns(a); j++) {
                    res[i][j] = a[rows(a) - i - 1][columns(a) - j - 1];
                }
            }
            break;
        }
        case 270: {
            res = make_matrix<T>(columns(a), rows(a));
            for (int i = 0; i < columns(a); i++) {
                for (int j = 0; j < rows(a); j++) {
                    res[i][j] = a[j][columns(a) - i - 1];
                }
            }
            break;
        }
    }
}

```

```

    }
    default: {
        res = a;
    }
}
return res;
}

template<typename T>
Matrix<T> &rotate_in_place(Matrix<T> &a, int degrees = 90) {
    assert(degrees % 90 == 0);
    assert(degrees % 180 == 0 || rows(a) == columns(a));
    if (degrees < 0) {
        degrees = 360 - ((-degrees) % 360);
    }
    int n = rows(a);
    switch (degrees % 360) {
        case 90: {
            transpose_in_place(a);
            for (auto &row : a) {
                std::reverse(row.begin(), row.end());
            }
            break;
        }
        case 180: {
            // A 180-degree rotation is a horizontal flip (reverse each row) followed by a vertical flip
            // (reverse the row order). This works for non-square matrices too, unlike a row/column index
            // scheme keyed on a single dimension.
            for (auto &row : a) {
                std::reverse(row.begin(), row.end());
            }
            std::reverse(a.begin(), a.end());
            break;
        }
        case 270: {
            transpose_in_place(a);
            for (int j = 0; j < n; j++) {
                for (int i = 0, k = columns(a) - 1; i < k; i++, k--) {
                    std::swap(a[i][j], a[k][j]);
                }
            }
            break;
        }
    }
    return a;
}

```

Example Usage

```

using namespace std;

int main() {
    using matrix = Matrix<int>;
    matrix a{{1, 2, 3}, {4, 5, 6}};
    matrix a90{{4, 1}, {5, 2}, {6, 3}};
    matrix a180{{6, 5, 4}, {3, 2, 1}};
    matrix a270{{3, 6}, {2, 5}, {1, 4}};
    assert(rotate(a, -270) == a90);
    assert(rotate(a, -180) == a180);
    assert(rotate(a, -90) == a270);
    assert(rotate(a, 0) == a);
    assert(rotate(a, 90) == a90);
    assert(rotate(a, 180) == a180);
    assert(rotate(a, 270) == a270);
    assert(rotate(a, 360) == a);
    assert(rotate(a, 450) == a90);
}

```

```

matrix b{{1, 2, 3}, {4, 5, 6}, {7, 8, 9}};
for (int d = -360; d <= 360; d += 90) {
    matrix m = b;
    assert(rotate_in_place(m, d) == rotate(b, d));
}

matrix m = make_matrix(5, 5, 10) + 10;
vector<int> v{1, 2, 3, 4, 5};
matrix mv{{300}, {300}, {300}, {300}, {300}};
assert(m * v == mv);

m[0][0] += 5;
assert(m[0][0] == 25 && m[1][1] == 20);
assert((m ^ 0) == identity_matrix(5));
assert((identity_matrix<int>(3) ^ 5) == identity_matrix<int>(3));
assert(power_sum(m, 3) == m + m * m + (m ^ 3));

Matrix<double> d = make_matrix<double>(2, 2, 0.5);
assert(rows(d) == 2 && columns(d) == 2);
assert((d + 0.25)[0][0] == 0.75);
return 0;
}

```

6.5.2 Row Reduction

Converts a matrix to reduced row echelon form using Gaussian elimination to solve a system of linear equations and compute rank. Each round finds a row with a nonzero entry in the current leading column, swaps the largest-magnitude candidate into place for partial pivoting, normalizes that row, and subtracts multiples of it from every other row to clear the column. In practice, this method can still be prone to rounding error on certain matrices. For a more accurate algorithm for solving systems of linear equations, LU decomposition with row partial pivoting should be used.

- `row_reduce(a)` assigns the matrix `a` to its reduced row echelon form, returning a reference to the modified argument itself.
- `matrix_rank(a)` returns the rank of matrix `a`, i.e. the number of nonzero rows after row reduction.
- `solve_system(a, b, &x)` solves the system of linear equations $ax = b$ given an m by n matrix `a` of real values, and a length m vector `b`, returning 0 if there is one solution, -1 if there are zero solutions, or -2 if there are infinite solutions. If there is exactly one solution, then the output pointer `x` is populated with the solution vector of length n .

Time Complexity:

- $O(m \cdot n \cdot \min(m, n))$ per call to `row_reduce()` and `matrix_rank()`, where m and n are the numbers of rows and columns of `a`, respectively.
- $O(m \cdot n \cdot \min(m, n))$ per call to `solve_system()`.

Space Complexity:

- $O(1)$ auxiliary for `row_reduce()`.
- $O(m \cdot n)$ auxiliary for `matrix_rank()` and `solve_system()`.

```

#include <cassert>
#include <cmath>
#include <utility>
#include <vector>

const double EPS = 1e-9;

template<typename Matrix>
Matrix &row_reduce(Matrix &a) {
    if (a.empty()) {

```



```

    return a;
}
int rows = static_cast<int>(a.size()), cols = static_cast<int>(a[0].size());
for (int r = 0, lead = 0; r < rows && lead < cols; lead++) {
    int pivot = r;
    for (int i = r + 1; i < rows; i++) {
        if (fabs(a[i][lead]) > fabs(a[pivot][lead])) {
            pivot = i;
        }
    }
    if (fabs(a[pivot][lead]) < EPS) {
        continue;
    }
    std::swap(a[pivot], a[r]);
    auto lv = a[r][lead];
    for (int j = 0; j < cols; j++) {
        a[r][j] /= lv;
    }
    for (int i = 0; i < rows; i++) {
        if (i != r) {
            lv = a[i][lead];
            for (int j = 0; j < cols; j++) {
                a[i][j] -= lv * a[r][j];
            }
        }
    }
    for (int j = 0; j < lead; j++) {
        a[r][j] = 0;
    }
    a[r][lead] = 1;
    r++;
}
return a;
}

template<typename Matrix>
int matrix_rank(Matrix a) {
    row_reduce(a);
    int rank = 0;
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        for (int j = 0; j < static_cast<int>(a[i].size()); j++) {
            if (fabs(a[i][j]) > EPS) {
                rank++;
                break;
            }
        }
    }
    return rank;
}

template<typename Matrix, typename T>
int solve_system(const Matrix &a, const std::vector<T> &b, std::vector<T> *x) {
    if (x == nullptr || a.empty() || a.size() != b.size()) {
        return -1;
    }
    int rows = static_cast<int>(a.size()), cols = static_cast<int>(a[0].size());
    Matrix m(a);
    for (int i = 0; i < rows; i++) {
        m[i].push_back(b[i]);
    }
    row_reduce(m);
    int rank = 0;
    for (int i = 0; i < rows; i++) {
        int lead = -1;
        for (int j = 0; j < cols && lead < 0; j++) {
            if (fabs(m[i][j]) > EPS) {
                lead = j;
            }
        }
    }
}

```

```

    }
    if (lead < 0) {
        // A zero coefficient row with a nonzero constant means 0 = nonzero (no solution).
        if (fabs(m[i][cols]) > EPS) {
            return -1;
        }
    } else {
        rank++;
    }
}
// A unique solution requires a pivot in every column; fewer means free variables remain. This
// also catches trailing free columns, which a per-row "lead > i" test alone would miss.
if (rank < cols) {
    return -2;
}
x->resize(cols);
for (int i = 0; i < cols; i++) {
    (*x)[i] = m[i][cols];
}
return 0;
}

```

Example Usage

```

using namespace std;

int main() {
    vector<vector<double>> a{{-1, 2, 5}, {1, 0, -6}, {-4, 2, 2}};
    vector<double> b{3, 1, -2};
    vector<double> x;
    assert(solve_system(a, b, &x) == 0);
    assert(matrix_rank(a) == 3);
    assert(matrix_rank(vector<vector<double>>{{1, 2, 3}, {2, 4, 6}, {0, 1, 1}}) == 2);
    assert(solve_system(vector<vector<double>>{{0, 0}}, vector<double>{1}, &x) == -1);
    assert(solve_system(vector<vector<double>>{{1, 0}}, vector<double>{1}, &x) == -2);
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        double sum = 0;
        for (int j = 0; j < static_cast<int>(a[i].size()); j++) {
            sum += a[i][j] * x[j];
        }
        assert(fabs(sum - b[i]) < EPS);
    }
    return 0;
}

```

6.5.3 Determinant and Inverse

Computes the determinant and inverse of a square matrix using Gaussian elimination. The inverse of a matrix a is another matrix b such that ab equals the identity matrix. The inverse of a exists if and only if the determinant of a is nonzero. In this case, a is called invertible or non-singular. The determinant falls out of elimination as the product of the pivots, negated once per row swap. The inverse is found by appending the identity matrix and row-reducing the combined matrix: the operations that turn a into the identity turn the identity into the inverse. In practice, simple Gaussian elimination is prone to rounding error on certain matrices. For a more accurate algorithm for solving systems of linear equations, use LU decomposition with row partial pivoting. For an integer matrix whose determinant is wanted exactly, the Bareiss algorithm runs a fraction-free elimination: every intermediate value is itself a determinant of a submatrix, so the divisions are always exact and the arithmetic stays in integers.

- `det_naive(a)` returns the determinant of an n by n matrix `a`, using the classic divide-and-conquer algorithm by Laplace expansions. It is division-free and computes in the matrix's element type, so an integer matrix

yields an exact integer determinant; but at $O(n!)$ it is only practical for tiny `n`, so prefer `det_bareiss` for exact integer determinants.

- `det(a, eps = 1e-10)` returns the determinant of an n by n matrix `a` using Gaussian elimination, treating pivots within `eps` of zero as singular.
- `det_bareiss(a)` returns the exact determinant of an integer matrix `a` using fraction-free elimination, with no rounding error. Stored entries are minors of `a`, but the products in each update can be larger. They use 128-bit intermediates when available; otherwise those products must fit in `int64_t`.
- `invert(a, eps = 1e-10)` returns 0 after assigning the n by n matrix `a` to its inverse. It must contain real values. The function leaves `a` unchanged and returns -1 if a pivot is within `eps` of zero.

Time Complexity:

- $O(n!)$ per call to `det_naive()`, where n is the dimension of the matrix.
- $O(n^3)$ per call to `det()`, `det_bareiss()`, and `invert()` where n is the dimension of the matrix.

Space Complexity:

- $O(n)$ auxiliary stack space and $O(n^3)$ auxiliary heap space for `det_naive()`, where n is the dimension of the matrix.
- $O(n^2)$ auxiliary heap space for `det()`, `det_bareiss()`, and `invert()`.

```
#include <cmath>
#include <cstdint>
#include <type_traits>
#include <utility>
#include <vector>

template<typename SquareMatrix>
auto det_naive(const SquareMatrix &a) {
    int n = static_cast<int>(a.size());
    using T = std::decay_t<decltype(a[0][0])>;
    if (n == 0) {
        return T(1);
    }
    if (n == 1) {
        return a[0][0];
    }
    if (n == 2) {
        return a[0][0] * a[1][1] - a[0][1] * a[1][0];
    }
    T res = 0;
    SquareMatrix temp(n - 1, typename SquareMatrix::value_type(n - 1));
    for (int p = 0; p < n; p++) {
        int h = 0, k = 0;
        for (int i = 1; i < n; i++) {
            for (int j = 0; j < n; j++) {
                if (j == p) {
                    continue;
                }
                temp[h][k++] = a[i][j];
                if (k == n - 1) {
                    h++;
                    k = 0;
                }
            }
        }
        res += (p % 2 == 0 ? 1 : -1) * a[0][p] * det_naive(temp);
    }
    return res;
}

template<typename SquareMatrix>
double det(const SquareMatrix &a, double eps = 1e-10) {
    int n = static_cast<int>(a.size());
    SquareMatrix b(a);
    double res = 1.0;
```

```

for (int i = 0; i < n; i++) {
    // Partial pivoting: bring the largest-magnitude entry in column i to the diagonal. Each row
    // swap negates the determinant, so the swap sign must be tracked. Selecting pivot rows out of
    // order without this sign correction can return |det| with the wrong sign.
    int p = i;
    for (int r = i + 1; r < n; r++) {
        if (fabs(b[r][i]) > fabs(b[p][i])) {
            p = r;
        }
    }
    if (fabs(b[p][i]) < eps) {
        return 0;
    }
    if (p != i) {
        std::swap(b[p], b[i]);
        res = -res;
    }
    res *= b[i][i];
    for (int j = i + 1; j < n; j++) {
        double z = b[j][i] / b[i][i];
        for (int k = i; k < n; k++) {
            b[j][k] -= z * b[i][k];
        }
    }
}
return res;
}

int64_t det_bareiss(std::vector<std::vector<int64_t>> a) {
    int n = static_cast<int>(a.size());
    int64_t prev = 1, sign = 1;
#ifdef __SIZEOF_INT128__
    __extension__ typedef __int128 int128_t;
#endif
    for (int k = 0; k < n; k++) {
        if (a[k][k] == 0) { // Swap in a nonzero pivot; each swap flips the sign.
            int p = -1;
            for (int r = k + 1; r < n; r++) {
                if (a[r][k] != 0) {
                    p = r;
                    break;
                }
            }
            if (p == -1) {
                return 0;
            }
            std::swap(a[k], a[p]);
            sign = -sign;
        }
        for (int i = k + 1; i < n; i++) {
            for (int j = k + 1; j < n; j++) {
#ifdef __SIZEOF_INT128__
                int128_t numerator =
                    static_cast<int128_t>(a[i][j]) * a[k][k] - static_cast<int128_t>(a[i][k]) * a[k][j];
                a[i][j] = static_cast<int64_t>(numerator / prev);
#else
                // Overflow warning.
                a[i][j] = (a[i][j] * a[k][k] - a[i][k] * a[k][j]) / prev;
#endif
            }
        }
        prev = a[k][k];
    }
    return n == 0 ? 1 : sign * a[n - 1][n - 1];
}

template<typename SquareMatrix>
int invert(SquareMatrix &a, double eps = 1e-10) {

```

```

int n = static_cast<int>(a.size());
SquareMatrix b(a);
for (int i = 0; i < n; i++) {
    b[i].resize(2 * n);
    for (int j = n; j < n * 2; j++) {
        b[i][j] = (i == j - n ? 1 : 0);
    }
}
for (int i = 0; i < n; i++) {
    // Partial pivoting: without it a zero on the diagonal divides by zero and fails even for an
    // invertible matrix (e.g. a permutation matrix). Picking the largest-magnitude pivot also
    // improves numerical stability.
    int p = i;
    for (int r = i + 1; r < n; r++) {
        if (fabs(b[r][i]) > fabs(b[p][i])) {
            p = r;
        }
    }
    if (fabs(b[p][i]) < eps) {
        return -1;
    }
    std::swap(b[p], b[i]);
    double z = b[i][i];
    for (int j = i; j < n * 2; j++) {
        b[i][j] /= z;
    }
    for (int j = 0; j < n; j++) {
        if (i != j) {
            double z = b[j][i];
            for (int k = 0; k < n * 2; k++) {
                b[j][k] -= z * b[i][k];
            }
        }
    }
}
for (int i = 0; i < n; i++) {
    b[i].erase(b[i].begin(), b[i].begin() + n);
}
a = std::move(b);
return 0;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<vector<double>> a{{6, 1, 1}, {4, -2, 5}, {2, 8, 7}};
    int n = static_cast<int>(a.size());
    vector<vector<double>> res(n, vector<double>(n, 0));
    double d = det(a);
    assert(fabs(d - det_naive(a)) < 1e-10);

    // Bareiss gives the determinant of an integer matrix exactly, with no rounding.
    vector<vector<int64_t>> ai{{6, 1, 1}, {4, -2, 5}, {2, 8, 7}};
    assert(det_naive(ai) == -306); // Division-free, so also exact in the int64_t element type.
    assert(det_naive(vector<vector<int64_t>>{}) == 1);
    assert(det_bareiss(ai) == -306);
    assert(det_bareiss({{2, 0, 0}, {0, 3, 0}, {0, 0, 5}}) == 30);
    assert(det_bareiss({{1, 2}, {2, 4}}) == 0); // Singular.
    vector<vector<double>> inv = a;
    assert(invert(inv) == 0);
    vector<vector<double>> singular{{1, 2}, {2, 4}}, unchanged = singular;
    assert(invert(singular) == -1 && singular == unchanged);
    for (int i = 0; i < n; i++) {

```

```

    for (int j = 0; j < n; j++) {
        for (int k = 0; k < n; k++) {
            res[i][j] += a[i][k] * inv[k][j];
        }
    }
}
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        assert(fabs(res[i][j] - (i == j ? 1 : 0)) < 1e-10);
    }
}
return 0;
}

```

6.5.4 LU Decomposition

The LU decomposition of a matrix a with row-partial pivoting is a factorization of a (after some rows are possibly permuted by a permutation matrix p) as a product of a lower triangular matrix l and an upper triangular matrix u . This factorization can be used to tackle many common problems in linear algebra such as solving systems of linear equations and computing determinants. An improvement on basic row reduction, LU decomposition by row-partial pivoting keeps the relative magnitude of matrix values small, thus reducing the relative error due to rounding in computed solutions.

- `lu_decompose(a, &picol, eps = 1e-10)` assigns the m by n matrix `a` to merged LU decomposition matrix `lu`, returning either 0 or 1 denoting the “sign” of the permutation parity (0 if the number of overall row swaps performed is even, or 1 if it is odd), or -1 denoting a degenerate matrix (i.e. singular for square matrices). The merged matrix `lu` has `lu[i][j] = l[i][j]` for $i > j$ and `lu[i][j] = u[i][j]` for $i \leq j$. Note that the algorithm always yields a unit lower triangular matrix for which the diagonal entries `l[i][i]` are always equal to 1, so this is not explicitly stored in the resulting merged matrix. For general i and j , the values of the lower and upper triangular matrices should be accessed via the `getl(lu, i, j)` and `getu(lu, i, j)` functions. Optionally, a `vector<int>` pointer `picol` may be passed to return the permutation vector `picol` where `picol[i]` stores the only column that is equal to 1 in row i of the permutation matrix p (all other columns in row i of p are implicitly 0). The permutation matrix p corresponding to `picol` satisfies $pa = lu$.
- `solve_system(a, b, &x, eps = 1e-10)` solves the system of linear equations $ax = b$ given an m by n matrix `a` of real values, and a length m vector `b`, returning 0 if there is one solution or -1 if there are zero or infinite solutions. If there is exactly one solution, then the output pointer `x` is populated with the solution vector of length n .
- `det(a)` returns the determinant of an n by n matrix `a` using LU decomposition.
- `invert(a)` assigns the n by n matrix `a` to its inverse (if it exists), returning 0 if the inversion was successful or -1 if `a` has no inverse.

Time Complexity:

- $O(m \cdot n \cdot \min(m, n))$ per call to `lu_decompose(a)`, where m and n are the numbers of rows and columns of `a`, respectively.
- $O(m \cdot n^2)$ per call to `solve_system()`, which requires $m \geq n$.
- $O(n^3)$ per call to `det(a)` and `invert(a)`, where n is the length of square matrix `a`.

Space Complexity:

- $O(1)$ auxiliary for `lu_decompose()`.
- $O(n^2)$ auxiliary for `det()` and `invert()`.
- $O(m \cdot n)$ auxiliary for `solve_system()`.

```

#include <algorithm>
#include <cmath>
#include <numeric>
#include <vector>

```

```

template<typename Matrix>
int lu_decompose(Matrix &a, std::vector<int> *p1col = nullptr, const double eps = 1e-10) {
    int rows = static_cast<int>(a.size());
    int cols = a.empty() ? 0 : static_cast<int>(a[0].size());
    int parity = 0;
    if (p1col != nullptr) {
        p1col->resize(rows);
        std::iota(p1col->begin(), p1col->end(), 0);
    }
    for (int i = 0; i < rows && i < cols; i++) {
        int pi = i;
        for (int k = i + 1; k < rows; k++) {
            if (fabs(a[k][i]) > fabs(a[pi][i])) {
                pi = k;
            }
        }
        if (fabs(a[pi][i]) < eps) {
            return -1;
        }
        if (pi != i) {
            if (p1col != nullptr) {
                std::iter_swap(p1col->begin() + i, p1col->begin() + pi);
            }
            std::iter_swap(a.begin() + i, a.begin() + pi);
            parity = 1 - parity;
        }
        for (int j = i + 1; j < rows; j++) {
            a[j][i] /= a[i][i];
            for (int k = i + 1; k < cols; k++) {
                a[j][k] -= a[j][i] * a[i][k];
            }
        }
    }
    return parity;
}

template<typename Matrix>
double getl(const Matrix &lu, int i, int j) {
    return i > j ? lu[i][j] : (i < j ? 0 : 1);
}

template<typename Matrix>
double getu(const Matrix &lu, int i, int j) {
    return i <= j ? lu[i][j] : 0;
}

template<typename Matrix, typename T>
int solve_system(
    const Matrix &a, const std::vector<T> &b, std::vector<T> *x, const double eps = 1e-10
) {
    int rows = static_cast<int>(a.size());
    int cols = a.empty() ? 0 : static_cast<int>(a[0].size());
    if (x == nullptr || a.empty() || a.size() != b.size() || rows < cols) {
        return -1;
    }
    x->resize(cols);
    std::vector<int> p1col;
    Matrix lu;
    int status = lu_decompose(lu = a, &p1col, eps);
    if (status < 0) {
        return status;
    }
    for (int i = 0; i < cols; i++) {
        (*x)[i] = b[p1col[i]];
        for (int k = 0; k < i; k++) {
            (*x)[i] -= getl(lu, i, k) * (*x)[k];
        }
    }
}

```

```

for (int i = cols - 1; i >= 0; i--) {
    for (int k = i + 1; k < cols; k++) {
        (*x)[i] -= getu(lu, i, k) * (*x)[k];
    }
    (*x)[i] /= getu(lu, i, i);
}
for (int i = 0; i < rows; i++) {
    double val = 0;
    for (int j = 0; j < cols; j++) {
        val += a[i][j] * (*x)[j];
    }
    // Mixed absolute/relative tolerance: dividing by b[i] alone would skip the check for negative
    // b[i] (ratio goes negative) and divide by zero when b[i] == 0.
    if (fabs(val - b[i]) > eps * (1.0 + fabs(b[i]))) {
        return -1;
    }
}
return 0;
}

template<typename T>
double det(const T &a) {
    int n = static_cast<int>(a.size());
    T lu;
    int status = lu_decompose(lu = a);
    if (status < 0) {
        return 0;
    }
    double res = 1;
    for (int i = 0; i < n; i++) {
        res *= lu[i][i];
    }
    return status == 0 ? res : -res;
}

template<typename T>
int invert(T &a) {
    int n = static_cast<int>(a.size());
    std::vector<int> p1col;
    int status = lu_decompose(a, &p1col);
    if (status < 0) {
        return status;
    }
    T ia(n, typename T::value_type(n, 0));
    for (int j = 0; j < n; j++) {
        for (int i = 0; i < n; i++) {
            if (p1col[i] == j) {
                ia[i][j] = 1.0;
            } else {
                ia[i][j] = 0.0;
            }
            for (int k = 0; k < i; k++) {
                ia[i][j] -= getl(a, i, k) * ia[k][j];
            }
        }
        for (int i = n - 1; i >= 0; i--) {
            for (int k = i + 1; k < n; k++) {
                ia[i][j] -= getu(a, i, k) * ia[k][j];
            }
            ia[i][j] /= getu(a, i, i);
        }
    }
    a.swap(ia);
    return 0;
}

```


Example Usage

```
#include <cassert>
using namespace std;

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-9;
}

int main() {
    { // Solve a system.
        vector<vector<double>> a{{-1, 2, 5}, {1, 0, -6}, {-4, 2, 2}};
        vector<double> b{3, 1, -2};
        vector<double> x;
        assert(solve_system(a, b, &x) == 0);
        for (int i = 0; i < static_cast<int>(a.size()); i++) {
            double sum = 0;
            for (int j = 0; j < static_cast<int>(a[0].size()); j++) {
                sum += a[i][j] * x[j];
            }
            assert(EQ(sum, b[i]));
        }
    }
    { // Find the determinant.
        vector<vector<double>> a{{1, 3, 5}, {2, 4, 7}, {1, 1, 0}};
        assert(EQ(det(a), 4));
    }
    { // Find the inverse.
        vector<vector<double>> a{{6, 1, 1}, {4, -2, 5}, {2, 8, 7}};
        vector<vector<double>> inv = a;
        int n = static_cast<int>(a.size());
        vector<vector<double>> res(n, vector<double>(n, 0));
        assert(invert(inv) == 0);
        for (int i = 0; i < n; i++) {
            for (int j = 0; j < n; j++) {
                for (int k = 0; k < n; k++) {
                    res[i][j] += a[i][k] * inv[k][j];
                }
            }
        }
        for (int i = 0; i < n; i++) {
            for (int j = 0; j < n; j++) {
                assert(EQ(res[i][j], i == j ? 1 : 0));
            }
        }
    }
    return 0;
}
```

6.5.5 Characteristic Polynomial

Computes the characteristic polynomial of a square matrix over a prime field: the coefficients of $\det(xI - A)$, a monic degree- n polynomial whose roots are the eigenvalues of A . Its constant term recovers the determinant (up to sign), its second-highest coefficient is the negated trace, and evaluating it at a scalar gives $\det(xI - A)$ for that x . It is the standard tool behind eigenvalue counting, matrix-power recurrences via Cayley-Hamilton, and problems that reduce to a determinant in an unknown.

A direct expansion would cost a determinant per coefficient. Instead the matrix is first reduced to upper Hessenberg form, where every entry more than one place below the diagonal is zero. This reduction uses similarity transforms (a row operation paired with the inverse column operation), which leave the characteristic polynomial unchanged. The characteristic polynomial of a Hessenberg matrix then follows from a short recurrence over its leading principal submatrices, adding one row and column at a time, for an overall cost of $O(n^3)$. The arithmetic

is exact modulo a prime, since the Hessenberg reduction divides by pivots and is not numerically stable over the reals.

- `characteristic_polynomial(a)` returns a vector c of length $n + 1$ with $\det(xI - \mathbf{a}) = c_0 + c_1x + \dots + c_nx^n$ modulo `MOD`, where $c_n = 1$. The matrix `a` must be square; its entries are reduced modulo `MOD` internally.

Time Complexity: $O(n^3)$ per call, where n is the dimension of the matrix.

Space Complexity: $O(n^2)$ auxiliary.

```
#include <stdint>
#include <utility>
#include <vector>

const int64_t MOD = 998244353;

int64_t powmod(int64_t b, int64_t e, int64_t m) {
    int64_t res = 1;
    for (b %= m; e > 0; e >>= 1) {
        if (e & 1) {
            res = res * b % m;
        }
        b = b * b % m;
    }
    return res;
}

int64_t inv(int64_t a) {
    return powmod((a % MOD + MOD) % MOD, MOD - 2, MOD);
}

void to_hessenberg(std::vector<std::vector<int64_t>> &m) {
    int n = static_cast<int>(m.size());
    for (int j = 0; j + 2 < n; j++) {
        int pivot = -1;
        for (int i = j + 1; i < n; i++) {
            if (m[i][j] != 0) {
                pivot = i;
                break;
            }
        }
        if (pivot == -1) {
            continue;
        }
        if (pivot != j + 1) { // Move the pivot to the subdiagonal via a similarity swap.
            std::swap(m[pivot], m[j + 1]);
            for (int r = 0; r < n; r++) {
                std::swap(m[r][pivot], m[r][j + 1]);
            }
        }
        int64_t pinv = inv(m[j + 1][j]);
        for (int k = j + 2; k < n; k++) {
            int64_t u = m[k][j] * pinv % MOD;
            if (u == 0) {
                continue;
            }
            for (int c = 0; c < n; c++) { // Row k -= u * row (j + 1).
                m[k][c] = (m[k][c] - u * m[j + 1][c] % MOD + MOD) % MOD;
            }
            for (int r = 0; r < n; r++) { // Inverse column operation keeps it a similarity.
                m[r][j + 1] = (m[r][j + 1] + u * m[r][k]) % MOD;
            }
        }
    }
}

std::vector<int64_t> characteristic_polynomial(std::vector<std::vector<int64_t>> a) {
    int n = static_cast<int>(a.size());
```

```

for (auto &row : a) {
    for (int64_t &x : row) {
        x = (x % MOD + MOD) % MOD;
    }
}
to_hessenberg(a);
std::vector<std::vector<int64_t>> p(n + 1);
p[0] = {1};
for (int k = 0; k < n; k++) {
    p[k + 1].assign(k + 2, 0);
    for (int i = 0; i < static_cast<int>(p[k].size()); i++) {
        p[k + 1][i + 1] = (p[k + 1][i + 1] + p[k][i]) % MOD;           // x * p[k].
        p[k + 1][i] = (p[k + 1][i] - a[k][k] * p[k][i] % MOD + MOD) % MOD; // -a[k][k] * p[k].
    }
    int64_t t = 1;
    for (int i = 0; i < k; i++) {
        t = t * a[k - i][k - i - 1] % MOD; // Running product of subdiagonal entries.
        int64_t coef = t * a[k - i - 1][k] % MOD;
        for (int j = 0; j < static_cast<int>(p[k - i - 1].size()); j++) {
            p[k + 1][j] = (p[k + 1][j] - coef * p[k - i - 1][j] % MOD + MOD) % MOD;
        }
    }
}
return p[n];
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // [[1, 2], [3, 4]]: det(xI - A) = x^2 - 5x - 2.
    vector<vector<int64_t>> a{{1, 2}, {3, 4}};
    vector<int64_t> c = characteristic_polynomial(a);
    assert((c == vector<int64_t>{MOD - 2, MOD - 5, 1})); // -2, -5, 1.

    // For any matrix, c_{n-1} = -trace and c_0 = (-1)^n det.
    vector<vector<int64_t>> b{{2, 0, 1}, {0, 3, 0}, {4, 1, 5}};
    vector<int64_t> cb = characteristic_polynomial(b);
    assert(cb.size() == 4 && cb[3] == 1);
    assert(cb[2] == (MOD - (2 + 3 + 5)) % MOD); // Negated trace.
    // det(B) = 2(15) - 0 + 1(0 - 12) = 18, so c_0 = (-1)^3 * 18 = -18.
    assert(cb[0] == (MOD - 18) % MOD);
    return 0;
}

```

6.5.6 Matrix-Tree Theorem

Kirchhoff's matrix-tree theorem counts the spanning trees of an undirected graph. Form the Laplacian matrix $L = D - A$, where D is the diagonal matrix of degrees and A is the adjacency matrix (with multiplicities for parallel edges). Every cofactor of L is equal, and that common value is the number of spanning trees. Equivalently, delete any one row and the matching column from L and take the determinant of what remains.

The determinant is evaluated modulo a prime by Gaussian elimination, so the count is returned modulo `MOD`. Spanning-tree counts grow astronomically, so a modular count is the usual contest form; run the routine under a few different primes and combine with the Chinese remainder theorem to recover a large exact value. The same construction handles weighted graphs (sum edge weights into the Laplacian to get the weighted tree count) and, with the directed Laplacian and a fixed deleted row, counts rooted spanning arborescences.

- `count_spanning_trees(n, edges)` returns the number of spanning trees of the undirected graph on nodes $[0, n)$ with the given edge list, modulo `MOD`. Parallel edges are honored; self-loops do not affect the count. The result is 0 when the graph is disconnected, and 1 when `n` is 1.

Time Complexity: $O(n^3 + m)$ per call, where n is the number of nodes and m is the number of edges.

Space Complexity: $O(n^2)$ auxiliary.

```
#include <stdint>
#include <utility>
#include <vector>

const int64_t MOD = 998244353;

int64_t powmod(int64_t b, int64_t e, int64_t m) {
    int64_t res = 1;
    for (b %= m; e > 0; e >= 1) {
        if (e & 1) {
            res = res * b % m;
        }
        b = b * b % m;
    }
    return res;
}

int64_t count_spanning_trees(int n, const std::vector<std::pair<int, int>> &edges) {
    if (n <= 1) {
        return n == 1 ? 1 : 0;
    }
    // Build the Laplacian, then drop the last row and column to form an (n-1) by (n-1) cofactor.
    int m = n - 1;
    std::vector<std::vector<int64_t>> lap(m, std::vector<int64_t>(m, 0));
    for (const auto &e : edges) {
        int u = e.first, v = e.second;
        if (u == v) {
            continue; // Self-loops contribute nothing to the Laplacian.
        }
        if (u < m) {
            lap[u][u] = (lap[u][u] + 1) % MOD;
        }
        if (v < m) {
            lap[v][v] = (lap[v][v] + 1) % MOD;
        }
        if (u < m && v < m) {
            lap[u][v] = (lap[u][v] - 1 + MOD) % MOD;
            lap[v][u] = (lap[v][u] - 1 + MOD) % MOD;
        }
    }
    // Determinant modulo MOD by Gaussian elimination.
    int64_t det = 1;
    for (int col = 0; col < m; col++) {
        int pivot = -1;
        for (int r = col; r < m; r++) {
            if (lap[r][col] != 0) {
                pivot = r;
                break;
            }
        }
        if (pivot == -1) {
            return 0; // Singular cofactor: the graph is disconnected.
        }
        if (pivot != col) {
            std::swap(lap[pivot], lap[col]);
            det = (MOD - det) % MOD;
        }
        det = det * lap[col][col] % MOD;
        int64_t pinv = powmod(lap[col][col], MOD - 2, MOD);
        for (int r = col + 1; r < m; r++) {

```

```

    int64_t factor = lap[r][col] * pinv % MOD;
    if (factor == 0) {
        continue;
    }
    for (int k = col; k < m; k++) {
        lap[r][k] = (lap[r][k] - factor * lap[col][k]) % MOD;
        if (lap[r][k] < 0) {
            lap[r][k] += MOD;
        }
    }
}
}
return det;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // A path 0-1-2 and a single triangle each have a known number of spanning trees.
    assert(count_spanning_trees(3, {{0, 1}, {1, 2}}) == 1); // A tree has exactly one.
    assert(count_spanning_trees(3, {{0, 1}, {1, 2}, {2, 0}}) == 3); // Triangle: drop any one edge.

    // Complete graphs: Cayley's formula gives  $n^{n-2}$  spanning trees.
    auto complete = [](int n) {
        vector<pair<int, int>> e;
        for (int i = 0; i < n; i++) {
            for (int j = i + 1; j < n; j++) {
                e.emplace_back(i, j);
            }
        }
        return count_spanning_trees(n, e);
    };
    assert(complete(4) == 16); //  $4^2$ .
    assert(complete(5) == 125); //  $5^3$ .

    // Disconnected graphs have no spanning tree.
    assert(count_spanning_trees(4, {{0, 1}, {2, 3}}) == 0);
    return 0;
}

```

6.5.7 Sparse Matrix

A sparse matrix stores only nonzero entries, which is useful when the matrix is large but most values are zero. This implementation keeps both row maps and column maps in sync: `row(i)` can iterate all nonzeros in row `i`, while `column(j)` can iterate all nonzeros in column `j`. That bidirectional storage is especially useful for sparse elimination or graph-like matrix operations, where both row updates and column pivot lookups are common.

- `SparseMatrix<T>(h, w)` constructs an `h` by `w` sparse matrix.
- `height()` returns the number of rows.
- `width()` returns the number of columns.
- `nonzeros()` returns the number of stored nonzero entries.
- `get(i, j)` returns the value at row `i`, column `j`, or 0 if the entry is not stored.
- `set(i, j, value)` assigns an entry. Assigning zero erases it from both maps.
- `add(i, j, delta)` adds `delta` to an entry. If the result becomes zero, the entry is erased.
- `row(i)` returns a map from column index to value for the nonzero entries in row `i`.
- `column(j)` returns a map from row index to value for the nonzero entries in column `j`.
- `swap_rows(i, j)` swaps two rows while keeping the column maps synchronized.

- `transpose()` transposes the matrix in place.
- `multiply_vector(x)` returns the matrix-vector product with vector `x`.
- `row_reduce(a, limit)` converts columns $[0, \text{limit})$ of `a` to sparse row echelon form, returning the rank found in those columns.
- `det(a)` returns the determinant of a square sparse matrix.
- `sparse_rank(a)` returns the rank of a sparse matrix.
- `solve_system(a, b, &x)` solves the system $ax = b$, returning 0 for one solution, -1 for no solution, or -2 for infinitely many solutions. When one solution exists, `x` is populated.

The class assumes that `T{}` is the additive zero and that `value == T{} is a valid zero test. The elimination routines also require field-like arithmetic because they divide by pivots, so use types such as double, rational numbers, or modular integers. For floating-point matrices, replace is_zero() with an epsilon comparison if small roundoff values should be erased.`

Sparse elimination chooses, in each pivot column, the available pivot row with the fewest stored entries. This heuristic can reduce fill-in, although sparse elimination can still become dense on unfriendly inputs.

Time Complexity:

- $O(\log d)$ per call to `get()`, `set()`, and `add()`, where d is the number of nonzeros in the touched row or column.
- $O(1)$ per call to `transpose()` and $O(n)$ per call to `multiply_vector()`, where n is the number of stored nonzero entries.
- $O((r_i + r_j) \cdot \log d)$ per call to `swap_rows(i, j)`, where r_i and r_j are the sizes of the two rows.
- $O(f \cdot \log d)$ for sparse elimination, where f is the number of entry updates performed after fill-in and d is a touched row or column size. In the worst case this is still cubic.
- $O(f \cdot \log d)$ per call to `det()`, `sparse_rank()`, and `solve_system()`.
- $O(1)$ per call to `height()`, `width()`, `nonzeros()`, `row()`, and `column()`.

Space Complexity:

- $O(n)$ storage.
- $O(1)$ auxiliary for `get()`, `set()`, `add()`, `row()`, `column()`, and `transpose()`.
- $O(r_i + r_j)$ auxiliary for `swap_rows(i, j)`.
- $O(h)$ auxiliary for `multiply_vector()`.
- $O(n + h)$ auxiliary for elimination, in addition to the copied matrix used by `det()`, `sparse_rank()`, and `solve_system()`.

```
#include <cassert>
#include <cstdint>
#include <map>
#include <unordered_set>
#include <utility>
#include <vector>

template<typename T>
class SparseMatrix {
    int h = 0, w = 0, stored = 0;
    std::vector<std::map<int, T>> rows, cols;

    bool is_zero(const T &value) const { return value == T{}; }

public:
    SparseMatrix(int h_, int w_) : h(h_), w(w_), rows(h_), cols(w_) {}

    int height() const { return h; }
    int width() const { return w; }
    int nonzeros() const { return stored; }

    T get(int i, int j) const {
        auto it = rows[i].find(j);
        return it == rows[i].end() ? T{} : it->second;
    }
}
```

```

void set(int i, int j, const T &value) {
    auto it = rows[i].find(j);
    if (is_zero(value)) {
        if (it != rows[i].end()) {
            rows[i].erase(it);
            cols[j].erase(i);
            stored--;
        }
        return;
    }
    if (it == rows[i].end()) {
        stored++;
    }
    rows[i][j] = value;
    cols[j][i] = value;
}

void add(int i, int j, const T &delta) {
    if (is_zero(delta)) {
        return;
    }
    set(i, j, get(i, j) + delta);
}

const std::map<int, T> &row(int i) const { return rows[i]; }
const std::map<int, T> &column(int j) const { return cols[j]; }

void swap_rows(int i, int j) {
    if (i == j) {
        return;
    }
    std::unordered_set<int> touched;
    for (const auto &[col, value] : rows[i]) {
        touched.insert(col);
    }
    for (const auto &[col, value] : rows[j]) {
        touched.insert(col);
    }
    for (int col : touched) {
        T x = get(i, col), y = get(j, col);
        set(i, col, y);
        set(j, col, x);
    }
}

void transpose() {
    std::swap(h, w);
    std::swap(rows, cols);
}

std::vector<T> multiply_vector(const std::vector<T> &x) const {
    assert(static_cast<int>(x.size()) == w);
    std::vector<T> res(h);
    for (int i = 0; i < h; i++) {
        for (const auto &[j, value] : rows[i]) {
            res[i] += value * x[j];
        }
    }
    return res;
}

};

template<typename T>
int row_reduce(SparseMatrix<T> &a, int limit) {
    assert(0 <= limit && limit <= a.width());
    int r = 0;
    for (int c = 0; c < limit && r < a.height(); c++) {

```

```

    int pivot = -1;
    for (const auto &[i, value] : a.column(c)) {
        if (i >= r && (pivot == -1 || a.row(i).size() < a.row(pivot).size())) {
            pivot = i;
        }
    }
    if (pivot == -1) {
        continue;
    }
    a.swap_rows(r, pivot);
    T inv_pivot = T{1} / a.get(r, c);
    std::vector<std::pair<int, T>> pivot_row(a.row(r).begin(), a.row(r).end());
    std::vector<int> touched_rows;
    for (const auto &[i, value] : a.column(c)) {
        if (i > r) {
            touched_rows.push_back(i);
        }
    }
    for (int i : touched_rows) {
        T coeff = -a.get(i, c) * inv_pivot;
        for (const auto &[j, value] : pivot_row) {
            if (j != c) {
                a.add(i, j, coeff * value);
            }
        }
        a.set(i, c, T{});
    }
    r++;
}
return r;
}

```

```

template<typename T>
T det(SparseMatrix<T> a) {
    assert(a.height() == a.width());
    int swaps = 0, r = 0, n = a.height();
    T det = 1;
    for (int c = 0; c < n && r < n; c++) {
        int pivot = -1;
        for (const auto &[i, value] : a.column(c)) {
            if (i >= r && (pivot == -1 || a.row(i).size() < a.row(pivot).size())) {
                pivot = i;
            }
        }
        if (pivot == -1) {
            return T{};
        }
        if (pivot != r) {
            a.swap_rows(r, pivot);
            swaps++;
        }
        T pivot_value = a.get(r, c);
        det *= pivot_value;
        T inv_pivot = T{1} / pivot_value;
        std::vector<std::pair<int, T>> pivot_row(a.row(r).begin(), a.row(r).end());
        std::vector<int> touched_rows;
        for (const auto &[i, value] : a.column(c)) {
            if (i > r) {
                touched_rows.push_back(i);
            }
        }
        for (int i : touched_rows) {
            T coeff = -a.get(i, c) * inv_pivot;
            for (const auto &[j, value] : pivot_row) {
                if (j != c) {
                    a.add(i, j, coeff * value);
                }
            }
        }
    }
}

```



```

        a.set(i, c, T{});
    }
    r++;
}
return (r == n ? (swaps % 2 == 0 ? det : -det) : T{});
}

template<typename T>
int sparse_rank(SparseMatrix<T> a) {
    return row_reduce(a, a.width());
}

template<typename T>
int solve_system(SparseMatrix<T> a, const std::vector<T> &b, std::vector<T> *x) {
    if (x == nullptr || a.height() != static_cast<int>(b.size())) {
        return -1;
    }
    int h = a.height(), w = a.width();
    SparseMatrix<T> aug(h, w + 1);
    for (int i = 0; i < h; i++) {
        for (const auto &[j, value] : a.row(i)) {
            aug.set(i, j, value);
        }
        aug.set(i, w, b[i]);
    }
    row_reduce(aug, w);
    std::vector<int> pivot_col;
    for (int i = 0; i < h; i++) {
        int lead = -1;
        for (const auto &[j, value] : aug.row(i)) {
            if (j < w) {
                lead = j;
                break;
            }
        }
        if (lead == -1) {
            if (aug.get(i, w) != T{}) {
                return -1;
            }
        } else {
            pivot_col.push_back(lead);
        }
    }
    if (static_cast<int>(pivot_col.size()) < w) {
        return -2;
    }
    x->assign(w, T{});
    for (int row = static_cast<int>(pivot_col.size()) - 1; row >= 0; row--) {
        int c = pivot_col[row];
        T rhs = aug.get(row, w);
        for (const auto &[j, value] : aug.row(row)) {
            if (j > c && j < w) {
                rhs -= value * (*x)[j];
            }
        }
        (*x)[c] = rhs / aug.get(row, c);
    }
    return 0;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {

```

```

SparseMatrix<int64_t> a(3, 4);
a.set(0, 1, 5);
a.set(2, 3, 7);
a.add(0, 1, -2);
a.add(1, 2, 4);

assert(a.height() == 3 && a.width() == 4);
assert(a.nonzeros() == 3);
assert(a.get(0, 1) == 3);
assert(a.get(0, 0) == 0);
assert(a.row(0).begin()->first == 1);
assert(a.column(2).begin()->first == 1);

vector<int64_t> x{10, 20, 30, 40};
assert((a.multiply_vector(x) == vector<int64_t>{60, 120, 280}));

a.add(0, 1, -3); // Entries that become zero are removed from both row and column maps.
assert(a.get(0, 1) == 0);
assert(a.nonzeros() == 2);
assert(a.row(0).empty());
assert(a.column(1).empty());

a.transpose();
assert(a.height() == 4 && a.width() == 3);
assert(a.get(3, 2) == 7);
assert(a.get(2, 1) == 4);

SparseMatrix<double> d(3, 3);
d.set(0, 0, 2);
d.set(0, 1, 1);
d.set(1, 1, 3);
d.set(2, 0, 1);
d.set(2, 2, 4);
double determinant = det(d);
assert(determinant > 24 - 1e-9 && determinant < 24 + 1e-9);
assert(sparse_rank(d) == 3);

SparseMatrix<double> sys(3, 3);
sys.set(0, 0, -1);
sys.set(0, 1, 2);
sys.set(0, 2, 5);
sys.set(1, 0, 1);
sys.set(1, 2, -6);
sys.set(2, 0, -4);
sys.set(2, 1, 2);
sys.set(2, 2, 2);
vector<double> rhs{3, 1, -2}, sol;
assert(solve_system(sys, rhs, &sol) == 0);
for (int i = 0; i < sys.height(); i++) {
    double sum = 0;
    for (const auto &[j, value] : sys.row(i)) {
        sum += value * sol[j];
    }
    assert(sum > rhs[i] - 1e-9 && sum < rhs[i] + 1e-9);
}
return 0;
}

```

6.6 Polynomials

6.6.1 Number Theoretic Transform

The number theoretic transform (NTT) is the discrete Fourier transform carried out in modular arithmetic instead of over the complex numbers. It multiplies two polynomials (equivalently, computes the convolution of two sequences) modulo a prime in $O(n \log n)$ time, with no floating point error. The modulus must be of the form $c \cdot 2^k + 1$ so that the ring of integers modulo it contains a 2^k -th root of unity, which limits transform sizes to powers of two up to 2^k . The code below uses the common prime $998244353 = 119 \cdot 2^{23} + 1$ with primitive root 3, supporting transforms of any power-of-two length up to 2^{23} .

The transform replaces a coefficient vector with its evaluations at the powers of a root of unity; pointwise multiplying two such evaluation vectors and transforming back gives their convolution. This is the same divide-and-conquer butterfly as the fast Fourier transform, with the root of unity and all arithmetic taken modulo the prime.

- `ntt(a, invert)` transforms the vector `a` in place, whose length must be a power of two. The forward transform uses `invert = false`; the inverse uses `invert = true`.
- `convolve(a, b)` returns the convolution of `a` and `b` modulo the prime, that is, the coefficients of the product of the two polynomials whose coefficients are the inputs. The result has length `a.size() + b.size() - 1`, or is empty if either input is empty. Input values are assumed to lie in $[0, \text{MOD})$. Small inputs use direct multiplication to avoid NTT overhead.

Time Complexity:

- $O(n \log n)$ per call to `ntt()`, where n is the length of the vector.
- $O(|a||b|)$ per call to `convolve()` on small inputs and $O(n \log n)$ otherwise, where n is the padded transform length.

Space Complexity:

- $O(1)$ auxiliary for `ntt()`, transforming in place.
- $O(n)$ auxiliary for `convolve()`.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <utility>
#include <vector>

const int64_t MOD = 998244353;
const int64_t ROOT = 3;
const int MAX_POWER_OF_TWO = 23;
const int NAIVE_CUTOFF = 150;

int64_t powmod(int64_t b, int64_t e, int64_t m) {
    int64_t res = 1;
    for (b %= m; e > 0; e >>= 1) {
        if (e & 1) {
            res = res * b % m;
        }
        b = b * b % m;
    }
    return res;
}

void ntt(std::vector<int64_t> &a, bool invert) {
    int n = static_cast<int>(a.size());
    assert(n > 0 && (n & (n - 1)) == 0 && n <= (1 << MAX_POWER_OF_TWO));
    for (int i = 1, j = 0; i < n; i++) {
        int bit = n >> 1;
        for (; j & bit; bit >>= 1) {
            j ^= bit;
        }
    }
}
```

```

    }
    j ^= bit;
    if (i < j) {
        std::swap(a[i], a[j]);
    }
}
for (int len = 2; len <= n; len <<= 1) {
    int64_t root = powmod(ROOT, (MOD - 1) / len, MOD);
    if (invert) {
        root = powmod(root, MOD - 2, MOD);
    }
    for (int i = 0; i < n; i += len) {
        int64_t w = 1;
        for (int k = 0; k < len / 2; k++) {
            int64_t u = a[i + k], v = a[i + k + len / 2] * w % MOD;
            a[i + k] = (u + v) % MOD;
            a[i + k + len / 2] = (u - v + MOD) % MOD;
            w = w * root % MOD;
        }
    }
}
if (invert) {
    int64_t n_inv = powmod(n, MOD - 2, MOD);
    for (int64_t &x : a) {
        x = x * n_inv % MOD;
    }
}
}

std::vector<int64_t> convolve(std::vector<int64_t> a, std::vector<int64_t> b) {
    if (a.empty() || b.empty()) {
        return {};
    }
    int result_size = a.size() + b.size() - 1;
    if (std::min(a.size(), b.size()) < NAIVE_CUTOFF) {
        std::vector<int64_t> res(result_size);
        for (int i = 0; i < static_cast<int>(a.size()); i++) {
            for (int j = 0; j < static_cast<int>(b.size()); j++) {
                res[i + j] = (res[i + j] + a[i] * b[j]) % MOD;
            }
        }
        return res;
    }
    int n = 1;
    while (n < result_size) {
        n <<= 1;
    }
    assert(n <= (1 << MAX_POWER_OF_TWO));
    a.resize(n);
    b.resize(n);
    ntt(a, false);
    ntt(b, false);
    for (int i = 0; i < n; i++) {
        a[i] = a[i] * b[i] % MOD;
    }
    ntt(a, true);
    a.resize(result_size);
    return a;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {

```

```

// (1 + 2x + 3x^2)(4 + 5x + 6x^2) = 4 + 13x + 28x^2 + 27x^3 + 18x^4.
vector<int64_t> a{1, 2, 3}, b{4, 5, 6};
vector<int64_t> c = convolve(a, b);
assert((c == vector<int64_t>{4, 13, 28, 27, 18}));
assert(convolve({}, b).empty());
assert((convolve(vector<int64_t>{5}, vector<int64_t>{7}) == vector<int64_t>{35}));
vector<int64_t> ones(160, 1);
vector<int64_t> d = convolve(ones, ones);
assert(d[0] == 1 && d[159] == 160 && d[318] == 1);
return 0;
}

```

6.6.2 Fast Fourier Transform

The fast Fourier transform (FFT) multiplies two polynomials, or equivalently convolves two sequences, in $O(n \log n)$ time by evaluating them at the complex roots of unity. Like the number theoretic transform of the previous section, it replaces a coefficient vector with its evaluations at the powers of a root of unity, multiplies two such evaluation vectors pointwise, and transforms back. The difference is the arithmetic: the FFT works over the complex numbers with $e^{2\pi i/n}$ as the root of unity, so it convolves arbitrary real coefficients rather than residues modulo a prime.

The trade-off is precision. Complex arithmetic in double precision accumulates rounding error, so the result is recovered by rounding each output to the nearest integer; this is reliable as long as the true coefficients stay well within the roughly 15 significant decimal digits a `double` can hold. When exact results modulo a prime are required, or coefficients are large, prefer the number theoretic transform; use the FFT for real-valued convolution or big-integer multiplication.

- `fft(a, invert)` transforms the complex vector `a` in place, whose length must be a power of two. The forward transform uses `invert = false`; the inverse uses `invert = true`.
- `convolve(a, b)` returns the convolution of two integer sequences `a` and `b`, that is, the coefficients of the product of the two polynomials whose coefficients are the inputs. The result has length `a.size() + b.size() - 1`, or is empty if either input is empty, with each entry rounded to the nearest integer. Small inputs use direct multiplication to avoid FFT overhead. The exact coefficients and the direct branch's intermediate products and sums must fit in `int64_t`.

Time Complexity:

- $O(n \log n)$ per call to `fft()`, where n is the length of the vector.
- $O(|a||b|)$ per call to `convolve()` on small inputs and $O(n \log n)$ otherwise, where n is the padded transform length.

Space Complexity:

- $O(1)$ auxiliary for `fft()`, transforming in place.
- $O(n)$ auxiliary for `convolve()`.

```

#include <algorithm>
#include <cassert>
#include <cmath>
#include <complex>
#include <cstdint>
#include <utility>
#include <vector>

typedef std::complex<double> cd;
const double PI = acos(-1.0);
const int NAIVE_CUTOFF = 150;

void fft(std::vector<cd> &a, bool invert) {
    int n = static_cast<int>(a.size());
    assert(n > 0 && (n & (n - 1)) == 0);

```

```

for (int i = 1, j = 0; i < n; i++) {
    int bit = n >> 1;
    for (; j & bit; bit >>= 1) {
        j ^= bit;
    }
    j ^= bit;
    if (i < j) {
        std::swap(a[i], a[j]);
    }
}
for (int len = 2; len <= n; len <= 1) {
    double angle = 2 * PI / len * (invert ? -1 : 1);
    cd root(cos(angle), sin(angle));
    for (int i = 0; i < n; i += len) {
        cd w(1);
        for (int k = 0; k < len / 2; k++) {
            cd u = a[i + k], v = a[i + k + len / 2] * w;
            a[i + k] = u + v;
            a[i + k + len / 2] = u - v;
            w *= root;
        }
    }
}
if (invert) {
    for (cd &x : a) {
        x /= n;
    }
}
}

std::vector<int64_t> convolve(const std::vector<int64_t> &a, const std::vector<int64_t> &b) {
    if (a.empty() || b.empty()) {
        return {};
    }
    int result_size = a.size() + b.size() - 1;
    if (std::min(a.size(), b.size()) < NAIVE_CUTOFF) {
        std::vector<int64_t> res(result_size);
        for (int i = 0; i < static_cast<int>(a.size()); i++) {
            for (int j = 0; j < static_cast<int>(b.size()); j++) {
                res[i + j] += a[i] * b[j]; // Overflow warning.
            }
        }
        return res;
    }
    int n = 1;
    while (n < result_size) {
        n <= 1;
    }
    std::vector<cd> fa(a.begin(), a.end()), fb(b.begin(), b.end());
    fa.resize(n);
    fb.resize(n);
    fft(fa, false);
    fft(fb, false);
    for (int i = 0; i < n; i++) {
        fa[i] *= fb[i];
    }
    fft(fa, true);
    std::vector<int64_t> res(result_size);
    for (int i = 0; i < result_size; i++) {
        res[i] = llround(fa[i].real());
    }
    return res;
}

```

Example Usage

```
using namespace std;

int main() {
    // (1 + 2x + 3x^2)(4 + 5x + 6x^2) = 4 + 13x + 28x^2 + 27x^3 + 18x^4.
    vector<int64_t> a{1, 2, 3}, b{4, 5, 6};
    vector<int64_t> c = convolve(a, b);
    assert((c == vector<int64_t>{4, 13, 28, 27, 18}));
    assert(convolve({}, b).empty());
    assert((convolve(vector<int64_t>{5}, vector<int64_t>{7}) == vector<int64_t>{35}));
    vector<int64_t> ones(160, 1);
    vector<int64_t> d = convolve(ones, ones);
    assert(d[0] == 1 && d[159] == 160 && d[318] == 1);
    return 0;
}
```

6.6.3 Polynomial Operations

Polynomial operations treat a vector `a` as coefficients of $a_0 + a_1x + \dots + a_{n-1}x^{n-1}$ over a field. This section implements the core algebra needed for formal power series and polynomial division modulo the prime 998244353. Multiplication uses the number theoretic transform, so inverse and division are fast enough for large polynomials while still keeping the API close to the underlying coefficient vectors.

The code aliases one modular field element as `Coeff` and one coefficient vector as `Poly`. Input coefficients must lie in $[0, \text{MOD})$.

- `trim(a)` removes trailing zero coefficients from `a`.
- `add(a, b)` returns `a + b`.
- `subtract(a, b)` returns `a - b`.
- `multiply(a, b)` returns `a * b`.
- `derivative(a)` returns the formal derivative of `a`.
- `integral(a)` returns the formal antiderivative of `a` with constant term zero.
- `inverse(a, n)` returns the first `n` coefficients of `1 / a`, requiring `a[0]` to be nonzero.
- `divide(a, b)` returns the polynomial quotient of `a / b`, requiring `b` to be nonzero.
- `modulo(a, b)` returns the polynomial remainder of `a / b`, requiring `b` to be nonzero.

The inverse is a formal power series inverse modulo x^n : it finds b such that $a \cdot b \equiv 1 \pmod{x^n}$. Division uses the standard reversal trick: reverse both polynomials, compute a truncated series inverse of the reversed divisor, multiply, truncate, and reverse back.

Time Complexity:

- $O(n)$ per call to `trim()`, `add()`, `subtract()`, `derivative()`, and `integral()`.
- $O(|a||b|)$ per call to `multiply()` on small inputs and $O(n \log n)$ otherwise, where n is the padded transform length.
- $O(n \log n)$ per call to `inverse(a, n)`, where n is the requested length.
- $O(n \log n)$ per call to `divide()` and `modulo()`, where n is the padded multiplication length.

Space Complexity: $O(n)$ auxiliary.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <utility>
#include <vector>

using Coeff = int64_t;
using Poly = std::vector<Coeff>;
```

```

const Coeff MOD = 998244353;
const Coeff ROOT = 3;
const int MAX_POWER_OF_TWO = 23;
const int NAIVE_CUTOFF = 150;

Coeff powmod(Coeff b, Coeff e, Coeff m = MOD) {
    Coeff res = 1;
    for (b %= m; e > 0; e >>= 1) {
        if (e & 1) {
            res = res * b % m;
        }
        b = b * b % m;
    }
    return res;
}

void ntt(Poly &a, bool invert) {
    int n = static_cast<int>(a.size());
    assert(n > 0 && (n & (n - 1)) == 0 && n <= (1 << MAX_POWER_OF_TWO));
    for (int i = 1, j = 0; i < n; i++) {
        int bit = n >> 1;
        for (; j & bit; bit >>= 1) {
            j ^= bit;
        }
        j ^= bit;
        if (i < j) {
            std::swap(a[i], a[j]);
        }
    }
    for (int len = 2; len <= n; len <<= 1) {
        Coeff root = powmod(ROOT, (MOD - 1) / len);
        if (invert) {
            root = powmod(root, MOD - 2);
        }
        for (int i = 0; i < n; i += len) {
            Coeff w = 1;
            for (int k = 0; k < len / 2; k++) {
                Coeff u = a[i + k], v = a[i + k + len / 2] * w % MOD;
                a[i + k] = (u + v) % MOD;
                a[i + k + len / 2] = (u - v + MOD) % MOD;
                w = w * root % MOD;
            }
        }
    }
    if (invert) {
        Coeff n_inv = powmod(n, MOD - 2);
        for (Coeff &x : a) {
            x = x * n_inv % MOD;
        }
    }
}

void trim(Poly &a) {
    while (!a.empty() && a.back() == 0) {
        a.pop_back();
    }
}

Poly add(const Poly &a, const Poly &b) {
    Poly res(std::max(a.size(), b.size()));
    for (int i = 0; i < static_cast<int>(res.size()); i++) {
        Coeff x = i < static_cast<int>(a.size()) ? a[i] : 0;
        Coeff y = i < static_cast<int>(b.size()) ? b[i] : 0;
        res[i] = (x + y) % MOD;
    }
    trim(res);
    return res;
}

```



```

Poly subtract(const Poly &a, const Poly &b) {
    Poly res(std::max(a.size(), b.size()));
    for (int i = 0; i < static_cast<int>(res.size()); i++) {
        Coeff x = i < static_cast<int>(a.size()) ? a[i] : 0;
        Coeff y = i < static_cast<int>(b.size()) ? b[i] : 0;
        res[i] = (x - y + MOD) % MOD;
    }
    trim(res);
    return res;
}

Poly multiply(Poly a, Poly b) {
    if (a.empty() || b.empty()) {
        return {};
    }
    int result_size = a.size() + b.size() - 1;
    if (std::min(a.size(), b.size()) < NAIVE_CUTOFF) {
        Poly res(result_size);
        for (int i = 0; i < static_cast<int>(a.size()); i++) {
            for (int j = 0; j < static_cast<int>(b.size()); j++) {
                res[i + j] = (res[i + j] + a[i] * b[j]) % MOD;
            }
        }
        trim(res);
        return res;
    }
    int n = 1;
    while (n < result_size) {
        n <<= 1;
    }
    assert(n <= (1 << MAX_POWER_OF_TWO));
    a.resize(n);
    b.resize(n);
    ntt(a, false);
    ntt(b, false);
    for (int i = 0; i < n; i++) {
        a[i] = a[i] * b[i] % MOD;
    }
    ntt(a, true);
    a.resize(result_size);
    trim(a);
    return a;
}

Poly derivative(const Poly &a) {
    if (a.size() <= 1) {
        return {};
    }
    Poly res(a.size() - 1);
    for (int i = 1; i < static_cast<int>(a.size()); i++) {
        res[i - 1] = a[i] * i % MOD;
    }
    trim(res);
    return res;
}

Poly integral(const Poly &a) {
    Poly res(a.size() + 1);
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        res[i + 1] = a[i] * powmod(i + 1, MOD - 2) % MOD;
    }
    return res;
}

Poly inverse(const Poly &a, int n) {
    assert(n >= 0 && !a.empty() && a[0] != 0);
    if (n == 0) {

```

```

    return {};
}
Poly res{powmod(a[0], MOD - 2)};
while (static_cast<int>(res.size()) < n) {
    int len = static_cast<int>(res.size()) << 1;
    Poly cut(a.begin(), a.begin() + std::min(static_cast<int>(a.size()), len));
    Poly prod = multiply(multiply(res, res), cut);
    res.resize(len);
    for (int i = len / 2; i < std::min(len, static_cast<int>(prod.size())); i++) {
        res[i] = (MOD - prod[i]) % MOD;
    }
}
res.resize(n);
return res;
}

Poly divide(Poly a, Poly b) {
    trim(a);
    trim(b);
    assert(!b.empty());
    if (a.size() < b.size()) {
        return {};
    }
    int quotient_size = static_cast<int>(a.size() - b.size() + 1);
    std::reverse(a.begin(), a.end());
    std::reverse(b.begin(), b.end());
    b.resize(quotient_size);
    Poly q = multiply(a, inverse(b, quotient_size));
    q.resize(quotient_size);
    std::reverse(q.begin(), q.end());
    trim(q);
    return q;
}

Poly modulo(const Poly &a, const Poly &b) {
    Poly q = divide(a, b);
    Poly res = subtract(a, multiply(q, b));
    if (res.size() >= b.size()) {
        res.resize(b.size() - 1);
    }
    trim(res);
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // (1 + 2x + 3x^2)(4 + 5x + 6x^2) = 4 + 13x + 28x^2 + 27x^3 + 18x^4.
    Poly a{1, 2, 3}, b{4, 5, 6};
    assert((add(a, b) == Poly{5, 7, 9}));
    assert((subtract(b, a) == Poly{3, 3, 3}));
    assert((multiply(a, b) == Poly{4, 13, 28, 27, 18}));

    // d/dx (5 + 7x + 11x^2) = 7 + 22x; integrating back chooses constant term 0.
    Poly p{5, 7, 11};
    assert((derivative(p) == Poly{7, 22}));
    assert((integral(derivative(p)) == Poly{0, 7, 11}));

    // (1 - x)^-1 = 1 + x + x^2 + x^3 + ... modulo x^6.
    Poly inv = inverse({1, MOD - 1}, 6);
    assert((inv == Poly{1, 1, 1, 1, 1, 1}));

    // (x^3 - 1) / (x - 1) = x^2 + x + 1, remainder 0.
}

```

```

Poly dividend{MOD - 1, 0, 0, 1};
Poly divisor{MOD - 1, 1};
assert((divide(dividend, divisor) == Poly{1, 1, 1}));
assert(modulo(dividend, divisor).empty());

Poly ones(160, 1);
Poly square = multiply(ones, ones);
assert(square[0] == 1 && square[159] == 160 && square[318] == 1);
return 0;
}

```

6.6.4 Berlekamp-Massey

The Berlekamp-Massey algorithm finds the shortest linear recurrence that generates a given sequence over a field. Given the first terms of a sequence taken modulo a prime, it returns coefficients $\Lambda_0, \dots, \Lambda_{L-1}$ of the minimal recurrence $s_i = \Lambda_0 s_{i-1} + \Lambda_1 s_{i-2} + \dots + \Lambda_{L-1} s_{i-L} \pmod{p}$, holding for every $i \geq L$. It scans the sequence once, maintaining a current candidate recurrence and correcting it whenever a term fails to match the prediction, using the last position where a correction was needed. The order L never decreases, and $2L$ terms suffice to determine a recurrence of order L . This is the standard tool for guessing a linear recurrence from sampled values, after which the sequence can be extended or its n -th term found by fast methods.

An order- L recurrence generally needs at least $2L$ terms to be determined, so the result is most trustworthy when $\mathbf{s}$ is comfortably longer than twice the returned length. If the returned L is close to $n/2$, the sequence is under-determined and the recurrence may be a spurious overfit.

The code operates modulo the prime 998244353; change `MOD` to use a different prime field.

- `berlekamp_massey(s)` returns the coefficients of the shortest recurrence that the sequence $\mathbf{s}$ satisfies, as described above. The returned length is the order L of the recurrence. An empty vector is returned when $\mathbf{s}$ is entirely zero (no recurrence is needed). Entries are reduced modulo `MOD` internally.

Time Complexity: $O(n^2 + L \log p)$ per call, where n is the length of $\mathbf{s}$, L is the order of the recurrence found, and p is the modulus.

Space Complexity: $O(n)$ auxiliary.

```

#include <stdint>
#include <vector>

const int64_t MOD = 998244353;

int64_t powmod(int64_t b, int64_t e, int64_t m) {
    int64_t res = 1;
    for (b %= m; e > 0; e >>= 1) {
        if (e & 1) {
            res = res * b % m;
        }
        b = b * b % m;
    }
    return res;
}

std::vector<int64_t> berlekamp_massey(const std::vector<int64_t> &s) {
    int n = static_cast<int>(s.size()), len = 0, m = 0;
    if (n == 0) {
        return {};
    }
    std::vector<int64_t> cur(n), last(n), prev;
    cur[0] = last[0] = 1;
    int64_t inv_last_delta = 1;
    for (int i = 0; i < n; i++) {
        m++;

```

```

int64_t delta = (s[i] % MOD + MOD) % MOD;
for (int j = 1; j <= len; j++) {
    int64_t value = (s[i - j] % MOD + MOD) % MOD;
    delta = (delta + cur[j] * value) % MOD;
}
if (delta == 0) {
    continue;
}
prev = cur;
int64_t coef = delta * inv_last_delta % MOD;
for (int j = m; j < n; j++) {
    cur[j] = (cur[j] - coef * last[j - m] % MOD + MOD) % MOD;
}
if (2 * len > i) {
    continue;
}
len = i + 1 - len;
last = prev;
inv_last_delta = powmod(delta, MOD - 2, MOD); // Refresh only when the length increases.
m = 0;
}
cur.resize(len + 1);
cur.erase(cur.begin());
for (int64_t &x : cur) {
    x = (MOD - x) % MOD;
}
return cur;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    assert(berlekamp_massey({}).empty());
    // Fibonacci: s[i] = s[i-1] + s[i-2].
    vector<int64_t> fib{1, 1, 2, 3, 5, 8, 13, 21};
    assert((berlekamp_massey(fib) == vector<int64_t>{1, 1}));

    // Geometric: s[i] = 2 * s[i-1].
    vector<int64_t> geo{1, 2, 4, 8, 16, 32};
    assert((berlekamp_massey(geo) == vector<int64_t>{2}));
    assert((berlekamp_massey(vector<int64_t>{1, -1, 1, -1, 1, -1}) == vector<int64_t>{MOD - 1}));

    // Tribonacci: s[i] = s[i-1] + s[i-2] + s[i-3].
    vector<int64_t> trib{0, 1, 1, 2, 4, 7, 13, 24, 44};
    assert((berlekamp_massey(trib) == vector<int64_t>{1, 1, 1}));
    return 0;
}

```

6.6.5 Linear Recurrence (Kitamasa)

Computes the n -th term of a linear recurrence modulo a prime in logarithmic time. Given a recurrence of order L , $s_i = c_0 s_{i-1} + c_1 s_{i-2} + \dots + c_{L-1} s_{i-L}$, together with the first L terms, the naive approach unrolls n steps. Kitamasa's method instead jumps directly to index n in $O(L^2 \log n)$ by working with the characteristic polynomial. The coefficient layout matches the output of Berlekamp-Massey, so the two compose directly: guess the recurrence from sampled values, then jump to any index.

The key identity is that the n -th term is a fixed linear combination of the first L terms, with weights given by x^n reduced modulo the characteristic polynomial $f(x) = x^L - c_0 x^{L-1} - \dots - c_{L-1}$. The reduction $x^L = c_0 x^{L-1} + \dots + c_{L-1}$ lets any product of two degree- $< L$ polynomials be folded back to degree $< L$, so x^n

$\text{mod } f$ is obtained by binary exponentiation, and the n -th term is the dot product of its coefficients with the initial terms.

- `kth_term(rec, init, n)` returns s_n modulo `MOD`, where `rec` holds the coefficients $c_0, \dots, c_{L-1}$ and `init` holds at least L initial terms $s_0, \dots, s_{L-1}$. Indexing is 0-based and `n` ≥ 0 . Values are reduced modulo `MOD` internally.

Time Complexity: $O(L^2 \log n)$ per call.

Space Complexity: $O(L)$ auxiliary.

```
#include <cassert>
#include <cstdint>
#include <vector>

const int64_t MOD = 998244353;

void reduce(std::vector<int64_t> &p, const std::vector<int64_t> &c) {
    int L = static_cast<int>(c.size());
    for (int d = static_cast<int>(p.size()) - 1; d >= L; d--) {
        if (p[d] != 0) {
            for (int j = 0; j < L; j++) {
                p[d - 1 - j] = (p[d - 1 - j] + p[d] * c[j]) % MOD;
            }
            p[d] = 0;
        }
    }
    p.resize(L);
}

std::vector<int64_t> combine(
    const std::vector<int64_t> &a, const std::vector<int64_t> &b, const std::vector<int64_t> &c
) {
    std::vector<int64_t> res(a.size() + b.size() - 1, 0);
    for (int i = 0; i < static_cast<int>(a.size()); i++) {
        for (int j = 0; j < static_cast<int>(b.size()); j++) {
            res[i + j] = (res[i + j] + a[i] * b[j]) % MOD;
        }
    }
    reduce(res, c);
    return res;
}

int64_t kth_term(const std::vector<int64_t> &rec, const std::vector<int64_t> &init, int64_t n) {
    assert(n >= 0 && init.size() >= rec.size());
    std::vector<int64_t> c(rec);
    for (int64_t &x : c) {
        x = (x % MOD + MOD) % MOD;
    }
    int L = static_cast<int>(c.size());
    if (n < static_cast<int64_t>(init.size())) {
        return ((init[n] % MOD) + MOD) % MOD;
    }
    if (L == 0) {
        return 0;
    }
    std::vector<int64_t> result(L, 0), base(L, 0);
    result[0] = 1; // The polynomial 1 = x^0.
    if (L == 1) {
        base[0] = c[0]; // x is congruent to c_0 modulo f.
    } else {
        base[1] = 1; // The polynomial x.
    }
    for (int64_t e = n; e > 0; e >>= 1) {
        if (e & 1) {
            result = combine(result, base, c);
        }
    }
}
```

```

    base = combine(base, base, c);
}
int64_t ans = 0;
for (int j = 0; j < L; j++) {
    ans = (ans + result[j] * (((init[j] % MOD) + MOD) % MOD)) % MOD;
}
return ans;
}

```

Example Usage

```

using namespace std;

int main() {
    // Fibonacci s_i = s_{i-1} + s_{i-2} with s_0 = s_1 = 1: 1, 1, 2, 3, 5, 8, ...
    vector<int64_t> fib_rec{1, 1}, fib_init{1, 1};
    assert(kth_term(fib_rec, fib_init, 0) == 1);
    assert(kth_term(fib_rec, fib_init, 9) == 55);
    assert(kth_term(fib_rec, fib_init, 10) == 89);

    // Tribonacci s_i = s_{i-1} + s_{i-2} + s_{i-3} with s_0,s_1,s_2 = 0,1,1.
    vector<int64_t> trib_rec{1, 1, 1}, trib_init{0, 1, 1};
    assert(kth_term(trib_rec, trib_init, 8) == 44);

    // Geometric s_i = 3 s_{i-1}, s_0 = 1: matches 3^n modulo MOD.
    vector<int64_t> geo_rec{3}, geo_init{1};
    int64_t pw = 1;
    for (int i = 0; i < 50; i++) {
        assert(kth_term(geo_rec, geo_init, i) == pw);
        pw = pw * 3 % MOD;
    }
    assert(kth_term({-1}, {1}, 5) == MOD - 1);
    return 0;
}

```

6.6.6 Polynomial Interpolation

Recovers the unique polynomial of degree less than n that passes through n given points, working modulo a prime. The x coordinates must be distinct; the y coordinates may be anything. Two operations are provided: reconstructing the full coefficient vector, and evaluating the interpolating polynomial at a single point without ever forming the coefficients.

Coefficient recovery uses Newton's divided differences. It first replaces the sample values by their successive divided differences in place, then expands the Newton basis $1, (x - x_0), (x - x_0)(x - x_1), \dots$ into ordinary coefficients by multiplying through one factor at a time. Point evaluation instead applies the Lagrange formula directly, summing each sample weighted by the product of $(t - x_j)/(x_k - x_j)$ over the other nodes. Coefficient recovery performs $O(n^2)$ modular inverses, while point evaluation performs n .

- `interpolate(x, y)` returns the coefficients $a_0, \dots, a_{n-1}$ (constant term first) of the polynomial P with $P(\mathbf{x}[i]) = \mathbf{y}[i]$, modulo `MOD`. The entries of `x` must be distinct modulo `MOD`.
- `interpolate_at(x, y, t)` returns $P(\mathbf{t})$ modulo `MOD` for that same polynomial, without building its coefficients.

Time Complexity: $O(n^2 \log p)$ per call to `interpolate()` and $O(n^2 + n \log p)$ per call to `interpolate_at()`, where n is the number of points and p is `MOD`.

Space Complexity: $O(n)$ auxiliary.

```

#include <cassert>
#include <cstdint>

```

```

#include <vector>

const int64_t MOD = 998244353;

int64_t powmod(int64_t b, int64_t e, int64_t m) {
    int64_t res = 1;
    for (b %= m; e > 0; e >>= 1) {
        if (e & 1) {
            res = res * b % m;
        }
        b = b * b % m;
    }
    return res;
}

int64_t modnorm(int64_t a) { return (a %= MOD) < 0 ? a + MOD : a; }
int64_t addmod(int64_t a, int64_t b) { return modnorm(modnorm(a) + modnorm(b)); }
int64_t submod(int64_t a, int64_t b) { return modnorm(modnorm(a) - modnorm(b)); }
int64_t inv(int64_t a) { return powmod(modnorm(a), MOD - 2, MOD); }

std::vector<int64_t> interpolate(std::vector<int64_t> x, std::vector<int64_t> y) {
    int n = static_cast<int>(x.size());
    assert(x.size() == y.size());
    for (int i = 0; i < n; i++) {
        x[i] = modnorm(x[i]);
        y[i] = modnorm(y[i]);
    }
    // Replace y by its divided differences in place.
    for (int k = 0; k < n; k++) {
        for (int i = k + 1; i < n; i++) {
            y[i] = submod(y[i], y[k]) * inv(submod(x[i], x[k])) % MOD;
        }
    }
    // Expand the Newton basis into ordinary coefficients.
    std::vector<int64_t> res(n, 0), temp(n, 0);
    temp[0] = 1;
    for (int k = 0; k < n; k++) {
        int64_t last = 0;
        for (int i = 0; i < n; i++) {
            res[i] = addmod(res[i], y[k] * temp[i] % MOD);
            int64_t old = temp[i];
            temp[i] = last;
            last = old;
            temp[i] = submod(temp[i], last * x[k] % MOD);
        }
    }
    return res;
}

int64_t interpolate_at(const std::vector<int64_t> &x, const std::vector<int64_t> &y, int64_t t) {
    int n = static_cast<int>(x.size());
    assert(x.size() == y.size());
    t = modnorm(t);
    int64_t res = 0;
    for (int k = 0; k < n; k++) {
        int64_t num = 1, den = 1;
        for (int j = 0; j < n; j++) {
            if (j != k) {
                num = num * submod(t, x[j]) % MOD;
                den = den * submod(x[k], x[j]) % MOD;
            }
        }
        res = addmod(res, modnorm(y[k]) * num % MOD * inv(den) % MOD);
    }
    return res;
}

```

Example Usage

```
#include <cassert>
using namespace std;

int main() {
    // P(x) = x^2 + 2x + 3 sampled at x = 0, 1, 2 gives y = 3, 6, 11.
    vector<int64_t> x{0, 1, 2}, y{3, 6, 11};
    vector<int64_t> coef = interpolate(x, y);
    assert((coef == vector<int64_t>{3, 2, 1}));
    assert(interpolate_at(x, y, 5) == 38); // 25 + 10 + 3.

    // Non-consecutive nodes for P(x) = 2x^3 - x + 4.
    auto poly = [](int64_t t) { return modnorm(2 * t * t * t - t + 4); };
    vector<int64_t> xs{2, 5, 7, 10}, ys;
    for (int64_t v : xs) {
        ys.push_back(poly(v));
    }
    vector<int64_t> c = interpolate(xs, ys);
    assert((c == vector<int64_t>{4, MOD - 1, 0, 2})); // 4 - x + 2x^3.
    assert(interpolate_at(xs, ys, 13) == poly(13));
    return 0;
}
```

6.7 Game Theory

6.7.1 Nim

Solves normal-play nim and related impartial games where each position is the disjoint sum of independent piles. In normal play, the player who makes the last move wins. A nim position is losing exactly when the XOR of all pile sizes is 0.

This is the base case of the Sprague-Grundy theorem: every finite impartial game position is equivalent to a Nim pile of size equal to its Grundy number, and sums of games combine by XOR.

- `nim_sum(piles)` returns the XOR of all pile sizes.
- `first_player_wins(piles)` returns whether the player to move has a winning strategy.
- `winning_move(piles)` returns a move (`pile`, `new_size`) that moves to XOR 0, or $(-1, -1)$ if the position is losing.

Time Complexity: $O(n)$ per call, where n is the number of piles.

Space Complexity: $O(1)$ auxiliary.

```
#include <utility>
#include <vector>

int nim_sum(const std::vector<int> &piles) {
    int x = 0;
    for (int p : piles) {
        x ^= p;
    }
    return x;
}

bool first_player_wins(const std::vector<int> &piles) {
    return nim_sum(piles) != 0;
}

std::pair<int, int> winning_move(const std::vector<int> &piles) {
    int x = nim_sum(piles);
```



```

    if (x == 0) {
        return {-1, -1};
    }
    for (int i = 0; i < static_cast<int>(piles.size()); i++) {
        if (int target = piles[i] ^ x; target < piles[i]) {
            return {i, target};
        }
    }
    return {-1, -1};
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> piles{3, 4, 5};
    // 3 ^ 4 ^ 5 = 2, so the position is winning.
    assert(nim_sum(piles) == 2);
    assert(first_player_wins(piles));
    auto [pile_idx, new_size] = winning_move(piles);
    piles[pile_idx] = new_size;
    // A winning move always leaves XOR 0 for the opponent.
    assert(nim_sum(piles) == 0);

    vector<int> losing{1, 2, 3};
    assert(!first_player_wins(losing));
    assert(winning_move(losing).first == -1);
    return 0;
}

```

6.7.2 Nim Product

Nimbers are the values of impartial games, added by XOR. Under the additional operation of nim multiplication they form a field: the nonnegative integers below 2^{2^k} are closed under both operations, with XOR as addition and nim product as multiplication. Writing $\otimes$ for the nim product and $\oplus$ for XOR, nim multiplication is commutative, associative, and distributive over $\oplus$, and is defined recursively by $a \otimes b = \text{mex}\{(a' \otimes b) \oplus (a \otimes b') \oplus (a' \otimes b') : 0 \leq a' < a, 0 \leq b' < b\}$. Its main use is multiplying the Grundy values of independent “coin-turning” games whose move sets combine as a product, such as two-dimensional turning games built from one-dimensional ones.

Rather than evaluate the recursive definition directly, this implementation precomputes the nim product of every pair of single bits $2^i \otimes 2^j$ for $i, j < 64$. Each such product follows from the Fermat-like rule $2^{2^k} \otimes 2^{2^k} = 2^{2^k} \oplus 2^{2^k-1}$ and the fact that distinct Fermat powers multiply like ordinary powers of two. A full product then XORs together the bit-pair products selected by the set bits of the operands. Because every 64-bit value lies in the field of size 2^{64} , the operation is total on `uint64_t` and the nonzero values have multiplicative inverses.

- `nim_product(x, y)` returns the nim product of `x` and `y`.
- `nim_pow(b, e)` returns `b` raised to the `e`-th power under nim multiplication.
- `nim_inverse(b)` returns the multiplicative inverse of a nonzero `b`, so that `nim_product(b, nim_inverse(b)) == 1`.

Time Complexity:

- $O(1)$ table construction on first use (a fixed 64 by 64 table).
- $O(64^2)$ per call to `nim_product()`; $O(64^2 \log e)$ per call to `nim_pow()` and `nim_inverse()`.

Space Complexity: $O(1)$ for a fixed 64 by 64 table of 64-bit products.

```

#include <cassert>

```

```

#include <cstdint>

uint64_t bit_prod[64][64];

// Build nim products for pairs of single-bit values. Each entry depends only on entries with a
// smaller first index, so a single ascending pass fills the table.
const bool nim_product_init = [] {
    for (int i = 0; i < 64; i++) {
        for (int j = 0; j < 64; j++) {
            if ((i & j) == 0) {
                bit_prod[i][j] = uint64_t(1) << (i | j); // Distinct Fermat powers multiply like 2^(i|j).
            } else {
                int a = (i & j) & ~(i & j); // Lowest Fermat power shared by both exponents.
                bit_prod[i][j] = bit_prod[i ^ a][j] ^ bit_prod[(i ^ a) | (a - 1)][(j ^ a) | (i & (a - 1))];
            }
        }
    }
    return true;
}();

uint64_t nim_product(uint64_t x, uint64_t y) {
    uint64_t res = 0;
    for (int i = 0; i < 64 && (x >> i) != 0; i++) {
        if ((x >> i) & 1) {
            for (int j = 0; j < 64 && (y >> j) != 0; j++) {
                if ((y >> j) & 1) {
                    res ^= bit_prod[i][j];
                }
            }
        }
    }
    return res;
}

uint64_t nim_pow(uint64_t b, uint64_t e) {
    uint64_t res = 1;
    for (; e > 0; e >>= 1) {
        if (e & 1) {
            res = nim_product(res, b);
        }
        b = nim_product(b, b);
    }
    return res;
}

uint64_t nim_inverse(uint64_t b) {
    assert(b != 0);
    // The multiplicative group of the size-2^64 field has order 2^64 - 1, so b^(2^64 - 2) = b^{-1}.
    return nim_pow(b, static_cast<uint64_t>(-2));
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    // Smallest nontrivial products: {0,1,2,3} form the field GF(4) under XOR and nim product.
    assert(nim_product(2, 2) == 3);
    assert(nim_product(2, 3) == 1);
    assert(nim_product(3, 3) == 2);
    assert(nim_product(1, 12345) == 12345); // 1 is the multiplicative identity.
    assert(nim_product(0, 12345) == 0);

    // Field axioms on a range of values: commutativity, distributivity over XOR, and inverses.
    for (uint64_t a = 0; a < 64; a++) {

```

```

for (uint64_t b = 0; b < 64; b++) {
    assert(nim_product(a, b) == nim_product(b, a));
    for (uint64_t c = 0; c < 8; c++) {
        assert(nim_product(a, b ^ c) == (nim_product(a, b) ^ nim_product(a, c)));
    }
}
if (a != 0) {
    assert(nim_product(a, nim_inverse(a)) == 1);
}
}
assert(nim_product(0x1234567890abcdefULL, nim_inverse(0x1234567890abcdefULL)) == 1);
return 0;
}

```

6.7.3 Sprague-Grundy

Computes Grundy numbers for finite impartial games. The Grundy number of a game position is the minimum excluded nonnegative integer (MEX) of the Grundy numbers reachable in one move. A position is losing exactly when its Grundy number is 0.

For a sum of independent impartial games, the combined Grundy number is the XOR of the component Grundy numbers. This is the Sprague-Grundy theorem.

- `subtraction_game_grundy(max_stones, moves)` computes Grundy numbers for a game where a move subtracts one positive value in `moves` from the pile. Nonpositive values are ignored.
- `sum_grundy(grundies)` returns the XOR of component Grundy numbers.

Time Complexity:

- $O(n \cdot m)$ per call to `subtraction_game_grundy(max_stones, moves)`, where n is `max_stones` and m is the number of allowed moves.
- $O(k)$ per call to `sum_grundy()`, where k is the input length.

Space Complexity:

- $O(n)$ output and $O(m)$ auxiliary for `subtraction_game_grundy(max_stones, moves)`.
- $O(1)$ auxiliary for `sum_grundy(grundies)`.

```

#include <vector>

int mex(const std::vector<int> &values) {
    std::vector<char> seen(values.size() + 1, false);
    for (int v : values) {
        if (0 <= v && v < static_cast<int>(seen.size())) {
            seen[v] = true;
        }
    }
    for (int i = 0; i < static_cast<int>(seen.size()); i++) {
        if (!seen[i]) {
            return i;
        }
    }
    return static_cast<int>(seen.size());
}

std::vector<int> subtraction_game_grundy(int max_stones, const std::vector<int> &moves) {
    std::vector<int> grundy(max_stones + 1, 0);
    for (int stones = 1; stones <= max_stones; stones++) {
        std::vector<int> reachable;
        for (int m : moves) {
            if (0 < m && m <= stones) {
                reachable.push_back(grundy[stones - m]);
            }
        }
    }
}

```

```

    }
    Grundy[stones] = mex(reachable);
}
return Grundy;
}

int sum_Grundy(const std::vector<int> &grundies) {
    int x = 0;
    for (int g : grundies) {
        x ^= g;
    }
    return x;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<int> moves{1, 3, 4};
    auto g = subtraction_game_Grundy(10, moves);
    // With moves 1, 3, 4, pile size 2 is losing and pile size 4 has mex {0, 1} = 2.
    assert(g[0] == 0);
    assert(g[1] == 1);
    assert(g[2] == 0);
    assert(g[4] == 2);
    assert(g[7] == 0);

    vector<int> heaps{g[5], g[7], g[10]};
    // Independent games combine by xor of their Grundy numbers.
    assert(sum_Grundy(heaps) == (g[5] ^ g[7] ^ g[10]));
    return 0;
}

```

6.7.4 Grundy Numbers on DAG

Computes Grundy numbers for impartial games whose positions form a directed acyclic graph. Each node is a game position, and each outgoing edge is a legal move. Terminal nodes have Grundy number 0; all other nodes take the MEX of their successors' Grundy numbers.

The implementation uses memoized DFS. It assumes the graph is acyclic; if cycles are present, the usual finite impartial-game Grundy recurrence is not directly valid without additional analysis.

- `Grundy_on_dag(g)` returns the Grundy number of every node in graph `g`, where `g[u]` contains all positions reachable from position `u` in one move.

Time Complexity: $O(n + m)$ per call, where n is the number of nodes and m is the number of edges.

Space Complexity: $O(d)$ auxiliary stack space and $O(n + m)$ auxiliary heap space, where d is DFS depth.

```

#include <vector>

int Grundy_dfs(int u, const std::vector<std::vector<int>> &g, std::vector<int> *memo) {
    if ((*memo)[u] != -1) {
        return (*memo)[u];
    }
    std::vector<char> seen(g[u].size() + 1, false);
    for (int v : g[u]) {
        int x = Grundy_dfs(v, g, memo);
        if (x < static_cast<int>(seen.size())) {
            seen[x] = true;
        }
    }
    int mex = 0;
    while (seen[mex]) mex++;
    (*memo)[u] = mex;
    return mex;
}

```

```

    }
}
int res = 0;
while (res < static_cast<int>(seen.size()) && seen[res]) {
    res++;
}
(*memo)[u] = res;
return res;
}

std::vector<int> GrundyOnDag(const std::vector<std::vector<int>> &g) {
    std::vector<int> memo(g.size(), -1);
    for (int u = 0; u < static_cast<int>(g.size()); u++) {
        GrundyDFS(u, g, &memo);
    }
    return memo;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<vector<int>> g{{1, 2}, {3}, {3, 4}, {}, {}};
    auto Grundy = GrundyOnDag(g);
    // Terminal nodes have Grundy 0; their predecessors have mex {0} = 1.
    assert(Grundy[3] == 0);
    assert(Grundy[4] == 0);
    assert(Grundy[1] == 1);
    assert(Grundy[2] == 1);
    // Node 0 can move to two positions with Grundy 1, so mex {1} = 0.
    assert(Grundy[0] == 0);
    return 0;
}

```

6.7.5 Minimax and Alpha-Beta Pruning

Searches a finite two-player, zero-sum, perfect-information game tree using minimax and alpha-beta pruning.

Minimax assumes both players play optimally: the maximizing player chooses the child with largest value, and the minimizing player chooses the child with smallest value.

Alpha-beta pruning computes the same value as minimax, but avoids exploring branches that cannot affect the final answer. It is most effective when good moves are searched first.

The example game is “take 1 or 2 stones”: starting with `stones` stones, players alternate taking 1 or 2 stones, and the player who takes the last stone wins. The evaluation returns 1 when the maximizing player wins and -1 when the minimizing player wins.

- `minimax(stones, maximizing)` returns the exact game-tree value, where `maximizing` indicates whose turn it is.
- `alpha_beta(stones, maximizing, alpha = -2, beta = 2)` returns the same value, with pruning using the current lower bound `alpha` and upper bound `beta`.
- `best_take(stones)` returns an optimal first move using alpha-beta pruning, either 1 or 2, for positive `stones`.

Time Complexity: $O(b^d)$ per call in the worst case, where b is branching factor and d is search depth.

Alpha-beta can achieve $O(b^{d/2})$ with optimal move ordering.

Space Complexity: $O(d)$ auxiliary recursion stack space.

```

#include <algorithm>

```

```

int minimax(int stones, bool maximizing) {
    if (stones == 0) {
        return maximizing ? -1 : 1;
    }
    if (maximizing) {
        int best = -2;
        for (int take = 1; take <= 2 && take <= stones; take++) {
            best = std::max(best, minimax(stones - take, false));
        }
        return best;
    } else {
        int best = 2;
        for (int take = 1; take <= 2 && take <= stones; take++) {
            best = std::min(best, minimax(stones - take, true));
        }
        return best;
    }
}

int alpha_beta(int stones, bool maximizing, int alpha = -2, int beta = 2) {
    if (stones == 0) {
        return maximizing ? -1 : 1;
    }
    if (maximizing) {
        int value = -2;
        for (int take = 1; take <= 2 && take <= stones; take++) {
            value = std::max(value, alpha_beta(stones - take, false, alpha, beta));
            alpha = std::max(alpha, value);
            if (alpha >= beta) {
                break;
            }
        }
        return value;
    } else {
        int value = 2;
        for (int take = 1; take <= 2 && take <= stones; take++) {
            value = std::min(value, alpha_beta(stones - take, true, alpha, beta));
            beta = std::min(beta, value);
            if (alpha >= beta) {
                break;
            }
        }
        return value;
    }
}

int best_take(int stones) {
    int best_move = 1, best_value = -2;
    for (int take = 1; take <= 2 && take <= stones; take++) {
        int value = alpha_beta(stones - take, false);
        if (value > best_value) {
            best_value = value;
            best_move = take;
        }
    }
    return best_move;
}

```

Example Usage

```

#include <cassert>

int main() {
    assert(minimax(1, true) == 1);
    assert(minimax(3, true) == -1);
    for (int stones = 1; stones <= 10; stones++) {

```

```
    assert(minimax(stones, true) == alpha_beta(stones, true));  
}  
assert(best_take(4) == 1); // Move to losing state with 3 stones.  
assert(best_take(5) == 2); // Move to losing state with 3 stones.  
return 0;  
}
```

Chapter 7

Geometry

7.1 Geometry Primitives

7.1.1 Point

A 2D point class templated on the coordinate type `T`. Integer types support some exact operations, while metric operations (`norm`, `arg`, `normalize`, `rotateCW/CCW`, `reflect` across a line) return coordinates of type `fp_t`, which is `T` when `T` is floating-point and `double` otherwise.

Exact operations (return `TPoint<T>` or `T`, no precision lost for integers) include: element-wise arithmetic, `dot()`, `cross()`, `sqnorm()`, cardinal rotations, `reflect(point)`, equality, and lexicographic ordering. `operator==` and `operator<` are exact for every coordinate type, as required by standard algorithms and containers. `EQ(p, q)` compares floating-point point coordinates using `EPS` and remains exact for other coordinate types.

Overflow warning: the exact products `dot()`, `cross()`, and `sqnorm()` grow like the squared coordinate magnitude. With `TPoint<int>` these overflow a 32-bit `int` once coordinates exceed roughly a few tens of thousands, so use `PointL` (`TPoint<long long>`) for larger integer coordinates.

Metric operations (return `TPoint<fp_t>` or `fp_t`):

- `norm()`, `arg()`, `proj()`, `normalize()`, `rotateCW()`, `rotateCCW()`, `reflect(line)`.
- `operator/` also promotes to `fp_t`.

Implicit conversion from `TPoint<U>` to `TPoint<T>` is provided when `U` is integral and `T` is floating-point, so `PointI` can be passed wherever `PointD` is expected.

Non-floating-point coordinate types such as `Modular` or `Rational` compose for the exact operations (arithmetic, `dot`, `cross`, `sqnorm`, comparisons, and `EQ`) using the coordinate type's own operators. Their metric operations convert coordinates explicitly to `double`; for `Modular`, this uses the stored representative and is not finite-field geometry.

Type aliases:

- `PointI = TPoint<int>`: exact integer geometry (small coordinates only; see overflow warning)
- `PointL = TPoint<long long>`: exact integer geometry for large coordinates
- `PointD = TPoint<double>`: standard floating-point
- `PointLD = TPoint<long double>`: extra precision
- `Point = PointD`: default point type is double

Time Complexity: $O(1)$ per call to the constructor and all other operations.

Space Complexity:

- $O(1)$ for storage of the point.

- $O(1)$ auxiliary for all operations.

```

#include <cmath>
#include <ostream>
#include <tuple>
#include <type_traits>
#include <utility>

const double EPS = 1e-9;

// Epsilon-aware for floating-point coordinates; exact for int and for coordinate types like Modular
// or Rational, which therefore compose for all of the predicates below.
template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T>
struct TPoint {
    // Metric operations preserve floating-point coordinates and otherwise promote to double.
    using fp_t = std::conditional_t<std::is_floating_point<T>::value, T, double>;

    T x, y;

    TPoint() : x(0), y(0) {}
    TPoint(T x, T y) : x(x), y(y) {}
    explicit TPoint(const std::pair<T, T> &p) : x(p.first), y(p.second) {}

    // Implicit conversion from integral to floating-point point (optional, can just use to_double()).
    template<
        typename U,
        typename = std::enable_if_t<std::is_integral<U>::value && std::is_floating_point<T>::value>>
    TPoint(const TPoint<U> &p) : x(static_cast<T>(p.x)), y(static_cast<T>(p.y)) {}

    friend bool EQ(const TPoint &a, const TPoint &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }

    bool operator==(const TPoint &p) const { return x == p.x && y == p.y; }
    bool operator!=(const TPoint &p) const { return !(*this == p); }
    bool operator<(const TPoint &p) const { return std::tie(x, y) < std::tie(p.x, p.y); }
    bool operator>(const TPoint &p) const { return p < *this; }
    bool operator<=(const TPoint &p) const { return !(*this > p); }
    bool operator>=(const TPoint &p) const { return !(*this < p); }
    TPoint operator+(const TPoint &p) const { return {x + p.x, y + p.y}; }
    TPoint operator-(const TPoint &p) const { return {x - p.x, y - p.y}; }
    TPoint operator+(T v) const { return {x + v, y + v}; }
    TPoint operator-(T v) const { return {x - v, y - v}; }
    TPoint operator*(T v) const { return {x * v, y * v}; }

    // Division always promotes to fp_t to avoid integer truncation.
    TPoint<fp_t> operator/(fp_t v) const { return {(fp_t)x / v, (fp_t)y / v}; }

    TPoint &operator+=(const TPoint &p) { x += p.x; y += p.y; return *this; }
    TPoint &operator-=(const TPoint &p) { x -= p.x; y -= p.y; return *this; }
    TPoint &operator+=(T v) { x += v; y += v; return *this; }
    TPoint &operator-=(T v) { x -= v; y -= v; return *this; }
    TPoint &operator*=(T v) { x *= v; y *= v; return *this; }
    friend TPoint operator+(T v, const TPoint &p) { return p + v; }
    friend TPoint operator*(T v, const TPoint &p) { return p * v; }

    // --- Exact operations: return T or TPoint<T>, work for any coordinate type ---

```

```

T sqnorm() const { return x * x + y * y; } // Overflow warning!
T dot(const TPoint &p) const { return x * p.x + y * p.y; } // Overflow warning!
T cross(const TPoint &p) const { return x * p.y - y * p.x; } // Overflow warning!

// --- Floating-point operations: return fp_t or TPoint<fp_t> ---

fp_t norm() const { return hypot(static_cast<fp_t>(x), static_cast<fp_t>(y)); }
fp_t arg() const { return atan2(static_cast<fp_t>(y), static_cast<fp_t>(x)); }

fp_t proj(const TPoint &p) const {
    return (static_cast<fp_t>(x) * p.x + static_cast<fp_t>(y) * p.y) / p.norm();
}

TPoint<fp_t> normalize() const {
    fp_t n = norm();
    if (n < static_cast<fp_t>(1e-30)) { // guard against dividing by a near-zero norm, not a
        return {fp_t(0), fp_t(0)}; // geometric tolerance (EPS is too coarse here)
    }
    return {static_cast<fp_t>(x) / n, static_cast<fp_t>(y) / n};
}

// --- Cardinal rotations: exact for all coordinate types including int ---

// Returns (x, y) rotated 90/180/270 degrees counter-clockwise about the origin.
TPoint rotate90() const { return {-y, x}; }
TPoint rotate180() const { return {-x, -y}; }
TPoint rotate270() const { return {y, -x}; }

// Returns (x, y) rotated 90/180/270 degrees counter-clockwise about point p.
TPoint rotate90(const TPoint &p) const { return (*this - p).rotate90() + p; }
TPoint rotate180(const TPoint &p) const { return (*this - p).rotate180() + p; }
TPoint rotate270(const TPoint &p) const { return (*this - p).rotate270() + p; }

// --- Arbitrary-angle rotations: always return floating-point ---

// Returns (x, y) rotated t radians clockwise about the origin.
TPoint<fp_t> rotateCW(fp_t t) const {
    fp_t fx = (fp_t)x, fy = (fp_t)y;
    return {fx * cos(t) + fy * sin(t), fy * cos(t) - fx * sin(t)};
}

// Returns (x, y) rotated t radians counter-clockwise about the origin.
TPoint<fp_t> rotateCCW(fp_t t) const {
    fp_t fx = (fp_t)x, fy = (fp_t)y;
    return {fx * cos(t) - fy * sin(t), fx * sin(t) + fy * cos(t)};
}

// Returns (x, y) rotated t radians clockwise about point p.
TPoint<fp_t> rotateCW(const TPoint &p, fp_t t) const {
    return TPoint<fp_t>{(fp_t)x - p.x, (fp_t)y - p.y}.rotateCW(t) +
        TPoint<fp_t>{(fp_t)p.x, (fp_t)p.y};
}

// Returns (x, y) rotated t radians counter-clockwise about point p.
TPoint<fp_t> rotateCCW(const TPoint &p, fp_t t) const {
    return TPoint<fp_t>{(fp_t)x - p.x, (fp_t)y - p.y}.rotateCCW(t) +
        TPoint<fp_t>{(fp_t)p.x, (fp_t)p.y};
}

// --- Reflections ---

// Returns (x, y) reflected across point p. Exact for any coordinate type.
TPoint reflect(const TPoint &p) const { return {2 * p.x - x, 2 * p.y - y}; }

// Returns (x, y) reflected across the line containing points p and q.
// Always returns floating-point coordinates.
TPoint<fp_t> reflect(const TPoint &p, const TPoint &q) const {
    TPoint<fp_t> fp{(fp_t)p.x, (fp_t)p.y};

```

```

    if (EQ(p, q)) {
        return TPoint<fp_t>{(fp_t)x, (fp_t)y}.reflect(fp);
    }
    TPoint<fp_t> r{(fp_t)x - p.x, (fp_t)y - p.y}, s{(fp_t)q.x - p.x, (fp_t)q.y - p.y};
    fp_t ssq = s.sqnorm();
    r = TPoint<fp_t>{(r.x * s.x + r.y * s.y) / ssq, (r.x * s.y - r.y * s.x) / ssq};
    return TPoint<fp_t>{r.x * s.x - r.y * s.y + fp.x, r.x * s.y + r.y * s.x + fp.y};
}

// --- Explicit type conversions ---

TPoint<double> to_double() const { return {static_cast<double>(x), static_cast<double>(y)}; }

TPoint<long double> to_ldouble() const {
    return {static_cast<long double>(x), static_cast<long double>(y)};
}

// --- Friend free-function versions ---

friend T sqnorm(const TPoint &p) { return p.sqnorm(); }
friend fp_t norm(const TPoint &p) { return p.norm(); }
friend fp_t arg(const TPoint &p) { return p.arg(); }
friend T dot(const TPoint &p, const TPoint &q) { return p.dot(q); }
friend T cross(const TPoint &p, const TPoint &q) { return p.cross(q); }
friend fp_t proj(const TPoint &p, const TPoint &q) { return p.proj(q); }
friend TPoint<fp_t> normalize(const TPoint &p) { return p.normalize(); }
friend TPoint rotate90(const TPoint &p) { return p.rotate90(); }
friend TPoint rotate180(const TPoint &p) { return p.rotate180(); }
friend TPoint rotate270(const TPoint &p) { return p.rotate270(); }
friend TPoint rotate90(const TPoint &p, const TPoint &q) { return p.rotate90(q); }
friend TPoint rotate180(const TPoint &p, const TPoint &q) { return p.rotate180(q); }
friend TPoint rotate270(const TPoint &p, const TPoint &q) { return p.rotate270(q); }
friend TPoint<fp_t> rotateCW(const TPoint &p, fp_t t) { return p.rotateCW(t); }
friend TPoint<fp_t> rotateCCW(const TPoint &p, fp_t t) { return p.rotateCCW(t); }
friend TPoint<fp_t> rotateCW(const TPoint &p, const TPoint &q, fp_t t) { return p.rotateCW(q, t); }
friend TPoint<fp_t> rotateCCW(const TPoint &p, const TPoint &q, fp_t t) { return p.rotateCCW(q, t); }
friend TPoint reflect(const TPoint &p, const TPoint &q) { return p.reflect(q); }
friend TPoint<fp_t> reflect(const TPoint &p, const TPoint &a, const TPoint &b) { return p.reflect(a, b); }

friend std::ostream &operator<<(std::ostream &out, const TPoint &p) {
    if constexpr (std::is_floating_point<T>::value) {
        return out << "(" << (fabs(p.x) < EPS ? 0 : p.x) << "," << (fabs(p.y) < EPS ? 0 : p.y) << ")";
    }
    return out << "(" << p.x << "," << p.y << ")";
}
};

using PointI = TPoint<int>;
using PointL = TPoint<long long>;
using PointD = TPoint<double>;
using PointLD = TPoint<long double>;
using Point = PointD; // Default point type is double.

// Can compose with numerical types from chapter 6:
// using PointB = TPoint<BigInt>;
// using PointR = TPoint<Rational<int64_t>>;
// using PointM = TPoint<Modular<1000000007>>;

```

Example Usage

```

#include <cassert>

const double PI = acos(-1.0);

int main() {
    Point p(-10, 3);

```

```

assert(EQ(Point(-18, 29), p + Point(-3, 9) * 6.0 / 2.0 - Point(-1, 1)));
assert(EQ(109, p.sqnorm()));
assert(EQ(10.44030650891, p.norm()));
assert(EQ(2.850135859112, p.arg()));
assert(EQ(0, p.dot(Point(3, 10))));
assert(EQ(0, p.cross(Point(10, -3))));
assert(EQ(10, p.proj(Point(-10, 0))));
assert(EQ(1, p.normalize().norm()));
assert(EQ(Point(-3, -10), p.rotate90()));
assert(EQ(Point(10, -3), p.rotate180()));
assert(EQ(Point(3, 10), p.rotate270()));
assert(EQ(Point(3, 12), p.rotateCW(Point(1, 1), PI / 2)));
assert(EQ(Point(1, -10), p.rotateCCW(Point(2, 2), PI / 2)));
assert(EQ(Point(10, -3), p.reflect(Point(0, 0))));
assert(EQ(Point(-10, -3), p.reflect(Point(-2, 0), Point(5, 0))));

// Integer point - exact arithmetic, float-only ops return PointD.
PointI a(3, 4), b(1, 0);
assert(a + b == PointI(4, 4));
assert(a - b == PointI(2, 4));
assert(a * 2 == PointI(6, 8));
assert(a.sqnorm() == 25);
assert(a.dot(b) == 3);
assert(a.cross(b) == -4);
assert(EQ(a.norm(), 5.0));
assert(a.rotate90() == PointI(-4, 3));
assert(a.rotate180() == PointI(-3, -4));
assert(a.rotate270() == PointI(4, -3));
PointD anorm = a.normalize(); // returns PointD
assert(EQ(anorm.norm(), 1.0));
PointD adiv = a / 2.0; // returns PointD (division always promotes)
assert(EQ(adiv.x, 1.5) && EQ(adiv.y, 2.0));
auto rotated = PointI(1, 0).rotateCCW(PI / 4); // returns PointD (arbitrary rotation promotes)
double root2over2 = sqrt(2.0) / 2.0;
assert(EQ(rotated.x, root2over2) && EQ(rotated.y, root2over2));

// Implicit conversion PointI -> PointD.
PointD promoted = PointI(1, 2);
assert(promoted == PointD(1, 2));
return 0;
}

```

7.1.2 Line

A straight line in two dimensions, templated on the coordinate type `T`. The line $ax + by + c = 0$ is canonicalized by default. With `EXACT = true`, coefficients are divided by their Euclidean GCD and given a fixed sign; this supports built-in integer types as well as `BigInt` and `Rational` from previous chapters. With `EXACT = false`, coefficients are scaled to $b = 1$ when b is nonzero, or to $a = 1$ for vertical lines; this supports floating-point and `Modular` coefficients. Thus proportional valid coefficient vectors have the same canonical representation, up to floating-point rounding.

Set the template flag `CANONICALIZE = false` to preserve the supplied coefficients. `EXACT` defaults to `true` only for built-in integral types; set it explicitly for custom types such as `BigInt` and `Rational`. The internal Euclidean GCD only requires `%`, comparisons, and basic arithmetic. Calling `canonicalized()` returns a canonical copy. `operator==` always compares stored coefficients exactly, while `EQ(l1, l2)` canonicalizes raw lines if necessary and uses `EPS` for floating-point coefficients.

Storage and the predicates `contains()`, `is_parallel()`, and `is_perpendicular()` stay exact for integral type `T`. The inherently-fractional operations `slope()`, `x()`, and `y()` return `fp_t`, which is `T` when `T` is floating-point and `double` otherwise.

Operations include `horizontal()`, `vertical()`, `slope()`, evaluating y at some x with `y()` (and vice versa with `x()`), checking if a point falls on the line with `contains()`, checking if another line is parallel or perpendicular with `is_parallel()` or `is_perpendicular()`, and finding the `parallel()` or `perpendicular()` line through a point.

Overflow warning: the exact predicates form products of coefficients (e.g. cross terms in `is_parallel()`/`is_perpendicular()` and `a*p.x + b*p.y` in `contains()`), which grow like the squared coordinate magnitude. For integral type `T`, use `LineL` (`TLine<int64_t>`) once coordinates reach a few tens of thousands, or the 32-bit products may overflow.

Type aliases:

- `LineI = TLine<int>`: exact integer-coefficient lines (small values only; see overflow warning)
- `LineL = TLine<int64_t>`: exact integer-coefficient lines for large coordinates
- `LineD = TLine<double>`: standard floating-point
- `LineLD = TLine<long double>`: extra precision
- `Line = LineD`: default line type is double

Time Complexity:

- $O(\log C)$ per integral canonicalizing constructor or call to `canonicalize()`, where C is the largest absolute coefficient.
- $O(1)$ per call to all other operations.

Space Complexity:

- $O(1)$ for storage of the line.
- $O(1)$ auxiliary for all operations.

```
#include <cmath>
#include <cstdint>
#include <limits>
#include <ostream>
#include <type_traits>
#include <utility>

const double EPS = 1e-9;
const double M_NAN = std::numeric_limits<double>::quiet_NaN();

// Epsilon-aware for floating-point coordinates; exact for int and for coordinate types like Modular
// or Rational, which therefore compose for all of the predicates below.
template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

// SFINAE guard: valid only when Pt exposes numeric .x/.y members, so templated point constructors
// don't hijack scalar-argument calls.
template<typename Pt>
using if_point = decltype(std::declval<const Pt &>().x, std::declval<const Pt &>().y, void());

template<typename T, bool CANONICALIZE = true, bool EXACT = std::is_integral<T>::value>
struct TLine {
    // Slope and solve operations preserve floating-point coordinates and otherwise promote to double.
    using fp_t = std::conditional_t<std::is_floating_point<T>::value, T, double>;

    T a, b, c;

    void canonicalize() {
        if constexpr (EXACT) {
            auto gcd = [](T x, T y) {
                if (x < T(0)) x = -x;
                if (y < T(0)) y = -y;
                while (y != T(0)) {
                    x %= y;
                    std::swap(x, y);
                }
            };
            gcd(a, b);
            gcd(a, c);
            gcd(b, c);
        }
    }
};
```

```

    }
    return x;
};
T g = gcd(a, gcd(b, c));
if (g != 0) {
    a /= g;
    b /= g;
    c /= g;
}
if (a < 0 || (a == 0 && b < 0)) {
    a = -a;
    b = -b;
    c = -c;
}
} else if (!EQ(b, T(0))) {
    a /= b;
    c /= b;
    b = T(1);
} else if (!EQ(a, T(0))) {
    c /= a;
    a = T(1);
    b = T(0);
}
}

TLine() : a(0), b(0), c(0) {} // Invalid or uninitialized line.

TLine(T a, T b, T c) : a(a), b(b), c(c) {
    if constexpr (CANONICALIZE) canonicalize();
}

// Line through two points. Coefficients are exact (type T) for integer points.
template<typename Pt, typename = if_point<Pt>>
// Overflow warning!
TLine(const Pt &p, const Pt &q) : a(q.y - p.y), b(p.x - q.x), c(-(a * p.x + b * p.y)) {
    if constexpr (CANONICALIZE) canonicalize();
}

// Line of a given slope through a point. Intended for floating-point T (a generic slope
// produces non-integer coefficients), so this constructor is disabled for integer T.
template<
    class Pt, typename = if_point<Pt>, typename U = T,
    class = std::enable_if_t<std::is_floating_point<U>::value>>
TLine(fp_t slope, const Pt &p) : a(-slope), b(1), c(slope * p.x - p.y) {
    if constexpr (CANONICALIZE) canonicalize();
}

TLine<T, true, EXACT> canonicalized() const { return TLine<T, true, EXACT>(a, b, c); }

bool operator==(const TLine &l) const { return a == l.a && b == l.b && c == l.c; }
bool operator!=(const TLine &l) const { return !(*this == l); }

friend bool EQ(const TLine &l1, const TLine &l2) {
    if constexpr (CANONICALIZE) {
        return EQ(l1.a, l2.a) && EQ(l1.b, l2.b) && EQ(l1.c, l2.c);
    }
    return EQ(l1.canonicalized(), l2.canonicalized());
}

// Returns whether the line is a valid line (not both a and b zero).
bool valid() const { return !(EQ(a, 0) && EQ(b, 0)); }
bool horizontal() const { return valid() && EQ(a, 0); }
bool vertical() const { return valid() && EQ(b, 0); }

// Slope -a/b, or NaN if the line is vertical or invalid.
fp_t slope() const { return (!valid() || EQ(b, 0)) ? M_NAN : -(fp_t)a / b; }

// Solve for x at a given y (NaN if horizontal or invalid).

```

```

fp_t x(fp_t y) const { return (!valid() || EQ(a, 0)) ? M_NAN : -((fp_t)b * y + c) / a; }

// Solve for y at a given x (NaN if vertical or invalid).
fp_t y(fp_t x) const { return (!valid() || EQ(b, 0)) ? M_NAN : -((fp_t)a * x + c) / b; }

// Whether point p lies on the line. Exact for integer T and integer p.
template<typename Pt>
bool contains(const Pt &p) const {
    return valid() && EQ(a * p.x + b * p.y + c, 0); // Overflow warning!
}

// Parallel iff the normals are parallel (cross == 0); perpendicular iff normals are perpendicular
// (dot == 0). Both exact for integer T. Overflow warning!
bool is_parallel(const TLine &l) const { return valid() && l.valid() && EQ(a * l.b, l.a * b); }

bool is_perpendicular(const TLine &l) const {
    return valid() && l.valid() && EQ(a * l.a, -(b * l.b));
}

// Parallel/perpendicular line through point p. Exact for integer T and integer p.
template<typename Pt, typename = if_point<Pt>>
TLine parallel(const Pt &p) const {
    return TLine(a, b, -(a * p.x + b * p.y)); // Overflow warning!
}

template<typename Pt, typename = if_point<Pt>>
TLine perpendicular(const Pt &p) const {
    return TLine(-b, a, b * p.x - a * p.y); // Overflow warning!
}

friend std::ostream &operator<<(std::ostream &out, const TLine &l) {
    auto clean = [](T v) -> double {
        return fabs(static_cast<double>(v)) < EPS ? 0 : static_cast<double>(v);
    };
    return out << clean(l.a) << "x" << std::showpos << clean(l.b) << "y" << clean(l.c) << "=0"
        << std::noshowpos;
}

};

using LineI = TLine<int>;
using LineL = TLine<int64_t>;
using LineD = TLine<double>;
using LineLD = TLine<long double>;
using Line = LineD; // Default line type is double.

// Can compose with numerical types from chapter 6:
// using LineB = TLine<BigInt, true, true>;
// using LineR = TLine<Rational<int64_t>, true, true>;

```

Example Usage

```

#include <cassert>

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
};

int main() {
    Line invalid;
    assert(!invalid.contains(Point(0, 0)) && !invalid.is_parallel(invalid));
}

```

```

// Floating-point line.
Line l(2, -5, -8);
Line para = Line(2, -5, -8).parallel(Point(-6, -2));
Line perp = Line(2, -5, -8).perpendicular(Point(-6, -2));
assert(l.is_parallel(para) && l.is_perpendicular(perp));
assert(l.slope() == 0.4);
assert(EQ(para, Line(-0.4, 1, -0.4)));
assert(EQ(perp, Line(2.5, 1, 17)));
assert(l.contains(Point(1.5, -1))); // 2(1.5) - 5(-1) - 8 = 0.

// Integer line through integer points - exact canonical coefficients and predicates.
LineI il(PointI(0, 0), PointI(2, 1)); // a=1, b=-2, c=0 -> x - 2y = 0
assert(il.a == 1 && il.b == -2 && il.c == 0);
assert(il.contains(PointI(4, 2))); // exact integer check
assert(!il.contains(PointI(4, 3)));
assert(il.is_parallel(LineI(PointI(1, 1), PointI(3, 2)))); // both slope 1/2
assert(il.is_perpendicular(LineI(PointI(0, 0), PointI(1, -2)))); // dot of normals = 0
assert(il.slope() == 0.5); // fp_t result

LineI reduced(2, -4, -6);
assert(reduced == LineI(1, -2, -3));

using RawLine = TLine<double, false>;
RawLine raw(2, -5, -8), scaled(-0.4, 1, 1.6);
assert(raw != scaled); // Stored coefficients differ.
assert(EQ(raw, scaled)); // Canonicalized copies represent the same line.
RawLine canonical = raw;
canonical.canonicalize();
assert(canonical == scaled);
assert(EQ(raw.canonicalized(), Line(-0.4, 1, 1.6)));
return 0;
}

```

7.1.3 Circle

A circle in two dimensions with epsilon-aware operations. The circle centered at (h, k) is represented by the relation $(x - h)^2 + (y - k)^2 = r^2$, where the radius r is normalized to a nonnegative number.

This implementation stores and computes circle parameters in `double`. Point-accepting constructors and predicates are templated on the point type `Pt` and only read `.x`/`.y`, so they accept `Point`, `PointD`, or `PointI` from 7.1.1 or any struct with numeric `.x` and `.y` fields.

- `Circle()` constructs the zero-radius circle centered at the origin.
- `Circle(r)` constructs a circle of radius `abs(r)` centered at the origin.
- `Circle(h, k, r)` constructs a circle of radius `abs(r)` centered at (h, k) .
- `Circle(o, r)` constructs a circle of radius `abs(r)` centered at point `o`.
- `Circle(a, b)` constructs the circle whose diameter is segment `a-b`.
- `Circle(a, b, c)` constructs the circumcircle through non-collinear points `a`, `b`, and `c`, throwing if the points are collinear.
- `Circle(a, b, r)` constructs one circle of radius `abs(r)` passing through points `a` and `b`, throwing if no finite circle is determined by the inputs.
- `operator==` and `operator!=` compare stored centers and radii exactly; `EQ(a, b)` compares them using `EPS`.
- `contains(p)` returns whether point `p` lies inside or on the circle.
- `on_edge(p)` returns whether point `p` lies on the circle boundary.
- `incircle(a, b, c)` returns the circle inscribed in triangle `abc`.
- `in_circumcircle(a, b, c, d)` returns 1 if `d` lies strictly inside the circle through `a`, `b`, and `c`, 0 if it lies on the circle, or -1 if it lies outside. This has exact reasoning for integral points by evaluating the determinant directly without constructing a `Circle` or taking square roots. Points `a`, `b`, and `c` must not be collinear. For integer inputs, the sign test is exact provided the chosen intermediate type does not overflow.

The floating-point constructors and predicates use an absolute tolerance `EPS`.

Overflow warning: the determinant in `in_circumcircle()` grows like the fourth power of the coordinate magnitude. For large integer coordinates, use a point or intermediate type wide enough for that determinant.

Time Complexity: $O(1)$ per call to the constructors and all other operations.

Space Complexity:

- $O(1)$ for storage of the circle.
- $O(1)$ auxiliary for all operations.

```
#include <cmath>
#include <cstdint>
#include <ostream>
#include <stdexcept>
#include <type_traits>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U>
bool LE(T a, U b) {
    return !LT(b, a);
}

// SFINAE guard: valid only when Pt exposes numeric .x/.y members. This keeps the templated point
// constructors from hijacking calls like Circle(double, double, double).
template<typename Pt>
using if_point = decltype(std::declval<const Pt &>().x, std::declval<const Pt &>().y, void());

struct Circle {
    double h, k, r;

    Circle() : h(0), k(0), r(0) {}
    explicit Circle(double r) : h(0), k(0), r(fabs(r)) {}
    Circle(double h, double k, double r) : h(h), k(k), r(fabs(r)) {}

    template<typename Pt, typename = if_point<Pt>>
    Circle(const Pt &o, double r) : h(o.x), k(o.y), r(fabs(r)) {}

    template<typename Pt, typename = if_point<Pt>>
    Circle(const Pt &a, const Pt &b) {
        h = (a.x + b.x) / 2.0;
        k = (a.y + b.y) / 2.0;
        r = hypot(a.x - h, a.y - k);
    }

    template<typename Pt, typename = if_point<Pt>>
    Circle(const Pt &a, const Pt &b, const Pt &c) {
        double ax = a.x, ay = a.y, bx = b.x, by = b.y, cx = c.x, cy = c.y;
        double asq = ax * ax + ay * ay, bsq = bx * bx + by * by, csq = cx * cx + cy * cy;
        double d = 2.0 * (ax * (by - cy) + bx * (cy - ay) + cx * (ay - by));
        if (EQ(d, 0)) {
            throw std::runtime_error("No circumcircle from collinear points.");
        }
        h = (asq * (by - cy) + bsq * (cy - ay) + csq * (ay - by)) / d;
```

```

    k = (asq * (cx - bx) + bsq * (ax - cx) + csq * (bx - ax)) / d;
    r = hypot(ax - h, ay - k);
}

template<typename Pt, typename = if_point<Pt>>
Circle(const Pt &a, const Pt &b, double r) : r(fabs(r)) {
    if (EQ(this->r, 0) && EQ(a.x, b.x) && EQ(a.y, b.y)) {
        h = a.x;
        k = a.y;
        return;
    }
    double d = hypot(b.x - a.x, b.y - a.y);
    if (EQ(d, 0)) {
        throw std::runtime_error("Identical points, infinite circles.");
    }
    if (LT(this->r * 2.0, d)) {
        throw std::runtime_error("Points too far away to make Circle.");
    }
    double z = this->r * this->r - d * d / 4.0;
    double v = sqrt(z < 0 ? 0 : z) / d;
    double mx = (a.x + b.x) / 2.0, my = (a.y + b.y) / 2.0;
    h = mx + v * (a.y - b.y);
    k = my + v * (b.x - a.x);
    // The other answer is (h, k) = (mx - v*(a.y - b.y), my - v*(b.x - a.x)).
}

bool operator==(const Circle &c) const { return h == c.h && k == c.k && r == c.r; }
bool operator!=(const Circle &c) const { return !(*this == c); }

friend bool EQ(const Circle &a, const Circle &b) {
    return EQ(a.h, b.h) && EQ(a.k, b.k) && EQ(a.r, b.r);
}

template<typename Pt>
bool contains(const Pt &p) const {
    return LE(hypot(p.x - h, p.y - k), r);
}

template<typename Pt>
bool on_edge(const Pt &p) const {
    return EQ(hypot(p.x - h, p.y - k), r);
}

friend std::ostream &operator<<(std::ostream &out, const Circle &c) {
    return out << std::showpos << "(x" << -(fabs(c.h) < EPS ? 0 : c.h) << ")^2+"
        << "(y" << -(fabs(c.k) < EPS ? 0 : c.k) << ")^2" << std::noshowpos << "="
        << (fabs(c.r) < EPS ? 0 : c.r * c.r);
}

};

template<typename Pt>
Circle incircle(const Pt &a, const Pt &b, const Pt &c) {
    double al = hypot(b.x - c.x, b.y - c.y);
    double bl = hypot(a.x - c.x, a.y - c.y);
    double cl = hypot(a.x - b.x, a.y - b.y);
    double l = al + bl + cl;
    double px = a.x - c.x, py = a.y - c.y, qx = b.x - c.x, qy = b.y - c.y;
    return EQ(l, 0) ? Circle(a.x, a.y, 0)
        : Circle(
            (al * a.x + bl * b.x + cl * c.x) / l, (al * a.y + bl * b.y + cl * c.y) / l,
            fabs(px * qy - py * qx) / l
        );
}

template<typename Pt>
int in_circumcircle(const Pt &a, const Pt &b, const Pt &c, const Pt &d) {
    using T = decltype(a.x + a.x);
    using W = std::conditional_t<

```

```

    std::is_floating_point<T>::value, long double,
    std::conditional_t<std::is_integral<T>::value, int64_t, T>>;
W adx = (W)a.x - d.x, ady = (W)a.y - d.y;
W bdx = (W)b.x - d.x, bdy = (W)b.y - d.y;
W cdx = (W)c.x - d.x, cdy = (W)c.y - d.y;
// Overflow warning!
W det = (adx * adx + ady * ady) * (bdx * cdy - cdx * bdy) +
        (bdx * bdx + bdy * bdy) * (cdx * ady - adx * cdy) +
        (cdx * cdx + cdy * cdy) * (adx * bdy - bdx * ady);
W orient = ((W)b.x - a.x) * ((W)c.y - a.y) - ((W)b.y - a.y) * ((W)c.x - a.x);
W val = W(0) < orient ? det : -det;
return EQ(val, W(0)) ? 0 : (W(0) < val ? 1 : -1);
}

```

Example Usage

```

#include <cassert>

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
};

int main() {
    Circle c(-2, 5, sqrt(10));
    assert(EQ(c, Circle(Point(-2, 5), sqrt(10))));
    assert(EQ(c, Circle(Point(1, 6), Point(-5, 4))));
    assert(EQ(c, Circle(Point(-3, 2), Point(-3, 8), Point(-1, 8))));
    assert(EQ(Circle(Point(0, 0), Point(0, 2), -1), Circle(0, 1, 1)));
    assert(EQ(c, incircle(Point(-12, 5), Point(3, 0), Point(0, 9))));
    assert(c.contains(Point(-2, 8)) && !c.contains(Point(-2, 9)));
    assert(c.on_edge(Point(-1, 2)) && !c.on_edge(Point(-1.01, 2)));

    // Integer-coordinate points are accepted; the Circle is computed in double.
    assert(EQ(c, Circle(PointI(1, 6), PointI(-5, 4))));
    assert(EQ(c, Circle(PointI(-3, 2), PointI(-3, 8), PointI(-1, 8))));
    assert(EQ(c, incircle(PointI(-12, 5), PointI(3, 0), PointI(0, 9))));
    assert(c.contains(PointI(-2, 8)) && !c.contains(PointI(-2, 9)));
    assert(c.on_edge(PointI(-1, 2)));

    // Exact integer in-circle predicate: unit circle through (1,0), (0,1), (-1,0).
    assert(in_circumcircle(PointI(1, 0), PointI(0, 1), PointI(-1, 0), PointI(0, 0)) == 1); // inside
    assert(
        in_circumcircle(PointI(1, 0), PointI(0, 1), PointI(-1, 0), PointI(2, 0)) == -1
    ); // outside
    assert(
        in_circumcircle(PointI(1, 0), PointI(0, 1), PointI(-1, 0), PointI(0, -1)) == 0
    ); // on edge
    // Orientation-independent: reversing a, b, c gives the same answer.
    assert(in_circumcircle(PointI(-1, 0), PointI(0, 1), PointI(1, 0), PointI(0, 0)) == 1);
    return 0;
}

```

7.1.4 Triangle

Common triangle calculations in two dimensions. The functions are templated on the point type `Pt`, which should accept `Point`, `PointD`, or `PointI` from 7.1.1 or any struct with numeric `.x` and `.y` fields. `triangle_area()` returns

`double` regardless of input type (the area is fractional via the `/2`). `same_side()` and `point_in_triangle()` do all arithmetic in the point's own coordinate type and are exact for integer points: they reduce each cross product to a sign before combining, so no precision is lost and cross products are never multiplied together.

- `triangle_area(a, b, c)` returns the area of the triangle with vertices `a`, `b`, and `c`.
- `triangle_area_sides(s1, s2, s3)` returns the area of a triangle with side lengths `s1`, `s2`, and `s3`. The given lengths must be nonnegative and form a valid triangle.
- `triangle_area_medians(m1, m2, m3)` returns the area of a triangle with medians of lengths `m1`, `m2`, and `m3`. The median of a triangle is a line segment joining a vertex to the midpoint of the opposing edge.
- `triangle_area_altitudes(h1, h2, h3)` returns the area of a triangle with altitudes `h1`, `h2`, and `h3`. An altitude of a triangle is the shortest line between a vertex and the infinite line that is extended from its opposite edge.
- `same_side(p1, p2, a, b, edge_is_inside = true)` returns whether points `p1` and `p2` lie on the same side of the line containing points `a` and `b`. If one or both points lie exactly on the line, then the result will depend on `edge_is_inside`.
- `point_in_triangle(p, a, b, c, edge_is_inside = true)` returns whether point `p` lies within the triangle with vertices `a`, `b`, and `c`. If the point lies on or close to an edge (by roughly `EPS`), then the result will depend on `edge_is_inside`.

Use an integral point type for integer coordinates, and a floating-point type for fractional coordinates.

Overflow warning: each individual cross product is still on the order of the squared coordinate magnitude, so for integer point types use a 64-bit coordinate type (e.g. `PointL` from 7.1.1) once coordinates reach a few tens of thousands.

Time Complexity: $O(1)$ per operation.

Space Complexity: $O(1)$ auxiliary for all operations.

```
#include <cmath>
#include <type_traits>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename Pt>
double triangle_area(const Pt &a, const Pt &b, const Pt &c) {
    double acx = static_cast<double>(a.x) - c.x, acy = static_cast<double>(a.y) - c.y;
    double bcx = static_cast<double>(b.x) - c.x, bcy = static_cast<double>(b.y) - c.y;
    return fabs(acx * bcy - acy * bcx) / 2.0;
}

double triangle_area_sides(double s1, double s2, double s3) {
    double s = (s1 + s2 + s3) / 2.0;
    return sqrt(s * (s - s1) * (s - s2) * (s - s3));
}

double triangle_area_medians(double m1, double m2, double m3) {
    return 4.0 * triangle_area_sides(m1, m2, m3) / 3.0;
}

double triangle_area_altitudes(double h1, double h2, double h3) {
    if (EQ(h1, 0) || EQ(h2, 0) || EQ(h3, 0)) {
        return 0;
    }
}
```

```

}
double x = h1 * h1, y = h2 * h2, z = h3 * h3;
double v = 2.0 / (x * y) + 2.0 / (x * z) + 2.0 / (y * z);
return 1.0 / sqrt(v - 1.0 / (x * x) - 1.0 / (y * y) - 1.0 / (z * z));
}

template<typename Pt>
bool same_side(const Pt &p1, const Pt &p2, const Pt &a, const Pt &b, bool edge_is_inside = true) {
    // Cross products in the point's native coordinate type (exact for integer points).
    auto c1 = (b.x - a.x) * (p1.y - a.y) - (b.y - a.y) * (p1.x - a.x);
    auto c2 = (b.x - a.x) * (p2.y - a.y) - (b.y - a.y) * (p2.x - a.x);
    // Compare the signs, not the product c1 * c2, to avoid integer overflow.
    int s1 = LT(0, c1) ? 1 : (LT(c1, 0) ? -1 : 0);
    int s2 = LT(0, c2) ? 1 : (LT(c2, 0) ? -1 : 0);
    return edge_is_inside ? (s1 * s2 >= 0) : (s1 * s2 > 0);
}

template<typename Pt>
bool point_in_triangle(
    const Pt &p, const Pt &a, const Pt &b, const Pt &c, bool edge_is_inside = true
) {
    return same_side(p, a, b, c, edge_is_inside) && same_side(p, b, a, c, edge_is_inside) &&
        same_side(p, c, a, b, edge_is_inside);
}

```

Example Usage

```

#include <cassert>

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
};

int main() {
    assert(EQ(6, triangle_area(Point(0, -1), Point(4, -1), Point(0, -4))));
    assert(EQ(6, triangle_area_sides(3, 4, 5)));
    assert(EQ(6, triangle_area_medians(3.605551275, 2.5, 4.272001873)));
    assert(EQ(6, triangle_area_altitudes(3, 4, 2.4)));

    // Fractional coordinates require a floating-point point type.
    assert(point_in_triangle(Point(0, 0), Point(-1, 0), Point(0, -2), Point(4, 0)));
    assert(!point_in_triangle(Point(0, 1), Point(-1, 0), Point(0, -2), Point(4, 0)));
    assert(point_in_triangle(Point(-2.44, 0.82), Point(-1, 0), Point(-3, 1), Point(4, 0)));
    assert(!point_in_triangle(Point(-2.44, 0.7), Point(-1, 0), Point(-3, 1), Point(4, 0)));

    // Integer coordinates: same_side / point_in_triangle are computed exactly in int.
    assert(EQ(8, triangle_area(PointI(0, 0), PointI(4, 0), PointI(0, 4))));
    assert(point_in_triangle(PointI(1, 1), PointI(0, 0), PointI(4, 0), PointI(0, 4)));
    assert(!point_in_triangle(PointI(5, 5), PointI(0, 0), PointI(4, 0), PointI(0, 4)));
    assert(point_in_triangle(PointI(2, 2), PointI(0, 0), PointI(4, 0), PointI(0, 4))); // on edge
    assert(!point_in_triangle(PointI(2, 2), PointI(0, 0), PointI(4, 0), PointI(0, 4), false));
    return 0;
}

```

7.1.5 Rectangle

Common axis-aligned rectangle calculations in two dimensions. The functions are templated on the point type `Pt`, which should accept `Point`, `PointD`, or `PointI` from 7.1.1, or any struct with numeric `.x` and `.y` fields.

- `point_in_rectangle(p, v, w, h, edge_is_inside = true)` returns whether point `p` lies inside the axis-aligned rectangle whose opposite corners are `v` and `(v.x + w, v.y + h)`. Negative widths and heights are supported. The `edge_is_inside` flag controls whether boundary points count.
- `point_in_rectangle(p, a, b, edge_is_inside = true)` returns whether point `p` lies inside the axis-aligned rectangle with opposite corners `a` and `b`.
- `rectangle_intersection(a1, b1, a2, b2, &p, &q, edge_is_inside = true)` computes the intersection of two axis-aligned rectangles, where `a1/b1` and `a2/b2` are opposite-corner pairs. Returns `-1` if the rectangles are disjoint, `0` if they partially intersect, `1` if the first rectangle is completely inside the second, and `2` if the second rectangle is completely inside the first. If an intersection exists, its lower-left and upper-right corners are stored in `p` and `q` when those pointers are non-null. The `edge_is_inside` flag controls whether boundary-only contact counts as intersecting.

For integer-coordinate inputs, comparisons are exact as long as intermediate coordinate additions, subtractions, and min/max values do not overflow. Floating-point inputs use EPS-based comparisons.

Time Complexity: $O(1)$ per operation.

Space Complexity: $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cmath>
#include <type_traits>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U> bool GT(T a, U b) { return LT(b, a); }
template<typename T, typename U> bool LE(T a, U b) { return !LT(b, a); }
template<typename T, typename U> bool GE(T a, U b) { return !LT(a, b); }

template<typename Pt, typename T>
bool point_in_rectangle(
    const Pt &p, const Pt &v, const T &w, const T &h, const bool edge_is_inside = true
) {
    if (w < 0) {
        return point_in_rectangle(p, Pt(v.x + w, v.y), -w, h, edge_is_inside);
    }
    if (h < 0) {
        return point_in_rectangle(p, Pt(v.x, v.y + h), w, -h, edge_is_inside);
    }
    return edge_is_inside ? (GE(p.x, v.x) && LE(p.x, v.x + w) && GE(p.y, v.y) && LE(p.y, v.y + h))
        : (GT(p.x, v.x) && LT(p.x, v.x + w) && GT(p.y, v.y) && LT(p.y, v.y + h));
}

template<typename Pt>
bool point_in_rectangle(const Pt &p, const Pt &a, const Pt &b, const bool edge_is_inside = true) {
    auto xl = std::min(a.x, b.x), yl = std::min(a.y, b.y);
    auto xh = std::max(a.x, b.x), yh = std::max(a.y, b.y);
    return point_in_rectangle(p, Pt(xl, yl), xh - xl, yh - yl, edge_is_inside);
}
```

```

}

template<typename Pt>
int rectangle_intersection(
    const Pt &a1, const Pt &b1, const Pt &a2, const Pt &b2, Pt *p = nullptr, Pt *q = nullptr,
    const bool edge_is_inside = true
) {
    // Normalize each rectangle to [x1, xh] x [y1, yh] with coordinate-wise min/max only.
    // (Using auto keeps the native coordinate type, so integer rectangles stay exact.)
    auto x11 = std::min(a1.x, b1.x), xh1 = std::max(a1.x, b1.x);
    auto y11 = std::min(a1.y, b1.y), yh1 = std::max(a1.y, b1.y);
    auto x12 = std::min(a2.x, b2.x), xh2 = std::max(a2.x, b2.x);
    auto y12 = std::min(a2.y, b2.y), yh2 = std::max(a2.y, b2.y);
    // The intersection is the overlap of the two coordinate intervals.
    auto ixl = std::max(x11, x12), ixh = std::min(xh1, xh2);
    auto iyl = std::max(y11, y12), iyh = std::min(yh1, yh2);
    bool disjoint = edge_is_inside ? (GT(ixl, ixh) || GT(iyl, iyh)) : (GE(ixl, ixh) || GE(iyl, iyh));
    if (disjoint) {
        return -1; // Completely disjoint.
    }
    // Opposing corners of the overlap: bottom-left in p, top-right in q.
    if (p != nullptr) {
        *p = Pt(ixl, iyl);
    }
    if (q != nullptr) {
        *q = Pt(ixh, iyh);
    }
    // The overlap equals a rectangle exactly when that rectangle is contained in the other.
    if (EQ(ixl, x11) && EQ(ixh, xh1) && EQ(iyl, y11) && EQ(iyh, yh1)) {
        return 1; // Rectangle 1 completely inside 2.
    }
    if (EQ(ixl, x12) && EQ(ixh, xh2) && EQ(iyl, y12) && EQ(iyh, yh2)) {
        return 2; // Rectangle 2 completely inside 1.
    }
    return 0; // Partial intersection.
}

```

Example Usage

```

#include <cassert>

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
};

int main() {
    assert(point_in_rectangle(Point(0, -1), Point(0, -3), 3, 2));
    assert(!point_in_rectangle(Point(0, -1), Point(0, -3), 3, 2, false));
    assert(point_in_rectangle(Point(2, -2), Point(3, -3), -3, 2));
    assert(!point_in_rectangle(Point(0, 0), Point(3, -1), -3, -2));
    assert(point_in_rectangle(Point(2, -2), Point(3, -3), Point(0, -1)));
    assert(!point_in_rectangle(Point(0, -1), Point(3, -3), Point(0, -1), false));
    assert(!point_in_rectangle(Point(-1, -2), Point(3, -3), Point(0, -1)));

    Point p, q;
    // Output coordinates come directly from input coordinates via min/max, so exact equality is safe.
    assert(-1 == rectangle_intersection(Point(0, 0), Point(1, 1), Point(2, 2), Point(3, 3)));
    assert(0 == rectangle_intersection(Point(1, 1), Point(7, 7), Point(5, 5), Point(0, 0), &p, &q));
    assert(p == Point(1, 1) && q == Point(5, 5));
    assert(1 == rectangle_intersection(Point(1, 1), Point(0, 0), Point(0, 0), Point(1, 10), &p, &q));
    assert(p == Point(0, 0) && q == Point(1, 1));
    assert(2 == rectangle_intersection(Point(0, 5), Point(5, 7), Point(1, 6), Point(2, 5), &p, &q));
    assert(p == Point(1, 5) && q == Point(2, 6));
    assert(

```

```

-1 ==
rectangle_intersection(Point(0, 0), Point(1, 1), Point(1, 1), Point(2, 2), &p, &q, false)
);
return 0;
}

```

7.2 Angles, Distances, and Intersections

7.2.1 Angles

Angle calculations in two dimensions. All operations are inherently floating-point (`atan2`, `acos`, trigonometry), so these functions use the local double-coordinate `PointD` type. Convert integer points to `PointD` when asking for angles. The constants `DEG` and `RAD` may be used as multipliers to convert between degrees and radians. For example, if `t` is a value in degrees, then `t * DEG` is the equivalent angle in radians; if `t` is in radians, then `t * RAD` is the equivalent angle in degrees.

- `reduce_deg(t)` takes an angle `t` degrees and returns an equivalent angle in the range $[0, 360)$ degrees (e.g. -630 becomes 90).
- `reduce_rad(t)` takes an angle `t` radians and returns an equivalent angle in the range $[0, 2\pi)$ radians (e.g. 720.5 becomes 0.5).
- `polar_point(r, t)` returns a two-dimensional Cartesian point given radius `r` and angle `t` radians in polar coordinates (see `std::polar()`).
- `polar_angle(p)` returns the angle in radians of the line segment from $(0, 0)$ to point `p`, relative counterclockwise to the positive x -axis.
- `angle(a, o, b)` returns the smallest angle in radians formed by the points `a`, `o`, `b` with vertex at point `o`.
- `angle_between(a, b)` returns the directed counter-clockwise angle in $[0, 2\pi)$ from vector `a` to vector `b`, treating both points as vectors from the origin.
- `angle_between(a1, b1, a2, b2)` returns the smaller angle in radians between two lines $a_1x + b_1y + c_1 = 0$ and $a_2x + b_2y + c_2 = 0$, limited to $[0, \pi/2]$.
- `cross(a, b, o = Pt(0, 0))` returns the signed z -component of the three-dimensional cross product between points `a` and `b`, where the z -coordinates are implicitly zero and the origin is shifted to point `o`. This is also double the signed area of the triangle from these three points.
- `turn(a, o, b)` returns 1 if the path `a` $\rightarrow$ `o` $\rightarrow$ `b` forms a left turn on the plane, 0 if the path forms a straight line segment, or -1 if it forms a right turn.

The points passed to `angle()` must differ from its vertex `o`, vectors passed to `angle_between()` must be nonzero, and coefficient pairs passed to its line overload must represent valid lines.

Overflow warning: `cross()` and `turn()` preserve the point's coordinate type, so their products may overflow for integral points. Use a 64-bit coordinate type for large integer coordinates.

Time Complexity: $O(1)$ per operation.

Space Complexity: $O(1)$ auxiliary for all operations.

```

#include <algorithm>
#include <cmath>
#include <type_traits>

const double EPS = 1e-9;
const double PI = acos(-1.0);
const double DEG = PI / 180;
const double RAD = 180 / PI;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;

```



```

    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

double reduce_deg(double t) {
    if (t < -360) {
        return reduce_deg(fmod(t, 360));
    }
    if (t < 0) {
        return t + 360;
    }
    return (t >= 360) ? fmod(t, 360) : t;
}

double reduce_rad(double t) {
    if (t < -2 * PI) {
        return reduce_rad(fmod(t, 2 * PI));
    }
    if (t < 0) {
        return t + 2 * PI;
    }
    return (t >= 2 * PI) ? fmod(t, 2 * PI) : t;
}

template<typename OutPt>
OutPt polar_point(double r, double t) {
    return OutPt(r * cos(t), r * sin(t));
}

template<typename Pt>
double polar_angle(const Pt &p) {
    double t = atan2(static_cast<double>(p.y), static_cast<double>(p.x));
    return (t < 0) ? (t + 2 * PI) : t;
}

template<typename Pt>
double angle(const Pt &a, const Pt &o, const Pt &b) {
    double ux = o.x - a.x, uy = o.y - a.y;
    double vx = o.x - b.x, vy = o.y - b.y;
    double cosine = (ux * vx + uy * vy) / (std::hypot(ux, uy) * std::hypot(vx, vy));
    return acos(std::max(-1.0, std::min(1.0, cosine)));
}

template<typename Pt>
double angle_between(const Pt &a, const Pt &b) {
    double t = atan2(a.x * b.y - a.y * b.x, a.x * b.x + a.y * b.y);
    return (t < 0) ? (t + 2 * PI) : t;
}

double angle_between(const double &a1, const double &b1, const double &a2, const double &b2) {
    double t = atan2(a1 * b2 - a2 * b1, a1 * a2 + b1 * b2);
    if (t < 0) {
        t += PI;
    }
    return LT(PI / 2, t) ? (PI - t) : t;
}

template<typename Pt>
auto cross(const Pt &a, const Pt &b, const Pt &o = Pt(0, 0)) {
    return (a.x - o.x) * (b.y - o.y) - (a.y - o.y) * (b.x - o.x);
}

template<typename Pt>

```

```
int turn(const Pt &a, const Pt &o, const Pt &b) {
    auto c = cross(a, b, o);
    return LT(c, 0) ? 1 : (LT(0, c) ? -1 : 0);
}
```

Example Usage

```
struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
};

#include <cassert>

int main() {
    assert(EQ(123, reduce_deg(-8 * 360 + 123)));
    assert(EQ(1.2345, reduce_rad(2 * PI * 8 + 1.2345)));
    assert(EQ(polar_point<Point>(4, PI), Point(-4, 0)));
    assert(EQ(polar_point<Point>(4, -PI / 2), Point(0, -4)));
    assert(EQ(45, polar_angle(Point(5, 5)) * RAD));
    assert(EQ(135 * DEG, polar_angle(Point(-4, 4))));
    assert(EQ(90 * DEG, angle(Point(5, 0), Point(0, 5), Point(-5, 0))));
    assert(EQ(225 * DEG, angle_between(Point(0, 5), Point(5, -5))));
    assert(-1 == cross(Point(0, 1), Point(1, 0), Point(0, 0)));
    assert(-1 == turn(Point(0, 1), Point(0, 0), Point(-5, -5)));
    return 0;
}
```

7.2.2 Distances

Distance calculations in two dimensions for points, lines, and line segments. The functions are templated on the point type `Pt`: pass any struct with numeric `.x` and `.y` fields (for example `PointD`/`PointI` from 7.1.1, or the local example struct). `sqdist` preserves the coordinate type and is exact for integer points. Functions that return an actual distance use `double` for the final division/square root. `closest_point` returns a point of the same type as its inputs, so use a floating-point `Pt` there for a meaningful (non-truncated) result.

- `dist(a, b)` and `sqdist(a, b)` respectively return the distance and squared distance between points `a` and `b`.
- `line_dist(p, a, b, c)` returns the distance from point `p` to the line $ax + by + c = 0$.
- `line_dist(p, a, b)` returns the distance from point `p` to the infinite line containing points `a` and `b`. If `a = b`, the distance from `p` to `a` is returned.
- `line_dist(a1, b1, c1, a2, b2, c2)` returns the distance between two lines.
- `seg_dist(p, a, b)` returns the distance from point `p` to the line segment `a-b`.
- `seg_dist(a, b, c, d)` returns the minimum distance between line segments `a-b` and `c-d`, or 0 if the segments touch or intersect.
- `closest_point(a, b, p)` returns the point on segment `a-b` closest to point `p`.

Coefficient triples passed to `line_dist()` must represent valid lines. Use a floating-point point type with `closest_point()` when the closest point may have fractional coordinates.

Overflow warning: squared distances, dot products, and cross products are formed in the point's coordinate arithmetic type. For integer point types, use a 64-bit coordinate type when coordinates may exceed a few tens of thousands.

Time Complexity: $O(1)$ per operation.

Space Complexity: $O(1)$ auxiliary for all operations.

```
#include <algorithm>
```

```

#include <cmath>
#include <type_traits>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U> bool LE(T a, U b) { return !LT(b, a); }
template<typename T, typename U> bool GE(T a, U b) { return !LT(a, b); }

// sqdist returns the coordinate type (exact for integer points).
template<typename Pt>
auto sqdist(const Pt &a, const Pt &b) {
    auto dx = b.x - a.x, dy = b.y - a.y;
    return dx * dx + dy * dy; // Overflow warning!
}

template<typename Pt>
double dist(const Pt &a, const Pt &b) {
    return sqrt(static_cast<double>(sqdist(a, b)));
}

template<typename Pt, typename T>
double line_dist(const Pt &p, const T &a, const T &b, const T &c) {
    return fabs(static_cast<double>(a) * p.x + static_cast<double>(b) * p.y + c) /
        sqrt(static_cast<double>(a) * a + static_cast<double>(b) * b);
}

template<typename Pt>
double line_dist(const Pt &p, const Pt &a, const Pt &b) {
    if (EQ(a.x, b.x) && EQ(a.y, b.y)) {
        return dist(p, a);
    }
    auto n = sqdist(a, b);
    auto d = (p.x - a.x) * (b.x - a.x) + (p.y - a.y) * (b.y - a.y); // Overflow warning!
    double u = static_cast<double>(d) / n;
    double dx = a.x + u * (b.x - a.x) - p.x, dy = a.y + u * (b.y - a.y) - p.y;
    return sqrt(dx * dx + dy * dy);
}

template<typename T>
double line_dist(const T &a1, const T &b1, const T &c1, const T &a2, const T &b2, const T &c2) {
    if (EQ(static_cast<double>(a1) * b2, static_cast<double>(a2) * b1)) {
        double factor = EQ(b1, 0) ? (static_cast<double>(a1) / a2) : (static_cast<double>(b1) / b2);
        return EQ(c1, c2 * factor)
            ? 0
            : fabs(c2 * factor - c1) /
                sqrt(static_cast<double>(a1) * a1 + static_cast<double>(b1) * b1);
    }
    return 0;
}

template<typename Pt>
double seg_dist(const Pt &p, const Pt &a, const Pt &b) {
    if (EQ(a.x, b.x) && EQ(a.y, b.y)) {
        return dist(p, a);
    }
    auto n = sqdist(a, b);

```

```

auto d = (p.x - a.x) * (b.x - a.x) + (p.y - a.y) * (b.y - a.y); // Overflow warning!
if (LE(d, 0) || EQ(n, 0)) {
    return dist(p, a);
}
if (GE(d, n)) {
    return dist(p, b);
}
double t = static_cast<double>(d) / n;
double dx = a.x + t * (b.x - a.x) - p.x, dy = a.y + t * (b.y - a.y) - p.y;
return sqrt(dx * dx + dy * dy);
}

template<typename Pt>
double seg_dist(const Pt &a, const Pt &b, const Pt &c, const Pt &d) {
    if (EQ(a.x, b.x) && EQ(a.y, b.y)) {
        return seg_dist(a, c, d);
    }
    if (EQ(c.x, d.x) && EQ(c.y, d.y)) {
        return seg_dist(c, a, b);
    }
    auto ab_x = b.x - a.x, ab_y = b.y - a.y;
    auto ac_x = c.x - a.x, ac_y = c.y - a.y;
    auto cd_x = d.x - c.x, cd_y = d.y - c.y;
    auto c1 = ab_x * cd_y - ab_y * cd_x; // Overflow warning!
    auto c2 = ac_x * ab_y - ac_y * ab_x;
    if (EQ(c1, 0) && EQ(c2, 0)) {
        auto less = [](const Pt &u, const Pt &v) { return u.x != v.x ? u.x < v.x : u.y < v.y; };
        auto minp = [&](const Pt &u, const Pt &v) -> const Pt & { return less(v, u) ? v : u; };
        auto maxp = [&](const Pt &u, const Pt &v) -> const Pt & { return less(u, v) ? v : u; };
        const Pt &res1 = maxp(minp(a, b), minp(c, d));
        const Pt &res2 = minp(maxp(a, b), maxp(c, d));
        if (!less(res2, res1)) {
            return 0;
        }
    }
    else if (!EQ(c1, 0)) {
        auto t_num = ac_x * cd_y - ac_y * cd_x;
        bool c1_pos = c1 > 0;
        bool t_between_01 = c1_pos ? (LE(0, t_num) && LE(t_num, c1)) : (LE(t_num, 0) && LE(c1, t_num));
        bool u_between_01 = c1_pos ? (LE(0, c2) && LE(c2, c1)) : (LE(c2, 0) && LE(c1, c2));
        if (t_between_01 && u_between_01) {
            return 0;
        }
    }
    return std::min(
        std::min(seg_dist(a, c, d), seg_dist(b, c, d)), std::min(seg_dist(c, a, b), seg_dist(d, a, b))
    );
}

template<typename Pt>
Pt closest_point(const Pt &a, const Pt &b, const Pt &p) {
    Pt res{};
    if (EQ(a.x, b.x) && EQ(a.y, b.y)) {
        res.x = a.x;
        res.y = a.y;
        return res;
    }
    auto n = sqdist(a, b);
    auto d = (p.x - a.x) * (b.x - a.x) + (p.y - a.y) * (b.y - a.y); // Overflow warning!
    double t = static_cast<double>(d) / n;
    if (t <= 0) {
        res.x = a.x;
        res.y = a.y;
    }
    else if (t >= 1) {
        res.x = b.x;
        res.y = b.y;
    }
    else {
        res.x = a.x + t * (b.x - a.x);
        res.y = a.y + t * (b.y - a.y);
    }
}

```

```

    }
    return res;
}

```

Example Usage

```

#include <cassert>

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
};

int main() {
    assert(EQ(5, dist(Point(-1, -1), Point(2, 3))));
    assert(EQ(25, sqdist(Point(-1, -1), Point(2, 3))));
    assert(EQ(1.2, line_dist(Point(2, 1), -4, 3, -1)));
    assert(EQ(0.8, line_dist(Point(3, 3), Point(-1, -1), Point(2, 3))));
    assert(EQ(1.2, line_dist(Point(2, 1), Point(-1, -1), Point(2, 3))));
    assert(EQ(0, line_dist(-4, 3, -1, 8, 6, 2)));
    assert(EQ(0.8, line_dist(-4, 3, -1, -8, 6, -10)));
    assert(EQ(1.0, seg_dist(Point(3, 3), Point(-1, -1), Point(2, 3))));
    assert(EQ(1.2, seg_dist(Point(2, 1), Point(-1, -1), Point(2, 3))));
    assert(EQ(0, seg_dist(Point(0, 2), Point(3, 3), Point(-1, -1), Point(2, 3))));
    assert(EQ(0.6, seg_dist(Point(-1, 0), Point(-2, 2), Point(-1, -1), Point(2, 3))));

    // Integer points: sqdist is exact, dist returns double.
    PointI ia{-1, -1}, ib{2, 3};
    assert(sqdist(ia, ib) == 25);    // int result
    assert(EQ(dist(ia, ib), 5.0));   // double result
    return 0;
}

```

7.2.3 Line Intersection

Intersection calculations in two dimensions for straight lines and line segments. The functions are templated on the point type `Pt`: pass any struct with numeric `x` and `y` fields (for example `PointD`/`PointI` from 7.1.1, or the local example struct) for which `operator<` orders points lexicographically using exact coordinate comparisons. Point outputs are written through a caller-supplied pointer whose type is deduced, so the output point type may differ from the input – e.g. integer endpoints with a floating-point output point.

- `line_intersection(a1, b1, c1, a2, b2, c2, &p)` intersects lines $a_1x + b_1y + c_1 = 0$ and $a_2x + b_2y + c_2 = 0$, returning `-1` (parallel), `0` (one point, stored into `p`), or `1` (identical).
- `line_intersection(p1, p2, p3, p4, &p)` intersects the infinite lines determined by points `p1`, `p2`, `p3`, and `p4`, returning `-1` (parallel), `0` (one point, stored into `p`), or `1` (identical).
- `seg_intersection(a, b, c, d, &p, &q, touch_is_intersect = true)` intersects segments `a-b` and `c-d`, returning `-1` (none), `0` (one point), or `1` (overlapping segment). The `touch_is_intersect` flag controls whether segments that meet only at an endpoint count as intersecting. If the intersection is a point (and `p` is not `nullptr`), it will be stored into `p`. If the intersection is an overlapping segment (and `p` and `q` are not `nullptr`) then their endpoints will be stored into pointers `p` and `q`, which must reference a floating-point point type.
- `seg_intersection(a, b, c, d, touch_is_intersect = true)` is a simplified, detection-only version of the above. Both versions are exact for detection if the input point type is integral.
- `closest_point(a, b, c, p)` returns the closest point on line $ax + by + c = 0$ to point `p`.

Coefficient triples must represent valid lines, and each point pair used to define an infinite line must contain two distinct points.

Overflow warning: the exact integer paths multiply coordinate differences (cross products and squared lengths), which grow like the squared coordinate magnitude. With 32-bit `int` coordinates these overflow once coordinates exceed roughly a few tens of thousands, so for larger integer inputs use a point type with 64-bit (`int64_t`) coordinates.

Time Complexity: $O(1)$ per operation.

Space Complexity: $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cmath>
#include <type_traits>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U>
bool LE(T a, U b) {
    return !LT(b, a);
}

template<typename OutPt>
int line_intersection(double a1, double b1, double c1, double a2, double b2, double c2, OutPt *p) {
    // Cramer's rule for a1*x + b1*y = -c1, a2*x + b2*y = -c2.
    double det = a1 * b2 - a2 * b1;
    if (EQ(det, 0)) { // Parallel (lines need not be given in normalized form).
        // Identical iff the other two 2x2 determinants also vanish.
        return (EQ(a1 * c2 - a2 * c1, 0) && EQ(b1 * c2 - b2 * c1, 0)) ? 1 : -1;
    }
    p->x = (b1 * c2 - b2 * c1) / det;
    if (!EQ(b1, 0)) {
        p->y = -(a1 * p->x + c1) / b1;
    } else {
        p->y = -(a2 * p->x + c2) / b2;
    }
    return 0;
}

template<typename Pt, typename OutPt>
int line_intersection(const Pt &p1, const Pt &p2, const Pt &p3, const Pt &p4, OutPt *p) {
    auto a1 = p2.y - p1.y, b1 = p1.x - p2.x;
    auto c1 = -(p1.x * p2.y - p2.x * p1.y);
    auto a2 = p4.y - p3.y, b2 = p3.x - p4.x;
    auto c2 = -(p3.x * p4.y - p4.x * p3.y);
    auto x = -(c1 * b2 - c2 * b1), y = -(a1 * c2 - a2 * c1);
    auto det = a1 * b2 - a2 * b1;
    if (EQ(det, 0)) {
        return (EQ(x, 0) && EQ(y, 0)) ? 1 : -1;
    }
    p->x = static_cast<double>(x) / det;
    p->y = static_cast<double>(y) / det;
    return 0;
}
```

```

template<typename Pt>
bool point_on_segment(const Pt &p, const Pt &a, const Pt &b) {
    // Overflow risk for integer Pt: these products are ~0(max_coord^2); use int64_t if necessary.
    return EQ((p.x - a.x) * (b.y - a.y) - (p.y - a.y) * (b.x - a.x), 0) &&
        LE(std::min(a.x, b.x), p.x) && LE(p.x, std::max(a.x, b.x)) &&
        LE(std::min(a.y, b.y), p.y) && LE(p.y, std::max(a.y, b.y));
}

// Detection is exact for integer Pt. Intersection point output may use a separate float OutPt.
template<typename Pt, typename OutPt>
int seg_intersection(
    const Pt &a, const Pt &b, const Pt &c, const Pt &d, OutPt *p = nullptr, OutPt *q = nullptr,
    const bool touch_is_intersect = true
) {
    auto ab_x = b.x - a.x, ab_y = b.y - a.y;
    auto ac_x = c.x - a.x, ac_y = c.y - a.y;
    auto cd_x = d.x - c.x, cd_y = d.y - c.y;
    auto ab2 = ab_x * ab_x + ab_y * ab_y;
    auto cd2 = cd_x * cd_x + cd_y * cd_y;
    if (EQ(ab2, 0)) {
        if (touch_is_intersect && point_on_segment(a, c, d)) {
            if (p != nullptr) {
                p->x = static_cast<double>(a.x);
                p->y = static_cast<double>(a.y);
            }
            return 0; // Degenerate: first segment is a point intersecting the second segment.
        }
        return -1; // Degenerate: first segment is a point not intersecting the second segment.
    }
    if (EQ(cd2, 0)) {
        if (touch_is_intersect && point_on_segment(c, a, b)) {
            if (p != nullptr) {
                p->x = static_cast<double>(c.x);
                p->y = static_cast<double>(c.y);
            }
            return 0; // Degenerate: second segment is a point intersecting the first segment.
        }
        return -1; // Degenerate: second segment is a point not intersecting the first segment.
    }
    auto c1 = ab_x * cd_y - ab_y * cd_x;
    auto c2 = ac_x * ab_y - ac_y * ab_x;
    if (EQ(c1, 0) && EQ(c2, 0)) { // Collinear.
        Pt res1 = std::max(std::min(a, b), std::min(c, d));
        Pt res2 = std::min(std::max(a, b), std::max(c, d));
        bool overlap = touch_is_intersect ? !(res2 < res1) : (res1 < res2);
        if (overlap) {
            if (res1 == res2) {
                if (p != nullptr) {
                    p->x = static_cast<double>(res1.x);
                    p->y = static_cast<double>(res1.y);
                }
                return 0; // Collinear and overlapping: touch at one endpoint.
            }
            if (p != nullptr && q != nullptr) {
                p->x = static_cast<double>(res1.x);
                p->y = static_cast<double>(res1.y);
                q->x = static_cast<double>(res2.x);
                q->y = static_cast<double>(res2.y);
            }
            return 1; // Collinear and overlapping: overlap is a segment.
        }
        return -1; // Collinear and disjoint.
    }
    if (EQ(c1, 0)) {
        return -1; // Parallel and disjoint.
    }
    auto t_num = ac_x * cd_y - ac_y * cd_x;

```

```

bool c1_pos = c1 > 0;
bool t_ok = c1_pos ? (touch_is_intersect ? (LE(0, t_num) && LE(t_num, c1))
      : (LT(0, t_num) && LT(t_num, c1)))
      : (touch_is_intersect ? (LE(c1, t_num) && LE(t_num, 0))
      : (LT(c1, t_num) && LT(t_num, 0)));

bool u_ok = c1_pos ? (touch_is_intersect ? (LE(0, c2) && LE(c2, c1)) : (LT(0, c2) && LT(c2, c1)))
      : (touch_is_intersect ? (LE(c1, c2) && LE(c2, 0)) : (LT(c1, c2) && LT(c2, 0)));
if (t_ok && u_ok) {
    if (p != nullptr) {
        double t = static_cast<double>(t_num) / c1;
        p->x = static_cast<double>(a.x + t * ab_x);
        p->y = static_cast<double>(a.y + t * ab_y);
    }
    return 0; // Non-parallel with one intersection.
}
return -1; // Non-parallel with no intersection.
}

// Simplified detection-only version (no output points needed); exact for integer Pt.
// Alternatively, just call the version above and pass static_cast<Pt *>(nullptr) for p and q.
template<typename Pt>
int seg_intersection(
    const Pt &a, const Pt &b, const Pt &c, const Pt &d, const bool touch_is_intersect = true
) {
    auto ab_x = b.x - a.x, ab_y = b.y - a.y;
    auto ac_x = c.x - a.x, ac_y = c.y - a.y;
    auto cd_x = d.x - c.x, cd_y = d.y - c.y;
    auto ab2 = ab_x * ab_x + ab_y * ab_y;
    auto cd2 = cd_x * cd_x + cd_y * cd_y;
    if (EQ(ab2, 0)) {
        return (touch_is_intersect && point_on_segment(a, c, d)) ? 0 : -1;
    }
    if (EQ(cd2, 0)) {
        return (touch_is_intersect && point_on_segment(c, a, b)) ? 0 : -1;
    }
    auto c1 = ab_x * cd_y - ab_y * cd_x;
    auto c2 = ac_x * ab_y - ac_y * ab_x;
    if (EQ(c1, 0) && EQ(c2, 0)) { // Collinear.
        Pt res1 = std::max(std::min(a, b), std::min(c, d));
        Pt res2 = std::min(std::max(a, b), std::max(c, d));
        bool overlap = touch_is_intersect ? !(res2 < res1) : (res1 < res2);
        if (overlap) {
            return (res1 == res2) ? 0 : 1;
        }
        return -1;
    }
    if (EQ(c1, 0)) {
        return -1;
    }
    auto t_num = ac_x * cd_y - ac_y * cd_x;
    bool c1_pos = c1 > 0;
    bool t_ok = c1_pos ? (touch_is_intersect ? (LE(0, t_num) && LE(t_num, c1))
      : (LT(0, t_num) && LT(t_num, c1)))
      : (touch_is_intersect ? (LE(c1, t_num) && LE(t_num, 0))
      : (LT(c1, t_num) && LT(t_num, 0)));

    bool u_ok = c1_pos ? (touch_is_intersect ? (LE(0, c2) && LE(c2, c1)) : (LT(0, c2) && LT(c2, c1)))
      : (touch_is_intersect ? (LE(c1, c2) && LE(c2, 0)) : (LT(c1, c2) && LT(c2, 0)));
    return (t_ok && u_ok) ? 0 : -1;
}

template<typename Pt>
Pt closest_point(double a, double b, double c, const Pt &p) {
    Pt res{};
    if (EQ(a, 0)) {
        res.x = p.x;
        res.y = -c / b;
    } else if (EQ(b, 0)) {
        res.x = -c / a;
    }
}

```



```

    res.y = p.y;
} else {
    line_intersection(a, b, c, -b, a, b * p.x - a * p.y, &res);
}
return res;
}

```

Example Usage

```

#include <cassert>

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
    bool operator!=(const Point &p) const { return !(*this == p); }
    bool operator<(const Point &p) const { return x != p.x ? x < p.x : y < p.y; }
    bool operator>(const Point &p) const { return p < *this; }
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
    bool operator==(const PointI &p) const { return x == p.x && y == p.y; }
    bool operator!=(const PointI &p) const { return !(*this == p); }
    bool operator<(const PointI &p) const { return x != p.x ? x < p.x : y < p.y; }
    bool operator>(const PointI &p) const { return p < *this; }
};

int main() {
    Point p, q;

    assert(line_intersection(-1.0, 1.0, 0.0, 1.0, 1.0, -3.0, &p) == 0);
    assert(EQ(p, Point(1.5, 1.5)));
    assert(line_intersection(Point(0, 0), Point(1, 1), Point(0, 4), Point(4, 0), &p) == 0);
    assert(EQ(p, Point(2, 2)));

    auto test = [&](double a, double b, double c, double d, double e, double f, double g, double h) {
        return seg_intersection(Point(a, b), Point(c, d), Point(e, f), Point(g, h), &p, &q);
    };
    assert(0 == test(-4, 0, 4, 0, 0, -4, 0, 4) && EQ(p, Point(0, 0)));
    assert(0 == test(0, 0, 10, 10, 2, 2, 16, 4) && EQ(p, Point(2, 2)));
    assert(0 == test(-2, 2, -2, -2, -2, 0, 0, 0) && EQ(p, Point(-2, 0)));
    assert(0 == test(0, 4, 4, 4, 4, 0, 4, 8) && EQ(p, Point(4, 4)));
    assert(1 == test(10, 10, 0, 0, 2, 2, 6, 6));
    assert(EQ(p, Point(2, 2)) && EQ(q, Point(6, 6)));
    assert(-1 == test(6, 8, 8, 10, 12, 12, 4, 4));
    assert(-1 == test(-4, 2, -8, 8, 0, 0, -4, 6));

    assert(EQ(Point(2.5, 2.5), closest_point(-1.0, -1.0, 5.0, Point(0, 0))));
    assert(EQ(Point(3, 0), closest_point(1.0, 0.0, -3.0, Point(0, 0))));

    // Detection-only with integer coordinates (exact, no float needed).
    assert(seg_intersection(PointI(0, 0), PointI(4, 4), PointI(0, 4), PointI(4, 0)) == 0);
    assert(seg_intersection(PointI(0, 0), PointI(1, 0), PointI(2, 0), PointI(3, 0)) == -1);
    assert(seg_intersection(PointI(0, 0), PointI(4, 4), PointI(0, 4), PointI(4, 0), &p) == 0);
    assert(EQ(p, Point(2, 2)));

    // touch_is_intersect = false treats endpoint-only contact as non-intersecting. The T-junction
    // and shared-endpoint cases below intersect by default but are rejected when the flag is off.
    assert(0 == seg_intersection(Point(0, 4), Point(4, 4), Point(4, 0), Point(4, 8), &p, &q));
    assert(-1 == seg_intersection(Point(0, 4), Point(4, 4), Point(4, 0), Point(4, 8), &p, &q, false));
    assert(-1 == seg_intersection(PointI(0, 0), PointI(2, 2), PointI(2, 2), PointI(6, 6), false));
    return 0;
}

```

7.2.4 Circle Intersection

Circle tangent and intersection calculations in two dimensions. Such calculations are floating-point by nature, so this section mostly computes in `double`. The input point `p` to `tangent()` can be any type with numeric `.x` and `.y` fields, and intersection point outputs are written through a caller-supplied point type.

Circle radii are normalized to be nonnegative. `tangent()` requires a positive radius, and lines passed to `intersection()` must be valid.

- `tangent(c, p, l1 = nullptr, l2 = nullptr)` determines the line(s) tangent to circle `c` that pass through point `p`, returning `-1` if there is no tangent line because `p` is strictly inside `c`, `0` if there is exactly one tangent line because `p` is on the boundary of `c` (in which case the line will be stored into pointer `l1` if it's not `nullptr`), or `1` if there are two tangent lines because `p` is strictly outside of `c` (in which case the lines will be stored into pointers `l1` and `l2` if they are not `nullptr`).
- `intersection(c, l, p = nullptr, q = nullptr)` determines the intersection between the circle `c` and line `l`, returning `-1` if there is no intersection, `0` if the line has one intersection point because the line is tangent (in which case it will be stored into pointer `p` if it's not `nullptr`), or `1` if there are two intersection points because the line crosses through the circle (stored independently into `p` and `q` when those pointers are non-null).
- `intersection(c1, c2, p = nullptr, q = nullptr)` determines the intersection points between two circles `c1` and `c2`, returning `-2` if circle `c2` completely encloses circle `c1`, `-1` if circle `c1` completely encloses circle `c2`, `0` if the circles are completely disjoint, `1` if the circles are tangent with one intersection (stored in `p` if it's not `nullptr`), `2` if the circles intersect at two points (stored independently in `p` and `q` when those pointers are non-null), or `3` if the circles are equal and intersect at infinite points.
- `intersection_area(c1, c2)` returns the intersection area of circles `c1` and `c2`.

Time Complexity: $O(1)$ per operation.

Space Complexity: $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cmath>
#include <type_traits>

const double EPS = 1e-9, PI = acos(-1.0);

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U>
bool LE(T a, U b) {
    return !LT(b, a);
}

struct Circle {
    double h, k, r;

    Circle(double h, double k, double r) : h(h), k(k), r(fabs(r)) {}
};

struct Line {
    double a, b, c;

    Line() : a(0), b(0), c(0) {}
};
```

```

Line(double a, double b, double c) {
    if (!EQ(b, 0)) {
        this->a = a / b;
        this->c = c / b;
        this->b = 1;
    } else {
        this->c = c / a;
        this->a = 1;
        this->b = 0;
    }
}

Line(double px, double py, double qx, double qy) : a(0), b(0), c(0) {
    if (EQ(px, qx)) {
        if (!EQ(py, qy)) { // Vertical line.
            a = 1;
            b = 0;
            c = -px;
        } // Else, invalid line.
    } else {
        a = -static_cast<double>(py - qy) / (px - qx);
        b = 1;
        c = -(a * px) - (b * py);
    }
}

bool operator==(const Line &l1) const { return a == l1.a && b == l1.b && c == l1.c; }

friend bool EQ(const Line &l1, const Line &l2) {
    return EQ(l1.a, l2.a) && EQ(l1.b, l2.b) && EQ(l1.c, l2.c);
}
};

template<typename Pt>
int tangent(const Circle &c, const Pt &p, Line *l1 = nullptr, Line *l2 = nullptr) {
    auto sqnorm = [](double x, double y) { return x * x + y * y; };
    double vopx = p.x - c.h, vopy = p.y - c.k;
    if (EQ(sqnorm(vopx, vopy), c.r * c.r)) { // Point on an edge.
        if (l1 != nullptr) { // Tangent through p is perpendicular to the radius cp.
            *l1 = Line(vopx, vopy, -(vopx * p.x + vopy * p.y));
        }
        return 0;
    }
    if (LE(sqnorm(vopx, vopy), c.r * c.r)) {
        return -1; // Point inside Circle, no intersection.
    }
    double qx = vopx / c.r, qy = vopy / c.r;
    double n = sqnorm(qx, qy), d = qy * sqrt(n - 1.0);
    double t1x = (qx - d) / n, t1y = c.k;
    double t2x = (qx + d) / n, t2y = c.k;
    if (!EQ(qy, 0)) { // Common case.
        t1y += c.r * (1.0 - t1x * qx) / qy;
        t2y += c.r * (1.0 - t2x * qx) / qy;
    } else { // Point at center horizontal, y = 0.
        d = c.r * sqrt(1.0 - t1x * t1x);
        t1y += d;
        t2y -= d;
    }
    t1x = t1x * c.r + c.h;
    t2x = t2x * c.r + c.h;
    // Note: here, t1 and t2 are the two points of tangencies
    if (l1 != nullptr) {
        *l1 = Line(p.x, p.y, t1x, t1y);
    }
    if (l2 != nullptr) {
        *l2 = Line(p.x, p.y, t2x, t2y);
    }
}

```

```

    return 1;
}

template<typename OutPt>
int intersection(const Circle &c, const Line &l, OutPt *p = nullptr, OutPt *q = nullptr) {
    double v = c.h * l.a + c.k * l.b + l.c;
    double aabb = l.a * l.a + l.b * l.b;
    double disc = v * v / aabb - c.r * c.r;
    if (disc > EPS) {
        return -1;
    }
    double x0 = -l.a * v / aabb, y0 = -l.b * v / aabb;
    if (disc > -EPS) {
        if (p != nullptr) {
            p->x = x0 + c.h;
            p->y = y0 + c.k;
        }
        return 0;
    }
    double k = sqrt(std::max(0.0, disc / -aabb));
    if (p != nullptr) {
        p->x = x0 + k * l.b + c.h;
        p->y = y0 - k * l.a + c.k;
    }
    if (q != nullptr) {
        q->x = x0 - k * l.b + c.h;
        q->y = y0 + k * l.a + c.k;
    }
    return 1;
}

template<typename OutPt>
int intersection(const Circle &c1, const Circle &c2, OutPt *p = nullptr, OutPt *q = nullptr) {
    if (EQ(c1.h, c2.h) && EQ(c1.k, c2.k)) {
        return EQ(c1.r, c2.r) ? 3 : (c1.r > c2.r ? -1 : -2);
    }
    double d12x = c2.h - c1.h, d12y = c2.k - c1.k;
    double d = hypot(d12x, d12y);
    if (LT(c1.r + c2.r, d)) {
        return 0;
    }
    if (LT(d, fabs(c1.r - c2.r))) {
        return c1.r > c2.r ? -1 : -2;
    }
    double a = (c1.r * c1.r - c2.r * c2.r + d * d) / (2 * d);
    double x0 = c1.h + (d12x * a / d), y0 = c1.k + (d12y * a / d);
    double s = sqrt(std::max(0.0, c1.r * c1.r - a * a));
    double rx = -d12y * s / d, ry = d12x * s / d;
    if (EQ(rx, 0) && EQ(ry, 0)) {
        if (p != nullptr) {
            p->x = x0;
            p->y = y0;
        }
        return 1;
    }
    if (p != nullptr) {
        p->x = x0 - rx;
        p->y = y0 - ry;
    }
    if (q != nullptr) {
        q->x = x0 + rx;
        q->y = y0 + ry;
    }
    return 2;
}

double intersection_area(const Circle &c1, const Circle &c2) {
    double r = std::min(c1.r, c2.r), R = std::max(c1.r, c2.r);

```

```

double d = std::hypot(c2.h - c1.h, c2.k - c1.k);
if (LE(d, R - r)) {
    return PI * r * r;
}
if (LE(R + r, d)) {
    return 0;
}
double alpha = std::max(-1.0, std::min(1.0, (d * d + r * r - R * R) / 2 / d / r));
double beta = std::max(-1.0, std::min(1.0, (d * d + R * R - r * r) / 2 / d / R));
return r * r * acos(alpha) + R * R * acos(beta) -
    0.5 * sqrt((-d + r + R) * (d + r - R) * (d - r + R) * (d + r + R));
}

```

Example Usage

```

#include <cassert>

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
};

int main() {
    Line l1, l2;
    assert(-1 == tangent(Circle(0, 0, 4), Point(1, 1), &l1, &l2));
    assert(0 == tangent(Circle(0, 0, sqrt(2)), Point(1, 1), &l1, &l2));
    assert(EQ(l1, Line(-1, -1, 2)));
    assert(1 == tangent(Circle(0, 0, 2), Point(2, 2), &l1, &l2));
    assert(EQ(l1, Line(0, -2, 4)));
    assert(EQ(l2, Line(2, 0, -4)));

    Point p, q;
    assert(-1 == intersection(Circle(1, 1, 3), Line(5, 3, -30), &p, &q));
    assert(0 == intersection(Circle(1, 1, 3), Line(0, 1, -4), &p, &q));
    assert(EQ(p, Point(1, 4)));
    assert(1 == intersection(Circle(1, 1, 3), Line(0, 1, -1), &p, &q));
    assert(EQ(p, Point(4, 1)));
    assert(EQ(q, Point(-2, 1)));
    assert(1 == intersection(Circle(1, 1, 3), Line(0, 1, -1), &p));
    assert(EQ(p, Point(4, 1)));

    assert(-2 == intersection(Circle(1, 1, 1), Circle(0, 0, 3), &p, &q));
    assert(-1 == intersection(Circle(0, 0, 3), Circle(1, 1, 1), &p, &q));
    assert(0 == intersection(Circle(5, 0, 4), Circle(-5, 0, 4), &p, &q));
    assert(1 == intersection(Circle(-5, 0, 5), Circle(5, 0, 5), &p, &q));
    assert(EQ(p, Point(0, 0)));
    assert(2 == intersection(Circle(-0.5, 0, 1), Circle(0.5, 0, 1), &p, &q));
    assert(EQ(p, Point(0, -sqrt(3) / 2)));
    assert(EQ(q, Point(0, sqrt(3) / 2)));
    assert(
        2 == intersection(Circle(-0.5, 0, 1), Circle(0.5, 0, 1), static_cast<Point *>(nullptr), &q)
    );
    assert(EQ(q, Point(0, sqrt(3) / 2)));

    // Each circle passes through the other's center.
    double r = 3, a = intersection_area(Circle(-r / 2, 0, r), Circle(r / 2, 0, r));
    assert(EQ(a, r * r * (2 * PI / 3 - sqrt(3) / 2)));
    return 0;
}

```

7.2.5 Circle Tangents

Given two circles, compute their common tangent lines by returning the two points of tangency on each circle. External tangents touch both circles on the same side of the line joining their centers; internal tangents cross between the circles. A point can be treated as a zero-radius circle, so this also covers tangents from a point to a circle.

- `circle_tangents(c1, r1, c2, r2, internal = false)` returns 0, 1, or 2 tangent point pairs. Each pair (p, q) means the tangent line touches the first circle at p and the second circle at q . Set `internal = true` to get internal tangents; otherwise external tangents are returned.
- `all_circle_tangents(c1, r1, c2, r2)` returns all external and internal tangents.

If the two circles are identical, there are infinitely many tangents and the functions return an empty vector. If one circle strictly contains the other, there are no external tangents; if the circles overlap, there are no internal tangents.

Time Complexity: $O(1)$ per call.

Space Complexity: $O(1)$ auxiliary, excluding the returned vector.

```
#include <algorithm>
#include <cmath>
#include <type_traits>
#include <utility>
#include <vector>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
    Point operator+(const Point &p) const { return {x + p.x, y + p.y}; }
    Point operator-(const Point &p) const { return {x - p.x, y - p.y}; }
    Point operator*(double k) const { return {x * k, y * k}; }
    Point operator/(double k) const { return {x / k, y / k}; }
    double sqnorm() const { return x * x + y * y; }
    Point rotate90() const { return {-y, x}; }
};

template<typename PtA, typename PtB>
std::vector<std::pair<Point, Point>> circle_tangents(
    const PtA &c1, double r1, const PtB &c2, double r2, const bool internal = false
) {
    r1 = fabs(r1);
    r2 = fabs(r2);
    if (internal) {
        r2 = -r2;
    }
    Point d(c2.x - c1.x, c2.y - c1.y);
    double dr = r1 - r2, d2 = d.sqnorm(), h2 = d2 - dr * dr;
    if (EQ(d2, 0) || LT(h2, 0)) {
        return {};
    }
}
```

```

    }
    std::vector<std::pair<Point, Point>> res;
    for (double s : {-1.0, 1.0}) {
        Point v = (d * dr + d.rotate90() * (sqrt(std::max(0.0, h2)) * s)) / d2;
        res.push_back({Point(c1.x, c1.y) + v * r1, Point(c2.x, c2.y) + v * r2});
    }
    if (EQ(h2, 0)) {
        res.pop_back();
    }
    return res;
}

template<typename PtA, typename PtB>
std::vector<std::pair<Point, Point>> all_circle_tangents(
    const PtA &c1, double r1, const PtB &c2, double r2
) {
    auto res = circle_tangents(c1, r1, c2, r2);
    auto internal = circle_tangents(c1, r1, c2, r2, true);
    res.insert(res.end(), internal.begin(), internal.end());
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    Point a(0, 0), b(4, 0);
    auto ext = circle_tangents(a, 1, b, 1);
    assert(ext.size() == 2);
    assert(EQ(ext[0].first, Point(0, -1)) && EQ(ext[0].second, Point(4, -1)));
    assert(EQ(ext[1].first, Point(0, 1)) && EQ(ext[1].second, Point(4, 1)));

    auto in = circle_tangents(a, 1, b, 1, true);
    assert(in.size() == 2);
    assert(
        EQ(in[0].first, Point(0.5, -sqrt(3.0) / 2)) && EQ(in[0].second, Point(3.5, sqrt(3.0) / 2))
    );
    assert(
        EQ(in[1].first, Point(0.5, sqrt(3.0) / 2)) && EQ(in[1].second, Point(3.5, -sqrt(3.0) / 2))
    );

    auto point_to_circle = circle_tangents(Point(0, 0), 0, Point(2, 0), 1);
    assert(point_to_circle.size() == 2);
    assert(EQ(point_to_circle[0].first, Point(0, 0)));

    assert(circle_tangents(Point(0, 0), 5, Point(1, 0), 1).empty());           // contained
    assert(circle_tangents(Point(0, 0), 2, Point(3, 0), 2, true).empty());    // overlapping
    assert(all_circle_tangents(Point(0, 0), 1, Point(4, 0), 1).size() == 4);
    return 0;
}

```

7.3 Polygons and Point Sets

7.3.1 Polygon Sorting and Area

Given a list of distinct points in two dimensions, order them into a valid polygon and determine the area. The sorting comparators order two points by the sign of their cross product around a chosen center, that is, by angle. The area functions use the shoelace formula, summing the cross products of consecutive vertex pairs. The

functions are templated on the point type. The local `Point` struct (`double` coordinates) is the default; replace it with `Point`, `PointD`, or `PointI` from 7.1.1 or any struct with numeric `.x` and `.y` fields and `<` / `==` operators.

- `cw_comp(a, b, c)` returns whether point `a` compares clockwise before `b` about `c`.
- `polygon_area_2x(lo, hi)` returns exactly double the area of the polygon with vertices specified by the range `[lo, hi)` of points in either clockwise or counter-clockwise order. The return value is integral or floating-point, depending on the input point type. For integer vertices, divide by 2 in the caller if the exact area is needed.
- `polygon_area(lo, hi)` returns the area as `double`.
- `polygon_centroid(lo, hi)` returns the centroid (center of mass) of a non-degenerate simple polygon with vertices in boundary order as an `std::pair<double, double>`. It uses signed shoelace weights, so either clockwise or counter-clockwise inputs will work.

Overflow warning: `cw_comp` and `polygon_area_2x` form cross products that grow like the squared coordinate magnitude (and the shoelace sum accumulates over all vertices). For integer point types use a 64-bit coordinate type (e.g. `PointL` from 7.1.1) for large or numerous coordinates.

Time Complexity:

- $O(n)$ per call to `polygon_area_2x()`, `polygon_area()`, and `polygon_centroid()`, where n is the number of points.
- $O(1)$ per call to `cw_comp()` and the comparators.

Space Complexity: $O(1)$ auxiliary for all operations.

```
#include <algorithm>
#include <cassert>
#include <cmath>
#include <random>
#include <stdexcept>
#include <type_traits>
#include <utility>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

// Comparator (clockwise angular order about c). Exact: comparisons are pure sign tests on
// coordinate differences and the cross product, so it is a valid strict weak ordering and is
// exact for integer-coordinate points.
template<typename Pt>
bool cw_comp(const Pt &a, const Pt &b, const Pt &c) {
    if (a.x - c.x >= 0 && b.x - c.x < 0) {
        return true;
    }
    if (a.x - c.x < 0 && b.x - c.x >= 0) {
        return false;
    }
    if (a.x - c.x == 0 && b.x - c.x == 0) {
        if (a.y - c.y >= 0 || b.y - c.y >= 0) {
            return a.y > b.y;
        }
        return b.y > a.y;
    }
}

auto acx = a.x - c.x, acy = a.y - c.y;
auto bcx = b.x - c.x, bcy = b.y - c.y;
auto det = acx * bcy - acy * bcx; // Overflow warning!
if (det == 0) {
    auto acnorm = acx * acx + acy * acy; // Overflow warning!
    auto bcnorm = bcx * bcx + bcy * bcy;
    return acnorm > bcnorm;
}
return det < 0;
```



```

}

// Returns 2 * area. Result is exact (integer) for integer-coordinate points.
template<typename It>
auto polygon_area_2x(It lo, It hi) {
    using T = decltype(lo->x * lo->y);
    if (lo == hi) {
        return T(0);
    }
    T area = 0;
    for (It i = lo, j = hi - 1; i != hi; j = i++) {
        area += static_cast<T>(j->x - i->x) * static_cast<T>(j->y + i->y); // Overflow warning!
    }
    return area < T(0) ? -area : area;
}

template<typename It>
double polygon_area(It lo, It hi) {
    return static_cast<double>(polygon_area_2x(lo, hi)) / 2.0;
}

template<typename It>
std::pair<double, double> polygon_centroid(It lo, It hi) {
    assert(hi - lo >= 3);
    double cx = 0, cy = 0, area2 = 0;
    for (It i = lo, j = hi - 1; i != hi; j = i++) {
        double cross = static_cast<double>(j->x) * i->y - static_cast<double>(i->x) * j->y;
        cx += (static_cast<double>(j->x) + i->x) * cross;
        cy += (static_cast<double>(j->y) + i->y) * cross;
        area2 += cross;
    }
    if (EQ(area2, 0)) {
        throw std::runtime_error("Cannot compute centroid of zero-area polygon.");
    }
    return {cx / (3 * area2), cy / (3 * area2)};
}

```

Example Usage

```

#include <cassert>
#include <vector>
using namespace std;

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
    bool operator!=(const Point &p) const { return !(*this == p); }
    bool operator<(const Point &p) const { return x != p.x ? x < p.x : y < p.y; }
    bool operator>(const Point &p) const { return p < *this; }
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
    bool operator==(const PointI &p) const { return x == p.x && y == p.y; }
    bool operator!=(const PointI &p) const { return !(*this == p); }
    bool operator<(const PointI &p) const { return x != p.x ? x < p.x : y < p.y; }
    bool operator>(const PointI &p) const { return p < *this; }
};

template<typename It>
Point mean_center(It lo, It hi) {
    assert(lo != hi);
    double x_sum = 0, y_sum = 0, n = hi - lo;

```

```

for (It it = lo; it != hi; ++it) {
    x_sum += it->x;
    y_sum += it->y;
}
return Point(x_sum / n, y_sum / n);
}

int main() {
    vector<Point> points{Point(1, 3), Point(1, 2), Point(2, 1), Point(0, 0), Point(-1, 3)};
    vector<Point> v(points);
    std::mt19937 rng(1234567); // Fixed seed for reproducibility.
    std::shuffle(v.begin(), v.end(), rng);
    Point c = mean_center(v.begin(), v.end());
    assert(EQ(c, Point(0.6, 1.8)));
    sort(v.begin(), v.end(), [c](const Point &a, const Point &b) { return cw_comp(a, b, c); });
    assert((v == vector<Point>{{1, 3}, {1, 2}, {2, 1}, {0, 0}, {-1, 3}}));
    assert(EQ(polygon_area(v.begin(), v.end()), 5));

    sort(v.begin(), v.end(), [c](const Point &a, const Point &b) { return cw_comp(b, a, c); });
    assert((v == vector<Point>{{-1, 3}, {0, 0}, {2, 1}, {1, 2}, {1, 3}}));
    assert(EQ(polygon_area(v.begin(), v.end()), 5));

    // Integer points: polygon_area_2x is exact (no float arithmetic).
    vector<PointI> iv{{0, 0}, {4, 0}, {0, 3}}; // right triangle, area = 6
    assert(polygon_area_2x(iv.begin(), iv.end()) == 12); // exact int
    assert(EQ(polygon_area(iv.begin(), iv.end()), 6.0));
    auto centroid = polygon_centroid(iv.begin(), iv.end());
    assert(EQ(centroid.first, 4.0 / 3.0) && EQ(centroid.second, 1.0));
    return 0;
}

```

7.3.2 Point-in-Polygon

Given a point `p` and a polygon, determines whether `p` lies inside using ray casting. A ray cast from `p` must cross the polygon's boundary an odd number of times exactly when `p` is inside, so the function counts crossings between a horizontal ray and each polygon edge. The function is templated on the point type `Pt`. The local `Point` struct (double coordinates) is the default; replace it with `PointI`, `PointD`, or `PointI` from 7.1.1 or any struct with numeric `.x` and `.y` fields. All comparisons are exact (no epsilon): the ray-crossing parity needs a consistent comparison to be correct, and the result is exact for integer-coordinate points.

- `point_in_polygon(p, lo, hi, edge_is_inside = true)` returns whether `p` lies within the polygon with vertices specified by the range `[lo, hi)` of points in either clockwise or counter-clockwise order. The `edge_is_inside` flag controls whether points exactly on an edge or vertex count as inside.

Overflow warning: the edge orientation test forms a cross product that grows like the squared coordinate magnitude. For integer point types use a 64-bit coordinate type (e.g. `PointL` from 7.1.1) once coordinates exceed roughly a few tens of thousands.

Time Complexity: $O(n)$ per call, where n is the distance between `lo` and `hi`.

Space Complexity: $O(1)$ auxiliary.

```

template<typename Pt>
auto cross(const Pt &a, const Pt &b, const Pt &o) {
    // Overflow risk for integer Pt: these products are ~O(max_coord^2); use int64_t if necessary.
    return (a.x - o.x) * (b.y - o.y) - (a.y - o.y) * (b.x - o.x);
}

// Detection is exact for integer-coordinate Pt.
template<typename Pt, typename It>
bool point_in_polygon(const Pt &p, It lo, It hi, bool edge_is_inside = true) {
    if (lo == hi) {
        return false;
    }

```

```

}
bool res = false;
for (It i = lo, j = hi - 1; i != hi; j = i++) {
    if (i->y == p.y && (i->x == p.x || (j->y == p.y && (i->x <= p.x || j->x <= p.x)))) {
        return edge_is_inside;
    }
    if ((p.y < i->y) != (p.y < j->y)) {
        auto det = cross(*i, *j, p);
        if (det == 0) {
            return edge_is_inside;
        }
        if ((0 < det) != (i->y < j->y)) {
            res = !res;
        }
    }
}
return res;
}

```

Example Usage

```

#include <cassert>
#include <vector>
using namespace std;

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
};

int main() {
    vector<Point> poly{Point(-1, 3), Point(1, 3), Point(2, 1), Point(0, 0)};
    assert(point_in_polygon(Point(1, 2), poly.begin(), poly.end()));
    assert(point_in_polygon(Point(0, 3), poly.begin(), poly.end()));
    assert(!point_in_polygon(Point(0, 3), poly.begin(), poly.end(), false));
    assert(!point_in_polygon(Point(0, 3.01), poly.begin(), poly.end()));
    assert(!point_in_polygon(Point(2, 2), poly.begin(), poly.end()));

    // Integer-coordinate points: exact detection.
    vector<PointI> ipoly{{0, 0}, {4, 0}, {4, 4}, {0, 4}};
    assert(point_in_polygon(PointI(2, 2), ipoly.begin(), ipoly.end()));
    assert(!point_in_polygon(PointI(4, 2), ipoly.begin(), ipoly.end(), false));
    assert(!point_in_polygon(PointI(5, 2), ipoly.begin(), ipoly.end()));
    return 0;
}

```

7.3.3 Convex Hull and Diametral Pair

Given a list of points in two dimensions, computes the convex hull using the monotone chain algorithm and the diametral pair using rotating calipers. Monotone chain sorts the points lexicographically and builds the lower and upper hulls in one pass each, popping any point that would fail to create a counter-clockwise turn. Rotating calipers then walks two antipodal pointers around the hull, advancing whichever increases the separation, visiting every candidate diametral pair in linear time. The functions are templated on the iterator type; the value type of the iterator is used as the point type. Both functions accept either floating-point or integral coordinates, and use exact comparisons for integral points.

- `convex_hull(lo, hi)` returns the convex hull in counter-clockwise order for a range `[lo, hi)` of points. Duplicate input points and collinear points along hull edges are omitted. To instead return the hull points in clockwise order, replace the cross product comparisons `<= 0` with `>= 0`.
- `diametral_pair(lo, hi)` returns the maximum-distance pair of points.

Overflow warning: `cross()` and the squared distance in `diametral_pair()` grow like the squared coordinate magnitude. With 32-bit `int` coordinates they overflow once coordinates exceed a few tens of thousands; use a 64-bit (`int64_t`) coordinate type for larger integer inputs.

Time Complexity: $O(n \log n)$ per call, where n is the distance between `lo` and `hi`.

Space Complexity: $O(n)$ auxiliary for storage of the convex hull.

```
#include <algorithm>
#include <cmath>
#include <iterator>
#include <random>
#include <utility>
#include <vector>

template<typename Pt>
auto cross(const Pt &a, const Pt &b, const Pt &o) {
    // Overflow risk for integer Pt: these products are ~O(max_coord^2); use int64_t if necessary.
    return (a.x - o.x) * (b.y - o.y) - (a.y - o.y) * (b.x - o.x);
}

// Convex hull: exact for integer-coordinate points.
template<typename It>
auto convex_hull(It lo, It hi) {
    using Pt = typename std::iterator_traits<It>::value_type;
    std::vector<Pt> p(lo, hi);
    std::sort(p.begin(), p.end());
    auto same_point = [](const Pt &a, const Pt &b) { return !(a < b) && !(b < a); };
    p.erase(std::unique(p.begin(), p.end(), same_point), p.end());
    int n = static_cast<int>(p.size());
    if (n <= 1) {
        return p;
    }
    int k = 0;
    std::vector<Pt> res(2 * n);
    for (const Pt &q : p) {
        while (k >= 2 && cross(res[k - 1], q, res[k - 2]) <= 0) {
            k--;
        }
        res[k++] = q;
    }
    int t = k + 1;
    for (int i = n - 2; i >= 0; i--) {
        while (k >= t && cross(res[k - 1], p[i], res[k - 2]) <= 0) {
            k--;
        }
        res[k++] = p[i];
    }
    res.resize(k - 1);
    return res;
}

// Diametral pair: squared-distance comparisons are exact for integer-coordinate points.
template<typename It>
auto diametral_pair(It lo, It hi) {
    using Pt = typename std::iterator_traits<It>::value_type;
    auto h = convex_hull(lo, hi);
    int m = static_cast<int>(h.size());
    if (m == 0) {
        return std::pair<Pt, Pt>{};
    }
    if (m == 1) {
```

```

    return std::pair<Pt, Pt>{h[0], h[0]};
}
if (m == 2) {
    return std::pair<Pt, Pt>{h[0], h[1]};
}
int k = 1;
while (std::abs(cross(h[0], h[(k + 1) % m], h[m - 1])) > std::abs(cross(h[0], h[k], h[m - 1]))) {
    k++;
}
auto sqdist = [](const Pt &a, const Pt &b) {
    // Overflow risk for integer Pt: ~0(max_coord^2); use int64_t if necessary.
    auto dx = a.x - b.x, dy = a.y - b.y;
    return dx * dx + dy * dy;
};
auto maxsq = sqdist(h[0], h[0]);
std::pair<Pt, Pt> res{h[0], h[0]};
for (int i = 0, j = k; i <= k && j < m; i++) {
    auto d = sqdist(h[i], h[j]);
    if (d > maxsq) {
        maxsq = d;
        res = {h[i], h[j]};
    }
    while (j < m && std::abs(cross(h[(i + 1) % m], h[(j + 1) % m], h[i])) >
            std::abs(cross(h[(i + 1) % m], h[j], h[i]))) {
        d = sqdist(h[i], h[(j + 1) % m]);
        if (d > maxsq) {
            maxsq = d;
            res = {h[i], h[(j + 1) % m]};
        }
        j++;
    }
}
return res;
}

```

Example Usage

```

#include <cassert>
#include <vector>
using namespace std;

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
    bool operator==(const PointI &o) const { return x == o.x && y == o.y; }
    bool operator!=(const PointI &o) const { return !(*this == o); }
    bool operator<(const PointI &o) const { return x != o.x ? x < o.x : y < o.y; }
    bool operator>(const PointI &o) const { return o < *this; }
};

int main() {
    {
        vector<PointI> v{{1, 3}, {1, 2}, {2, 1}, {0, 0}, {-1, 3}};
        std::mt19937 rng(1234567); // Fixed seed for reproducibility.
        std::shuffle(v.begin(), v.end(), rng);
        vector<PointI> h{{-1, 3}, {0, 0}, {2, 1}, {1, 3}};
        assert(convex_hull(v.begin(), v.end()) == h);
    }
    {
        vector<PointI> v{{0, 0}, {3, 0}, {0, 3}, {1, 1}, {4, 4}};
        auto [p1, p2] = diametral_pair(v.begin(), v.end());
        assert(p1 == PointI(0, 0));
        assert(p2 == PointI(4, 4));
    }
    {
        vector<PointI> v{{0, 0}, {4, 0}, {4, 4}, {0, 4}, {2, 2}};
    }
}

```

```

auto h = convex_hull(v.begin(), v.end());
assert(h.size() == 4); // interior point (2,2) excluded
auto [p1, p2] = diametral_pair(v.begin(), v.end());
// diametral pair is exact
assert(
    (p1 == PointI(0, 0) && p2 == PointI(4, 4)) || (p1 == PointI(4, 4) && p2 == PointI(0, 0)) ||
    (p1 == PointI(4, 0) && p2 == PointI(0, 4)) || (p1 == PointI(0, 4) && p2 == PointI(4, 0))
);
}
{
    vector<PointI> v{{2, 3}, {2, 3}, {2, 3}};
    vector<PointI> expected{{2, 3}};
    assert(convex_hull(v.begin(), v.end()) == expected);
}
return 0;
}

```

7.3.4 Convex Polygon Queries

Given a convex polygon in counter-clockwise order, answer containment and tangent/intersection queries. Point containment is logarithmic by splitting the polygon into triangles sharing `poly[0]`. The line and tangent helpers below are written as simple linear scans; this is shorter and less assumption-heavy than the classic logarithmic versions, while still documenting the exact query semantics in one place.

- `point_in_convex_polygon(poly, p, edge_is_inside = true)` returns whether point `p` lies inside the convex polygon `poly`. The polygon must be in counter-clockwise order, with no repeated final vertex. Points on the boundary count as inside unless `edge_is_inside` is set to `false`.
- `line_convex_polygon_intersection(poly, a, b)` describes where the infinite directed line through `a` $\rightarrow$ `b` hits the polygon: $(-1, -1)$ for no intersection, $(i, -1)$ for touching vertex i , (i, i) for containing side $i \rightarrow i + 1$, or (i, j) for crossing sides $i \rightarrow i + 1$ and $j \rightarrow j + 1$, with side indices taken modulo the number of vertices. The polygon must have at least three vertices, and `a` and `b` must differ.
- `convex_polygon_tangents(poly, p)` returns the two tangent vertex indices from an outside point `p` as `(left, right)`. The left tangent has all polygon vertices on or to the left of the directed line `p` $\rightarrow$ `poly[left]`; the right tangent is analogous for the right side.

For integer-coordinate inputs, all tests are exact provided the cross products do not overflow.

Time Complexity:

- $O(\log n)$ per call to `point_in_convex_polygon()`, where n is the number of polygon vertices.
- $O(n)$ per call to `line_convex_polygon_intersection` and `convex_polygon_tangents`.

Space Complexity: $O(1)$ auxiliary.

```

#include <algorithm>
#include <cassert>
#include <cstdint>
#include <utility>
#include <vector>

template<typename Pt>
auto cross(const Pt &a, const Pt &b, const Pt &o) {
    // Overflow risk for integer Pt: these products are ~O(max_coord^2); use int64_t if necessary.
    return (a.x - o.x) * (b.y - o.y) - (a.y - o.y) * (b.x - o.x);
}

template<typename T>
int sgn(const T &x) {
    return (T(0) < x) - (x < T(0));
}

template<typename Pt>

```

```

bool on_segment(const Pt &p, const Pt &a, const Pt &b) {
    return cross(a, b, p) == 0 && std::min(a.x, b.x) <= p.x && p.x <= std::max(a.x, b.x) &&
        std::min(a.y, b.y) <= p.y && p.y <= std::max(a.y, b.y);
}

template<typename Pt>
bool point_in_convex_polygon(
    const std::vector<Pt> &poly, const Pt &p, const bool edge_is_inside = true
) {
    int n = static_cast<int>(poly.size());
    if (n == 0) {
        return false;
    }
    if (n == 1) {
        return edge_is_inside && p.x == poly[0].x && p.y == poly[0].y;
    }
    if (n == 2) {
        return edge_is_inside && on_segment(p, poly[0], poly[1]);
    }
    int left = sgn(cross(poly[1], p, poly[0]));
    int right = sgn(cross(poly[n - 1], p, poly[0]));
    if (left < 0 || right > 0) {
        return false;
    }
    if (left == 0) {
        return edge_is_inside && on_segment(p, poly[0], poly[1]);
    }
    if (right == 0) {
        return edge_is_inside && on_segment(p, poly[0], poly[n - 1]);
    }
    int lo = 1, hi = n - 1;
    while (hi - lo > 1) {
        int mid = (lo + hi) / 2;
        if (cross(poly[mid], p, poly[0]) >= 0) {
            lo = mid;
        } else {
            hi = mid;
        }
    }
    int s = sgn(cross(poly[lo + 1], p, poly[lo]));
    return edge_is_inside ? s >= 0 : s > 0;
}

template<typename Pt>
std::pair<int, int> line_convex_polygon_intersection(
    const std::vector<Pt> &poly, const Pt &a, const Pt &b
) {
    int n = static_cast<int>(poly.size());
    assert(n >= 3 && (a.x != b.x || a.y != b.y));
    std::vector<int> on, cross_edges;
    bool has_pos = false, has_neg = false;
    for (int i = 0; i < n; i++) {
        int s = sgn(cross(b, poly[i], a));
        has_pos |= s > 0;
        has_neg |= s < 0;
        if (s == 0) {
            on.push_back(i);
        }
    }
    if (!has_pos && !has_neg) {
        return n >= 2 ? std::make_pair(0, 0) : std::make_pair(-1, -1);
    }
    if (!(has_pos && has_neg)) {
        for (int i = 0; i < static_cast<int>(on.size()); i++) {
            int j = (i + 1) % static_cast<int>(on.size());
            if ((on[i] + 1) % n == on[j]) {
                return {on[i], on[j]};
            }
        }
    }
}

```

```

    }
    return on.empty() ? std::make_pair(-1, -1) : std::make_pair(on[0], -1);
}
for (int i = 0; i < n; i++) {
    int j = (i + 1) % n;
    int si = sgn(cross(b, poly[i], a)), sj = sgn(cross(b, poly[j], a));
    if (si == 0 || si * sj < 0) {
        cross_edges.push_back(i);
    }
}
if (cross_edges.size() == 1) {
    return {cross_edges[0], -1};
}
auto line_parameter = [&](int i) {
    int j = (i + 1) % n;
    auto ex = poly[j].x - poly[i].x, ey = poly[j].y - poly[i].y;
    auto ax = a.x, ay = a.y, dx = b.x - a.x, dy = b.y - a.y;
    auto det = dx * ey - dy * ex;
    return static_cast<double>((poly[i].x - ax) * ey - (poly[i].y - ay) * ex) / det;
};
if (line_parameter(cross_edges[1]) < line_parameter(cross_edges[0])) {
    std::swap(cross_edges[0], cross_edges[1]);
}
return {cross_edges[0], cross_edges[1]};
}

template<typename Pt>
std::pair<int, int> convex_polygon_tangents(const std::vector<Pt> &poly, const Pt &p) {
    int n = static_cast<int>(poly.size());
    std::pair<int, int> res{-1, -1};
    for (int i = 0; i < n; i++) {
        int prev = (i + n - 1) % n, next = (i + 1) % n;
        int s1 = sgn(cross(poly[i], poly[prev], p));
        int s2 = sgn(cross(poly[i], poly[next], p));
        bool left = s1 >= 0 && s2 >= 0;
        bool right = s1 <= 0 && s2 <= 0;
        if (left) {
            res.first = i;
        }
        if (right) {
            res.second = i;
        }
    }
    return res;
}

```

Example Usage

```

#include <algorithm>
using namespace std;

struct PointI {
    int64_t x, y;
    PointI(int64_t x = 0, int64_t y = 0) : x(x), y(y) {}
};

int main() {
    vector<PointI> square{{0, 0}, {4, 0}, {4, 4}, {0, 4}};
    assert(point_in_convex_polygon(square, PointI(2, 2)));
    assert(point_in_convex_polygon(square, PointI(4, 2)));
    assert(!point_in_convex_polygon(square, PointI(4, 2), false));
    assert(!point_in_convex_polygon(square, PointI(5, 2)));

    vector<PointI> pentagon{{0, 0}, {3, 0}, {5, 2}, {2, 5}, {-1, 3}};
    assert(point_in_convex_polygon(pentagon, PointI(2, 3)));
    assert(point_in_convex_polygon(pentagon, PointI(0, 0)));
}

```



```

assert(!point_in_convex_polygon(pentagon, PointI(0, 0), false));
assert(!point_in_convex_polygon(pentagon, PointI(5, 4)));

assert(line_convex_polygon_intersection(square, PointI(-1, 2), PointI(5, 2)) == make_pair(3, 1));
assert(line_convex_polygon_intersection(square, PointI(0, 0), PointI(4, 0)) == make_pair(0, 0));
assert(line_convex_polygon_intersection(square, PointI(5, 0), PointI(5, 4)) == make_pair(-1, -1));

// From the point to the right of the square, the upper-right and lower-right vertices are
// tangent.
auto tangents = convex_polygon_tangents(square, PointI(6, 2));
assert(tangents == make_pair(2, 1));
return 0;
}

```

7.3.5 Minimum Enclosing Circle

Given a set of points in two dimensions, finds the unique circle of minimum radius (equivalently, minimum area) containing all points.

This implementation uses Welzl's randomized incremental algorithm. The points are shuffled and processed one at a time while maintaining the current minimum enclosing circle. Whenever a point lies outside the current circle, the circle is rebuilt with that point constrained to lie on the boundary. The final circle is determined by at most three boundary points.

- `minimum_enclosing_circle(lo, hi)` returns the minimum enclosing circle for the range $[lo, hi)$ of points, where `lo` and `hi` must be random-access iterators. The input range is shuffled in place. The point type may be any type exposing numeric `.x` and `.y` members. The returned circle always uses `double` coordinates, so integer-coordinate inputs are accepted.

Time Complexity: $O(n)$ expected time per call, where n is the distance between `lo` and `hi`, or $O(n^3)$ in the worst-case for a particular shuffle order.

Space Complexity: $O(1)$ auxiliary.

```

#include <algorithm>
#include <cmath>
#include <random>
#include <stdexcept>
#include <type_traits>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U>
bool LE(T a, U b) {
    return !LT(b, a);
}

double sqnorm(double x, double y) {
    return x * x + y * y;
}

```

```

// SFINAE guard: valid only when Pt exposes numeric .x/.y members. This keeps the templated point
// constructors from hijacking calls like Circle(double, double, double).
template<typename Pt>
using if_point = decltype(std::declval<const Pt &>().x, std::declval<const Pt &>().y, void());

struct Circle {
    double h, k, r;

    Circle() : h(0), k(0), r(0) {}
    Circle(double h, double k, double r) : h(h), k(k), r(fabs(r)) {}

    template<typename Pt, typename = if_point<Pt>>
    Circle(const Pt &a, const Pt &b) {
        h = (a.x + b.x) / 2.0;
        k = (a.y + b.y) / 2.0;
        r = std::hypot(a.x - h, a.y - k);
    }

    template<typename Pt, typename = if_point<Pt>>
    Circle(const Pt &a, const Pt &b, const Pt &c) {
        double an = sqnorm(b.x - c.x, b.y - c.y);
        double bn = sqnorm(a.x - c.x, a.y - c.y);
        double cn = sqnorm(a.x - b.x, a.y - b.y);
        double wa = an * (bn + cn - an), wb = bn * (an + cn - bn), wc = cn * (an + bn - cn);
        double w = wa + wb + wc;
        if (EQ(w, 0)) {
            throw std::runtime_error("No circumcircle from collinear points.");
        }
        h = (wa * a.x + wb * b.x + wc * c.x) / w;
        k = (wa * a.y + wb * b.y + wc * c.y) / w;
        r = std::hypot(a.x - h, a.y - k);
    }

    template<typename Pt, typename = if_point<Pt>>
    bool contains(const Pt &p) const {
        return LE(sqnorm(p.x - h, p.y - k), r * r);
    }
};

// Input points can be any type with .x and .y fields; Circle output is always double.
template<typename It>
Circle minimum_enclosing_circle(It lo, It hi) {
    if (lo == hi) {
        return Circle(0, 0, 0);
    }
    if (lo + 1 == hi) {
        return Circle(lo->x, lo->y, 0);
    }
    static std::mt19937 rng(std::random_device{}());
    std::shuffle(lo, hi, rng);
    Circle res(*lo, *(lo + 1));
    for (It i = lo + 2; i != hi; ++i) {
        if (res.contains(*i)) {
            continue;
        }
        res = Circle(*lo, *i);
        for (It j = lo + 1; j != i; ++j) {
            if (res.contains(*j)) {
                continue;
            }
            res = Circle(*i, *j);
            for (It k = lo; k != j; ++k) {
                if (!res.contains(*k)) {
                    res = Circle(*i, *j, *k);
                }
            }
        }
    }
}

```

```
    return res;
}
```

Example Usage

```
#include <cassert>
#include <vector>
using namespace std;

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
};

int main() {
    vector<Point> empty;
    Circle none = minimum_enclosing_circle(empty.begin(), empty.end());
    assert(EQ(none.h, 0) && EQ(none.k, 0) && EQ(none.r, 0));

    vector<Point> one{{2, 3}};
    Circle single = minimum_enclosing_circle(one.begin(), one.end());
    assert(EQ(single.h, 2) && EQ(single.k, 3) && EQ(single.r, 0));

    vector<Point> two{{0, 0}, {4, 0}};
    Circle diameter = minimum_enclosing_circle(two.begin(), two.end());
    assert(EQ(diameter.h, 2) && EQ(diameter.k, 0) && EQ(diameter.r, 2));

    vector<Point> v{{0, 0}, {0, 1}, {1, 0}, {1, 1}};
    Circle res = minimum_enclosing_circle(v.begin(), v.end());
    assert(EQ(res.h, 0.5) && EQ(res.k, 0.5) && EQ(res.r, 1 / sqrt(2)));

    // Integer-coordinate input: Circle output is always double.
    vector<PointI> iv{{0, 0}, {0, 2}, {2, 0}, {2, 2}};
    Circle ic = minimum_enclosing_circle(iv.begin(), iv.end());
    assert(EQ(ic.h, 1.0) && EQ(ic.k, 1.0));
    return 0;
}
```

7.3.6 Closest Pair

Given a list of points in two dimensions, finds the closest pair using a divide-and-conquer algorithm. The points are split in half by x -coordinate and each half is solved recursively; the combining step then only needs to examine points within the best distance so far of the dividing line, where each point in this strip is compared to a constant number of y -ordered neighbors.

- `closest_pair(lo, hi, res = nullptr)` returns the minimum squared distance between any two points in the range $[lo, hi)$, where `lo` and `hi` must be random-access iterators. The input range is reordered during the computation. If `res` is non-null, one closest pair is stored there in lexicographic order. With fewer than two points, the maximum value of the squared-distance type is returned and `res` is unchanged. The function is templated on the point type and works with any type exposing numeric `.x` and `.y` members and a lexicographic `operator<`.

The returned distance preserves the coordinate arithmetic type. For integer-coordinate inputs, the result is therefore an exact squared distance provided intermediate products do not overflow. The returned pair contains the original point type.

Overflow warning: squared distances grow like the square of the coordinate magnitude. For integer point types, use a 64-bit coordinate type (e.g. `PointL` from 7.1.1) when coordinates may exceed a few tens of thousands.

Time Complexity: $O(n \log n)$ per call, where n is the distance between `lo` and `hi`.

Space Complexity:

- $O(n)$ auxiliary heap space.
- $O(\log n)$ auxiliary stack space.

```
#include <algorithm>
#include <cmath>
#include <iterator>
#include <limits>
#include <utility>
#include <vector>

template<typename Pt>
auto sqdist(const Pt &a, const Pt &b) {
    auto dx = a.x - b.x, dy = a.y - b.y;
    return dx * dx + dy * dy; // Overflow warning!
}

template<typename It, typename T, typename Pt = typename std::iterator_traits<It>::value_type>
void closest_pair_rec(It lo, It hi, std::vector<Pt> &tmp, T &best, std::pair<Pt, Pt> *res) {
    int n = hi - lo;
    if (n <= 3) {
        for (It i = lo; i != hi; ++i) {
            for (It j = i + 1; j != hi; ++j) {
                T d = sqdist(*i, *j);
                if (d < best) {
                    best = d;
                    if (res != nullptr) {
                        *res = std::minmax(*i, *j);
                    }
                }
            }
        }
    }
    std::sort(lo, hi, [](const Pt &a, const Pt &b) { return a.y != b.y ? a.y < b.y : a.x < b.x; });
    return;
}

It mid = lo + n / 2;
auto midx = mid->x;
closest_pair_rec(lo, mid, tmp, best, res);
closest_pair_rec(mid, hi, tmp, best, res);
// Each half is now y-sorted, so merge them before examining the center strip.
auto by_y = [](const Pt &a, const Pt &b) { return a.y != b.y ? a.y < b.y : a.x < b.x; };
tmp.clear();
std::merge(lo, mid, mid, hi, std::back_inserter(tmp), by_y);
std::move(tmp.begin(), tmp.end(), lo);
tmp.clear();
for (It it = lo; it != hi; ++it) {
    auto dx = it->x - midx;
    if (dx * dx < best) {
        tmp.push_back(*it);
    }
}
for (int i = 0; i < static_cast<int>(tmp.size()); i++) {
    for (int j = i + 1; j < static_cast<int>(tmp.size()); j++) {
        auto dy = tmp[j].y - tmp[i].y;
        if (dy * dy >= best) {
            break;
        }
    }
    T d = sqdist(tmp[i], tmp[j]);
    if (d < best) {
        best = d;
        if (res != nullptr) {
            *res = std::minmax(tmp[i], tmp[j]);
        }
    }
}
```

```

    }
    }
}
tmp.clear();
}

template<typename It, typename Pt = typename std::iterator_traits<It>::value_type>
auto closest_pair(It lo, It hi, std::pair<Pt, Pt> *res = nullptr) {
    using T = decltype(sqdist(*lo, *lo));
    T best = std::numeric_limits<T>::max();
    std::sort(lo, hi, [](const Pt &a, const Pt &b) { return a.x != b.x ? a.x < b.x : a.y < b.y; });
    std::vector<Pt> tmp;
    tmp.reserve(hi - lo);
    closest_pair_rec(lo, hi, tmp, best, res);
    return best;
}

```

Example Usage

```

#include <cassert>
#include <vector>
using namespace std;

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-9;
}

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    bool operator<(const Point &p) const { return x != p.x ? x < p.x : y < p.y; }
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
    bool operator==(const PointI &p) const { return x == p.x && y == p.y; }
    bool operator<(const PointI &o) const { return x != o.x ? x < o.x : y < o.y; }
};

int main() {
    vector<Point> v{{2, 3}, {12, 30}, {40, 50}, {5, 1}, {12, 10}, {3, 4}};
    pair<Point, Point> res;
    assert(EQ(closest_pair(v.begin(), v.end(), &res), 2));
    auto [p1, p2] = res;
    assert(p1 == Point(2, 3) && p2 == Point(3, 4));

    // Integer-coordinate input: exact pair selection, int squared distance returned.
    vector<PointI> iv{{0, 0}, {10, 10}, {3, 4}};
    pair<PointI, PointI> ires;
    assert(closest_pair(iv.begin(), iv.end(), &ires) == 25);
    auto [i1, i2] = ires;
    assert(i1 == PointI(0, 0) && i2 == PointI(3, 4));
    return 0;
}

```

7.3.7 Multiple Segment Intersection

Given a list of line segments in two dimensions, determine whether any pair of segments intersect using a sweep line algorithm. Endpoint events are processed from left to right while an ordered set maintains the segments

currently crossing the sweep line, sorted by y ; only segments that become adjacent in this set need to be tested, since any leftmost intersection must involve a pair that is adjacent just before it occurs. The input type `Segment<Pt>` is templated on the point type, so endpoints may be integer (`PointI`) or floating-point (`Point` or `PointD`), but must support `operator<` which orders points lexicographically. The cross-product sign tests in `seg_intersection` are exact for integer endpoints, so intersection detection is exact.

- `find_intersection(lo, hi, &res1, &res2)` returns whether any pair of segments intersect given a range $[lo, hi)$ of segments, where `lo` and `hi` are random-access iterators. If an intersection is found, then one such pair of segments will be stored into pointers `res1` and `res2`.

This implementation counts endpoint contacts as intersections. See 7.2.3 for pairwise intersection predicates when that distinction matters.

Overflow warning: `seg_intersection` forms the usual quadratic cross products, but the y -ordering cross-multiplication (`ay * bdx`) is cubic in the coordinate magnitude. With 32-bit `int` endpoints this overflows once coordinates reach the low thousands, so use a 64-bit (`int64_t`) coordinate type for any non-trivial integer inputs.

Time Complexity: $O(n \log n)$ per call, where n is the distance between `lo` and `hi`.

Space Complexity: $O(n)$ auxiliary, where n is the distance between `lo` and `hi`.

```
#include <algorithm>
#include <cassert>
#include <set>
#include <type_traits>
#include <utility>
#include <vector>

// Simplified detection-only version of seg_intersection() from 7.2.3 (exact for integral Pt).
template<typename Pt>
bool seg_intersection(const Pt &a, const Pt &b, const Pt &c, const Pt &d) {
    auto cross = [](const Pt &p, const Pt &q, const Pt &r) {
        return (q.x - p.x) * (r.y - p.y) - (q.y - p.y) * (r.x - p.x);
    };
    auto c1 = cross(a, b, c), c2 = cross(a, b, d);
    auto c3 = cross(c, d, a), c4 = cross(c, d, b);
    if (c1 == 0 && c2 == 0 && c3 == 0 && c4 == 0) {
        Pt p = std::max(std::min(a, b), std::min(c, d));
        Pt q = std::min(std::max(a, b), std::max(c, d));
        return !(q < p);
    }
    auto straddles = [](auto x, auto y) { return (x <= 0 && 0 <= y) || (y <= 0 && 0 <= x); };
    return straddles(c1, c2) && straddles(c3, c4);
}

// The point type Pt must support operator< (used to canonicalize endpoint order).
template<typename Pt>
struct Segment {
    Pt p, q;
    Segment() {}
    Segment(const Pt &p, const Pt &q) : p(std::min(p, q)), q(std::max(p, q)) {}
};

template<typename Pt>
bool intersects(const Segment<Pt> &s1, const Segment<Pt> &s2) {
    return seg_intersection(s1.p, s1.q, s2.p, s2.q);
}

// Seg is deduced from the output pointers (e.g. Segment<PointI>* or Segment<PointD>*).
template<typename It, typename Seg>
bool find_intersection(It lo, It hi, Seg *res1, Seg *res2) {
    using Pt = std::decay_t<decltype(lo->p)>;
    struct Event {
        Pt p;
        int type;
    };
    std::vector<Event> v;
    for (It it = lo; it != hi; ++it) {
        v.push_back({it->p, 1});
    }
    v.push_back({hi->p, 2});
    std::sort(v.begin(), v.end());
    std::set<Segment<Pt>> s;
    for (const Event &ev : v) {
        if (ev.type == 1) {
            s.insert(Segment<Pt>{ev.p, it->q});
        } else {
            auto it2 = s.lower_bound(Segment<Pt>{ev.p, ev.p});
            if (it2 != s.end()) {
                *res1 = *it2;
                *res2 = *it2;
                return true;
            }
            s.erase(it2);
        }
    }
    return false;
}
```

```

    It seg;
};
std::vector<Event> e;
for (It it = lo; it != hi; ++it) {
    if (it->p > it->q) {
        std::swap(it->p, it->q);
    }
    e.push_back(Event{it->p, 1, it});
    e.push_back(Event{it->q, -1, it});
}
std::sort(e.begin(), e.end(), [lo](const Event &a, const Event &b) {
    if (a.p.x != b.p.x) {
        return a.p.x < b.p.x;
    }
    if (a.type != b.type) {
        return b.type < a.type;
    }
    return a.p.y != b.p.y ? a.p.y < b.p.y : a.seg - lo < b.seg - lo;
});
// Compare y-values at sweep coordinate x without division: y = (p.y*dx + dy*(x - p.x)) / dx.
// Vertical segments use their lower endpoint.
auto ycmp = [](const auto &a, const auto &b, auto x) {
    // Overflow risk for integer Pt: the final cross-multiply is ~0(max_coord^3); use 64-bit
    // coordinates for non-trivial integer inputs.
    auto adx = a.q.x - a.p.x, bdx = b.q.x - b.p.x;
    auto ay = a.p.y * adx + (a.q.y - a.p.y) * (x - a.p.x);
    auto by = b.p.y * bdx + (b.q.y - b.p.y) * (x - b.p.x);
    if (adx == 0) {
        ay = a.p.y;
        adx = 1;
    }
    if (bdx == 0) {
        by = b.p.y;
        bdx = 1;
    }
    auto lhs = ay * bdx, rhs = by * adx;
    return (lhs < rhs) ? -1 : ((rhs < lhs) ? 1 : 0);
};
auto cmp = [lo, ycmp](It a, It b) {
    if (a == b) {
        return false;
    }
    auto x = std::max(a->p.x, b->p.x);
    int order = ycmp(*a, *b, x);
    return (order == 0) ? (a - lo < b - lo) : (order < 0);
};
using ActiveSet = std::set<It, decltype(cmp)>;
ActiveSet s(cmp);
std::vector<typename ActiveSet::iterator> position(hi - lo);
for (const auto &ev : e) {
    It seg = ev.seg;
    if (ev.type == 1) {
        auto it = s.insert(seg).first;
        position[seg - lo] = it;
        auto next = it;
        if (++next != s.end() && intersects(**next, *seg)) {
            *res1 = **next;
            *res2 = *seg;
            return true;
        }
        if (it != s.begin() && intersects(**--it, *seg)) {
            *res1 = **it;
            *res2 = *seg;
            return true;
        }
    } else {
        auto it = position[seg - lo];
        auto next = it;

```

```

    if (++next != s.end() && it != s.begin()) {
        auto prev = it;
        if (intersects(**next, **--prev)) {
            *res1 = **next;
            *res2 = **prev;
            return true;
        }
    }
    s.erase(it);
}
}
return false;
}

```

Example Usage

```

#include <vector>
using namespace std;

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    bool operator!=(const Point &p) const { return !(*this == p); }
    bool operator<(const Point &p) const { return x != p.x ? x < p.x : y < p.y; }
    bool operator>(const Point &p) const { return p < *this; }
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
    bool operator==(const PointI &p) const { return x == p.x && y == p.y; }
    bool operator!=(const PointI &p) const { return !(*this == p); }
    bool operator<(const PointI &p) const { return x != p.x ? x < p.x : y < p.y; }
    bool operator>(const PointI &p) const { return p < *this; }
};

int main() {
    { // Double-coordinate segments.
        vector<Segment<Point>> v{
            Segment<Point>(Point(0, 0), Point(2, 2)), Segment<Point>(Point(3, 0), Point(0, -1)),
            Segment<Point>(Point(0, 2), Point(2, -2)), Segment<Point>(Point(0, 3), Point(9, 0))
        };
        Segment<Point> res1, res2;
        assert(find_intersection(v.begin(), v.end(), &res1, &res2));
        assert(res1.p == Point(0, 0) && res1.q == Point(2, 2));
        assert(res2.p == Point(0, 2) && res2.q == Point(2, -2));
    }
    { // Integer-coordinate segments: detection is exact.
        vector<Segment<PointI>> v{
            Segment<PointI>({0, 0}, {2, 2}), Segment<PointI>({3, 0}, {0, -1}),
            Segment<PointI>({0, 2}, {2, -2}), Segment<PointI>({0, 3}, {9, 0})
        };
        Segment<PointI> res1, res2;
        assert(find_intersection(v.begin(), v.end(), &res1, &res2));

        vector<Segment<PointI>> disjoint{
            Segment<PointI>({0, 0}, {1, 0}), Segment<PointI>({0, 5}, {1, 5})
        };
        assert(!find_intersection(disjoint.begin(), disjoint.end(), &res1, &res2));
    }
    { // Shared endpoints count as intersections.
        vector<Segment<PointI>> shared{
            Segment<PointI>({0, 0}, {2, 2}), Segment<PointI>({2, 2}, {4, 0})
        };
        Segment<PointI> res1, res2;
    }
}

```



```

    assert(find_intersection(shared.begin(), shared.end(), &res1, &res2));
}
return 0;
}

```

7.3.8 Maximum Collinear Points

Given a set of points with integer coordinates, compute the maximum number of points that lie on one line. For every anchor point, normalize the direction vector to every other point by dividing by `gcd(abs(dx), abs(dy))` and choosing a canonical sign; equal normalized directions from the same anchor are collinear with it.

- `max_collinear_points(p)` returns the largest number of collinear points. Duplicate points are counted as separate input points.

Overflow warning: coordinate differences and their absolute values must fit in `int64_t`.

Time Complexity: $O(n^2 \log n)$ per call, where n is the number of points.

Space Complexity: $O(n)$ auxiliary.

```

#include <algorithm>
#include <cassert>
#include <cmath>
#include <cstdint>
#include <map>
#include <numeric>
#include <utility>
#include <vector>

struct Point {
    int64_t x, y;
    Point(int64_t x = 0, int64_t y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
};

int max_collinear_points(const std::vector<Point> &p) {
    int n = static_cast<int>(p.size()), ans = std::min(n, 1);
    for (int i = 0; i < n; i++) {
        std::map<std::pair<int64_t, int64_t>, int> cnt;
        int duplicates = 0, best = 0;
        for (int j = 0; j < n; j++) {
            if (i == j) {
                continue;
            }
            int64_t dx = p[j].x - p[i].x, dy = p[j].y - p[i].y; // Overflow warning!
            if (dx == 0 && dy == 0) {
                duplicates++;
                continue;
            }
            int64_t g = std::gcd(std::abs(dx), std::abs(dy));
            dx /= g;
            dy /= g;
            if (dx < 0 || (dx == 0 && dy < 0)) {
                dx = -dx;
                dy = -dy;
            }
            best = std::max(best, ++cnt[{dx, dy}]);
        }
        ans = std::max(ans, best + duplicates + 1);
    }
    return ans;
}

```

Example Usage

```
using namespace std;

int main() {
    vector<Point> p{{0, 0}, {1, 1}, {2, 2}, {3, 5}, {4, 6}, {3, 3}};
    assert(max_collinear_points(p) == 4);
    p.push_back(Point(1, 1));
    assert(max_collinear_points(p) == 5);
    assert(max_collinear_points({}) == 0);
    assert(max_collinear_points({Point(2, 3)}) == 1);
    return 0;
}
```

7.3.9 Manhattan MST

Given points in the plane, produce a small set of candidate edges guaranteed to contain a minimum spanning tree under Manhattan distance $|x_1 - x_2| + |y_1 - y_2|$. The plane is swept in four rotations/reflections; in each sweep, dominance by $x + y$ identifies the only nearby candidates that can matter. Run Kruskal on the returned edges to obtain the MST.

- `manhattan_mst_edges(p)` returns $O(n)$ candidate edges as tuples `(weight, u, v)`.
- `manhattan_mst_weight(p)` is a small Kruskal wrapper that returns the MST weight.

Overflow warning: the sweeps form coordinate sums, negations, and differences such as `x + y` and `dx + dy`. With `int64_t` coordinates, those intermediate values must fit in `int64_t`.

Time Complexity: $O(n \log n)$ per call, where n is the number of points.

Space Complexity: $O(n)$ auxiliary.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <map>
#include <numeric>
#include <tuple>
#include <utility>
#include <vector>

struct Point {
    int64_t x, y;
    Point(int64_t x = 0, int64_t y = 0) : x(x), y(y) {}
};

std::vector<std::tuple<int64_t, int, int>> manhattan_mst_edges(std::vector<Point> p) {
    int n = static_cast<int>(p.size());
    std::vector<int> id(n);
    std::iota(id.begin(), id.end(), 0);
    std::vector<std::tuple<int64_t, int, int>> edges;
    for (int rot = 0; rot < 4; rot++) {
        // Overflow warning!
        std::sort(id.begin(), id.end(), [&](int i, int j) {
            return p[i].x + p[i].y < p[j].x + p[j].y;
        });
        std::map<int64_t, int> sweep;
        for (int i : id) {
            for (auto it = sweep.lower_bound(-p[i].y); it != sweep.end(); it) {
                int j = it->second;
                int64_t dx = p[i].x - p[j].x, dy = p[i].y - p[j].y;
                if (dy > dx) {
                    break;
                }
            }
        }
    }
}
```

```

        edges.emplace_back(dx + dy, i, j);
        it = sweep.erase(it);
    }
    sweep[-p[i].y] = i;
}
for (Point &q : p) {
    if (rot & 1) {
        q.x = -q.x;
    } else {
        std::swap(q.x, q.y);
    }
}
}
return edges;
}

struct DSU {
    std::vector<int> parent, rank;

    explicit DSU(int n) : parent(n), rank(n, 0) { std::iota(parent.begin(), parent.end(), 0); }
    int find(int u) { return parent[u] == u ? u : parent[u] = find(parent[u]); }

    bool unite(int u, int v) {
        u = find(u);
        v = find(v);
        if (u == v) {
            return false;
        }
        if (rank[u] < rank[v]) {
            std::swap(u, v);
        }
        parent[v] = u;
        rank[u] += rank[v] == rank[v];
        return true;
    }
};

int64_t manhattan_mst_weight(const std::vector<Point> &p) {
    auto edges = manhattan_mst_edges(p);
    std::sort(edges.begin(), edges.end());
    DSU dsu(static_cast<int>(p.size()));
    int64_t res = 0;
    for (auto [w, u, v] : edges) {
        if (dsu.unite(u, v)) {
            res += w;
        }
    }
    return res;
}

```

Example Usage

```

using namespace std;

int main() {
    assert(manhattan_mst_weight({Point(5, -7)}) == 0);
    assert(manhattan_mst_weight({Point(0, 0), Point(3, -4)}) == 7);

    vector<Point> p{{0, 0}, {2, 1}, {3, 4}, {-1, 2}, {5, 0}};
    assert(manhattan_mst_weight(p) == 14);
    auto edges = manhattan_mst_edges(p);
    assert(edges.size() <= 4 * p.size());
    return 0;
}

```

7.4 Advanced Planar Geometry

7.4.1 Convex Polygon Cut

Given a convex polygon and a directed line from `p` to `q`, clips the polygon to the closed left half-plane of that line. The polygon is walked edge by edge: vertices on the left side of the line are kept, and whenever an edge crosses the line, the intersection point is appended to the output.

- `convex_cut(lo, hi, p, q)` returns the portion of the polygon lying on or to the left of the directed line from `p` to `q`. The input range `[lo, hi)` must contain the vertices of a convex polygon in boundary order, either clockwise or counterclockwise. The returned polygon preserves that boundary order. The points `p` and `q` must differ.

The function is templated on the input point type. Side classification is done with cross products. For integer-coordinate inputs, classification is exact only if the intermediate products do not overflow. Edge-line intersection points are computed in floating point, and the returned polygon uses `Point` with `double` coordinates.

Time Complexity: $O(n)$ per call, where n is the distance between `lo` and `hi`.

Space Complexity: $O(n)$ for the returned polygon and $O(1)$ auxiliary.

```
#include <cassert>
#include <cmath>
#include <type_traits>
#include <vector>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
};

template<typename PtA, typename PtB, typename PtO>
auto cross(const PtA &a, const PtB &b, const PtO &o) {
    return (a.x - o.x) * (b.y - o.y) - (a.y - o.y) * (b.x - o.x);
}

template<typename PtA, typename PtO, typename PtB>
int turn(const PtA &a, const PtO &o, const PtB &b) {
    auto c = cross(a, b, o);
    return LT(c, 0) ? 1 : (LT(0, c) ? -1 : 0);
}

template<typename PtA, typename PtB>
int line_intersection(
    const PtA &p1, const PtA &p2, const PtB &p3, const PtB &p4, Point *p = nullptr
) {
    double a1 = static_cast<double>(p2.y) - p1.y, b1 = static_cast<double>(p1.x) - p2.x;
    double c1 = -(static_cast<double>(p1.x) * p2.y - static_cast<double>(p2.x) * p1.y);
    double a2 = static_cast<double>(p4.y) - p3.y, b2 = static_cast<double>(p3.x) - p4.x;
```

```

double c2 = -(static_cast<double>(p3.x) * p4.y - static_cast<double>(p4.x) * p3.y);
double x = -(c1 * b2 - c2 * b1), y = -(a1 * c2 - a2 * c1);
double det = a1 * b2 - a2 * b1;
if (EQ(det, 0)) {
    return (EQ(x, 0) && EQ(y, 0)) ? 1 : -1;
}
if (p != nullptr) {
    *p = Point(x / det, y / det);
}
return 0;
}

template<typename It, typename Pt>
std::vector<Point> convex_cut(It lo, It hi, const Pt &p, const Pt &q) {
    assert(!EQ(p.x, q.x) || !EQ(p.y, q.y));
    if (lo == hi) {
        return {};
    }
    std::vector<Point> res;
    for (It i = lo, j = hi - 1; i != hi; j = i++) {
        Point pj(static_cast<double>(j->x), static_cast<double>(j->y));
        Point pi(static_cast<double>(i->x), static_cast<double>(i->y));
        int d1 = turn(q, p, *j), d2 = turn(q, p, *i);
        if (d1 <= 0) {
            res.push_back(pj);
        }
        if (d1 * d2 < 0) {
            Point r;
            line_intersection(p, q, *j, *i, &r);
            res.push_back(r);
        }
    }
    return res;
}

```

Example Usage

```

using namespace std;

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
};

bool EQ(const vector<Point> &a, const vector<Point> &b) {
    return a.size() == b.size() &&
        equal(a.begin(), a.end(), b.begin(), [](const Point &p, const Point &q) {
            return EQ(p, q);
        });
}

int main() {
    {
        vector<Point> v{{1, 3}, {2, 2}, {2, 1}, {0, 0}, {-1, 3}};
        // Cut using the vertical line through (0, 0).
        vector<Point> c{{-1, 3}, {0, 3}, {0, 0}};
        vector<Point> cut = convex_cut(v.begin(), v.end(), Point(0, 0), Point(0, 1));
        assert(EQ(cut, c));
    }
    { // On a non-convex input, the result may be multiple disjoint polygons!
        vector<Point> v{{0, 0}, {2, 2}, {0, 4}, {3, 4}, {3, 0}};
        vector<Point> c{{1, 0}, {0, 0}, {1, 1}, {1, 3}, {0, 4}, {1, 4}};
        assert(EQ(convex_cut(v.begin(), v.end(), Point(1, 0), Point(1, 4)), c));
    }
    {
        vector<PointI> v{{0, 0}, {4, 0}, {4, 4}, {0, 4}};
    }
}

```

```

    vector<Point> c{{0, 4}, {0, 0}, {2, 0}, {2, 4}};
    assert(EQ(convex_cut(v.begin(), v.end(), PointI(2, 0), PointI(2, 4)), c));
}
return 0;
}

```

7.4.2 Half-Plane Intersection

Given a list of directed lines, compute the convex polygon contained in every left half-plane. Half-plane intersection sorts the lines by direction, keeps only the tightest representative among parallel lines, and maintains a deque of lines whose pairwise intersections form the current boundary. Directions are ordered by half-plane and cross-product signs, without epsilon comparisons, to satisfy `std::sort`'s strict ordering requirement. `EPS` is then used outside the comparator to group parallel directions and retain only the tightest half-plane in each group.

- `half_plane_intersection(planes)` returns the polygon cut out by the half-planes in counter-clockwise order. Each half-plane is represented by `HalfPlane(p, q)`, meaning the closed region to the left of the directed line $p \rightarrow q$. This implementation is intended for intersections that form a bounded polygon with positive area; empty, unbounded, or degenerate inputs may return an empty vector. Add explicit bounding-box half-planes when a bounded polygon is required. The two points defining each half-plane must differ.

Time Complexity: $O(n \log n)$ per call, where n is the number of half-planes.

Space Complexity: $O(n)$ auxiliary.

```

#include <algorithm>
#include <cmath>
#include <deque>
#include <type_traits>
#include <vector>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
    Point operator+(const Point &p) const { return {x + p.x, y + p.y}; }
    Point operator-(const Point &p) const { return {x - p.x, y - p.y}; }
    Point operator*(double k) const { return {x * k, y * k}; }
    double cross(const Point &p) const { return x * p.y - y * p.x; }
};

struct HalfPlane {
    Point p, dir;

    HalfPlane() : p(), dir() {}
    HalfPlane(Point p, Point q) : p(p), dir(q - p) {}

    bool out(const Point &q) const { return LT(dir.cross(q - p), 0); }
}

```

```

Point intersection(const HalfPlane &h) const {
    double t = h.dir.cross(p - h.p) / dir.cross(h.dir);
    return p + dir * t;
}

bool operator<(const HalfPlane &h) const {
    auto half = [](const Point &d) { return d.y < 0 || (d.y == 0 && d.x < 0); };
    if (half(dir) != half(h.dir)) {
        return half(dir) < half(h.dir);
    }
    return dir.cross(h.dir) > 0;
}
};

std::vector<Point> half_plane_intersection(std::vector<HalfPlane> planes) {
    std::sort(planes.begin(), planes.end());
    std::vector<HalfPlane> unique;
    for (const HalfPlane &h : planes) {
        if (!unique.empty() && EQ(unique.back().dir.cross(h.dir), 0)) {
            if (unique.back().dir.cross(h.p - unique.back().p) > 0) {
                unique.back() = h;
            }
            continue;
        }
        unique.push_back(h);
    }
    std::deque<HalfPlane> dq;
    auto bad_back = [&](const HalfPlane &h) {
        return dq.size() >= 2 && h.out(dq[dq.size() - 2].intersection(dq.back()));
    };
    auto bad_front = [&](const HalfPlane &h) {
        return dq.size() >= 2 && h.out(dq[0].intersection(dq[1]));
    };
    for (const HalfPlane &h : unique) {
        while (bad_back(h)) {
            dq.pop_back();
        }
        while (bad_front(h)) {
            dq.pop_front();
        }
        dq.push_back(h);
    }
    while (dq.size() >= 3 && dq.front().out(dq[dq.size() - 2].intersection(dq.back()))) {
        dq.pop_back();
    }
    while (dq.size() >= 3 && dq.back().out(dq[0].intersection(dq[1]))) {
        dq.pop_front();
    }
    if (dq.size() < 3) {
        return {};
    }
    std::vector<Point> res;
    for (int i = 0; i < static_cast<int>(dq.size()); i++) {
        if (EQ(dq[i].dir.cross(dq[(i + 1) % dq.size()].dir), 0)) {
            return {};
        }
        res.push_back(dq[i].intersection(dq[(i + 1) % dq.size()]));
    }
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

```

```

double polygon_area(const vector<Point> &p) {
    double area = 0;
    for (int i = 0, j = static_cast<int>(p.size()) - 1; i < static_cast<int>(p.size()); j = i++) {
        area += p[j].cross(p[i]);
    }
    return fabs(area) / 2.0;
}

int main() {
    vector<HalfPlane> box{
        HalfPlane(Point(0, 0), Point(4, 0)),
        HalfPlane(Point(4, 0), Point(4, 4)),
        HalfPlane(Point(4, 4), Point(0, 4)),
        HalfPlane(Point(0, 4), Point(0, 0)),
    };
    auto square = half_plane_intersection(box);
    assert(square.size() == 4);
    assert(EQ(polygon_area(square), 16));

    auto tighter = box;
    tighter.push_back(HalfPlane(Point(0, 1), Point(4, 1))); // y >= 1
    assert(EQ(polygon_area(half_plane_intersection(tighter)), 12));

    box.push_back(HalfPlane(Point(2, 0), Point(2, 4))); // x <= 2
    auto rect = half_plane_intersection(box);
    assert(rect.size() == 4);
    assert(EQ(polygon_area(rect), 8));

    vector<HalfPlane> empty{
        HalfPlane(Point(0, 0), Point(1, 0)), // y >= 0
        HalfPlane(Point(1, 1), Point(0, 1)), // y <= 1
        HalfPlane(Point(0, 2), Point(1, 2)), // y >= 2
    };
    assert(half_plane_intersection(empty).empty());
    return 0;
}

```

7.4.3 Minkowski Sum

Given two convex polygons, compute their Minkowski sum: the set of all points $a + b$ where a comes from the first polygon and b comes from the second. After rotating both polygons to start at their lowest-leftmost vertex, the boundary edges of the sum are obtained by merging the two cyclic edge-angle lists. Minkowski sums are useful for collision/translation problems, convex polygon distance checks, and diameter computations via $P + (-P)$.

- `minkowski_sum(a, b)` returns the convex polygon representing the sum. Inputs must be convex polygons in boundary order, clockwise or counter-clockwise, with no repeated final vertex. Collinear consecutive output vertices are removed.

Time Complexity: $O(n + m)$ per call, where n and m are the polygon sizes.

Space Complexity: $O(n + m)$ auxiliary.

```

#include <algorithm>
#include <cmath>
#include <type_traits>
#include <utility>
#include <vector>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

```



```

}

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
    Point operator+(const Point &p) const { return {x + p.x, y + p.y}; }
    Point operator-(const Point &p) const { return {x - p.x, y - p.y}; }
    double cross(const Point &p) const { return x * p.y - y * p.x; }
};

double polygon_area_2x(const std::vector<Point> &p) {
    double area = 0;
    for (int i = 0, j = static_cast<int>(p.size()) - 1; i < static_cast<int>(p.size()); j = i++) {
        area += p[j].cross(p[i]);
    }
    return area;
}

std::vector<Point> normalize(std::vector<Point> p) {
    while (p.size() > 1 && EQ(p.front(), p.back())) {
        p.pop_back();
    }
    if (p.size() <= 1) {
        return p;
    }
    if (polygon_area_2x(p) < 0) {
        std::reverse(p.begin(), p.end());
    }
    int start = 0;
    for (int i = 1; i < static_cast<int>(p.size()); i++) {
        if (p[i].y < p[start].y || (EQ(p[i].y, p[start].y) && p[i].x < p[start].x)) {
            start = i;
        }
    }
    std::rotate(p.begin(), p.begin() + start, p.end());
    return p;
}

std::vector<Point> remove_collinear(std::vector<Point> p) {
    std::vector<Point> res;
    for (const Point &q : p) {
        while (res.size() >= 2 && EQ((res.back() - res[res.size() - 2]).cross(q - res.back()), 0)) {
            res.pop_back();
        }
        if (res.empty() || !EQ(q, res.back())) {
            res.push_back(q);
        }
    }
    while (res.size() >= 3 && EQ((res.back() - res[res.size() - 2]).cross(res[0] - res.back()), 0)) {
        res.pop_back();
    }
    while (res.size() >= 3 && EQ((res[0] - res.back()).cross(res[1] - res[0]), 0)) {
        res.erase(res.begin());
    }
    return res;
}

std::vector<Point> minkowski_sum(std::vector<Point> a, std::vector<Point> b) {
    a = normalize(std::move(a));
    b = normalize(std::move(b));
    if (a.empty() || b.empty()) {
        return {};
    }
    if (a.size() == 1) {
        for (Point &p : b) {
            p = p + a[0];
        }
    }
}

```

```

    }
    return b;
}
if (b.size() == 1) {
    for (Point &p : a) {
        p = p + b[0];
    }
    return a;
}
int n = static_cast<int>(a.size()), m = static_cast<int>(b.size());
a.push_back(a[0]);
a.push_back(a[1]);
b.push_back(b[0]);
b.push_back(b[1]);
std::vector<Point> res;
int i = 0, j = 0;
while (i < n || j < m) {
    res.push_back(a[i] + b[j]);
    if (i == n) {
        j++;
        continue;
    }
    if (j == m) {
        i++;
        continue;
    }
    double cr = (a[i + 1] - a[i]).cross(b[j + 1] - b[j]);
    if (cr > -EPS) {
        i++;
    }
    if (cr < EPS) {
        j++;
    }
}
return remove_collinear(res);
}

```

Example Usage

```

#include <cassert>
using namespace std;

int main() {
    vector<Point> empty;
    assert(minkowski_sum(empty, {Point(1, 2)}).empty());

    vector<Point> square{{0, 0}, {1, 0}, {1, 1}, {0, 1}};
    vector<Point> tri{{0, 0}, {2, 0}, {0, 2}};
    vector<Point> sum = minkowski_sum(square, tri);
    // The square thickens the triangle by one unit in the positive x/y directions.
    vector<Point> expected{{0, 0}, {3, 0}, {3, 1}, {1, 3}, {0, 3}};
    assert(sum == expected);

    vector<Point> shifted = minkowski_sum(square, {Point(2, 3)});
    // Adding a single point translates the polygon.
    vector<Point> expected_shifted{{2, 3}, {3, 3}, {3, 4}, {2, 4}};
    assert(shifted == expected_shifted);

    vector<Point> point_sum = minkowski_sum({Point(1, 2)}, {Point(3, 4)});
    assert(point_sum.size() == 1 && point_sum[0] == Point(4, 6));
    return 0;
}

```

7.4.4 Polygon Intersection and Union

Given two simple (non-self-intersecting) polygons, determine the areas of their intersection and union using a sweep line algorithm and the inclusion-exclusion principle. Every vertex and every pairwise edge intersection contributes its x -coordinate as a sweep boundary, so within each vertical slab between consecutive boundaries the two borders cannot cross, and the slab's overlap is accumulated exactly from trapezoid areas. The union then follows by inclusion-exclusion as the sum of the individual areas minus the intersection.

- `intersection_area(lo1, hi1, lo2, hi2)` returns the intersection area of two polygons respectively specified by two ranges `[lo1, hi1)` and `[lo2, hi2)` of vertices in boundary order, where `lo1`, `hi1`, `lo2`, and `hi2` must be random-access iterators.
- `union_area(lo1, hi1, lo2, hi2)` returns the union area of two polygons respectively specified by two ranges `[lo1, hi1)` and `[lo2, hi2)` of vertices in boundary order, where `lo1`, `hi1`, `lo2`, and `hi2` must be random-access iterators.

Overflow warning: the segment-intersection tests form cross products in the input point's coordinate type, which grow like the squared coordinate magnitude. For integral point types, use 64-bit coordinates (e.g. `PointL` from 7.1.1) if necessary.

Time Complexity: $O(n \cdot m \cdot (n + m) \cdot \log(n + m))$ per call to `intersection_area(lo1, hi1, lo2, hi2)` and `union_area(lo1, hi1, lo2, hi2)`, where n is the number of vertices in the first polygon and m is the number of vertices in the second polygon (or $O(N^3 \log N)$ for $N = n + m$).

Space Complexity: $O(n \cdot m)$ auxiliary for `intersection_area(lo1, hi1, lo2, hi2)` and `union_area(lo1, hi1, lo2, hi2)`, where n and m are the respective polygon sizes.

```
#include <algorithm>
#include <cmath>
#include <set>
#include <type_traits>
#include <vector>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U>
bool LE(T a, U b) {
    return !LT(b, a);
}

template<typename Pt>
bool point_on_segment(const Pt &p, const Pt &a, const Pt &b) {
    // Overflow risk for integer Pt: these products are ~O(max_coord^2); use int64_t if necessary.
    return EQ((p.x - a.x) * (b.y - a.y) - (p.y - a.y) * (b.x - a.x), 0) &&
        LE(std::min(a.x, b.x), p.x) && LE(p.x, std::max(a.x, b.x)) &&
        LE(std::min(a.y, b.y), p.y) && LE(p.y, std::max(a.y, b.y));
}

// Specialized version of seg_intersection() from 7.2.3, simplified for touch_is_intersect = true,
// returning -1 for no intersection, 0 for one intersection point, 1 for positive length overlap.
template<typename Pt>
int seg_intersection1(
    const Pt &a, const Pt &b, const Pt &c, const Pt &d, double *outx, double *outy
```

```

) {
    auto ab_x = b.x - a.x, ab_y = b.y - a.y;
    auto ac_x = c.x - a.x, ac_y = c.y - a.y;
    auto cd_x = d.x - c.x, cd_y = d.y - c.y;
    auto ab2 = ab_x * ab_x + ab_y * ab_y;
    auto cd2 = cd_x * cd_x + cd_y * cd_y;
    if (EQ(ab2, 0)) {
        if (point_on_segment(a, c, d)) {
            *outx = a.x;
            *outy = a.y;
            return 0;
        }
        return -1;
    }
    if (EQ(cd2, 0)) {
        if (point_on_segment(c, a, b)) {
            *outx = c.x;
            *outy = c.y;
            return 0;
        }
        return -1;
    }
    auto c1 = ab_x * cd_y - ab_y * cd_x;
    auto c2 = ac_x * ab_y - ac_y * ab_x;
    if (EQ(c1, 0) && EQ(c2, 0)) { // Collinear.
        Pt res1 = std::max(std::min(a, b), std::min(c, d));
        Pt res2 = std::min(std::max(a, b), std::max(c, d));
        if (!(res2 < res1)) {
            if (res1 == res2) {
                *outx = static_cast<double>(res1.x);
                *outy = static_cast<double>(res1.y);
                return 0;
            }
            return 1;
        }
        return -1;
    }
    if (EQ(c1, 0)) {
        return -1;
    }
    auto t_num = ac_x * cd_y - ac_y * cd_x;
    bool c1_pos = c1 > 0;
    bool t_ok = c1_pos ? (LE(0, t_num) && LE(t_num, c1)) : (LE(c1, t_num) && LE(t_num, 0));
    bool u_ok = c1_pos ? (LE(0, c2) && LE(c2, c1)) : (LE(c1, c2) && LE(c2, 0));
    if (t_ok && u_ok) {
        double t = static_cast<double>(t_num) / c1;
        *outx = static_cast<double>(a.x + t * ab_x);
        *outy = static_cast<double>(a.y + t * ab_y);
        return 0;
    }
    return -1;
}

// The two segments may use different point types (e.g. polygon edge vs. sweep line).
template<typename PtA, typename PtB>
int line_intersection1(
    const PtA &p1, const PtA &p2, const PtB &p3, const PtB &p4, double *outx, double *outy
) {
    double a1 = static_cast<double>(p2.y) - p1.y, b1 = static_cast<double>(p1.x) - p2.x;
    double c1 = -(a1 * p1.x + b1 * p1.y);
    double a2 = static_cast<double>(p4.y) - p3.y, b2 = static_cast<double>(p3.x) - p4.x;
    double c2 = -(a2 * p3.x + b2 * p3.y);
    double x = -(c1 * b2 - c2 * b1), y = -(a1 * c2 - a2 * c1);
    double det = a1 * b2 - a2 * b1;
    if (EQ(det, 0)) {
        return (EQ(x, 0) && EQ(y, 0)) ? 1 : -1;
    }
    if (outx != nullptr && outy != nullptr) {

```

```

    *outx = x / det;
    *outy = y / det;
}
return 0;
}

struct Event {
    double y;
    int mask_delta;
    Event(double y = 0, int mask_delta = 0) : y(y), mask_delta(mask_delta) {}
};

template<typename It>
double intersection_area(It lo1, It hi1, It lo2, It hi2) {
    if (lo1 == hi1 || lo2 == hi2) {
        return 0;
    }
    struct _PointD {
        double x, y;
    }; // For line intersection with the sweep line.
    std::vector<It> plo{lo1, lo2}, phi{hi1, hi2};
    std::set<double> xs;
    for (It it = lo1; it != hi1; ++it) {
        xs.insert(it->x);
    }
    for (It it = lo2; it != hi2; ++it) {
        xs.insert(it->x);
    }
    for (It i1 = lo1, j1 = hi1 - 1; i1 != hi1; j1 = i1++) {
        for (It i2 = lo2, j2 = hi2 - 1; i2 != hi2; j2 = i2++) {
            double outx, outy;
            if (seg_intersection1(*i1, *j1, *i2, *j2, &outx, &outy) == 0) {
                xs.insert(outx);
            }
        }
    }
    std::vector<double> xsa(xs.begin(), xs.end());
    double res = 0;
    for (int k = 0; k + 1 < static_cast<int>(xsa.size()); k++) {
        double x = (xsa[k] + xsa[k + 1]) / 2;
        _PointD sweep0{x, 0}, sweep1{x, 1};
        std::vector<Event> events;
        for (int poly = 0; poly < 2; poly++) {
            It lo = plo[poly], hi = phi[poly];
            double area = 0;
            for (It i = lo, j = hi - 1; i != hi; j = i++) {
                area += (static_cast<double>(j->x) - i->x) * (static_cast<double>(j->y) + i->y);
            }
            for (It j = lo, i = hi - 1; j != hi; i = j++) {
                double px, py;
                if (line_intersection1(*j, *i, sweep0, sweep1, &px, &py) == 0) {
                    double y = py, x0 = i->x, x1 = j->x;
                    int sgn_area = (area < 0 ? 1 : (area > 0 ? -1 : 0));
                    if (x0 < x && x1 > x) {
                        events.emplace_back(y, sgn_area * (1 << poly));
                    } else if (x0 > x && x1 < x) {
                        events.emplace_back(y, -sgn_area * (1 << poly));
                    }
                }
            }
        }
    }
    std::sort(events.begin(), events.end(), [](const Event &e1, const Event &e2) {
        if (e1.y != e2.y) {
            return e1.y < e2.y;
        }
        return e1.mask_delta < e2.mask_delta;
    });
    double a = 0;

```

```

    int cnt[2] = {0, 0};
    for (int j = 0; j < static_cast<int>(events.size()); j++) {
        if (cnt[0] != 0 && cnt[1] != 0) {
            a += events[j].y - events[j - 1].y;
        }
        int bit = std::abs(events[j].mask_delta);
        int poly = (bit == 1 ? 0 : 1);
        cnt[poly] += events[j].mask_delta / bit;
    }
    res += a * (xsa[k + 1] - xsa[k]);
}
return res;
}

template<typename It>
double polygon_area(It lo, It hi) {
    if (lo == hi) {
        return 0;
    }
    double area = 0;
    for (It i = lo, j = hi - 1; i != hi; j = i++) {
        area += (static_cast<double>(j->x) - i->x) * (static_cast<double>(j->y) + i->y);
    }
    return fabs(area / 2.0);
}

template<typename It>
double union_area(It lo1, It hi1, It lo2, It hi2) {
    return polygon_area(lo1, hi1) + polygon_area(lo2, hi2) - intersection_area(lo1, hi1, lo2, hi2);
}

```

Example Usage

```

#include <cassert>
using namespace std;

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    bool operator!=(const Point &p) const { return !(*this == p); }
    bool operator<(const Point &p) const { return x != p.x ? x < p.x : y < p.y; }
};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
    bool operator==(const PointI &o) const { return x == o.x && y == o.y; }
    bool operator!=(const PointI &p) const { return !(*this == p); }
    bool operator<(const PointI &p) const { return x != p.x ? x < p.x : y < p.y; }
};

int main() {
    vector<Point> p, s;
    // Irregular pentagon overlapping the square below in a triangle of area 1.5.
    p.emplace_back(1, 3);
    p.emplace_back(1, 2);
    p.emplace_back(2, 1);
    p.emplace_back(0, 0);
    p.emplace_back(-1, 3);
    // Square of area 12.5 in quadrant 2.
    s.emplace_back(0, 0);
    s.emplace_back(0, 3);
    s.emplace_back(-3, 3);
    s.emplace_back(-3, 0);
    assert(EQ(1.5, intersection_area(p.begin(), p.end(), s.begin(), s.end())));
}

```

```

assert(EQ(12.5, union_area(p.begin(), p.end(), s.begin(), s.end())));

// Integer-coordinate polygons are accepted (computation proceeds in double).
vector<PointI> ip{{1, 3}, {1, 2}, {2, 1}, {0, 0}, {-1, 3}};
vector<PointI> is{{0, 0}, {0, 3}, {-3, 3}, {-3, 0}};
assert(EQ(1.5, intersection_area(ip.begin(), ip.end(), is.begin(), is.end())));
assert(EQ(12.5, union_area(ip.begin(), ip.end(), is.begin(), is.end())));
return 0;
}

```

7.4.5 Delaunay Triangulation (Simple)

Given a set P of distinct two-dimensional points, a Delaunay triangulation of P is a triangulation of the convex hull of P such that no point of P lies strictly inside the circumcircle of any triangle in the triangulation. The Delaunay triangulation is not necessarily unique when four or more points are cocircular.

This implementation produces one valid triangulation using a simple brute-force algorithm. Each candidate triangle is tested against the empty-circumcircle condition, and candidates whose edges would properly cross already accepted triangles are rejected.

- `delaunay_triangulation(p)` returns a Delaunay triangulation of the distinct points in `p` as a vector of clockwise-oriented vertex triples, or an empty vector if no triangulation exists (e.g. fewer than three points, or all points collinear).

All arithmetic uses the point's own coordinate type, so integer inputs yield an exact triangulation.

Overflow warning: the lifted-paraboloid test grows like the fourth power of the coordinate magnitude. With 64-bit integer coordinates, it overflows once coordinates exceed roughly 17600. For floating-point coordinates the test is subject to the usual rounding error.

Time Complexity: $O(n^6)$ per call, where n is the number of points. There are $O(n^3)$ candidate triangles; each scans all points and, conservatively, up to $O(n^3)$ previously accepted candidates.

Space Complexity: $O(n)$ auxiliary, excluding the returned triangulation.

```

#include <cmath>
#include <tuple>
#include <type_traits>
#include <vector>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T, typename U>
bool LE(T a, U b) {
    return !LT(b, a);
}

// Specialized version of seg_intersection() from 7.2.3, simplified for touch_is_intersect = false,
// i.e. we're detecting proper crossings of two non-parallel segments whose interiors intersect.
template<typename Pt>
bool proper_seg_intersection(const Pt &a, const Pt &b, const Pt &c, const Pt &d) {

```

```

auto cross = [](const Pt &o, const Pt &p, const Pt &q) {
    return (p.x - o.x) * (q.y - o.y) - (p.y - o.y) * (q.x - o.x); // Overflow warning!
};
auto c1 = cross(a, b, c), c2 = cross(a, b, d), c3 = cross(c, d, a), c4 = cross(c, d, b);
return ((LT(c1, 0) && LT(0, c2)) || (LT(c2, 0) && LT(0, c1))) &&
        ((LT(c3, 0) && LT(0, c4)) || (LT(c4, 0) && LT(0, c3)));
}

template<typename Pt>
std::vector<std::tuple<Pt, Pt, Pt>> delaunay_triangulation(const std::vector<Pt> &p) {
    int n = static_cast<int>(p.size());
    std::vector<decltype(Pt::x)> z;
    for (const auto &[px, py] : p) {
        z.emplace_back(px * px + py * py); // Overflow warning!
    }
    std::vector<std::tuple<Pt, Pt, Pt>> res;
    std::vector<std::tuple<int, int, int>> res_idx;
    for (int i = 0; i < n - 2; i++) {
        for (int j = i + 1; j < n; j++) {
            for (int k = i + 1; k < n; k++) {
                if (j == k) {
                    continue;
                }
                // Overflow warning!
                auto nx = (p[j].y - p[i].y) * (z[k] - z[i]) - (p[k].y - p[i].y) * (z[j] - z[i]);
                auto ny = (p[k].x - p[i].x) * (z[j] - z[i]) - (p[j].x - p[i].x) * (z[k] - z[i]);
                auto nz = (p[j].x - p[i].x) * (p[k].y - p[i].y) - (p[k].x - p[i].x) * (p[j].y - p[i].y);
                if (LE(0, nz)) {
                    continue;
                }
                std::vector<Pt> s1{p[i], p[j], p[k], p[i]};
                for (int m = 0; m < n; m++) {
                    if (LT(0, nx * (p[m].x - p[i].x) + ny * (p[m].y - p[i].y) + nz * (z[m] - z[i]))) {
                        goto skip;
                    }
                }
                // Reject candidates whose edges properly cross already accepted triangles.
                for (auto [t0, t1, t2] : res_idx) {
                    std::vector<Pt> s2{p[t0], p[t1], p[t2], p[t0]};
                    for (int u = 0; u < 3; u++) {
                        for (int v = 0; v < 3; v++) {
                            if (proper_seg_intersection(s1[u], s1[u + 1], s2[v], s2[v + 1])) {
                                goto skip;
                            }
                        }
                    }
                }
                res.emplace_back(p[i], p[j], p[k]);
                res_idx.emplace_back(i, j, k);
                skip:;
            }
        }
    }
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

struct Point {
    double x, y;
    Point(double x = 0, double y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    bool operator<(const Point &p) const { return x != p.x ? x < p.x : y < p.y; }
}

```



```

};

struct PointI {
    int x, y;
    PointI(int x = 0, int y = 0) : x(x), y(y) {}
    bool operator==(const PointI &p) const { return x == p.x && y == p.y; }
    bool operator<(const PointI &p) const { return x != p.x ? x < p.x : y < p.y; }
};

int main() {
    vector<Point> v{{1, 3}, {1, 2}, {2, 1}, {0, 0}, {-1, 3}};
    vector<tuple<Point, Point, Point>> t{
        {Point{1, 3}, Point{1, 2}, Point{-1, 3}},
        {Point{1, 3}, Point{2, 1}, Point{1, 2}},
        {Point{1, 2}, Point{2, 1}, Point{0, 0}},
        {Point{1, 2}, Point{0, 0}, Point{-1, 3}}
    };
    assert(delaunay_triangulation(v) == t);

    // Integer-coordinate inputs are triangulated using exact integer arithmetic.
    vector<PointI> iv{{1, 3}, {1, 2}, {2, 1}, {0, 0}, {-1, 3}};
    vector<tuple<PointI, PointI, PointI>> ti{
        {PointI{1, 3}, PointI{1, 2}, PointI{-1, 3}},
        {PointI{1, 3}, PointI{2, 1}, PointI{1, 2}},
        {PointI{1, 2}, PointI{2, 1}, PointI{0, 0}},
        {PointI{1, 2}, PointI{0, 0}, PointI{-1, 3}}
    };
    assert(delaunay_triangulation(iv) == ti);
    return 0;
}

```

7.4.6 Delaunay Triangulation (Guibas-Stolfi)

Given a set P of distinct two-dimensional points, a Delaunay triangulation of P is a triangulation of the convex hull of P such that no point of P lies strictly inside the circumcircle of any triangle in the triangulation. The Delaunay triangulation is not necessarily unique when four or more points are cocircular.

This implementation produces one valid triangulation using the Guibas-Stolfi divide-and-conquer algorithm with a quad-edge data structure. The point set is split in half, each half is triangulated recursively, and the halves are stitched together from the bottom up by a sequence of cross edges, with circumcircle tests deciding each connecting edge and deleting invalidated ones.

- `delaunay_triangulation(p)` sorts `p` lexicographically and removes duplicate points in place (so the input vector is modified), then returns the triangles of one Delaunay triangulation as a vector of counterclockwise-oriented `Point<T>` triples (`std::tuple`). If fewer than three non-collinear unique points remain, the result is empty.

The point type must provide exact lexicographic `operator<`; sorting and duplicate removal requires this strict ordering rather than epsilon equality. All arithmetic uses the point's own coordinate type, so integer inputs yield an exact triangulation.

Overflow warning: the in-circle test grows like the fourth power of the coordinate magnitude and dominates the overflow budget. With 64-bit integer coordinates, it overflows once coordinates exceed roughly 14000. For floating-point coordinates, the predicates are subject to the usual rounding error.

Time Complexity: $O(n \log n)$ per call, where n is the number of input points.

Space Complexity:

- $O(n \log n)$ for allocated quad-edges in the worst case because deleted edges remain in the pool.
- $O(n)$ for the returned triangulation.

```

#include <algorithm>
#include <cmath>
#include <cstdint>
#include <tuple>
#include <type_traits>
#include <utility>
#include <vector>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool LT(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) < C(b) - EPS;
    return C(a) < C(b);
}

template<typename T>
struct Point {
    T x, y;
    Point(T x = 0, T y = 0) : x(x), y(y) {}
    bool operator==(const Point &p) const { return x == p.x && y == p.y; }
    friend bool EQ(const Point &a, const Point &b) { return EQ(a.x, b.x) && EQ(a.y, b.y); }
    bool operator<(const Point &p) const { return std::tie(x, y) < std::tie(p.x, p.y); }
};

template<typename T>
T cross(const Point<T> &a, const Point<T> &b, const Point<T> &o) {
    return (a.x - o.x) * (b.y - o.y) - (a.y - o.y) * (b.x - o.x); // Overflow warning!
}

template<typename T>
T incircle(const Point<T> &a, const Point<T> &b, const Point<T> &c, const Point<T> &d) {
    T adx = a.x - d.x, ady = a.y - d.y;
    T bdx = b.x - d.x, bdy = b.y - d.y;
    T cdx = c.x - d.x, cdy = c.y - d.y;
    T abdet = adx * bdy - bdx * ady; // Overflow warning!
    T bcdet = bdx * cdy - cdx * bdy;
    T cadet = cdx * ady - adx * cdy;
    T alift = adx * adx + ady * ady;
    T blift = bdx * bdx + bdy * bdy;
    T clift = cdx * cdx + cdy * cdy;
    return alift * bcdet + blift * cadet + clift * abdet;
}

template<typename T>
struct Edge {
    Point<T> origin;
    Edge *rot, *onext;
    bool primal, deleted, used;

    Edge *rev() const { return rot->rot; }
    Edge *lnext() const { return rot->rev()->onext->rot; }
    Edge *oprev() const { return rot->onext->rot; }
    Point<T> dest() const { return rev()->origin; }
};

template<typename T>
struct QuadEdgePool {
    std::vector<Edge<T>*> edges;

    QuadEdgePool() = default;
    QuadEdgePool(const QuadEdgePool &) = delete;

```

```

QuadEdgePool &operator=(const QuadEdgePool &) = delete;

~QuadEdgePool() {
    for (Edge<T> *e : edges) {
        delete e;
    }
}

Edge<T> *make_edge(const Point<T> &u, const Point<T> &v) {
    Edge<T> *e1 = new Edge<T>, *e2 = new Edge<T>, *e3 = new Edge<T>, *e4 = new Edge<T>;
    edges.push_back(e1);
    edges.push_back(e2);
    edges.push_back(e3);
    edges.push_back(e4);
    e1->origin = u;
    e2->origin = v;
    e1->rot = e3;
    e2->rot = e4;
    e3->rot = e2;
    e4->rot = e1;
    e1->onext = e1;
    e2->onext = e2;
    e3->onext = e4;
    e4->onext = e3;
    e1->primal = e2->primal = true;
    e3->primal = e4->primal = false;
    e1->deleted = e2->deleted = e3->deleted = e4->deleted = false;
    e1->used = e2->used = e3->used = e4->used = false;
    return e1;
}

void delete_edge(Edge<T> *e) {
    splice(e, e->oprev());
    splice(e->rev(), e->rev()->oprev());
    e->deleted = e->rot->deleted = e->rev()->deleted = e->rev()->rot->deleted = true;
}

Edge<T> *connect(Edge<T> *a, Edge<T> *b) {
    Edge<T> *e = make_edge(a->dest(), b->origin);
    splice(e, a->lnext());
    splice(e->rev(), b);
    return e;
}

static void splice(Edge<T> *a, Edge<T> *b) {
    std::swap(a->onext->rot->onext, b->onext->rot->onext);
    std::swap(a->onext, b->onext);
}

};

template<typename T>
bool left_of(const Point<T> &p, Edge<T> *e) {
    return LT(0, cross(e->origin, e->dest(), p));
}

template<typename T>
bool right_of(const Point<T> &p, Edge<T> *e) {
    return LT(cross(e->origin, e->dest(), p), 0);
}

template<typename T>
bool in_circle(const Point<T> &a, const Point<T> &b, const Point<T> &c, const Point<T> &d) {
    return LT(0, incircle(a, b, c, d));
}

template<typename T>
std::pair<Edge<T> *, Edge<T> *> build_triangulation(
    std::vector<Point<T>> &p, int lo, int hi, QuadEdgePool<T> &pool

```

```

) {
    if (hi - lo == 1) {
        Edge<T> *a = pool.make_edge(p[lo], p[hi]);
        return {a, a->rev()};
    }
    if (hi - lo == 2) {
        Edge<T> *a = pool.make_edge(p[lo], p[lo + 1]);
        Edge<T> *b = pool.make_edge(p[lo + 1], p[hi]);
        pool.splice(a->rev(), b);
        int side = LT(0, cross(p[lo], p[lo + 1], p[hi]))
            ? 1
            : (LT(cross(p[lo], p[lo + 1], p[hi]), 0) ? -1 : 0);
        if (side == 0) {
            return {a, b->rev()};
        }
        Edge<T> *c = pool.connect(b, a);
        if (side == 1) {
            return {a, b->rev()};
        }
        return {c->rev(), c};
    }
    int mid = lo + (hi - lo) / 2;
    auto [ldo, ldi] = build_triangulation(p, lo, mid, pool);
    auto [rdi, rdo] = build_triangulation(p, mid + 1, hi, pool);
    while (true) {
        if (left_of(rdi->origin, ldi)) {
            ldi = ldi->lnext();
        } else if (right_of(ldi->origin, rdi)) {
            rdi = rdi->rev()->onext;
        } else {
            break;
        }
    }
    Edge<T> *base = pool.connect(rdi->rev(), ldi);
    if (ldi->origin == ldo->origin) {
        ldo = base->rev();
    }
    if (rdi->origin == rdo->origin) {
        rdo = base;
    }
    while (true) {
        Edge<T> *lcand = base->rev()->onext;
        if (right_of(lcand->dest(), base)) {
            while (in_circle(base->dest(), base->origin, lcand->dest(), lcand->onext->dest())) {
                Edge<T> *next = lcand->onext;
                pool.delete_edge(lcand);
                lcand = next;
            }
        }
        Edge<T> *rcand = base->oprev();
        if (right_of(rcand->dest(), base)) {
            while (in_circle(base->dest(), base->origin, rcand->dest(), rcand->oprev()->dest())) {
                Edge<T> *prev = rcand->oprev();
                pool.delete_edge(rcand);
                rcand = prev;
            }
        }
        bool lvalid = right_of(lcand->dest(), base), rvalid = right_of(rcand->dest(), base);
        if (!lvalid && !rvalid) {
            break;
        }
        if (!lvalid ||
            (rvalid && in_circle(lcand->dest(), lcand->origin, rcand->origin, rcand->dest()))) {
            base = pool.connect(rcand, base->rev());
        } else {
            base = pool.connect(base->rev(), lcand->rev());
        }
    }
}

```

```

    return {ldo, rdo};
}

template<typename T>
auto delaunay_triangulation(std::vector<Point<T>> &p) {
    using Pt = Point<T>;
    using Triangle = std::tuple<Pt, Pt, Pt>;
    // Rotates triangle vertices so that the lexicographically smallest comes first, preserving the
    // cyclic (counterclockwise) order, to give each triangle a canonical, comparable representation.
    auto canonical = [](const Pt &a, const Pt &b, const Pt &c) -> Triangle {
        if (b < a && b < c) {
            return {b, c, a};
        }
        if (c < a && c < b) {
            return {c, a, b};
        }
        return {a, b, c};
    };
    std::sort(p.begin(), p.end());
    auto same_point = [](const Pt &a, const Pt &b) { return !(a < b) && !(b < a); };
    p.erase(std::unique(p.begin(), p.end(), same_point), p.end());
    std::vector<Triangle> res;
    if (p.size() < 3) {
        return res;
    }
    QuadEdgePool<T> pool;
    build_triangulation(p, 0, static_cast<int>(p.size()) - 1, pool);
    for (Edge<T> *start : pool.edges) {
        if (!start->primal || start->deleted || start->used) {
            continue;
        }
        std::vector<Pt> face;
        Edge<T> *curr = start;
        do {
            curr->used = true;
            face.push_back(curr->origin);
            curr = curr->lnext();
        } while (curr != start);
        if (face.size() == 3 && LT(0, cross(face[0], face[1], face[2]))) {
            res.push_back(canonical(face[0], face[1], face[2]));
        }
    }
    std::sort(res.begin(), res.end());
    return res;
}

```

Example Usage

```

#include <cassert>
using namespace std;

using PointL = Point<int64_t>;
using PointD = Point<double>;

int main() {
    vector<PointD> v{{1, 3}, {1, 2}, {2, 1}, {0, 0}, {-1, 3}};
    vector<tuple<PointD, PointD, PointD>> t{
        {PointD{-1, 3}, PointD{0, 0}, PointD{1, 2}},
        {PointD{-1, 3}, PointD{1, 2}, PointD{1, 3}},
        {PointD{0, 0}, PointD{2, 1}, PointD{1, 2}},
        {PointD{1, 2}, PointD{2, 1}, PointD{1, 3}}
    };
    assert(delaunay_triangulation(v) == t);

    // Integer-coordinate inputs are triangulated using exact integer arithmetic.
    vector<PointL> iv{{1, 3}, {1, 2}, {2, 1}, {0, 0}, {-1, 3}};
}

```

```

vector<tuple<PointL, PointL, PointL>> ti{
    {PointL{-1, 3}, PointL{0, 0}, PointL{1, 2}},
    {PointL{-1, 3}, PointL{1, 2}, PointL{1, 3}},
    {PointL{0, 0}, PointL{2, 1}, PointL{1, 2}},
    {PointL{1, 2}, PointL{2, 1}, PointL{1, 3}}
};
assert(delaunay_triangulation(iv) == ti);
return 0;
}

```

7.5 3D Geometry

7.5.1 3D Point

A 3D point class template supporting vector arithmetic, dot and cross products, distances, angles, unit vectors, normals, and rotation about an axis. Point equality and ordering are exact, while `EQ(p, q)` provides epsilon equality for floating-point coordinates. Integer coordinates are suitable for the exact algebraic operations `dot()`, `cross()`, and `sqnorm()` as long as intermediate products do not overflow; metric operations preserve floating-point coordinates and otherwise convert them to `double`.

`unit()` requires a nonzero point, `normal(p)` requires nonparallel vectors, and `rotate(angle, axis)` requires a nonzero axis.

Overflow warning: the exact products `dot()`, `cross()`, and `sqnorm()` grow like the squared coordinate magnitude. With `TPoint3<int>` these overflow a 32-bit `int` once coordinates exceed a few tens of thousands, so use `Point3L` (`TPoint3<int64_t>`) for larger integer coordinates.

Type aliases:

- `Point3I = TPoint3<int>`
- `Point3L = TPoint3<int64_t>`
- `Point3D = TPoint3<double>`
- `Point3LD = TPoint3<long double>`

Time Complexity: $O(1)$ per operation.

Space Complexity: $O(1)$ for storage and auxiliary.

```

#include <cassert>
#include <cmath>
#include <cstdint>
#include <tuple>
#include <type_traits>

const double EPS = 1e-9;

template<typename T, typename U, typename C = std::common_type_t<T, U>>
bool EQ(T a, U b) {
    if constexpr (std::is_floating_point_v<C>) return C(a) == C(b) || fabs(C(a) - C(b)) <= EPS;
    return C(a) == C(b);
}

template<typename T>
struct TPoint3 {
    using fp_t = std::conditional_t<std::is_floating_point<T>::value, T, double>;

    T x, y, z;

    TPoint3(T x = 0, T y = 0, T z = 0) : x(x), y(y), z(z) {}
    bool operator==(const TPoint3 &p) const { return std::tie(x, y, z) == std::tie(p.x, p.y, p.z); }
}

```

```

friend bool EQ(const TPoint3 &a, const TPoint3 &b) {
    return EQ(a.x, b.x) && EQ(a.y, b.y) && EQ(a.z, b.z);
}

bool operator!=(const TPoint3 &p) const { return !(*this == p); }
bool operator<(const TPoint3 &p) const { return std::tie(x, y, z) < std::tie(p.x, p.y, p.z); }
TPoint3 operator+(const TPoint3 &p) const { return {x + p.x, y + p.y, z + p.z}; }
TPoint3 operator-(const TPoint3 &p) const { return {x - p.x, y - p.y, z - p.z}; }
TPoint3 operator*(T k) const { return {x * k, y * k, z * k}; }
TPoint3<fp_t> operator/(fp_t k) const { return {(fp_t)x / k, (fp_t)y / k, (fp_t)z / k}; }
T dot(const TPoint3 &p) const { return x * p.x + y * p.y + z * p.z; } // Overflow warning!
T sqnorm() const { return x * x + y * y + z * z; } // Overflow warning!

fp_t norm() const {
    return hypot(hypot(static_cast<fp_t>(x), static_cast<fp_t>(y)), static_cast<fp_t>(z));
}

fp_t phi() const { return atan2(static_cast<fp_t>(y), static_cast<fp_t>(x)); }
TPoint3<fp_t> unit() const { return *this / norm(); }

TPoint3 cross(const TPoint3 &p) const {
    return {y * p.z - z * p.y, z * p.x - x * p.z, x * p.y - y * p.x}; // Overflow warning!
}

TPoint3<fp_t> normal(const TPoint3 &p) const { return cross(p).unit(); }

fp_t theta() const {
    return atan2(hypot(static_cast<fp_t>(x), static_cast<fp_t>(y)), static_cast<fp_t>(z));
}

TPoint3<fp_t> rotate(fp_t angle, const TPoint3 &axis) const {
    fp_t s = sin(angle), c = cos(angle);
    TPoint3<fp_t> u = axis.unit(), v((fp_t)x, (fp_t)y, (fp_t)z);
    return u * v.dot(u) * (1 - c) + v * c - v.cross(u) * s;
}
};

using Point3I = TPoint3<int>;
using Point3L = TPoint3<int64_t>;
using Point3D = TPoint3<double>;
using Point3LD = TPoint3<long double>;

```

Example Usage

```

int main() {
    Point3I a(1, 0, 0), b(0, 1, 0);
    assert(a.dot(b) == 0);
    assert(a.cross(b) == Point3I(0, 0, 1));
    assert(a.sqnorm() == 1);

    Point3D p(1, 0, 0);
    auto r = p.rotate(acos(-1.0) / 2, Point3D(0, 0, 1));
    assert(EQ(r, Point3D(0, 1, 0)));
    return 0;
}

```

7.5.2 Polyhedron Volume

Given an oriented triangular mesh of a closed polyhedron, compute its signed volume. Each triangle contributes the signed volume of the tetrahedron formed by the triangle abc and the origin $((a \times b) \cdot c)/6$. If faces are oriented outward, the result is positive for the orientation used by `convex_hull_3d()` in 7.5.3; take `abs()` if only the magnitude is needed.

- `signed_polyhedron_volume(p, faces)` returns the signed volume. Each face triplet (`a`, `b`, `c`) represents the triangle with vertices `p[a]`, `p[b]`, and `p[c]`.

Overflow warning: for integral point types, each scalar triple product must fit in the point's coordinate arithmetic type before it is converted to `double`.

Time Complexity: $O(f)$ per call, where f is the number of triangular faces.

Space Complexity: $O(1)$ auxiliary.

```
#include <cmath>
#include <tuple>
#include <vector>

struct Point3D {
    double x, y, z;
    Point3D(double x = 0, double y = 0, double z = 0) : x(x), y(y), z(z) {}
    double dot(const Point3D &p) const { return x * p.x + y * p.y + z * p.z; }

    Point3D cross(const Point3D &p) const {
        return {y * p.z - z * p.y, z * p.x - x * p.z, x * p.y - y * p.x};
    }
};

template<typename Pt, typename F>
double signed_polyhedron_volume(const std::vector<Pt> &p, const std::vector<F> &faces) {
    double volume = 0;
    for (const auto &[a, b, c] : faces) {
        volume += p[a].cross(p[b]).dot(p[c]); // Overflow warning!
    }
    return volume / 6.0;
}
```

Example Usage

```
#include <cassert>
using namespace std;

bool EQ(double a, double b) {
    return fabs(a - b) < 1e-9;
}

int main() {
    using Face = std::tuple<int, int, int>; // (a, b, c)
    vector<Point3D> p{{0, 0, 0}, {1, 0, 0}, {0, 1, 0}, {0, 0, 1}};
    vector<Face> faces{
        Face(0, 2, 1),
        Face(0, 1, 3),
        Face(0, 3, 2),
        Face(1, 2, 3),
    };
    assert(EQ(signed_polyhedron_volume(p, faces), 1.0 / 6.0));
    return 0;
}
```

7.5.3 3D Convex Hull

Given points in 3D, compute the triangular faces of their convex hull with an incremental algorithm. Start with a tetrahedron, delete every visible face for each new point, and stitch the horizon edges to that point. Returned faces are oriented outward.

- `convex_hull_3d(p)` returns triangular faces as `Face{a, b, c}` triples, where each face has vertices `p[a]`, `p[b]`, and `p[c]`.

Preconditions: at least four points, and no four points coplanar. For coplanar or highly degenerate inputs, perturbing points slightly or using a more robust library is recommended.

Time Complexity: $O(n^2)$ per call, where n is the number of points.

Space Complexity: $O(n^2)$ auxiliary.

```
#include <array>
#include <cassert>
#include <utility>
#include <vector>

const double EPS = 1e-9;

struct Point3D {
    double x, y, z;
    Point3D(double x = 0, double y = 0, double z = 0) : x(x), y(y), z(z) {}
    Point3D operator-(const Point3D &p) const { return {x - p.x, y - p.y, z - p.z}; }
    Point3D operator*(double k) const { return {x * k, y * k, z * k}; }
    double dot(const Point3D &p) const { return x * p.x + y * p.y + z * p.z; }

    Point3D cross(const Point3D &p) const {
        return {y * p.z - z * p.y, z * p.x - x * p.z, x * p.y - y * p.x};
    }
};

struct Face {
    int a, b, c;
    Face(int a = 0, int b = 0, int c = 0) : a(a), b(b), c(c) {}
};

std::vector<Face> convex_hull_3d(const std::vector<Point3D> &p) {
    int n = static_cast<int>(p.size());
    assert(n >= 4);
    std::vector<std::vector<std::array<int, 2>>> edge(
        n, std::vector<std::array<int, 2>>(n, {-1, -1})
    );
    std::vector<Face> faces;
    auto normal = [&](const Face &f) { return (p[f.b] - p[f.a]).cross(p[f.c] - p[f.a]); };
    auto visible = [&](const Face &f, int i) { return normal(f).dot(p[i] - p[f.a]) > EPS; };
    auto add_edge = [&](int a, int b, int c) {
        if (a > b) {
            std::swap(a, b);
        }
        auto &e = edge[a][b];
        (e[0] == -1 ? e[0] : e[1]) = c;
    };
    auto remove_edge = [&](int a, int b, int c) {
        if (a > b) {
            std::swap(a, b);
        }
        auto &e = edge[a][b];
        (e[0] == c ? e[0] : e[1]) = -1;
    };
    auto edge_count = [&](int a, int b) {
        if (a > b) {
            std::swap(a, b);
        }
        return (edge[a][b][0] != -1) + (edge[a][b][1] != -1);
    };
    auto make_face = [&](int a, int b, int c, int inside) {
        Face f(a, b, c);
        if (normal(f).dot(p[inside] - p[a]) > 0) {
            std::swap(f.b, f.c);
        }
    }
}
```

```

    add_edge(f.a, f.b, f.c);
    add_edge(f.b, f.c, f.a);
    add_edge(f.c, f.a, f.b);
    faces.push_back(f);
};
for (int i = 0; i < 4; i++) {
    for (int j = i + 1; j < 4; j++) {
        for (int k = j + 1; k < 4; k++) {
            make_face(i, j, k, 6 - i - j - k);
        }
    }
}
for (int i = 4; i < n; i++) {
    for (int j = 0; j < static_cast<int>(faces.size()); j++) {
        Face f = faces[j];
        if (visible(f, i)) {
            remove_edge(f.a, f.b, f.c);
            remove_edge(f.b, f.c, f.a);
            remove_edge(f.c, f.a, f.b);
            std::swap(faces[j--], faces.back());
            faces.pop_back();
        }
    }
    int old_faces = static_cast<int>(faces.size());
    for (int j = 0; j < old_faces; j++) {
        Face f = faces[j];
        if (edge_count(f.a, f.b) != 2) {
            make_face(f.a, f.b, i, f.c);
        }
        if (edge_count(f.b, f.c) != 2) {
            make_face(f.b, f.c, i, f.a);
        }
        if (edge_count(f.c, f.a) != 2) {
            make_face(f.c, f.a, i, f.b);
        }
    }
}
return faces;
}

```

Example Usage

```

using namespace std;

int main() {
    vector<Point3D> tetra{{0, 0, 0}, {1, 0, 0}, {0, 1, 0}, {0, 0, 1}};
    auto faces = convex_hull_3d(tetra);
    // A tetrahedron has four triangular faces.
    assert(faces.size() == 4);
    for (const Face &f : faces) {
        assert(f.a != f.b && f.b != f.c && f.c != f.a);
    }

    tetra.push_back(Point3D(1, 1, 1));
    faces = convex_hull_3d(tetra);
    // Adding the opposite point makes a triangular bipyramid with six triangular faces.
    assert(faces.size() == 6);
    for (const Face &f : faces) {
        assert(0 <= f.a && f.a < static_cast<int>(tetra.size()));
        assert(0 <= f.b && f.b < static_cast<int>(tetra.size()));
        assert(0 <= f.c && f.c < static_cast<int>(tetra.size()));
    }
    return 0;
}

```

Chapter 8

Miscellany

8.1 Contest Utilities

Small contest-template helpers that are useful across many algorithms. These snippets avoid algorithm-specific policy and are meant to be pasted near the top of a solution file.

- `Rep(i,N)`, `For(i,L,H)`, `Rev(i,N)`, `Dwn(i,H,L)`, and `Each(x,C)` are compact loop macros.
- `sz(c)` returns the signed size of a container.
- `ckmin(a, b)` assigns `a = b` and returns true if `b < a` before the assignment.
- `ckmax(a, b)` assigns `a = b` and returns true if `a < b` before the assignment.
- `between(x, a, b)` returns whether $a \leq x \leq b$.
- `clmp(x, a, b)` returns `x` clamped into the interval $[a, b]$. This is not needed on C++17 and later, where `std::clamp()` is available.
- `floor_div(a, b)` and `ceil_div(a, b)` divide signed integers with mathematical rounding toward negative or positive infinity. Requires `b != 0`.
- `sort_unique(v)` sorts a vector and removes duplicates.
- `erase_one(c, x)` erases one existing value from an associative container, asserting that it is present.
- `min_priority_queue<T>` is a min-heap alias.
- `RNG()` constructs a 64-bit Mersenne Twister seeded from `std::random_device`.
- `RNG(seed)` constructs a reproducible 64-bit Mersenne Twister.
- `rng.uniform_int(lo, hi)` returns a random integer in range $[lo, hi]$.
- `rng.uniform_real(lo = 0.0, hi = 1.0)` returns a random real number in the half-open range $[lo, hi)$.
- `rng.shuffle(lo, hi)` randomly shuffles the range $[lo, hi)$.
- `y_combinator(f)` wraps a recursive lambda so that the lambda can call itself as its first argument.

```
#include <algorithm>
#include <cassert>
#include <cstdint>
#include <functional>
#include <queue>
#include <random>
#include <type_traits>
#include <utility>
#include <vector>

#define Rep(i, N) for (int i = 0, _##i = (N); i < _##i; ++i)           // 0 to N-1
#define For(i, L, H) for (int i = (L), _##i = (H); i <= _##i; ++i)    // L to H
#define Rev(i, N) for (int i = (N); --i >= 0;)                        // N-1 to 0
#define Dwn(i, H, L) for (int i = (H), _##i = (L); i >= _##i; --i)    // H to L
#define Each(x, C) for (auto &x : (C))

template<typename C> int sz(const C &c) { return static_cast<int>(c.size()); }
```

```

template<typename T, typename U> bool ckmin(T &a, const U &b) { return b < a ? a = b, true : false; }
template<typename T, typename U> bool ckmax(T &a, const U &b) { return a < b ? a = b, true : false; }
template<typename T> bool between(const T &x, const T &a, const T &b) { return !(x < a) && !(b < x); }
template<typename T> T clmp(const T &x, const T &a, const T &b) { return std::min(std::max(x, a), b); }

template<typename T>
T floor_div(T a, T b) {
    assert(b != 0);
    T q = a / b, r = a % b;
    return q - (r != 0 && ((r < 0) != (b < 0)));
}

template<typename T>
T ceil_div(T a, T b) {
    assert(b != 0);
    T q = a / b, r = a % b;
    return q + (r != 0 && ((r < 0) == (b < 0)));
}

template<typename T>
void sort_unique(std::vector<T> &v) {
    std::sort(v.begin(), v.end());
    v.erase(std::unique(v.begin(), v.end()), v.end());
}

template<typename C, typename T>
void erase_one(C &c, const T &x) {
    auto it = c.find(x);
    assert(it != c.end());
    c.erase(it);
}

template<typename T>
using min_priority_queue = std::priority_queue<T, std::vector<T>, std::greater<T>>;

struct RNG {
    std::mt19937_64 gen;

    RNG() : RNG(std::random_device{}()) {}
    explicit RNG(uint64_t seed) : gen(seed) {}

    template<typename T>
    T uniform_int(T lo, T hi) {
        static_assert(std::is_integral<T>::value, "uniform_int() requires an integral type");
        return std::uniform_int_distribution<T>(lo, hi)(gen);
    }

    double uniform_real(double lo = 0.0, double hi = 1.0) {
        return std::uniform_real_distribution<double>(lo, hi)(gen);
    }

    template<typename It>
    void shuffle(It lo, It hi) {
        std::shuffle(lo, hi, gen);
    }
};

template<typename Fun>
class y_combinator_result {
    Fun fun;

public:
    template<typename T>
    explicit y_combinator_result(T &&fun_) : fun(std::forward<T>(fun_)) {}

    template<typename... Args>
    decltype(auto) operator()(Args &&...args) {
        return fun(std::ref(*this), std::forward<Args>(args)...);
    }
};

```

```

    }
};

template<typename Fun>
decltype(auto) y_combinator(Fun &&fun) {
    return y_combinator_result<std::decay_t<Fun>>>(std::forward<Fun>(fun));
}

```

Example Usage

```

#include <set>
using namespace std;

int main() {
    int tot = 0;
    Rep(i, 4) tot += i;      // 0 + 1 + 2 + 3 = 6
    For(i, 2, 5) tot += i;  // 6 + (2 + 3 + 4 + 5) = 20
    Rev(i, 3) tot += i;     // 20 + (2 + 1 + 0) = 23
    Dwn(i, 8, 6) tot += i;  // 23 + (8 + 7 + 6) = 44
    assert(tot == 44);

    int x = 10;
    assert(ckmin(x, 7) && x == 7);
    assert(!ckmin(x, 9) && x == 7);
    assert(ckmax(x, 11) && x == 11);

    vector<int> v = {1, 2, 3};
    assert(sz(v) == 3);
    assert(between(2, 1, 3));
    assert(between(1, 1, 3) && between(3, 1, 3));
    assert(!between(0, 1, 3));
    assert(clmp(-5, 1, 3) == 1);
    assert(clmp(10, 1, 3) == 3);
    assert(floor_div(7, -3) == -3);
    assert(floor_div(-7, 3) == -3);
    assert(ceil_div(7, -3) == -2);
    assert(ceil_div(-7, 3) == -2);
    v = {3, 1, 3, 2, 1};
    sort_unique(v);
    assert((v == vector<int>{1, 2, 3}));
    multiset<int> ms{1, 2, 2, 3};
    erase_one(ms, 2);
    assert(ms.count(2) == 1);
    min_priority_queue<int> pq;
    pq.push(3);
    pq.push(1);
    assert(pq.top() == 1);

    RNG rng(1234567); // Fixed seed for reproducibility.
    int r = rng.uniform_int(1, 6);
    assert(1 <= r && r <= 6);
    double d = rng.uniform_real();
    assert(0.0 <= d && d < 1.0);
    rng.shuffle(v.begin(), v.end());
    assert(sz(v) == 3);

    // Recursive lambda with y_combinator.
    auto fib =
        y_combinator([](auto fib, int n) -> int { return n < 2 ? n : fib(n - 1) + fib(n - 2); });
    assert(fib(10) == 55);

    // Recursive lambda without y_combinator: pass the lambda to itself as the first argument.
    vector<vector<int>> adj{{1, 2}, {0}, {0}};
    int seen = 0;
    auto dfs = [&](auto &&dfs, int u, int p) -> void {
        seen++;
    }
}

```

```

    Each(v, adj[u]) if (v != p) dfs(dfs, v, u);
};
dfs(dfs, 0, -1);
assert(seen == 3);

// Recursive lambdas are automatically supported in C++23.
#if __cplusplus >= 202302L
    auto fact = [&](this auto fact, int n) -> int { return n ? n * fact(n - 1) : 1; };
    assert(fact(5) == 120);
#endif
    return 0;
}

```

8.2 Fast IO

Fast input and output wrappers based on `fread()` and `fwrite()`. For most problems, `ios::sync_with_stdio(false); cin.tie(nullptr);` is simpler and fast enough. Use these classes when input is huge or when `iostream` overhead is measurable.

- `FastInput in(file = stdin)` reads tokens from a `FILE*`.
- `in >> x` reads a non-whitespace token into `char`, `std::string`, integral types, or floating point types.
- `FastOutput out(file = stdout)` writes to a `FILE*`.
- `out << x` writes `char`, C strings, `std::string`, integral types, or floating point types.
- `out.flush()` writes any buffered output immediately.
- `set_in(name)`, `set_out(name)`, and `set_io(iname, oname)` redirect standard input/output to files. For example, `set_io("task.in", "task.out")` is convenient for USACO-style problems.

The parser assumes valid input. Floating point input/output is provided for convenience, not as the main performance path.

```

#include <cassert>
#include <cctype>
#include <cstdio>
#include <cstdlib>
#include <cstring>
#include <limits>
#include <string>
#include <type_traits>

void set_in(const std::string &name) {
    FILE *result = freopen(name.c_str(), "r", stdin);
    assert(result != nullptr);
    (void)result;
}

void set_out(const std::string &name) {
    FILE *result = freopen(name.c_str(), "w", stdout);
    assert(result != nullptr);
    (void)result;
}

void set_io(const std::string &iname, const std::string &oname) {
    set_in(iname);
    set_out(oname);
}

struct FastInput {
    static constexpr int BUF_SIZE = 1 << 20;
    FILE *file;
    char buf[BUF_SIZE];
    int pos = 0, len = 0;
}

```

```

explicit FastInput(FILE *file_ = stdin) : file(file_) {}

char get_char() {
    if (pos == len) {
        len = static_cast<int>(fread(buf, 1, BUF_SIZE, file));
        pos = 0;
        if (len == 0) {
            return 0;
        }
    }
    return buf[pos++];
}

char peek_char() {
    char c = get_char();
    if (c) {
        --pos;
    }
    return c;
}

void skip_blanks() {
    while (true) {
        char c = get_char();
        if (!c) {
            return;
        }
        if (!std::isspace(static_cast<unsigned char>(c))) {
            --pos;
            return;
        }
    }
}

FastInput &operator>>(char &c) {
    skip_blanks();
    c = get_char();
    return *this;
}

FastInput &operator>>(std::string &s) {
    skip_blanks();
    s.clear();
    while (true) {
        char c = get_char();
        if (!c || std::isspace(static_cast<unsigned char>(c))) {
            break;
        }
        s += c;
    }
    return *this;
}

FastInput &operator>>(bool &x) {
    int val;
    *this >> val;
    x = val != 0;
    return *this;
}

template<typename T>
typename std::enable_if<
    std::is_integral<T>::value && !std::is_same<T, bool>::value, FastInput &>::type
operator>>(T &x) {
    skip_blanks();
    char c = get_char();
    bool neg = false;

```

```

    if (c == '-') {
        neg = true;
        c = get_char();
    }
    using U = typename std::make_unsigned<T>::type;
    U val = 0;
    while (c && !std::isspace(static_cast<unsigned char>(c))) {
        val = U(10) * val + U(c - '0');
        c = get_char();
    }
    if constexpr (std::is_signed<T>::value) {
        x = neg ? T(U(0) - val) : T(val);
    } else {
        x = T(val);
    }
    return *this;
}

template<typename T>
typename std::enable_if<std::is_floating_point<T>::value, FastInput &>::type operator>>(T &x) {
    std::string s;
    *this >> s;
    x = static_cast<T>(std::strtold(s.c_str(), nullptr));
    return *this;
}
};

struct FastOutput {
    static constexpr int BUF_SIZE = 1 << 20;
    FILE *file;
    char buf[BUF_SIZE];
    int pos = 0;

    explicit FastOutput(FILE *file_ = stdout) : file(file_) {}
    ~FastOutput() { flush(); }

    void flush() {
        if (pos) {
            fwrite(buf, 1, pos, file);
            pos = 0;
        }
    }

    void put_char(char c) {
        if (pos == BUF_SIZE) {
            flush();
        }
        buf[pos++] = c;
    }

    FastOutput &operator<<(char c) {
        put_char(c);
        return *this;
    }

    FastOutput &operator<<(const char *s) {
        while (*s) {
            put_char(*s++);
        }
        return *this;
    }

    FastOutput &operator<<(const std::string &s) {
        for (char c : s) {
            put_char(c);
        }
        return *this;
    }
}

```



```

FastOutput &operator<<(bool x) {
    put_char(x ? '1' : '0');
    return *this;
}

template<typename T>
typename std::enable_if<
    std::is_integral<T>::value && !std::is_same<T, bool>::value, FastOutput &>::type
operator<<(T x) {
    if (x == 0) {
        put_char('0');
        return *this;
    }
    using U = typename std::make_unsigned<T>::type;
    U val;
    if constexpr (std::is_signed<T>::value) {
        if (x < 0) {
            put_char('-');
            val = U(0) - U(x);
        } else {
            val = U(x);
        }
    } else {
        val = x;
    }
    char s[32];
    int n = 0;
    while (val > 0) {
        s[n++] = char('0' + val % 10);
        val /= 10;
    }
    while (n-- > 0) {
        put_char(s[n]);
    }
    return *this;
}

template<typename T>
typename std::enable_if<std::is_floating_point<T>::value, FastOutput &>::type operator<<(T x) {
    char s[64];
    std::snprintf(
        s, sizeof(s), "%.Lg", std::numeric_limits<T>::max_digits10, static_cast<long double>(x)
    );
    return *this << s;
}
};

```

Example Usage

```

int main() {
    // set_io("file.in", "file.out");

    FILE *input = tmpfile();
    std::fputs("42 hello 3.5 -2147483648 1 Z 1234567890123\n", input);
    std::rewind(input);

    FastInput in(input);
    int x;
    std::string s;
    double y;
    int z;
    bool flag;
    char ch;
    long long big;
    in >> x >> s >> y >> z >> flag >> ch >> big;
}

```

```

assert(x == 42 && s == "hello" && y == 3.5);
assert(z == -2147483648);
assert(flag && ch == 'Z' && big == 1234567890123LL);
std::fclose(input);

FILE *output = tmpfile();
{
    FastOutput out(output);
    out << x << ' ' << s << ' ' << y << ' ' << z << ' ' << flag << ' ' << ch << ' ' << big << '\n';
    out.flush();
}
std::rewind(output);
char buf[96] = {};
assert(std::fgets(buf, sizeof(buf), output));
assert(std::strcmp(buf, "42 hello 3.5 -2147483648 1 Z 1234567890123", 42) == 0);
std::fclose(output);
return 0;
}

```

8.3 Debugging

Local-only debug printing for contest code. The macro expands to a no-op unless `LOCAL` is defined (e.g. by passing `-DLOCAL` to the compiler flags), so it can remain in submitted code without producing output.

- `dbg_repr(x)` returns a string representation of arithmetic types, strings, pairs, tuples, iterable containers, and custom types that support insertion into `std::ostream`.
- `dbg(a, b, c)` prints the function name, line number, argument names, and values to `std::cerr` when `LOCAL` is defined.
- `pr(a, b, c)` prints only the values to `std::cerr` when `LOCAL` is defined.
- `dbg(...)` evaluates to no output when `LOCAL` is not defined.
- `pr(...)` evaluates to no output when `LOCAL` is not defined.

Avoid naming these helpers `to_string()` unless you want them mixed into overload resolution with `std::to_string()` and other user-defined `to_string()` functions. A separate name keeps debug formatting local to this snippet.

```

#include <cstdlib>
#include <iomanip>
#include <iostream>
#include <sstream>
#include <string>
#include <tuple>
#include <type_traits>
#include <utility>

template<typename T, typename = void>
struct is_iterable : std::false_type {};

template<typename T>
struct is_iterable<
    T, std::void_t<decltype(std::begin(std::declval<T>())), decltype(std::end(std::declval<T>()))>>
    : std::true_type {};

template<typename T, typename = void>
struct is_streamable : std::false_type {};

template<typename T>
struct is_streamable<
    T, std::void_t<decltype(std::declval<std::ostream &>() << std::declval<const T &>())>>
    : std::true_type {};

```

```

std::string dbg_repr(const std::string &s) { return '"' + s + '"'; }
std::string dbg_repr(const char *s) { return dbg_repr(std::string(s)); }
std::string dbg_repr(char c) { return std::string("\"") + c + "\""; }
std::string dbg_repr(bool b) { return b ? "true" : "false"; }

template<typename T>
typename std::enable_if<
    std::is_arithmetic<T>::value && !std::is_same<T, char>::value &&
    !std::is_same<T, bool>::value && !std::is_floating_point<T>::value,
    std::string>::type
dbg_repr(T x) {
    std::ostringstream out;
    out << x;
    return out.str();
}

template<typename T>
typename std::enable_if<std::is_floating_point<T>::value, std::string>::type dbg_repr(T x) {
    std::ostringstream out;
    out << std::setprecision(10) << x; // Set floating point dbg precision here.
    return out.str();
}

template<typename T>
typename std::enable_if<
    is_streamable<T>::value && !is_iterable<T>::value && !std::is_arithmetic<T>::value &&
    !std::is_convertible<T, std::string>::value,
    std::string>::type
dbg_repr(const T &x) {
    std::ostringstream out;
    out << x;
    return out.str();
}

template<typename A, typename B>
std::string dbg_repr(const std::pair<A, B> &p) {
    return "(" + dbg_repr(p.first) + ", " + dbg_repr(p.second) + ")";
}

template<typename Tuple, std::size_t... Is>
std::string dbg_tuple_repr(const Tuple &t, std::index_sequence<Is...>) {
    std::string res = "(";
    bool first = true;
    ((res += (first ? (first = false, "") : ", ") + dbg_repr(std::get<Is>(t))), ...);
    return res + ")";
}

template<typename... Ts>
std::string dbg_repr(const std::tuple<Ts...> &t) {
    return dbg_tuple_repr(t, std::index_sequence_for<Ts...>{});
}

template<typename T>
typename std::enable_if<
    is_iterable<T>::value && !std::is_convertible<T, std::string>::value, std::string>::type
dbg_repr(const T &v) {
    std::string res = "{";
    bool first = true;
    for (const auto &x : v) {
        if (!first) {
            res += ", ";
        }
        first = false;
        res += dbg_repr(x);
    }
    return res + "}";
}

```

```

void dbg_out(bool) {
    std::cerr << '\n';
}

template<typename Head, typename... Tail>
void dbg_out(bool leading_space, const Head &head, const Tail &...tail) {
    if (leading_space) {
        std::cerr << ' ';
    }
    std::cerr << dbg_repr(head);
    ((std::cerr << ' ' << dbg_repr(tail)), ...);
    std::cerr << '\n';
}

#ifdef LOCAL
#define dbg(...) \
    std::cerr << "[" << __func__ << ":" << __LINE__ << "]" << " " << #__VA_ARGS__ << ":", \
    dbg_out(true, __VA_ARGS__)
#define pr(...) dbg_out(false, __VA_ARGS__)
#else
#define dbg(...) ((void)0)
#define pr(...) ((void)0)
#endif

```

Example Usage

```

#include <cassert>
#include <map>
#include <set>
#include <vector>
using namespace std;

struct Point {
    int x, y;
};

// Can overload dbg_repr to dbg your custom structs.
string dbg_repr(const Point &p) {
    return "Point(" + dbg_repr(p.x) + ", " + dbg_repr(p.y) + ")";
}

struct Edge {
    int u, v, w;
};

string dbg_repr(const Edge &e) {
    return "Edge{" + dbg_repr(e.u) + "->" + dbg_repr(e.v) + ", w=" + dbg_repr(e.w) + "}";
}

struct Streamable {
    int value;

    friend ostream &operator<<(ostream &out, const Streamable &x) {
        return out << "Streamable(" << x.value << ")";
    }
};

int main() {
    vector<int> v = {1, 2, 3};
    auto p = make_pair(4, string("x"));
    auto t = make_tuple(1, true, string("ok"));
    map<string, int> mp = {"a", 1}, {"b", 2};
    Point pt{2, 3};
    Edge e{0, 1, 5};
    assert(dbg_repr(v) == "{1, 2, 3}");
    assert(dbg_repr(p) == "(4, \"x\")");
}

```

```

assert(dbg_repr(t) == "(1, true, \"ok\")");
assert(dbg_repr(mp) == "{(\"a\", 1), (\"b\", 2)}");
assert(dbg_repr(pt) == "Point(2, 3)");
assert(dbg_repr(e) == "Edge{0->1, w=5}");
assert(dbg_repr(Streamable{6}) == "Streamable(6)");
assert(dbg_repr(vector<Streamable>{{7}, {8}}) == "{Streamable(7), Streamable(8)}");
double d = 3.141592653589793;
dbg(v);
dbg(p, t, mp);
dbg(pt, e);
pr("plain output", d);
assert(dbg_repr(1.0 / 3.0) == "0.3333333333");
assert(dbg_repr(d) == "3.141592654");
return 0;
}

```

Example Output

```

[main:166] v: {1, 2, 3}
[main:167] p, t, mp: (4, "x") (1, true, "ok") {("a", 1), ("b", 2)}
[main:168] pt, e: Point(2, 3) Edge{0->1, w=5}
"plain output" 3.141592654

```

8.4 Timing and Benchmarking

Small timing helpers for local profiling and codebook examples. Use `std::chrono::steady_clock` instead of wall-clock time so elapsed durations are monotonic.

- `Timer()` starts a timer immediately.
- `timer.reset()` restarts the timer.
- `timer.elapsed()` returns the elapsed time in milliseconds as a double.
- `benchmark(name, iterations, f)` runs `f()` `iterations` times and returns the average elapsed milliseconds per call with the given name. `iterations` must be positive.

Benchmarks should be treated as local measurements, not proofs of asymptotic performance. Warm-up, CPU frequency scaling, compiler optimization, and input choice can dominate tiny timings.

```

#include <cassert>
#include <chrono>
#include <string>
#include <utility>
#include <vector>

class Timer {
    using Clock = std::chrono::steady_clock;
    Clock::time_point start_time;

public:
    Timer() { reset(); }

    void reset() { start_time = Clock::now(); }

    double elapsed() const {
        std::chrono::duration<double, std::milli> diff = Clock::now() - start_time;
        return diff.count();
    }
};

template<typename Fun>
std::pair<std::string, double> benchmark(const std::string &name, int iterations, Fun f) {
    assert(iterations > 0);
    Timer timer;

```

```

    for (int i = 0; i < iterations; ++i) {
        f();
    }
    return {name, timer.elapsed() / iterations};
}

```

Example Usage

```

#include <cassert>
#include <numeric>
#include <vector>
using namespace std;

int main() {
    Timer timer;
    std::vector<int> v(1000);
    std::iota(v.begin(), v.end(), 0);
    assert(timer.elapsed() >= 0.0);

    volatile long long sink = 0;
    auto [name, avg_time] = benchmark("sum", 5, [&] {
        long long cur = 0;
        for (int x : v) {
            cur += x;
        }
        sink += cur;
    });
    assert(name == "sum" && avg_time >= 0.0);
    assert(sink > 0);
    return 0;
}

```

8.5 Stress Testing

A simple stress-testing harness compares a trusted brute force implementation against an optimized implementation on many generated tests. This is most useful for finding small counterexamples while developing an algorithm.

- `stress(trials, generate, solve, brute)` runs `generate()` for each trial, compares `solve(test)` with `brute(test)`, and returns the first failing test if one is found.
- `stress(...)` returns `std::nullopt` if all generated tests pass.

The generated test type only needs to be copyable enough to store in `std::optional`, and both solutions must return comparable results.

```

#include <optional>

template<typename Generate, typename Solve, typename Brute>
auto stress(int trials, Generate generate, Solve solve, Brute brute)
    -> std::optional<decltype(generate())> {
    for (int tc = 0; tc < trials; ++tc) {
        auto test = generate();
        if (solve(test) != brute(test)) {
            return test;
        }
    }
    return std::nullopt;
}

```

Example Usage

```
#include <algorithm>
#include <cassert>
#include <random>
#include <vector>
using namespace std;

int main() {
    std::mt19937 rng(1234567); // Fixed seed for reproducibility.

    auto generate = [&] {
        int n = std::uniform_int_distribution<int>(1, 8)(rng);
        std::vector<int> a(n);
        for (int &x : a) {
            x = std::uniform_int_distribution<int>(-5, 5)(rng);
        }
        return a;
    };

    auto solve = [](std::vector<int> a) {
        std::sort(a.begin(), a.end());
        return a;
    };

    auto brute = [](std::vector<int> a) {
        do {
            if (std::is_sorted(a.begin(), a.end())) {
                return a;
            }
        } while (std::next_permutation(a.begin(), a.end()));
        std::sort(a.begin(), a.end());
        return a;
    };

    assert(!stress(100, generate, solve, brute));
    return 0;
}
```

8.6 Policy-Based Data Structures

GNU policy-based data structures provide non-standard associative containers and priority queues, including order-statistic trees and fast hash tables. They are not part of the C++ standard library, but they are available on many GNU C++ contest judges.

- `HashMap<K, V>` and `HashSet<K>` behave like faster, non-standard alternatives to `std::unordered_map<K, V>` and `std::unordered_set<K>`, using `gp_hash_table`. The default `SplitMix64Hasher` is for integer-like keys; pass a custom hash for other key types. It randomizes the hash seed to avoid weak adversarial integer hashes.
- `OrderedSet<T>` and `OrderedMap<K, V>` behave like `std::set<T>` and `std::map<K, V>` but also support `find_by_order(k)` and `order_of_key(x)`.
- `OrderedMultiset<T>` behaves like `std::multiset<T>` by storing duplicates as (value, unique ID) pairs.
- `OrderedMultimap<K, V>` behaves like `std::multimap<K, V>` by storing duplicate keys as (key, unique ID) pairs.
- `PairingHeap<T>` is a meldable priority queue; `push()` returns an iterator that can be passed to `modify()` or `erase()`.
- `BinaryHeap<T>`, `BinomialHeap<T>`, and `RcBinomialHeap<T>` are alternative GNU priority queue tags with the same interface.

The order-statistic operations use GNU PBDS conventions: `find_by_order(k)` is 0-based and `order_of_key(x)` returns the number of keys strictly less than `x`, even if `x` is absent. Raw PBDS ordered trees return iterators from `find_by_order(k)`; the multiset and multimap wrappers below unwrap that iterator into a value. The priority

queue `join()` operation melds another heap into the current one. Since this depends on `__gnu_pbds`, keep a standard library fallback in mind when portability matters.

Time Complexity:

- `HashMap`/`HashSet`: same expected bounds as `std::unordered_map`/`std::unordered_set` (expected $O(1)$, worst-case $O(n)$).
- `OrderedSet`/`OrderedMap`: same as `std::set`/`std::map`, with $O(\log n)$ for `find_by_order` and `order_of_key`.
- `OrderedMultiset`/`OrderedMultimap`: same as `std::multiset`/`std::multimap`, with $O(\log n)$ for `find_by_order` and `order_of_key`.
- $O(1)$ amortized per call to `push()` and `join()` for `PairingHeap`.
- $O(\log n)$ amortized per call to `pop()`, `modify()`, and `erase()` for `PairingHeap`.

Space Complexity: $O(n)$ for all containers.

```
#include <cassert>
#include <chrono>
#include <cstdint>
#include <cstdint>
#include <ext/pb_ds/assoc_container.hpp>
#include <ext/pb_ds/priority_queue.hpp>
#include <ext/pb_ds/tree_policy.hpp>
#include <functional>
#include <utility>

struct SplitMix64Hasher {
    // SplitMix64 mixer.
    static uint64_t mix64(uint64_t x) {
        x += 0x9e3779b97f4a7c15ULL;
        x = (x ^ (x >> 30)) * 0xbf58476d1ce4e5b9ULL;
        x = (x ^ (x >> 27)) * 0x94d049bb133111ebULL;
        return x ^ (x >> 31);
    }

    std::size_t operator()(uint64_t x) const {
        static const uint64_t FIXED_RANDOM =
            std::chrono::steady_clock::now().time_since_epoch().count();
        return mix64(x + FIXED_RANDOM);
    }
};

template<typename K, typename Hash = SplitMix64Hasher>
using HashSet = __gnu_pbds::gp_hash_table<K, __gnu_pbds::null_type, Hash>;

template<typename K, typename V, typename Hash = SplitMix64Hasher>
using HashMap = __gnu_pbds::gp_hash_table<K, V, Hash>;

template<typename K, typename V, typename Compare = std::less<K>>
using OrderedMap = __gnu_pbds::tree<
    K, V, Compare, __gnu_pbds::rb_tree_tag, __gnu_pbds::tree_order_statistics_node_update>;

template<typename T, typename Compare = std::less<T>>
using OrderedSet = OrderedMap<T, __gnu_pbds::null_type, Compare>;

template<typename T>
class OrderedMultiset {
    OrderedSet<std::pair<T, int>> tree;
    int timer = 0;

public:
    void insert(const T &x) { tree.insert({x, timer++}); }

    bool erase_one(const T &x) {
        auto it = tree.lower_bound({x, -1});
        if (it == tree.end() || it->first != x) {
            return false;
        }
    }
};
```



```

    }
    tree.erase(it);
    return true;
}

T find_by_order(int k) const {
    assert(0 <= k && k < size());
    return tree.find_by_order(k)->first;
}

int order_of_key(const T &x) const { return tree.order_of_key({x, -1}); }
int size() const { return static_cast<int>(tree.size()); }
};

template<typename K, typename V>
class OrderedMultimap {
    OrderedMap<std::pair<K, int>, V> tree;
    int timer = 0;

public:
    void insert(const K &key, const V &value) { tree.insert({key, timer++}, value); }

    bool erase_one(const K &key) {
        auto it = tree.lower_bound({key, -1});
        if (it == tree.end() || it->first.first != key) {
            return false;
        }
        tree.erase(it);
        return true;
    }

    std::pair<K, V> find_by_order(int k) const {
        assert(0 <= k && k < size());
        auto it = tree.find_by_order(k);
        return {it->first.first, it->second};
    }

    int order_of_key(const K &x) const { return tree.order_of_key({x, -1}); }
    int size() const { return static_cast<int>(tree.size()); }
};

template<typename T, typename Compare = std::less<T>>
using PairingHeap = __gnu_pbds::priority_queue<T, Compare, __gnu_pbds::pairing_heap_tag>;

template<typename T, typename Compare = std::less<T>>
using BinaryHeap = __gnu_pbds::priority_queue<T, Compare, __gnu_pbds::binary_heap_tag>;

template<typename T, typename Compare = std::less<T>>
using BinomialHeap = __gnu_pbds::priority_queue<T, Compare, __gnu_pbds::binomial_heap_tag>;

template<typename T, typename Compare = std::less<T>>
using RcBinomialHeap = __gnu_pbds::priority_queue<T, Compare, __gnu_pbds::rc_binomial_heap_tag>;

```

Example Usage

```

int main() {
    OrderedSet<int> s;
    s.insert(10);
    s.insert(20);
    s.insert(30);
    s.insert(20); // Duplicate ignored, like std::set.
    // Net new over std::set: order-statistic queries.
    assert(s.size() == 3);
    assert(*s.find_by_order(0) == 10);
    assert(*s.find_by_order(1) == 20);
    assert(s.order_of_key(25) == 2);
}

```

```

assert(s.order_of_key(100) == 3);

HashMap<int, int> hm;
hm[10] = 1;
hm[20] = 2;
assert(hm[10] == 1);
HashSet<int> hs;
hs.insert(7);
assert(hs.find(7) != hs.end());

OrderedMultiset<int> ms;
ms.insert(5);
ms.insert(5);
ms.insert(7);
assert(ms.size() == 3);
// Net new over std::multiset: order statistics with duplicates.
assert(ms.order_of_key(7) == 2);
assert(ms.find_by_order(1) == 5);
assert(ms.erase_one(5));
assert(ms.order_of_key(7) == 1);
assert(ms.erase_one(5));
assert(!ms.erase_one(5));
assert(ms.find_by_order(0) == 7);

OrderedMultimap<int, char> mp;
mp.insert(5, 'a');
mp.insert(5, 'b');
mp.insert(7, 'c');
// Net new over std::multimap: order statistics with duplicate keys.
assert(mp.order_of_key(7) == 2);
assert(mp.find_by_order(1).first == 5);
assert(mp.erase_one(5));
assert(mp.order_of_key(7) == 1);
assert(mp.erase_one(5));
assert(!mp.erase_one(5));
assert(mp.find_by_order(0) == std::make_pair(7, 'c'));

PairingHeap<int> h1, h2;
auto it = h1.push(10);
auto erased = h1.push(5);
h1.modify(it, 20); // Net new over std::priority_queue: update through a handle.
h1.erase(erased); // Also net new: remove an arbitrary heap item through a handle.
h2.push(7);
h1.join(h2); // Also net new: meld another heap into this one.
assert(h2.empty());
assert(h1.top() == 20);
h1.pop();
assert(h1.top() == 7);
return 0;
}

```

8.7 Multithreading

Basic multithreading can speed up independent local work such as stress tests, batch experiments, or offline preprocessing. It is rarely useful in standard contest submissions, and many judges restrict threads, so use it only when the environment allows it.

- `parallel_for(n, threads, f)` runs `f(i)` once for each $0 \leq i < n$, splitting the indices among at most `threads` worker threads.
- `parallel_cases(cases, threads, solve)` solves independent test cases in parallel and returns the outputs in the original case order.

Each call to `f(i)` must be safe to run concurrently with the others. If it writes shared state, use separate output slots or synchronization.

```
#include <algorithm>
#include <cassert>
#include <numeric>
#include <sstream>
#include <string>
#include <thread>
#include <vector>

template<typename Fun>
void parallel_for(int n, int threads, Fun f) {
    threads = std::max(1, std::min(threads, n));
    std::vector<std::thread> workers;
    workers.reserve(threads);
    for (int id = 0; id < threads; ++id) {
        workers.emplace_back([=, &f] {
            for (int i = id; i < n; i += threads) {
                f(i);
            }
        });
    }
    for (auto &worker : workers) {
        worker.join();
    }
}

template<typename Case, typename Solve>
std::vector<std::string> parallel_cases(const std::vector<Case> &cases, int threads, Solve solve) {
    std::vector<std::string> out(cases.size());
    parallel_for(static_cast<int>(cases.size()), threads, [&](int i) {
        std::ostringstream os;
        solve(cases[i], os);
        out[i] = os.str();
    });
    return out;
}
```

Example Usage

```
int main() {
    std::vector<long long> squares(1000);
    parallel_for(static_cast<int>(squares.size()), 4, [&](int i) { squares[i] = 1LL * i * i; });
    assert(squares[17] == 289);
    assert(std::accumulate(squares.begin(), squares.end(), 0LL) == 332833500);

    std::vector<int> cases = {3, 1, 4, 2};
    auto out =
        parallel_cases(cases, 2, [](int x, std::ostringstream &os) { os << x << "^2 = " << x * x; });
    assert((out == std::vector<std::string>{"3^2 = 9", "1^2 = 1", "4^2 = 16", "2^2 = 4"}));
    return 0;
}
```

8.8 Memory Allocation

Simple bump allocators for local contest code. These are useful when creating many small objects and individual `new/delete` overhead is measurable. They allocate from one fixed buffer and do not reclaim individual objects. The global `operator new` override sends all dynamic allocations through a static buffer. This is the shortest version, but it affects the whole program and should be used only when you are sure the fixed buffer is large enough.

- `BumpAllocator<T>` is an STL allocator for containers such as `std::vector<T, BumpAllocator<T>>`. It is less invasive than overriding global `new`, but still has no real deallocation.

These snippets are not portable production allocators: they are fixed-size, monotonic, and intended for local performance-sensitive contest solutions. The global override below covers ordinary single-object `new`; it does not attempt to be a complete replacement for all allocation forms such as `new[]` or over-aligned allocation.

Time Complexity: $O(1)$ per allocation.

Space Complexity: $O(\text{BUFFER_SIZE})$ static memory.

```
#include <cassert>
#include <cstdint>
#include <vector>

static constexpr std::size_t BUFFER_SIZE = 64 << 20;
alignas(std::max_align_t) static unsigned char bump_buffer[BUFFER_SIZE];
static std::size_t bump_pos = BUFFER_SIZE;

void *operator new(std::size_t size) {
    if (size == 0) {
        size = 1;
    }
    size = (size + alignof(std::max_align_t) - 1) & ~(alignof(std::max_align_t) - 1);
    assert(size <= bump_pos);
    bump_pos -= size;
    return bump_buffer + bump_pos;
}

void operator delete(void *) noexcept {}

void operator delete(void *, std::size_t) noexcept {}

template<typename T>
struct BumpAllocator {
    using value_type = T;

    BumpAllocator() = default;

    template<typename U>
    BumpAllocator(const BumpAllocator<U> &) {}

    T *allocate(std::size_t n) {
        static_assert(alignof(T) <= alignof(std::max_align_t), "over-aligned types are not supported");
        std::size_t size = n * sizeof(T);
        std::size_t align = alignof(T);
        assert(size <= bump_pos);
        bump_pos = (bump_pos - size) & ~(align - 1);
        return reinterpret_cast<T *>(bump_buffer + bump_pos);
    }

    void deallocate(T *, std::size_t) {}

    template<typename U>
    bool operator==(const BumpAllocator<U> &) const {
        return true;
    }

    template<typename U>
    bool operator!=(const BumpAllocator<U> &) const {
        return false;
    }
};
```

Example Usage

```
int main() {
    int *x = new int(42);
    assert(*x == 42);
    delete x;

    std::size_t before_vector = bump_pos;
    std::vector<int, BumpAllocator<int>> v;
    for (int i = 0; i < 1000; i++) {
        v.push_back(i);
    }
    assert(bump_pos < before_vector);
    assert(v.size() == 1000);
    assert(v[123] == 123);
    assert(v.back() == 999);
    return 0;
}
```